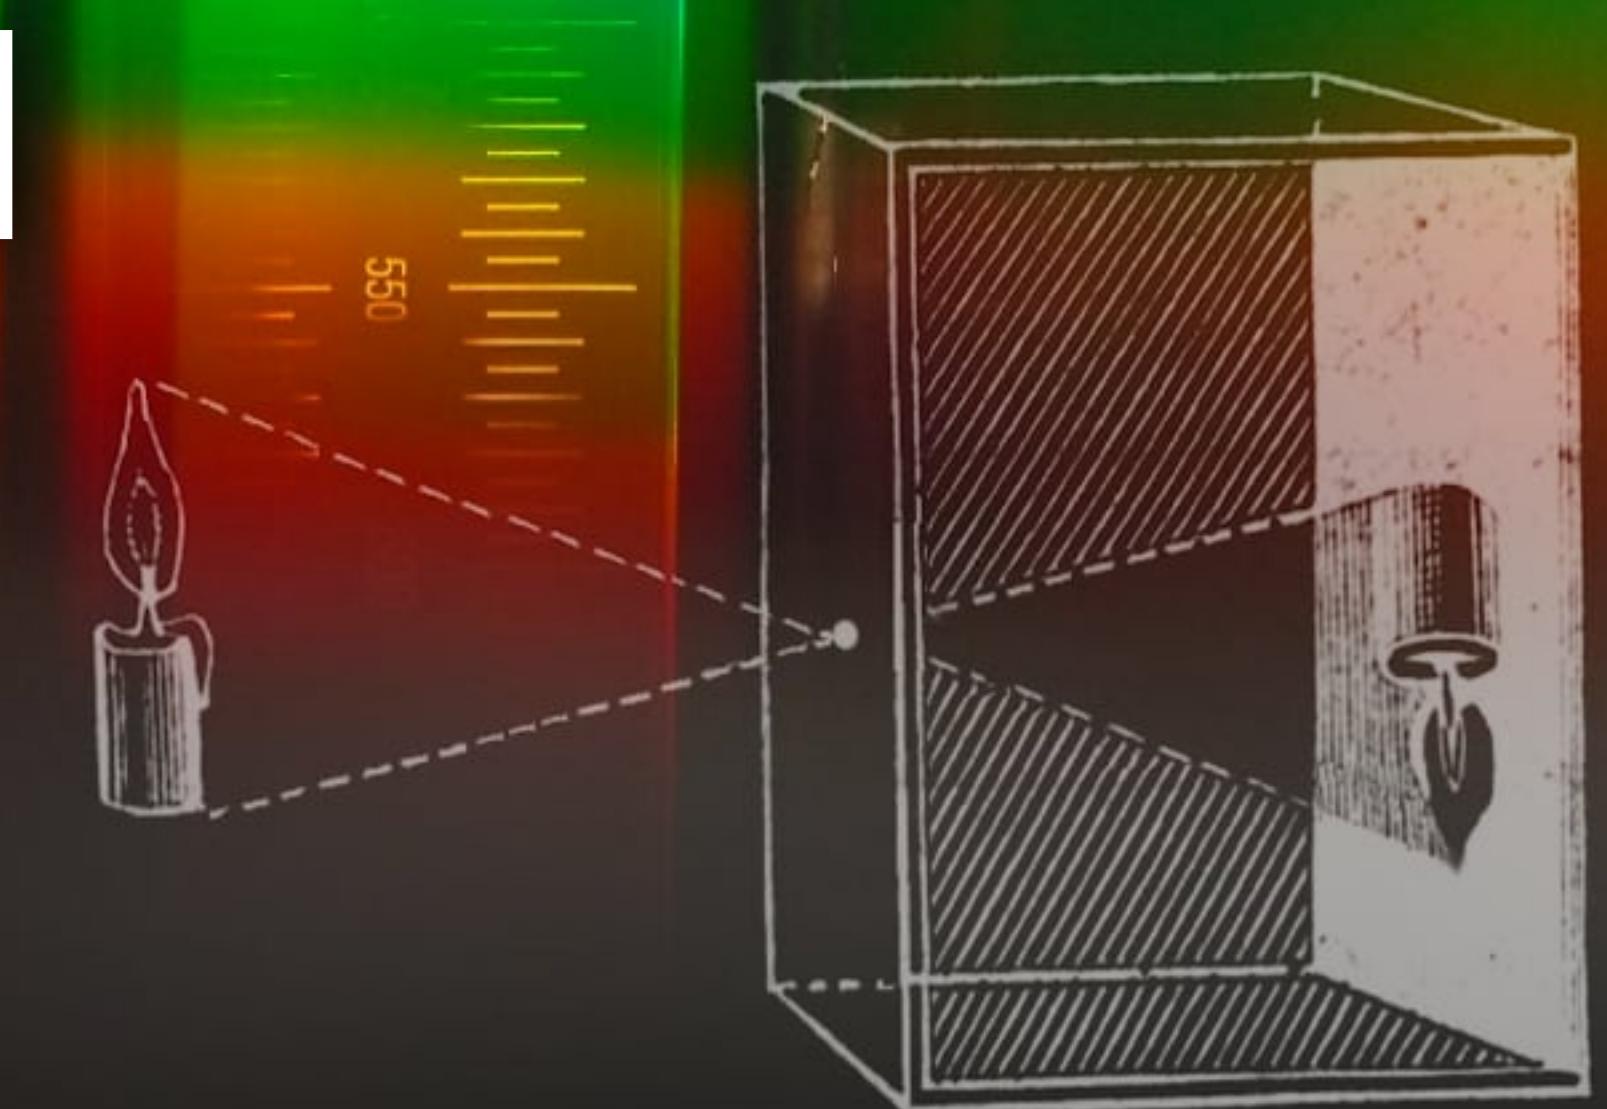


IMAGE FORMATION



Update on Enrollment

- **For-credit (40):** Will randomly sample from waitlist tonight to fill free spots. Notifications will be sent out tomorrow morning.
- **Listeners (20):** Will randomly sample from listener sign-ups. Notifications will be sent out tomorrow morning.
- Asked for a bigger room to accommodate more listeners, but none available :/
- Slides, recordings, assignments, and everything else will be made available publicly - so everyone has the same access as the “official” listeners!
- My goal for next year: Bigger room, course instead of seminar -> better accreditation.

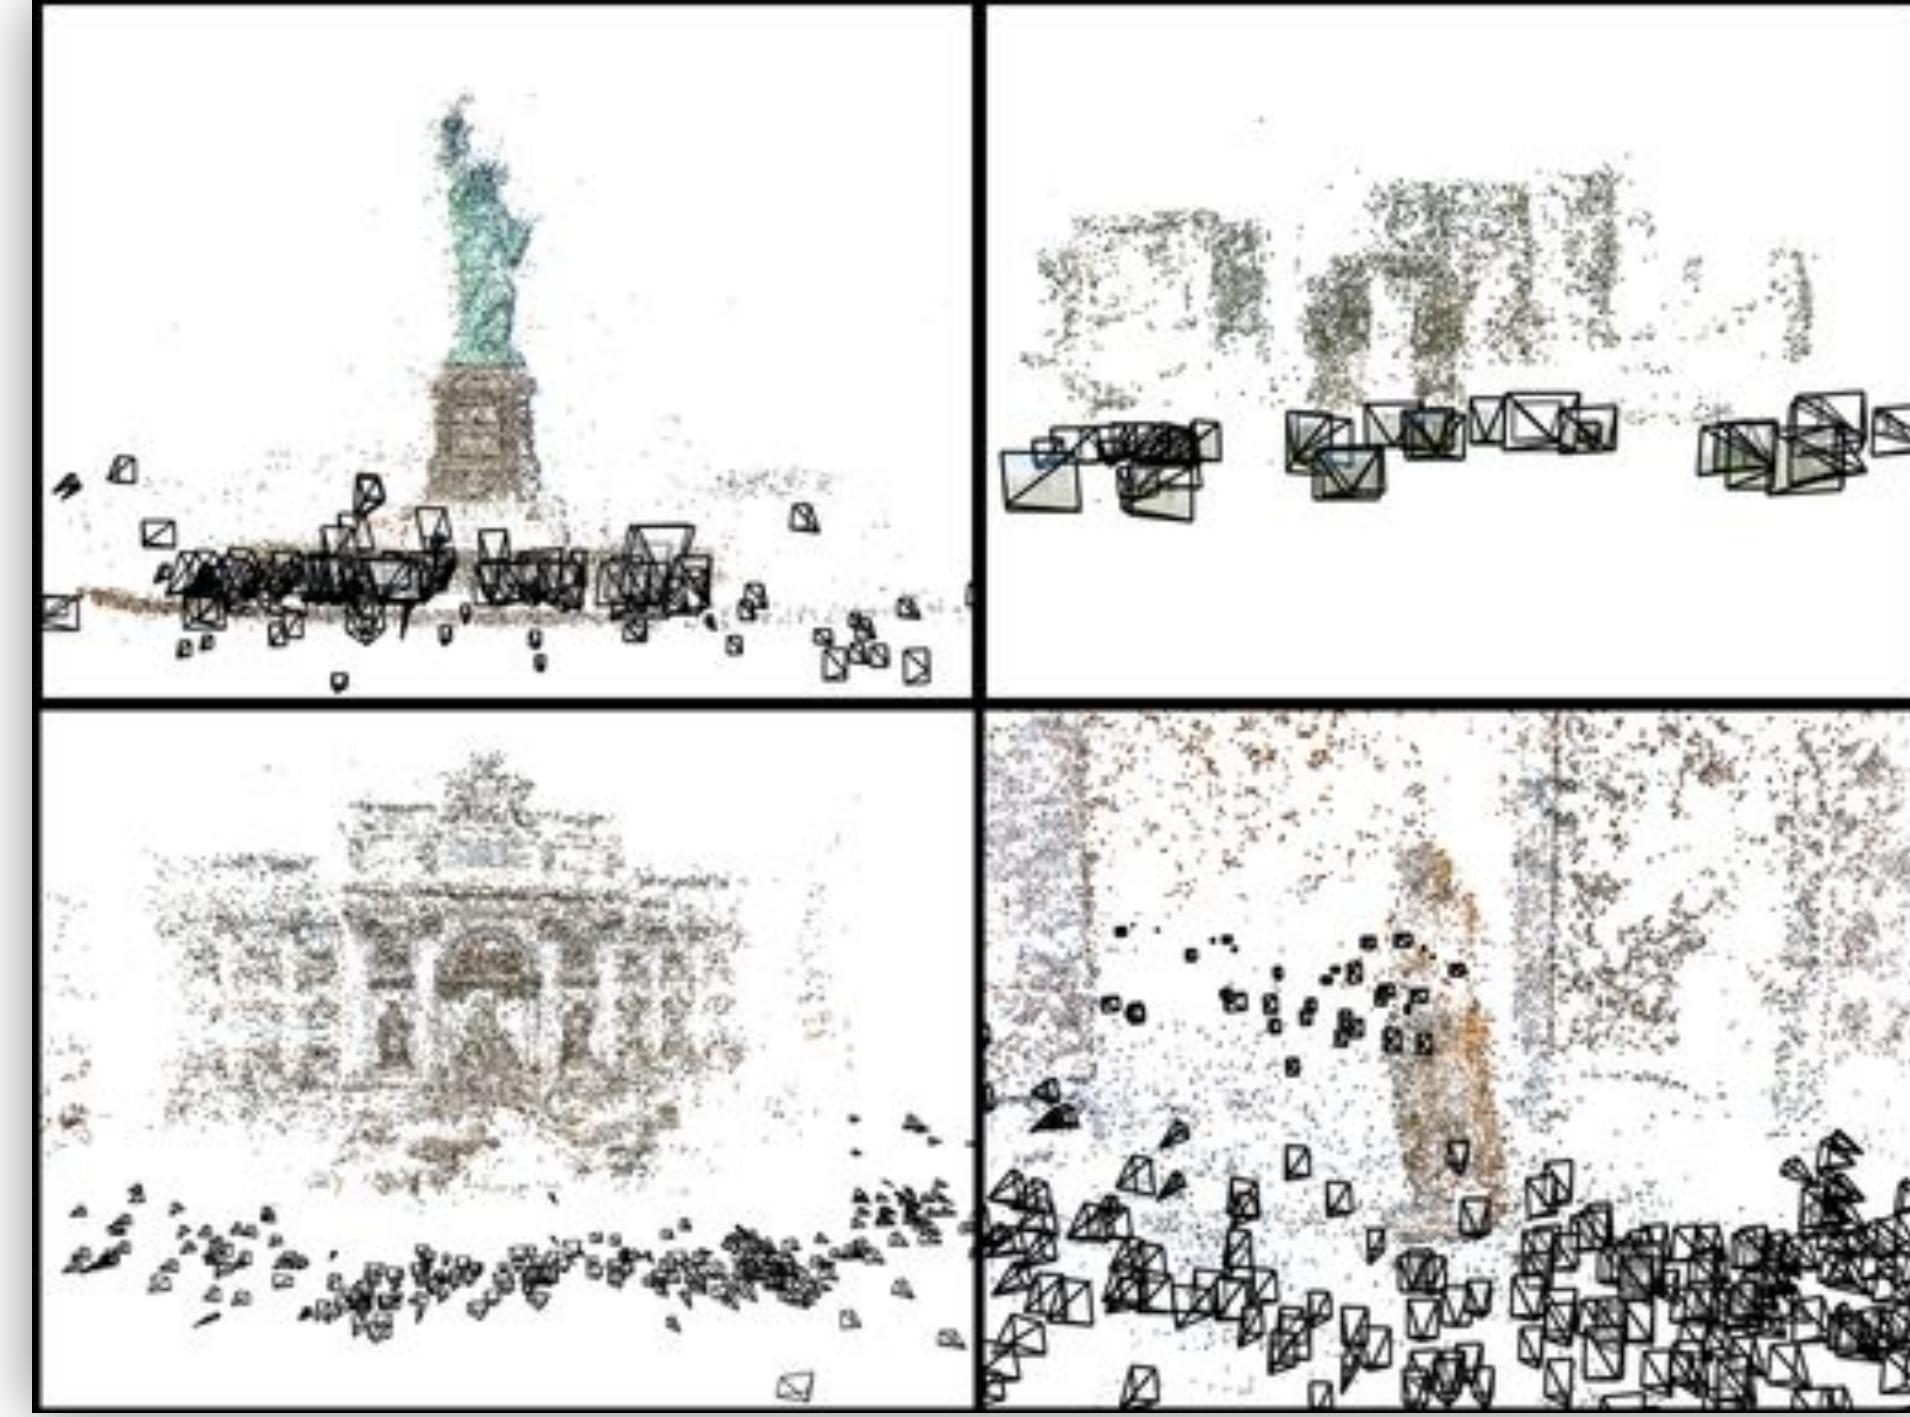
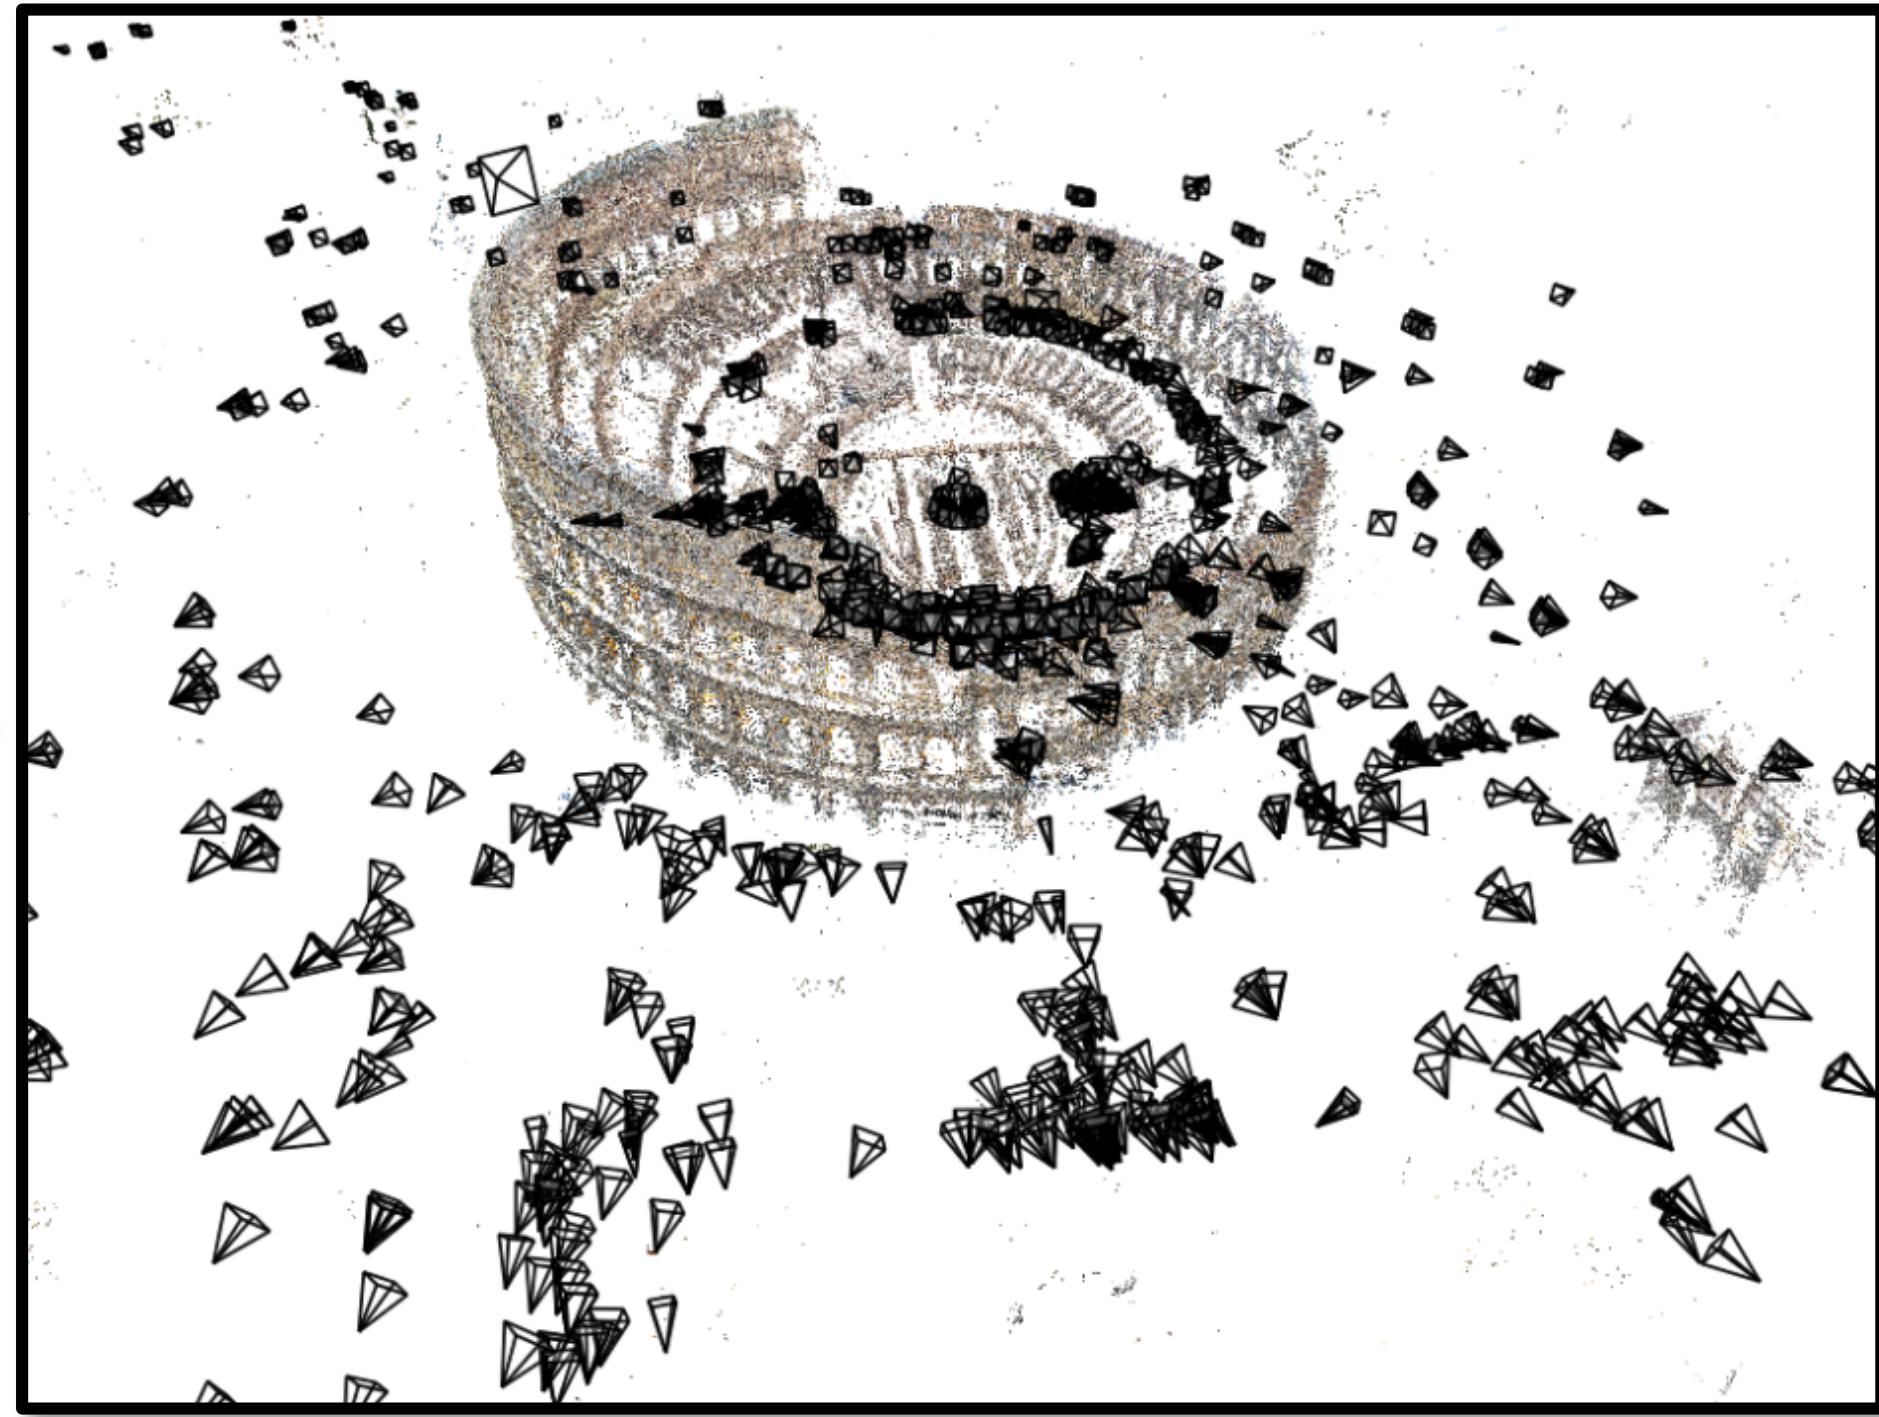
Other Admin Things

- Find calendar to subscribe to on course website.
- First lecture recording posted.
- Slides posted.
- First assignment posted later today!
Note: Very long, but not much code to write at all!

Course Project Clarifications

- **Teams of 1-3.**
- **Counts for 30% of final grade.**
 - Proposal (5%)
 - 3-minute Mid-term Presentation (5%)
 - Final Report + Presentation (20%)
 - Late days for proposal, but not for final report.
- **See course website for up-to-date deadlines. Summary:**
 - Proposal due Oct 18.
 - Project Mid-term Update Presentation (3 min per team) in Nov. 3rd lecture.
 - Final project presentations (~5-10 minutes per team) on Dec 8th and Dec 13th.

Module 1: Image Formation and Multi-View Geometry



Why?

We want to understand 3D world only from 2D observations (images). For that, we need to have a mathematical understanding of how they are connected.

What you'll learn.

Mathematical model of cameras. Reconstruct camera poses, approximate geometry, and camera parameters from 2D images of a scene.

Some Slides adapted from...

- CMU 16-889: Learning for 3D Vision

Prof. Shubham Tulsiani

- CMU 16-385: Computer Vision

Prof. Kris Kitani

- MIT 6.819/6.869: Advances in Computer Vision,

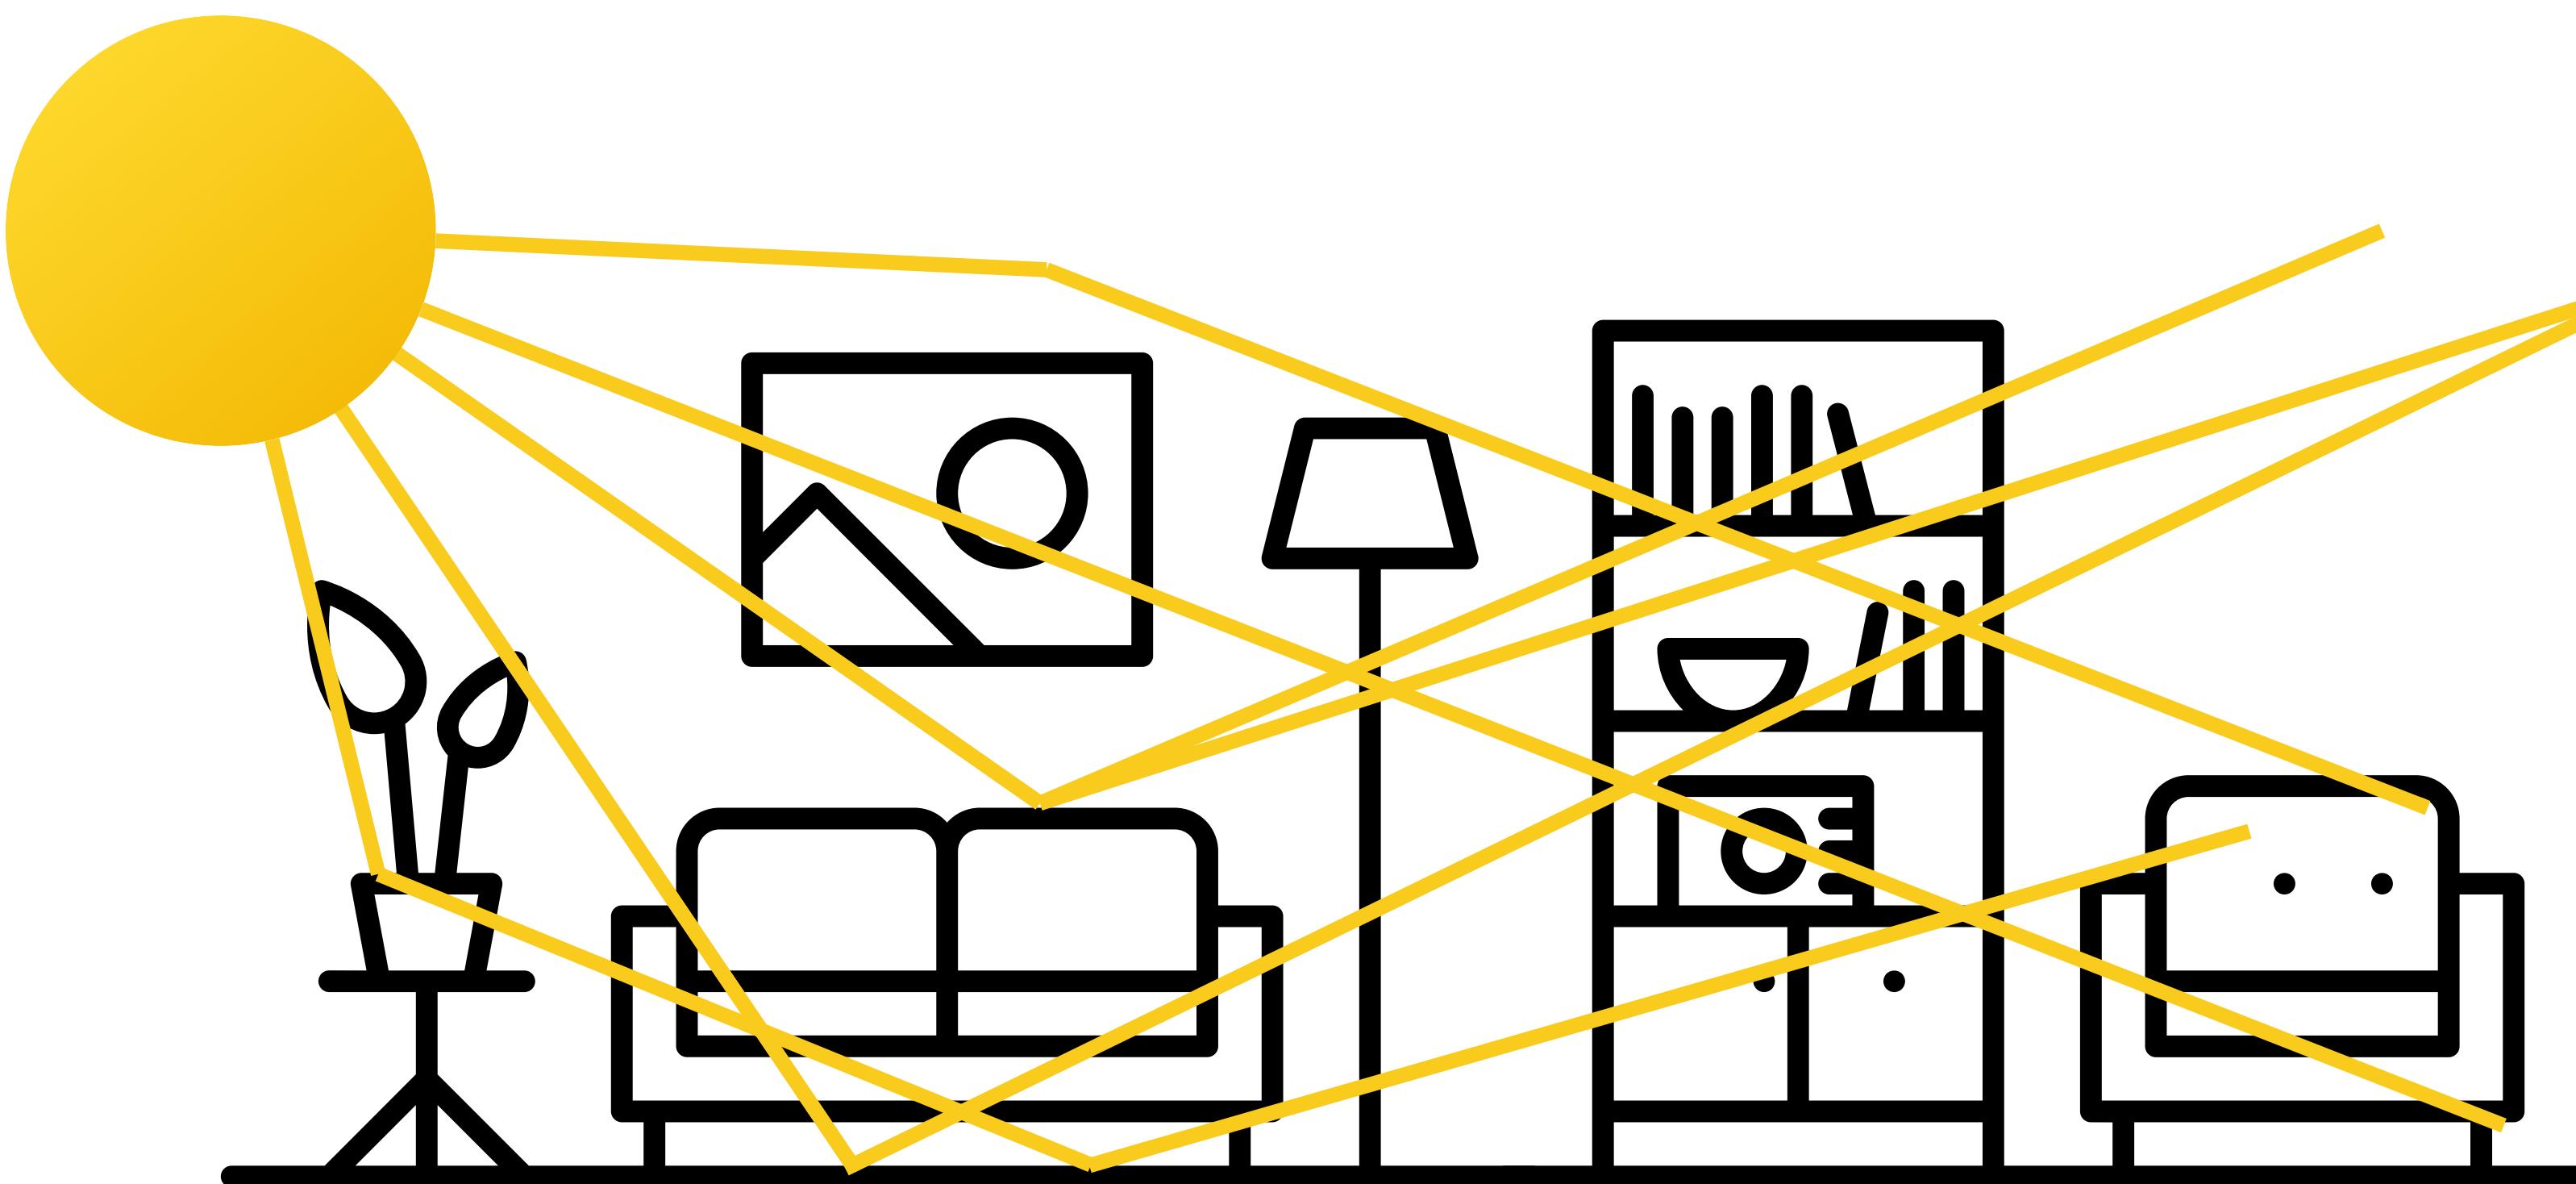
Profs. Bill Freeman, Phillip Isola, Antonio Torralba

What is a 3D scene?



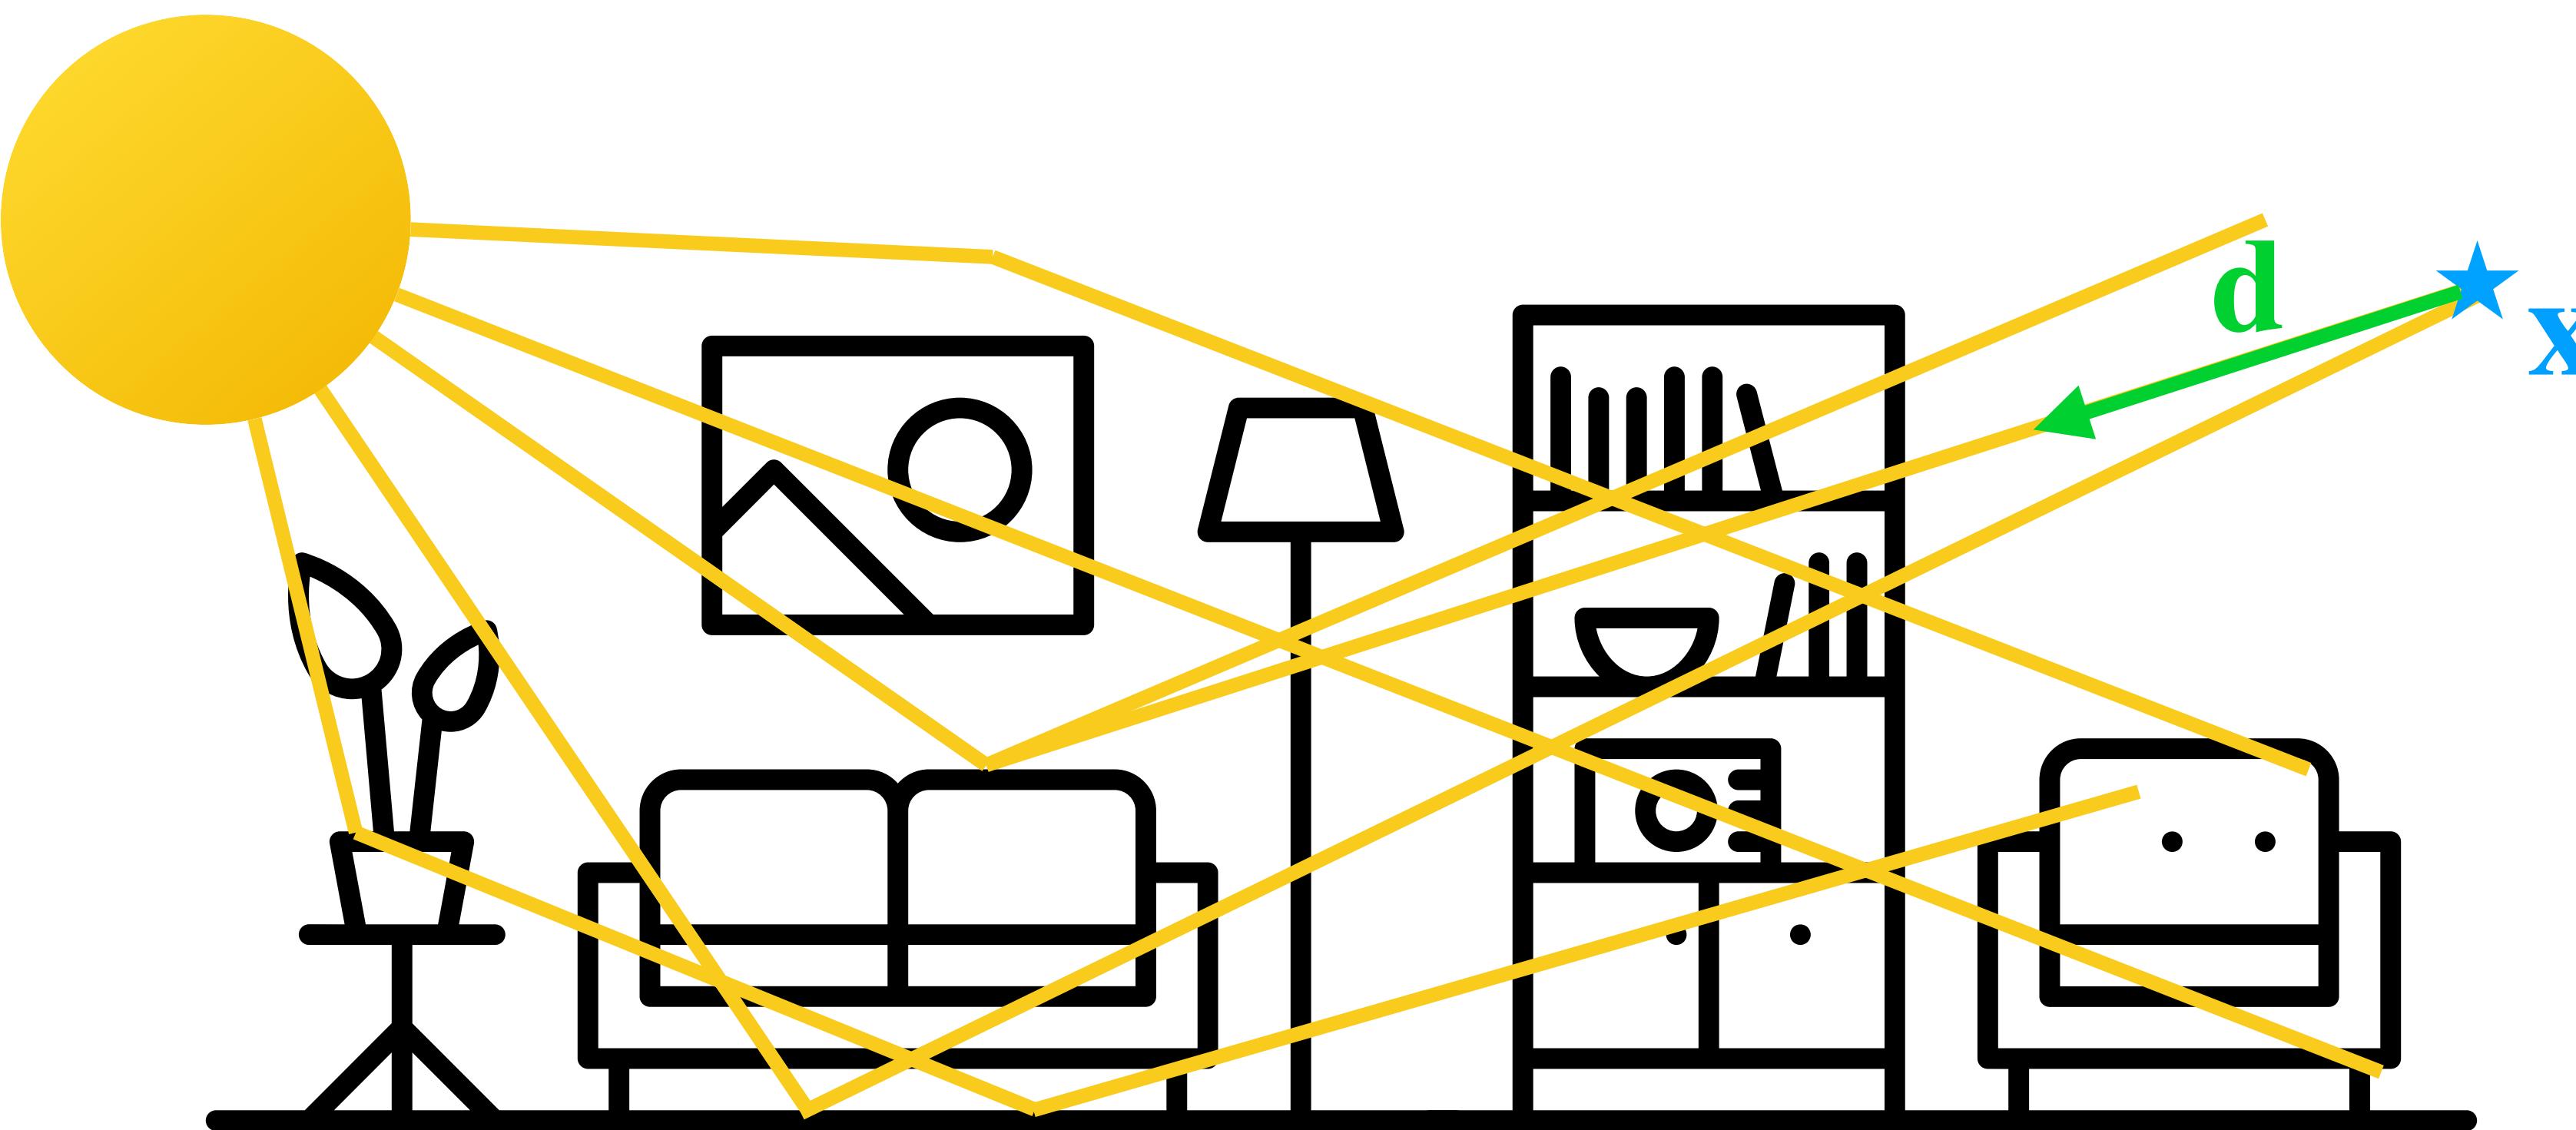
materials, light sources, 3D shape, color, weight, density, friction coefficients, etc

How do we observe scenes?



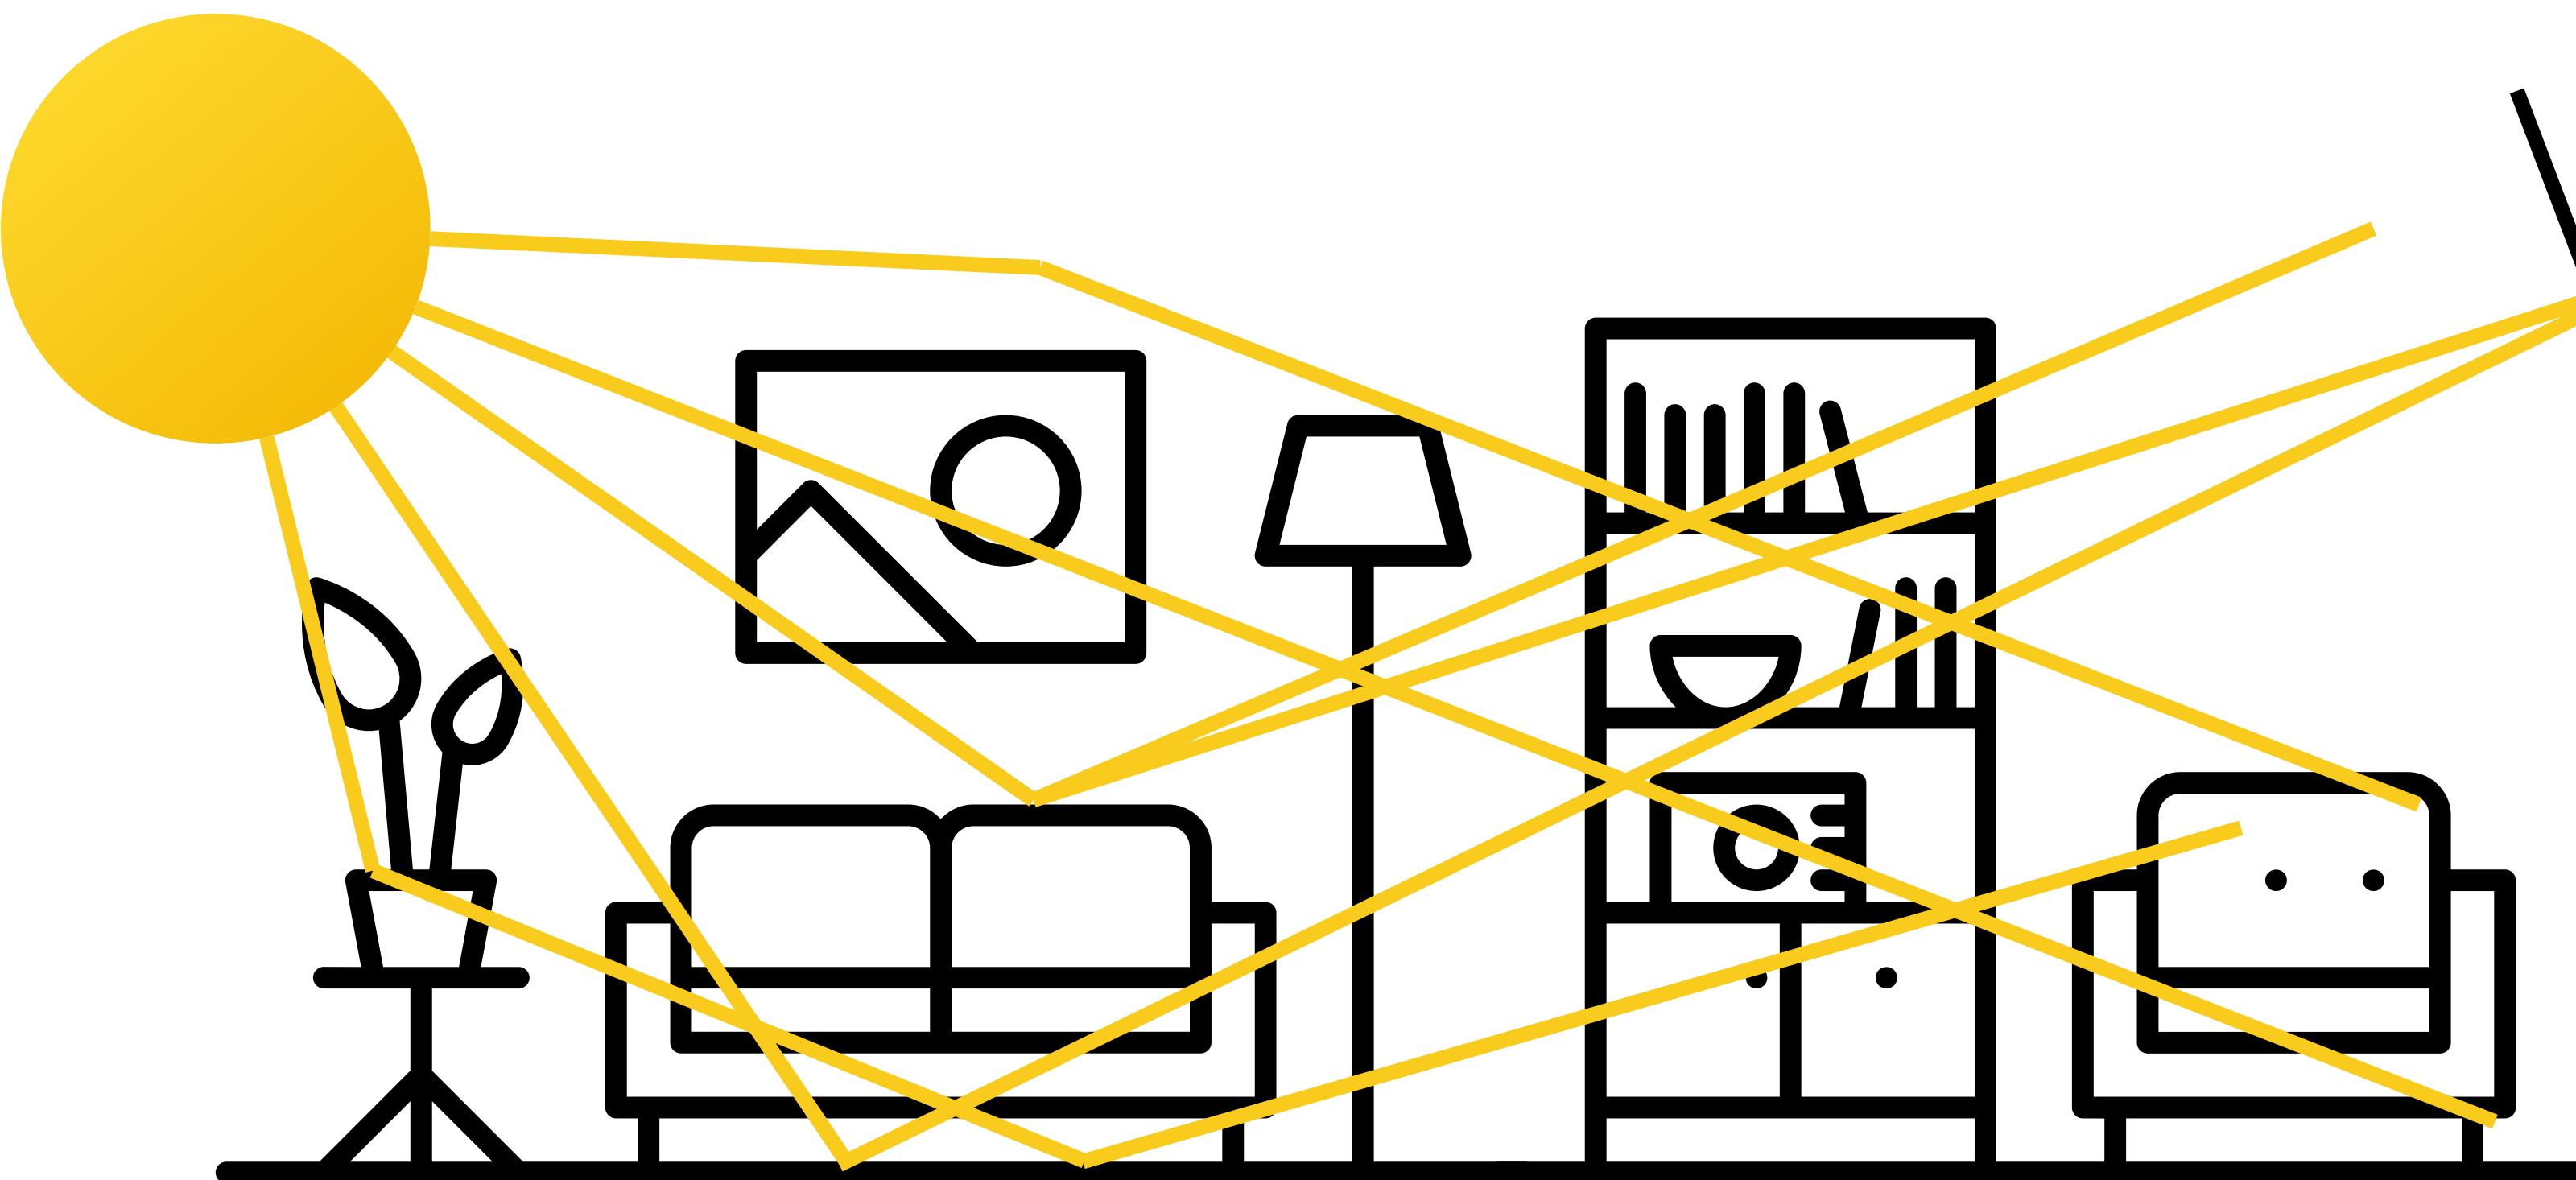
An eye (or a camera) observes a subset of all the light rays in a scene.

The Light Field: 3D coordinate plus ray direction is mapped to the color of that ray.

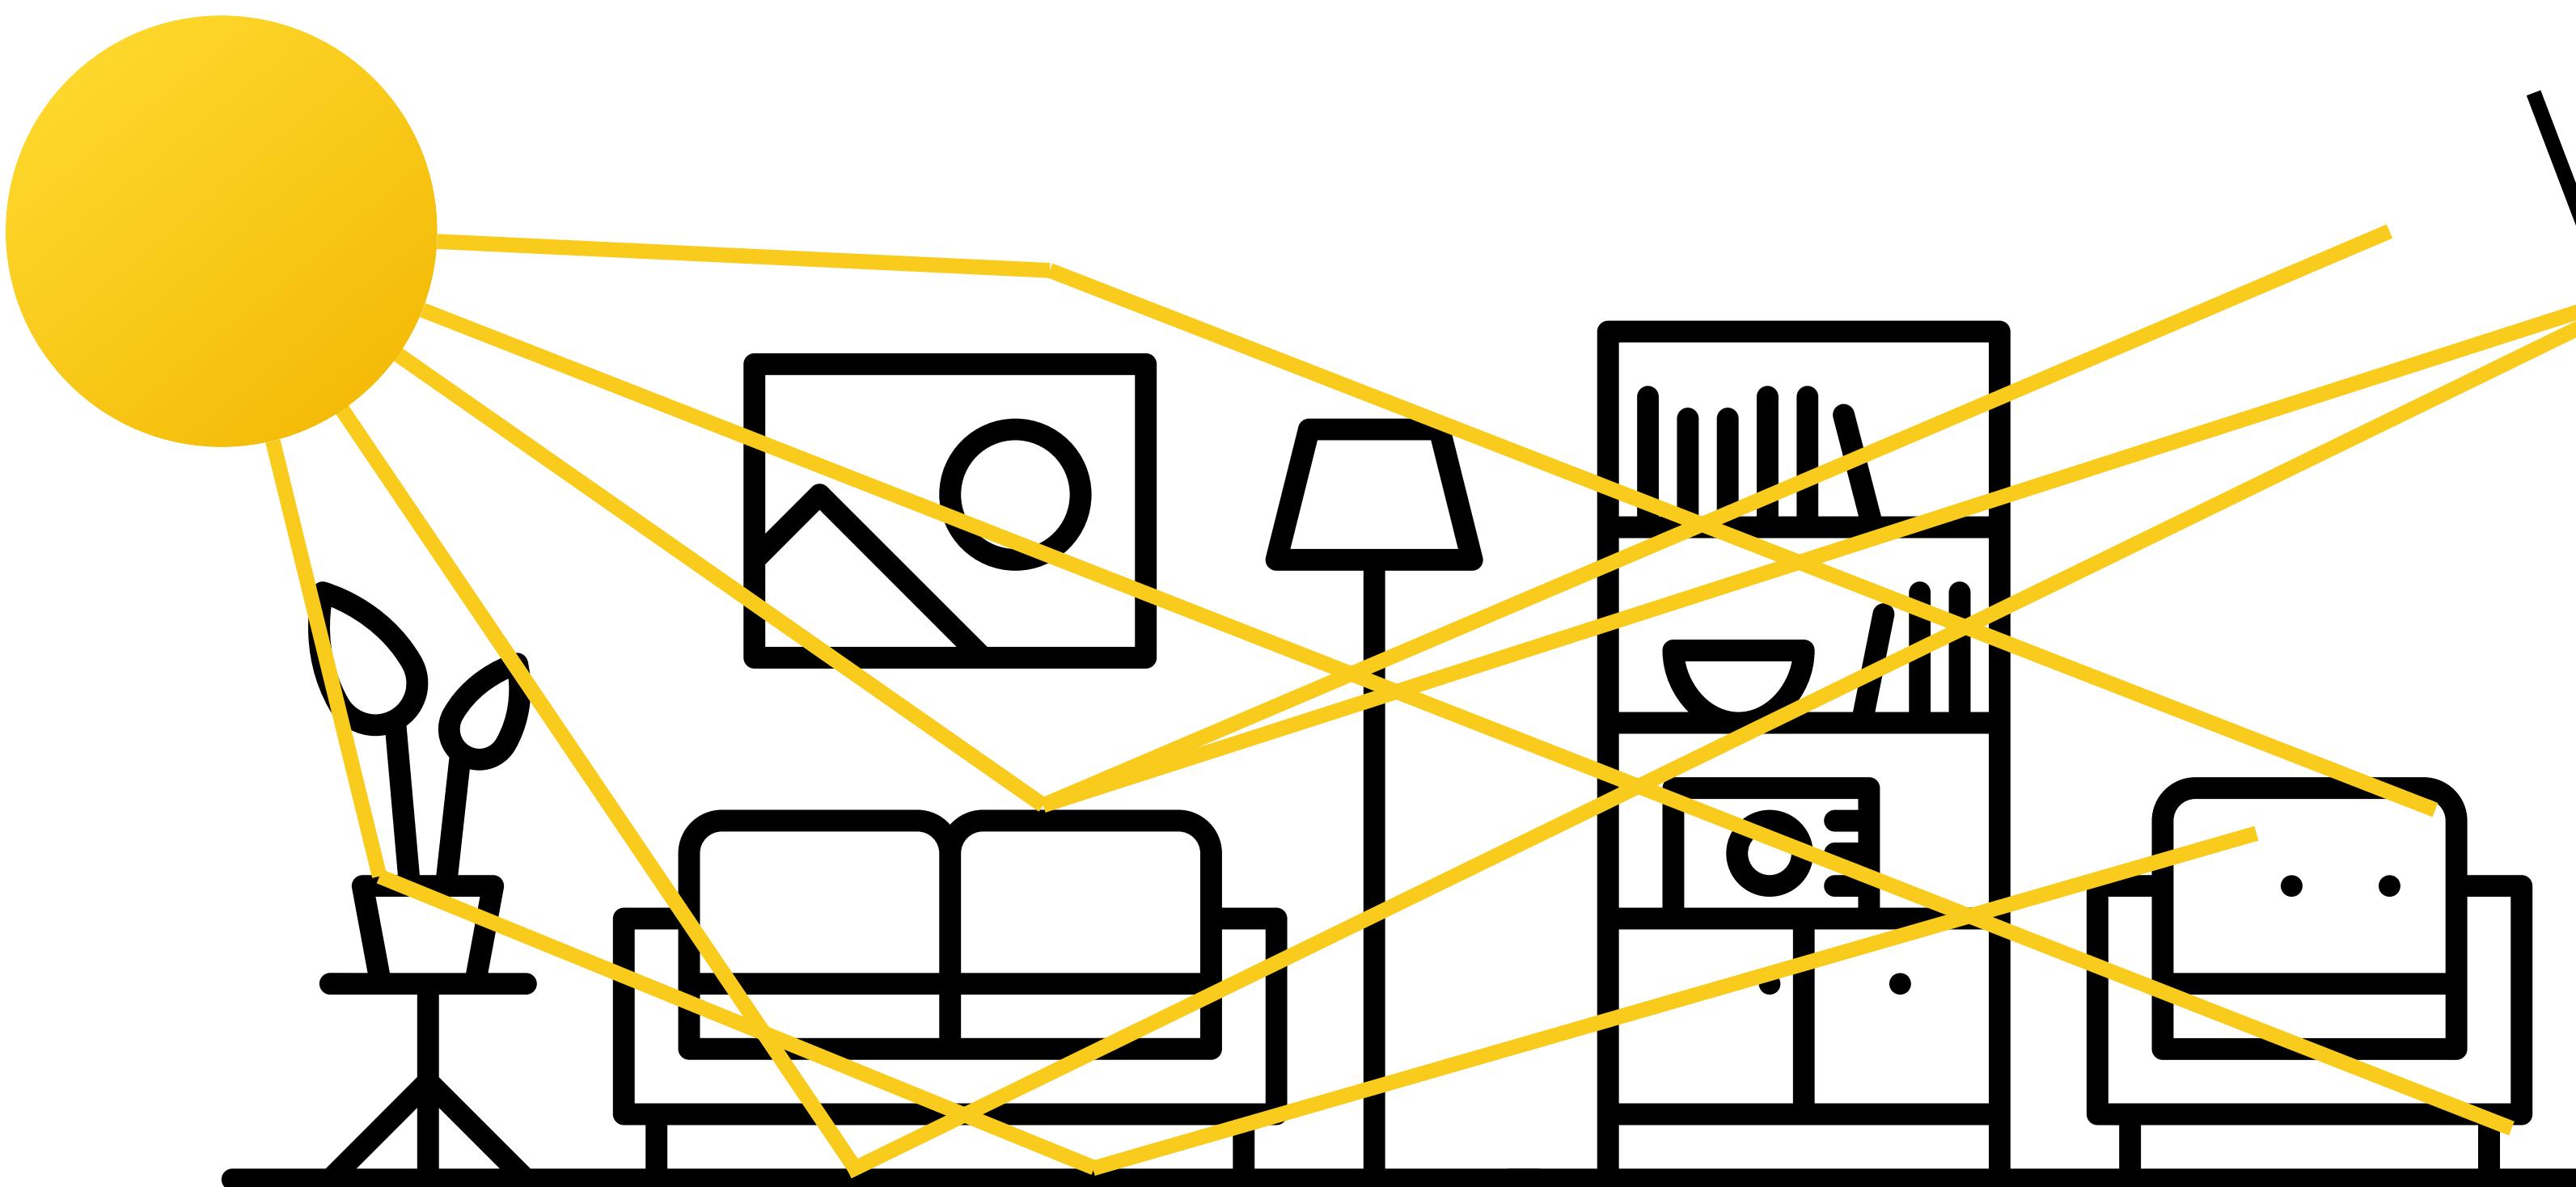


$$LF : \mathbb{R}^3 \times \mathbb{S}^2 \rightarrow \mathbb{R}^3, \quad LF(\mathbf{x}, \mathbf{d}) = \mathbf{c}$$

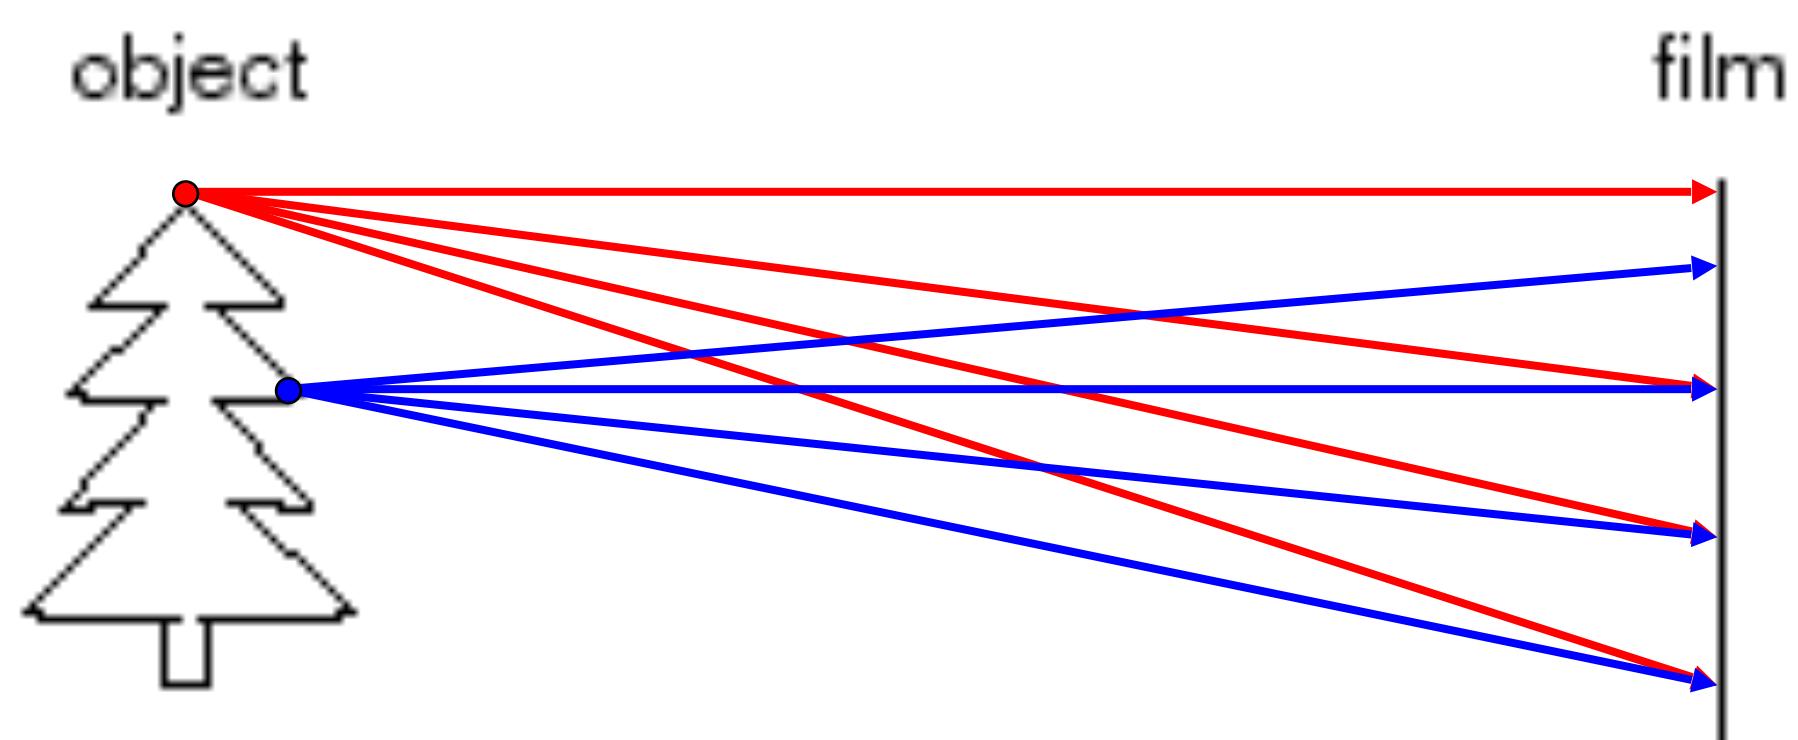
Why don't we get an image if we hold up a piece of paper?



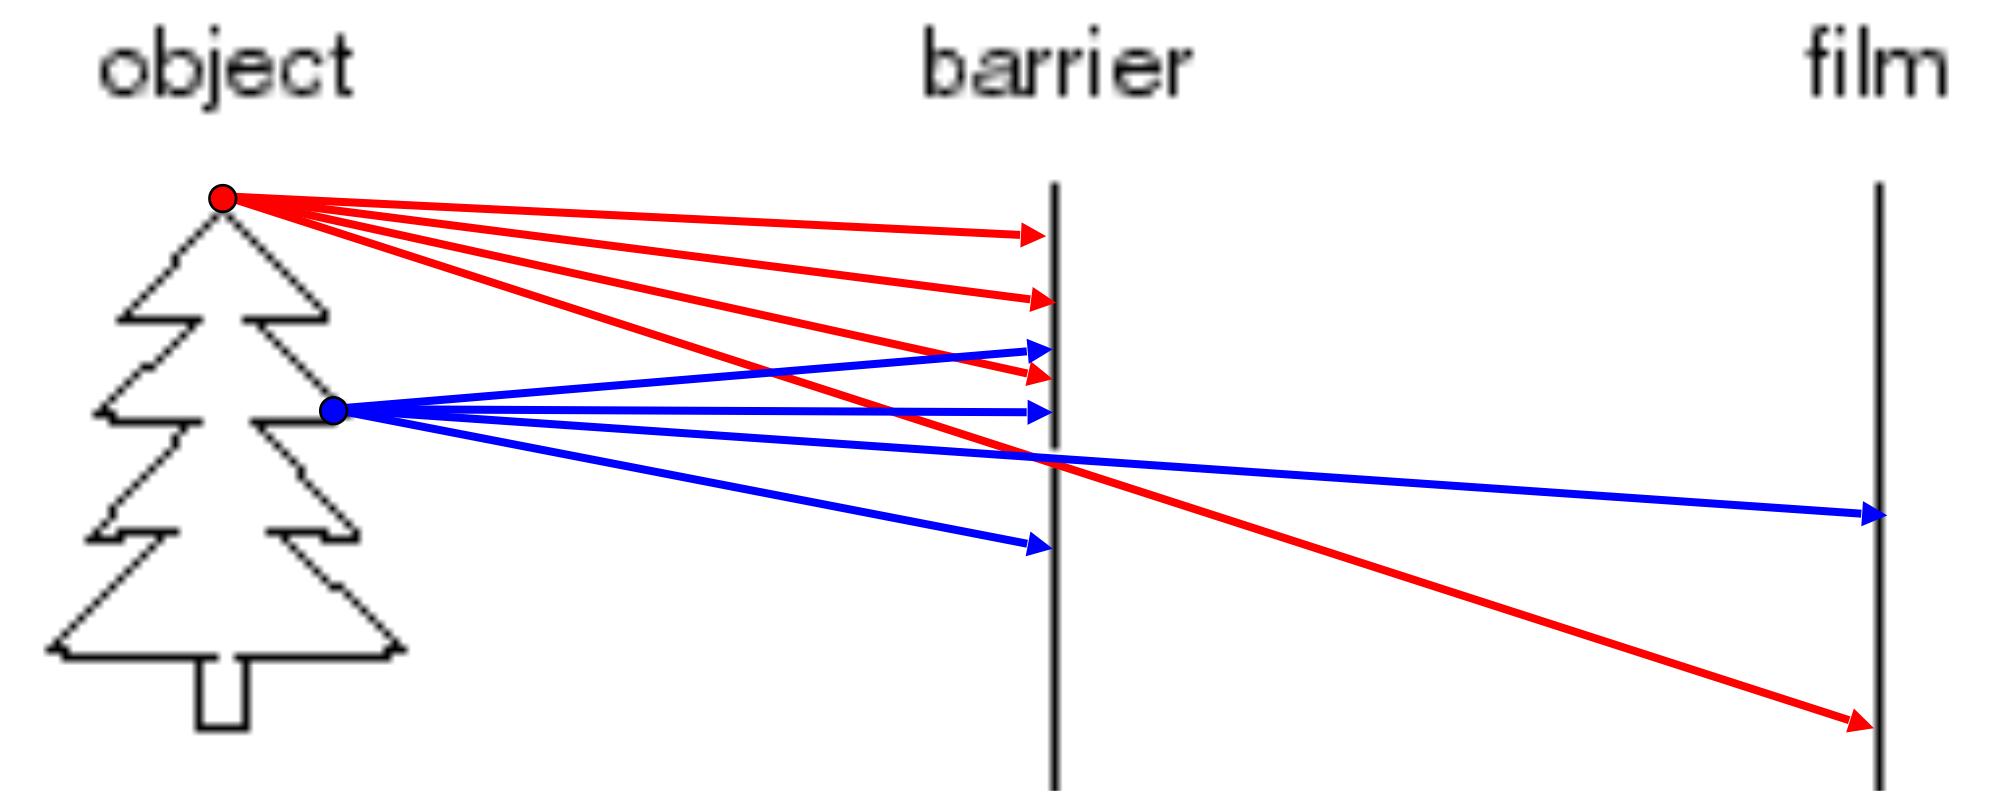
Every “pixel” on the paper is the average of all the rays in the scene.



Let's make a Camera!



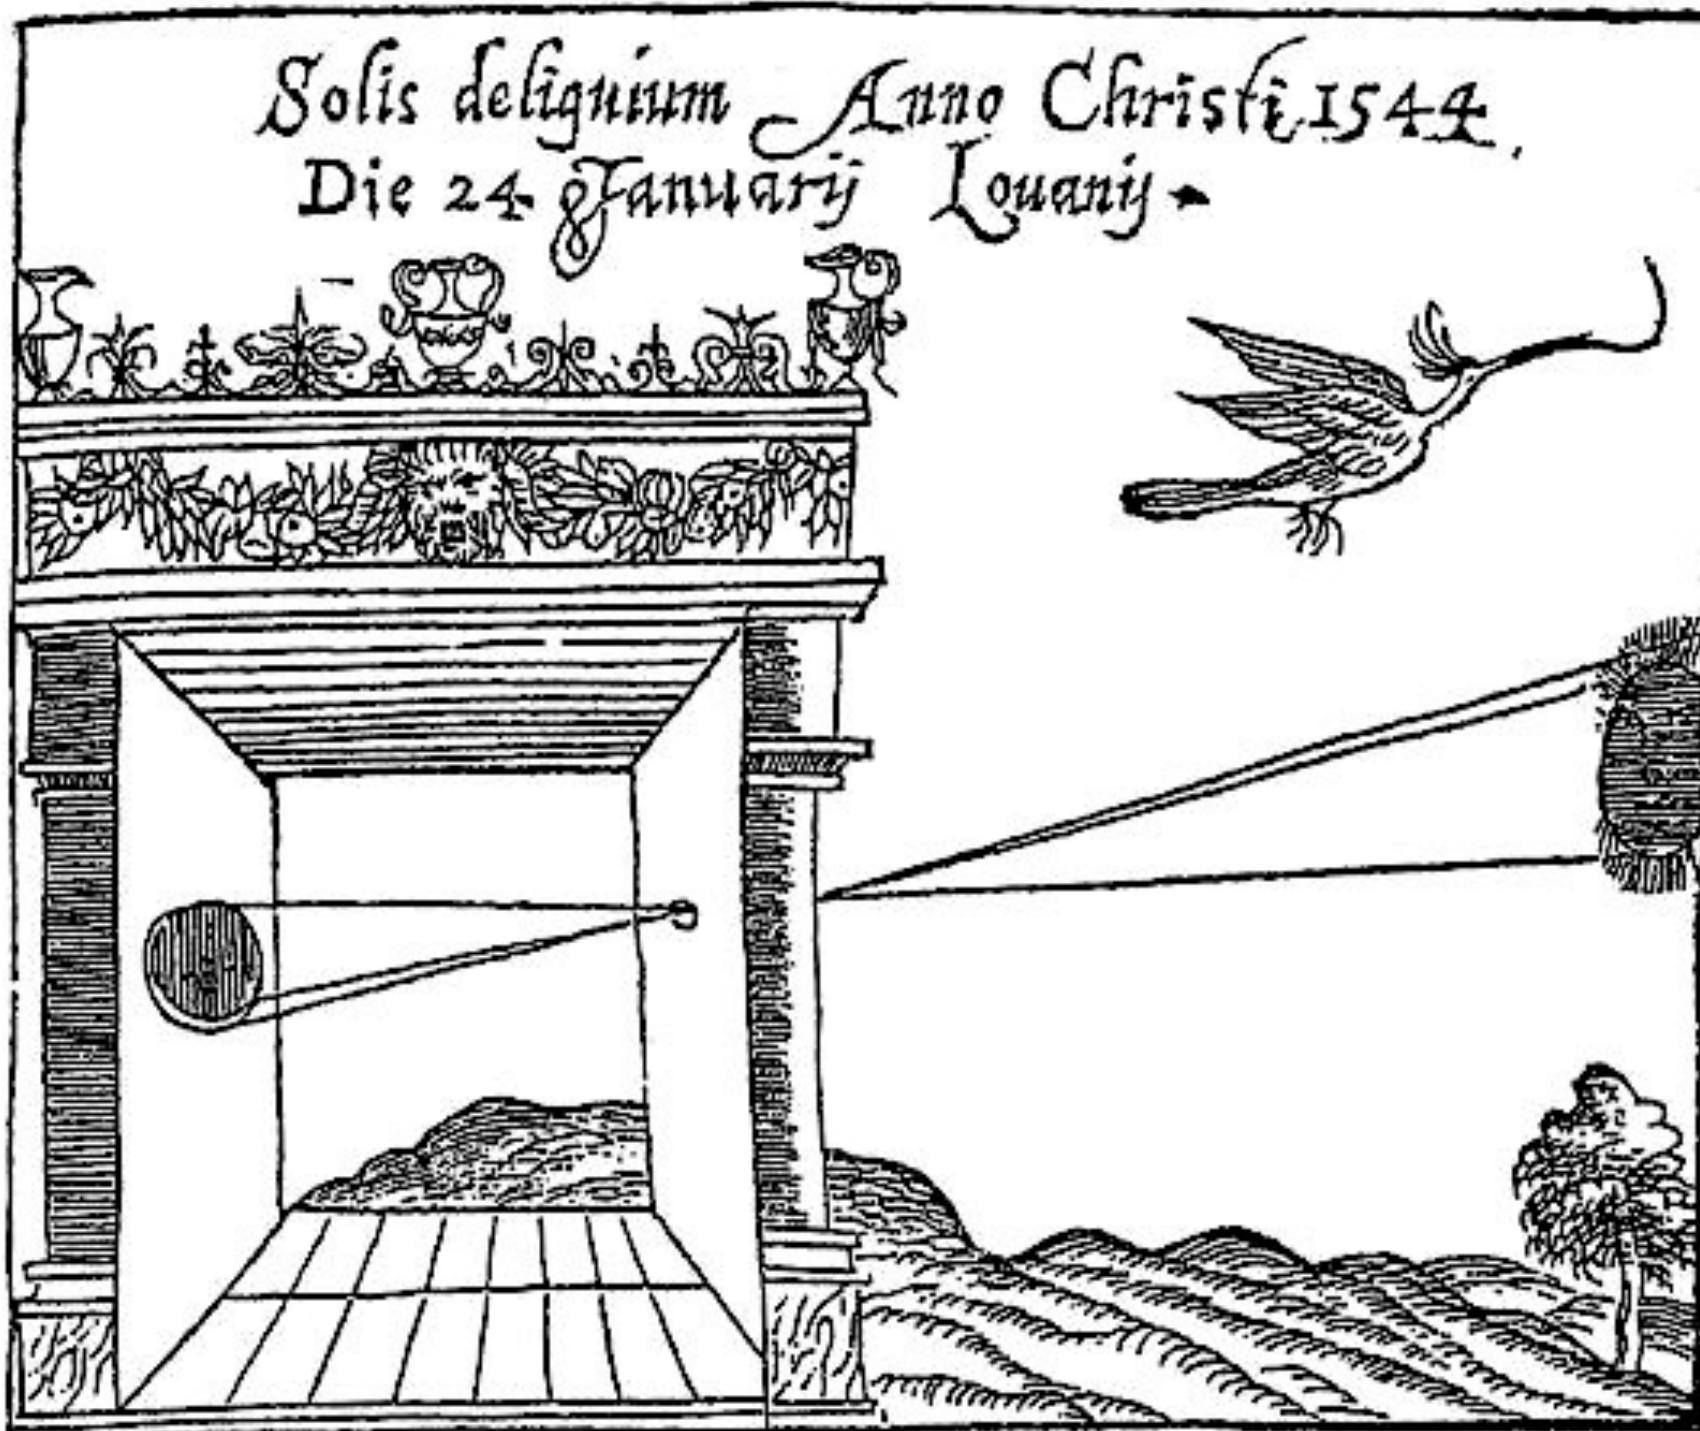
Idea 1: Put piece of film in front of an object



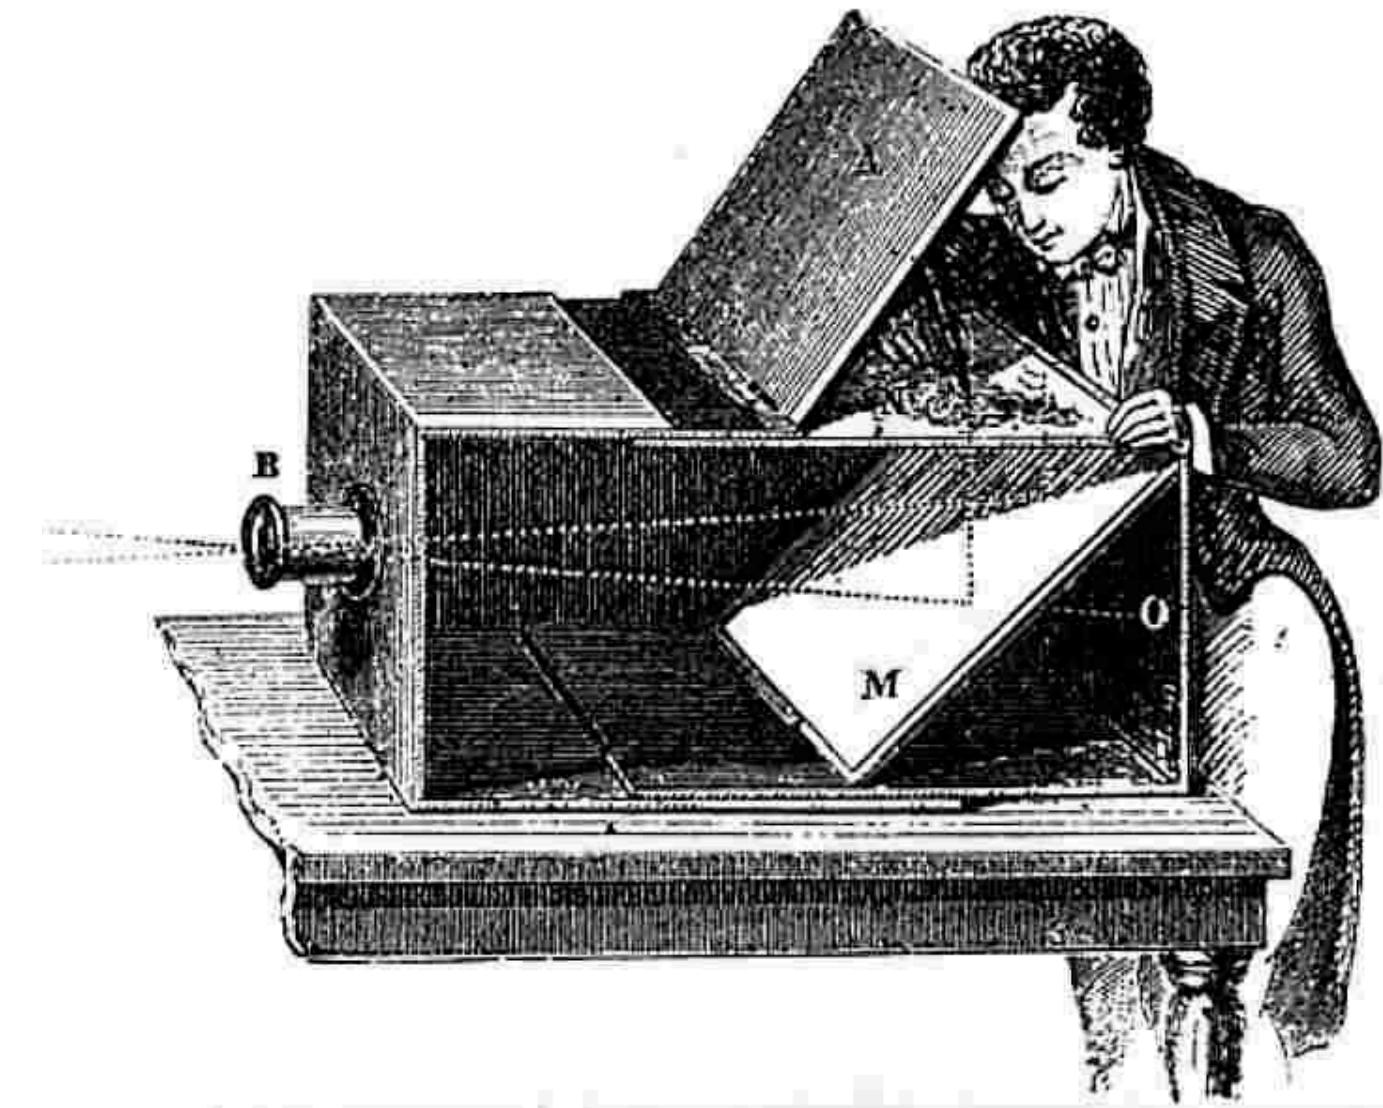
Add a barrier to block most rays.

Camera Obscura

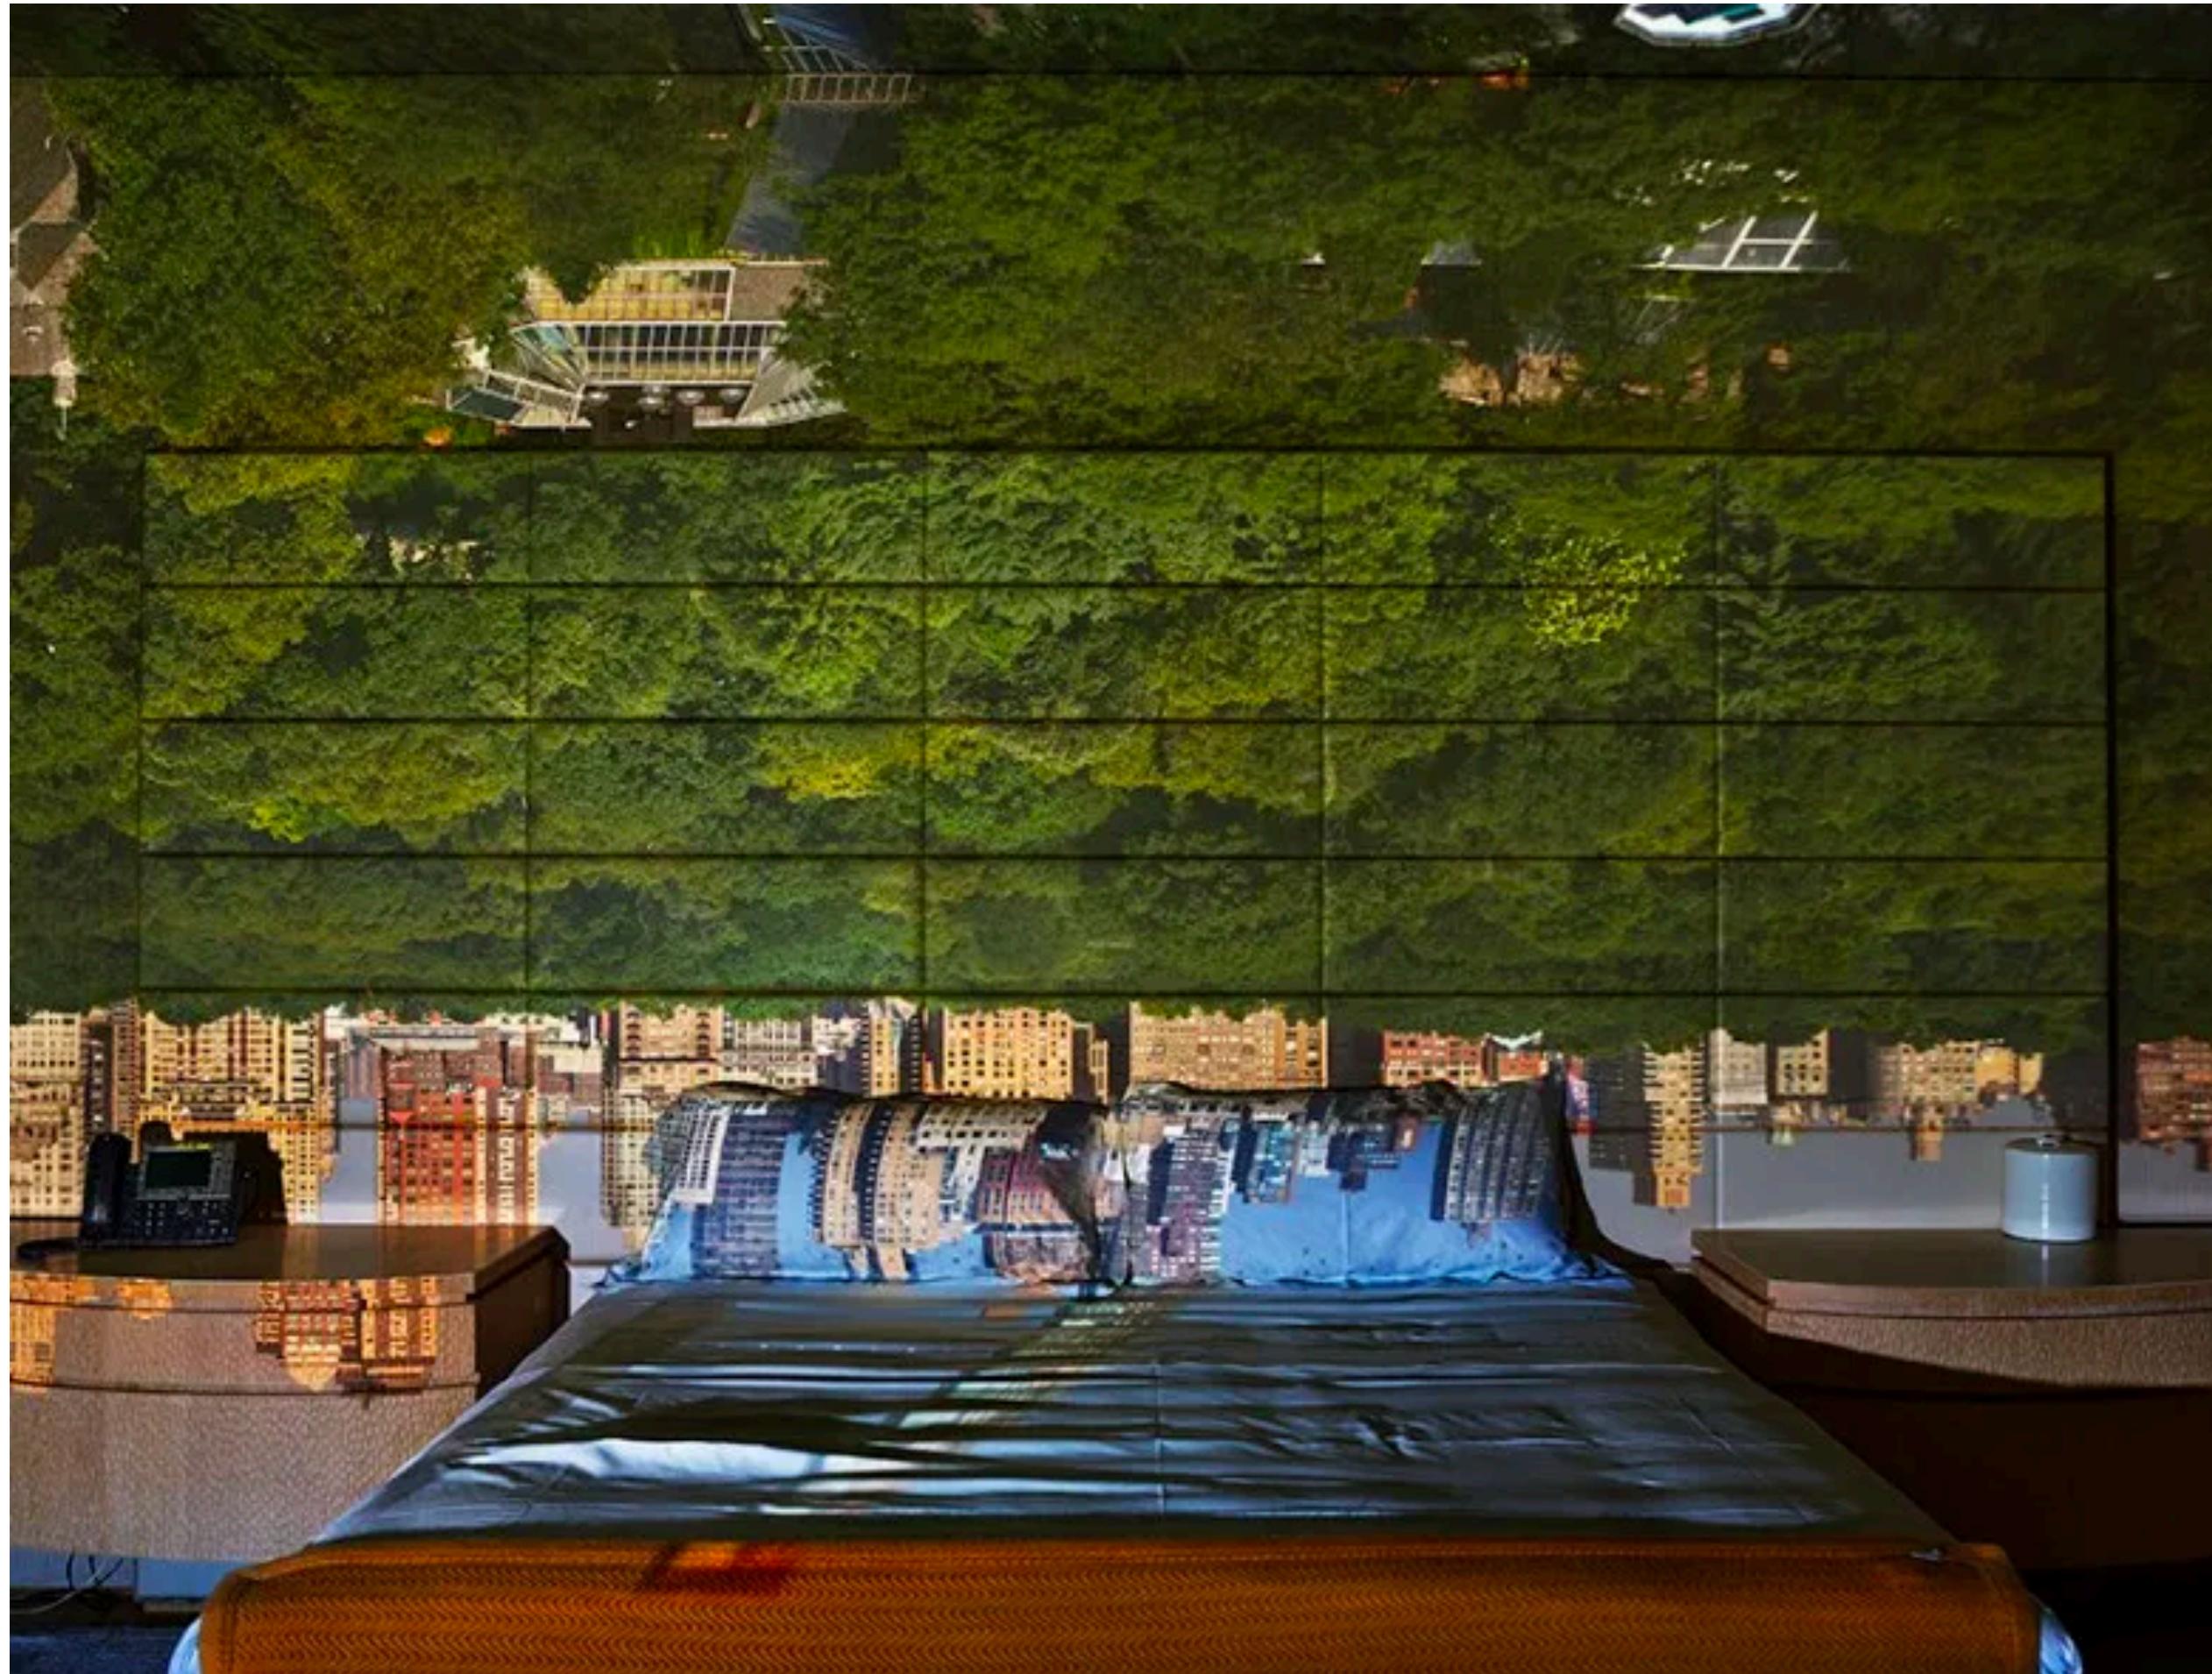
- Basic principle known to Mozi (470-390 BCE), Aristotle (384-322 BCE)
- Drawing aid for artists: described by Leonardo da Vinci (1452-1519)



Gemma Frisius, 1558



Camera Obscura (*Dark Room*)



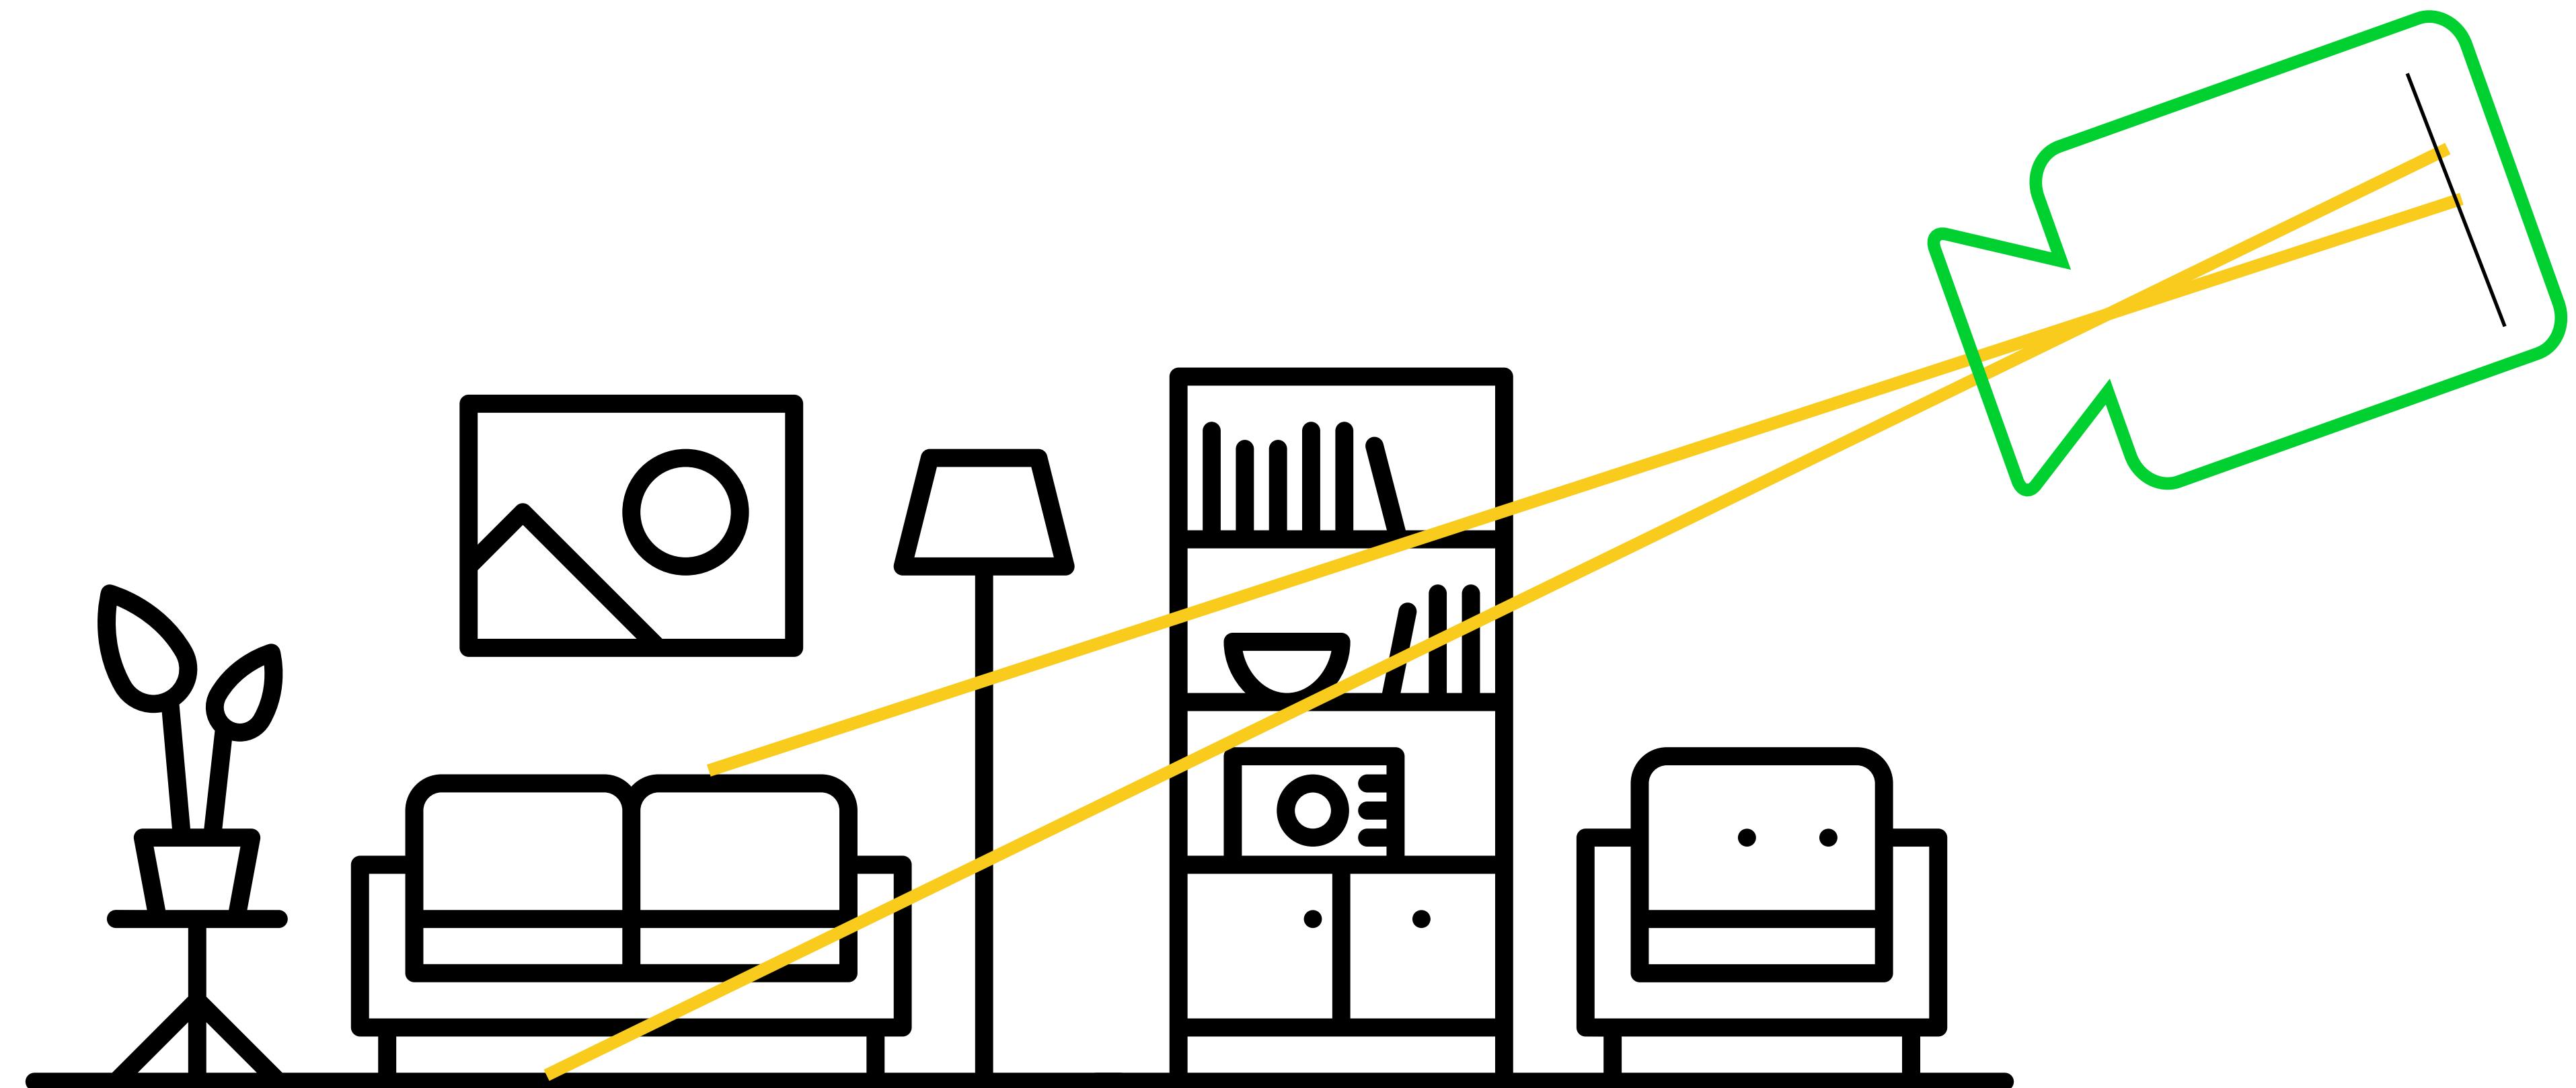
Camera Obscura: View of Central Park
Looking West in Bedroom. Summer, 2018

After scouting rooms and reserving one for at least a day, Morell masks the windows except for the aperture. He controls three elements: the size of the hole, with a smaller one yielding a sharper but dimmer image; the length of the exposure, usually eight hours; and the distance from the hole to the surface on which the outside image falls and which he will photograph. He used 4 x 5 and 8 x 10 view cameras and lenses ranging from 75 to 150 mm.

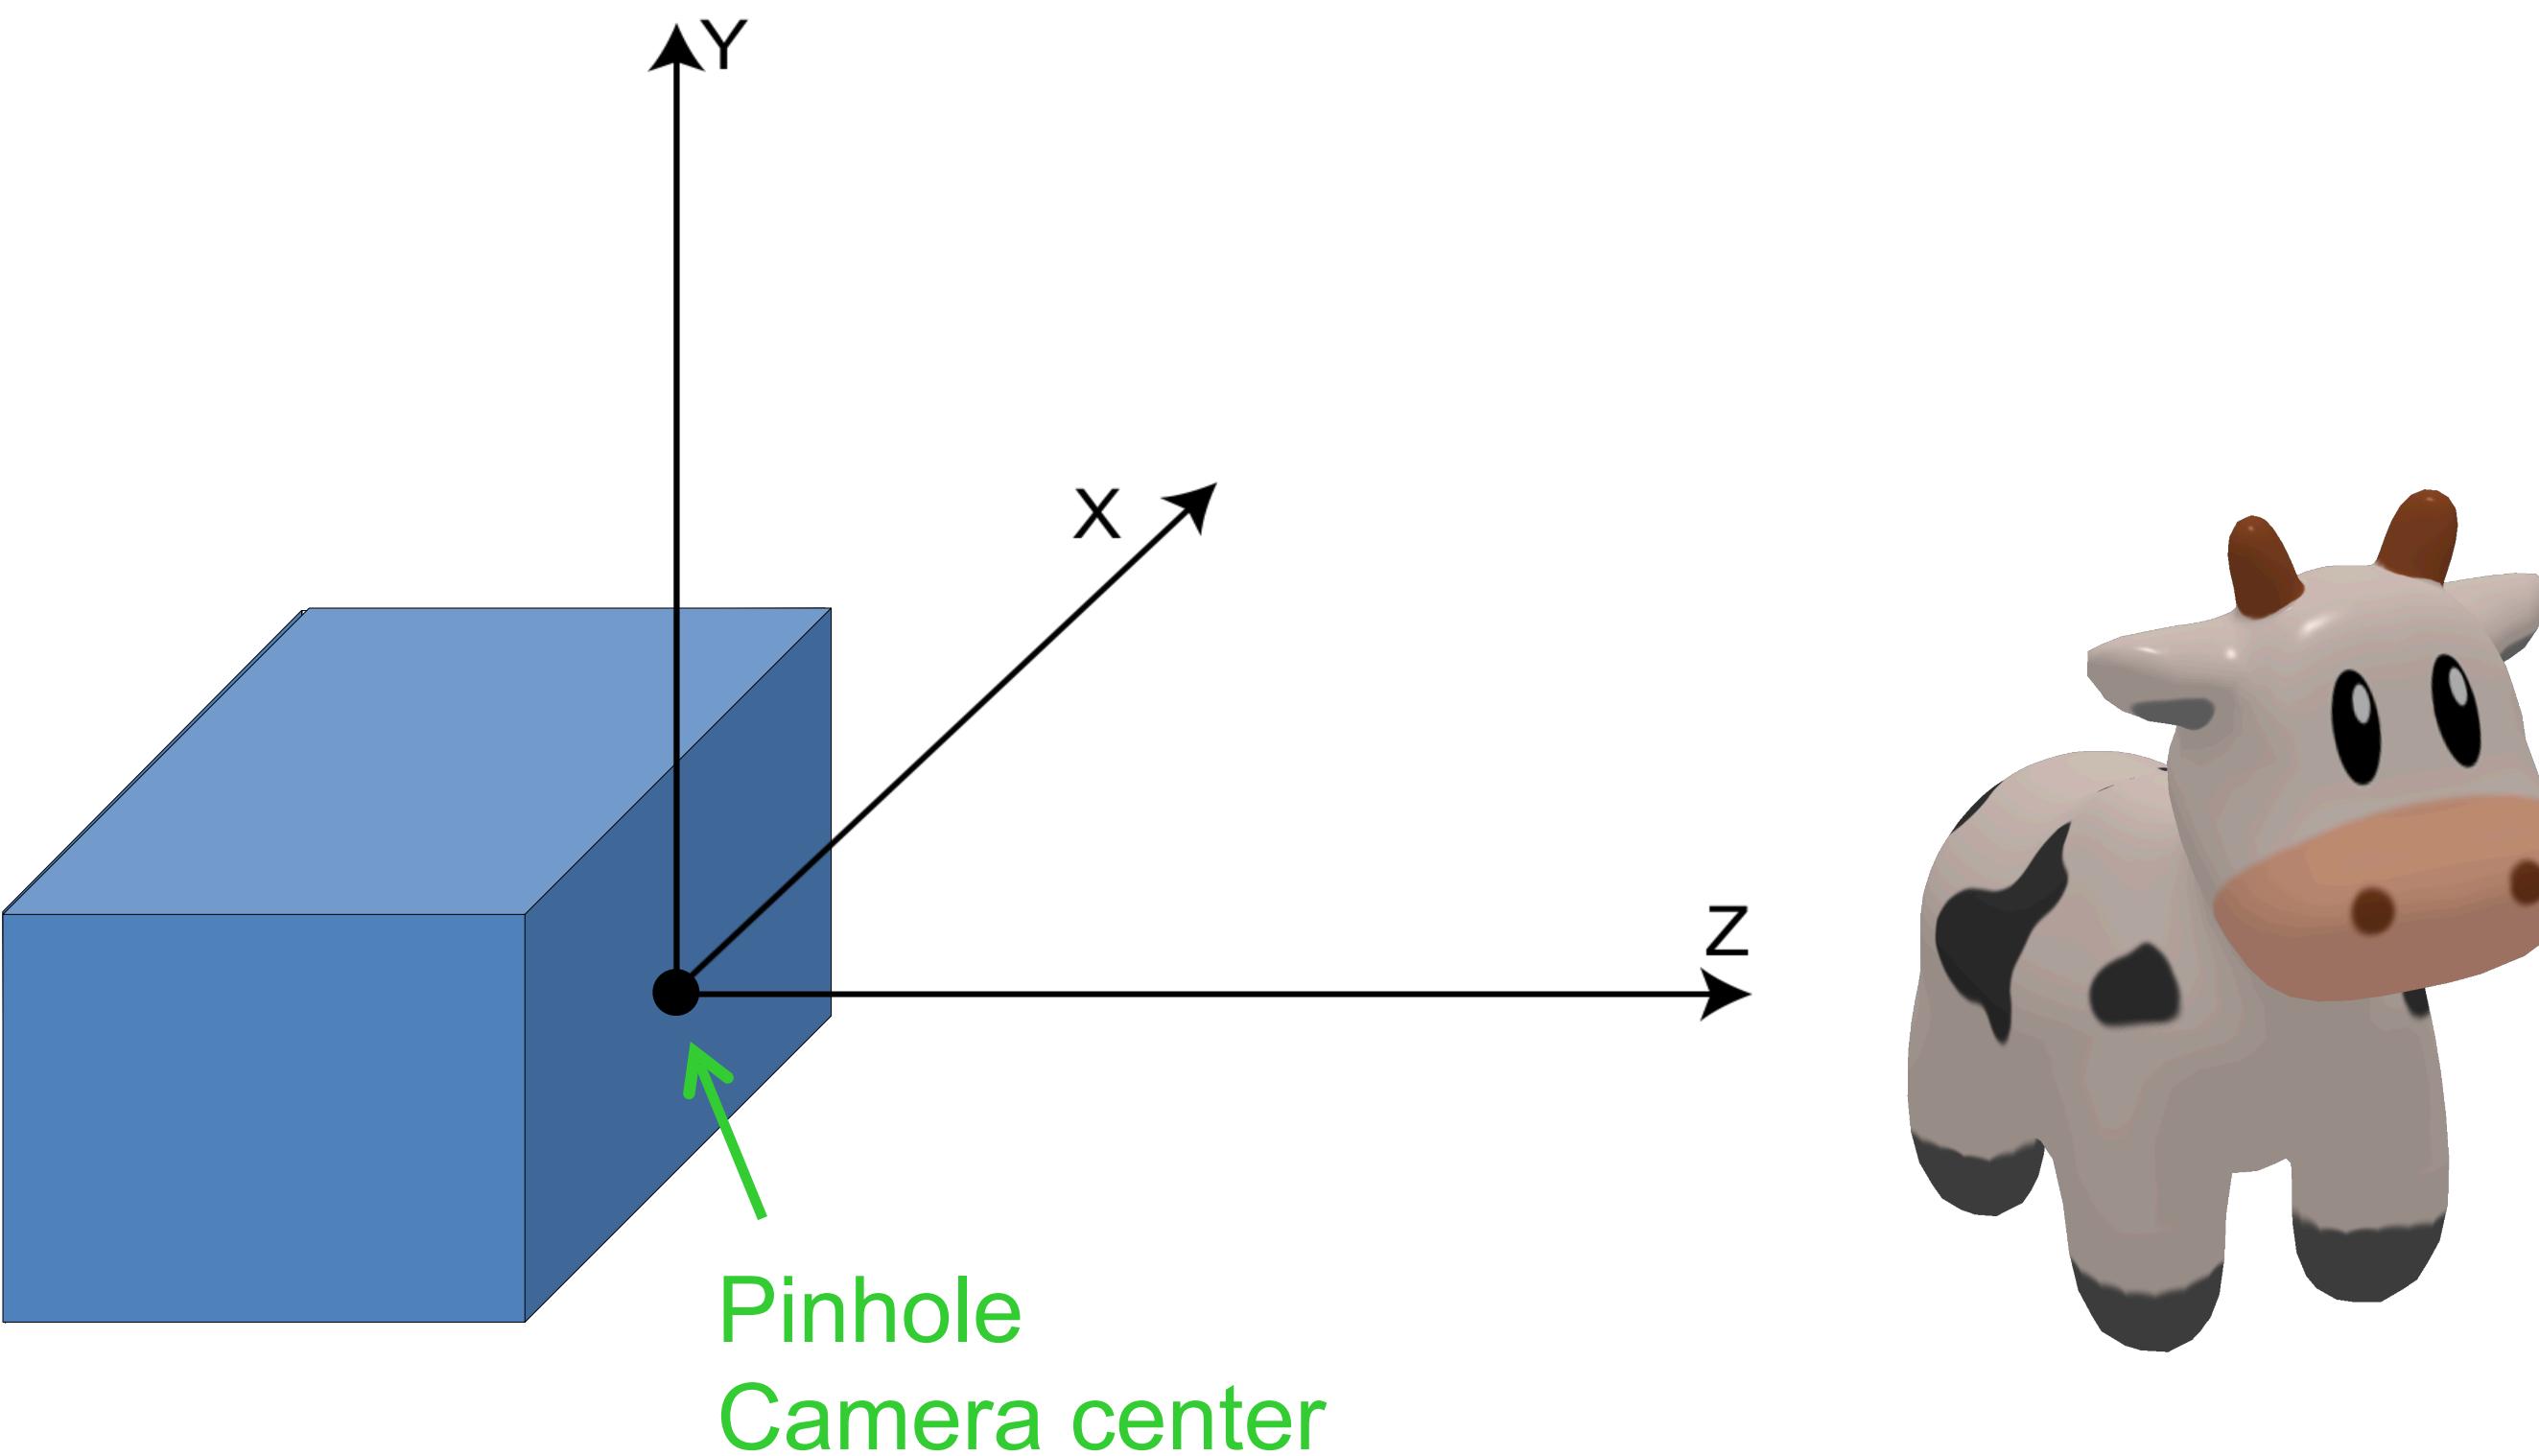
After he's done inside, it gets harder. "I leave the room and I am constantly checking the weather, I'm hoping the maid reads my note not to come in, I'm worrying that the sun will hit the plastic masking and it will fall down, or that I didn't trigger the lens."

[http://www.abelardomorell.net/
project/camera-obscura/](http://www.abelardomorell.net/project/camera-obscura/)

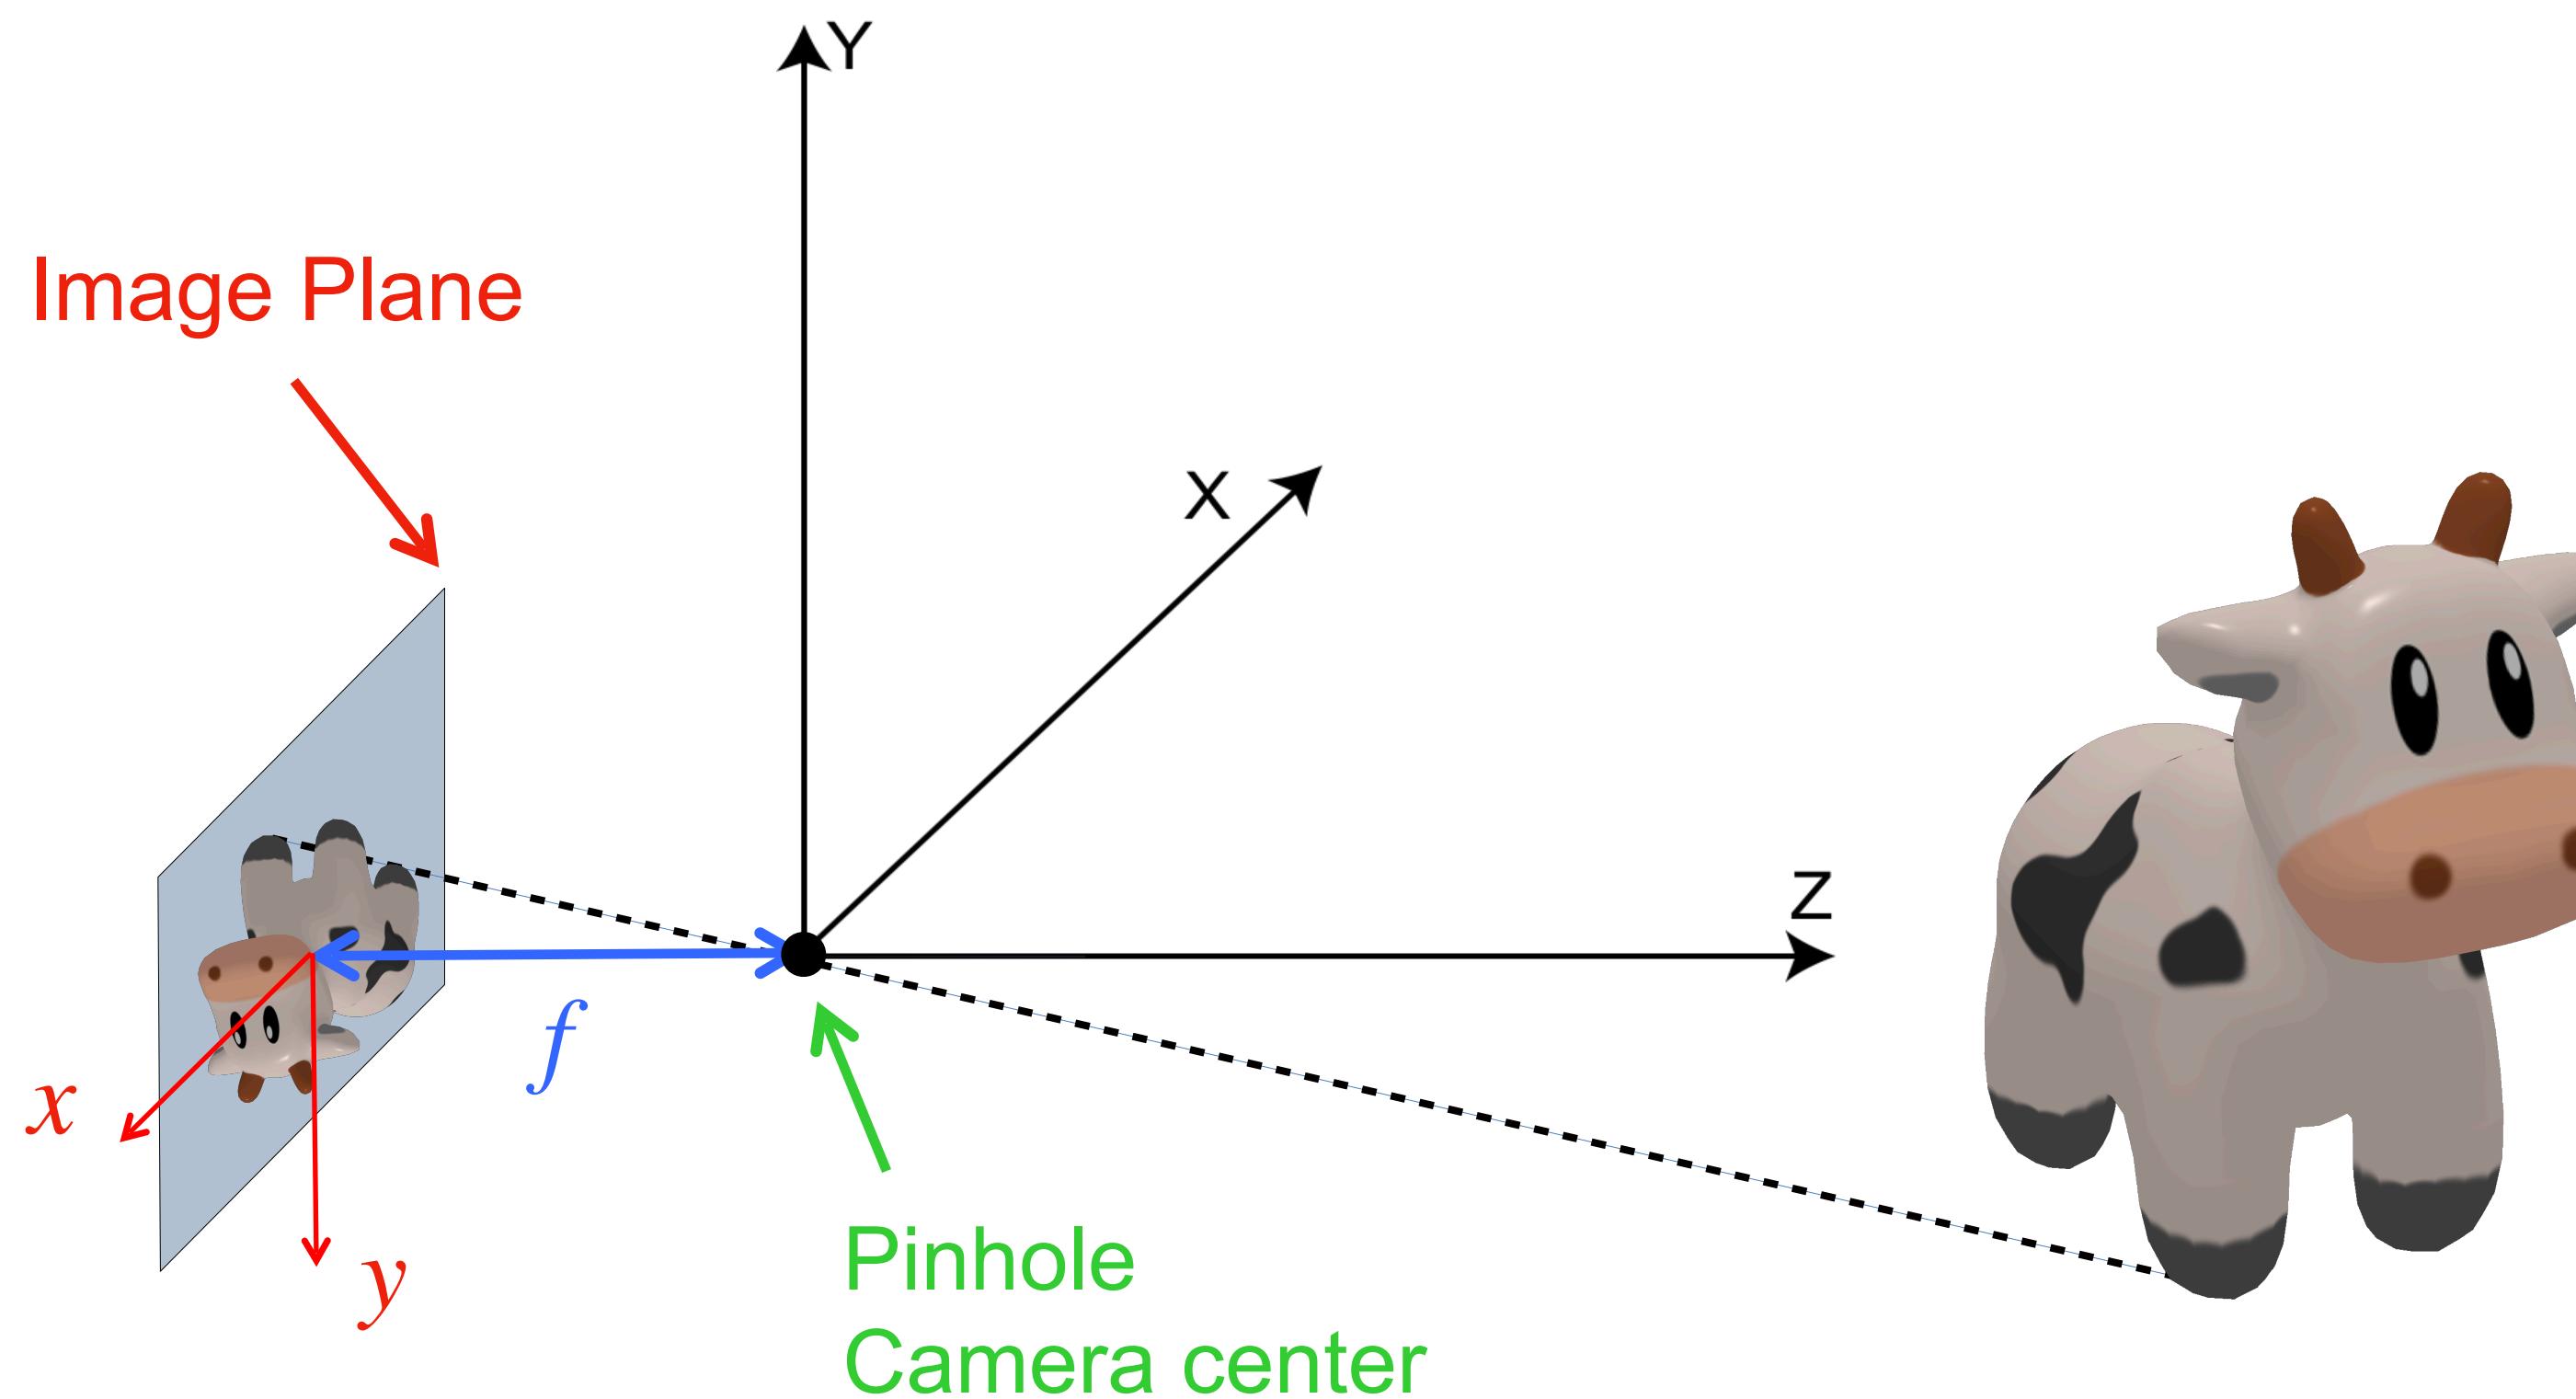
Today: Model the camera, describe the rays.



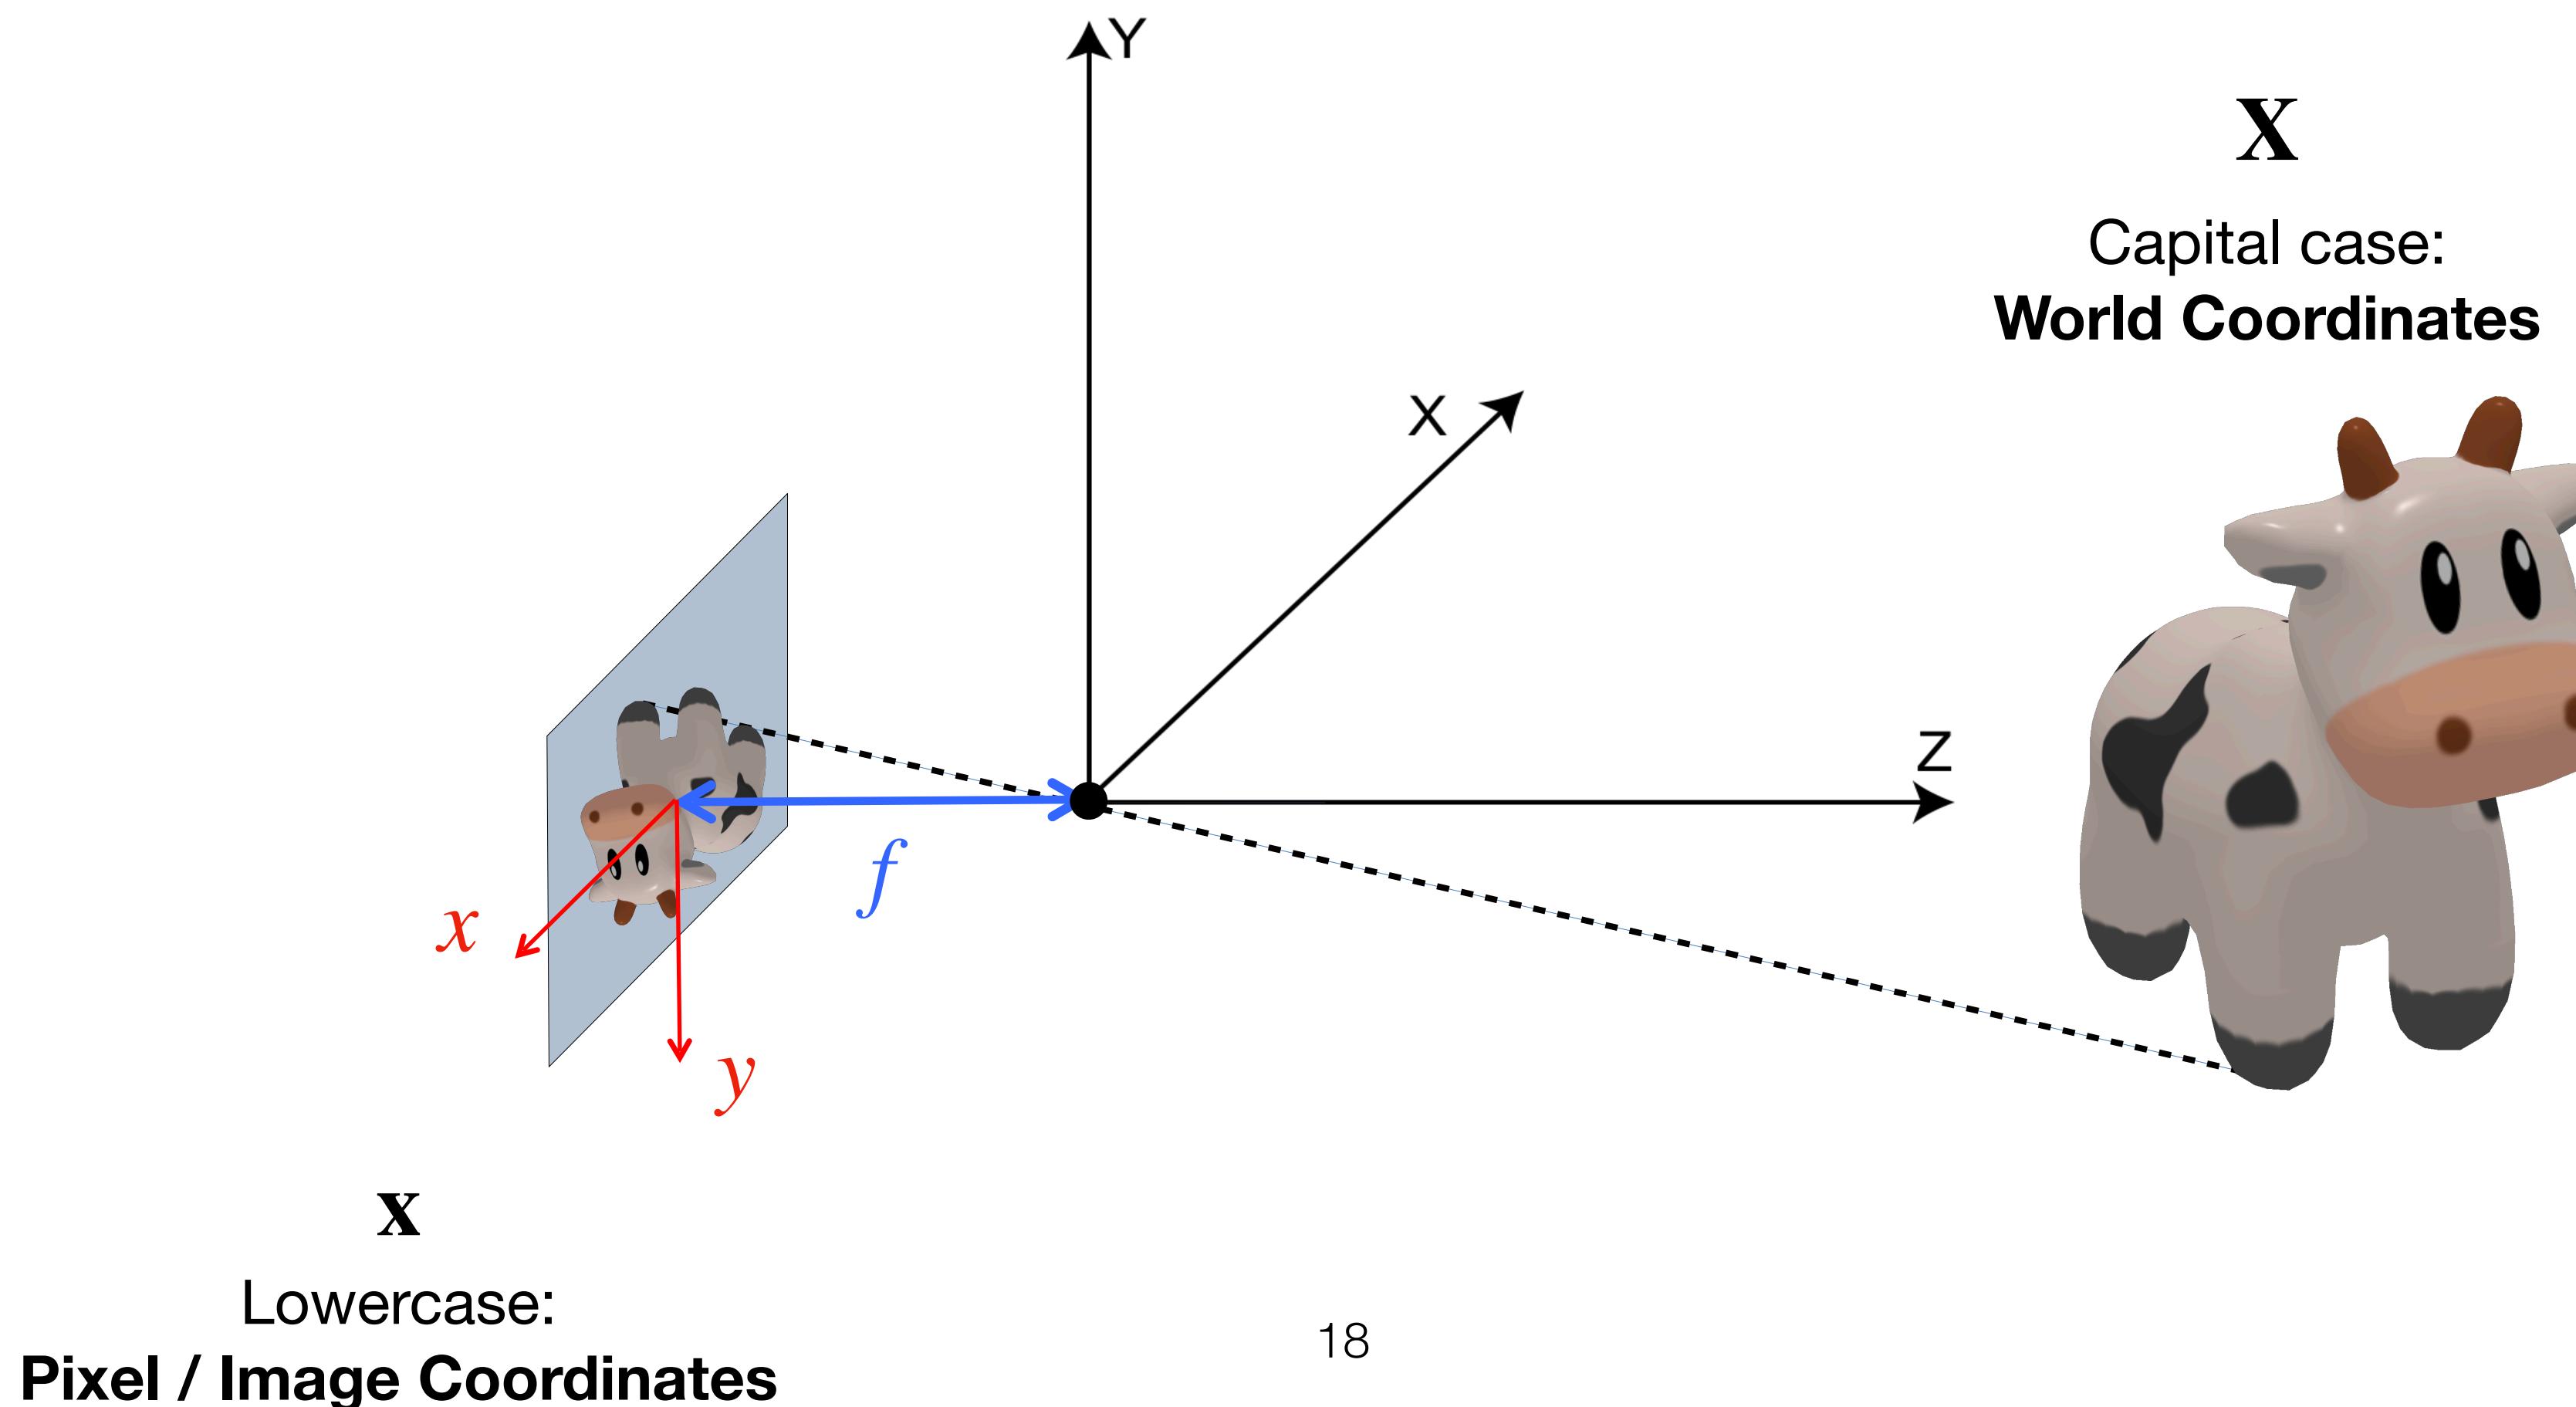
Perspective projection



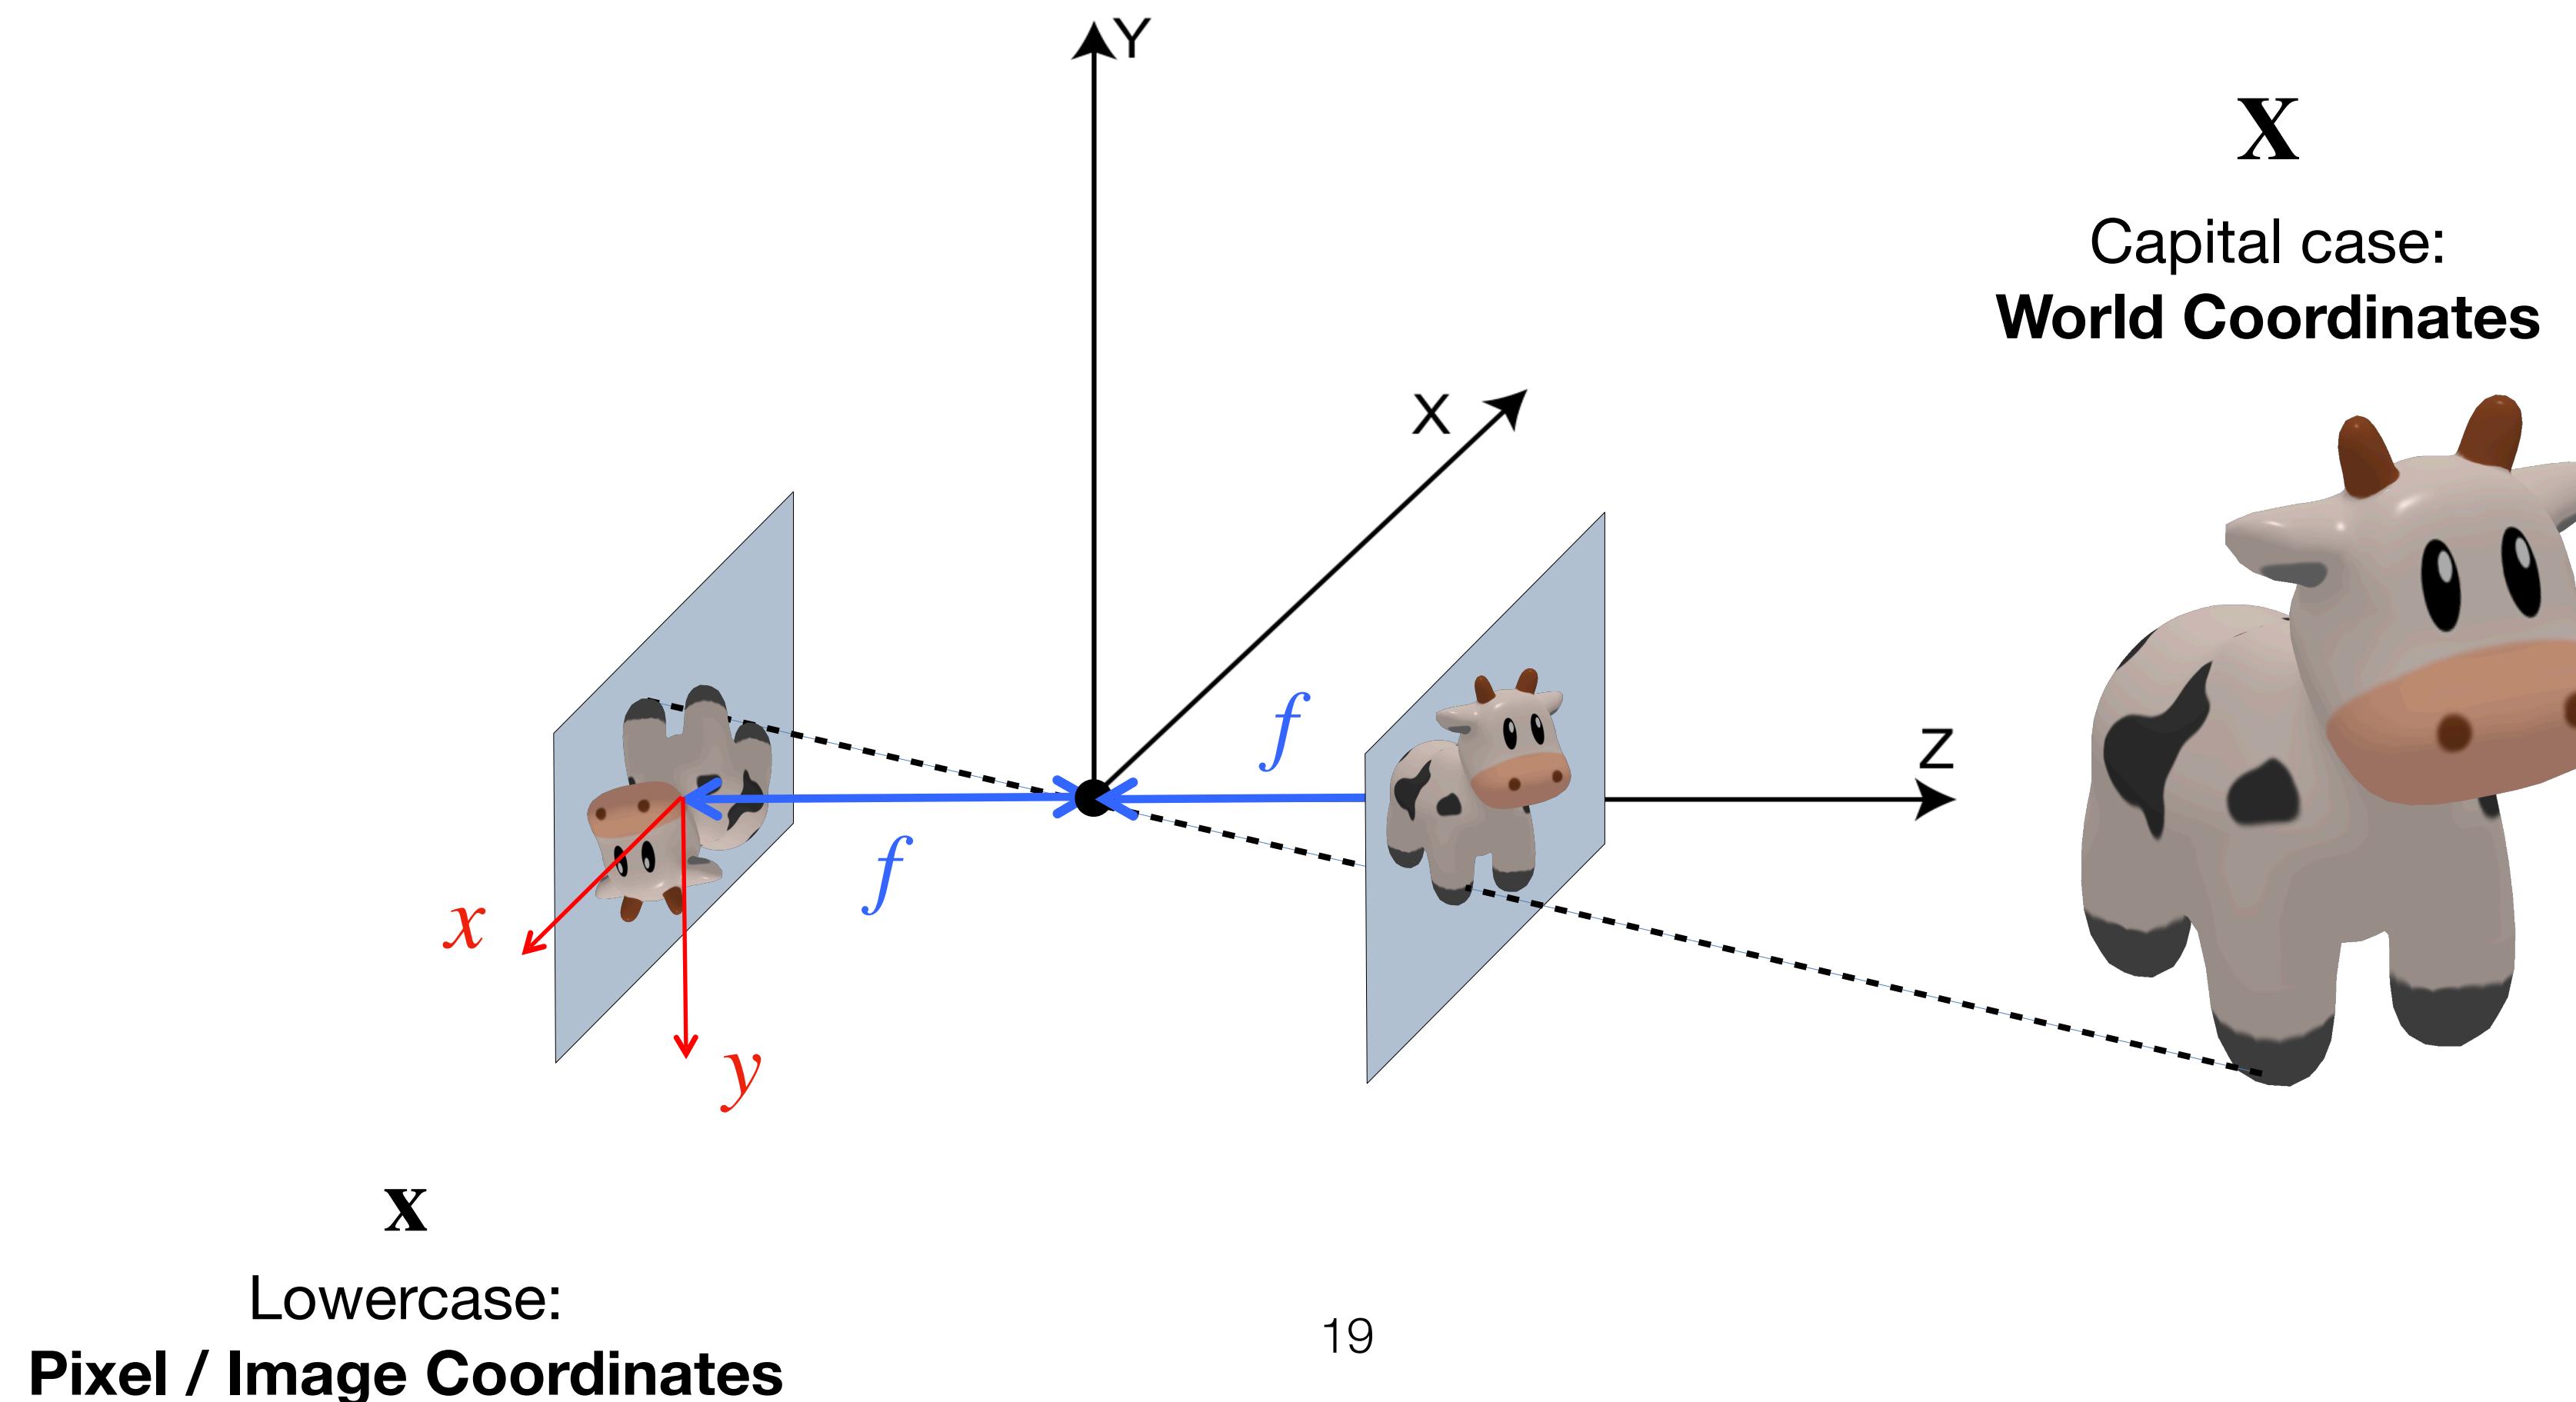
Perspective projection



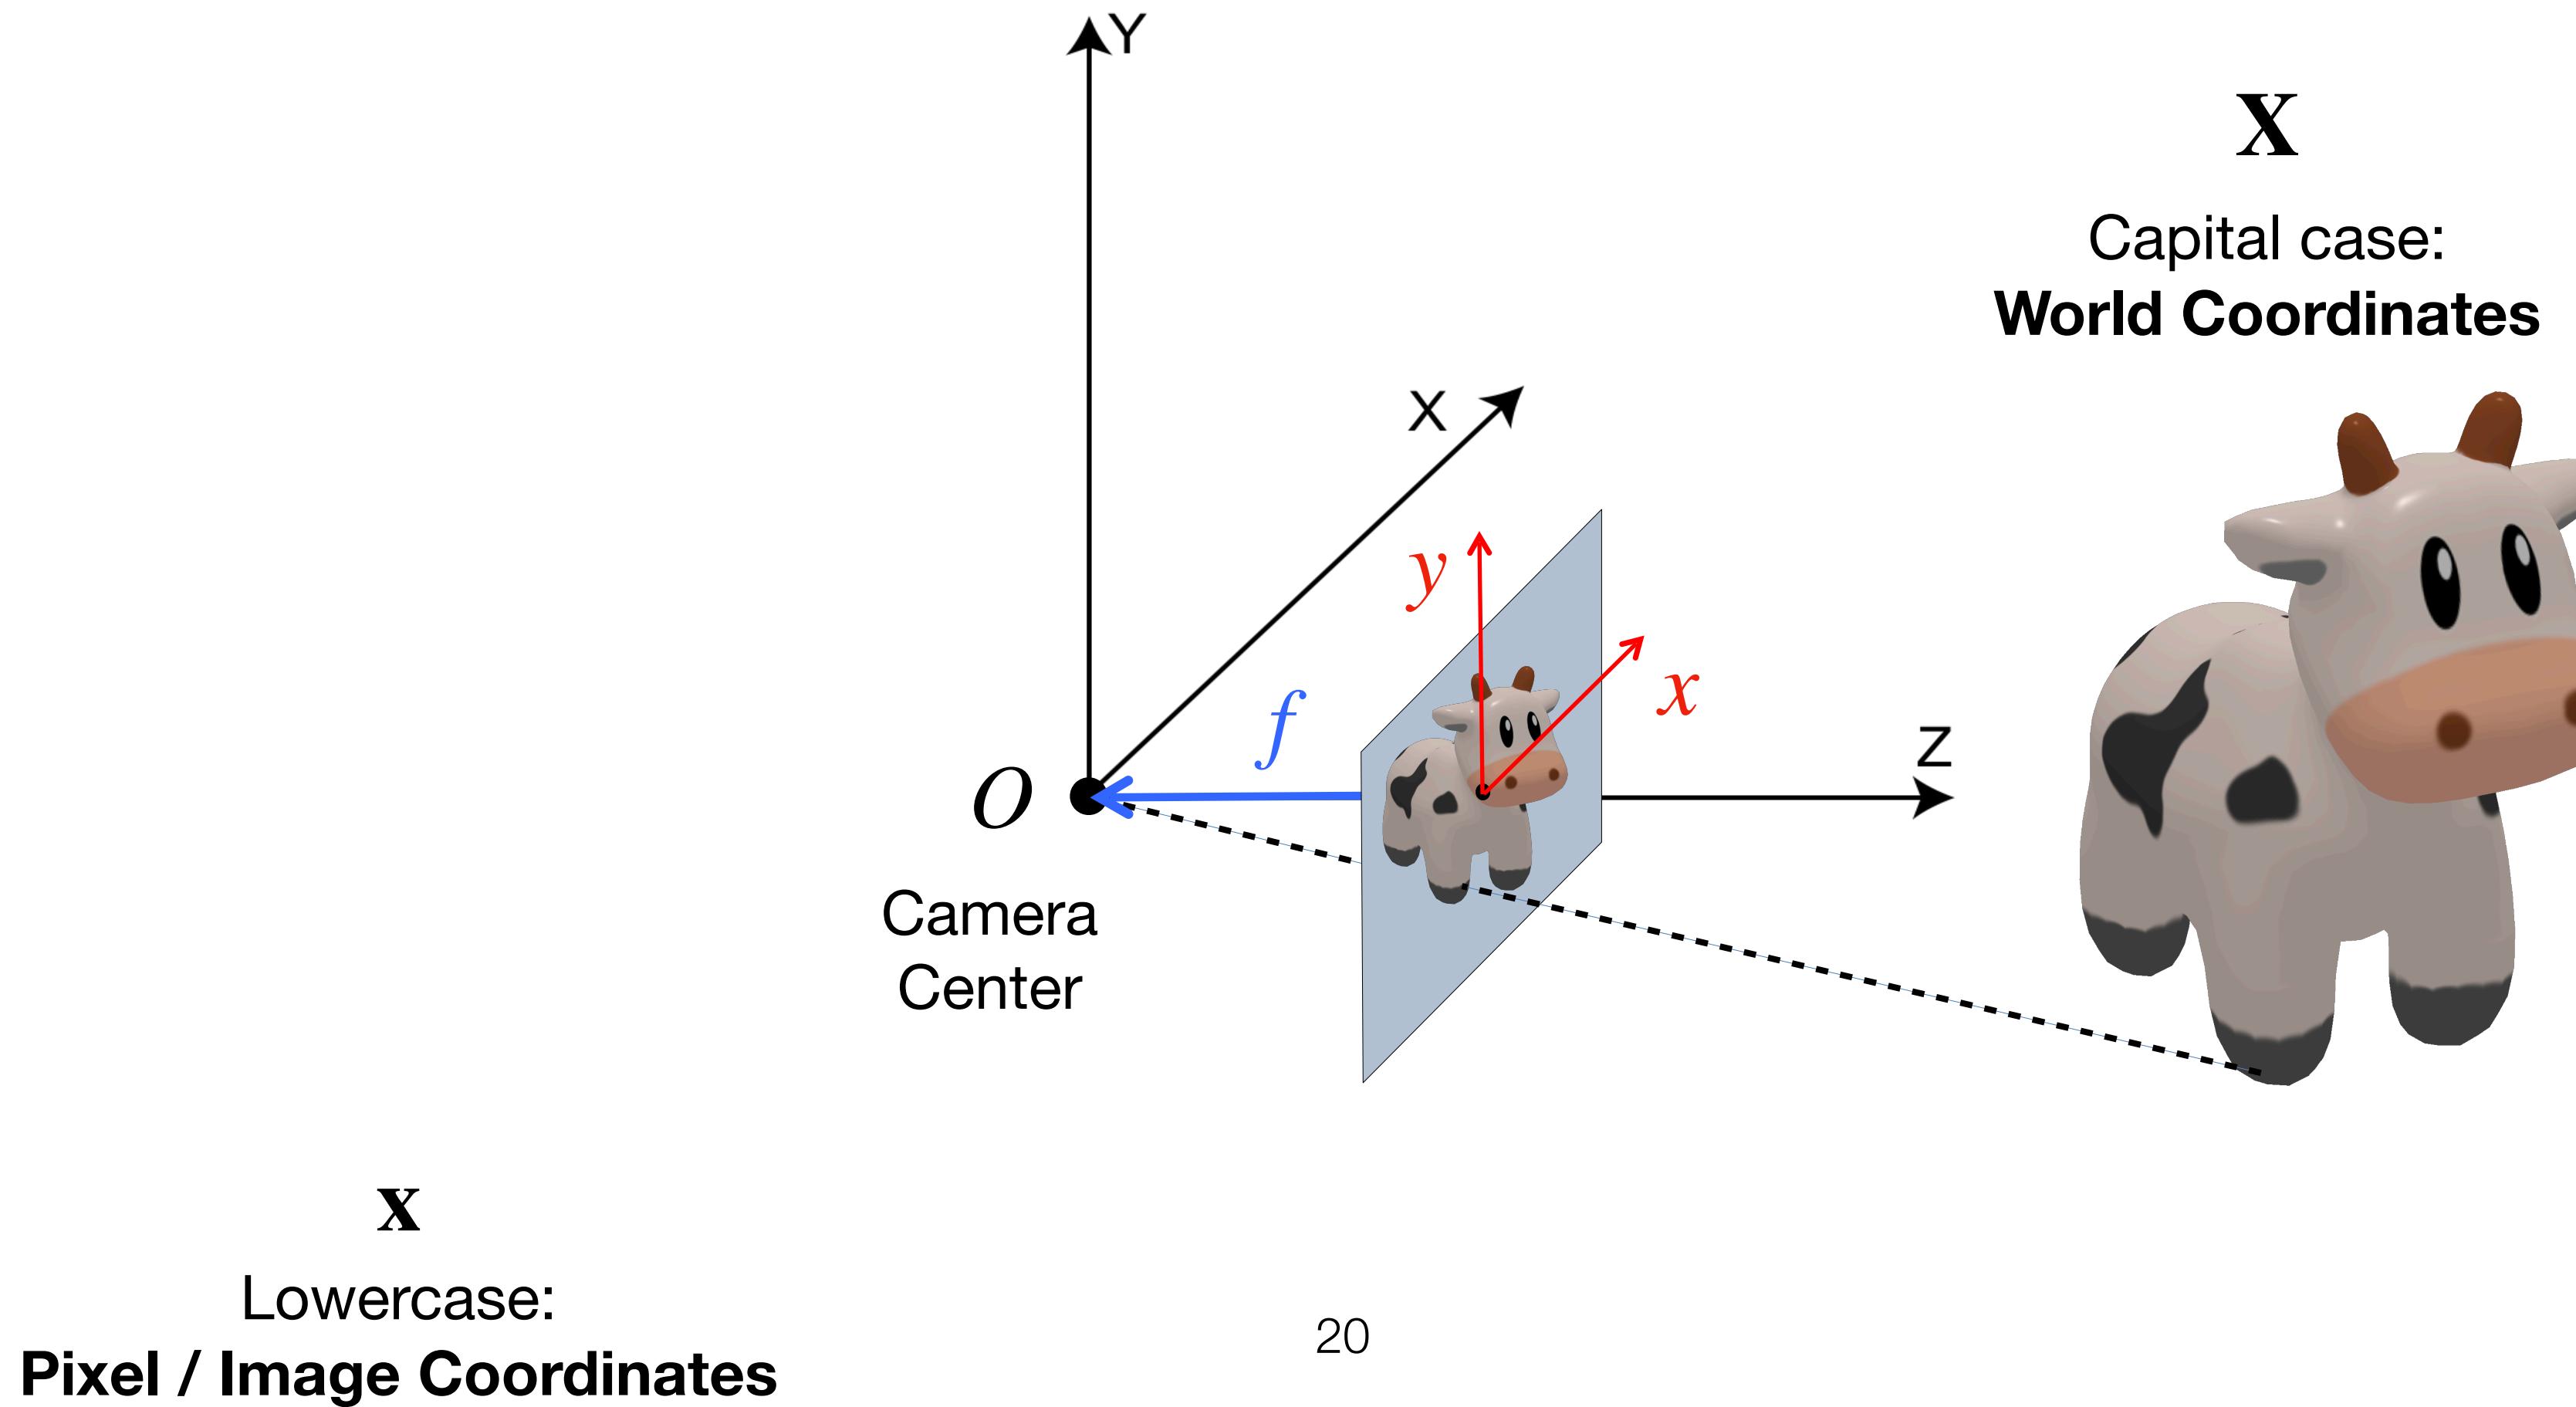
Perspective projection



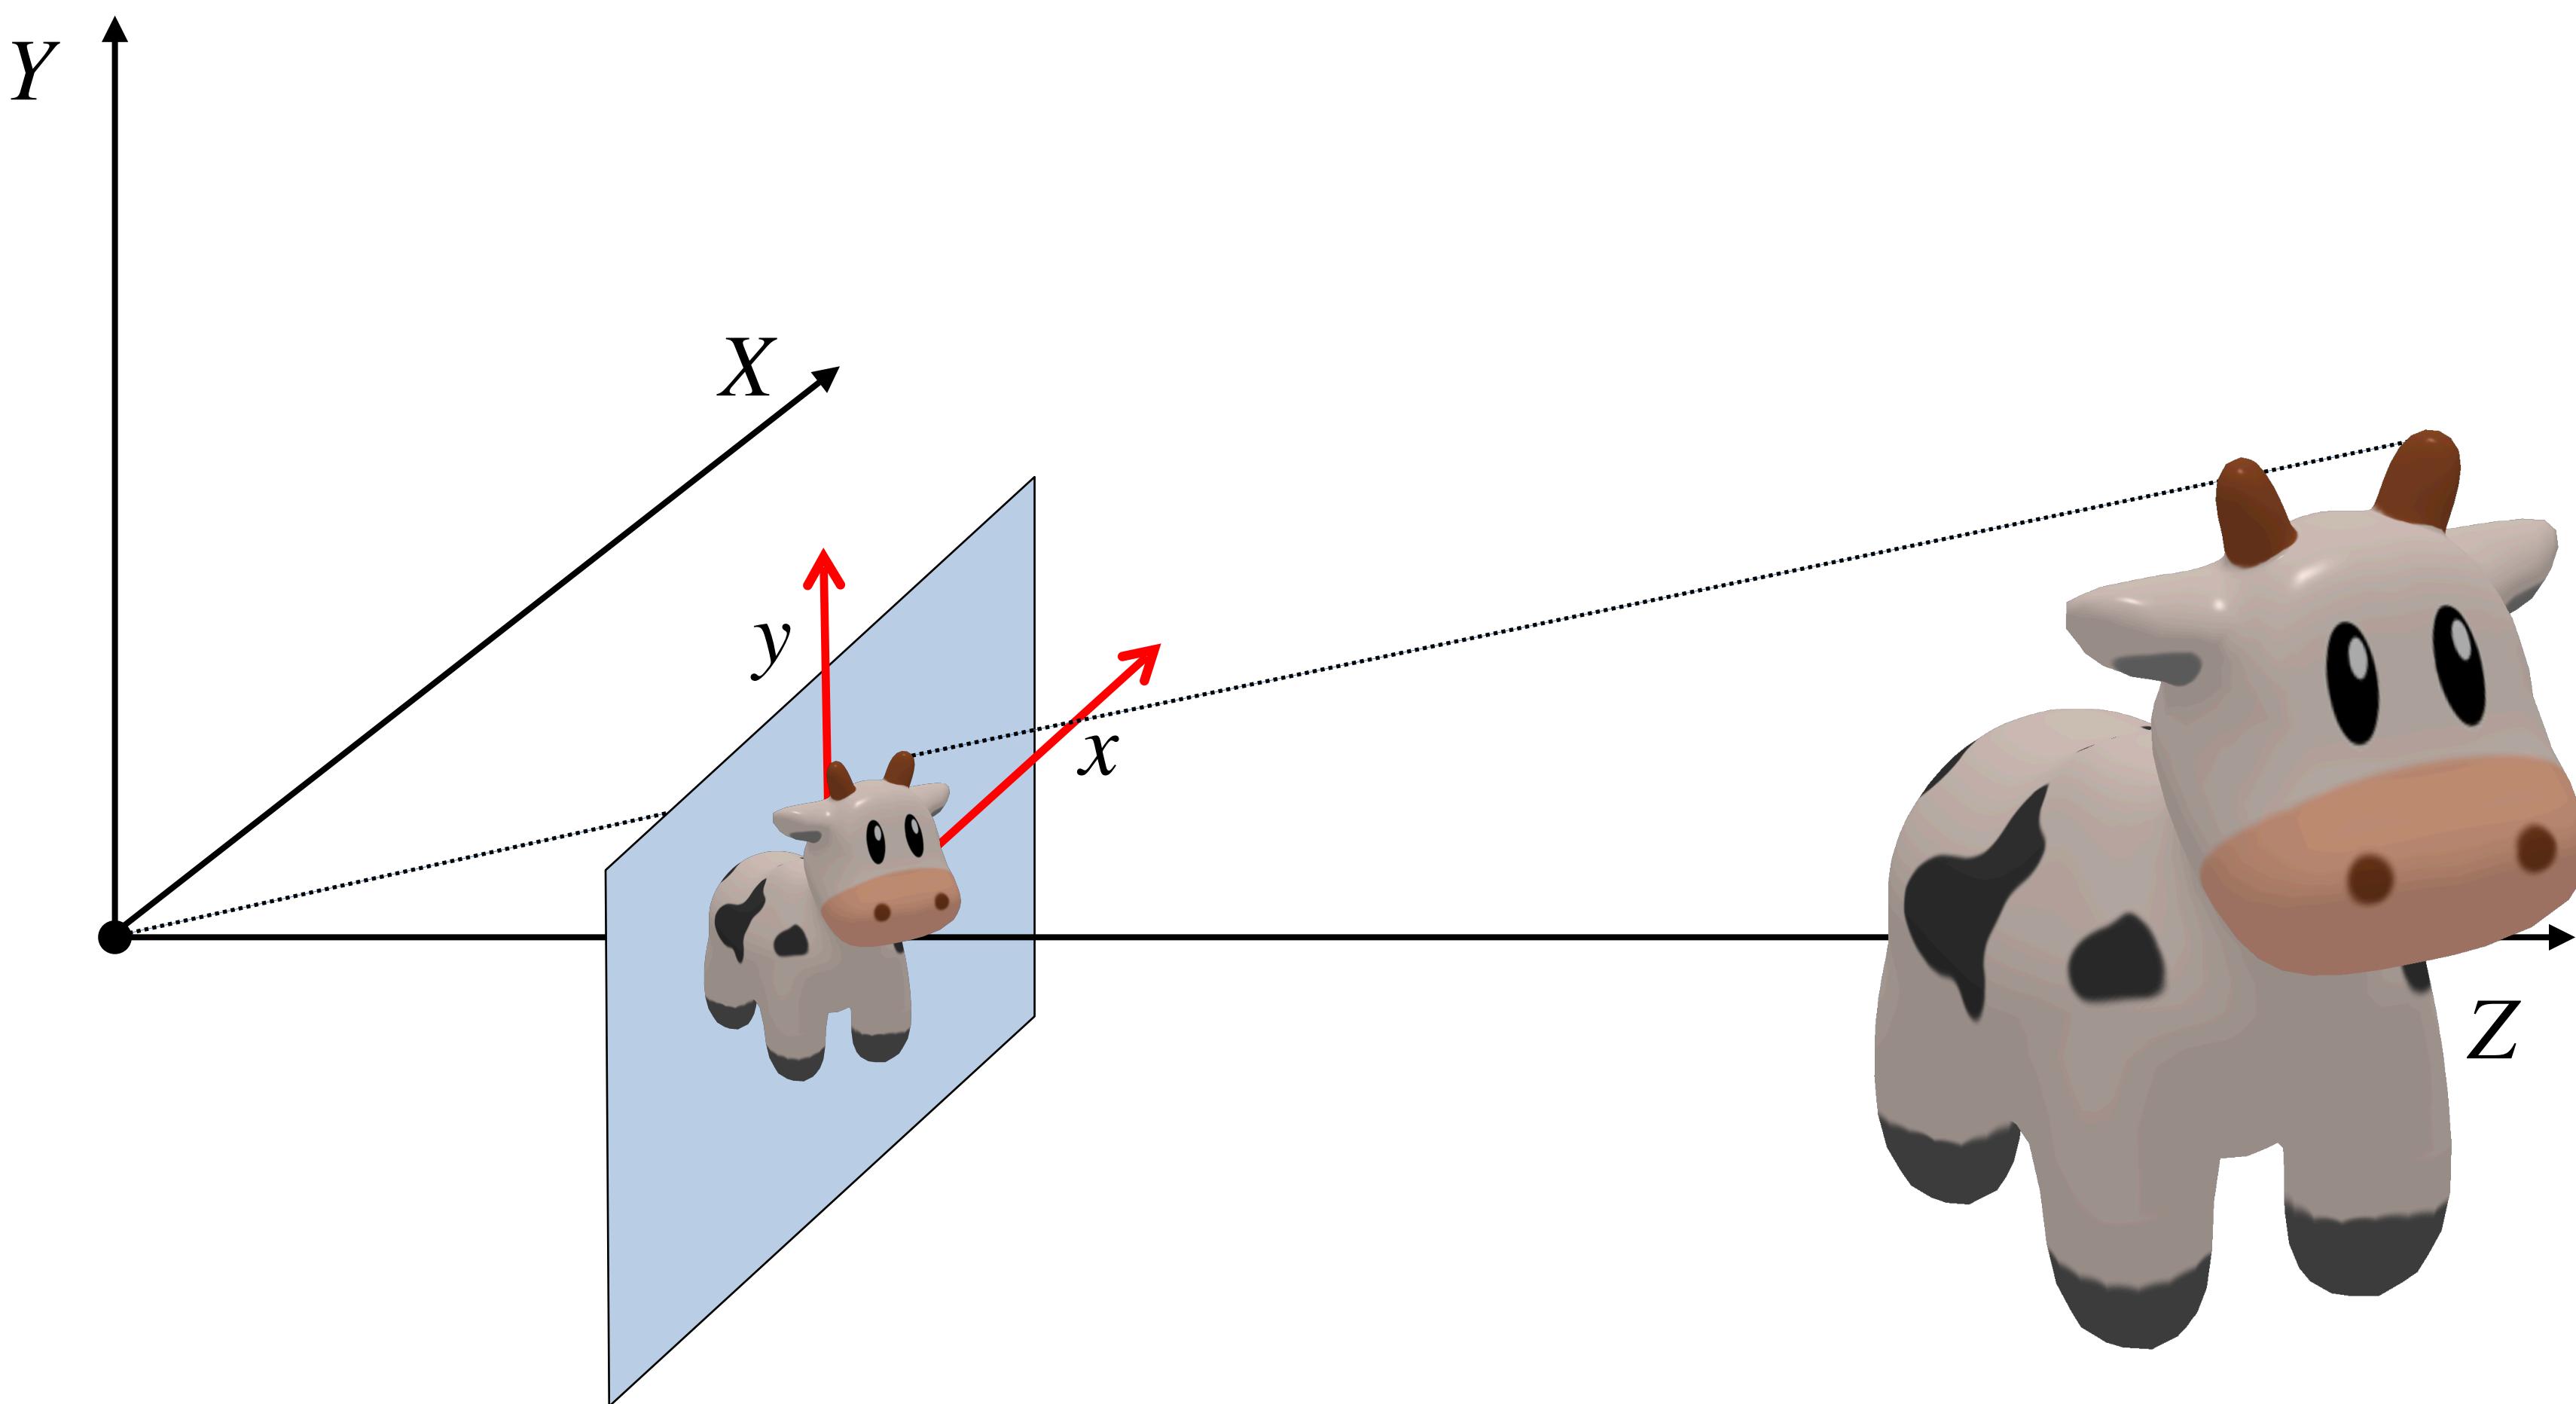
Perspective projection



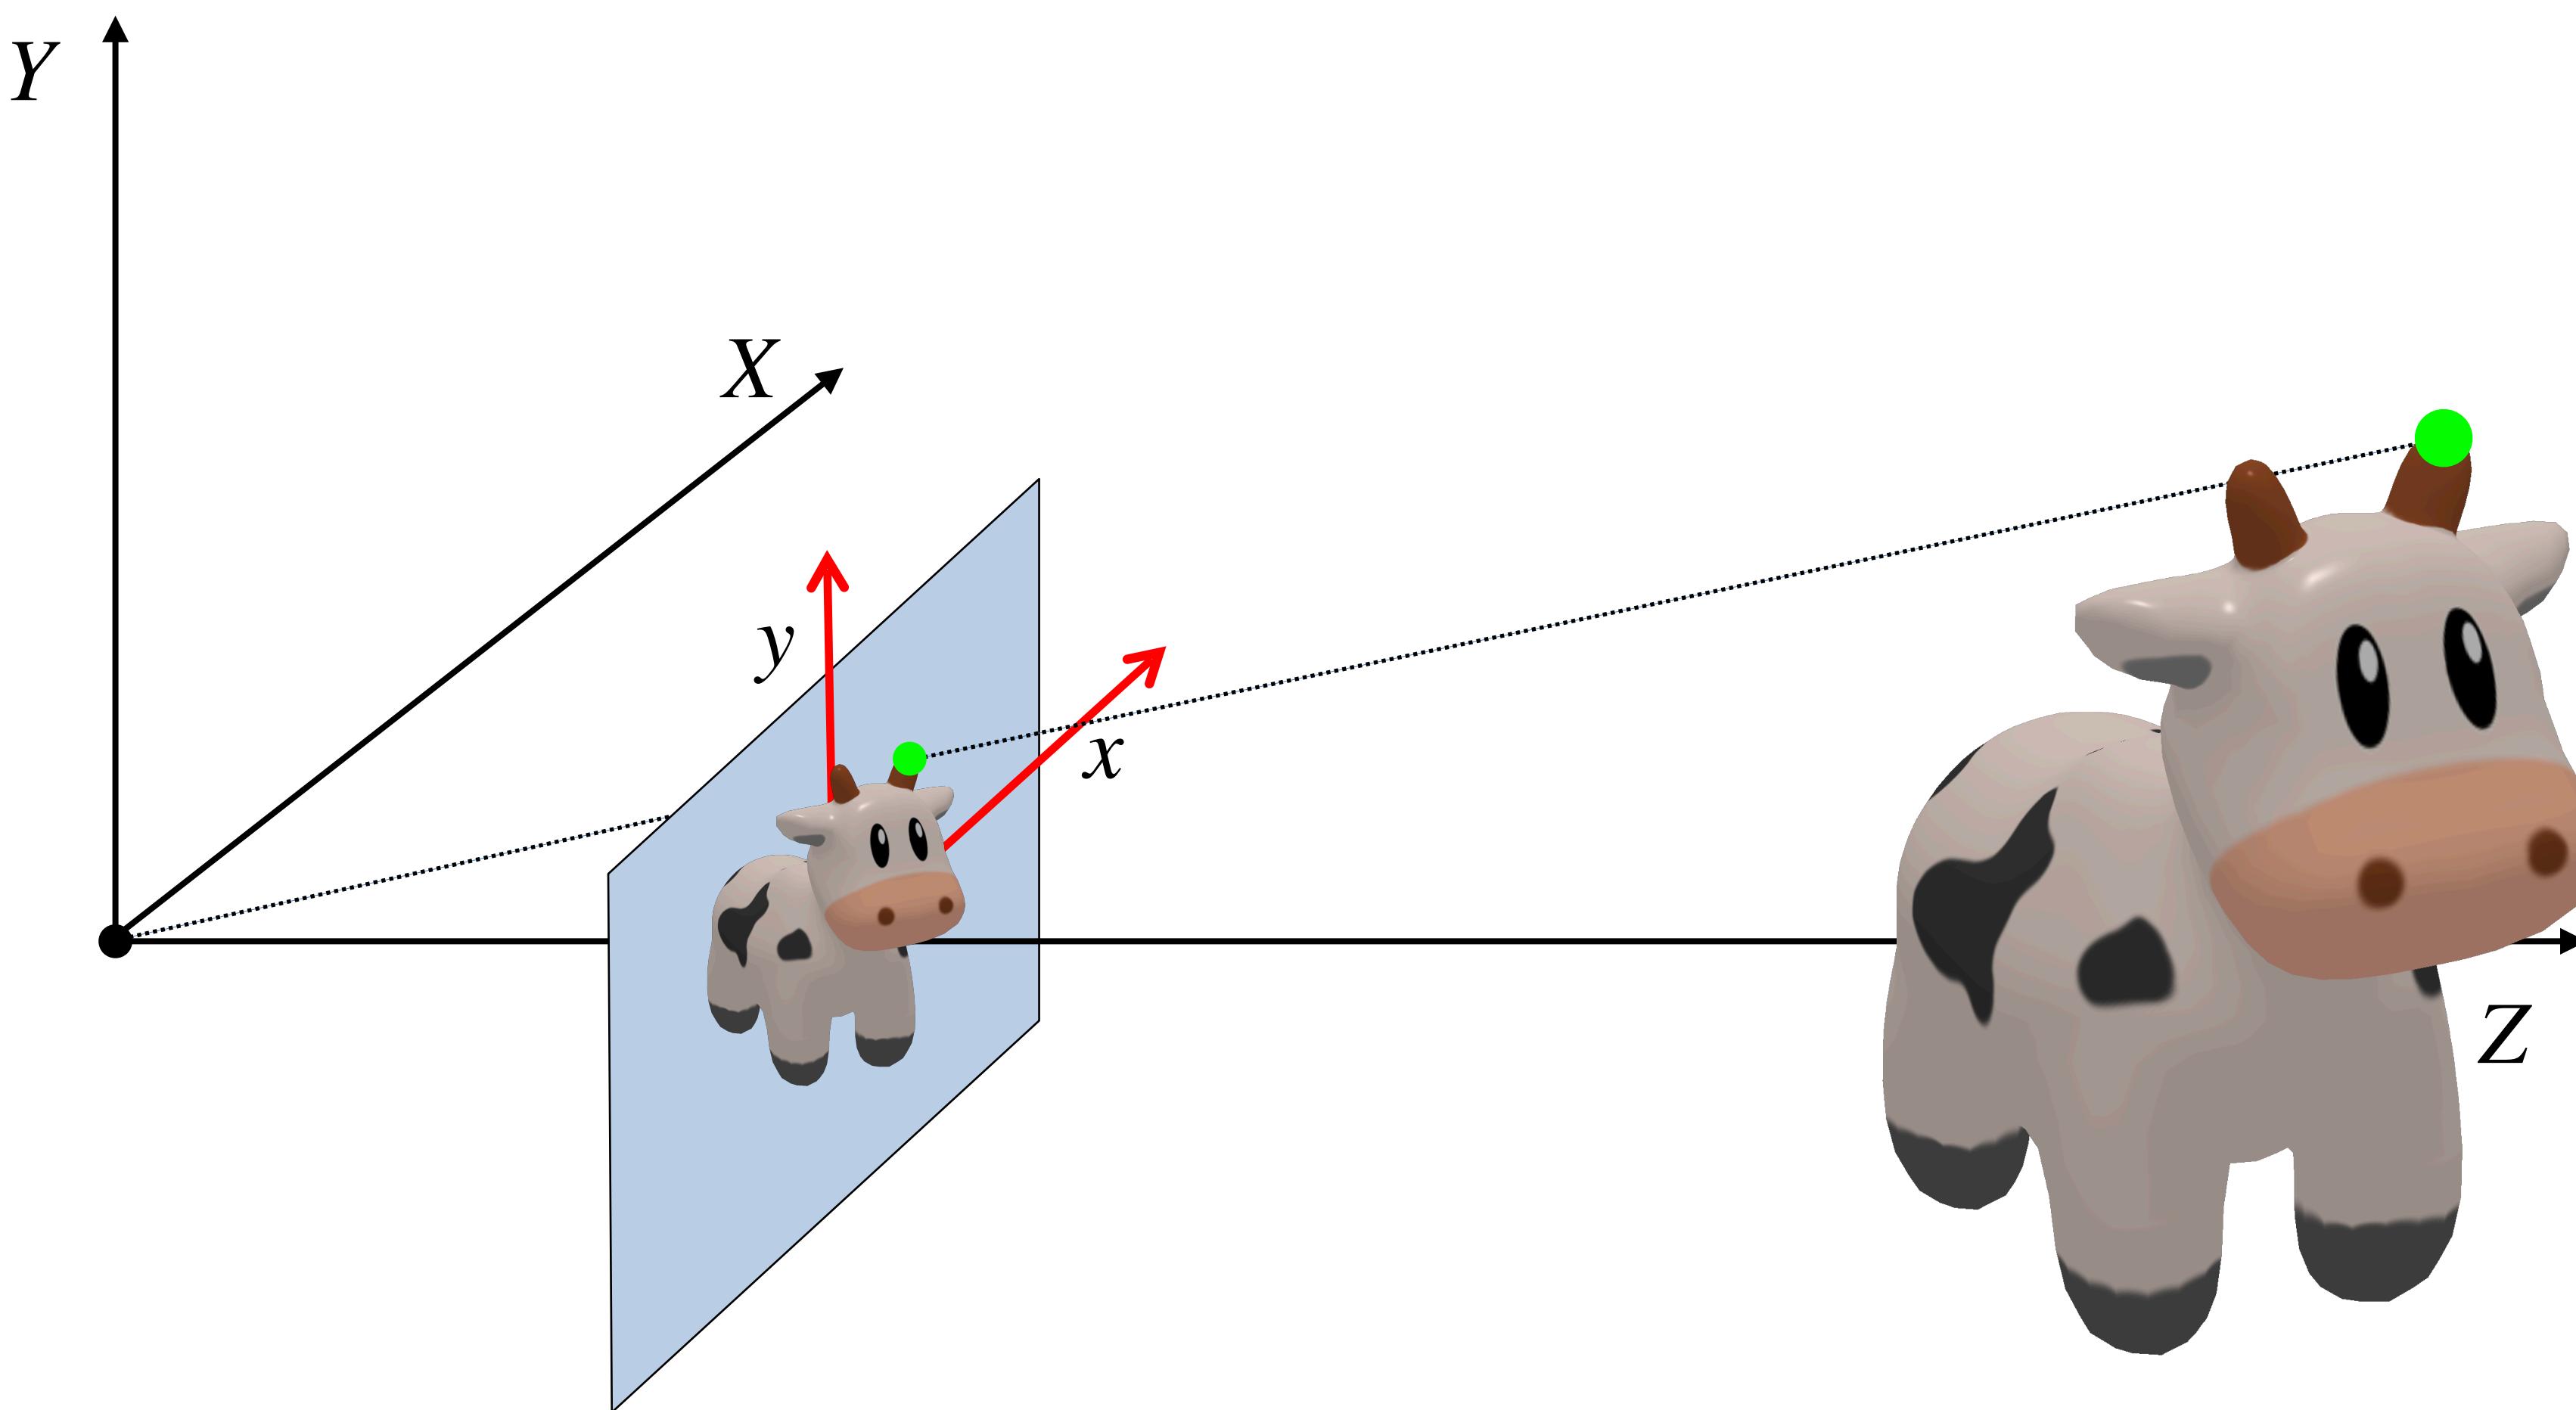
Perspective projection



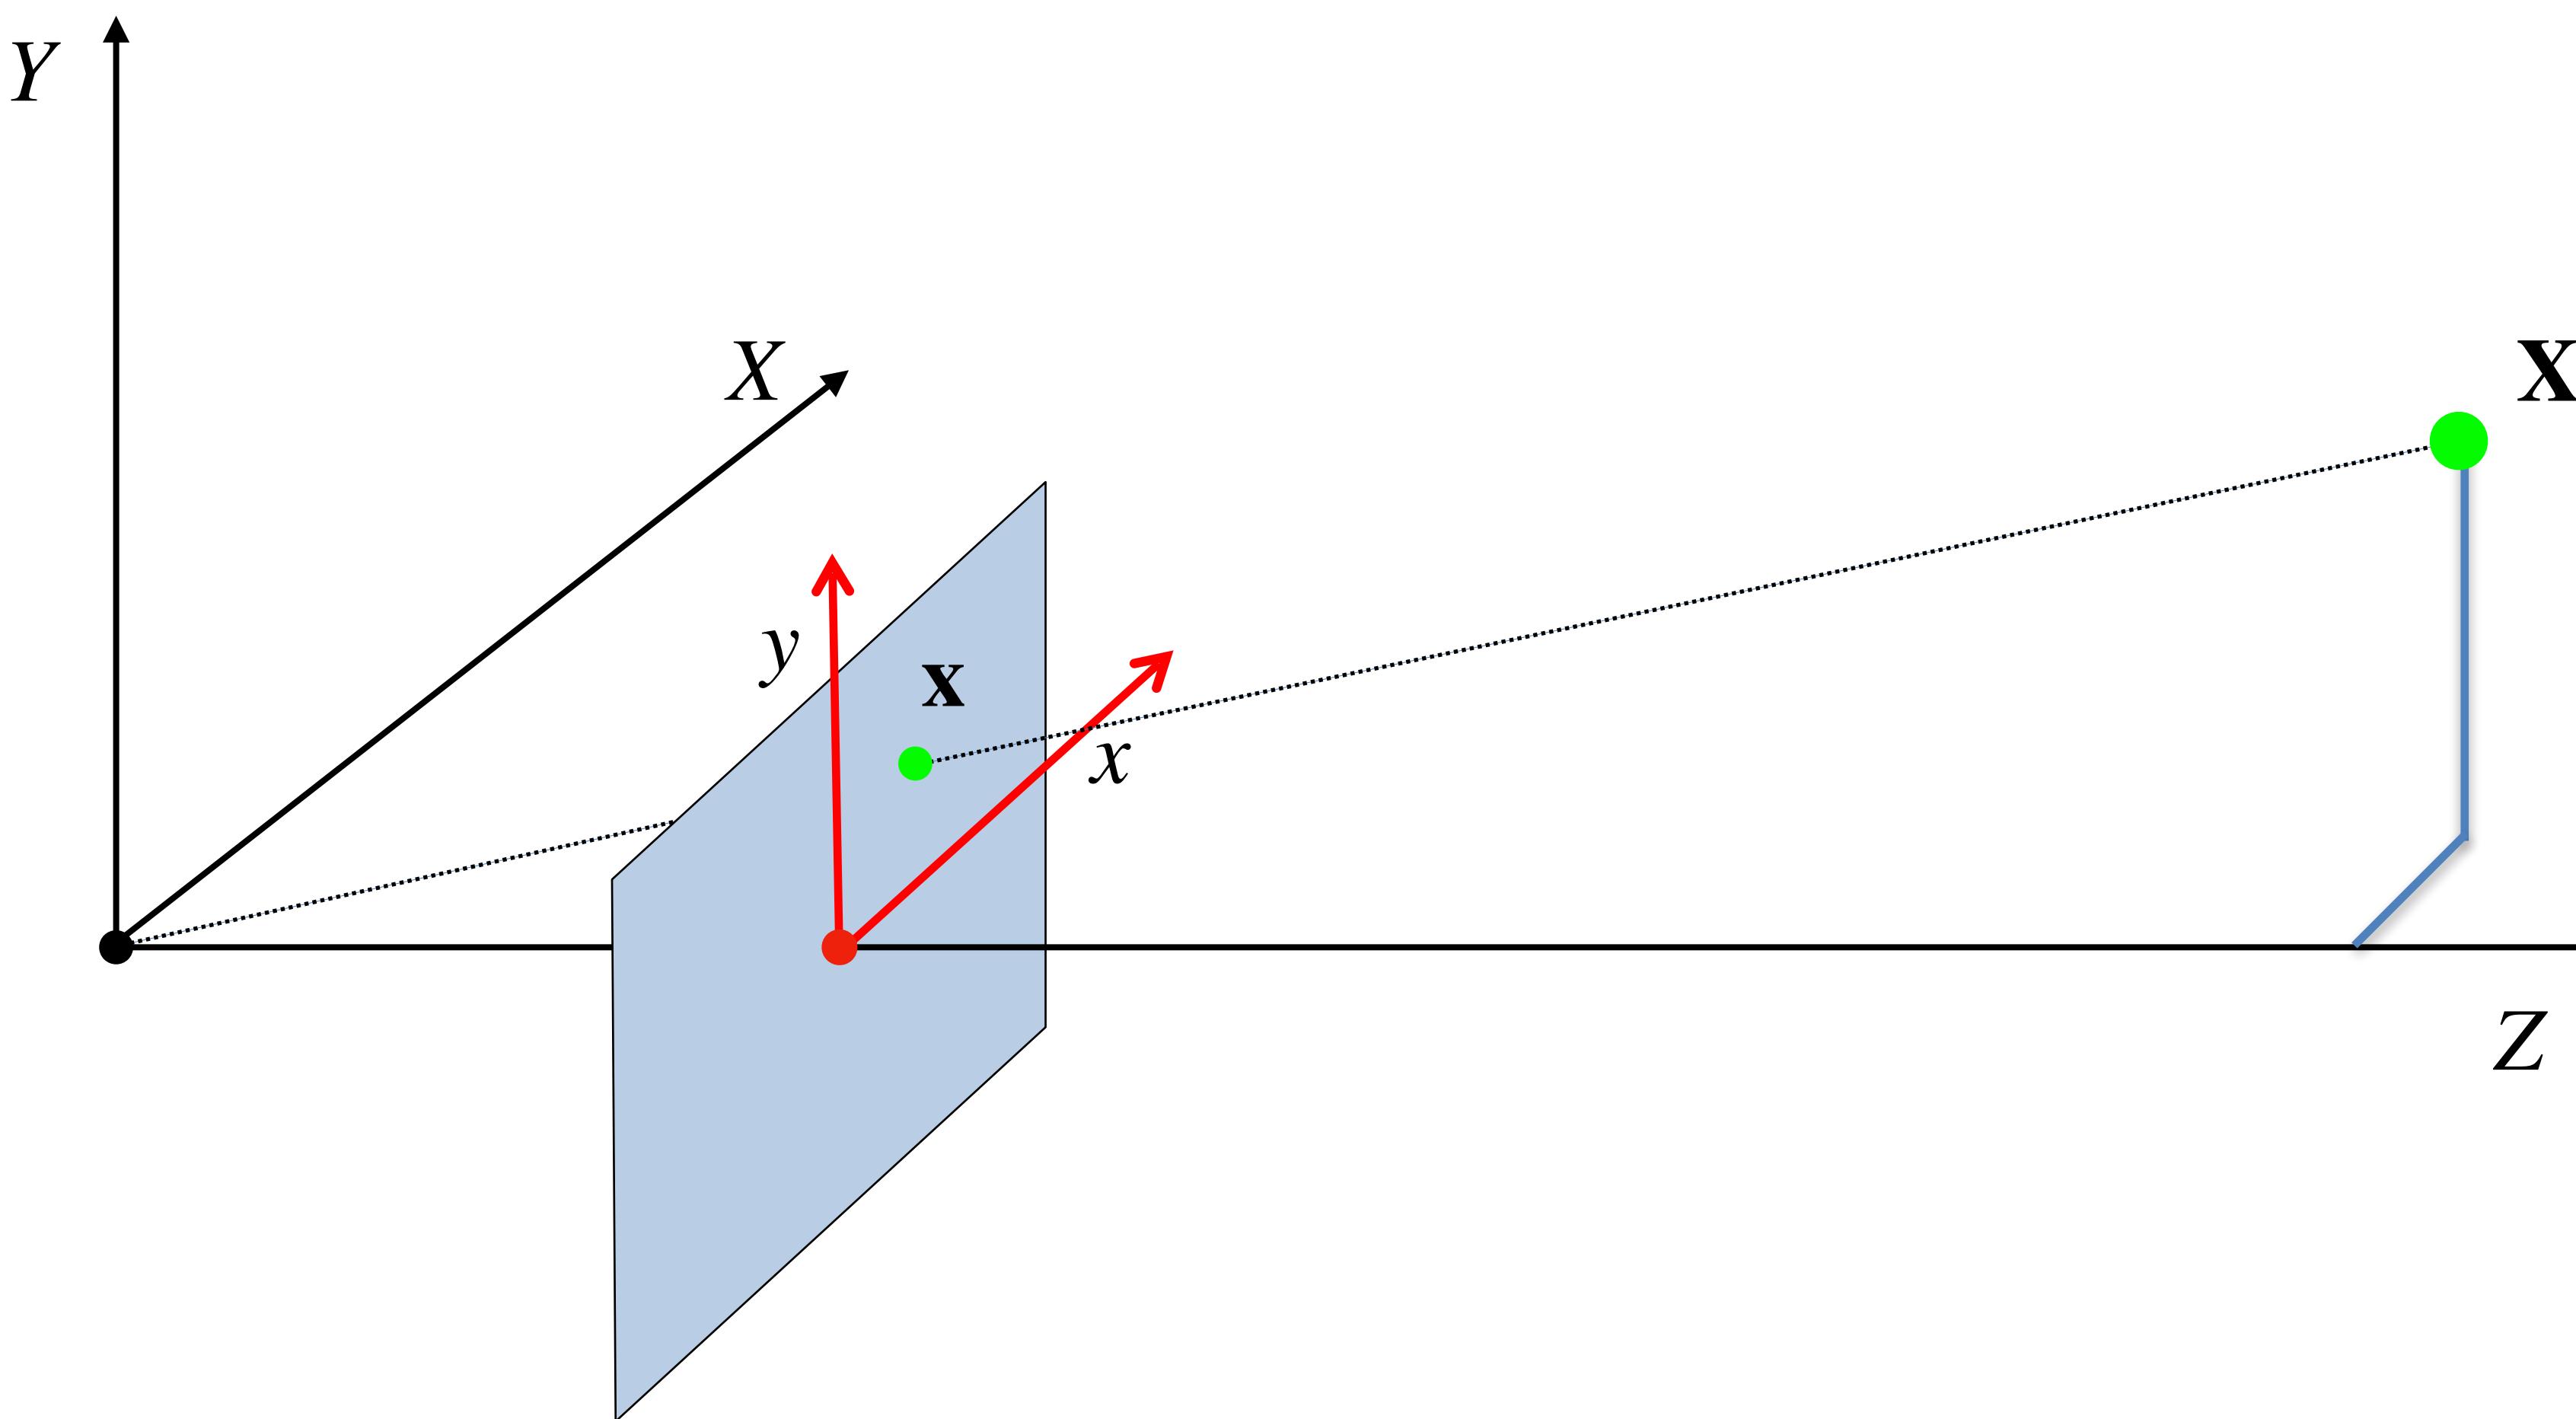
Perspective projection



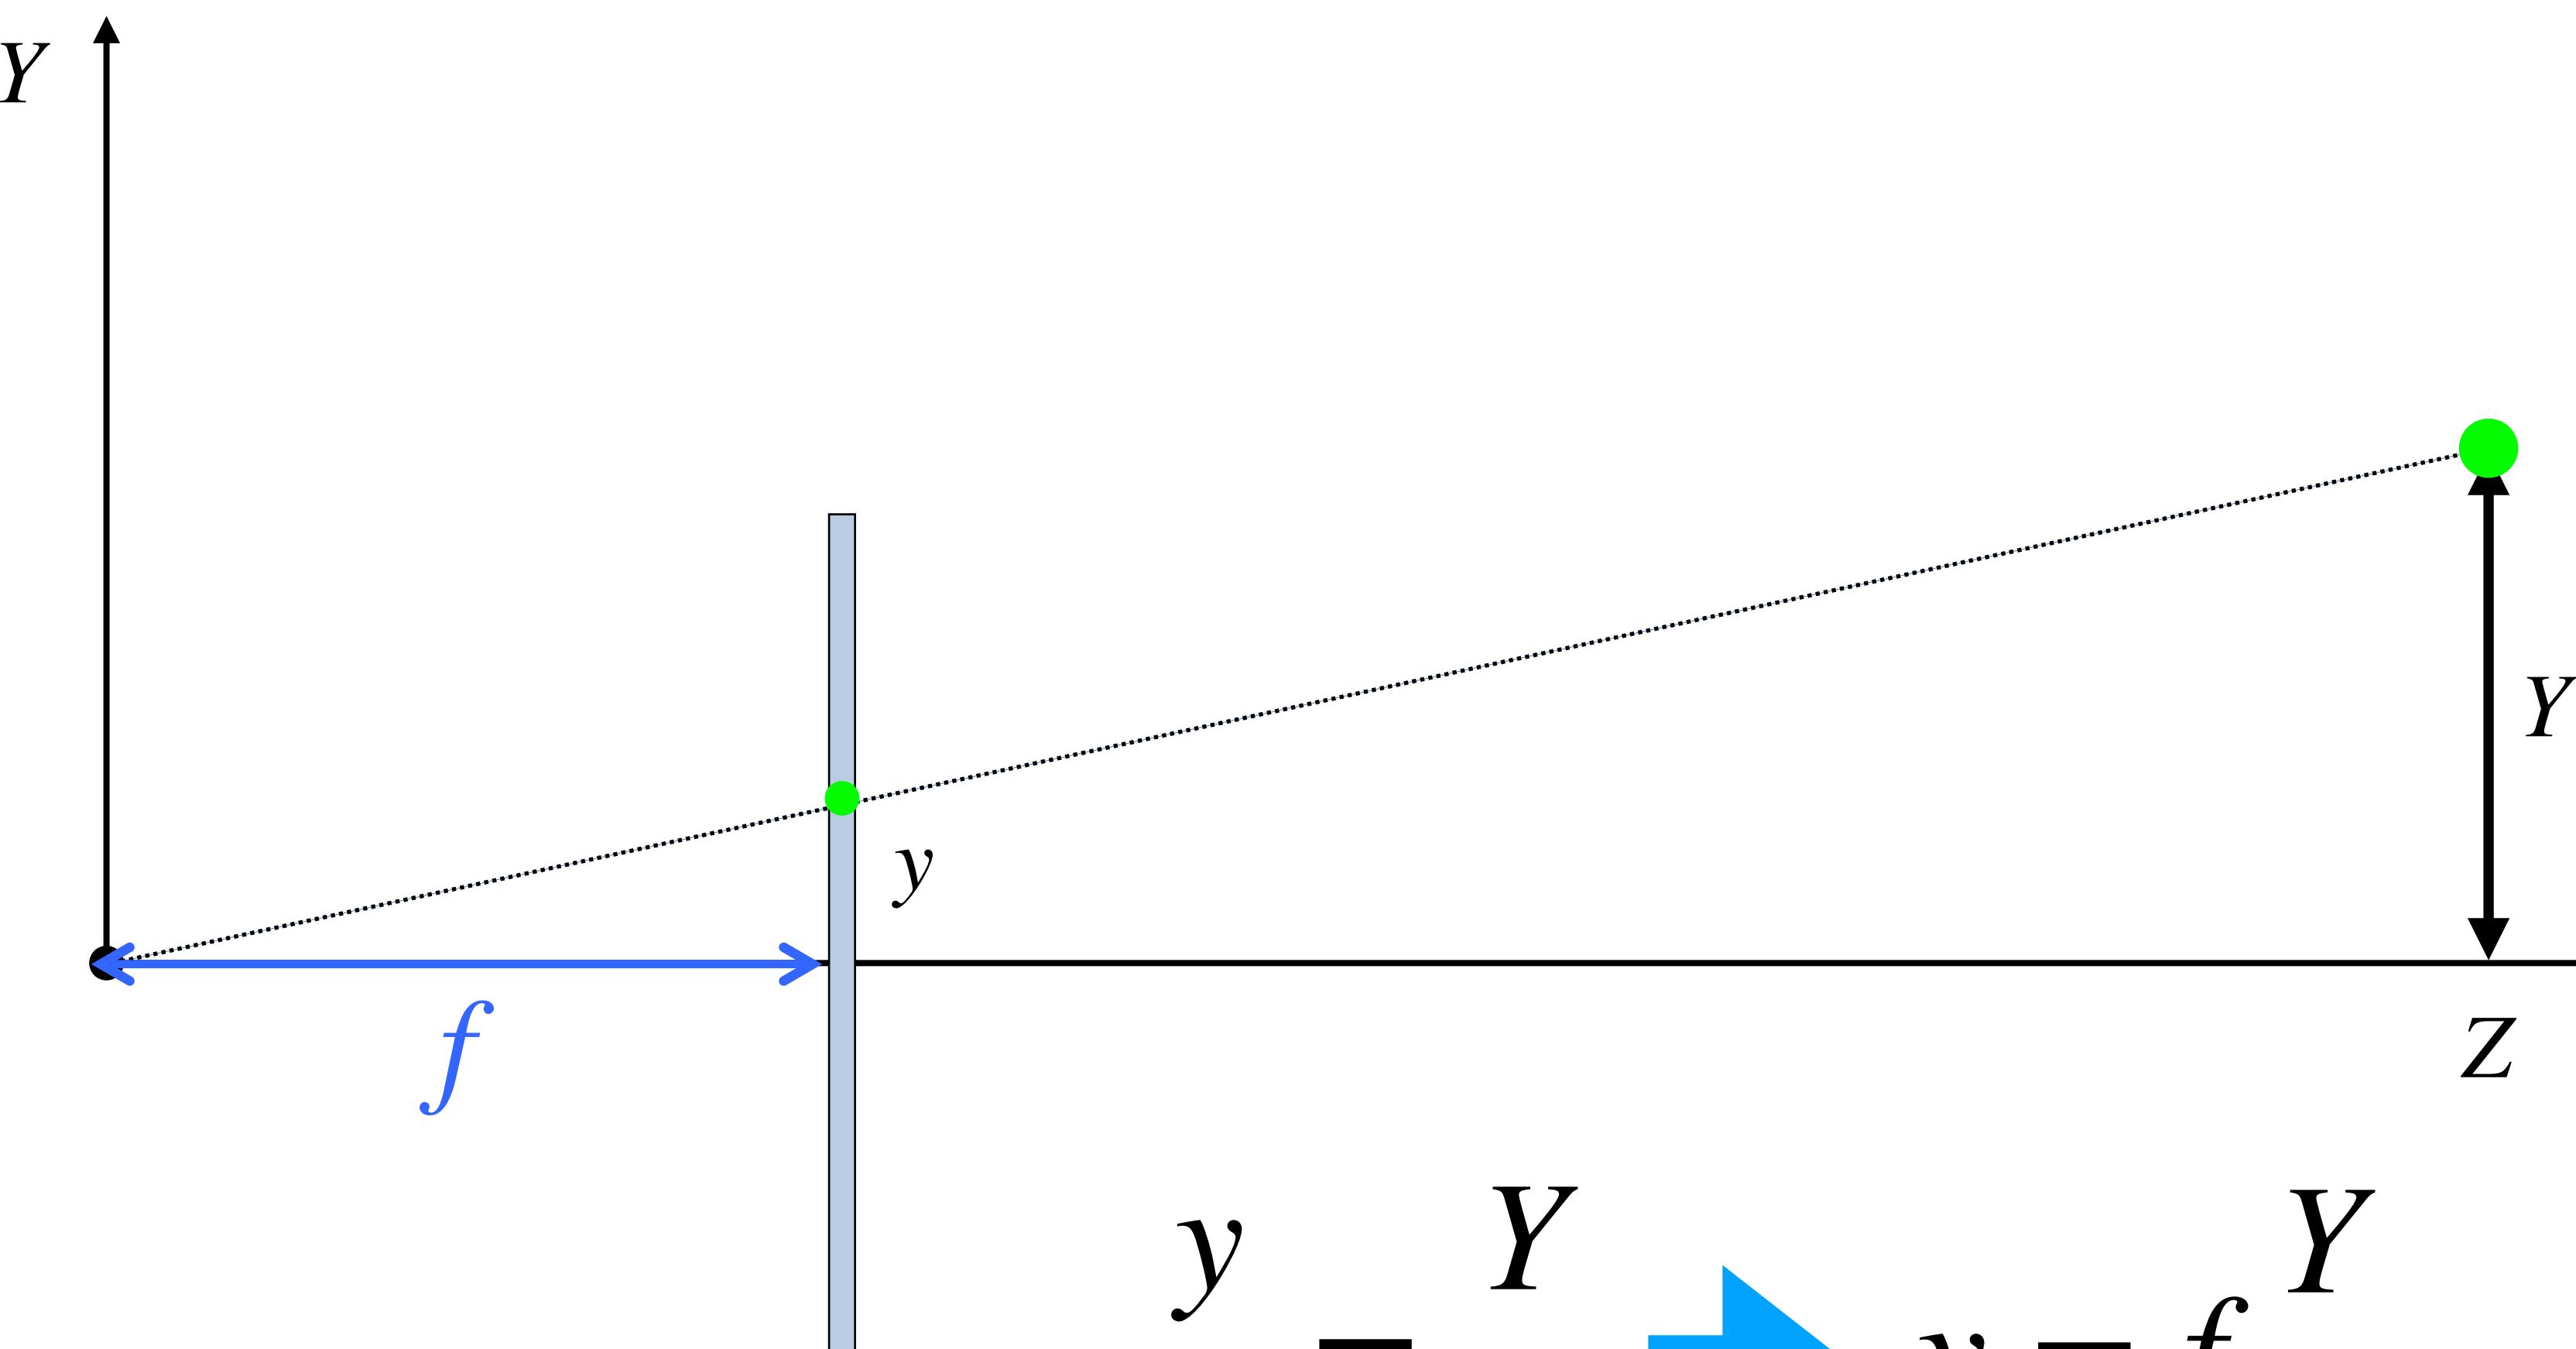
Perspective projection



Perspective projection

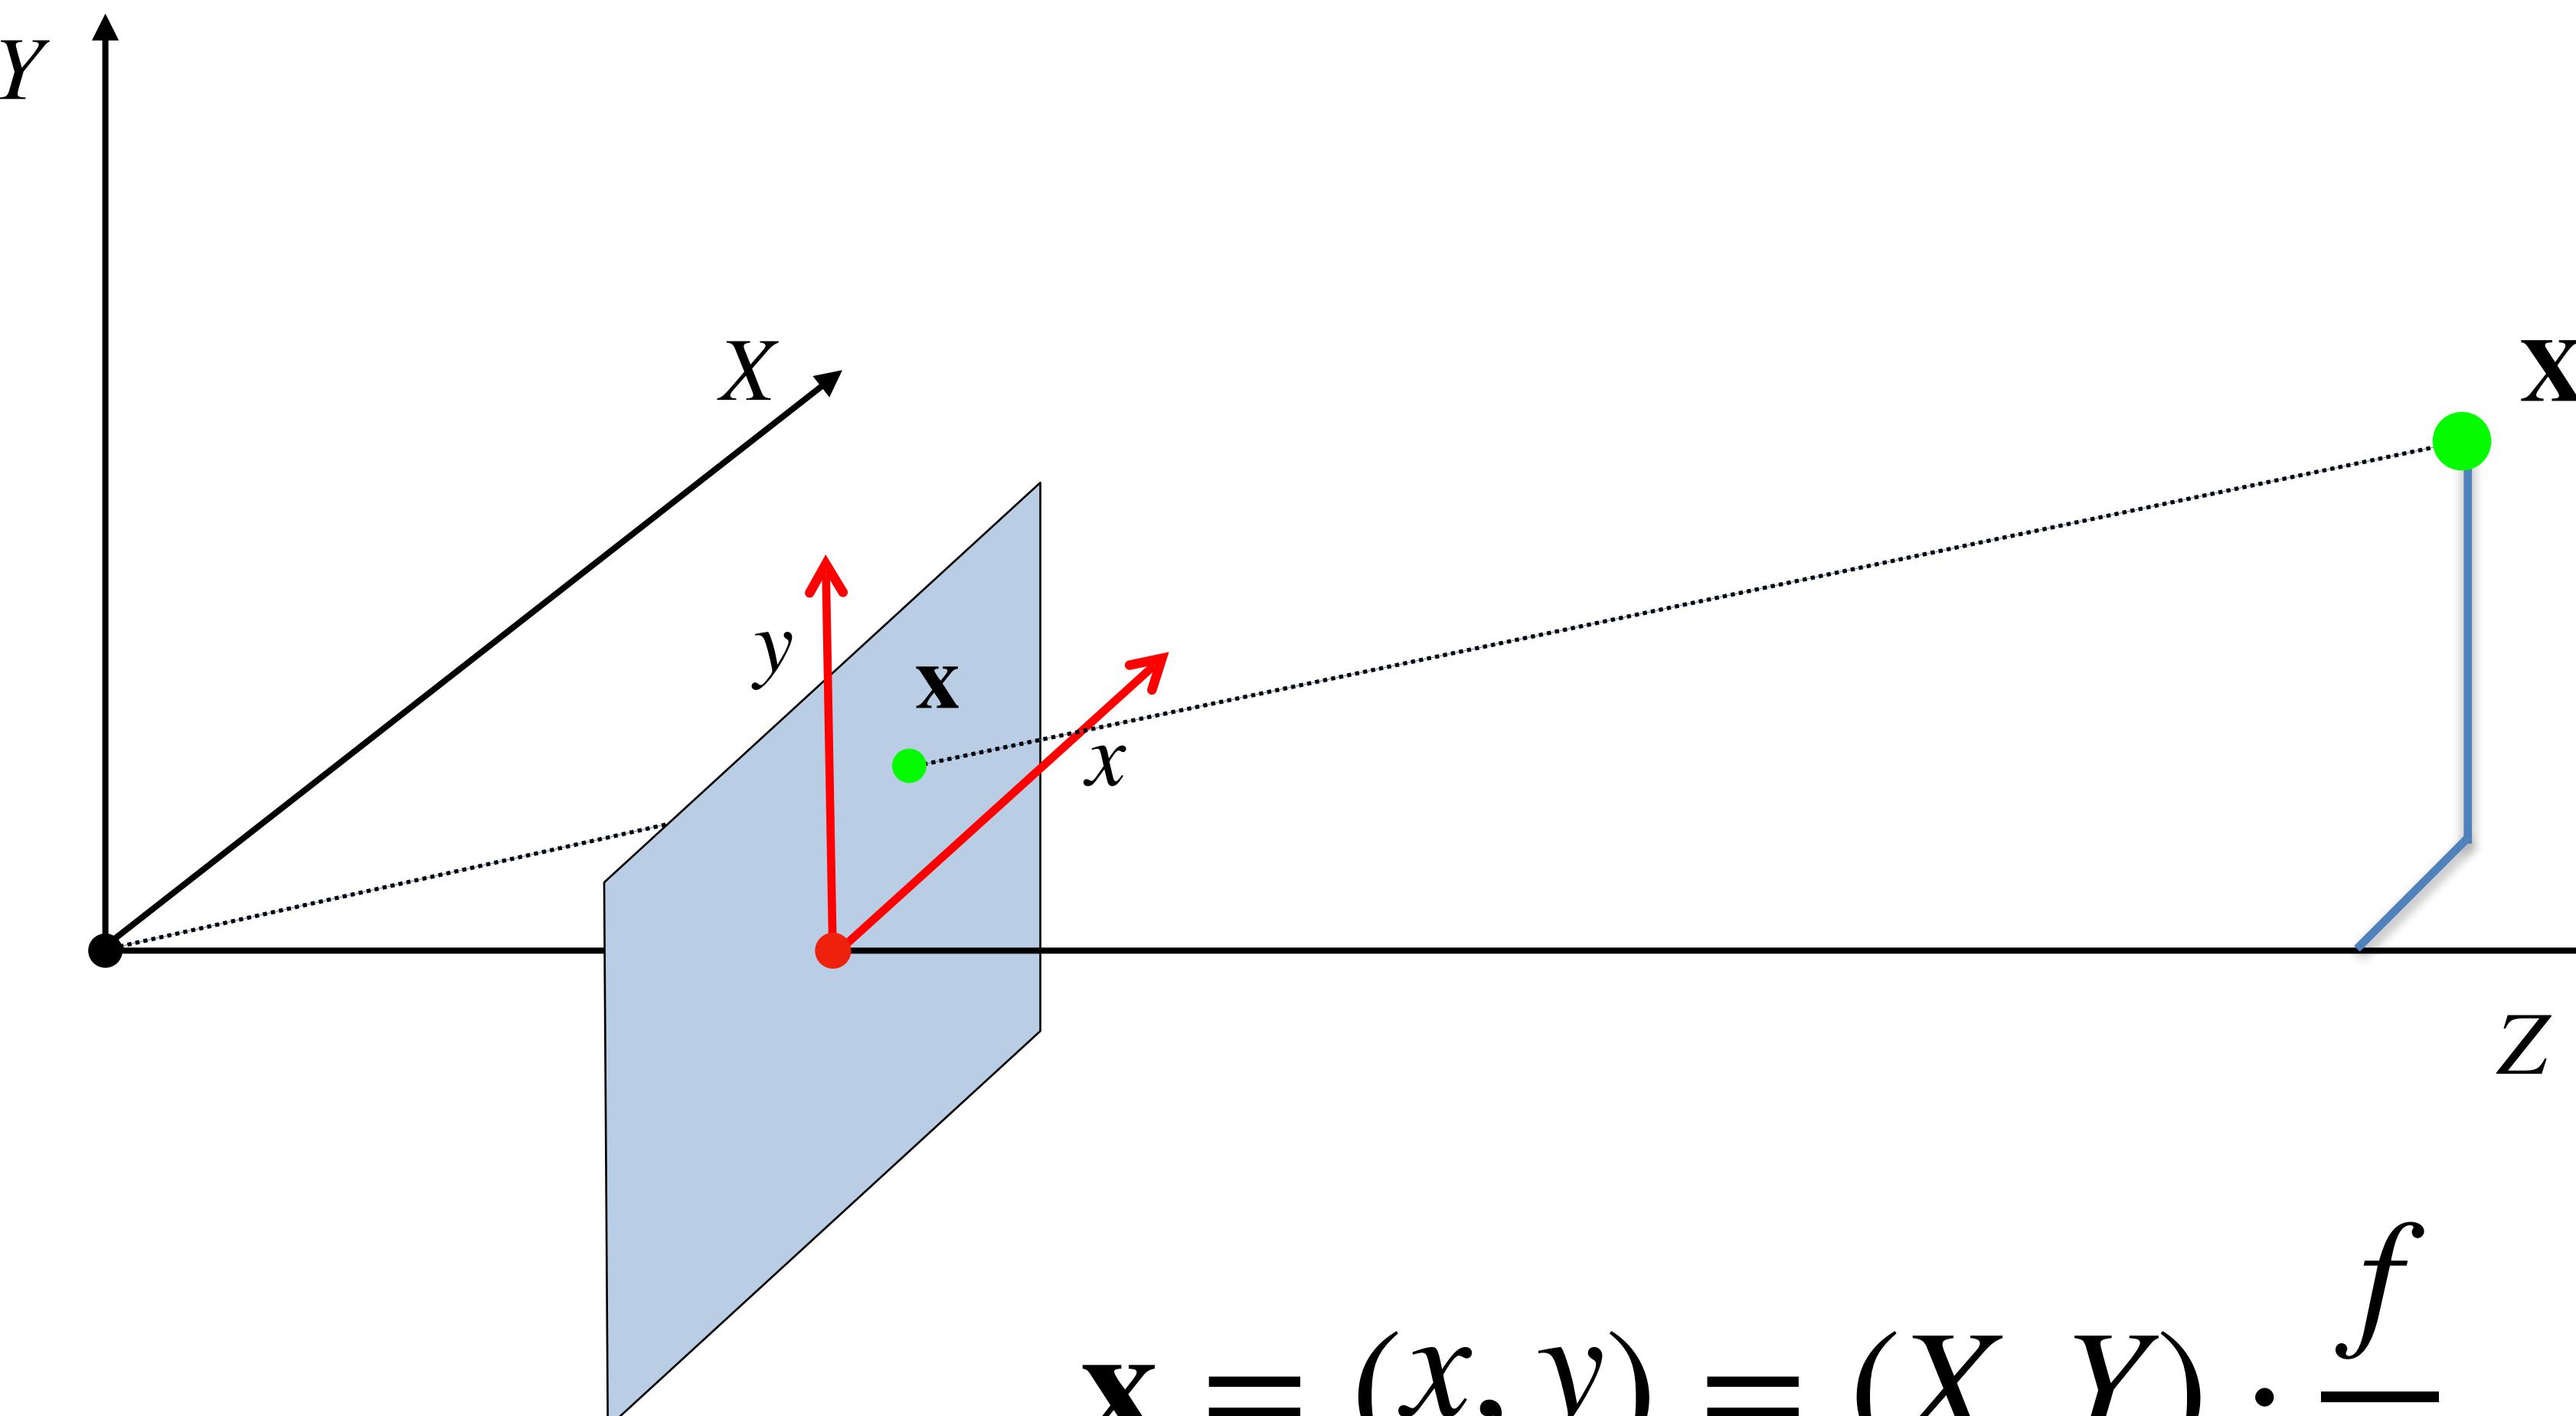


Perspective projection



$$\frac{y}{f} = \frac{Y}{Z} \rightarrow y = f \frac{Y}{Z}$$

Perspective projection

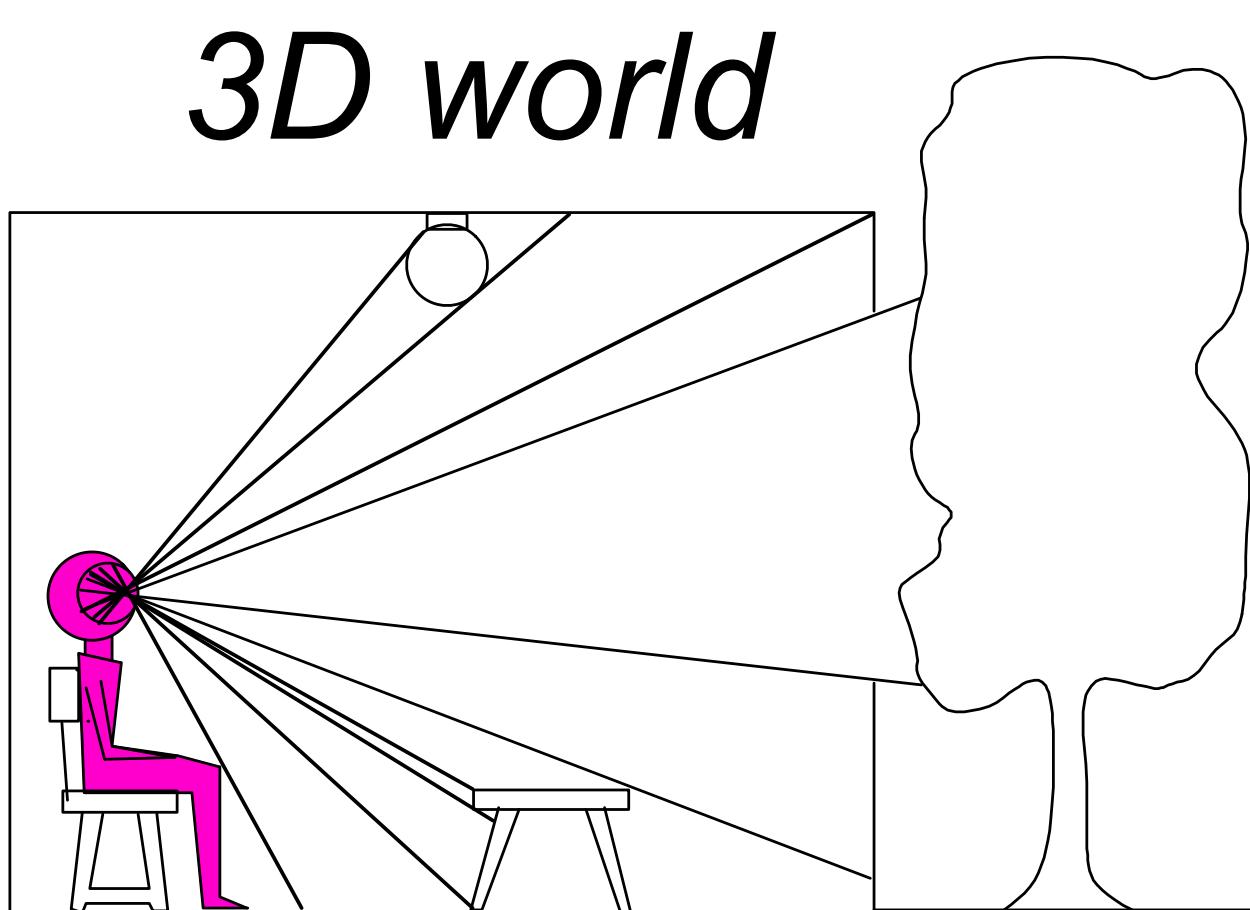


$$\mathbf{x} = (x, y) = (X, Y) \cdot \frac{f}{Z}$$

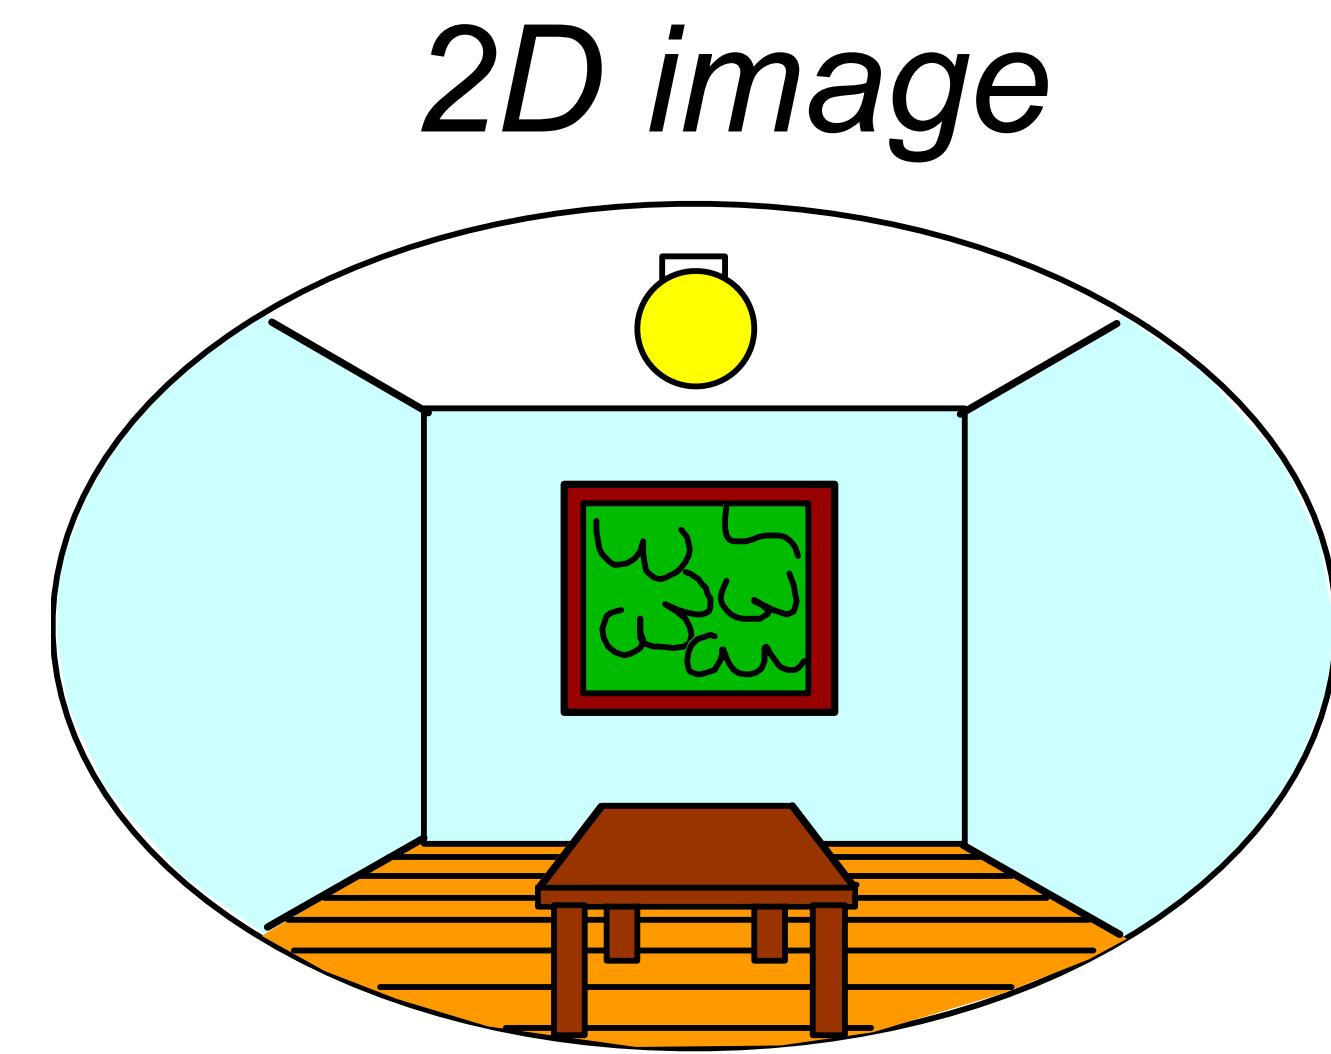
Image / Pixel
Coordinates

World Coordinates

Perspective Projection



Point of observation



What properties of the world are preserved?

- Straight lines, incidence

What properties are not preserved?

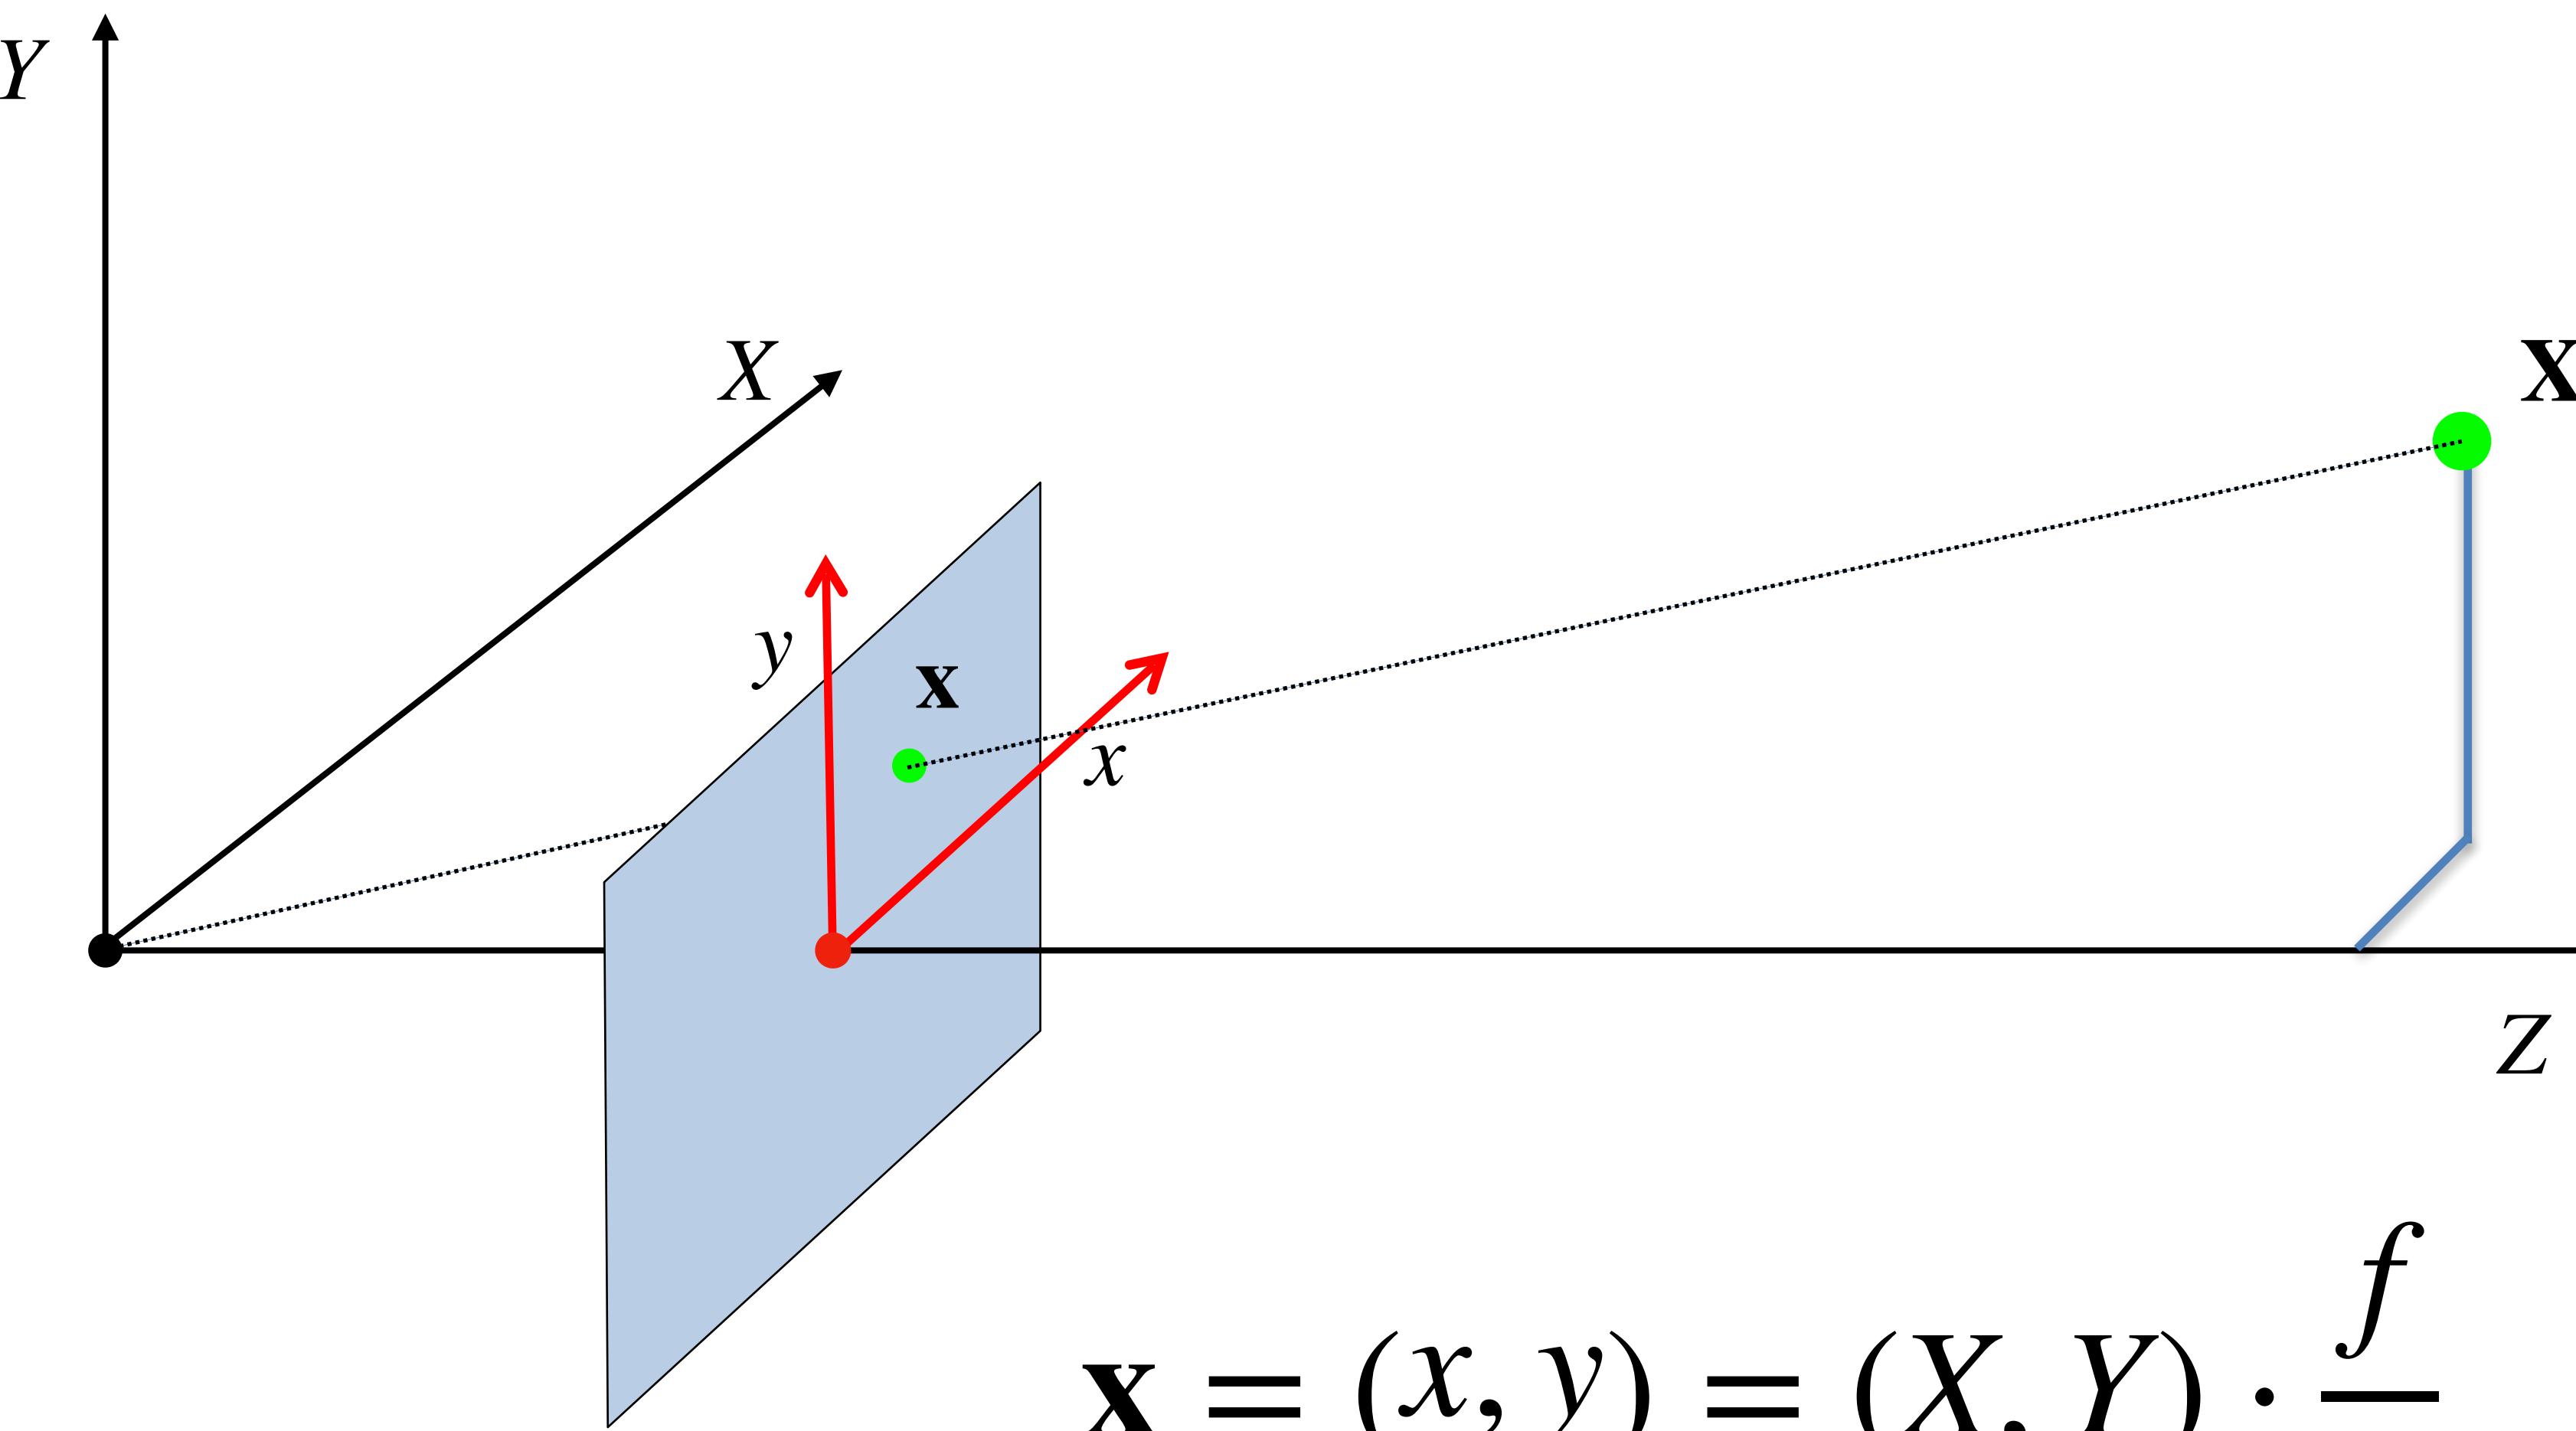
- Angles, lengths

Perspective Projection



Questions?

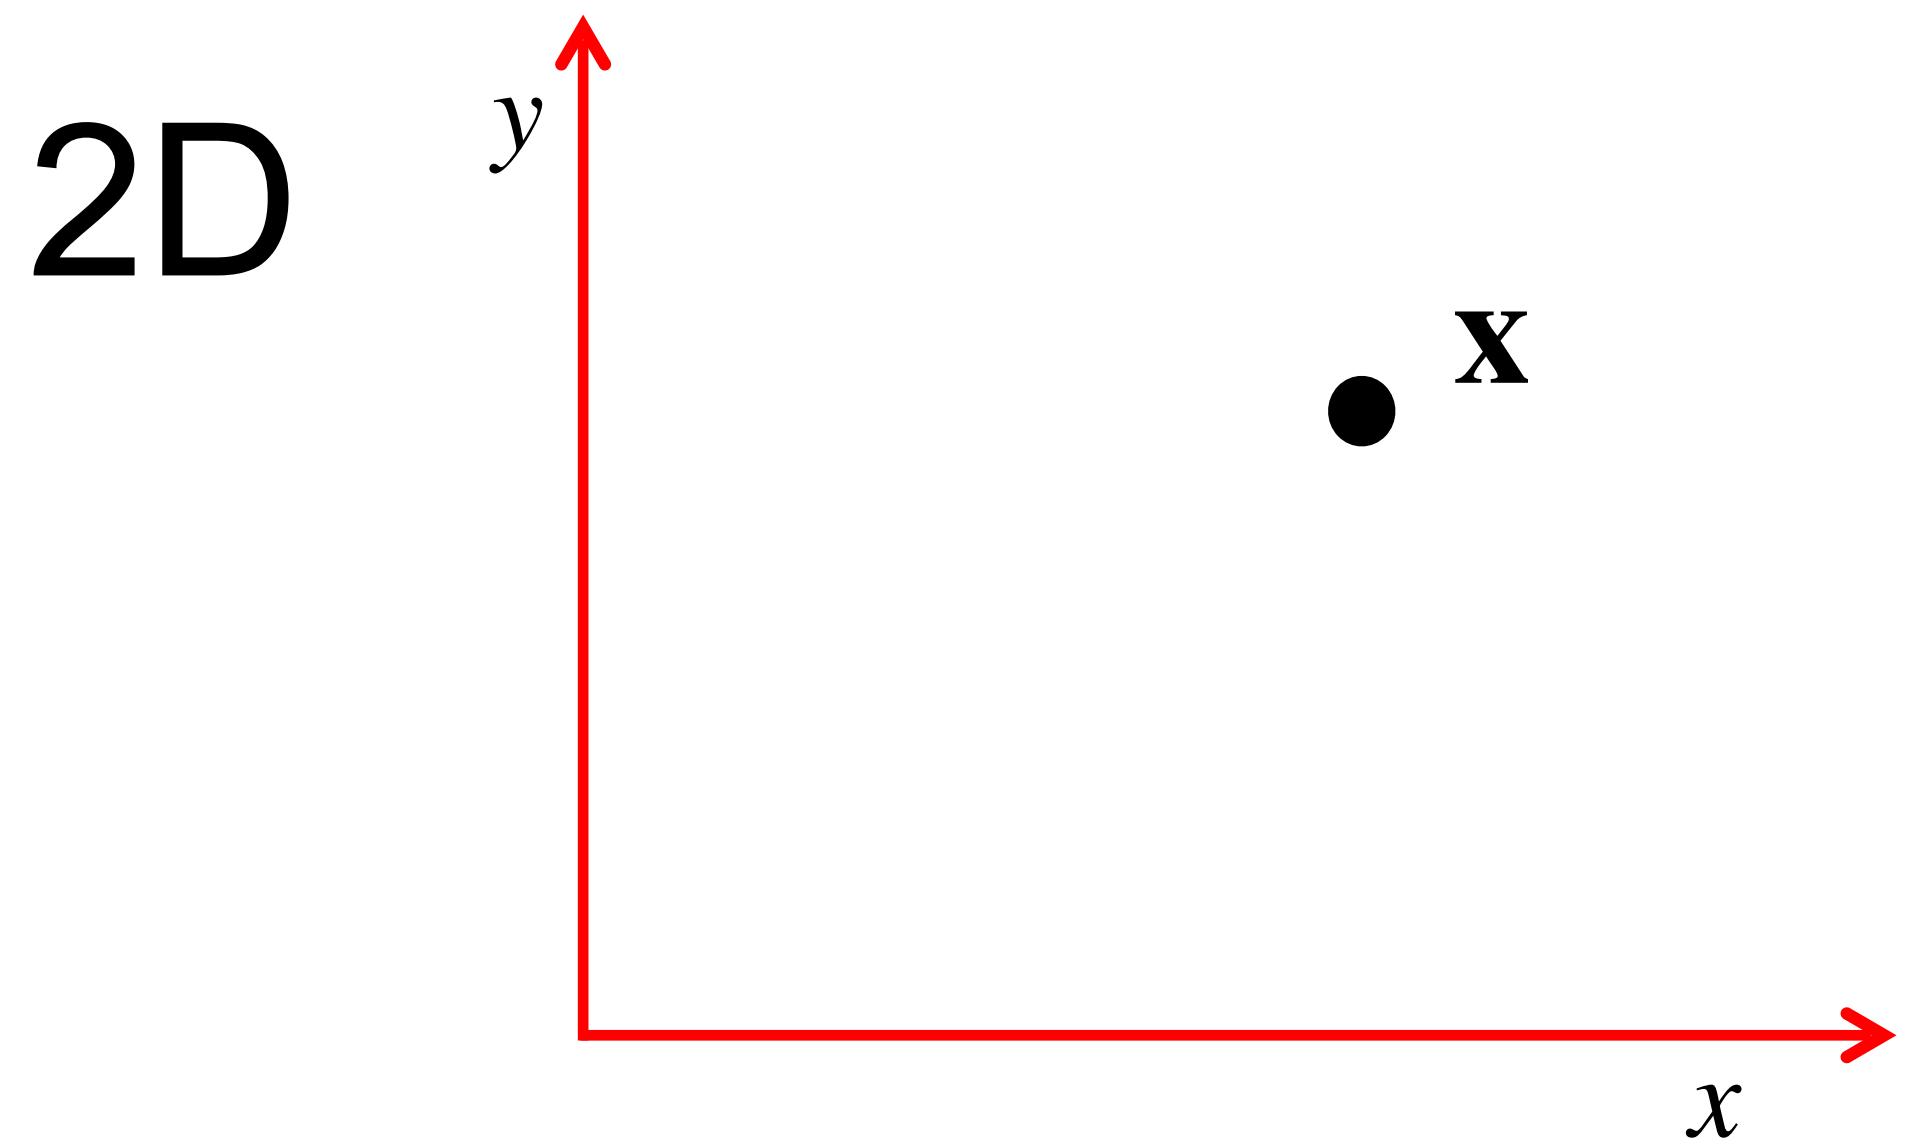
Perspective projection



$$\mathbf{x} = (x, y) = (X, Y) \cdot \frac{f}{Z}$$

This is awkward... Not a linear operator.

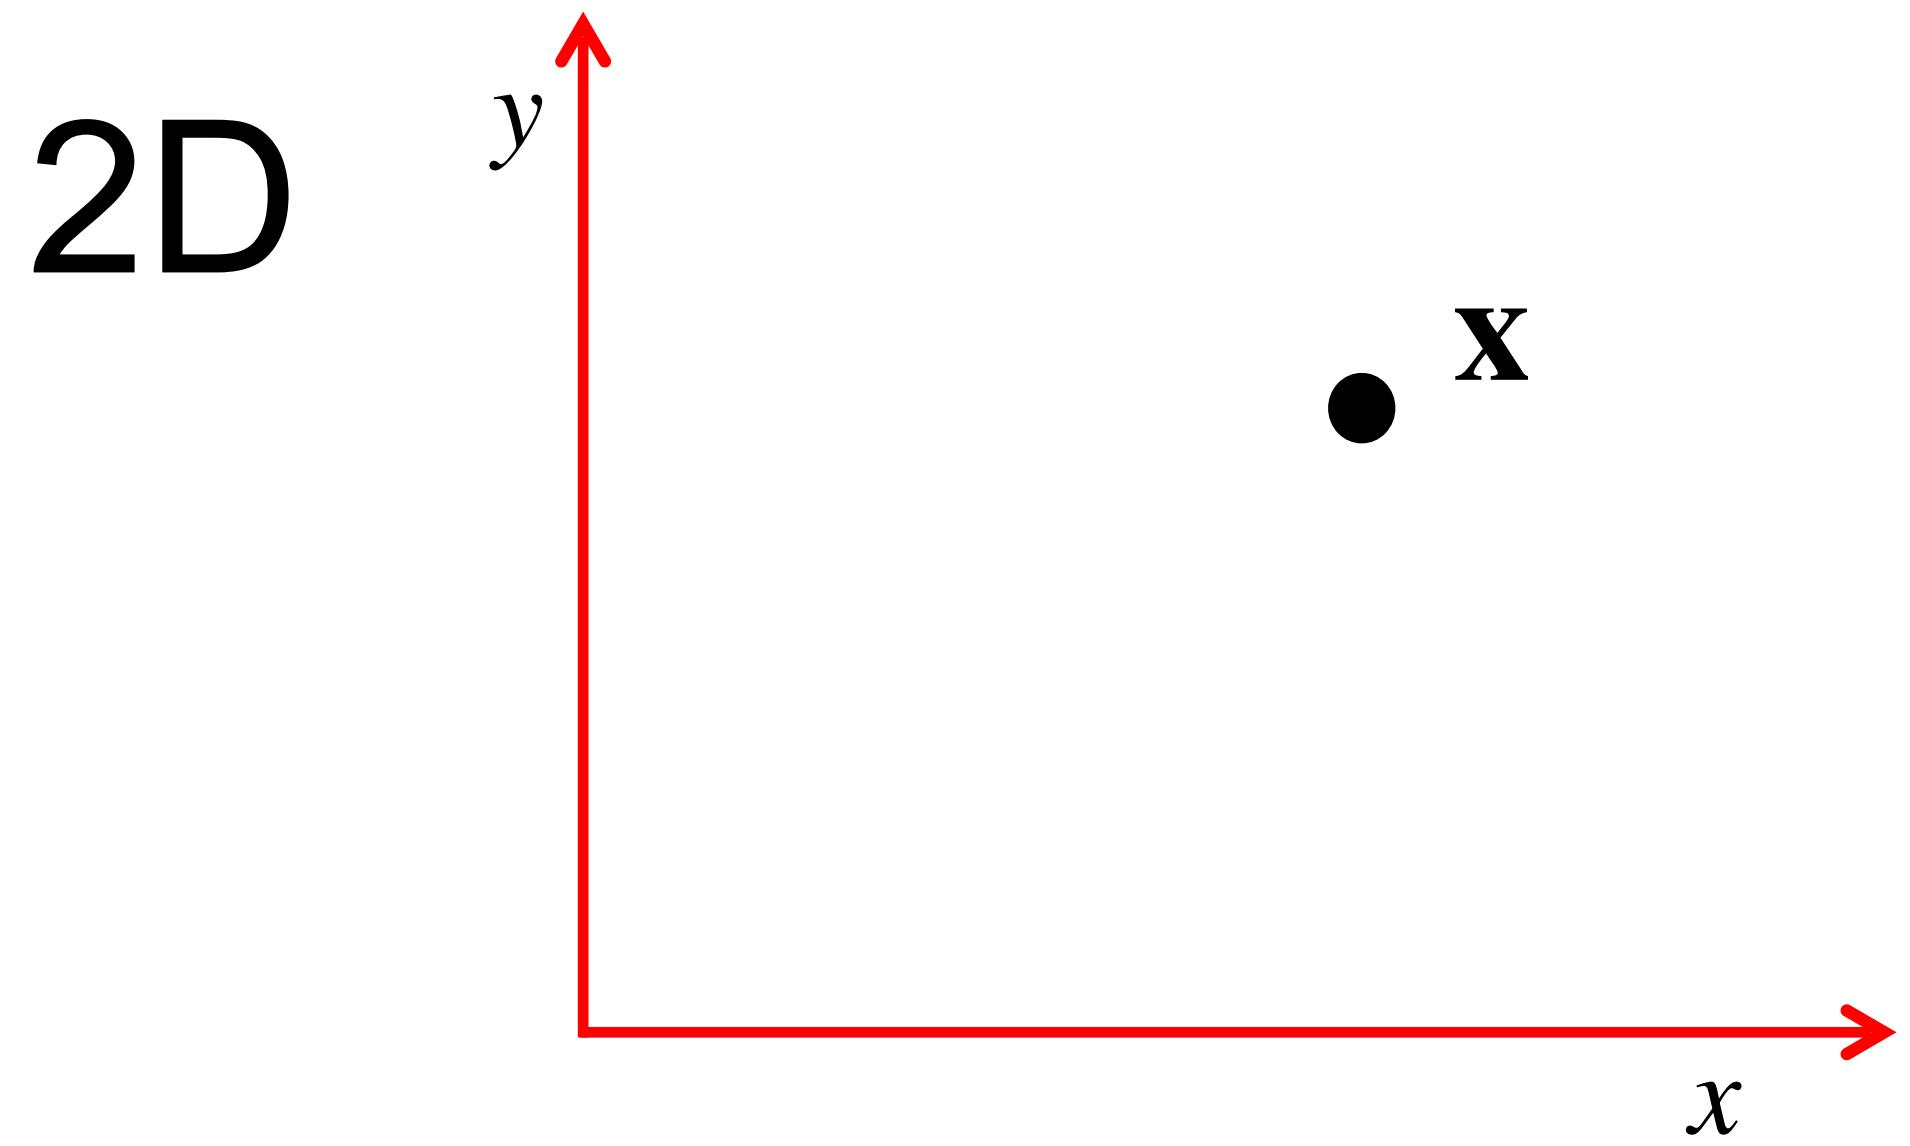
Homogeneous coordinates



$$\mathbf{x} = (x, y) \rightarrow \tilde{\mathbf{x}} = \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

heterogeneous coordinates homogeneous coordinates

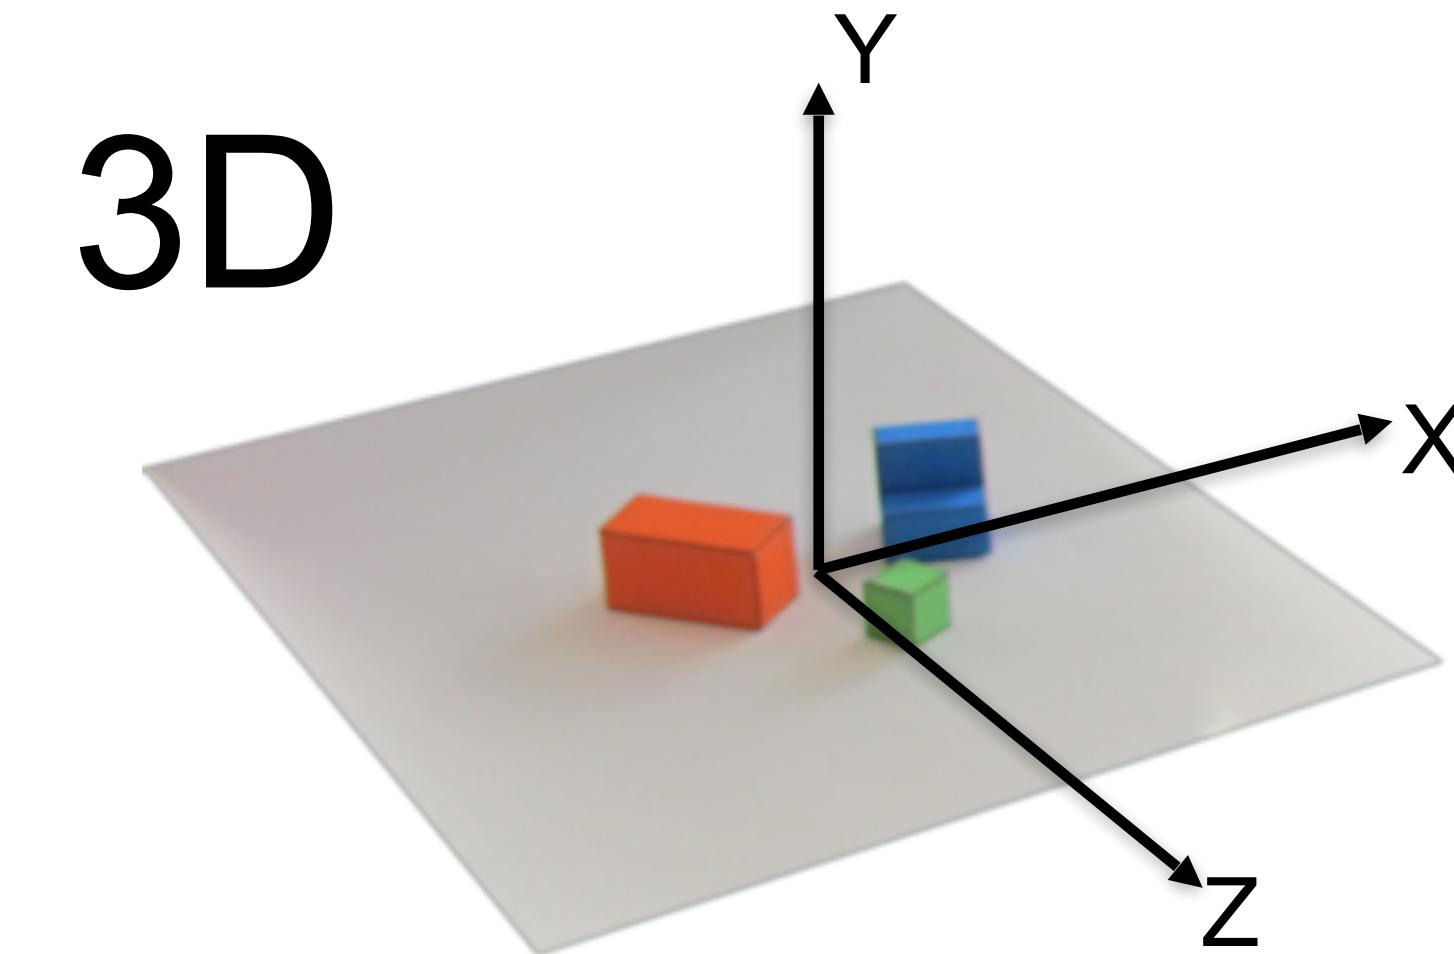
Homogeneous coordinates



$$\mathbf{x} = (x, y) \rightarrow \tilde{\mathbf{x}} = \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

heterogeneous
coordinates

homogeneous
coordinates



$$\mathbf{X} = (X, Y, Z) \rightarrow \tilde{\mathbf{X}} = \begin{bmatrix} X \\ Y \\ Z \\ 1 \end{bmatrix}$$

From heterogeneous to homogeneous:

$$\mathbf{x} = (x, y) \quad \rightarrow \quad \tilde{\mathbf{x}} = \begin{bmatrix} x \\ y \\ 1 \end{bmatrix} = \begin{bmatrix} w \cdot x \\ w \cdot y \\ w \cdot 1 \end{bmatrix}$$

From homogeneous to heterogeneous:

$$\tilde{\mathbf{x}} = \begin{bmatrix} w \cdot x \\ w \cdot y \\ w \end{bmatrix} \quad \rightarrow \quad \mathbf{x} = (x/w, y/w)$$

Homogeneous coordinates

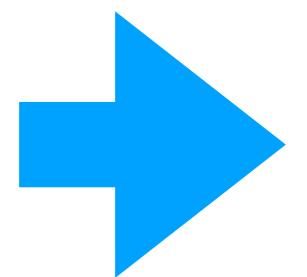
2D

$$\mathbb{R}^2$$

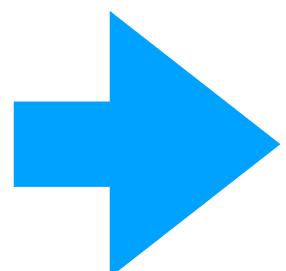
$$(x, y)$$

$$\mathbb{P}^2$$

$$\begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$



$$\begin{bmatrix} w \cdot x \\ w \cdot y \\ w \end{bmatrix}$$



$$(x/w, y/w)$$

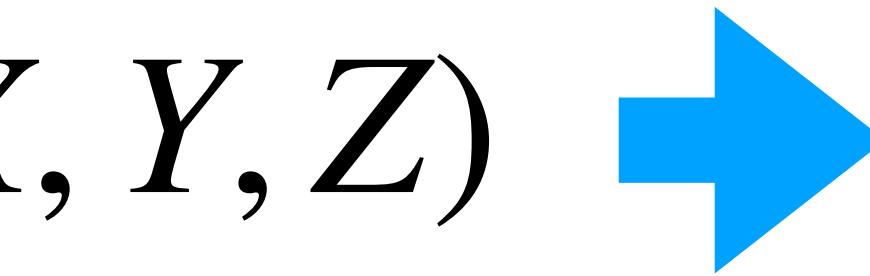
3D

$$\mathbb{R}^3$$

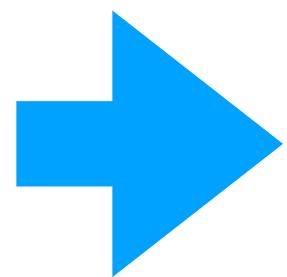
$$(X, Y, Z)$$

$$\mathbb{P}^3$$

$$\begin{bmatrix} X \\ Y \\ Z \\ W \end{bmatrix}$$



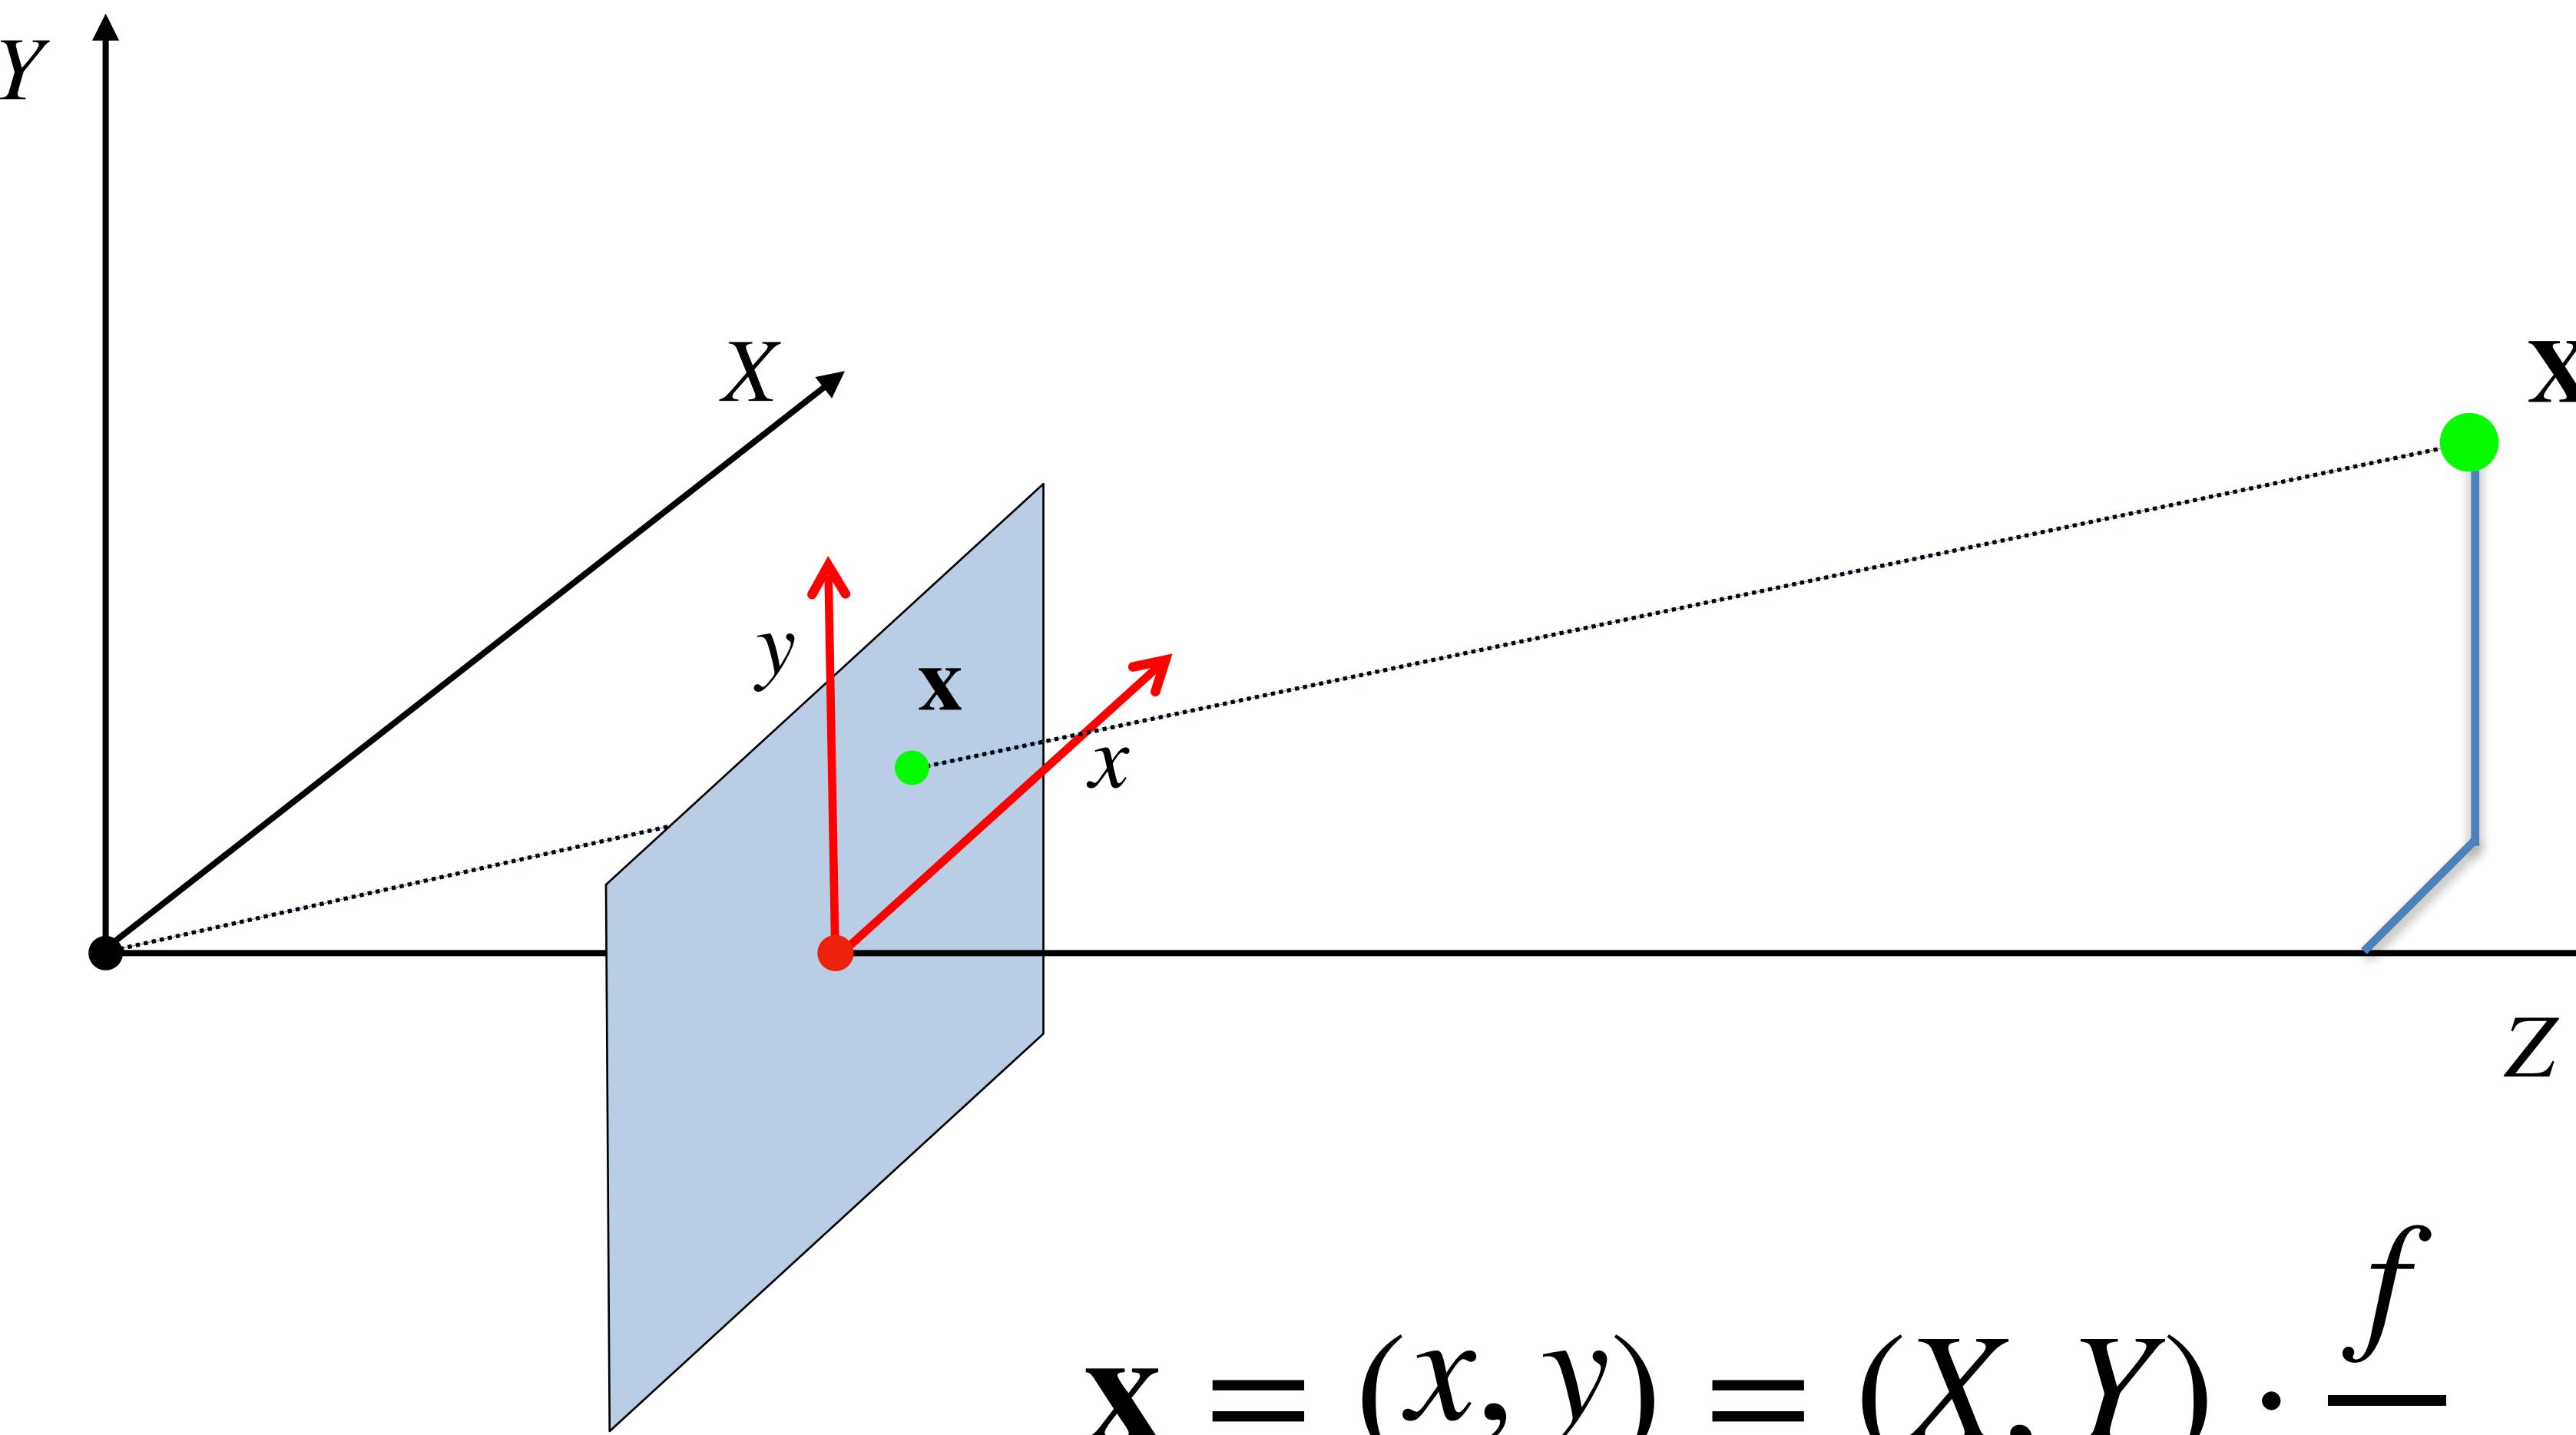
$$\begin{bmatrix} X \\ Y \\ Z \\ W \end{bmatrix}$$



$$(X/W, Y/W, Z/W)$$

Questions?

Perspective projection



$$\mathbf{x} = (x, y) = (X, Y) \cdot \frac{f}{Z}$$

This is awkward... Not a linear operator.

Perspective projection

Heterogeneous coordinates

$$\mathbf{x} = (x, y) = (X, Y) \cdot \frac{1}{Z}$$

Image / Pixel
Coordinates

World Coordinates

Homogeneous coordinates

$$\begin{bmatrix} x \\ y \\ w \end{bmatrix} = ? \begin{bmatrix} X \\ Y \\ Z \\ 1 \end{bmatrix}$$

Image / Pixel
Coordinates

World Coordinates

Perspective projection

Heterogeneous coordinates

$$\mathbf{x} = (x, y) = (X, Y) \cdot \frac{1}{Z}$$

World Coordinates

Image / Pixel
Coordinates

Homogeneous coordinates

$$\begin{bmatrix} x \\ y \\ w \end{bmatrix} = \begin{bmatrix} ? & ? & ? & ? \\ ? & ? & ? & ? \\ ? & ? & ? & ? \end{bmatrix} \begin{bmatrix} X \\ Y \\ Z \\ 1 \end{bmatrix}$$

Image / Pixel
Coordinates

World
coords.

Perspective projection

Heterogeneous coordinates

$$\mathbf{x} = (x, y) = (X, Y) \cdot \frac{1}{Z}$$

World Coordinates

Image / Pixel
Coordinates

Homogeneous coordinates

$$\begin{bmatrix} x \\ y \\ w \end{bmatrix} = \begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 1/f & 0 \end{bmatrix} \begin{bmatrix} X \\ Y \\ Z \\ 1 \end{bmatrix}$$

Image / Pixel
Coordinates

Projection Matrix

World
coords.

Perspective projection

Heterogeneous coordinates

$$\mathbf{x} = (x, y) = (X, Y) \cdot \frac{1}{Z}$$

Image / Pixel
Coordinates

World Coordinates

Homogeneous coordinates

$$\begin{bmatrix} x \\ y \\ w \end{bmatrix} = \begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 1/f & 0 \end{bmatrix} \begin{bmatrix} X \\ Y \\ Z \\ 1 \end{bmatrix} = \begin{bmatrix} X \\ Y \\ Z \\ Z/f \end{bmatrix}$$

Image / Pixel
Coordinates

Projection Matrix

World
coords.

Perspective projection

Homogeneous coordinates

$$\begin{bmatrix} x \\ y \\ w \end{bmatrix} = \begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 1/f & 0 \end{bmatrix} \begin{bmatrix} X \\ Y \\ Z \\ 1 \end{bmatrix} = \begin{bmatrix} X \\ Y \\ Z/f \end{bmatrix}$$

Image / Pixel
Coordinates Projection Matrix World
coords.

Perspective projection

$$\begin{bmatrix} x \\ y \\ w \end{bmatrix} = \begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 1/f & 0 \end{bmatrix} \begin{bmatrix} X \\ Y \\ Z \\ 1 \end{bmatrix} = \begin{bmatrix} X \\ Y \\ Z/f \end{bmatrix}$$

Image / Pixel
Coordinates Projection Matrix World
coords.

$$\begin{bmatrix} X \\ Y \\ Z/f \end{bmatrix} \xrightarrow{\text{blue arrow}} \left(f \frac{X}{Z}, f \frac{Y}{Z} \right)$$

Perspective projection

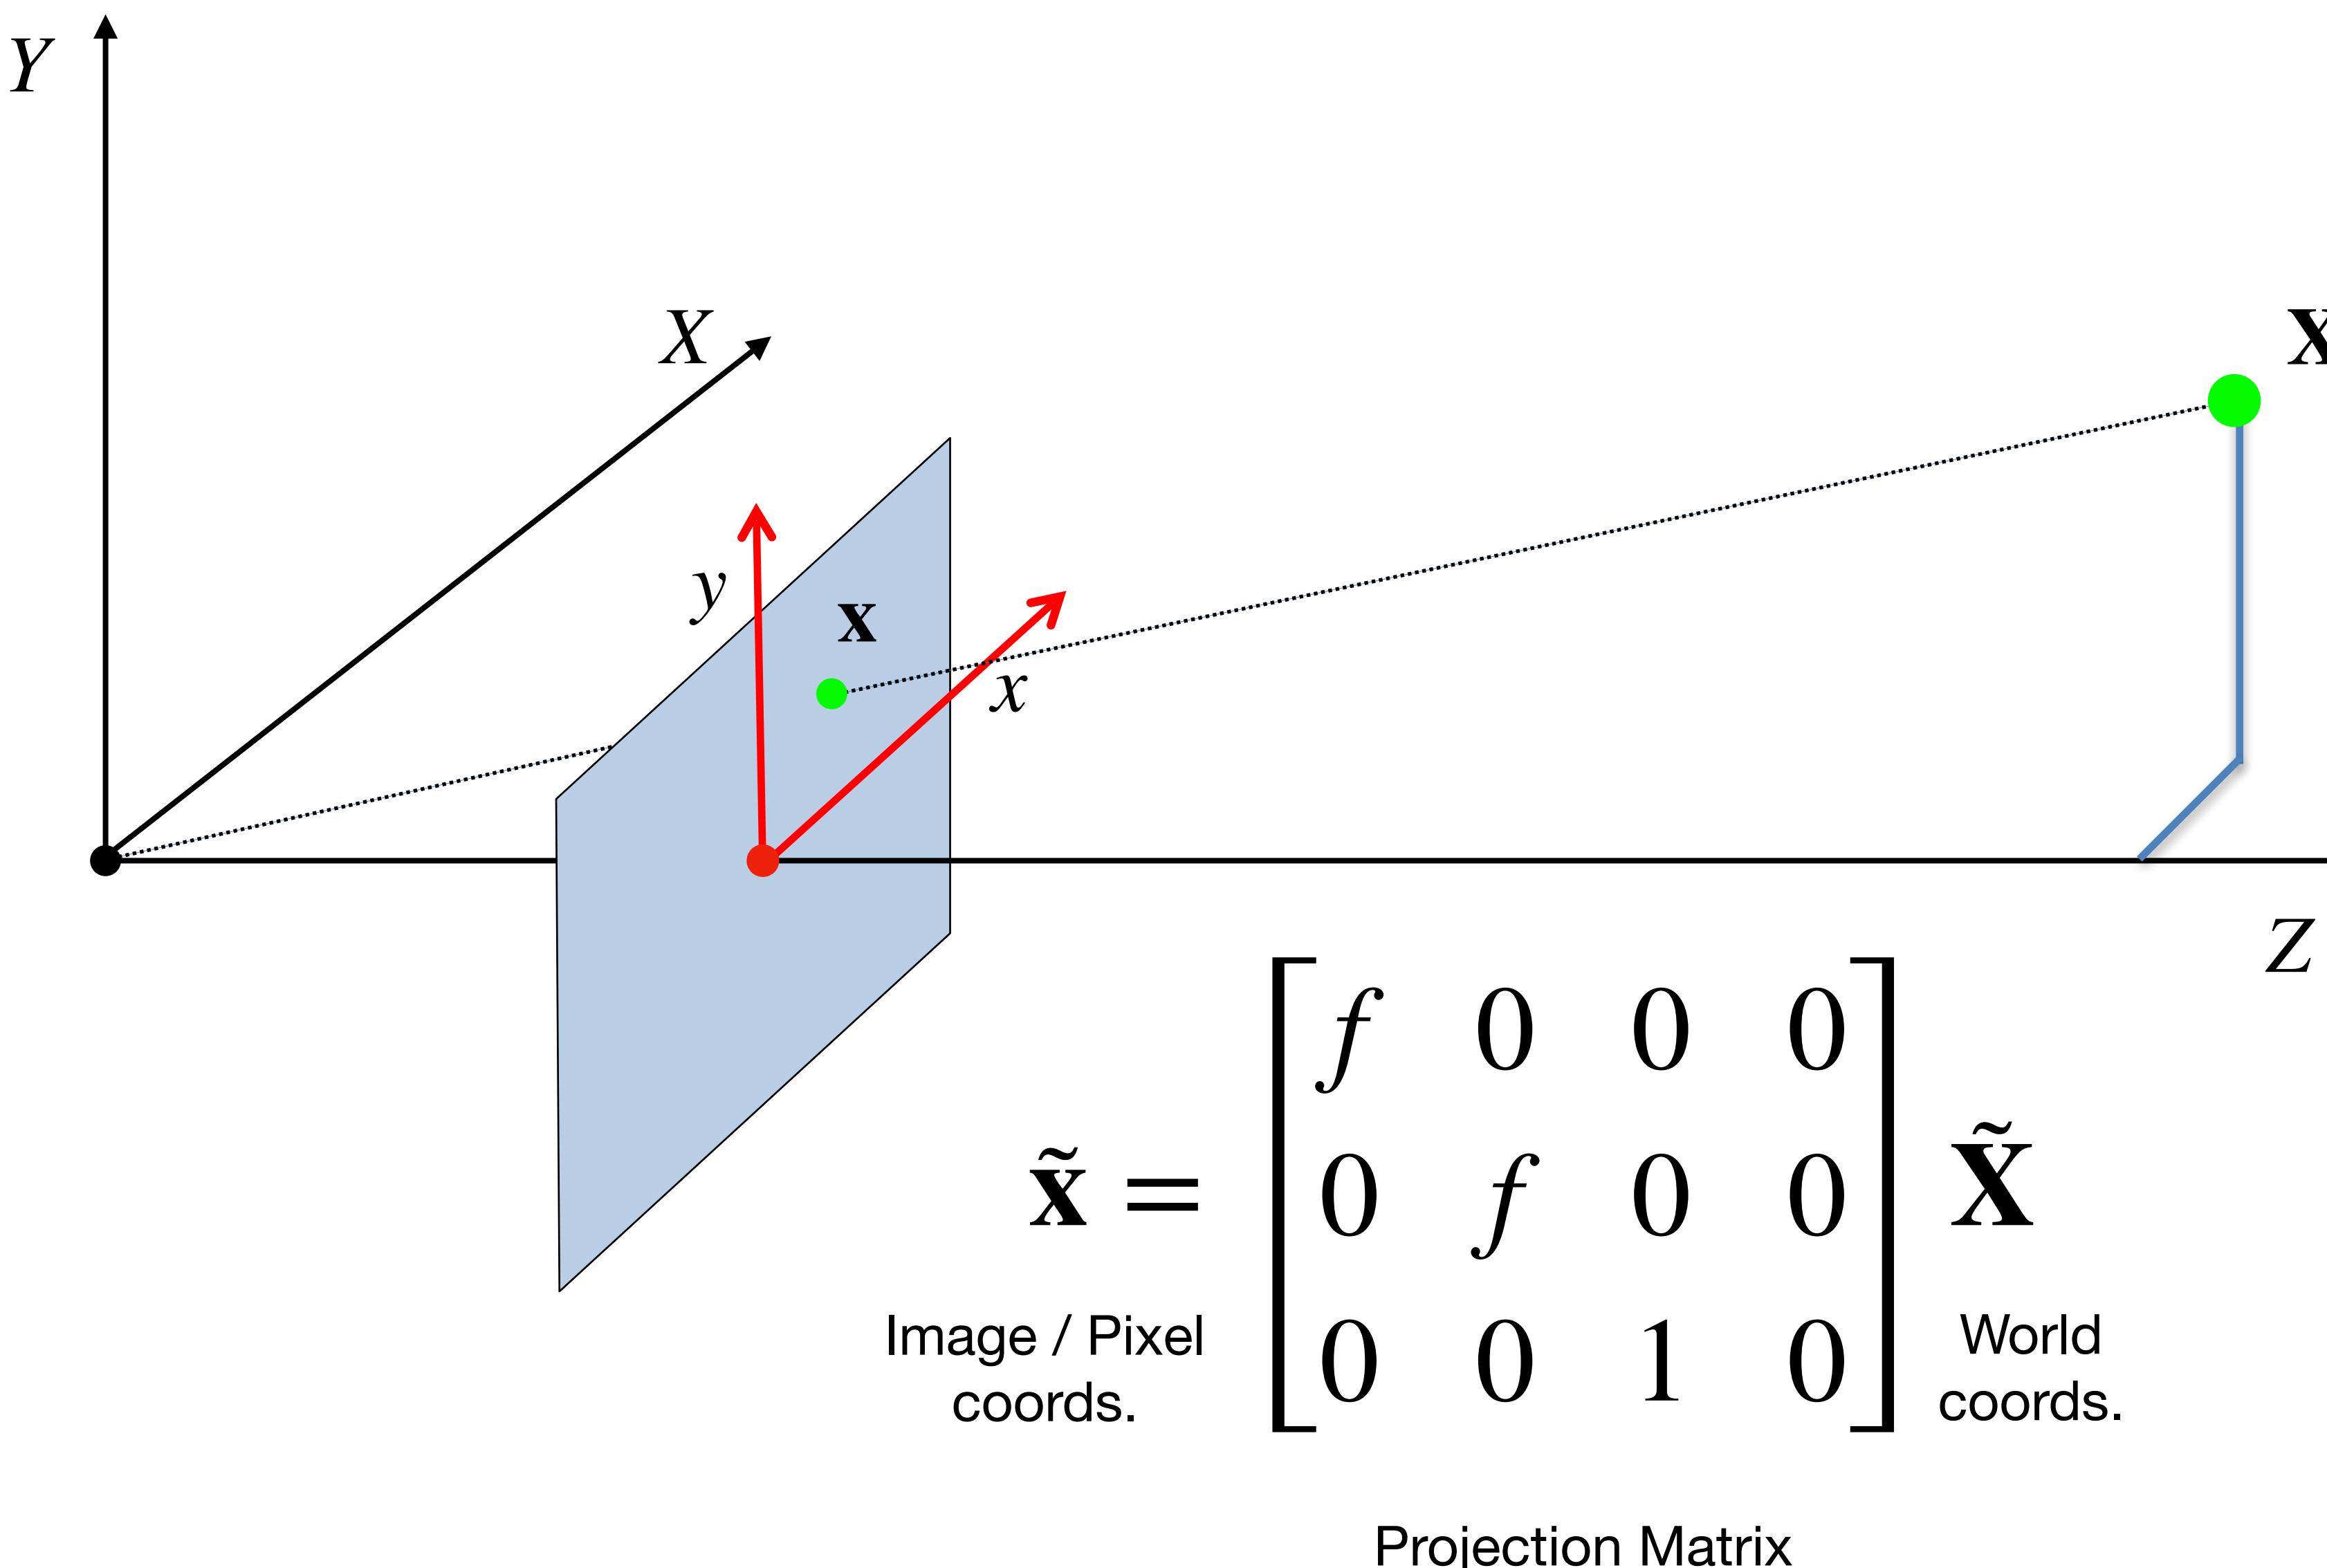
$$\tilde{\mathbf{x}} = \begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 1/f & 0 \end{bmatrix} \tilde{\mathbf{X}} = \begin{bmatrix} X \\ Y \\ Z/f \end{bmatrix} \rightarrow \left(f \frac{X}{Z}, f \frac{Y}{Z} \right)$$

$$\tilde{\mathbf{x}} = \begin{bmatrix} f & 0 & 0 & 0 \\ 0 & f & 0 & 0 \\ 0 & 0 & 1 & 0 \end{bmatrix} \tilde{\mathbf{X}} = \begin{bmatrix} f \cdot X \\ f \cdot Y \\ Z \end{bmatrix} \rightarrow \left(f \frac{X}{Z}, f \frac{Y}{Z} \right)$$

Perspective projection

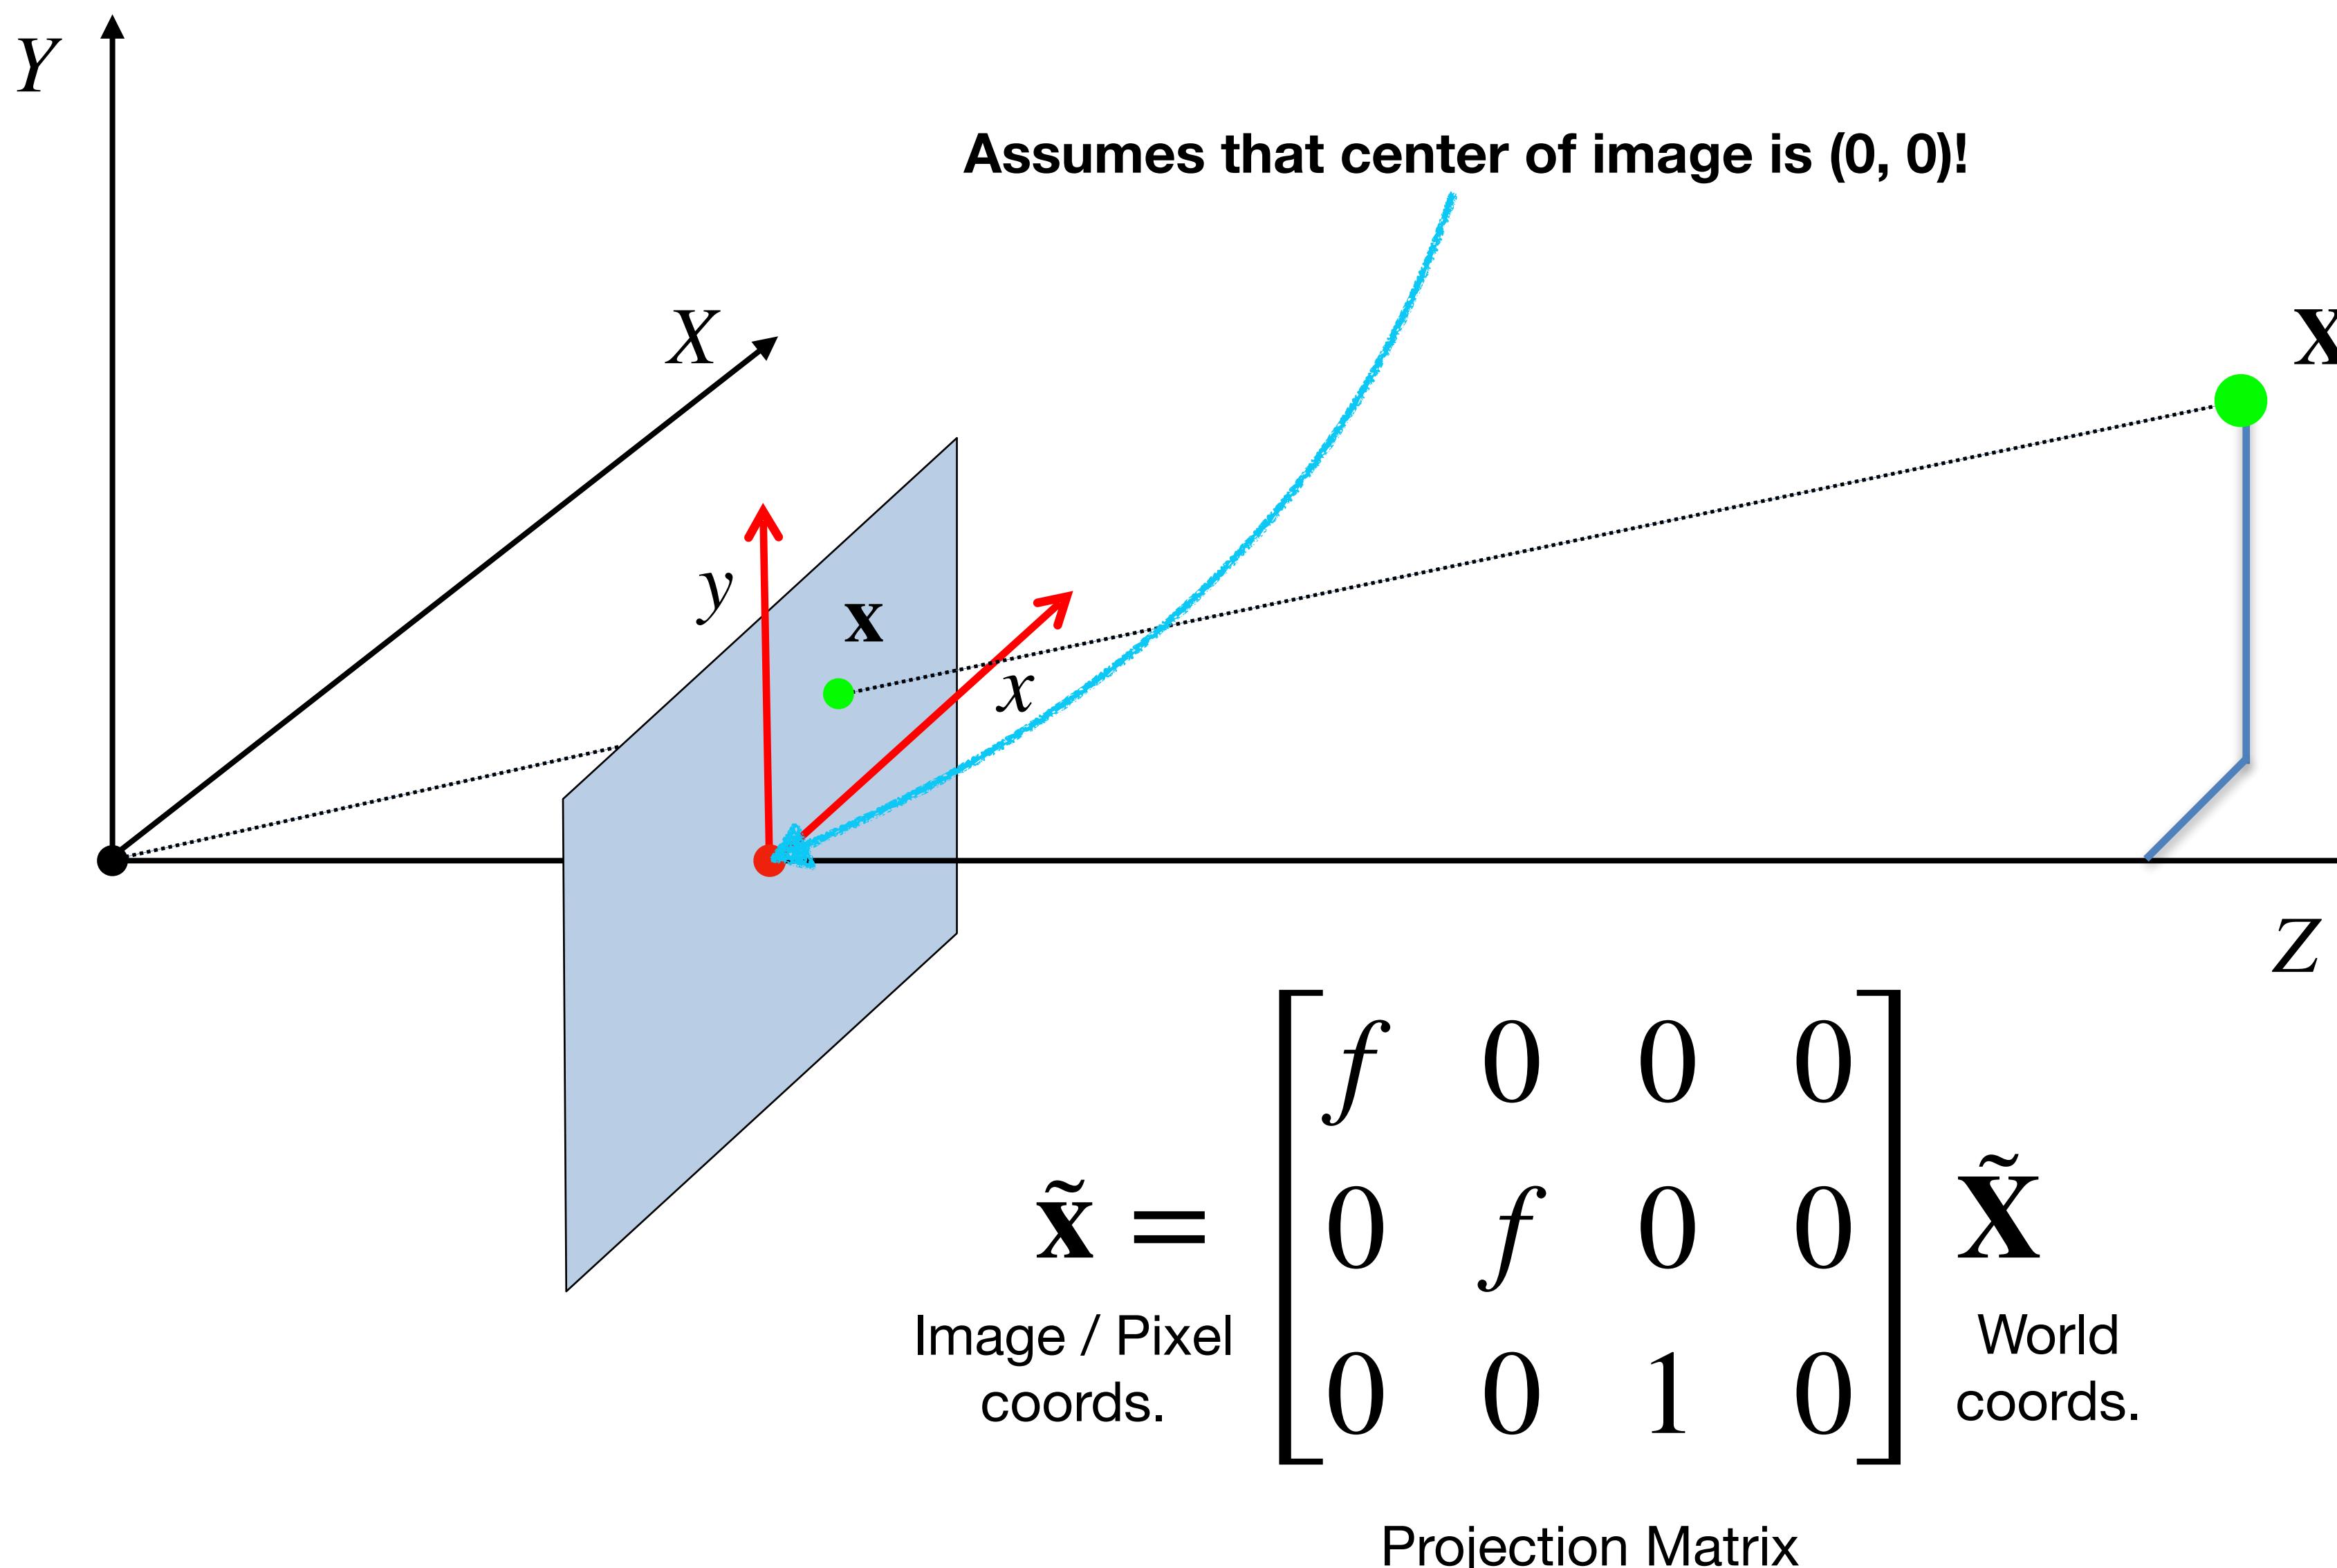
$$\tilde{\mathbf{x}} = \begin{bmatrix} f & 0 & 0 & 0 \\ 0 & f & 0 & 0 \\ 0 & 0 & 1 & 0 \end{bmatrix} \tilde{\mathbf{X}}$$

Perspective projection

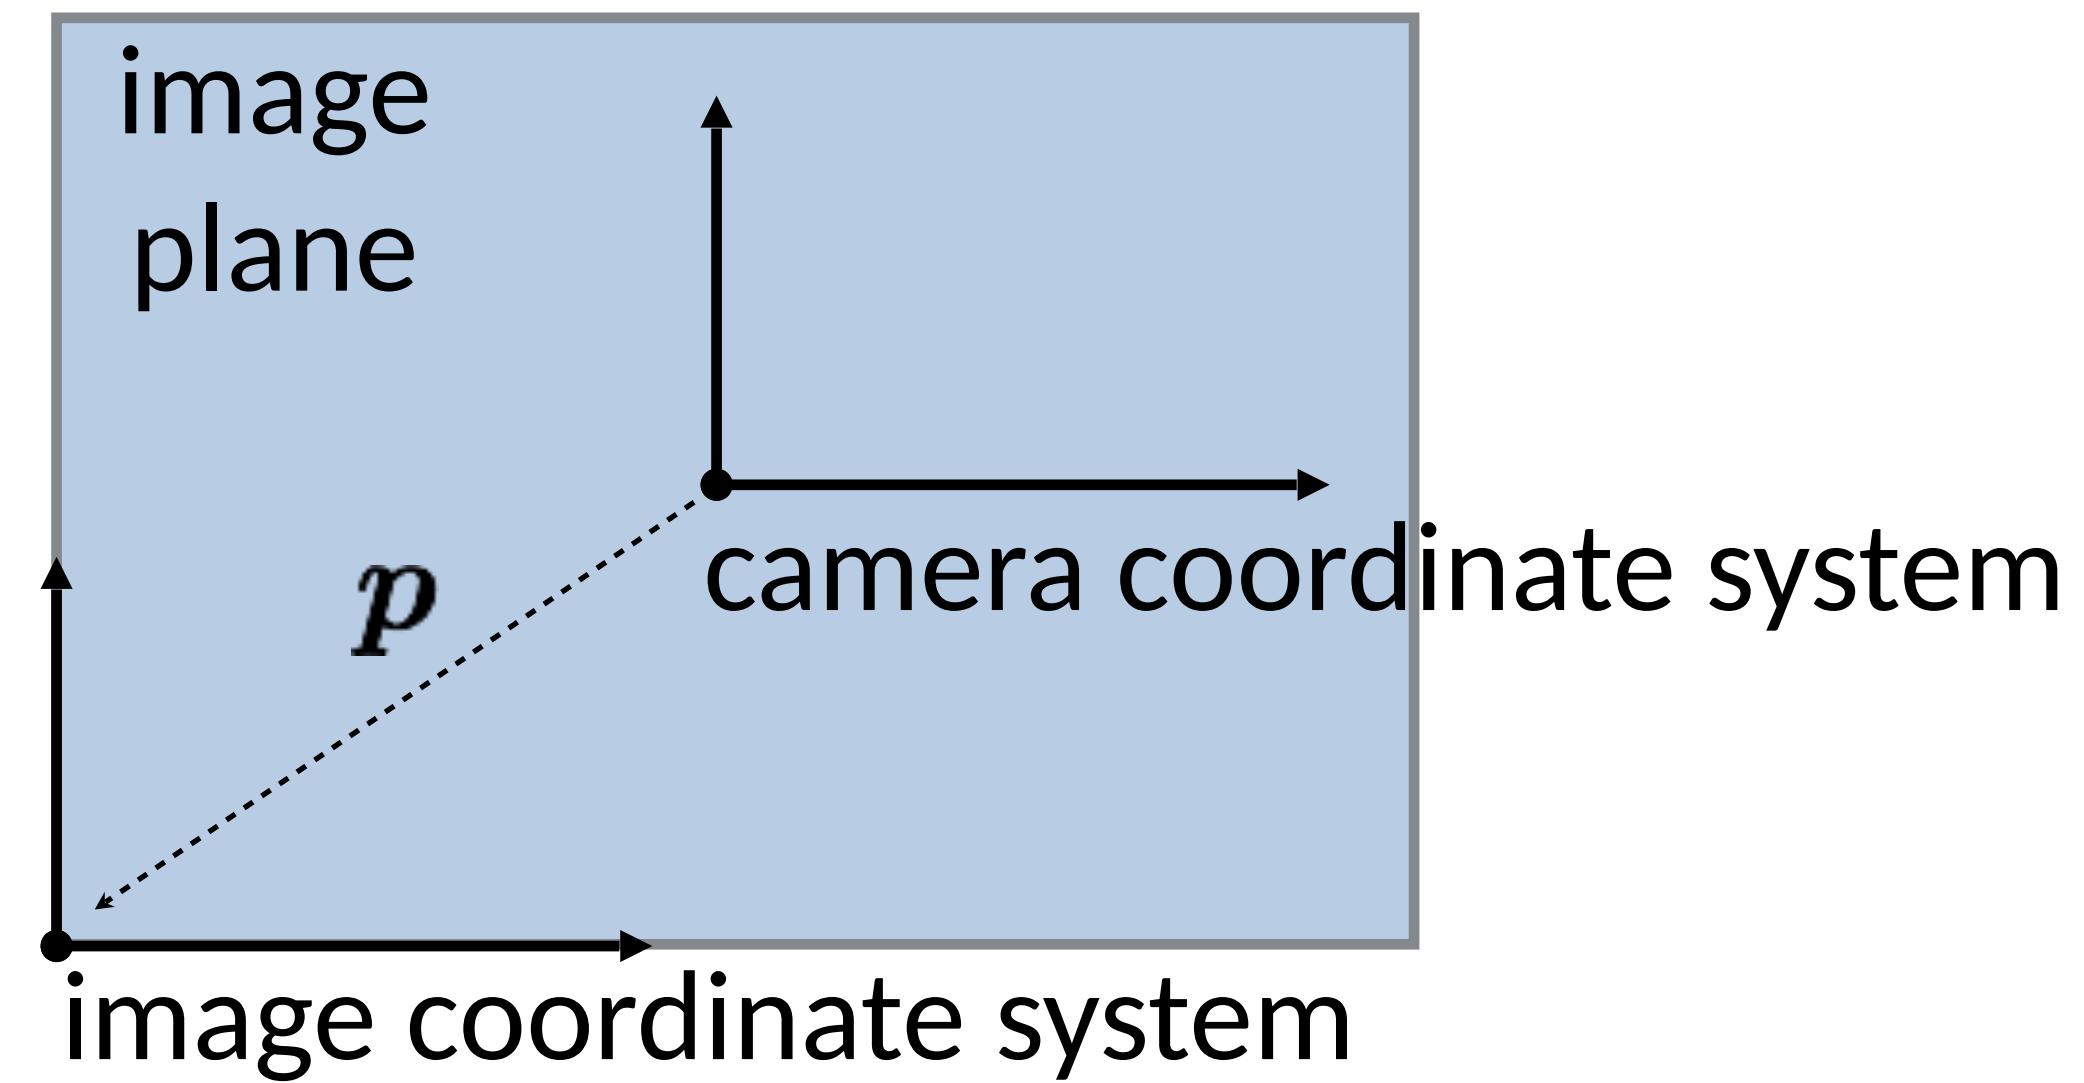


Questions?

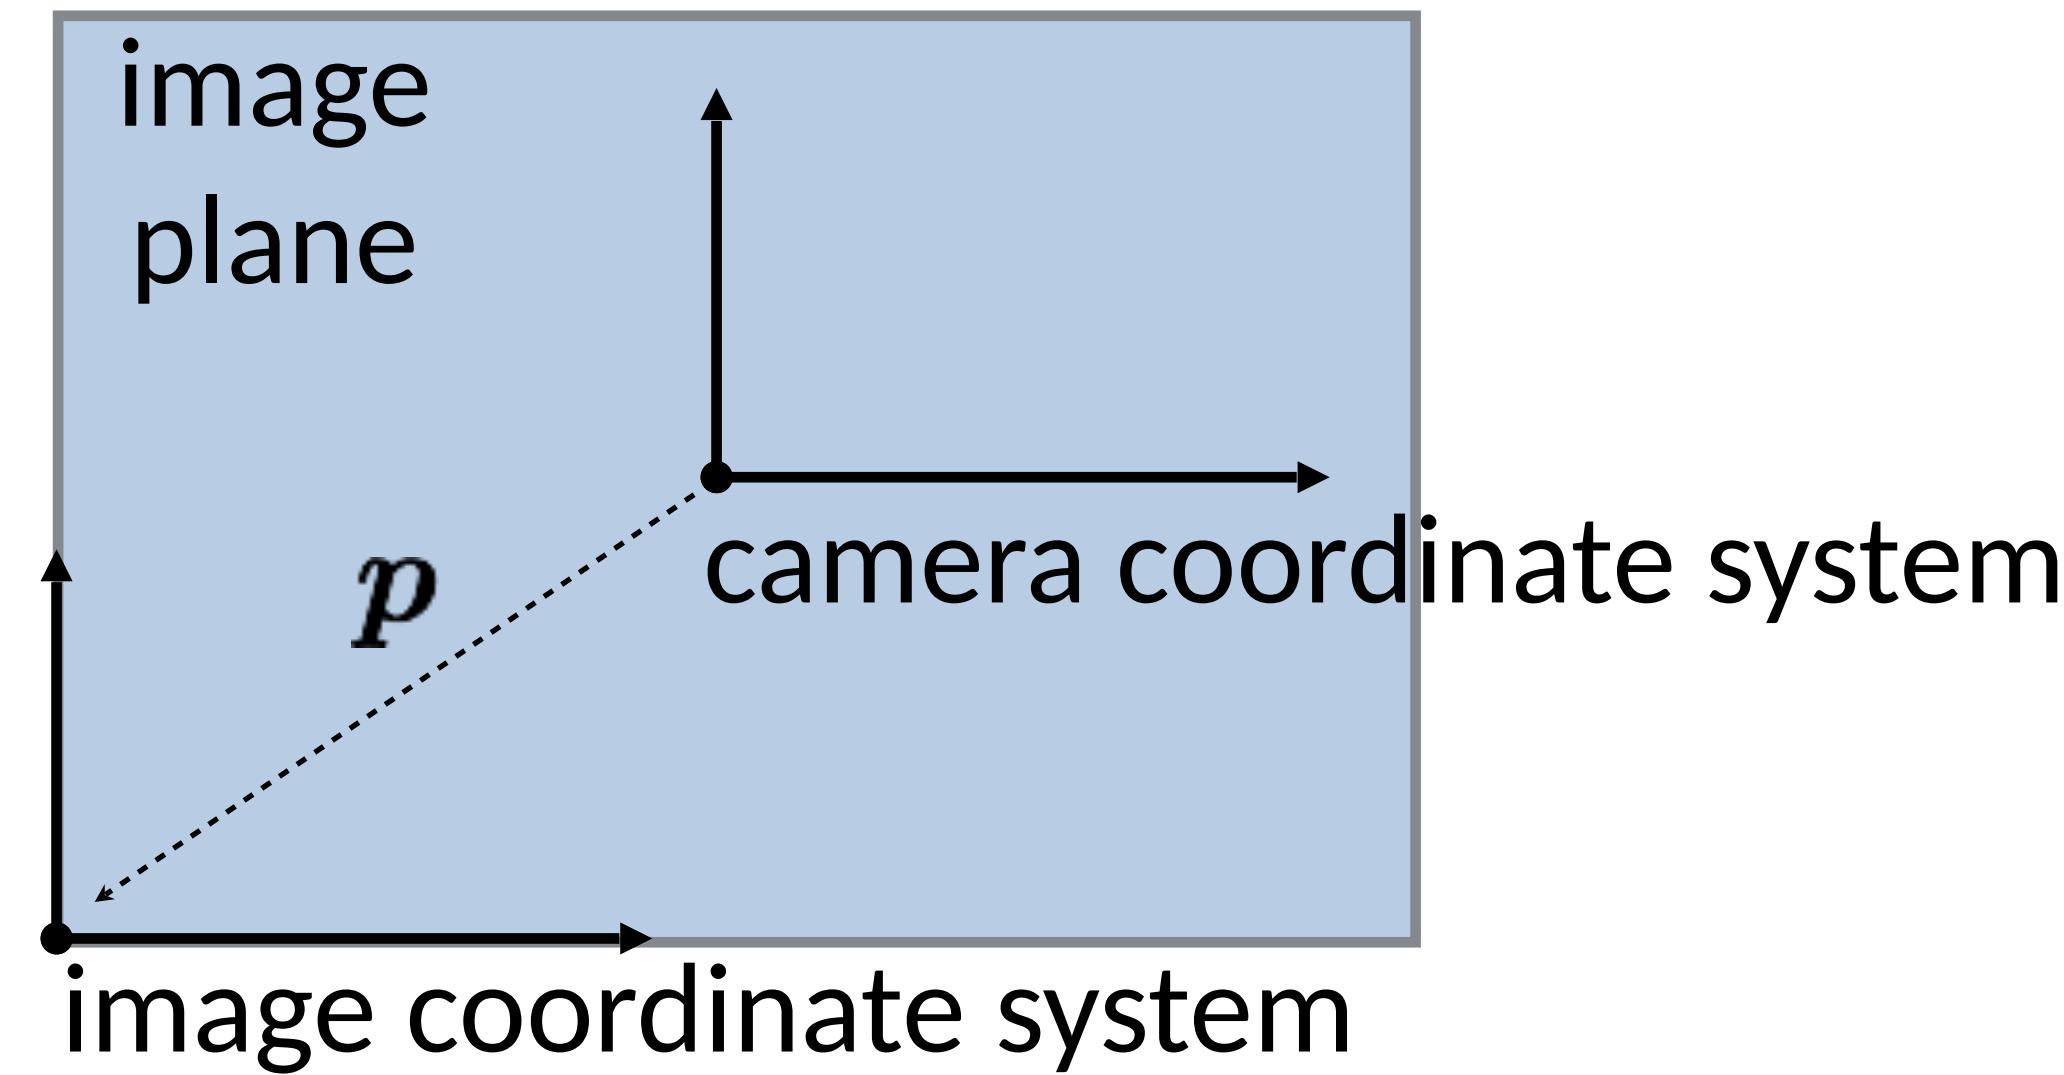
Perspective projection



More General Case: Arbitrary Image Centre



More General Case: Arbitrary Image Centre



How does the projection matrix change?

$$\begin{bmatrix} f & 0 & p_x & 0 \\ 0 & f & p_y & 0 \\ 0 & 0 & 1 & 0 \end{bmatrix}$$

shift vector
transforming
camera origin to
image origin

Decomposing the Projection Matrix

We can decompose the projection matrix like this:

$$\begin{bmatrix} f & 0 & p_x & 0 \\ 0 & f & p_y & 0 \\ 0 & 0 & 1 & 0 \end{bmatrix} = \begin{bmatrix} f & 0 & p_x \\ 0 & f & p_y \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 \end{bmatrix}$$

What does each part of the matrix represent?

Decomposing the Projection Matrix

We can decompose the projection matrix like this:

$$\begin{bmatrix} f & 0 & p_x & 0 \\ 0 & f & p_y & 0 \\ 0 & 0 & 1 & 0 \end{bmatrix} = \begin{bmatrix} f & 0 & p_x \\ 0 & f & p_y \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 \end{bmatrix}$$



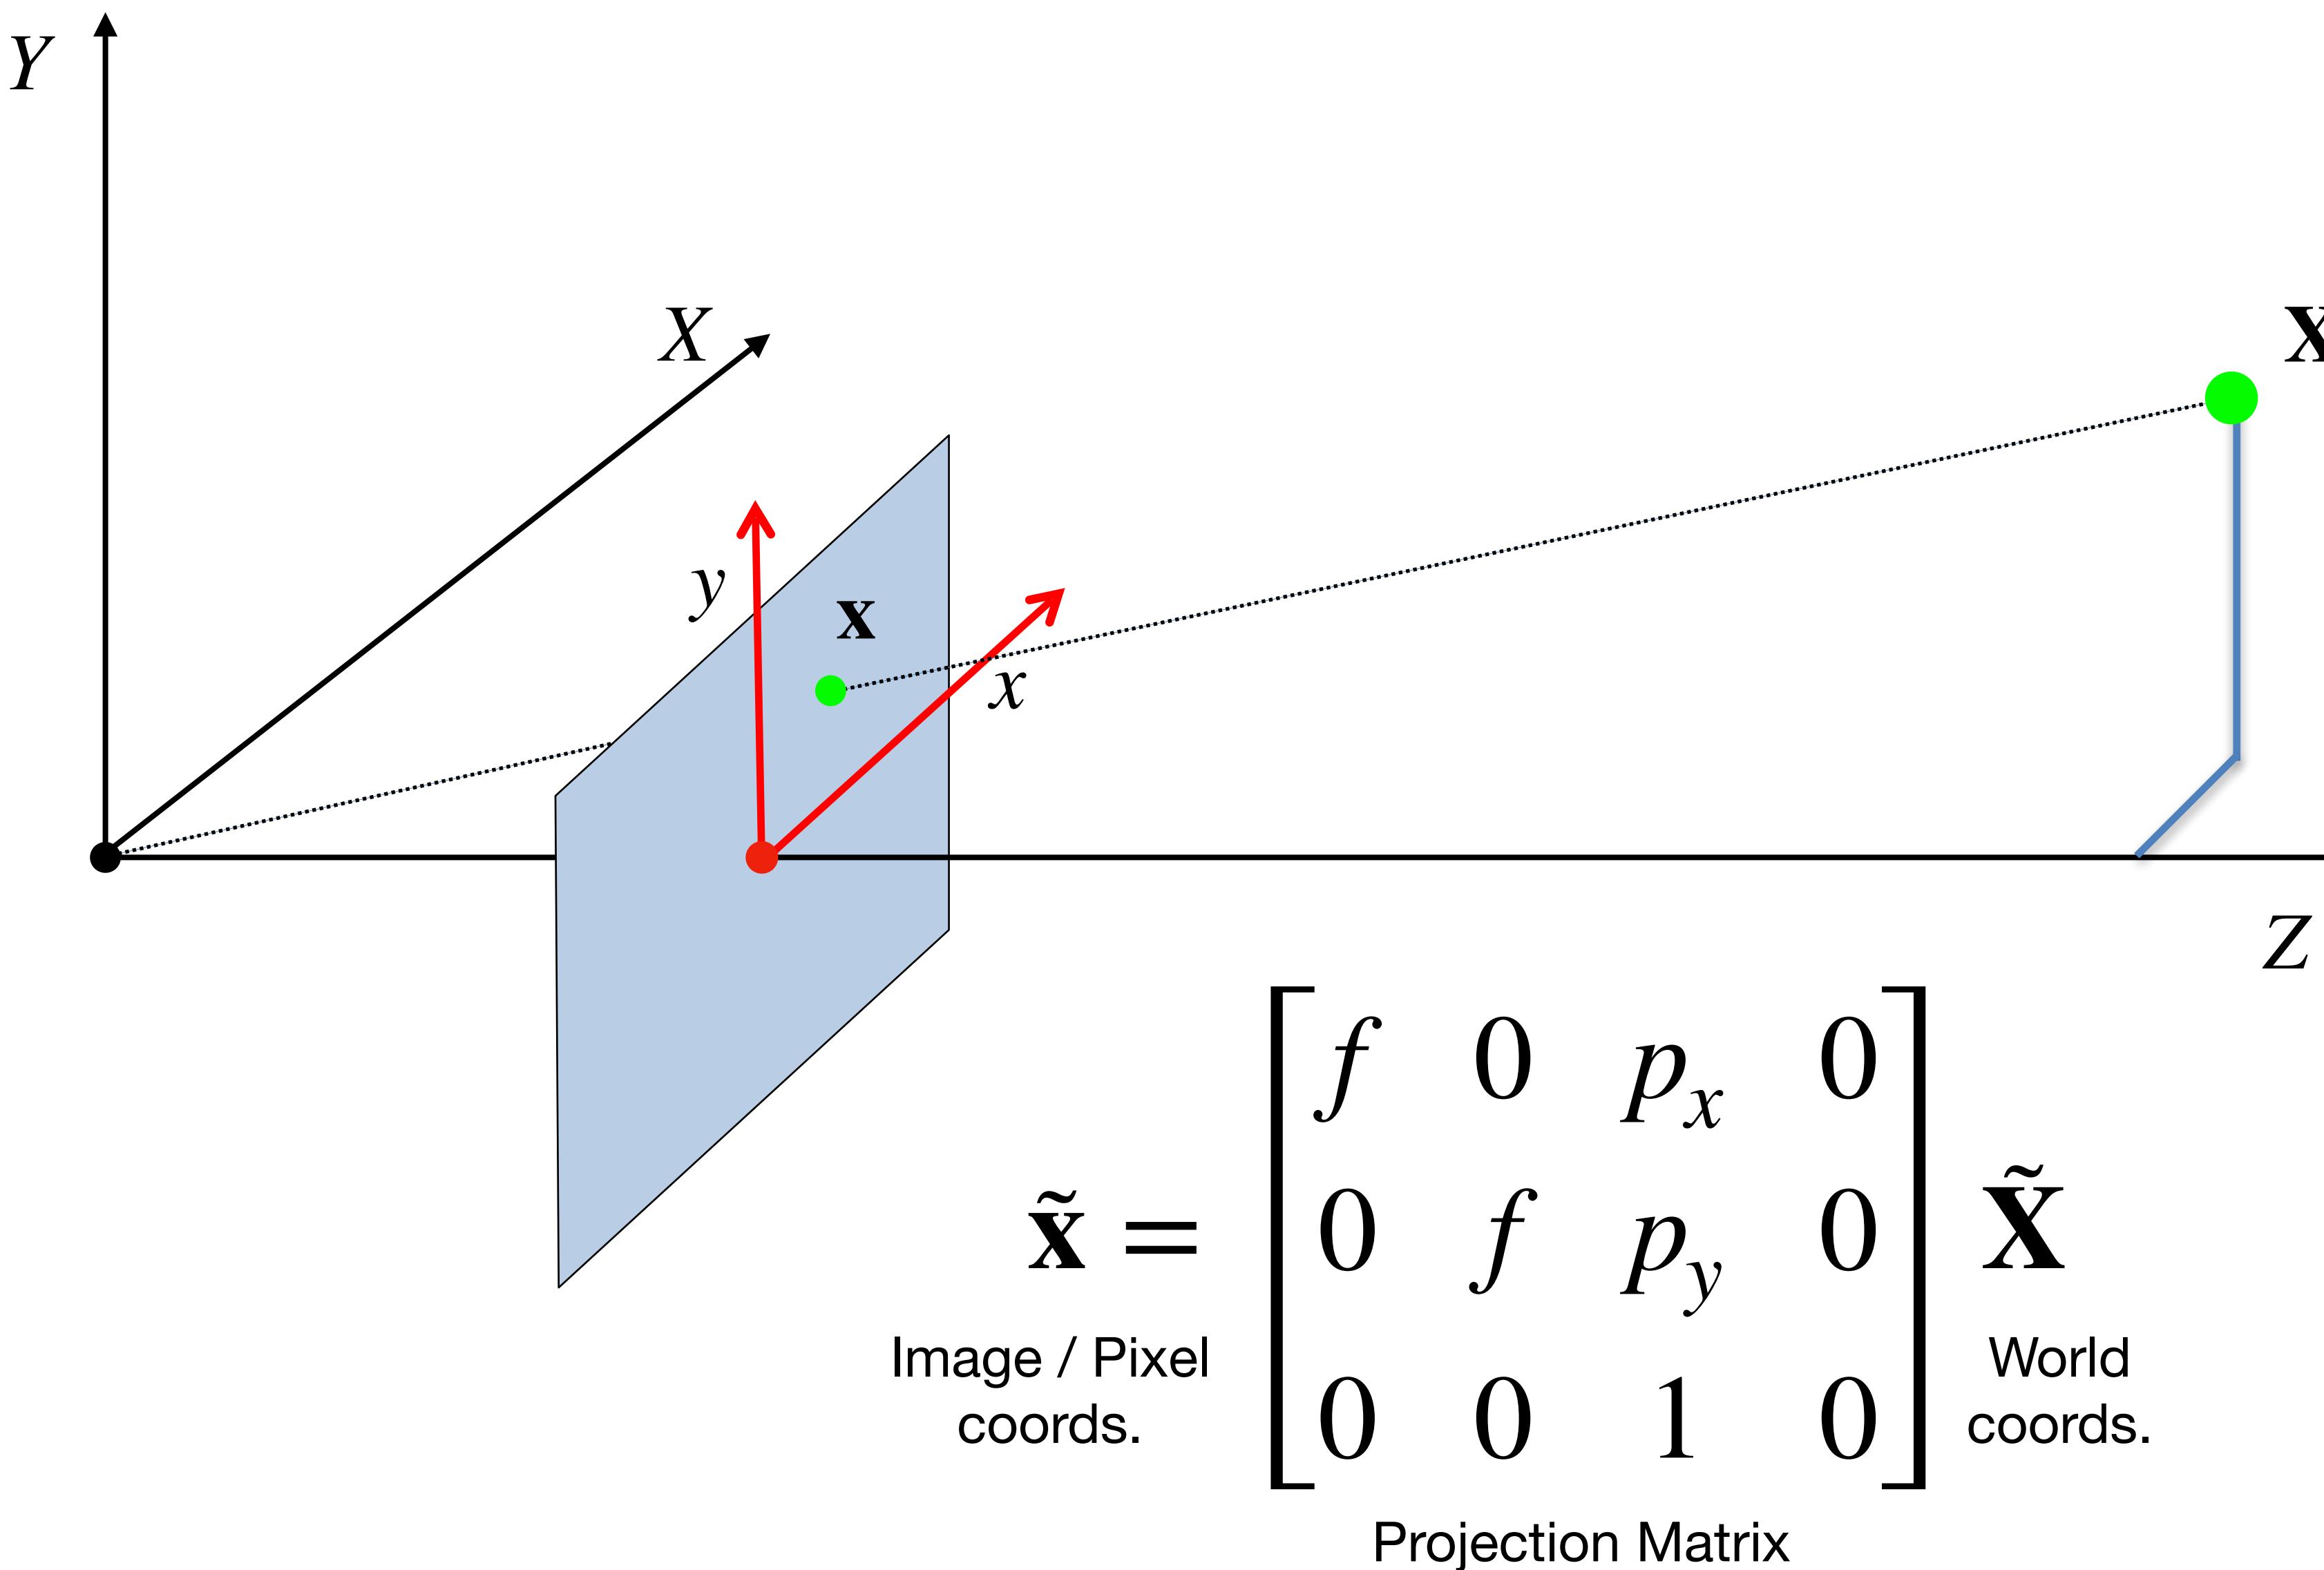
(homogeneous) transformation
from 2D to 2D, accounting for non
unit focal length and origin shift

(homogeneous) perspective projection
from 3D to 2D, assuming image plane at
 $z = 1$ and shared camera/image origin

Also written as: $\mathbf{K} [I | 0]$

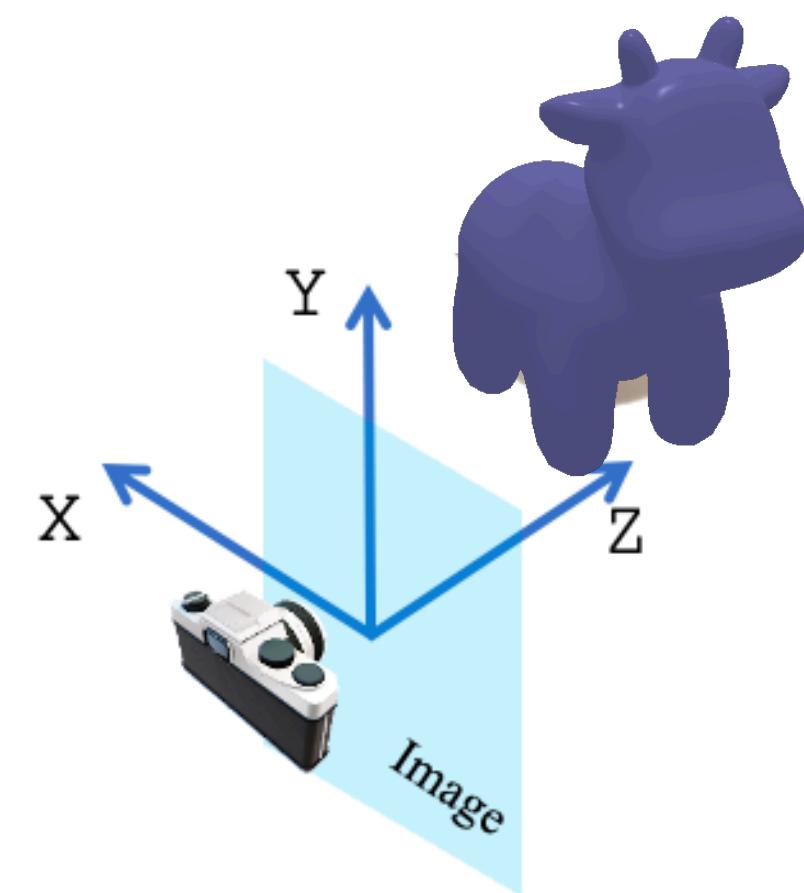
$$\text{where } \mathbf{K} = \begin{bmatrix} f & 0 & p_x \\ 0 & f & p_y \\ 0 & 0 & 1 \end{bmatrix}$$

Perspective projection

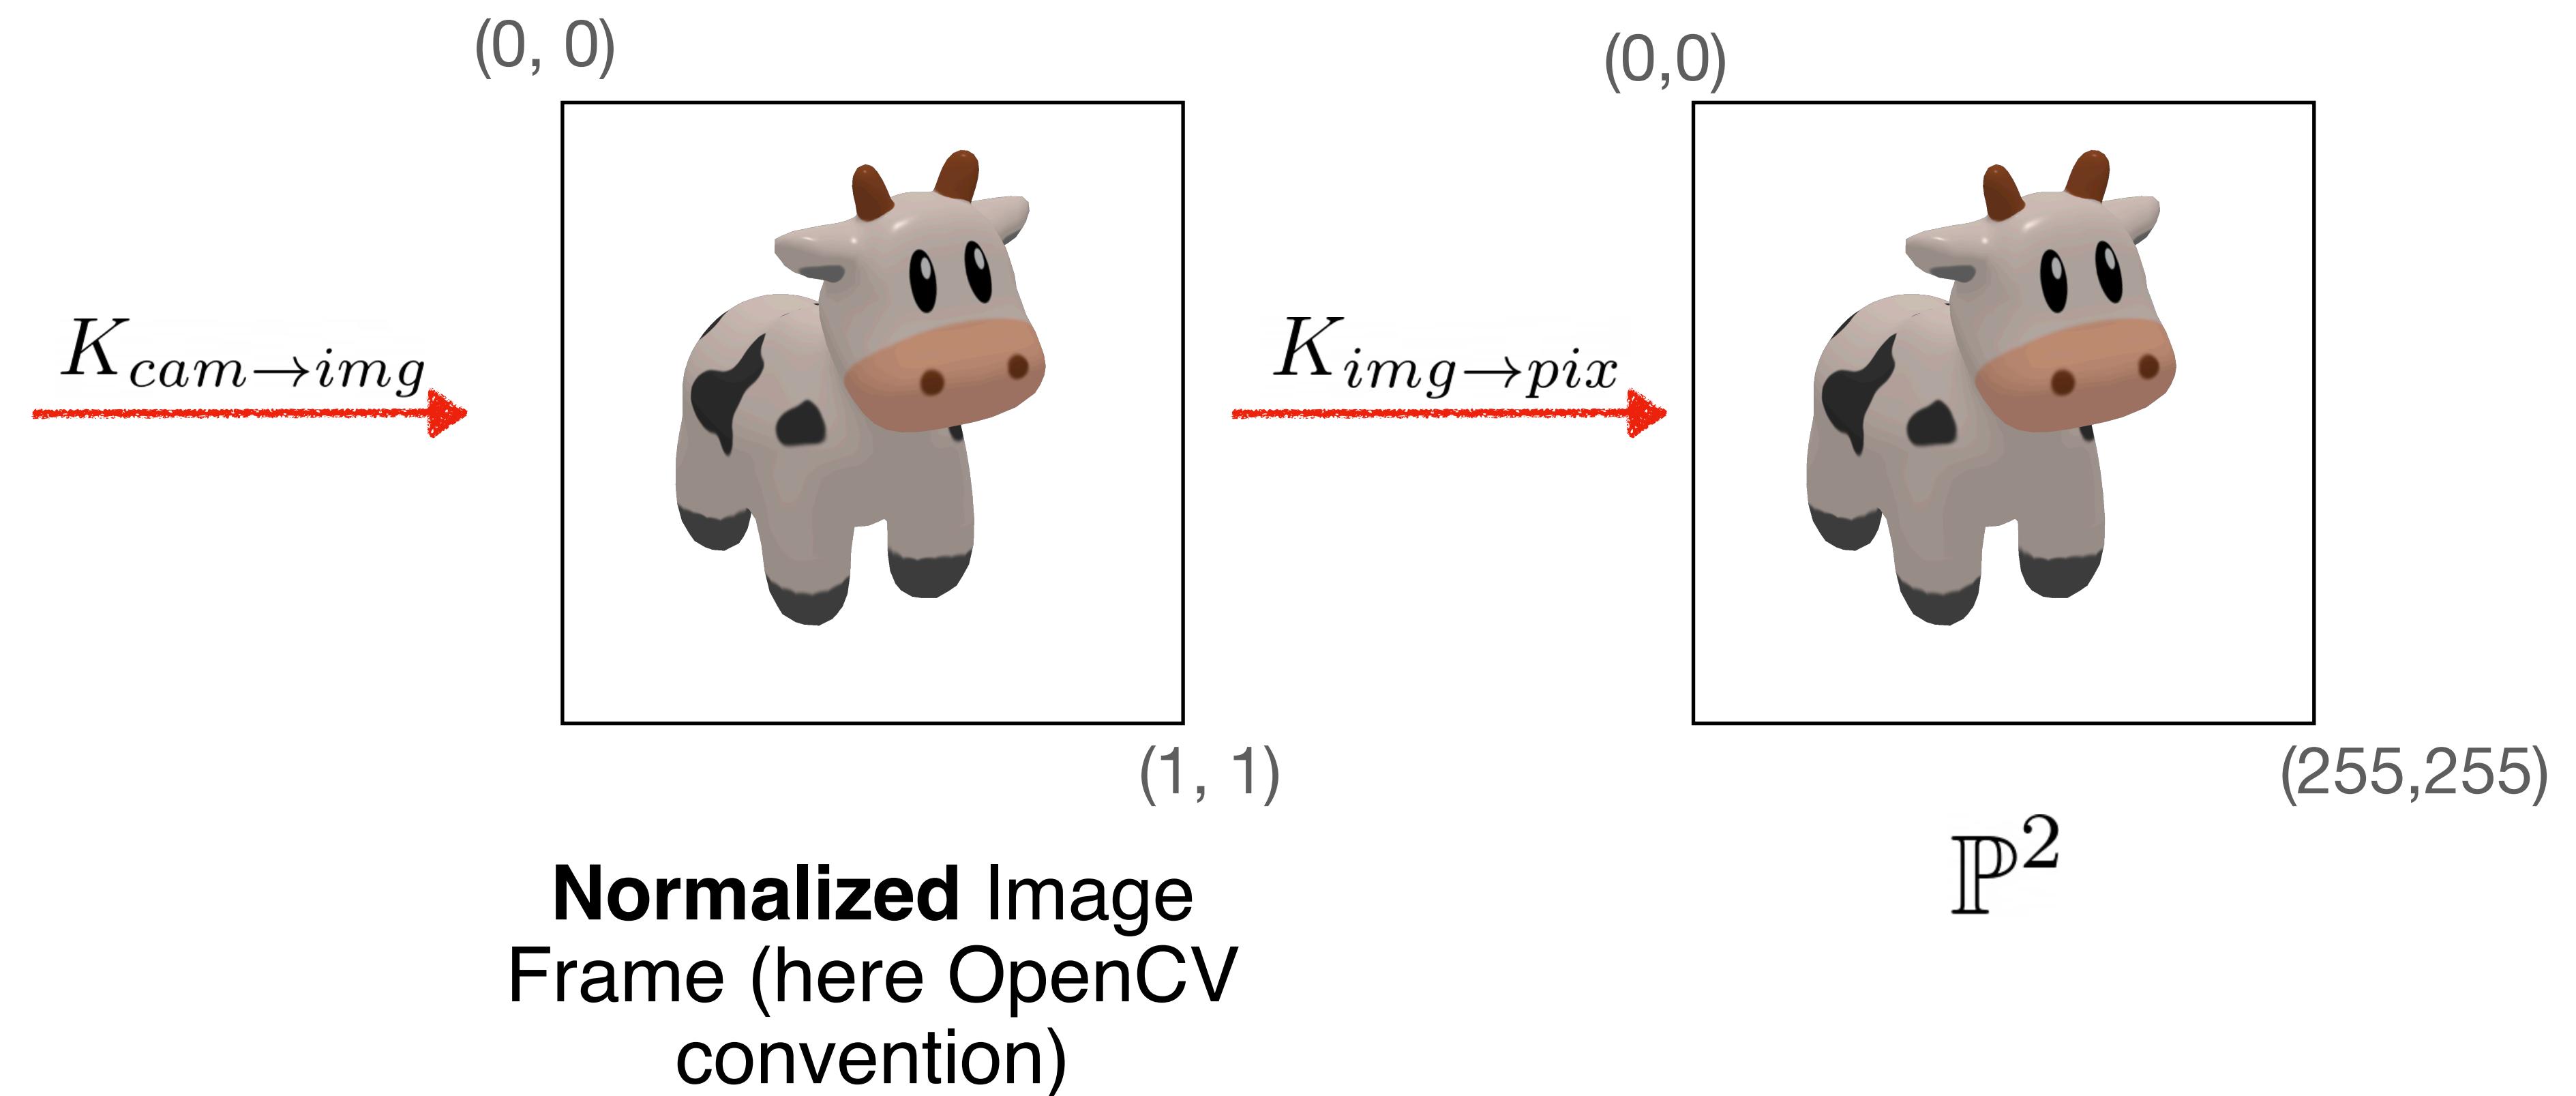


In practice: Decoupling Projection from Image Size

$$K \equiv K_{img \rightarrow pix} \ K_{cam \rightarrow img}$$

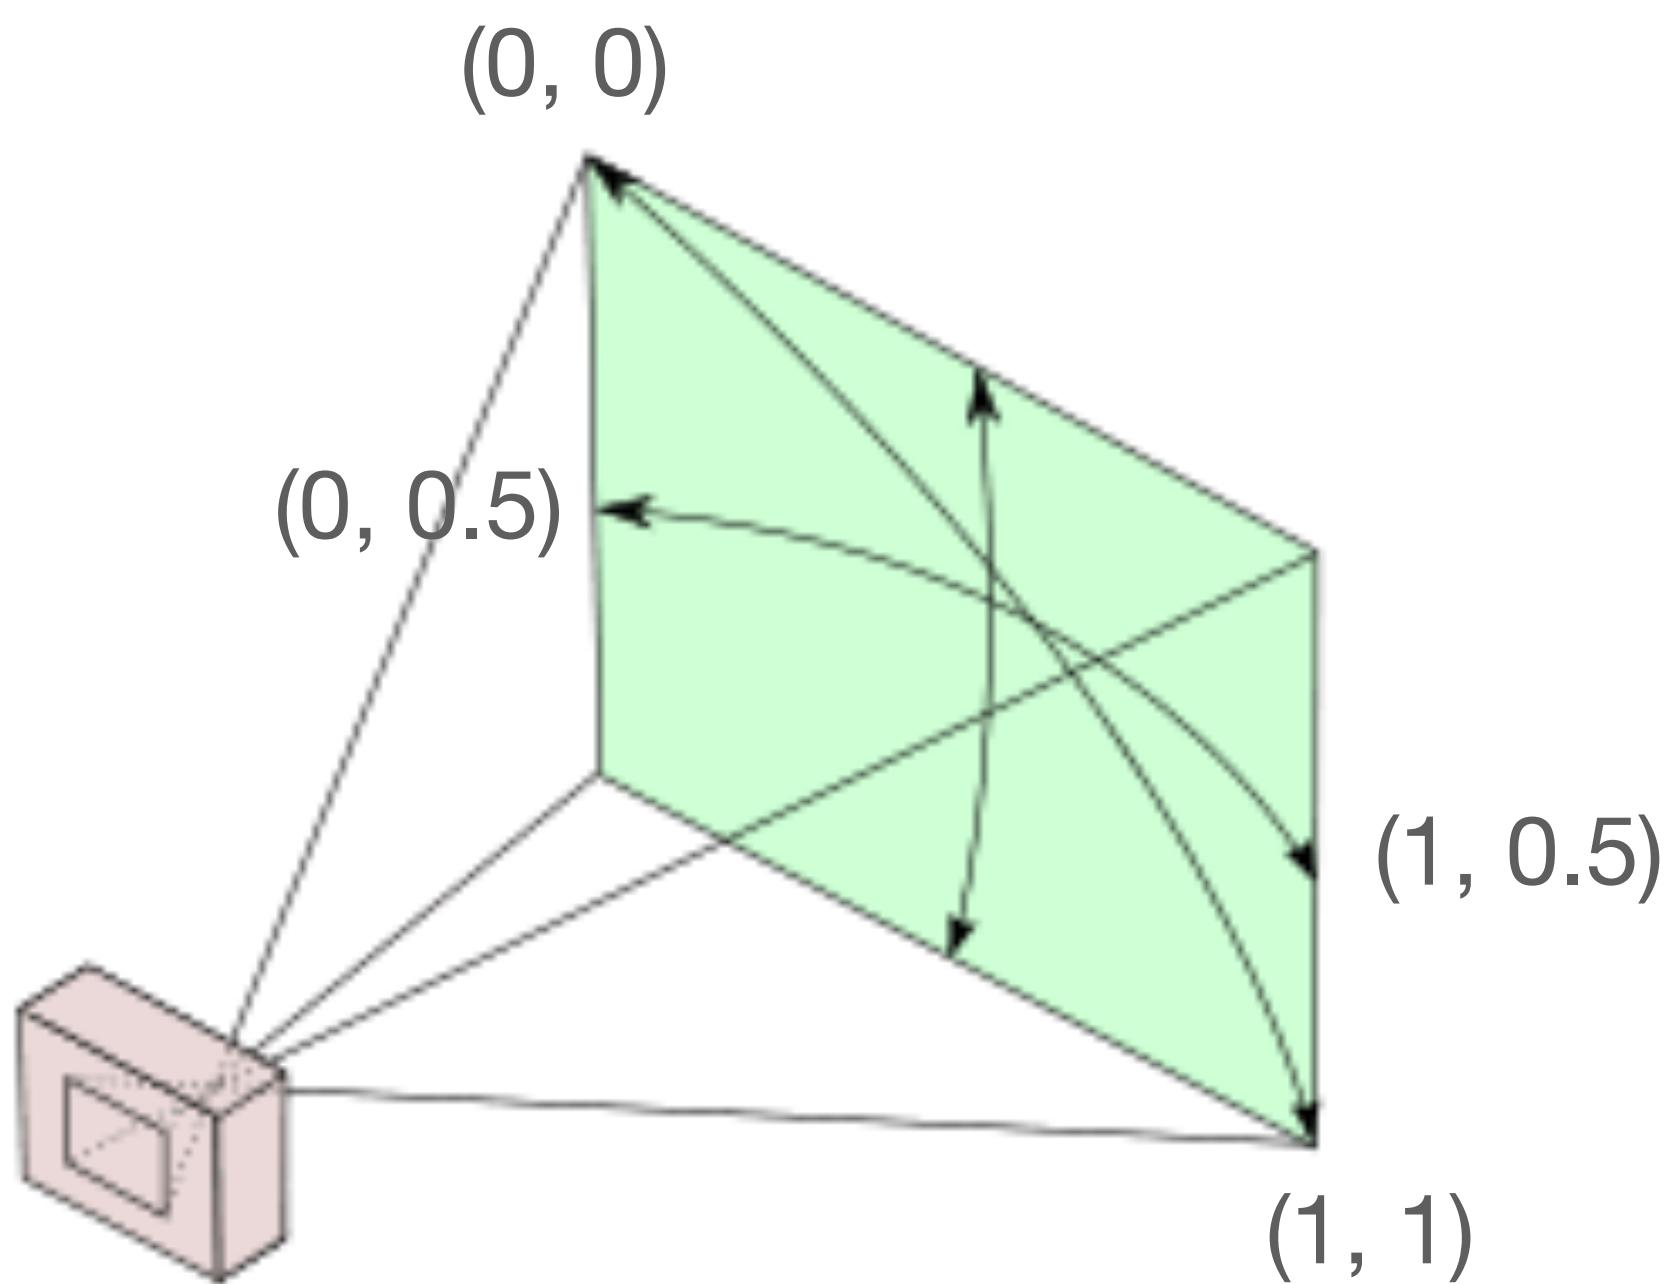


\mathbb{P}^3
3D in Camera
Frame

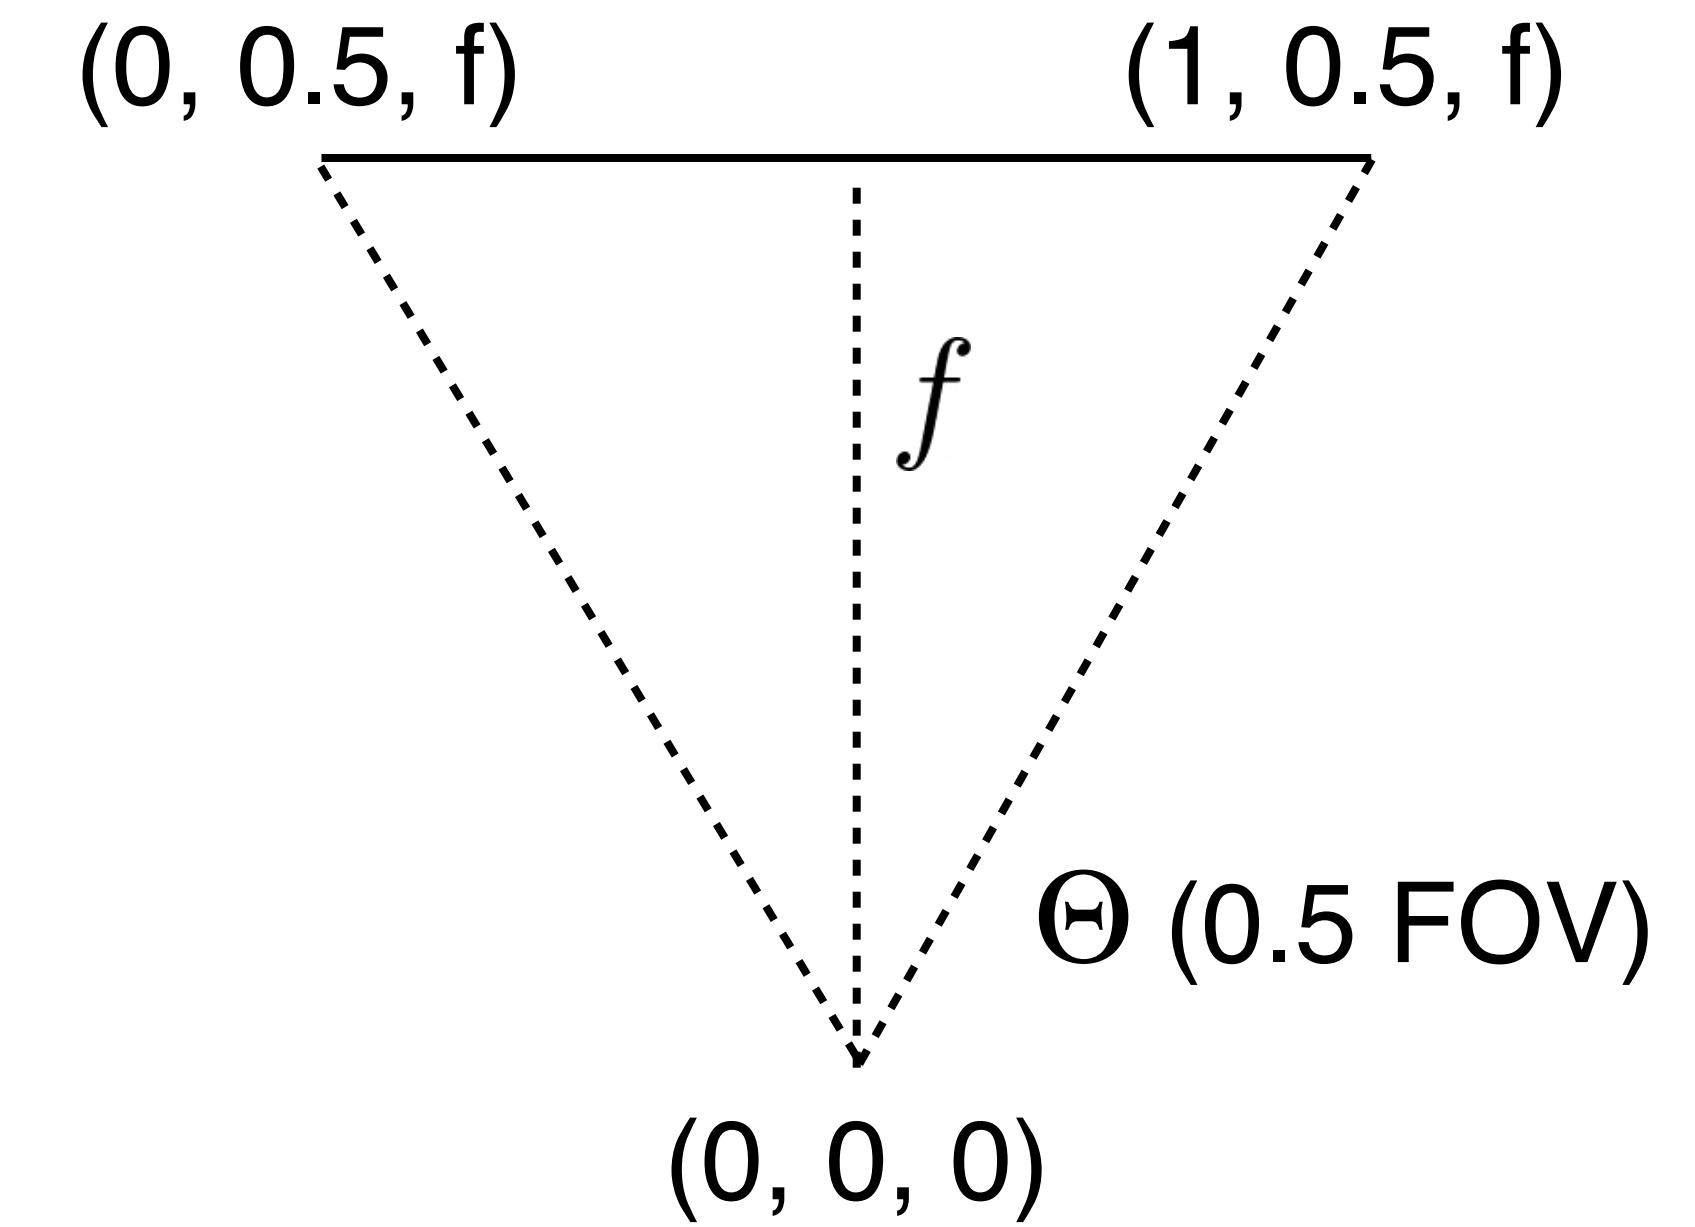


Exercise: (Horizontal) Field of view = 60 degrees

$$K_{cam \rightarrow img}$$

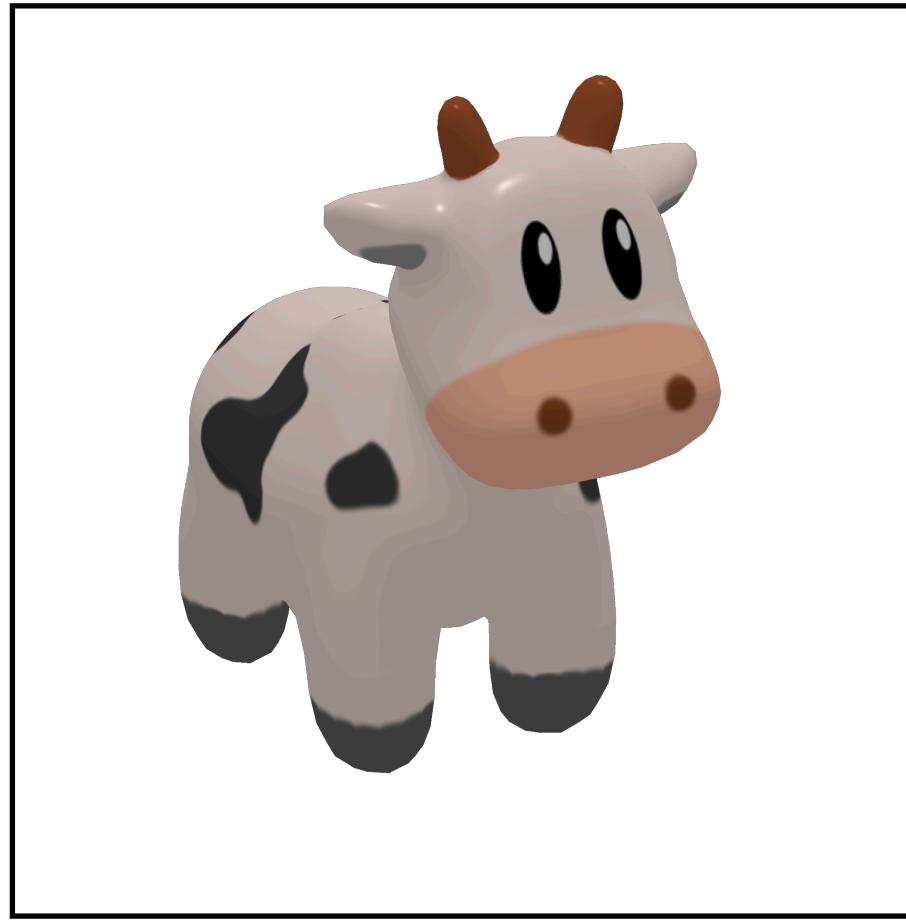


$$\begin{bmatrix} \frac{1}{2 \tan(\Theta)} & 0 & 0.5 & 0 \\ 0 & \frac{1}{2 \tan(\Theta)} & 0.5 & 0 \\ 0 & 0 & 1 & 0 \end{bmatrix}$$



Exercise: Cropping an Image

$K_{cam \rightarrow img}$

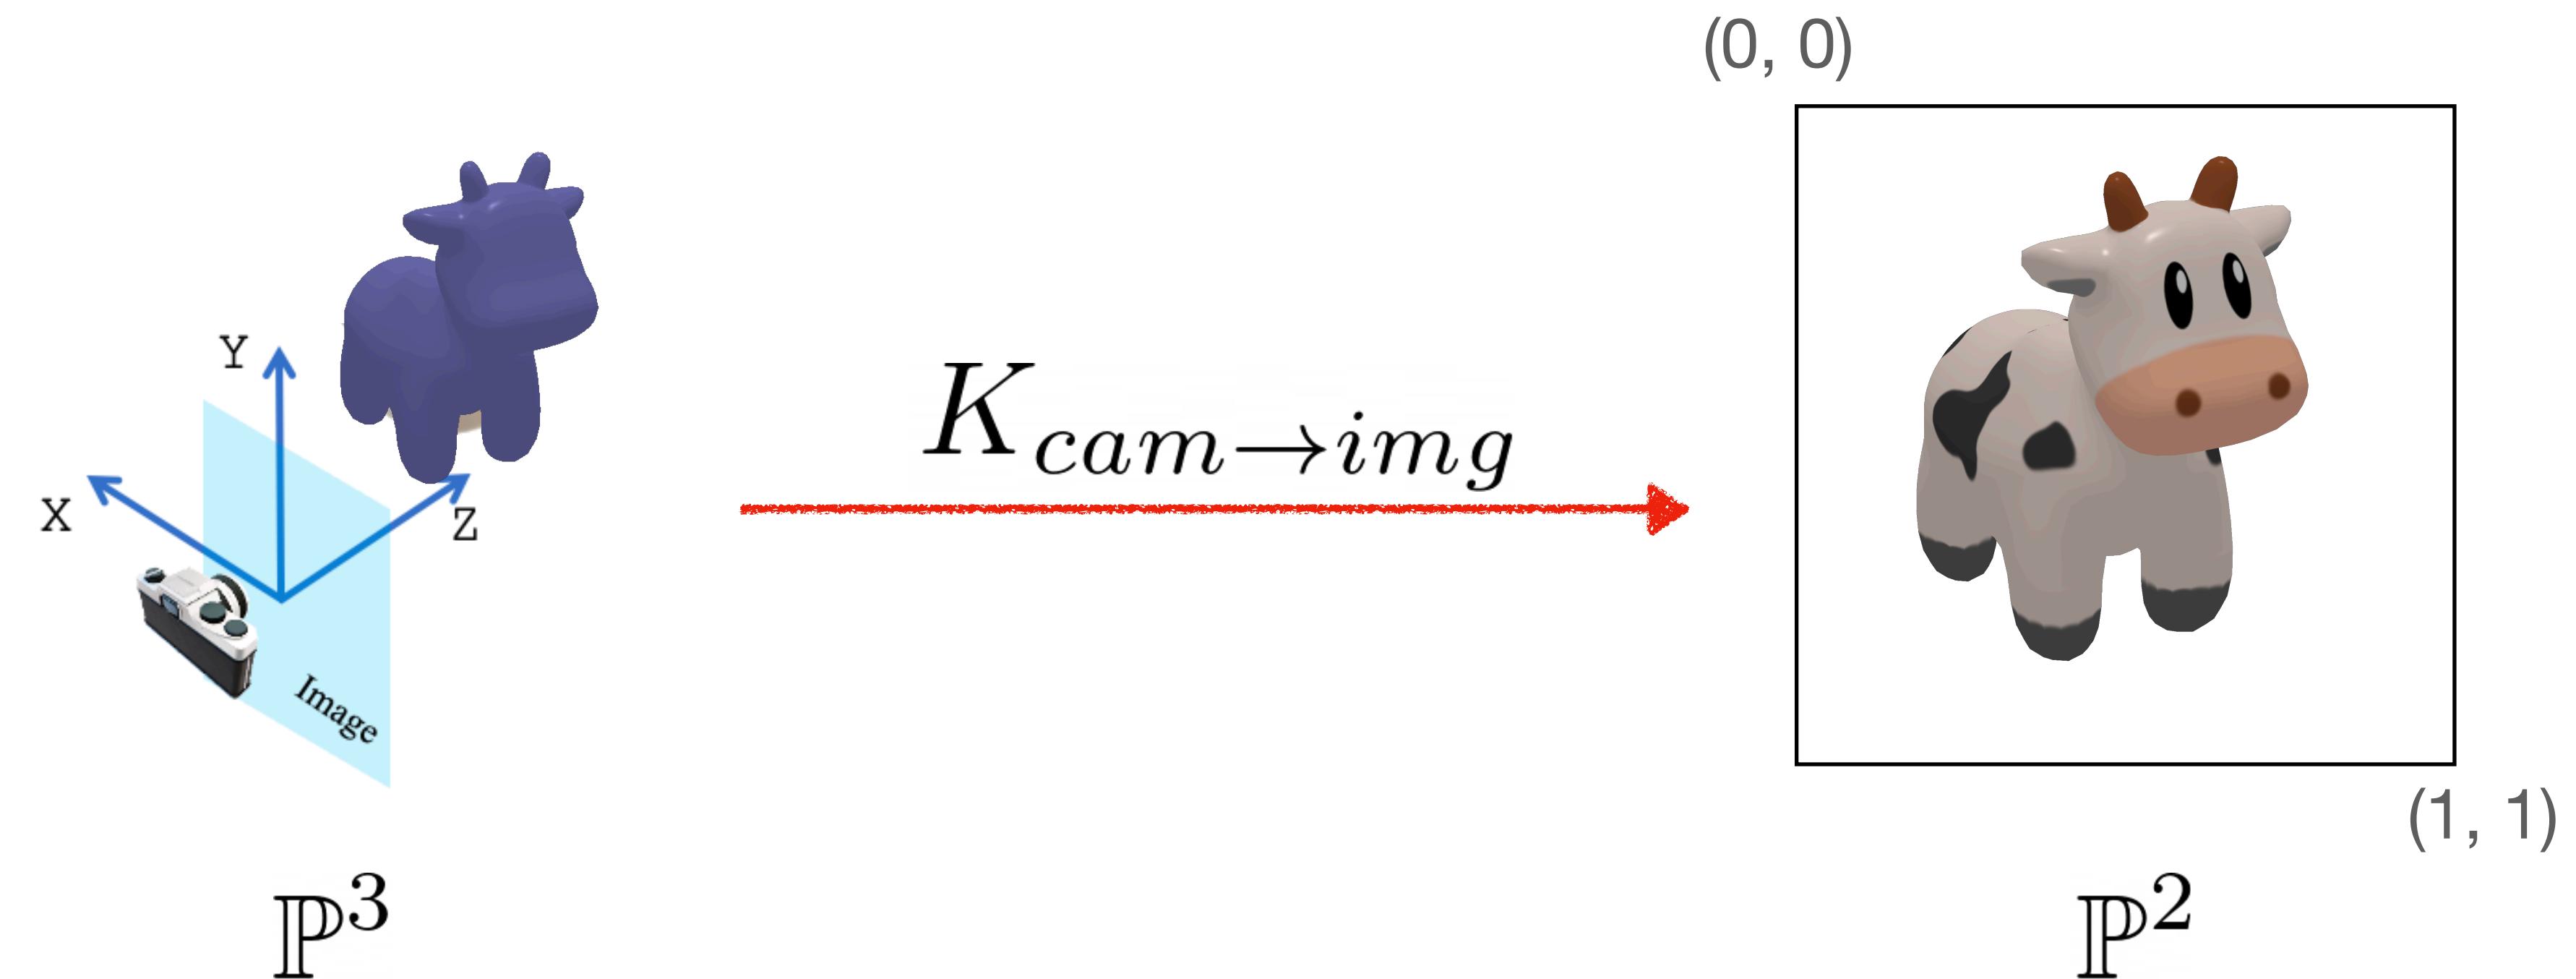


$K'_{cam \rightarrow img}$



$$K'_{cam \rightarrow img} = \begin{bmatrix} 2 & 0 & 0 \\ 0 & 2 & 0 \\ 0 & 0 & 1 \end{bmatrix} K_{cam \rightarrow img}$$

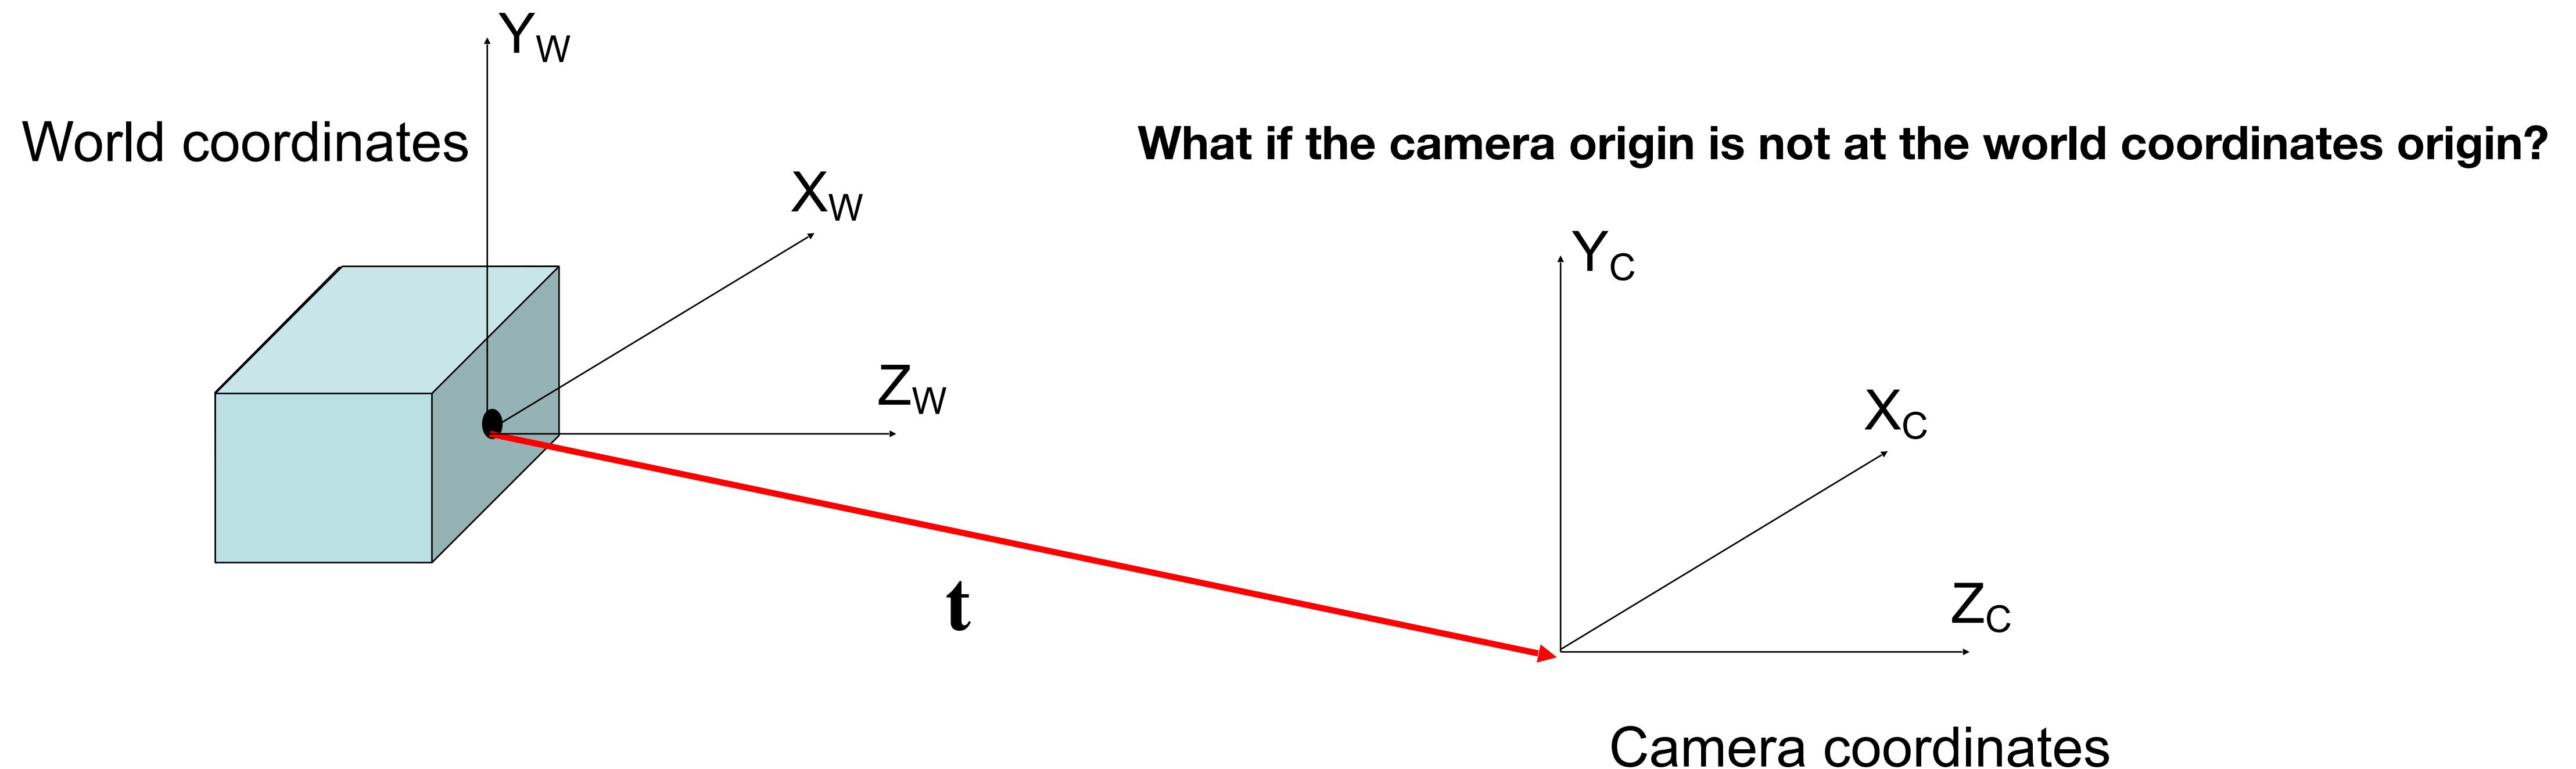
In practice: Decoupling Projection from Image Size



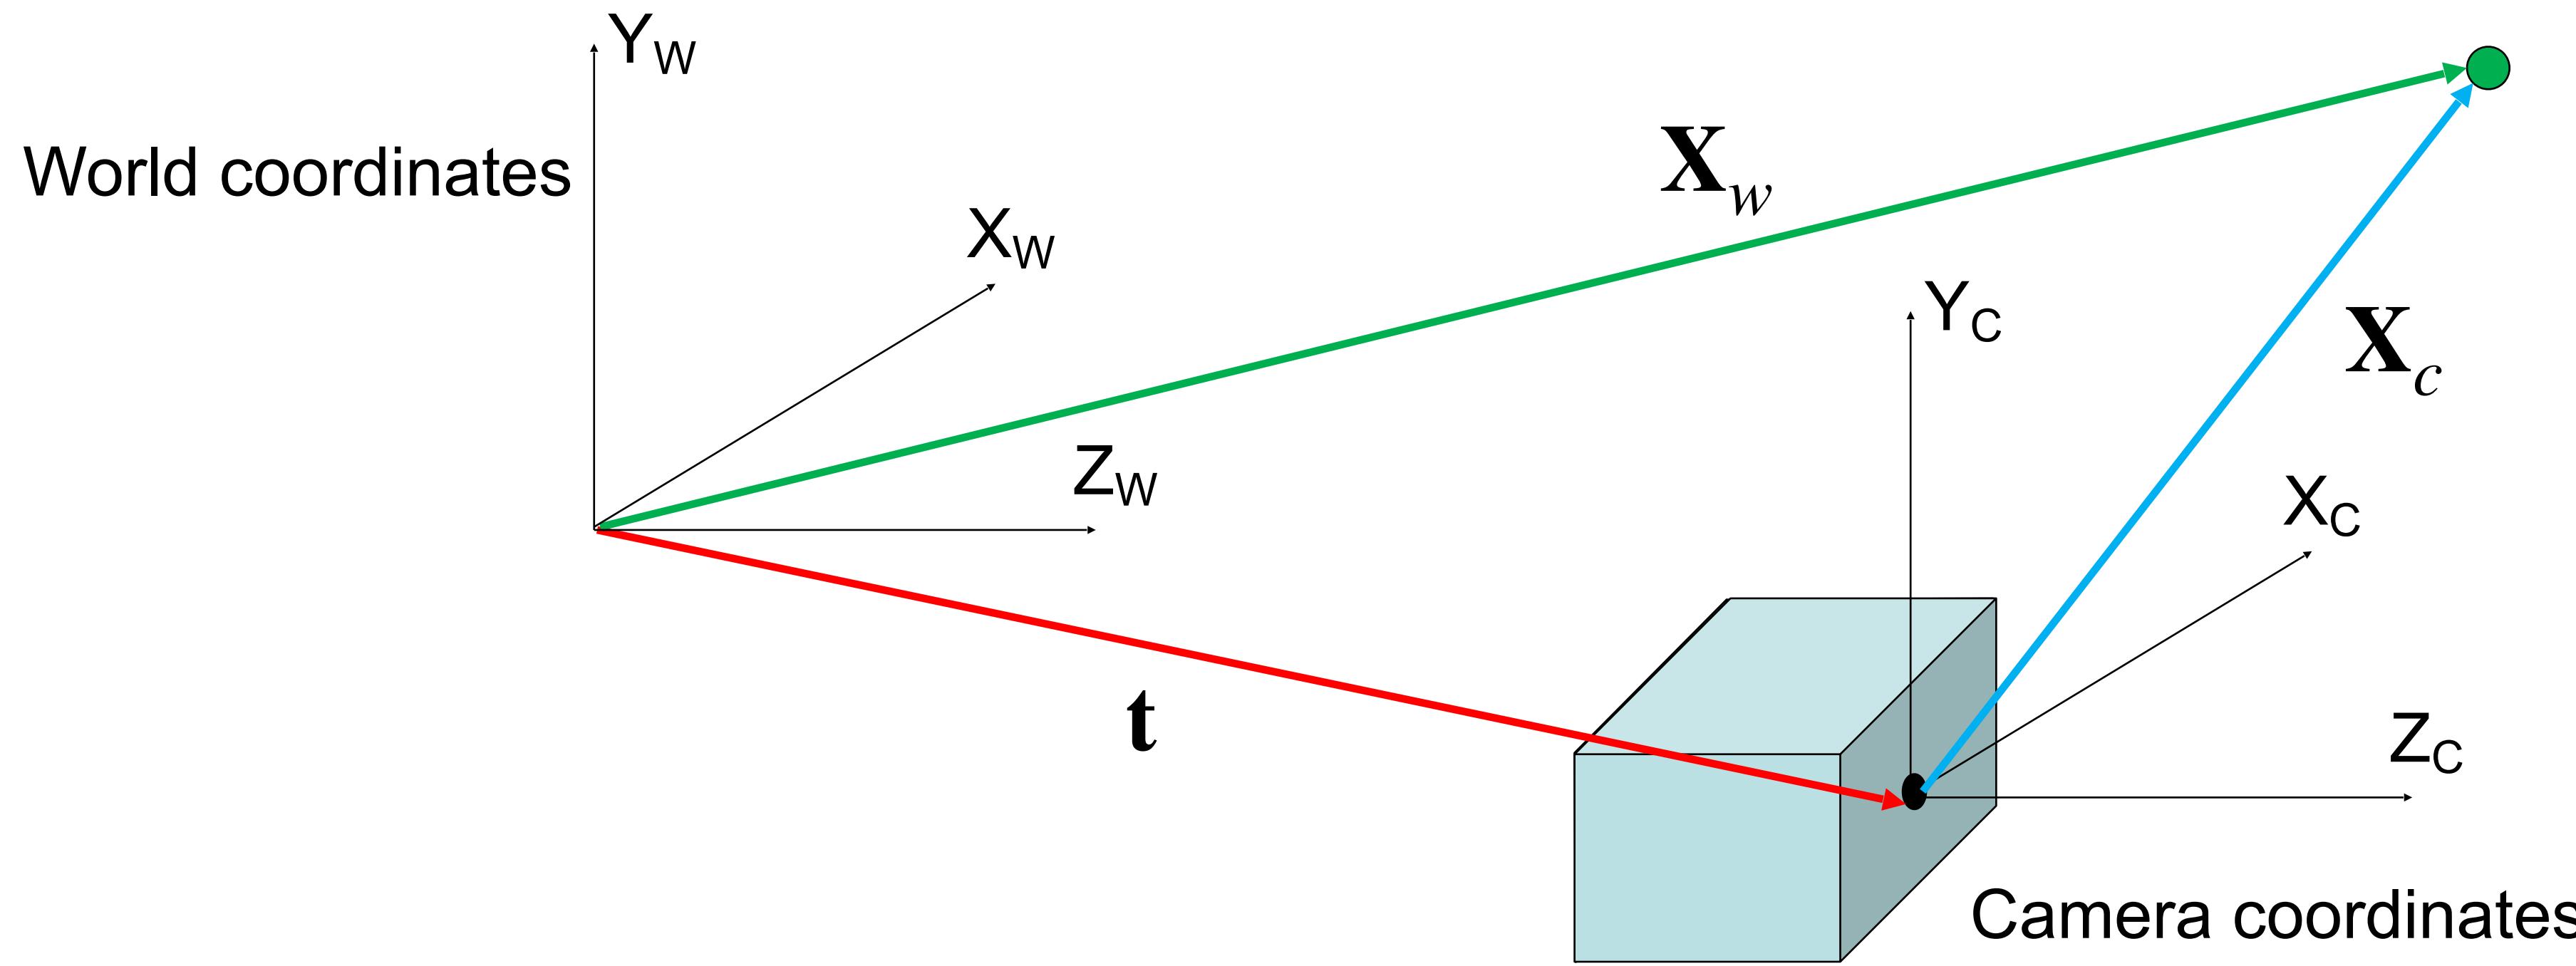
Helps (conceptually and in implementation) to reason in normalized image coordinates

Questions?

Camera parameters



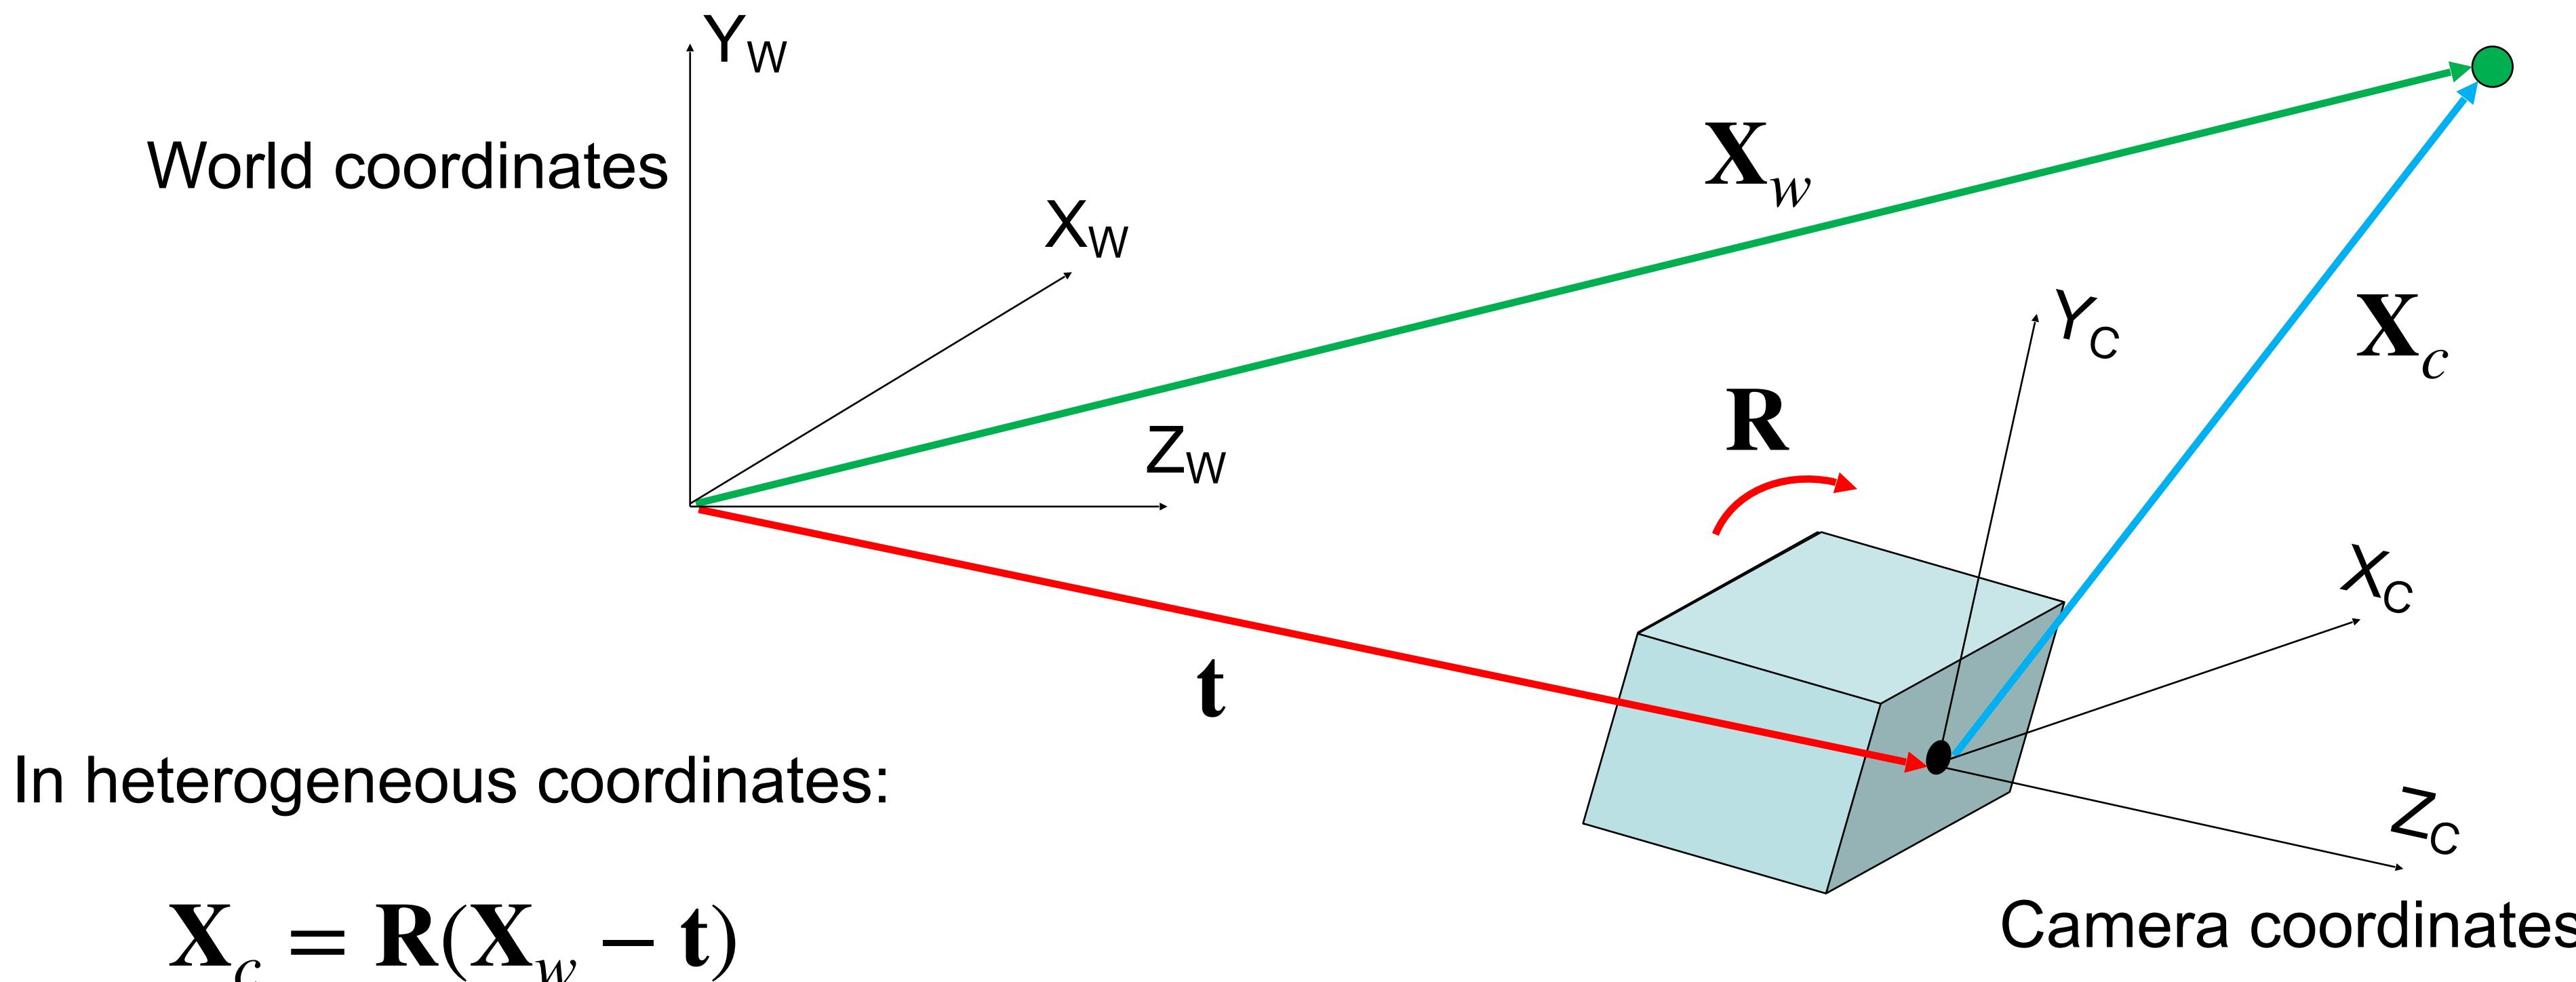
Camera parameters



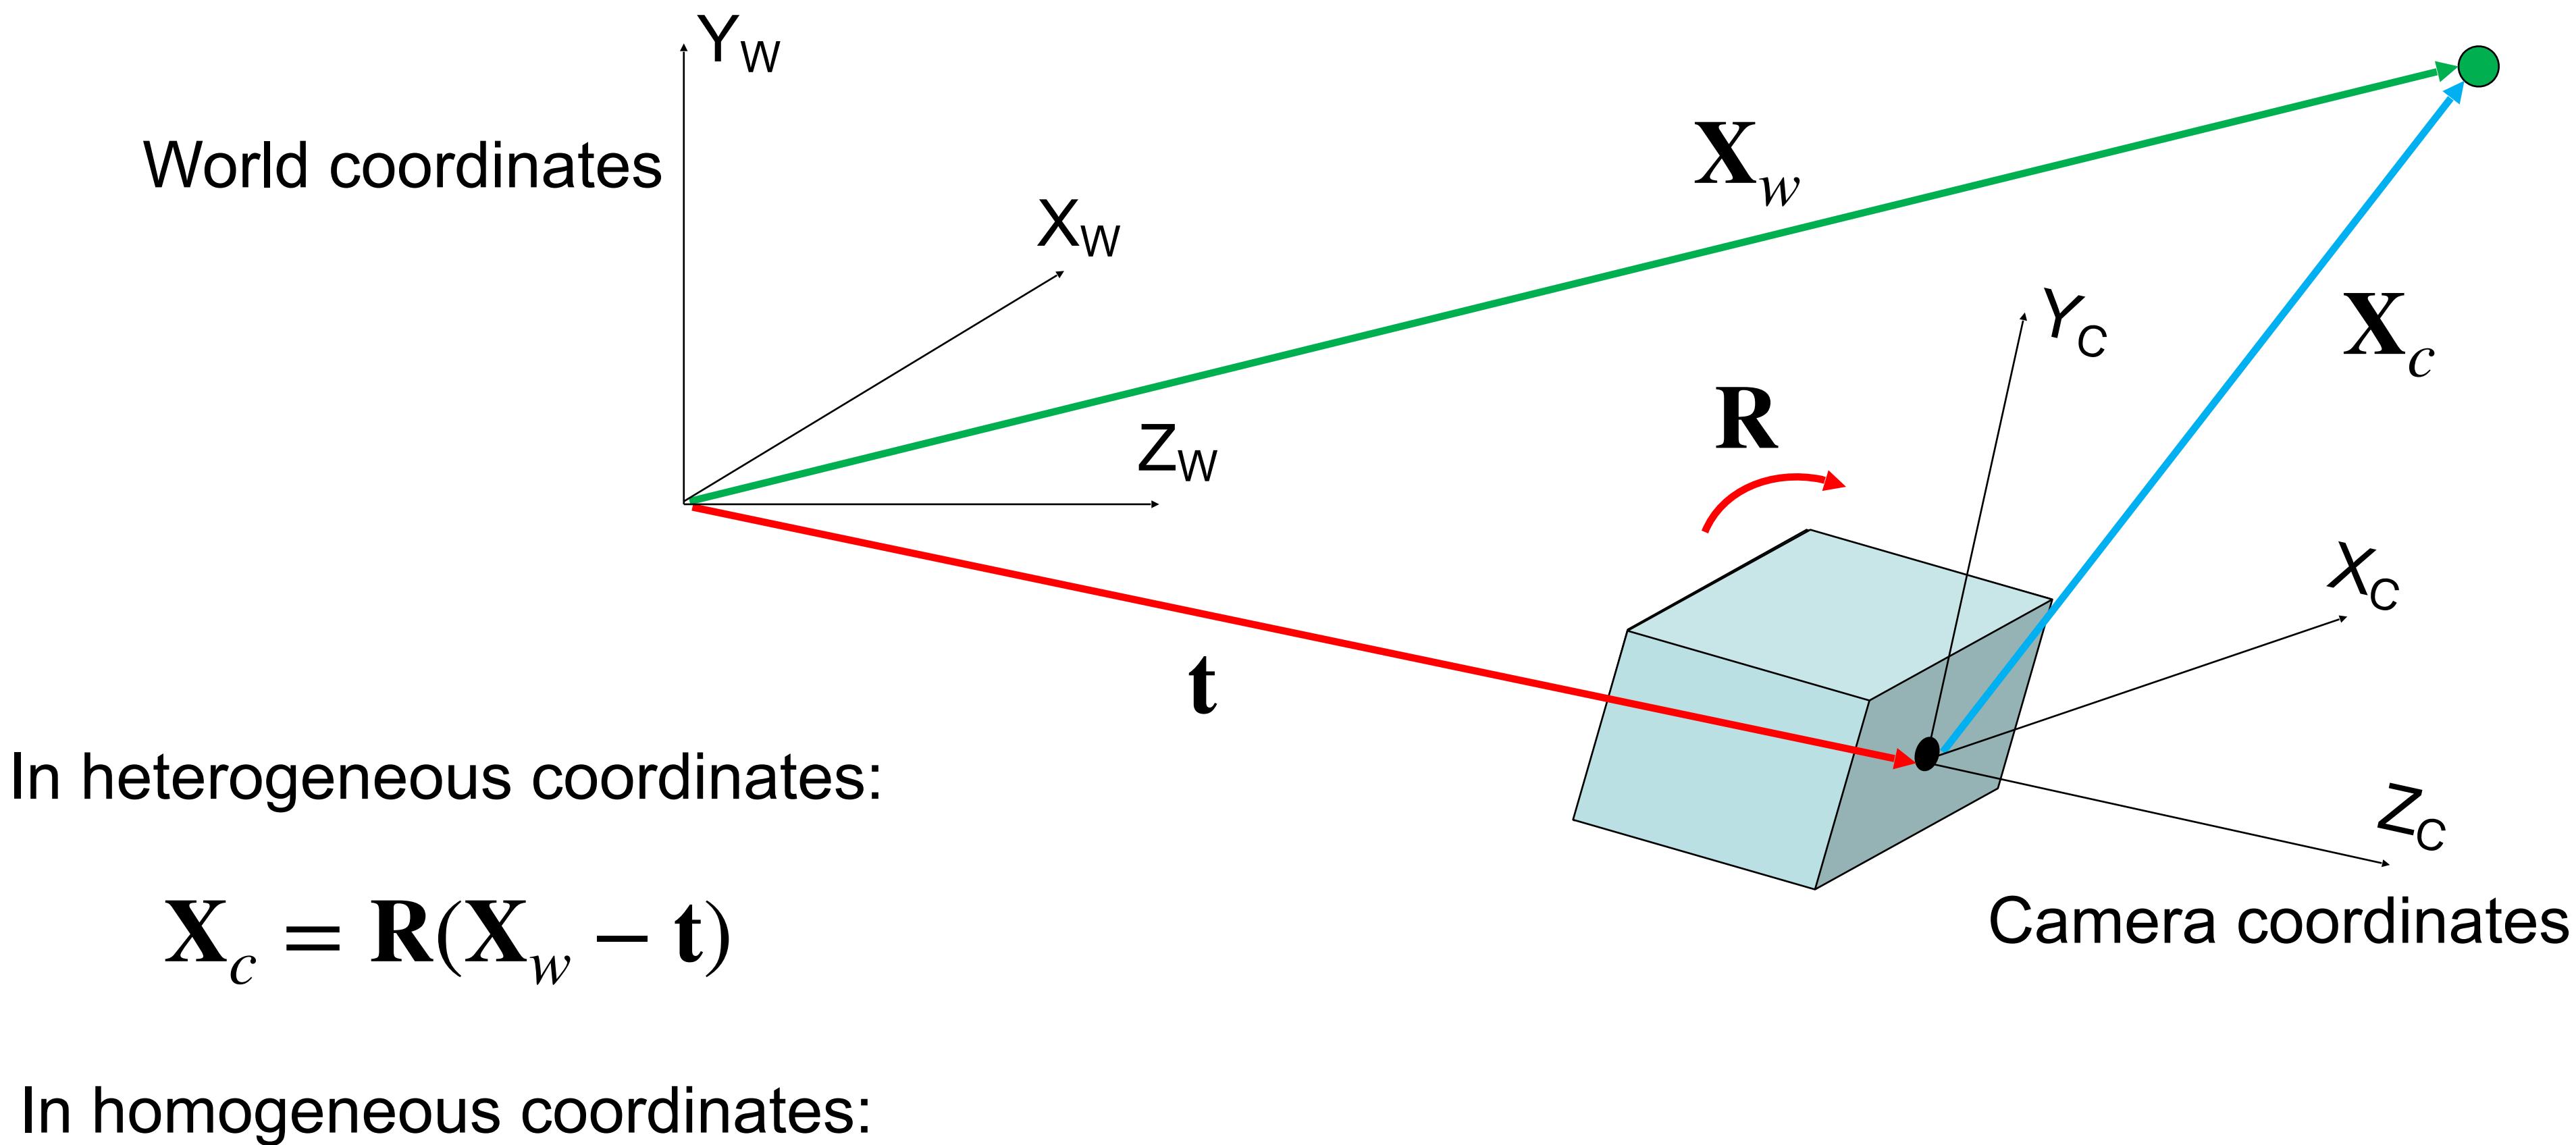
In heterogeneous coordinates:

$$\mathbf{X}_c = \mathbf{X}_w - t$$

Camera parameters

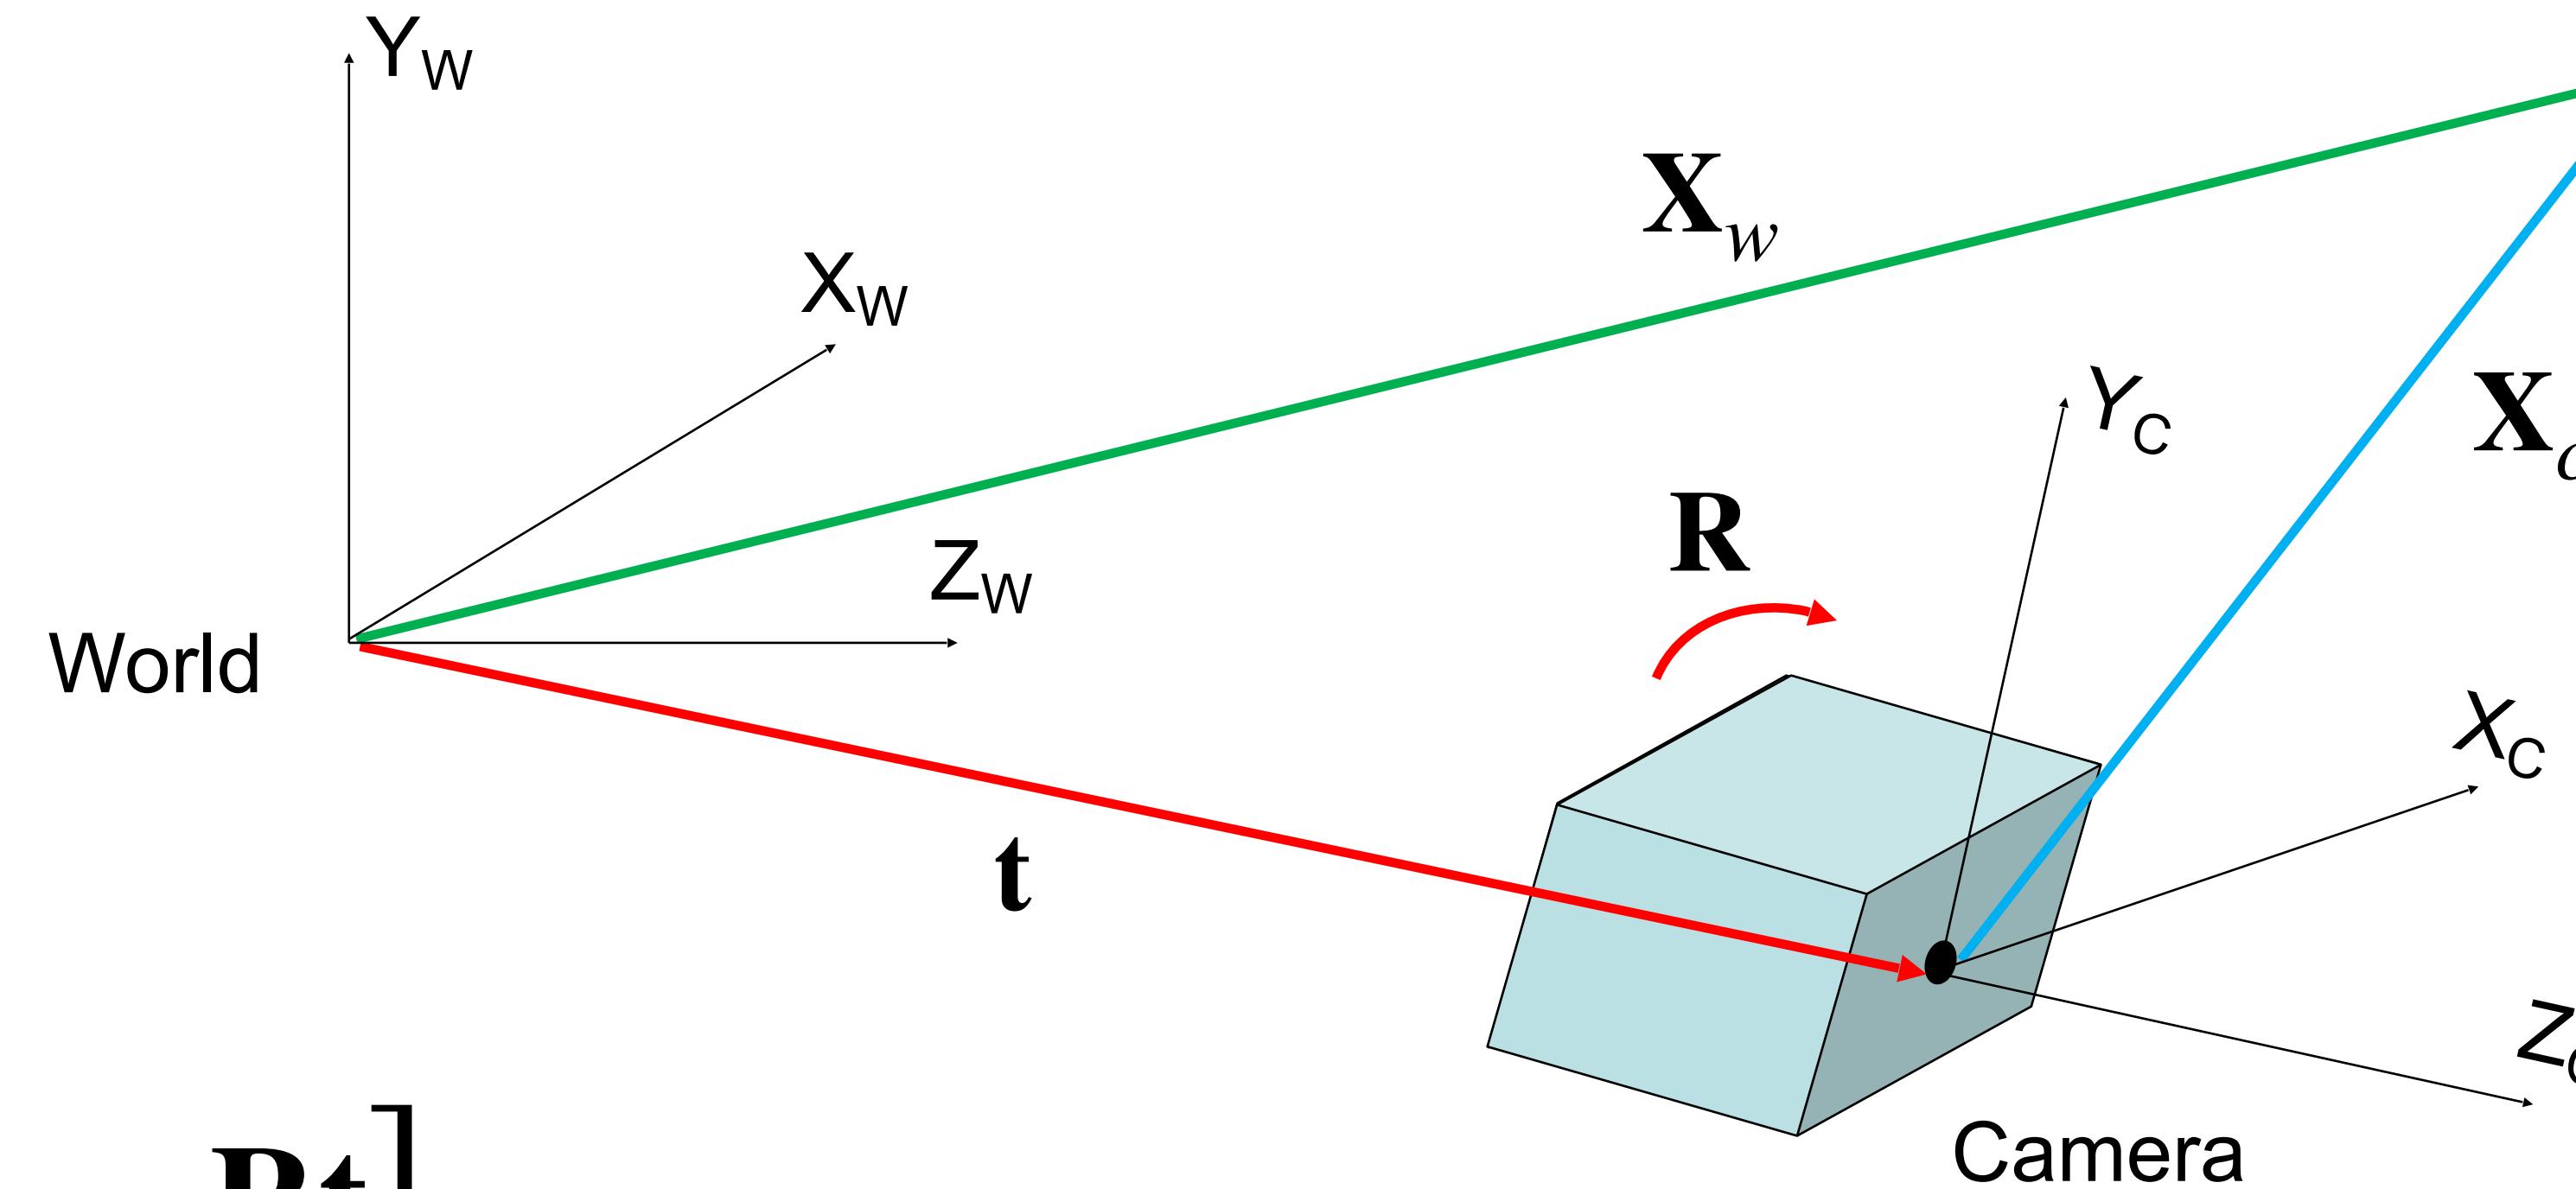


Camera parameters



$$\tilde{\mathbf{X}}_c = \begin{bmatrix} \mathbf{R} & -\mathbf{R}\mathbf{t} \\ \mathbf{0} & 1 \end{bmatrix} \tilde{\mathbf{X}}_w$$

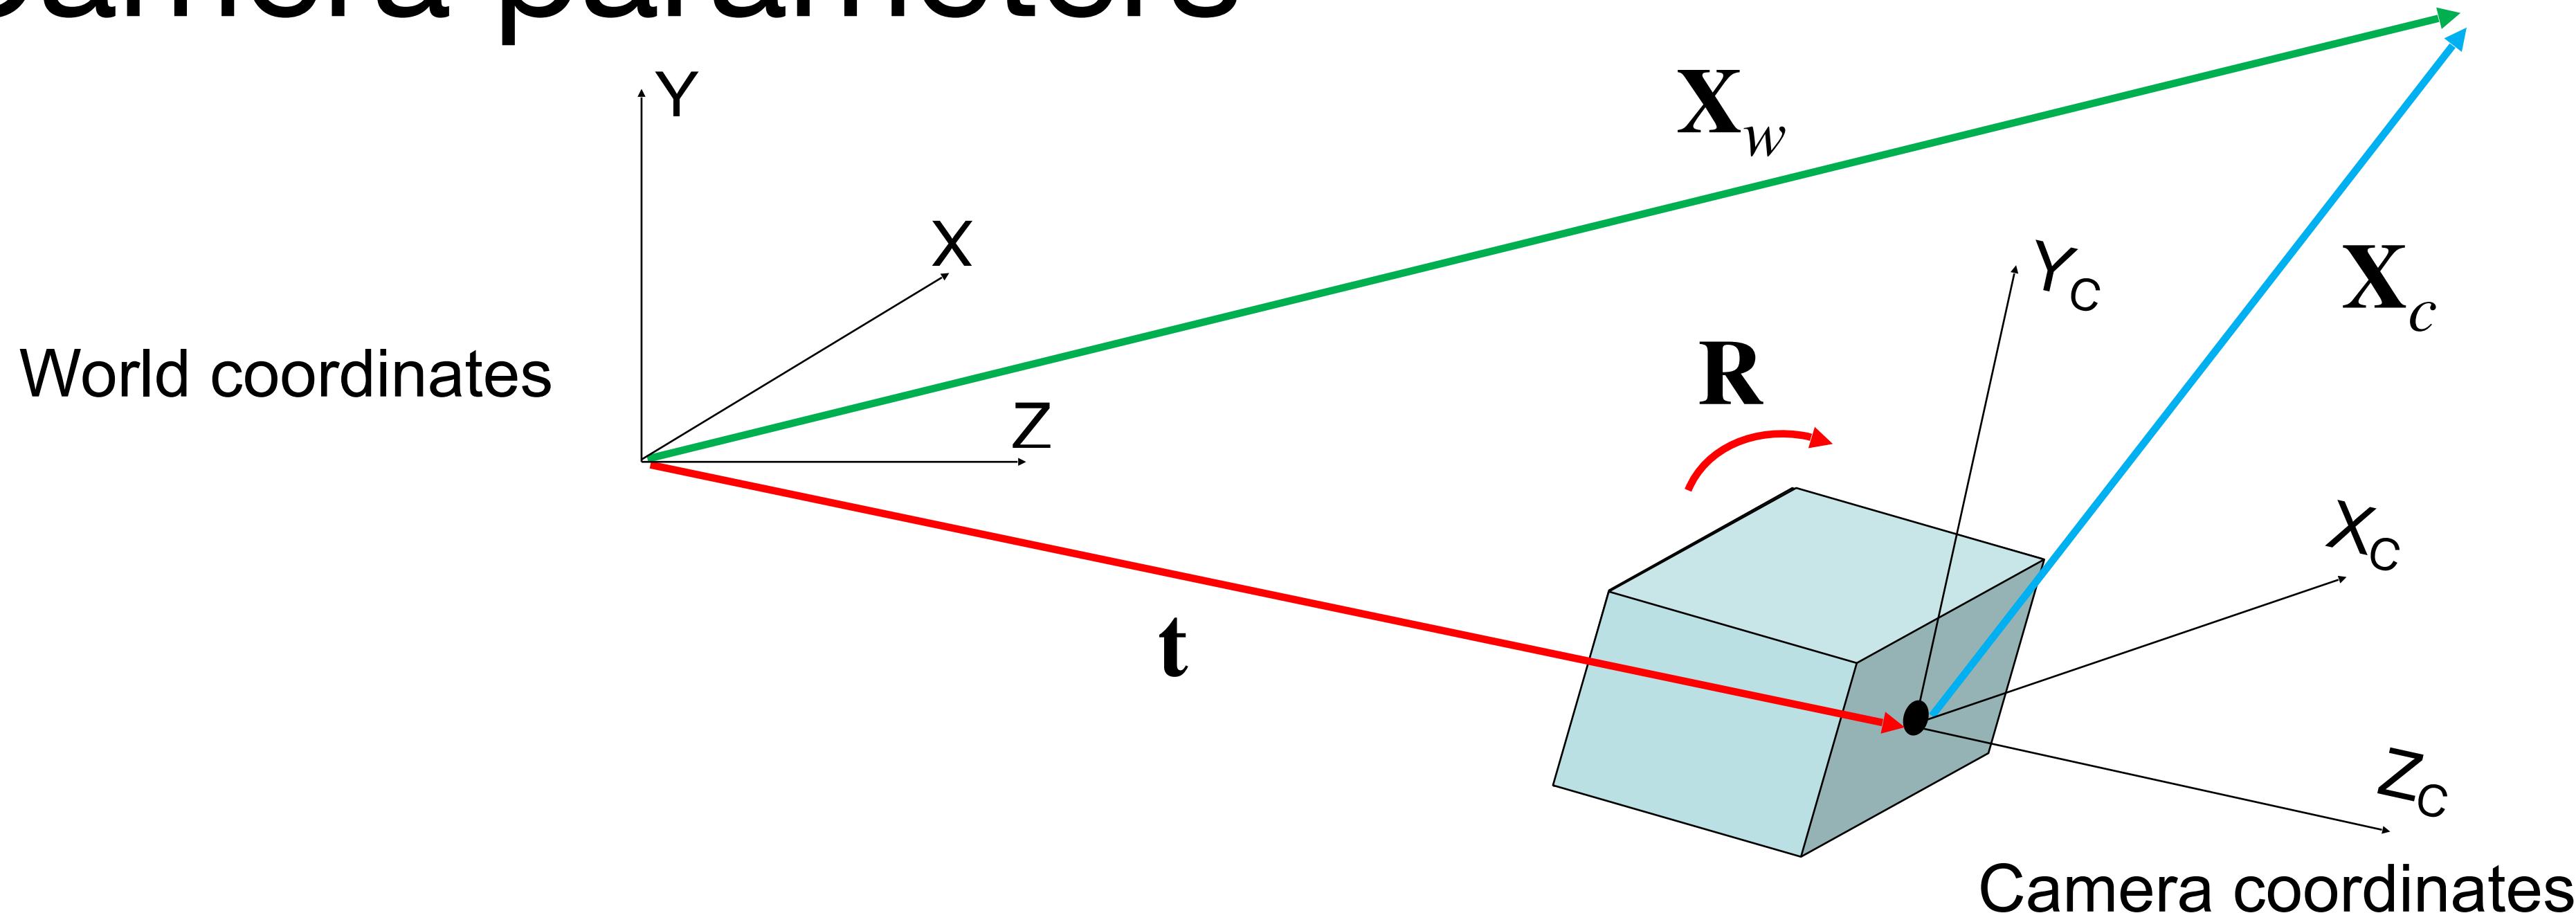
Cam2World vs. World2Cam extrinsic parameters



$$\tilde{\mathbf{X}}_c = \begin{bmatrix} \mathbf{R} & -\mathbf{R}\mathbf{t} \\ \mathbf{0} & 1 \end{bmatrix} \tilde{\mathbf{X}}_w = \mathbf{C}^{W2C} \tilde{\mathbf{X}}_w$$

$$\tilde{\mathbf{X}}_w = \begin{bmatrix} \mathbf{R}^T & \mathbf{t} \\ \mathbf{0} & 1 \end{bmatrix} \tilde{\mathbf{X}}_c = (\mathbf{C}^{W2C})^{-1} \tilde{\mathbf{X}}_c = \mathbf{C}^{C2W} \tilde{\mathbf{X}}_c$$

Camera parameters



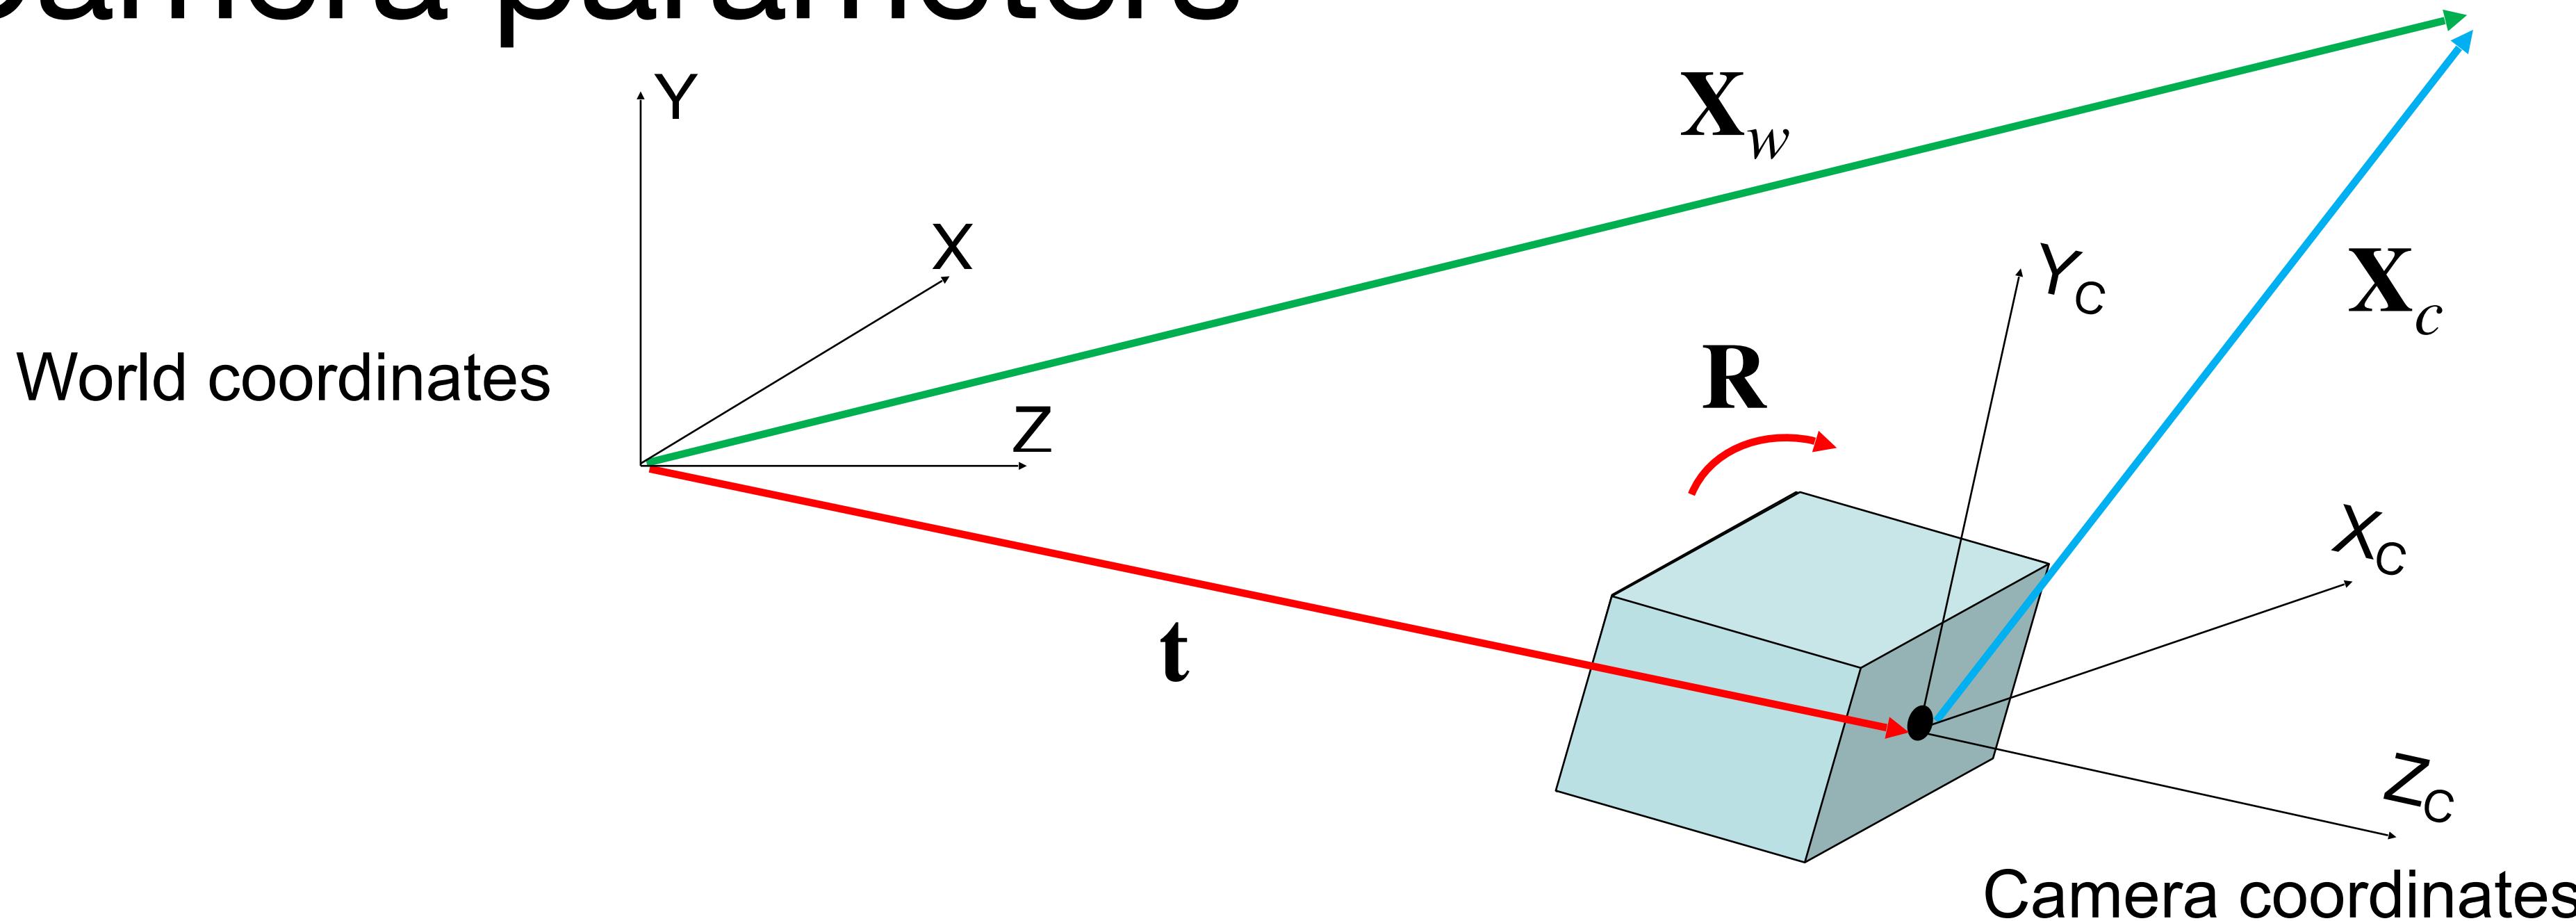
World coordinates to camera coordinates

$$\tilde{\mathbf{X}}_c = \begin{bmatrix} \mathbf{R} & -\mathbf{R}\mathbf{t} \\ \mathbf{0} & 1 \end{bmatrix} \tilde{\mathbf{X}}_w$$

Camera coordinates to image coordinates

$$\tilde{\mathbf{x}} = \begin{bmatrix} f & 0 & p_x & 0 \\ 0 & f & p_y & 0 \\ 0 & 0 & 1 & 0 \end{bmatrix} \tilde{\mathbf{X}}_c$$

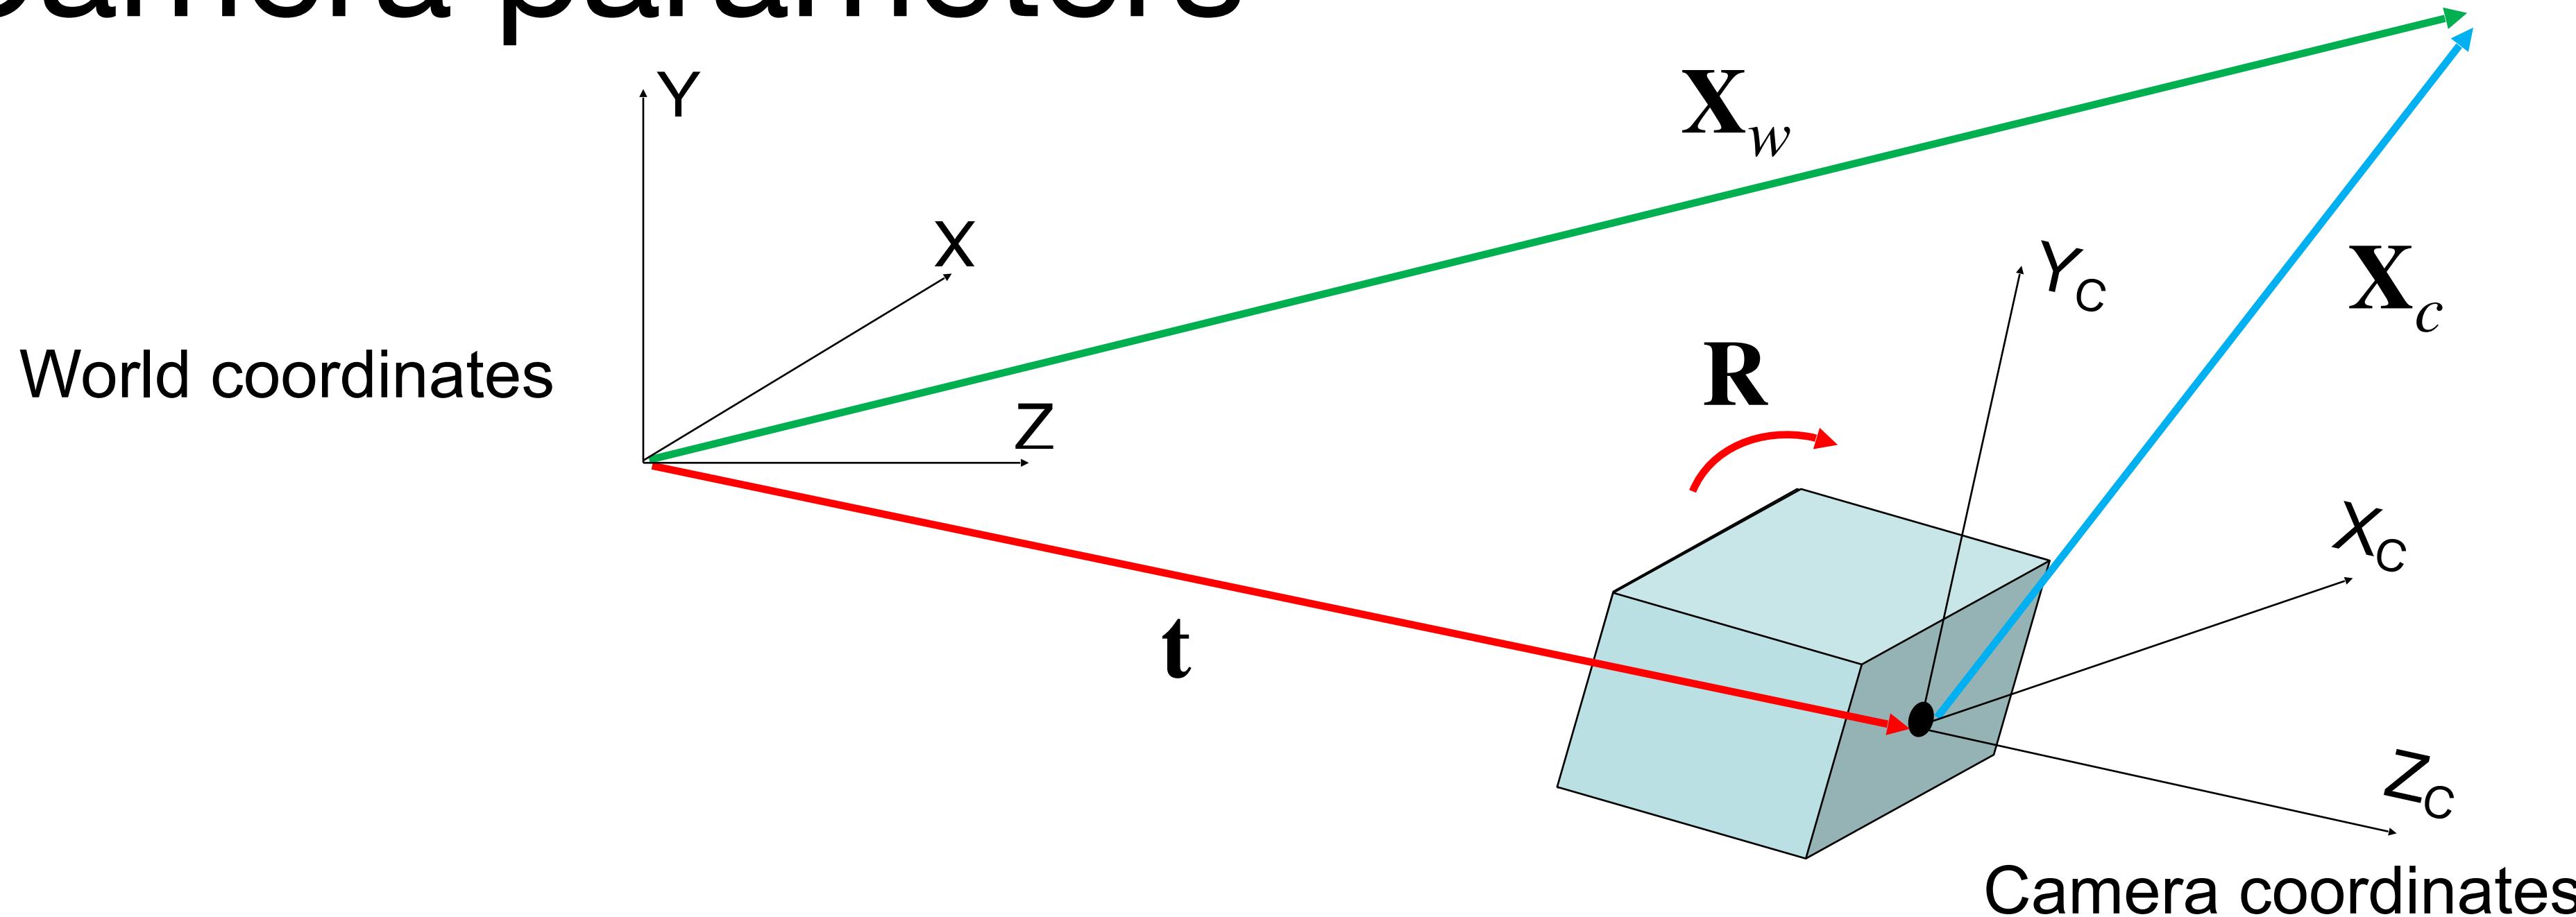
Camera parameters



World coordinates to image coordinates

$$\tilde{\mathbf{x}} = \begin{bmatrix} ? & ? & ? & ? \\ ? & ? & ? & ? \\ ? & ? & ? & ? \end{bmatrix} \tilde{\mathbf{X}}_w$$

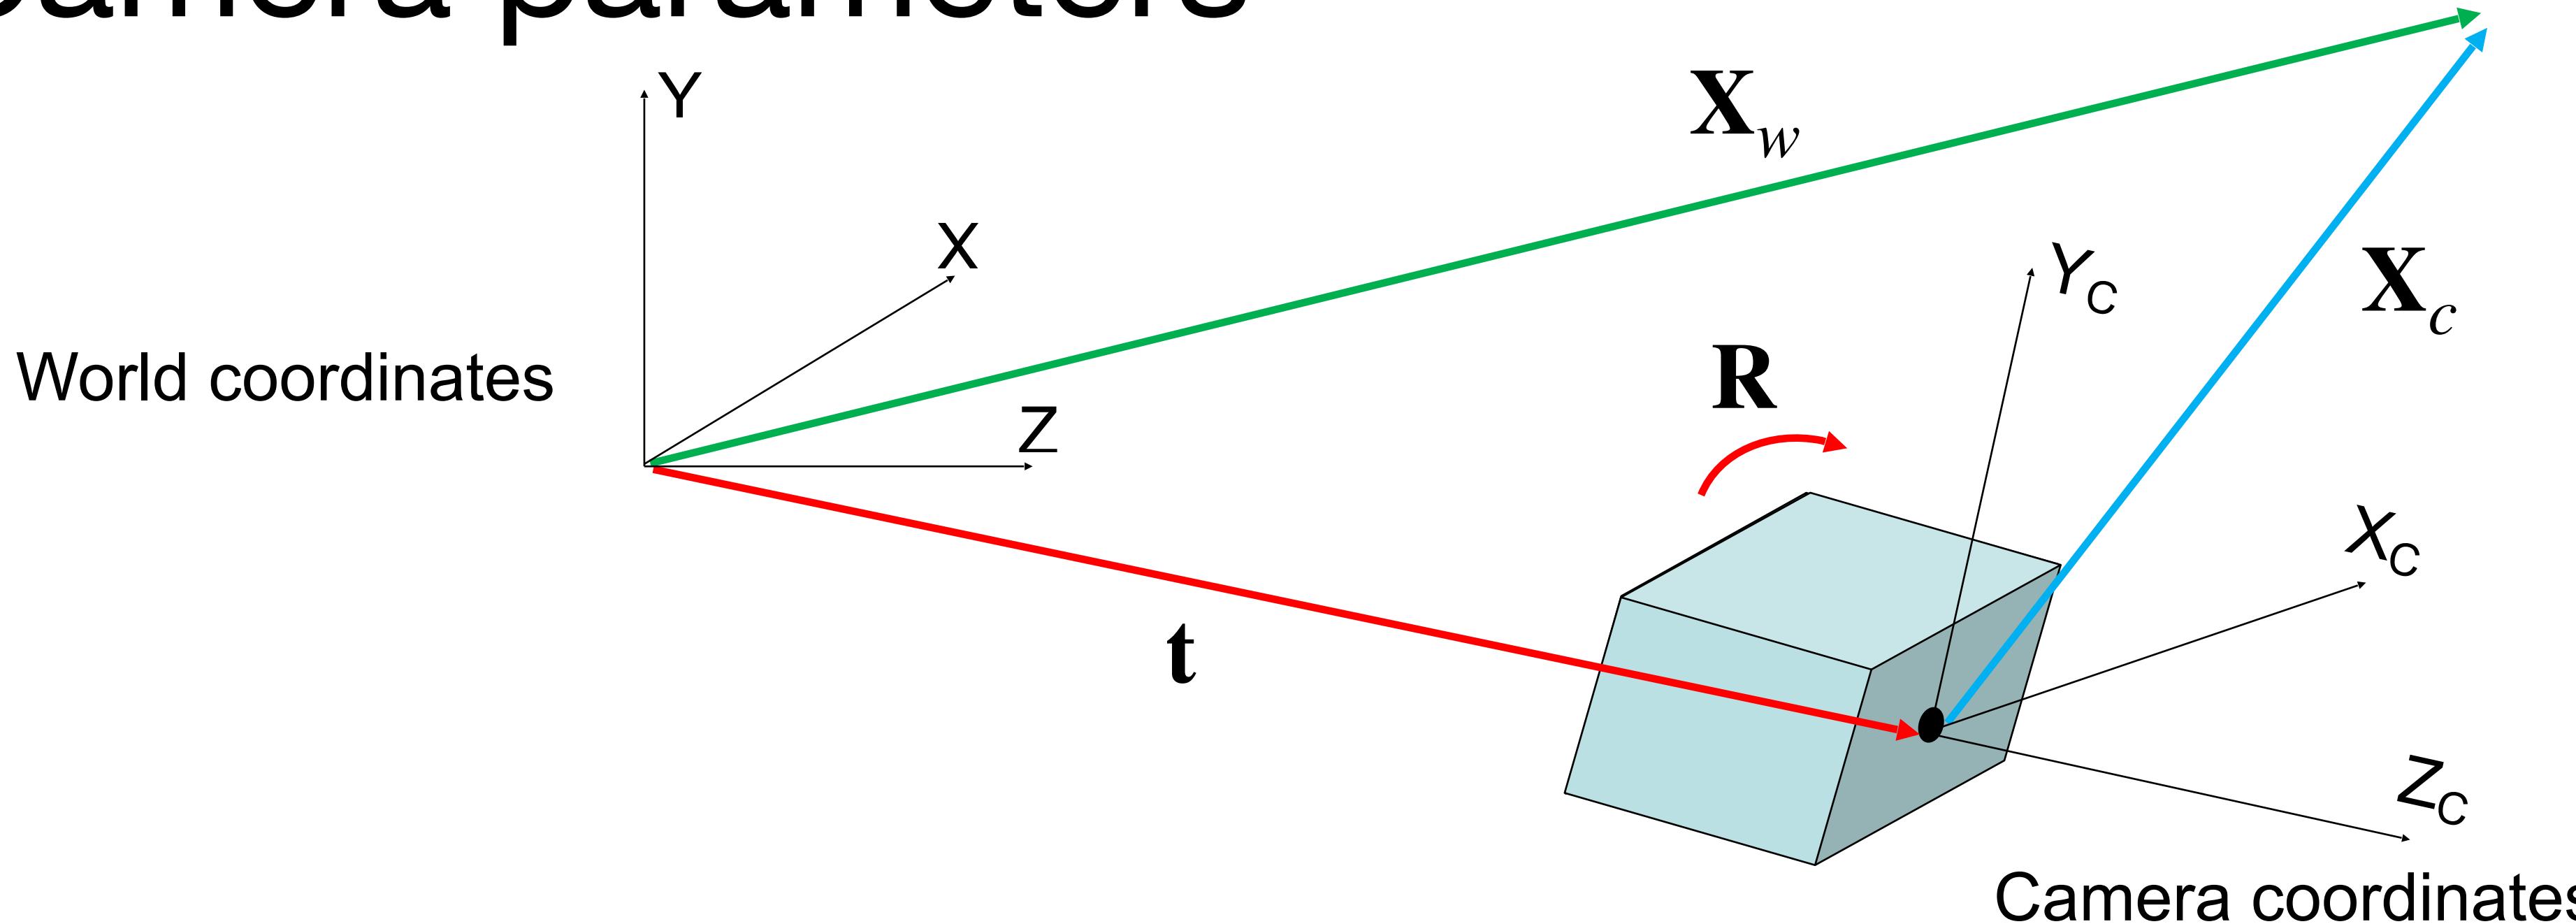
Camera parameters



World coordinates to image coordinates

$$\tilde{\mathbf{x}} = \begin{bmatrix} f & 0 & p_x & 0 \\ 0 & f & p_y & 0 \\ 0 & 0 & 1 & 0 \end{bmatrix} \begin{bmatrix} \mathbf{R} & -\mathbf{R}\mathbf{t} \\ 0 & 1 \end{bmatrix} \tilde{\mathbf{X}}_w$$

Camera parameters



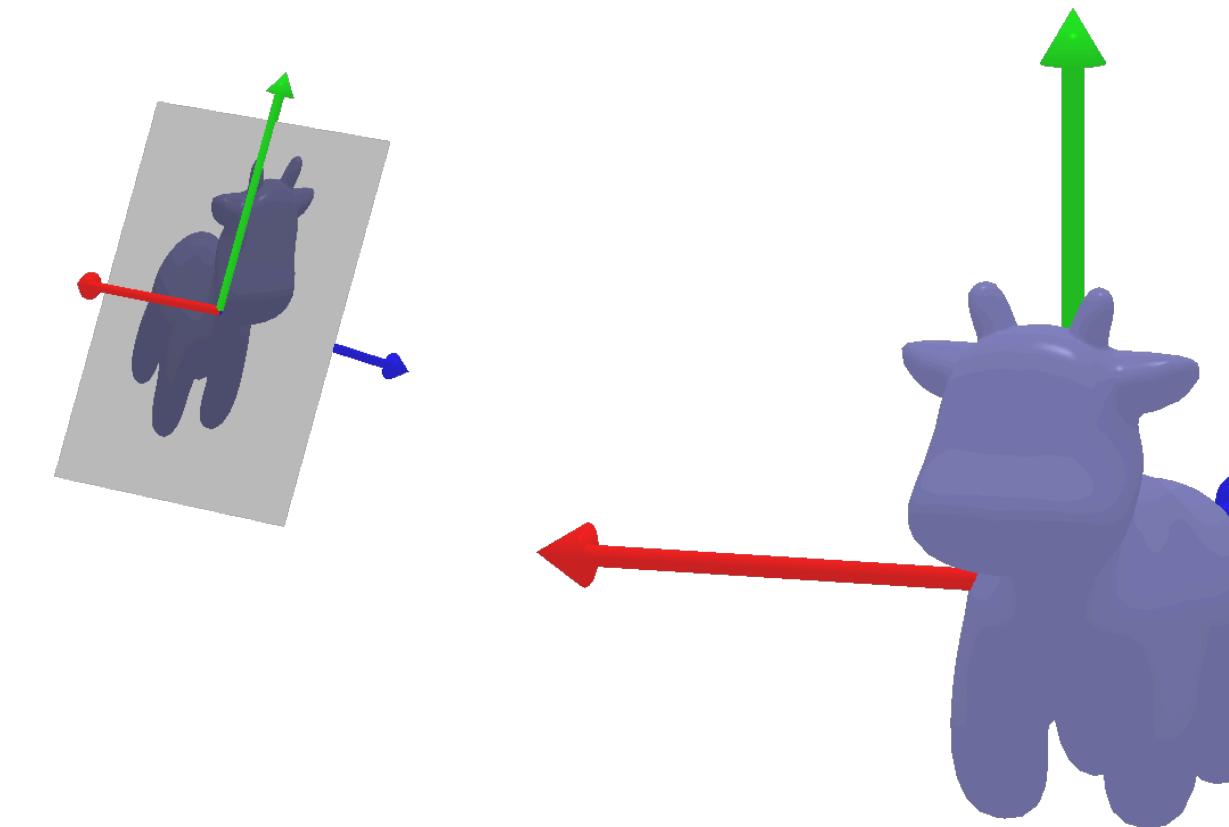
World coordinates to image coordinates

$$\tilde{\mathbf{x}} = \begin{bmatrix} f & 0 & p_x & 0 \\ 0 & f & p_y & 0 \\ 0 & 0 & 1 & 0 \end{bmatrix} \begin{bmatrix} \mathbf{R} & -\mathbf{R}\mathbf{t} \\ 0 & 1 \end{bmatrix} \tilde{\mathbf{X}}_w = \mathbf{P}\tilde{\mathbf{X}}_w$$

Intrinsic parameters **Extrinsic parameters**

World2Camera Transformations: Exercise

$$X_c = \begin{bmatrix} R & t \\ 0 & 1 \end{bmatrix} X_w; \quad t = -R\tilde{C}$$



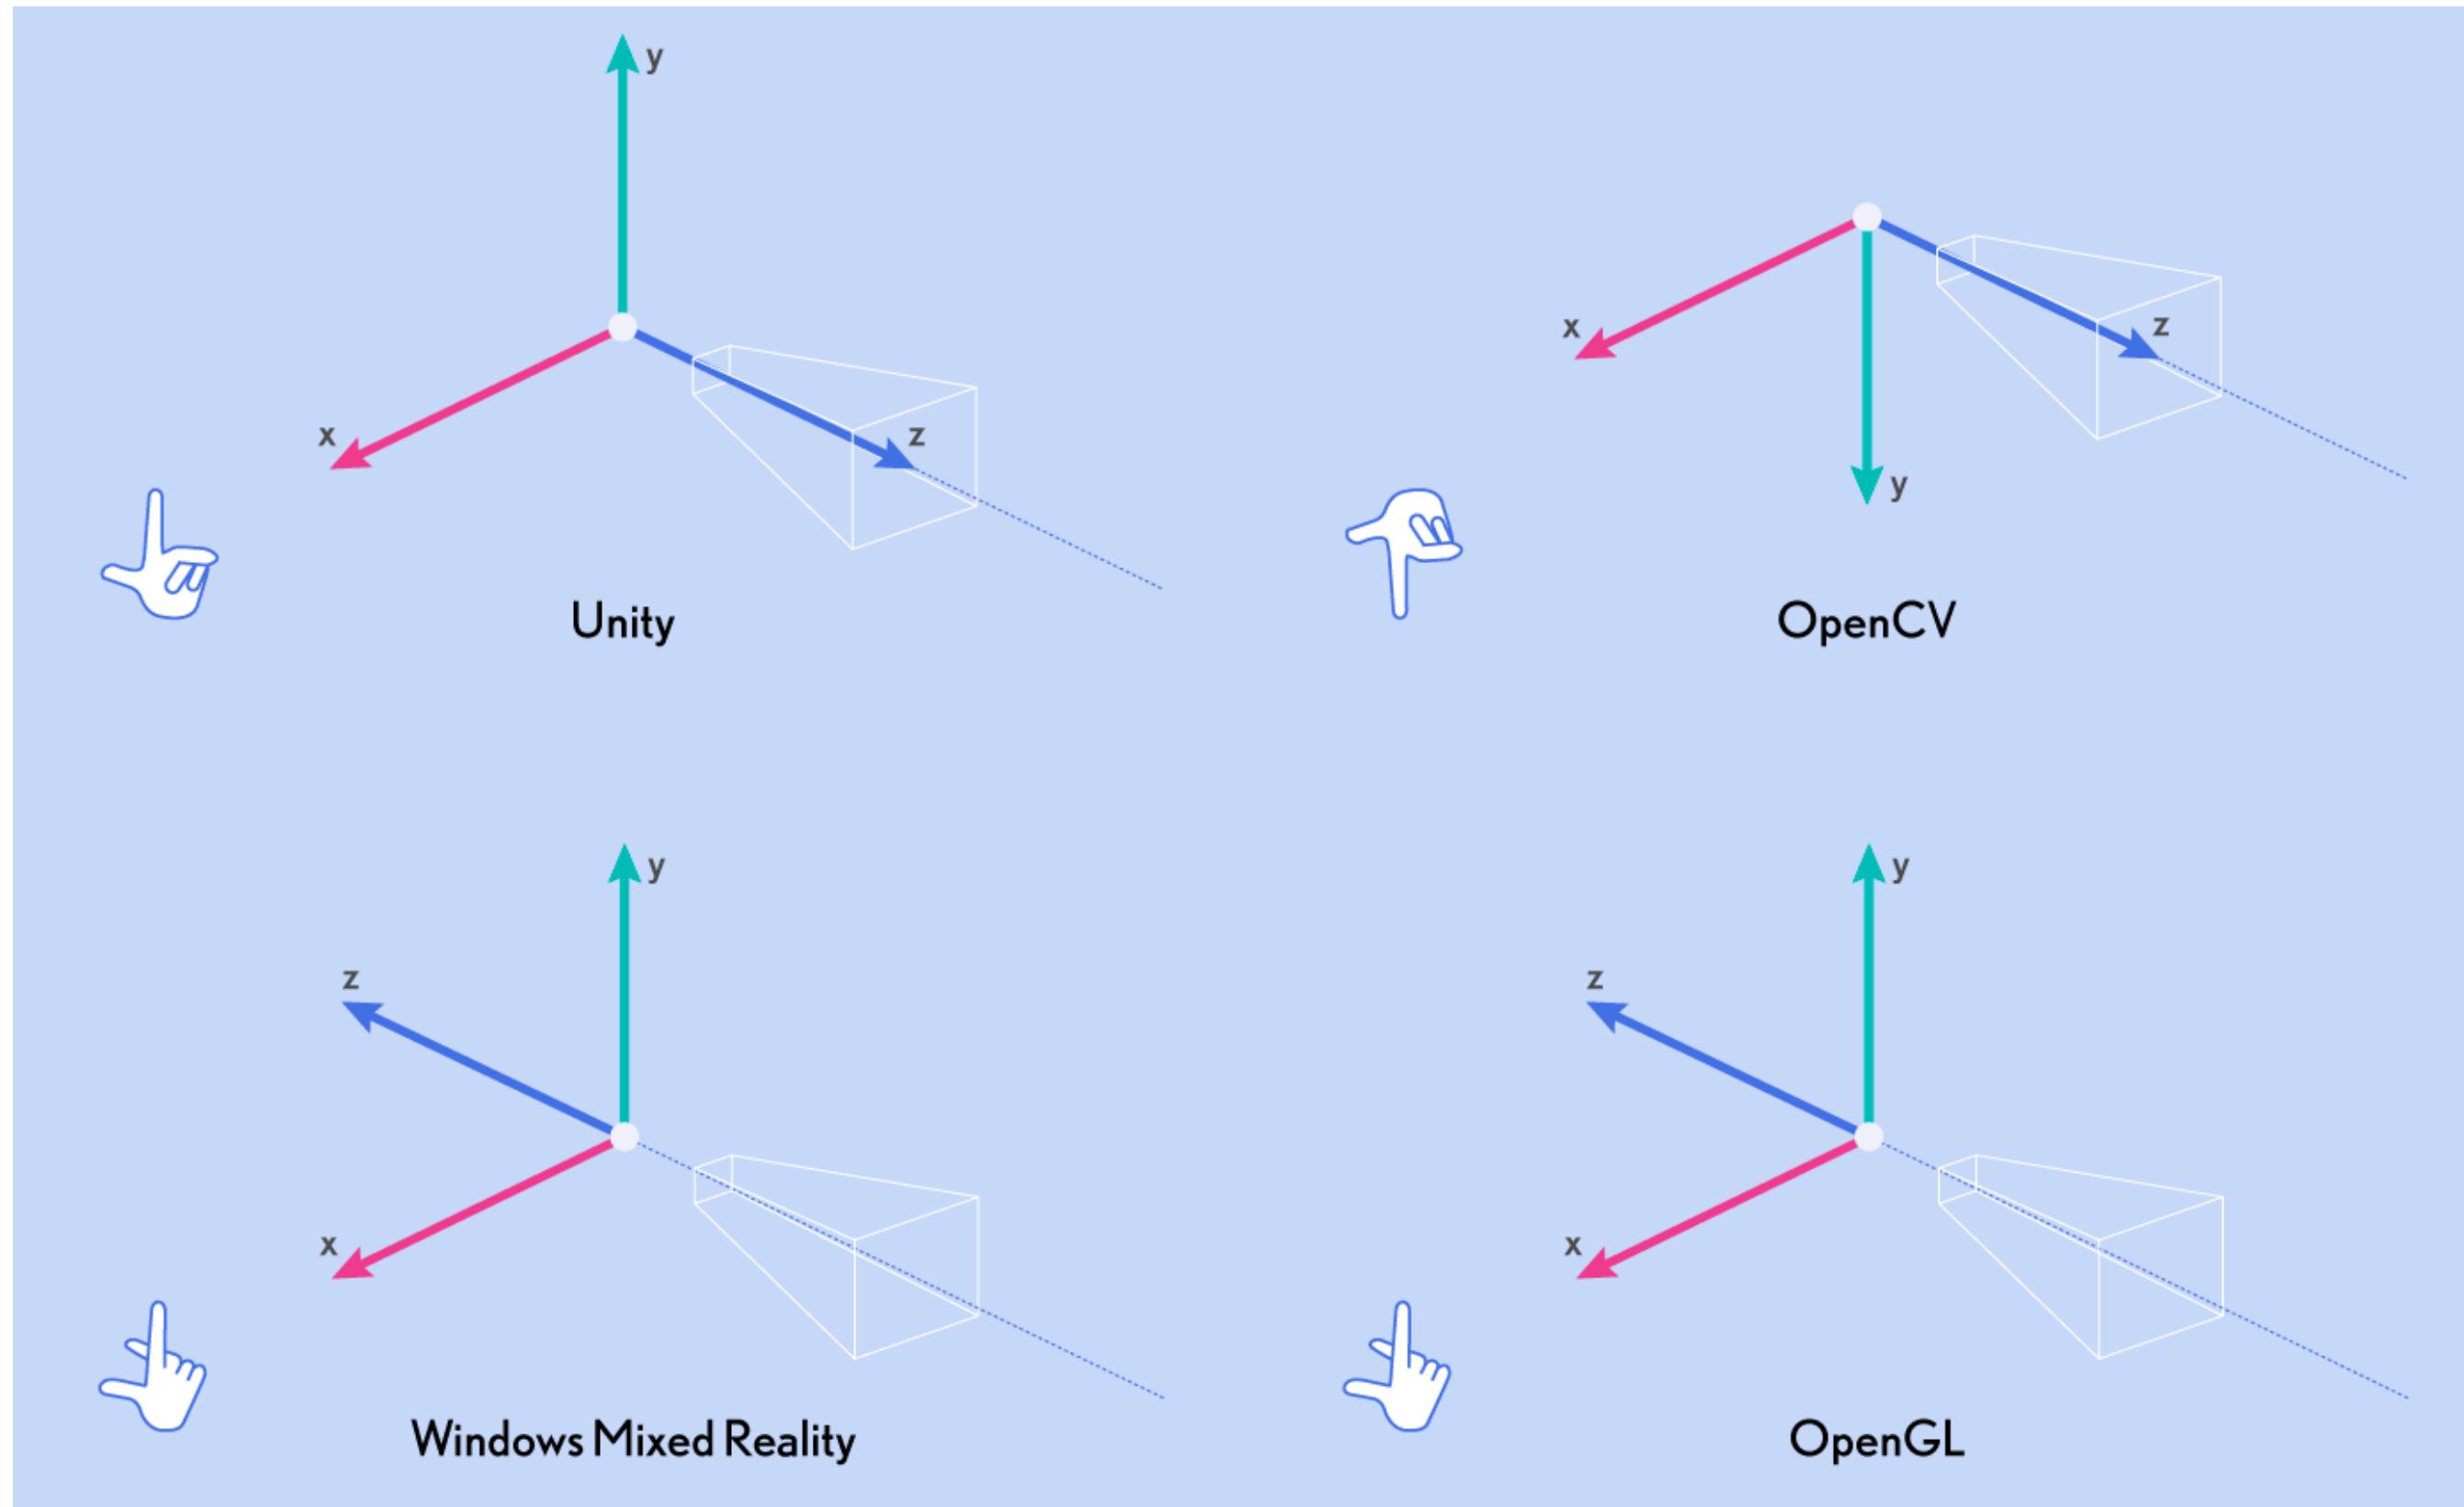
What are R , t for an upright (world and camera y-directions align) origin-facing camera 2m away from origin located at $(0,0,-2)$?

$$R = I; \quad t = \begin{bmatrix} 0 \\ 0 \\ 2 \end{bmatrix}$$

What are R , t for an upright (world and camera y-directions align) origin-facing camera 2m away from origin located at $(-2,0,0)$?

$$R = \begin{bmatrix} 0 & 0 & -1 \\ 0 & 1 & 0 \\ 1 & 0 & 0 \end{bmatrix}; \quad t = \begin{bmatrix} 0 \\ 0 \\ 2 \end{bmatrix}$$

Beware of Conventions



X=left or right? Y=up or down? etc..

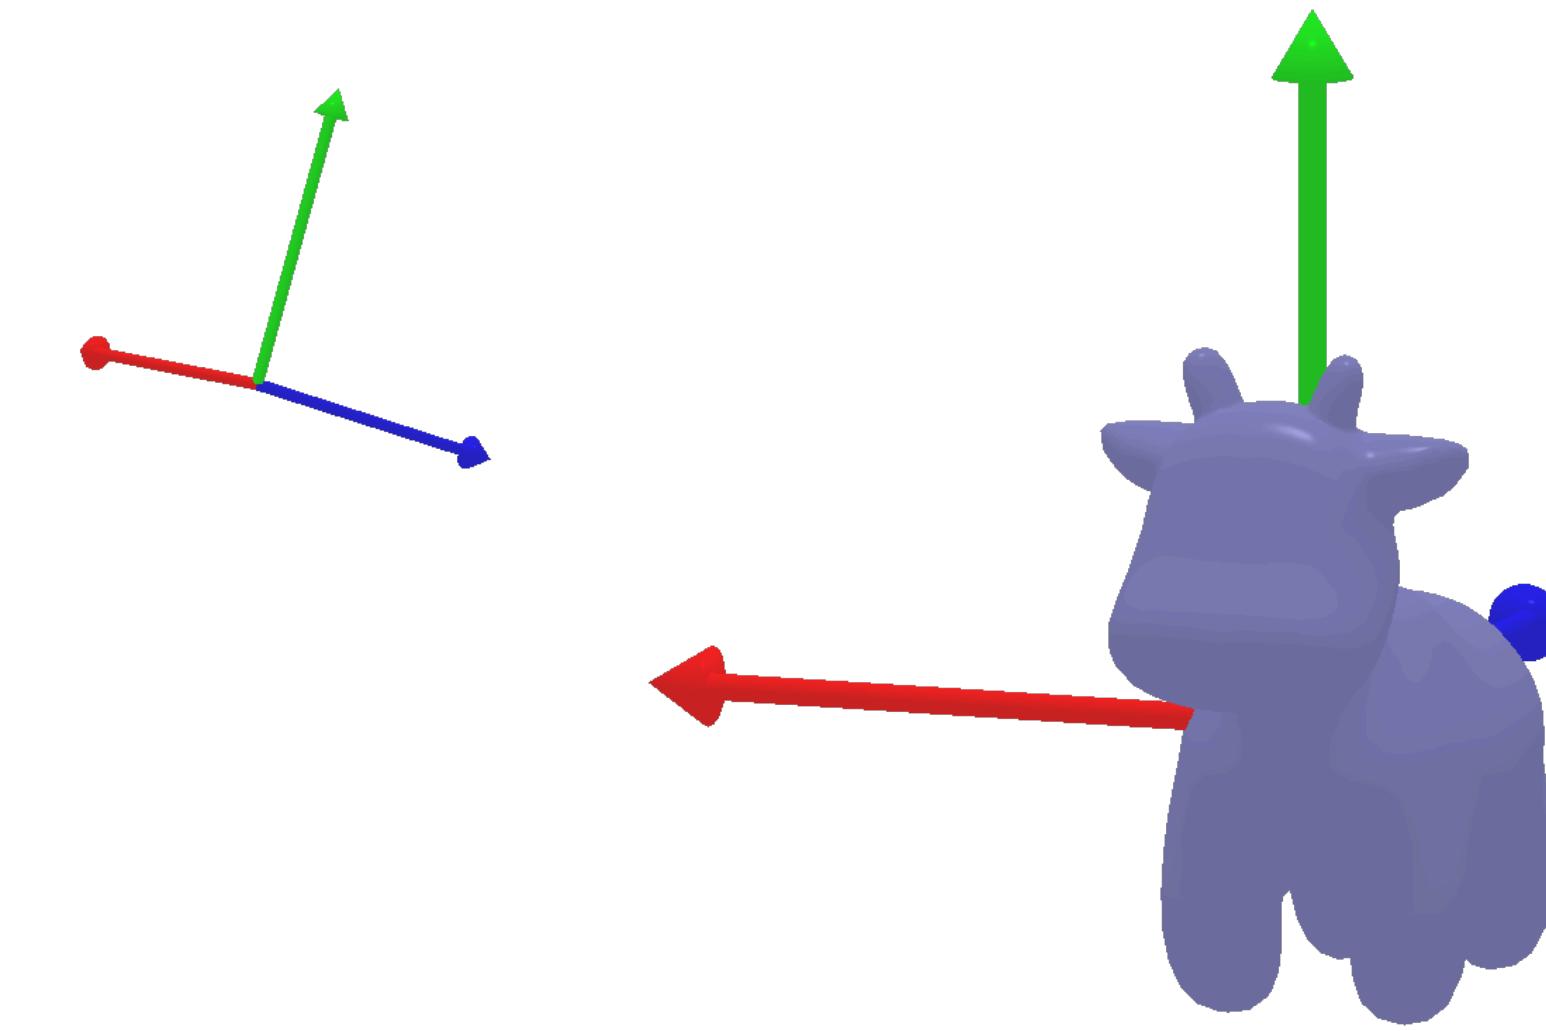
If you externally obtain transformation matrices (e.g. using someone else's code), make sure of convention compatibility (one of the *most common* sources of bugs!)

read <https://pytorch3d.org/docs/cameras>

Summary



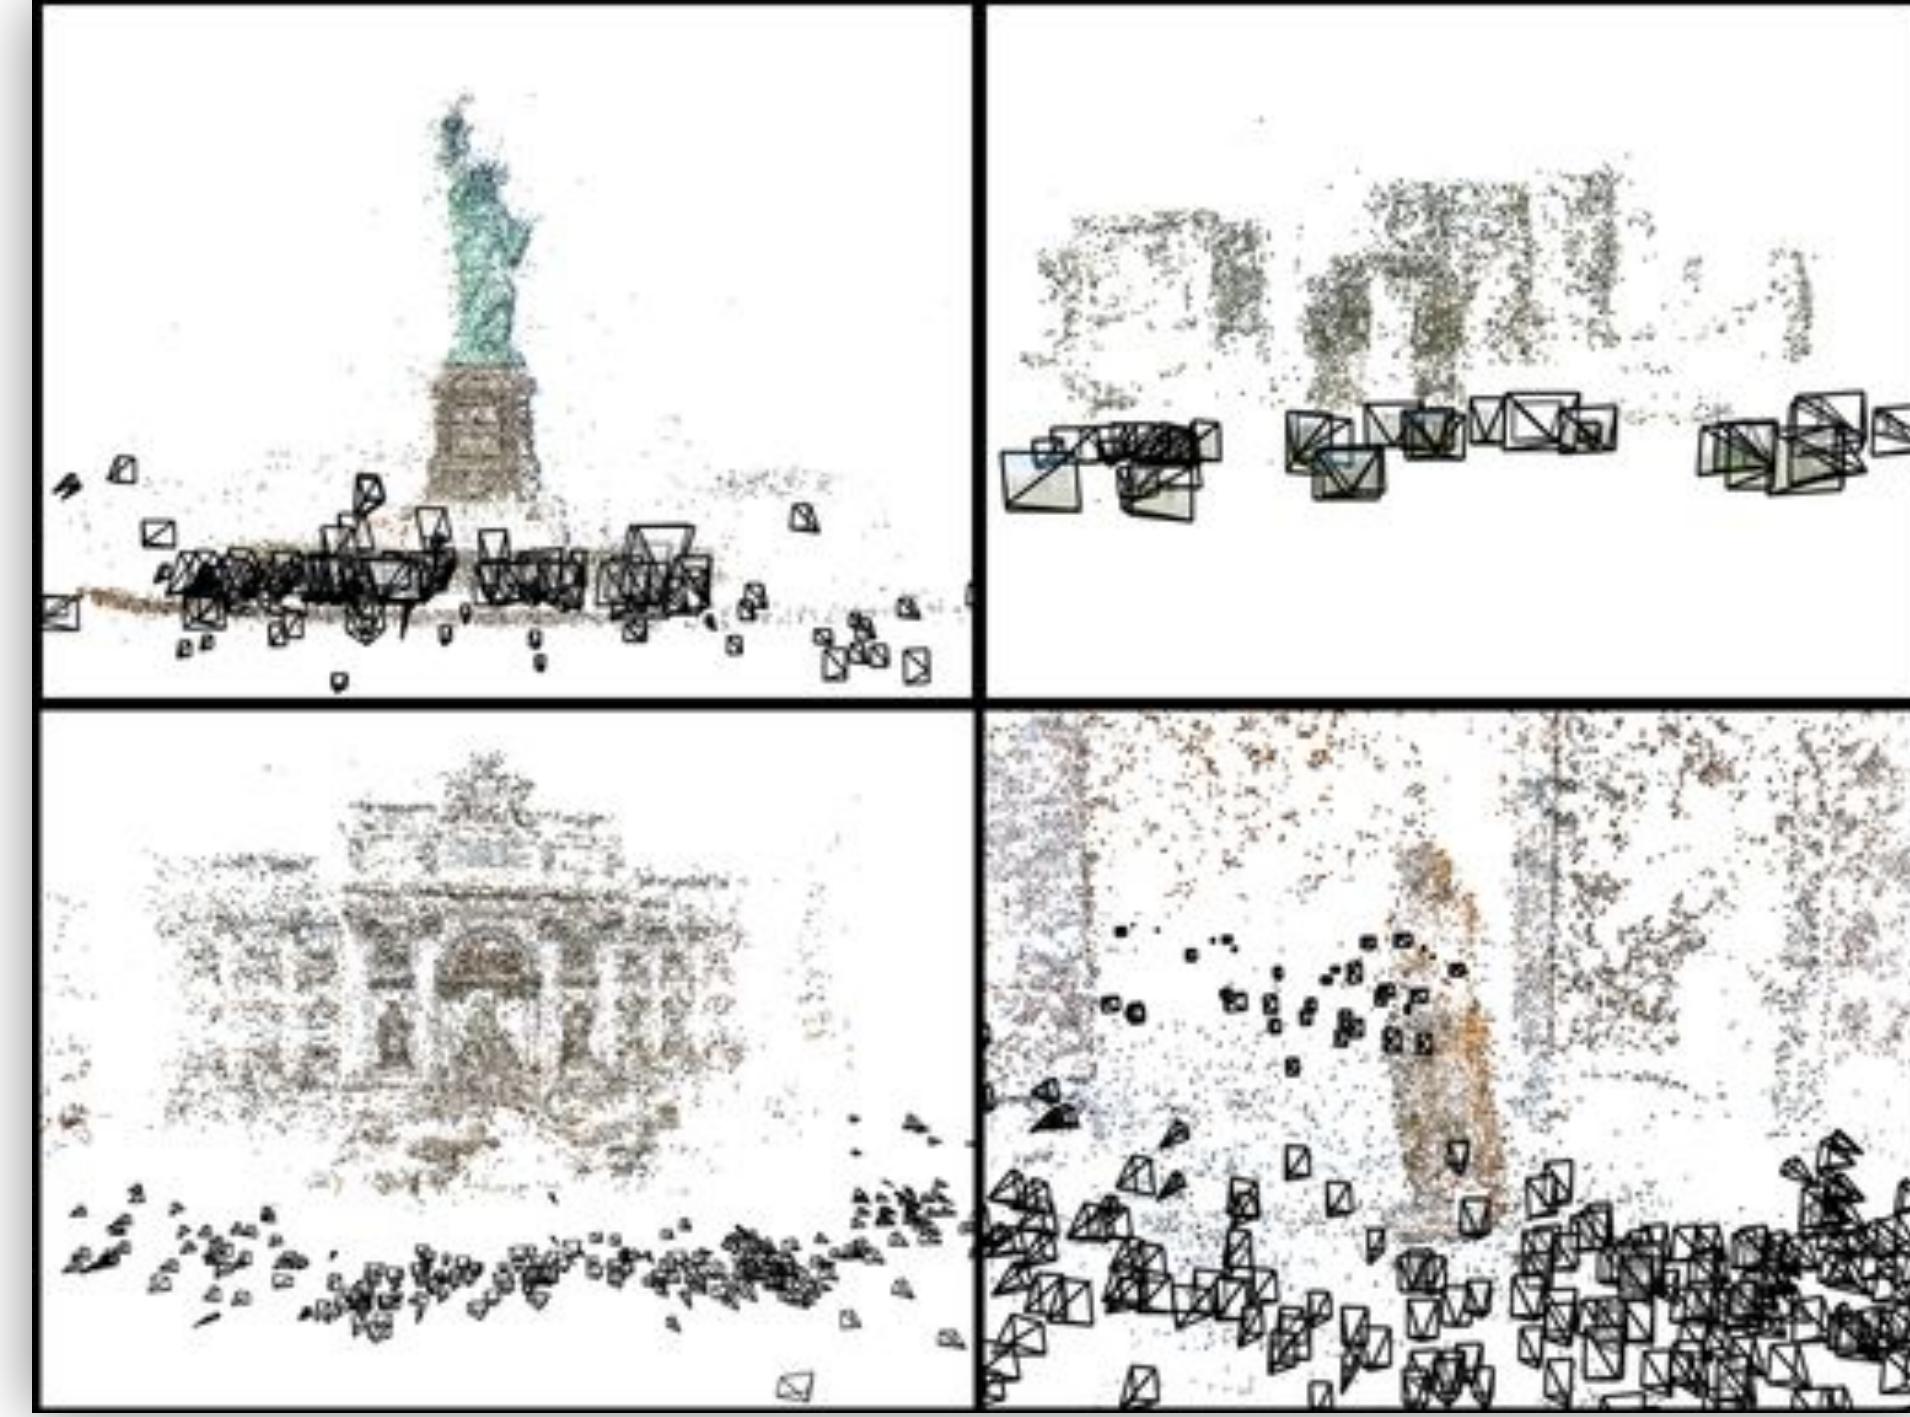
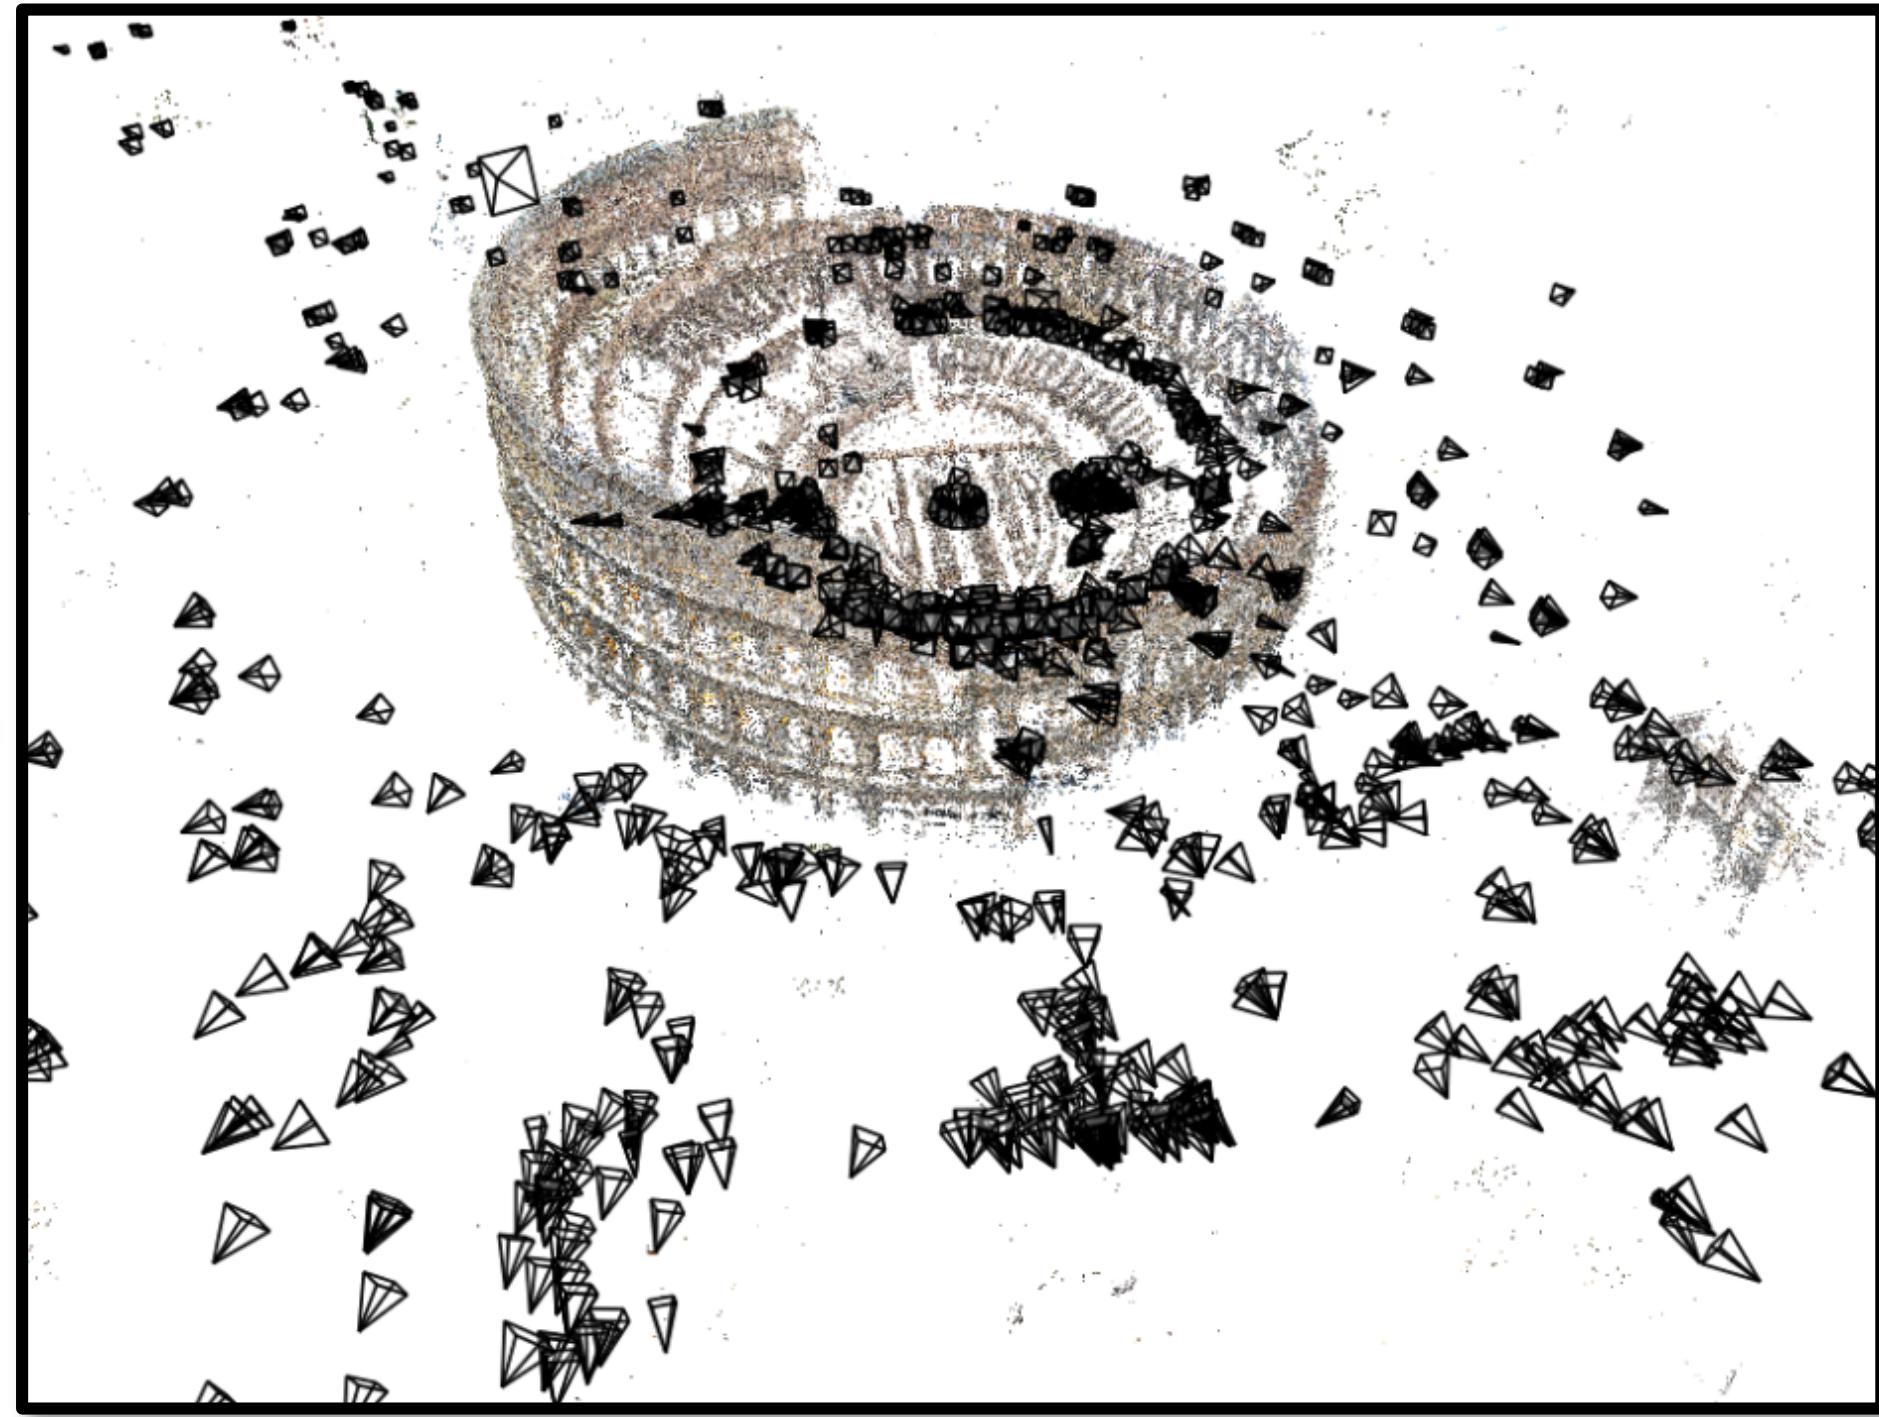
Projection — associating rays to points in a plane



Camera Transformation — from world to camera coordinates

Questions?

Module 1: Image Formation and Multi-View Geometry



Why?

We want to understand 3D world only from 2D observations (images). For that, we need to have a mathematical understanding of how they are connected.

What you'll learn.

Mathematical model of cameras. Reconstruct camera poses, approximate geometry, and camera parameters from 2D images of a scene.