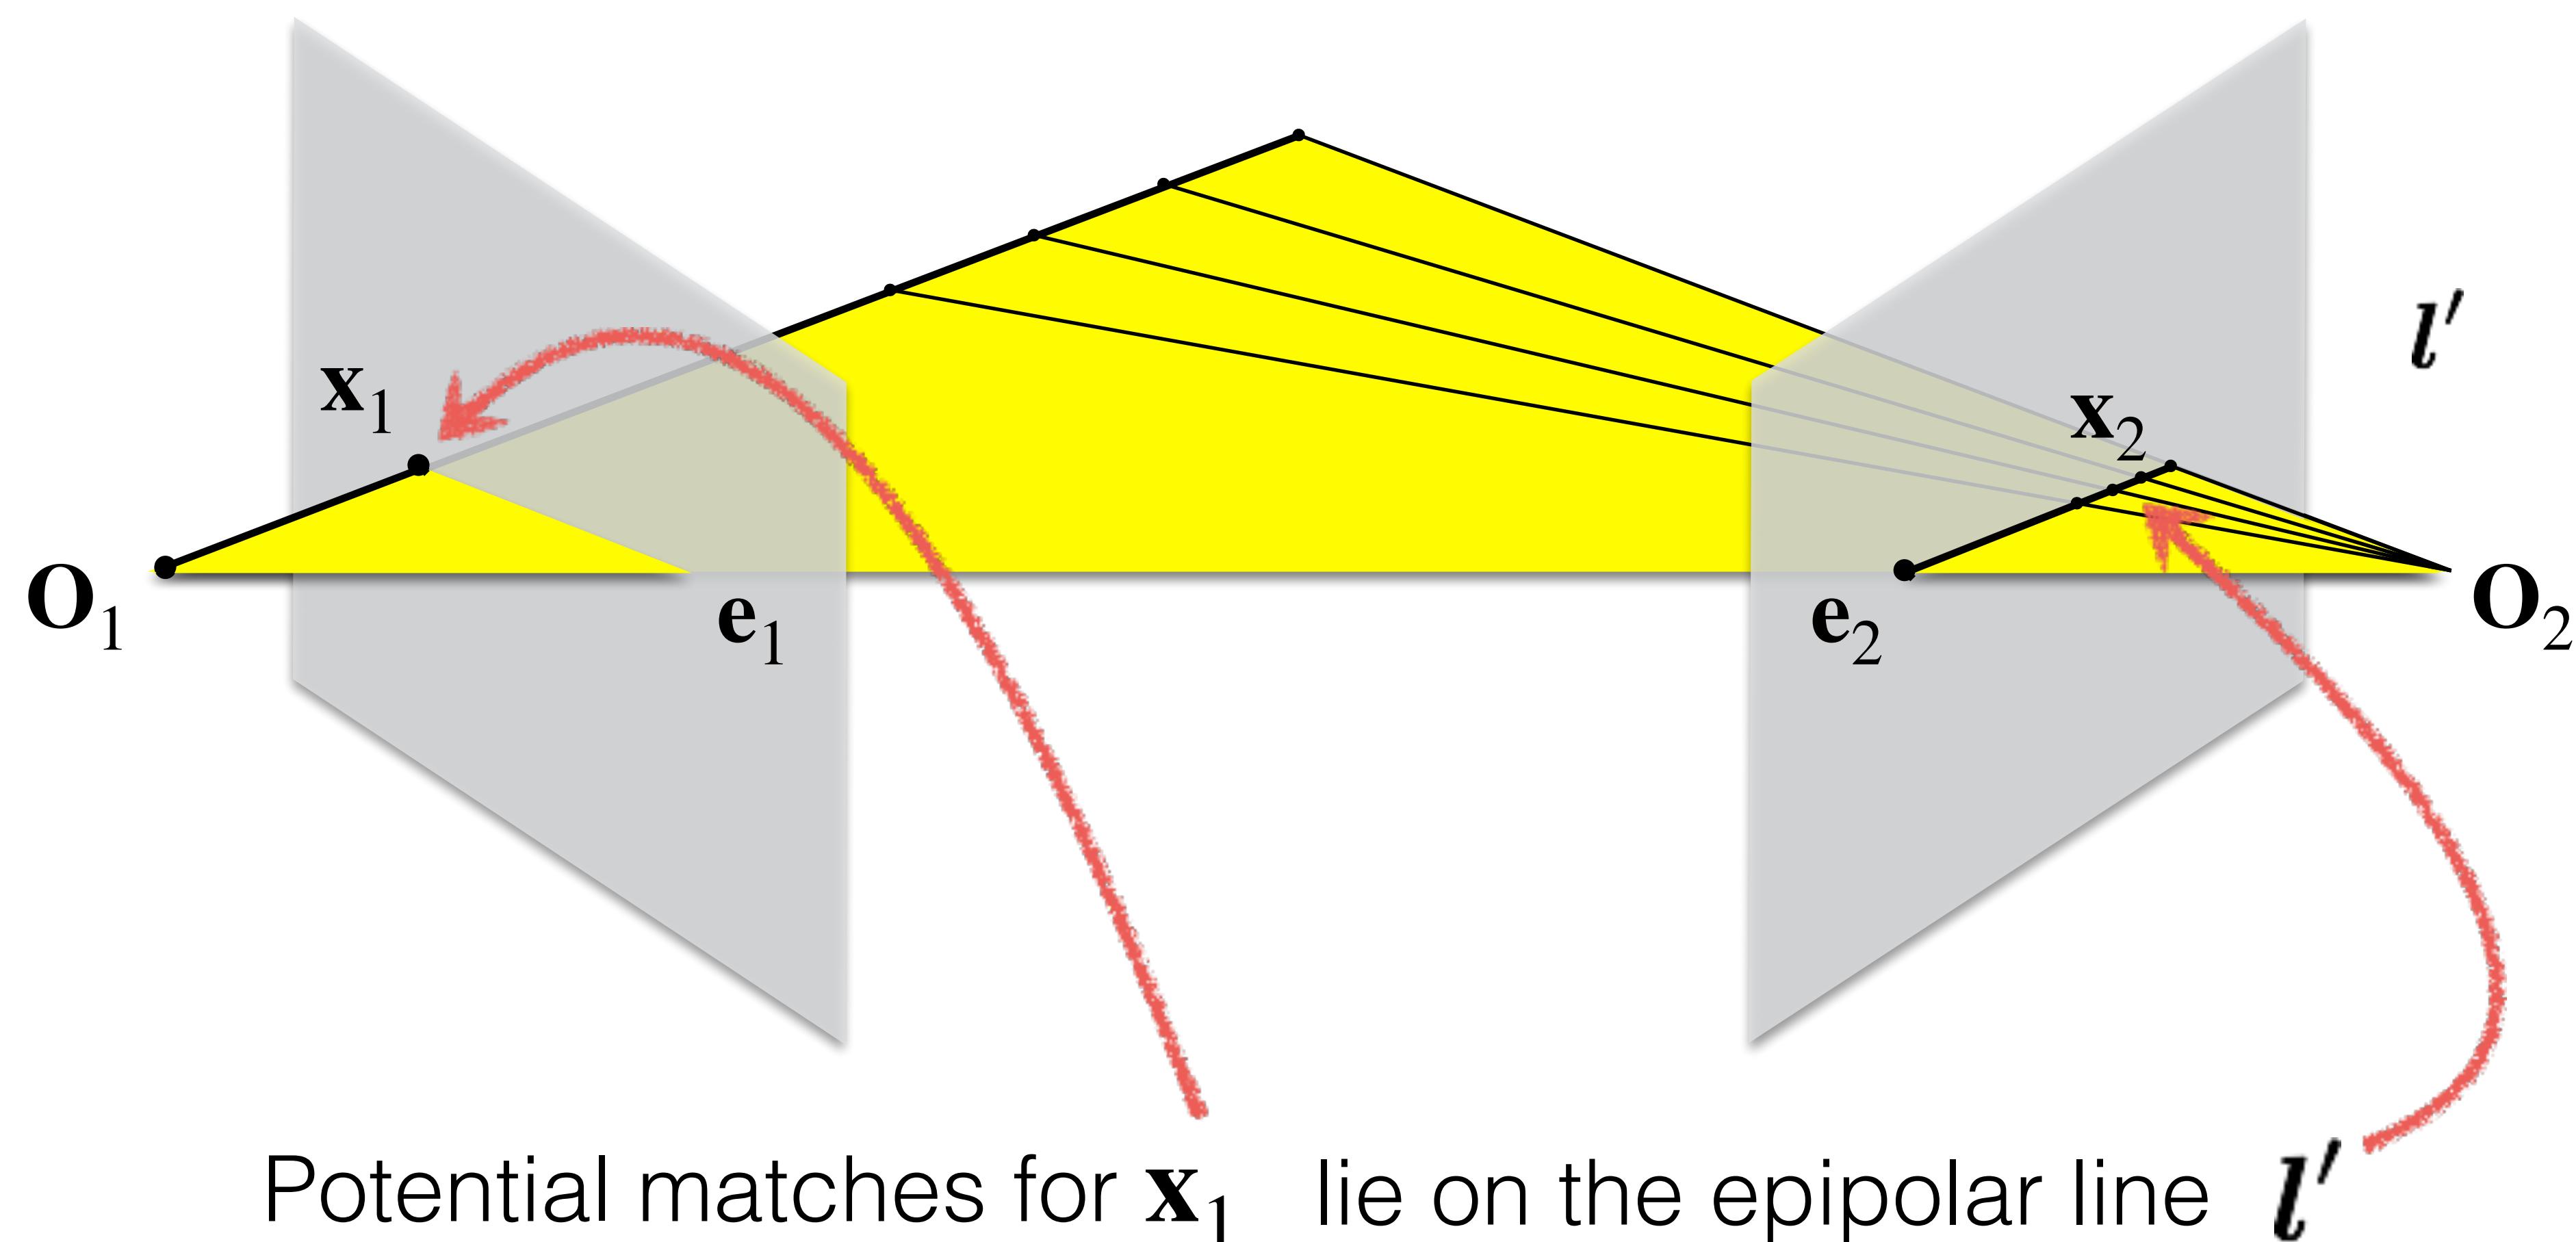


SCENE REPRESENTATIONS

Admin Things

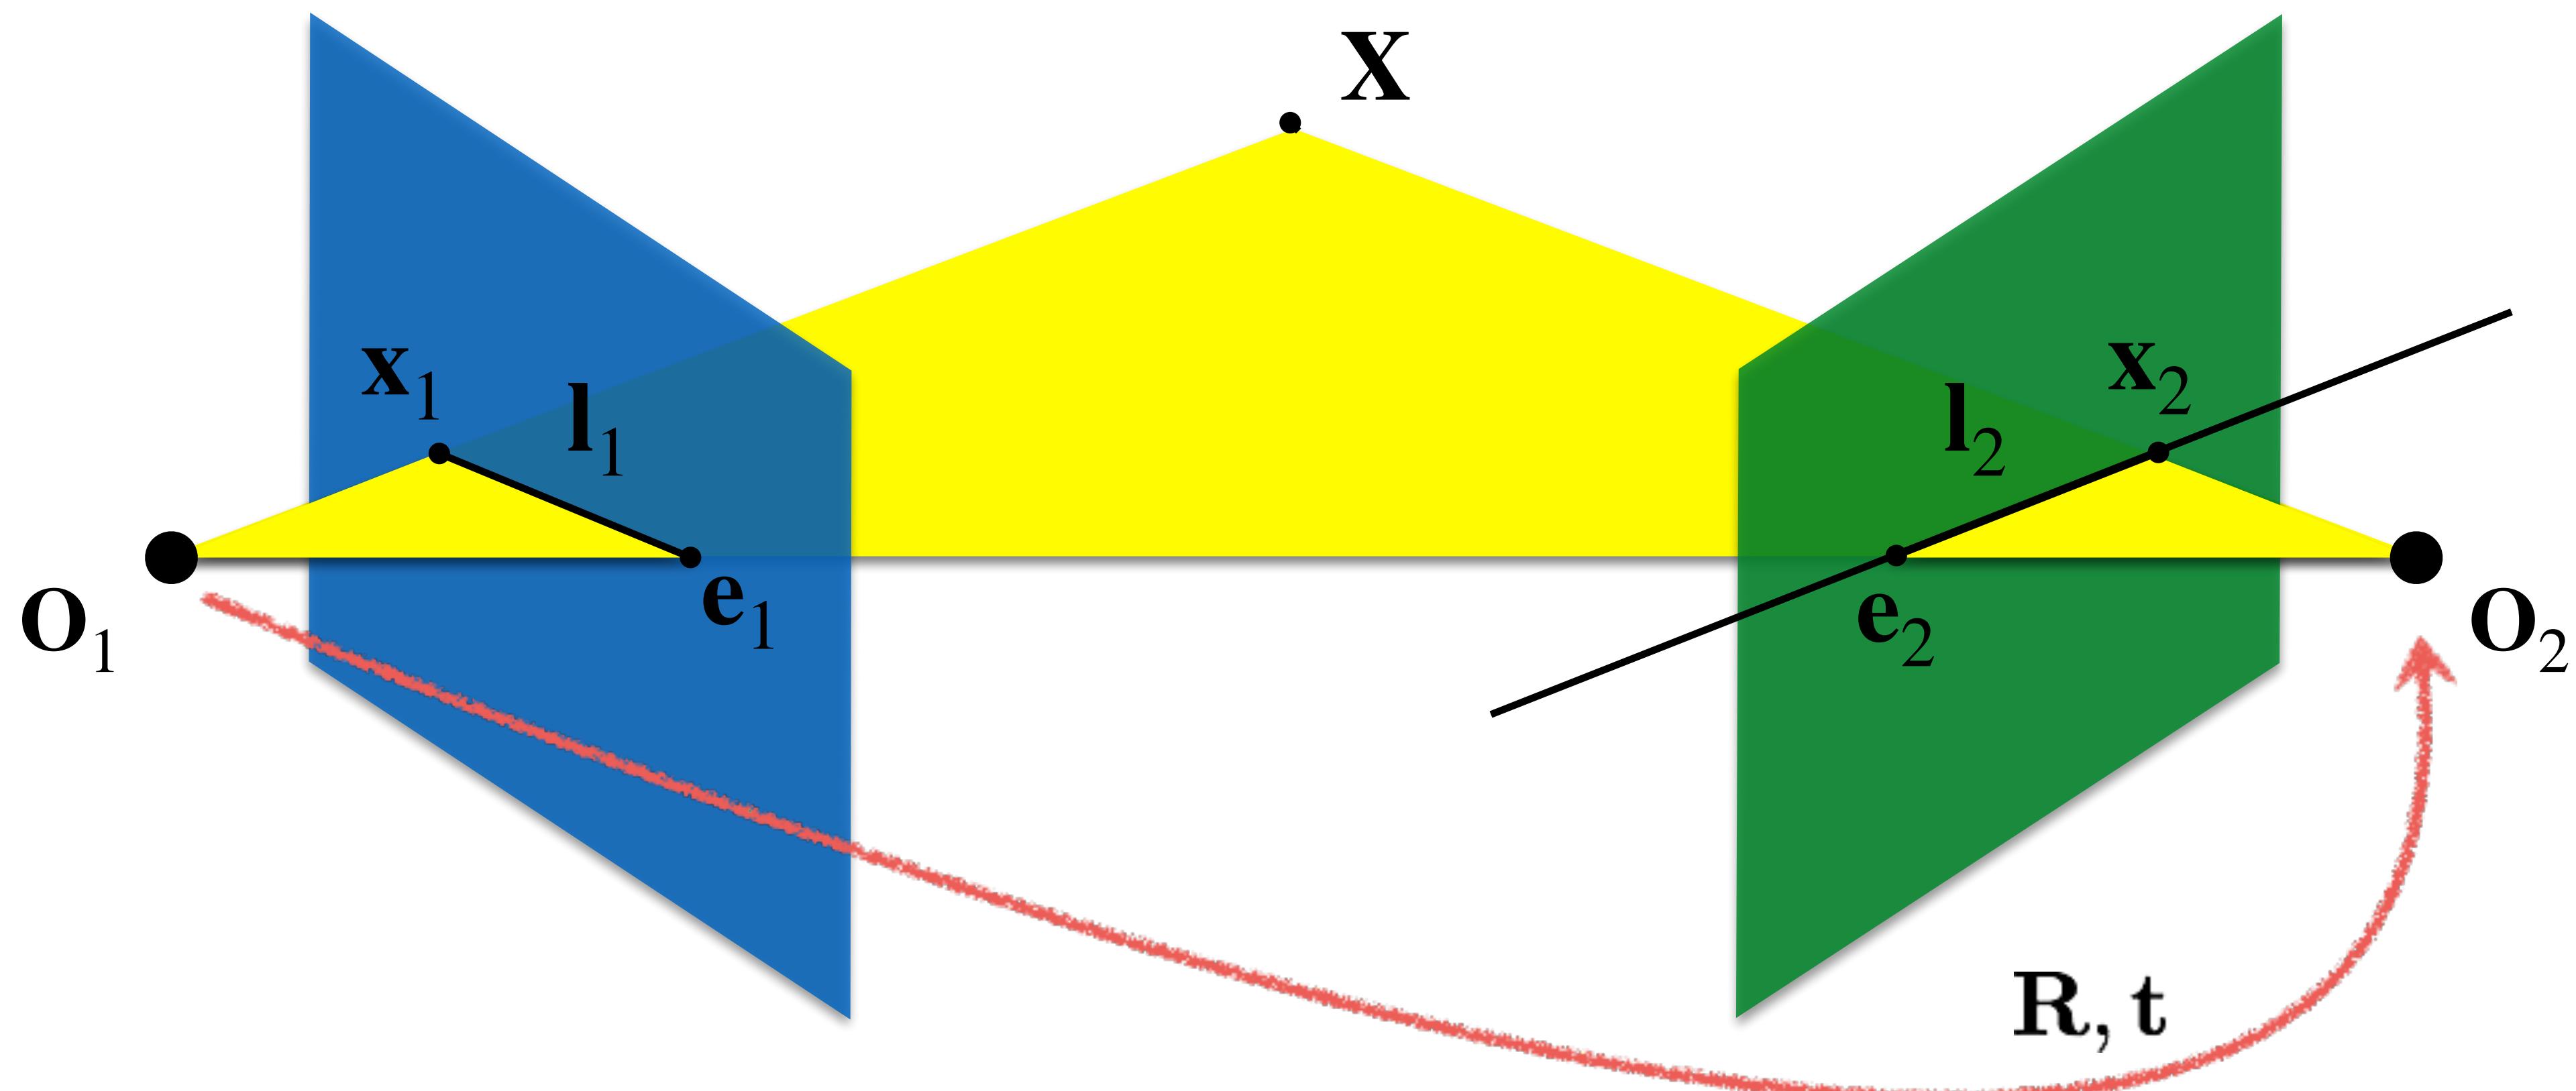
- Please get started with Assignment 1 if you haven't yet!
- Signup for Paper Sessions is online! Link:
https://docs.google.com/spreadsheets/d/1QdX7QJDVhRBA3J_8XN-IOkRRkqWOK0szOSqdPBA4BeE/edit?usp=sharing
- Some more guidance on paper sessions in next lecture. High-level:
 - Groups of 2-3 students.
 - You pick your paper! Can, but don't have to, pick from my suggestions.
 - 30 minutes total: ~7 minutes per student, ~9 minutes discussion.

Recap: Epipolar Geometry



$$\mathbf{F} = \mathbf{K}_2^{-1}(\mathbf{R}[t]_{\times})\mathbf{K}_1^{-1}$$

$$\mathbf{F}\tilde{\mathbf{x}}_1 = \mathbf{l}_2$$

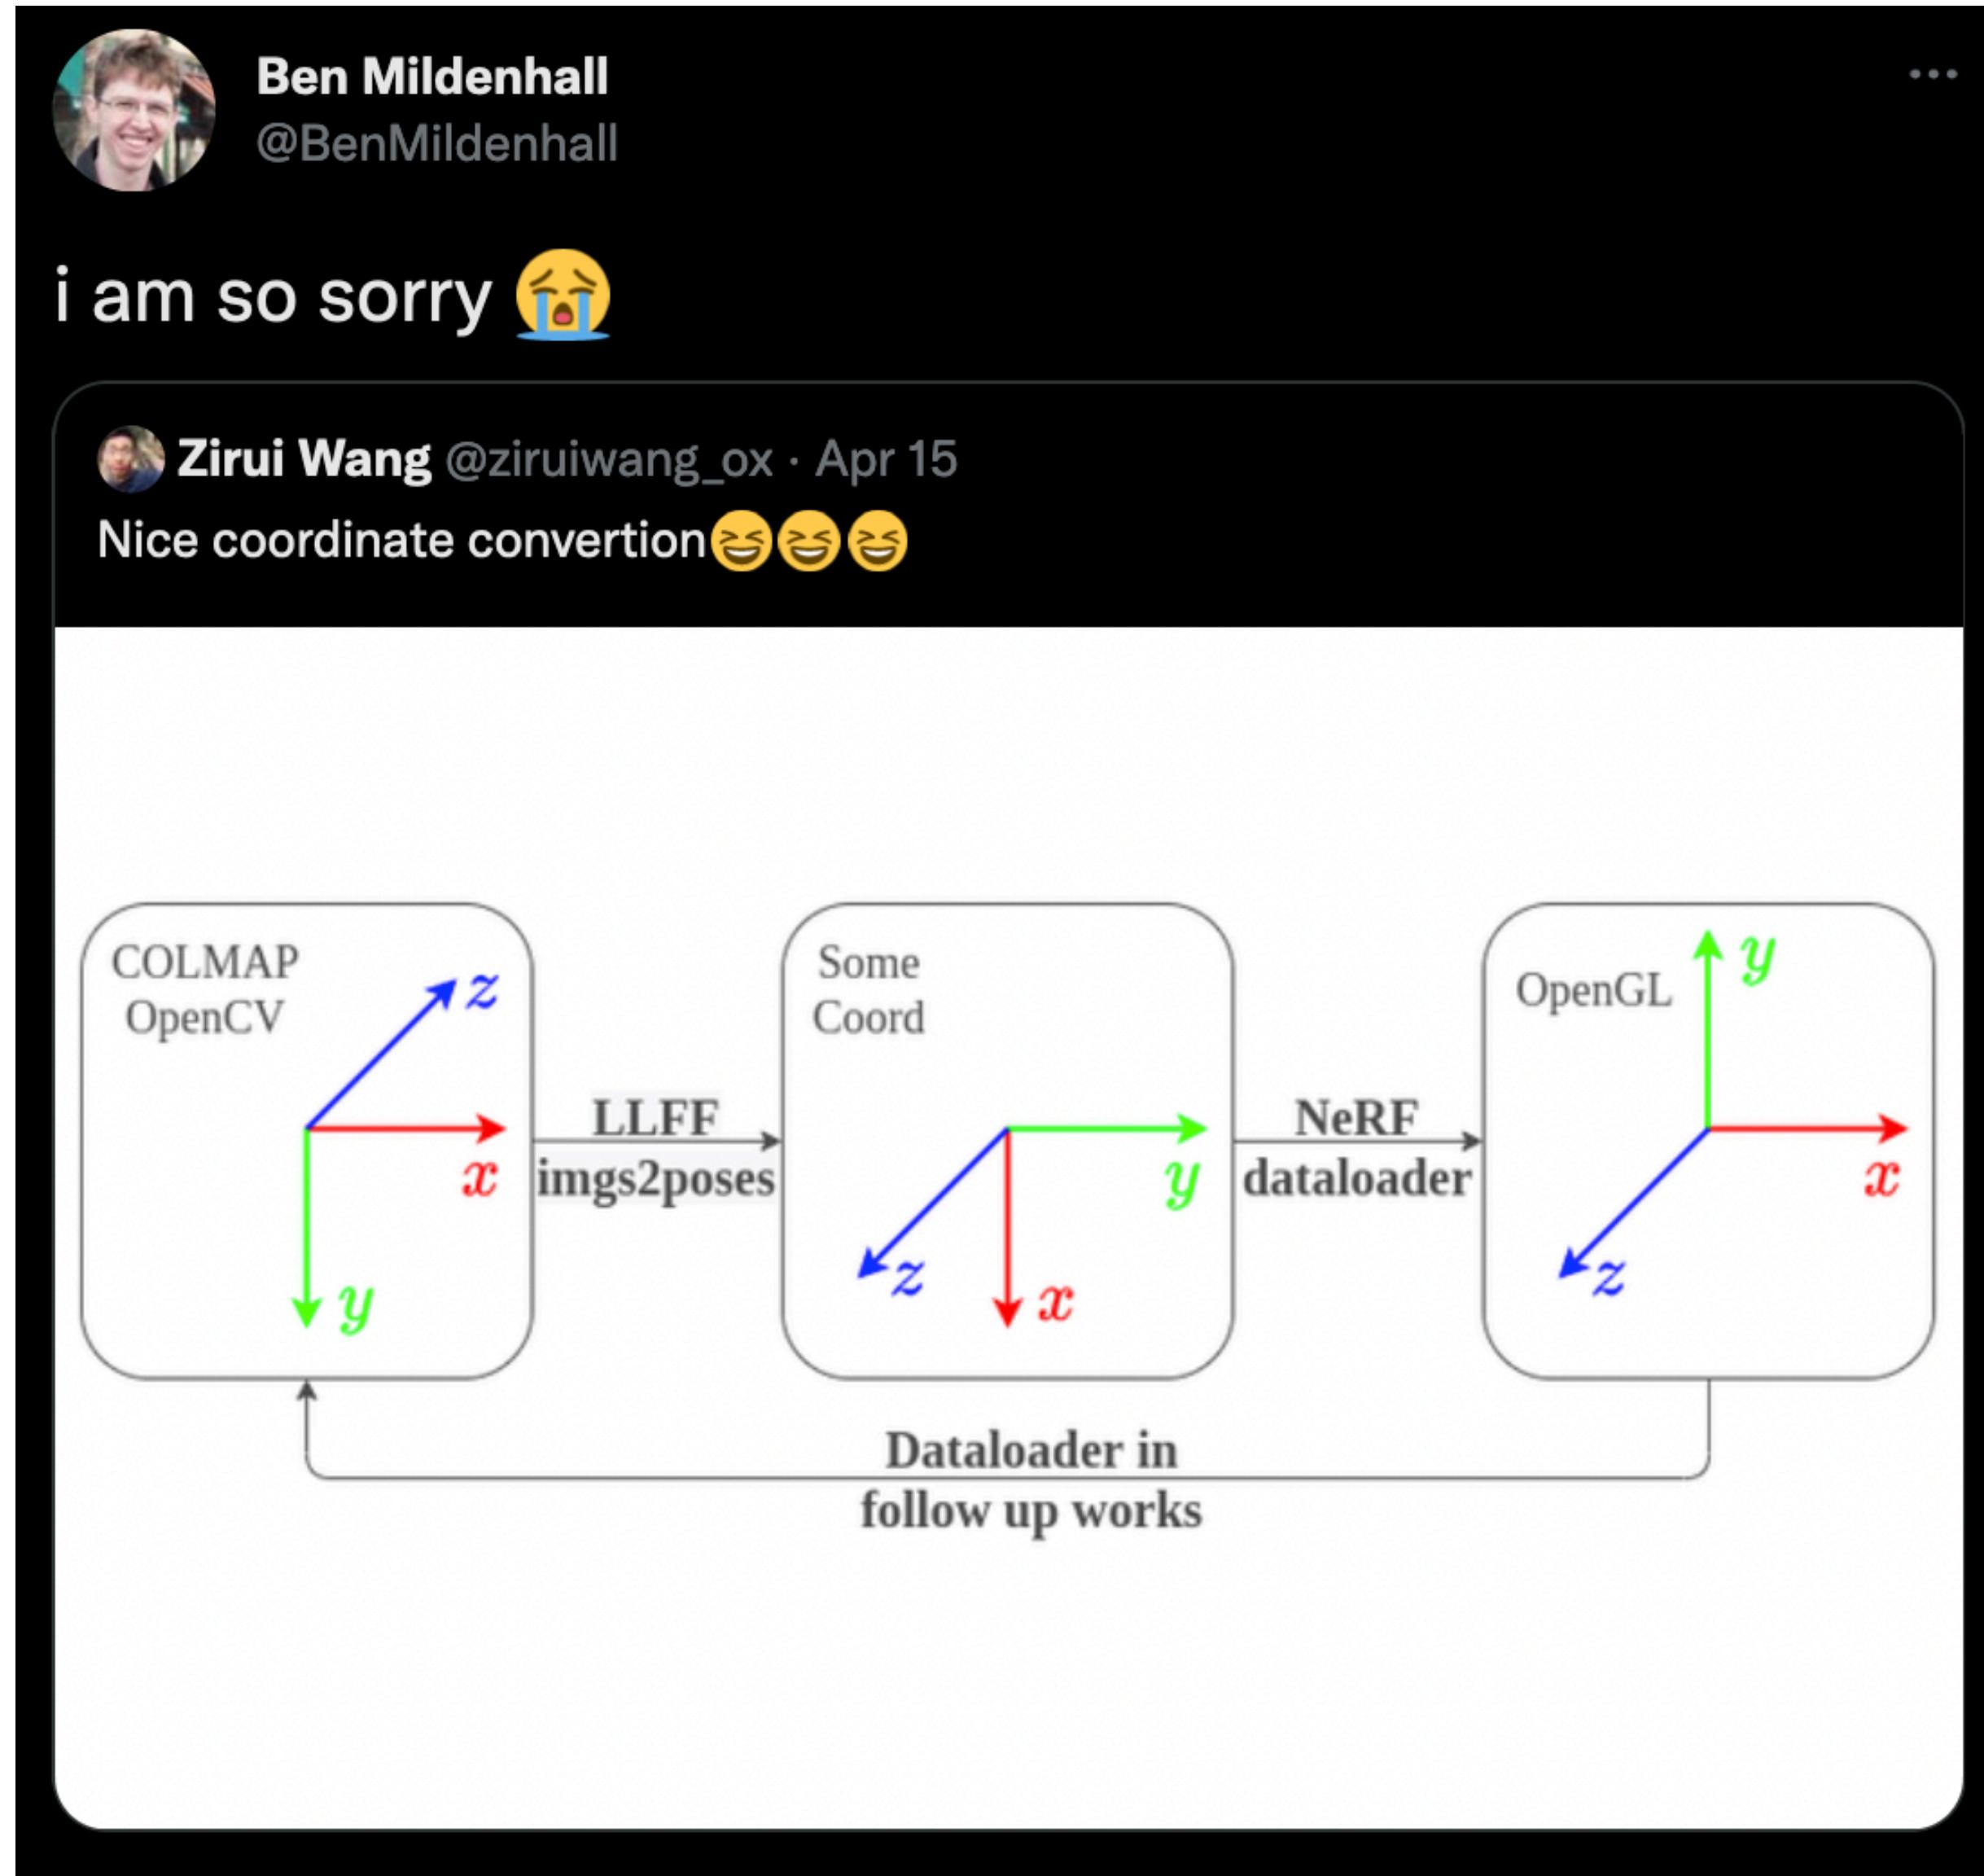


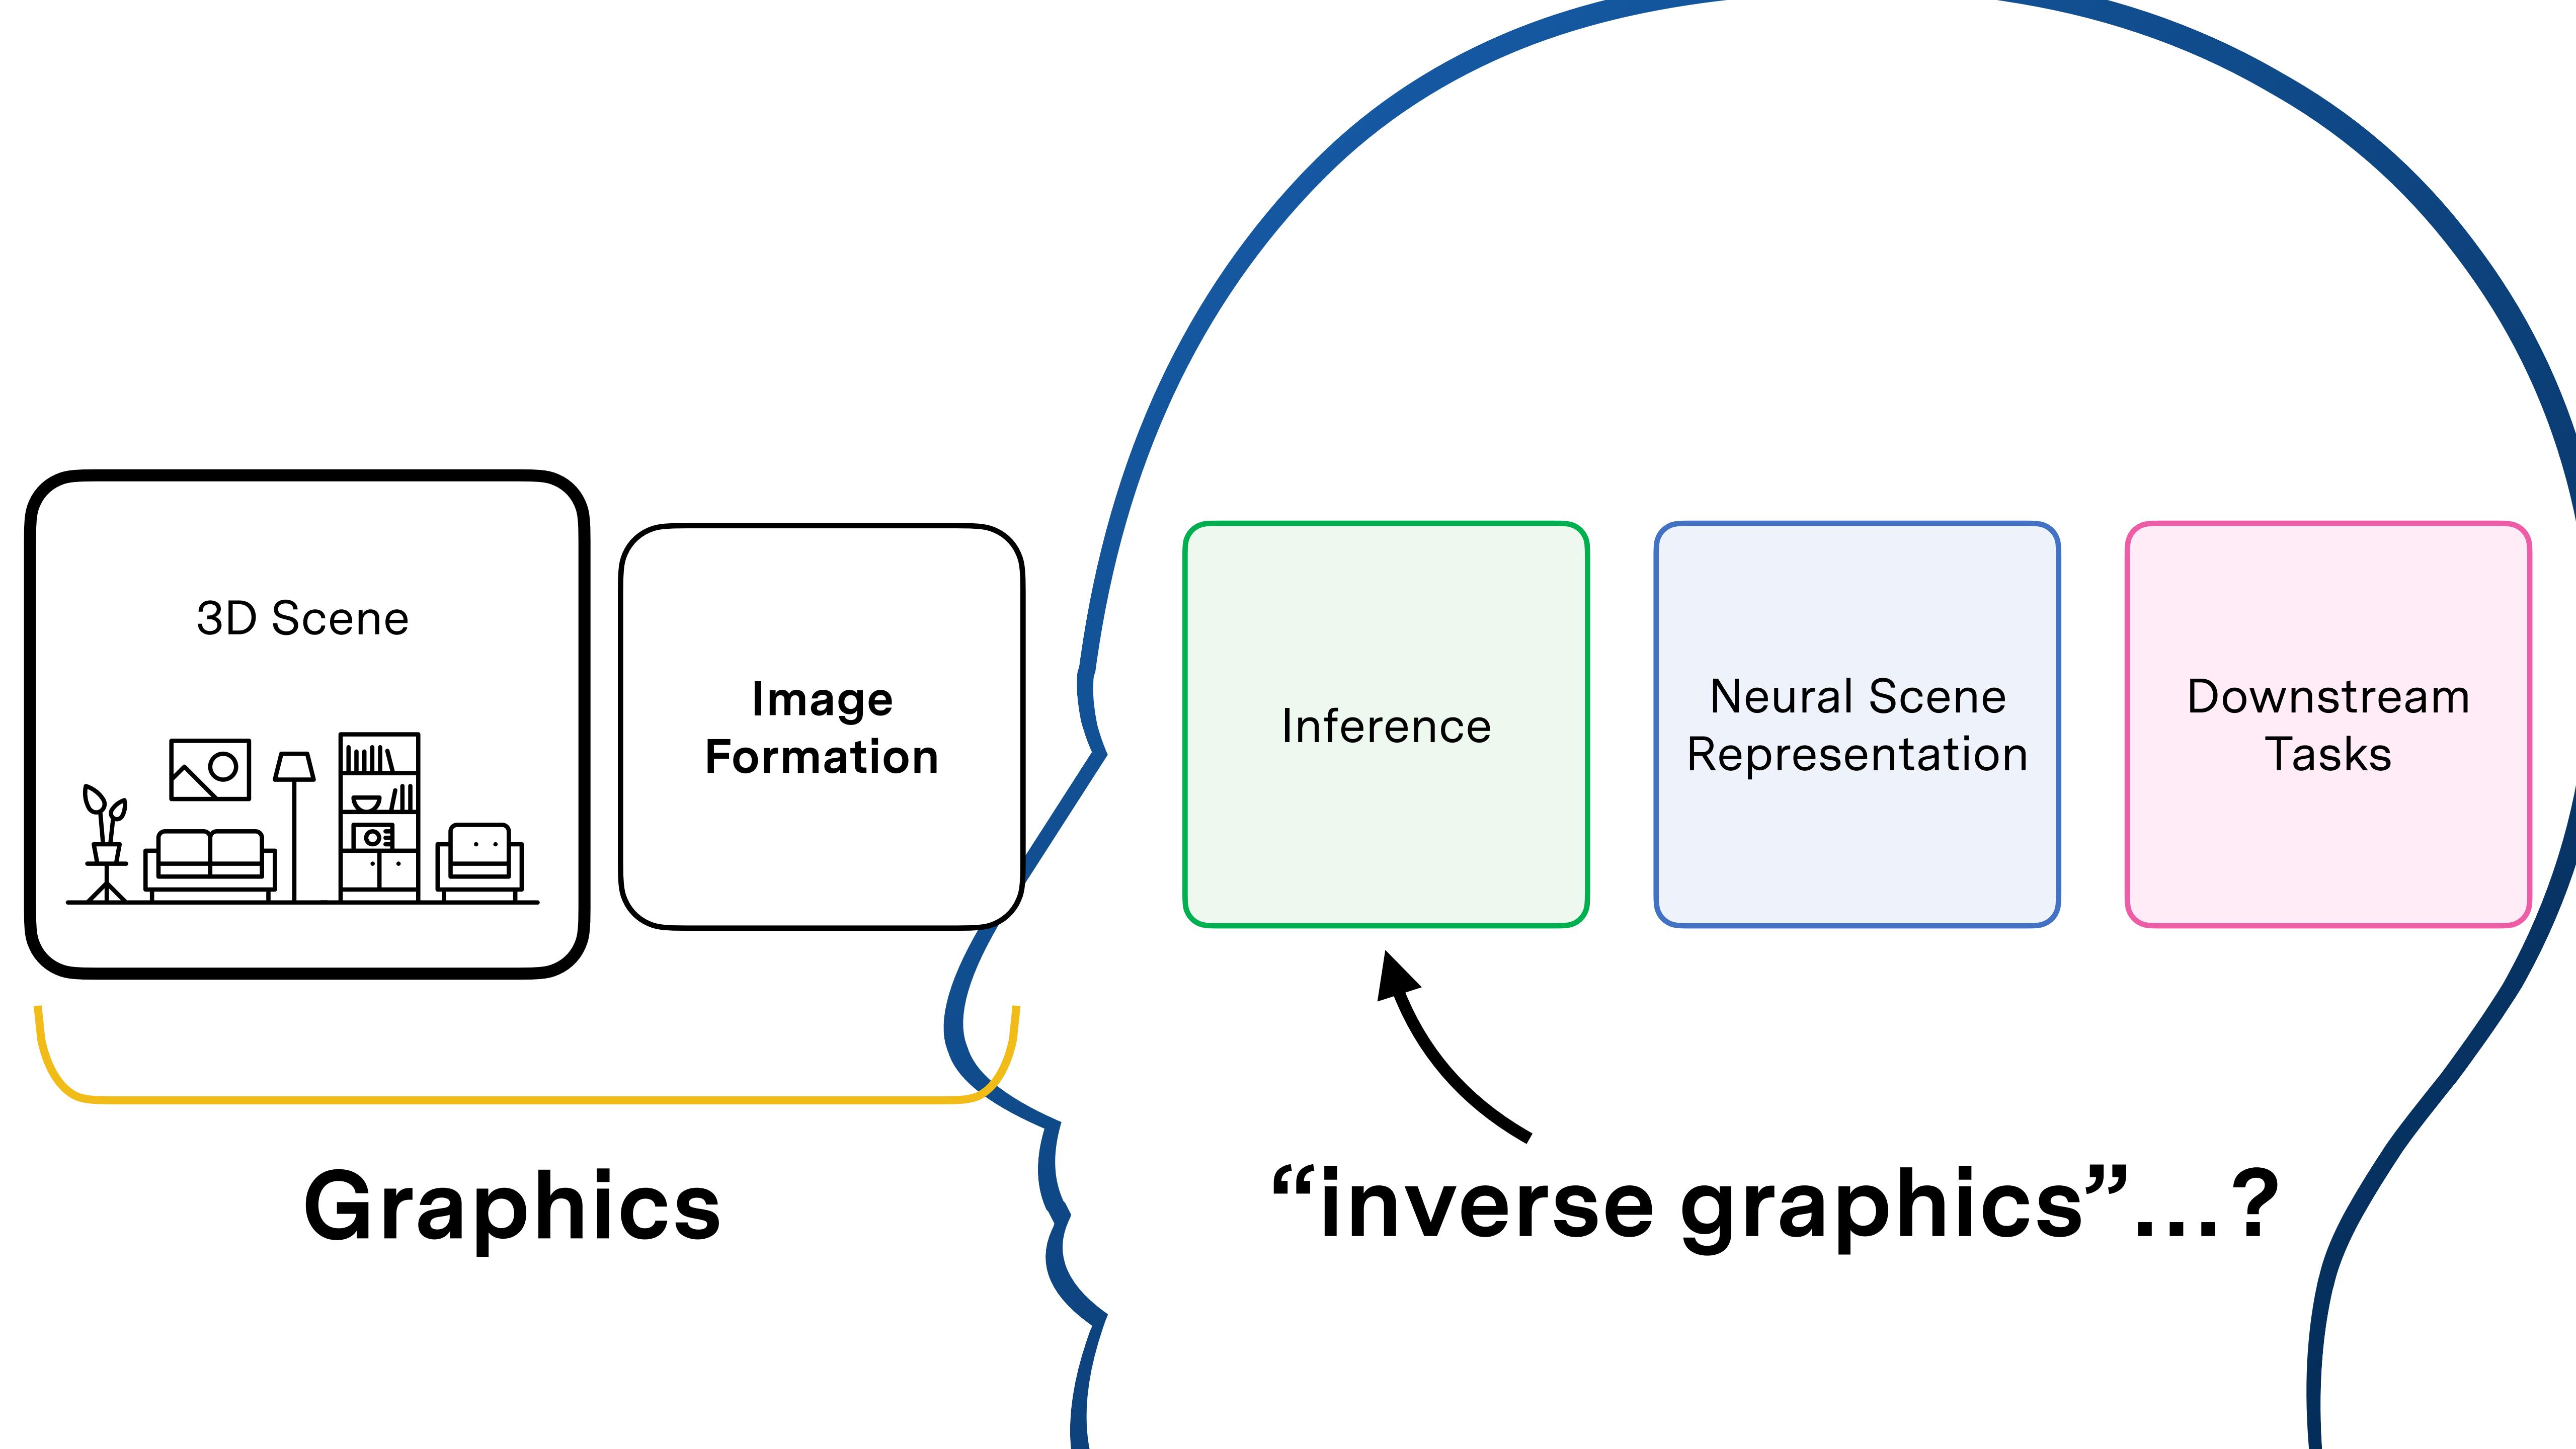
Recap: Bundle-Adjustment



$$\Pi^*, \mathcal{X}_w^* = \operatorname{argmin}_{\Pi, \mathcal{X}_w} \sum_{i=1}^N \sum_{p=1}^P w_{ip} \|\mathbf{x}_{ip}^s - \pi_i(\mathbf{x}_p^w)\|_2^2$$

On camera conventions





Today: How to computationally represent 3D scenes.

3D Scene

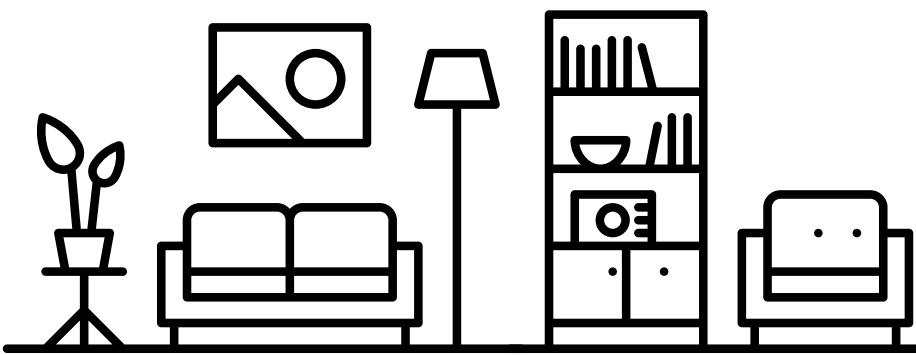


Image
Formation

Inference

Neural Scene
Representation

Downstream
Tasks

Why?

At the end of the day, we want to make predictions about 3D scenes.
For that, we need to know how we can represent 3D scenes computationally.

What you'll
learn.

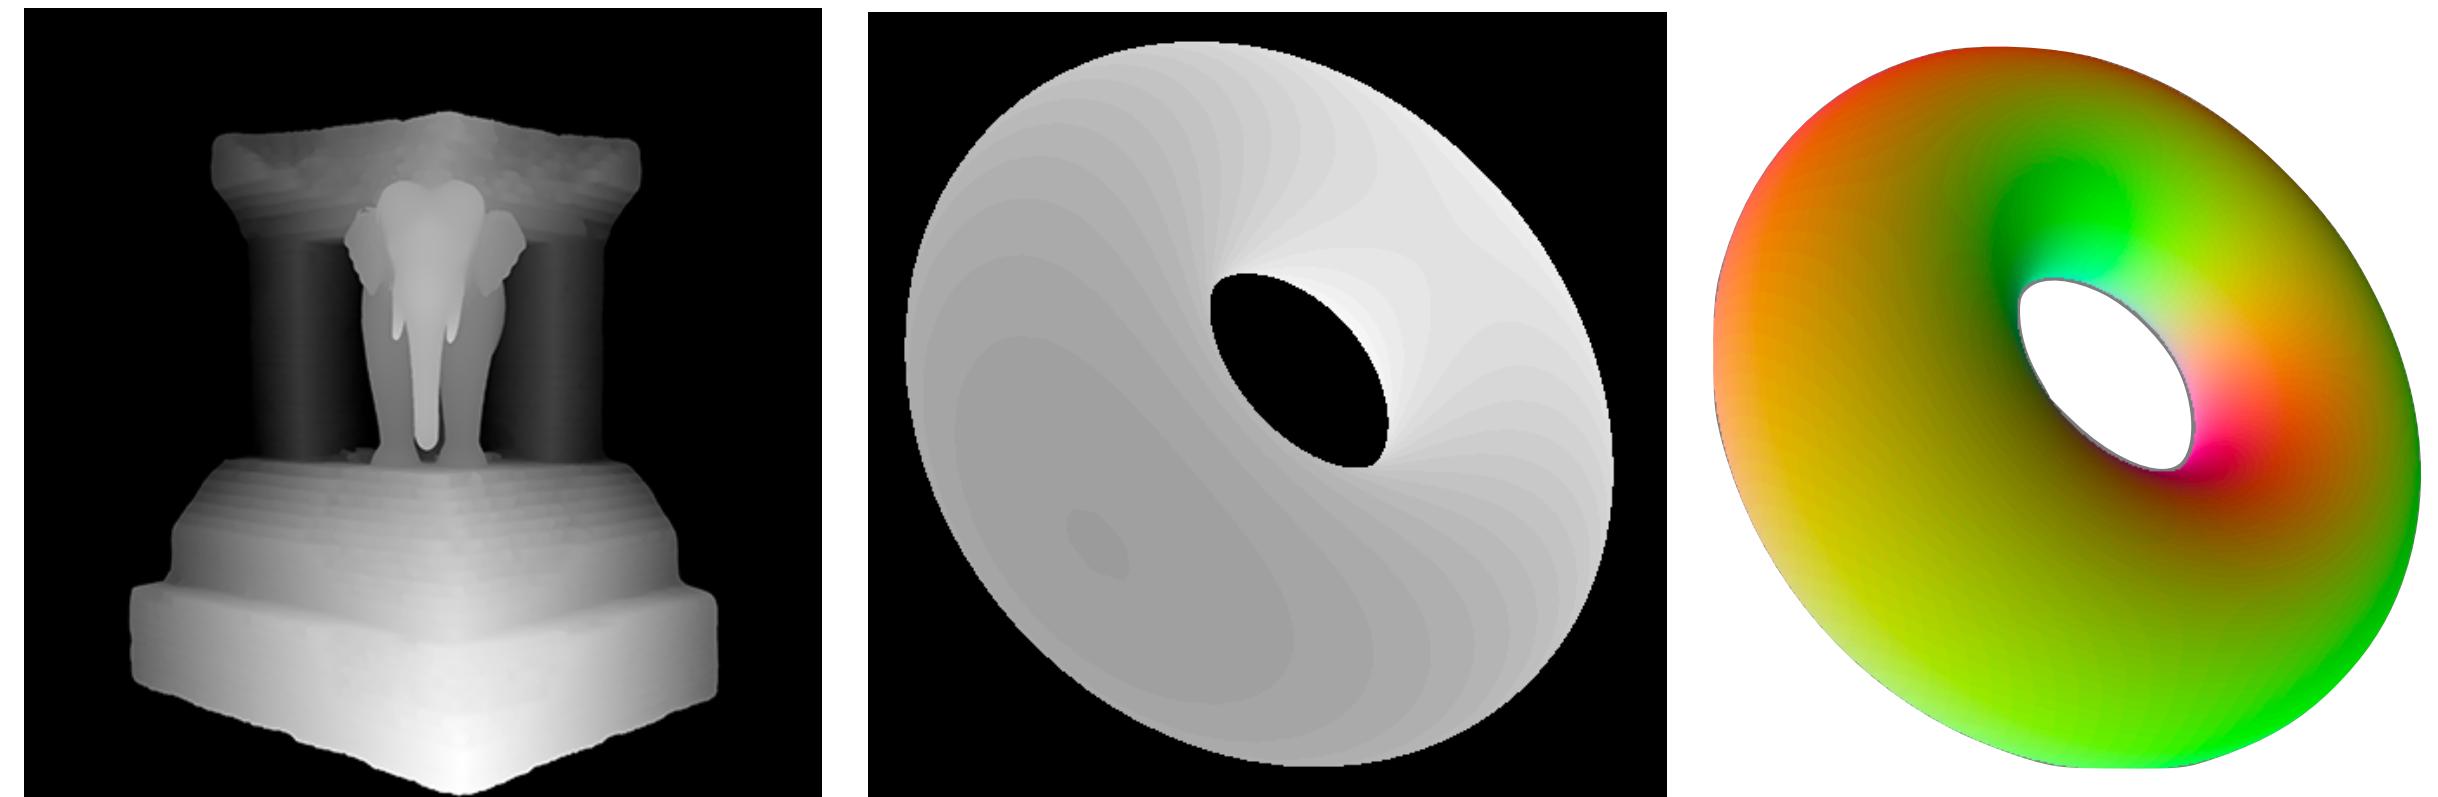
Surface-based representations, volumetric representations, discrete representations, continuous representations.

Some Slides adapted from...

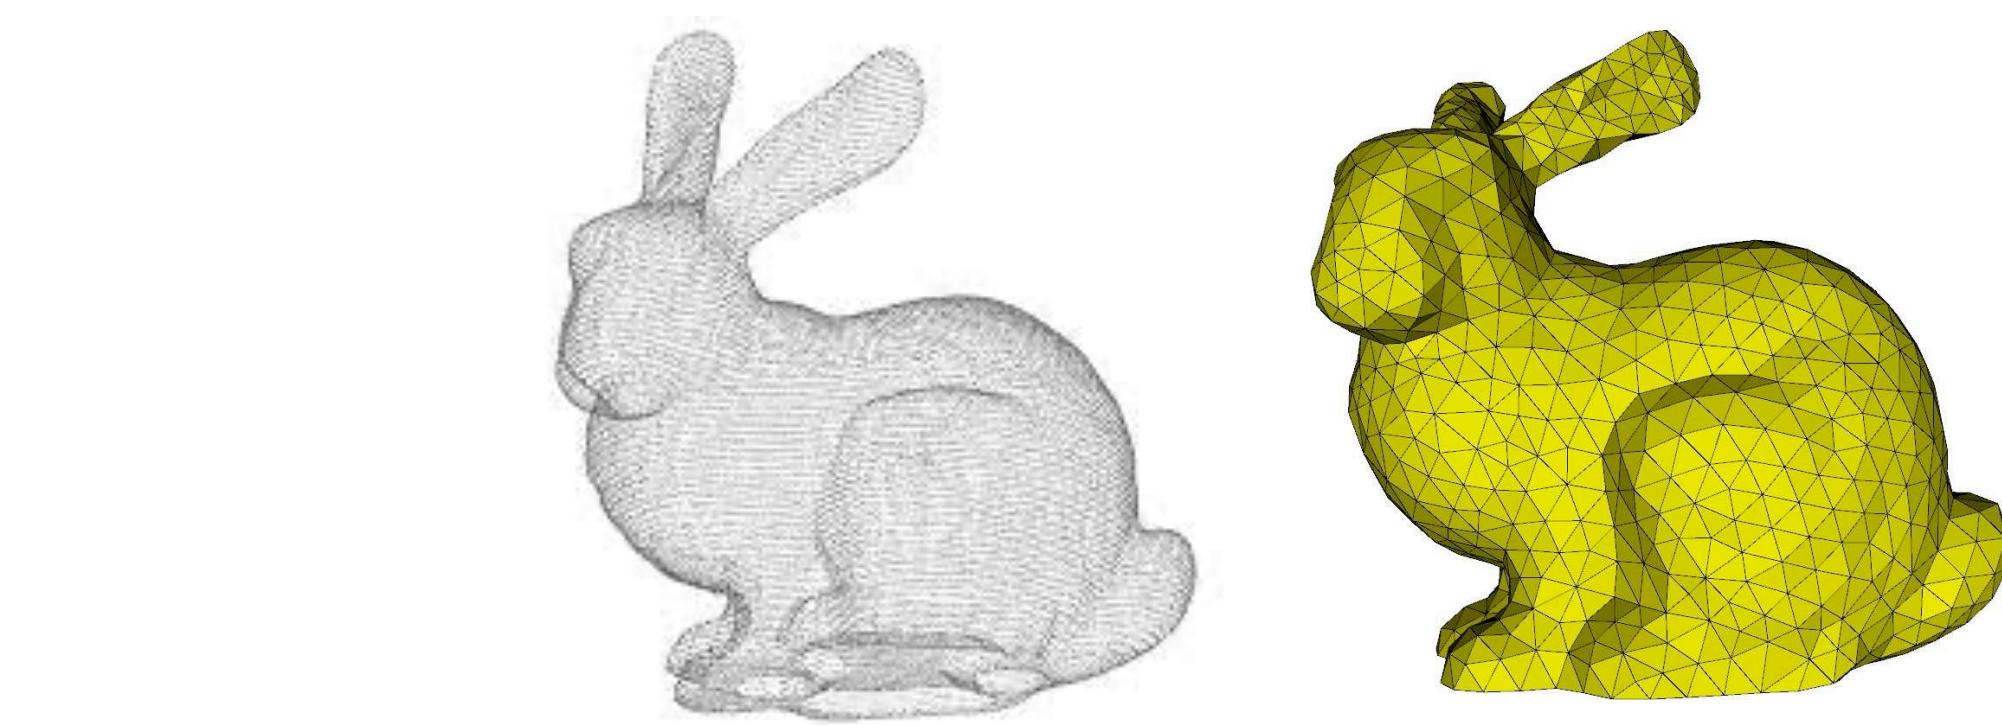
CMU 16-889: Learning for 3D Vision

Prof. Shubham Tulsiani

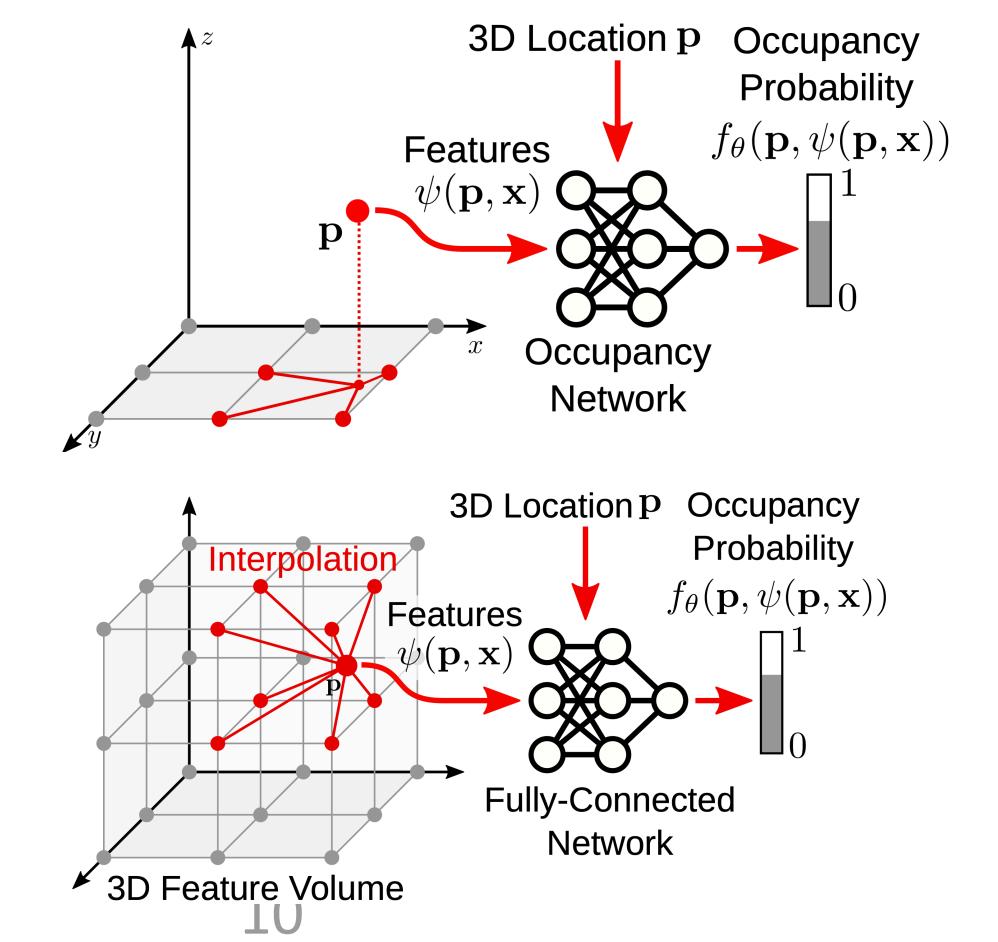
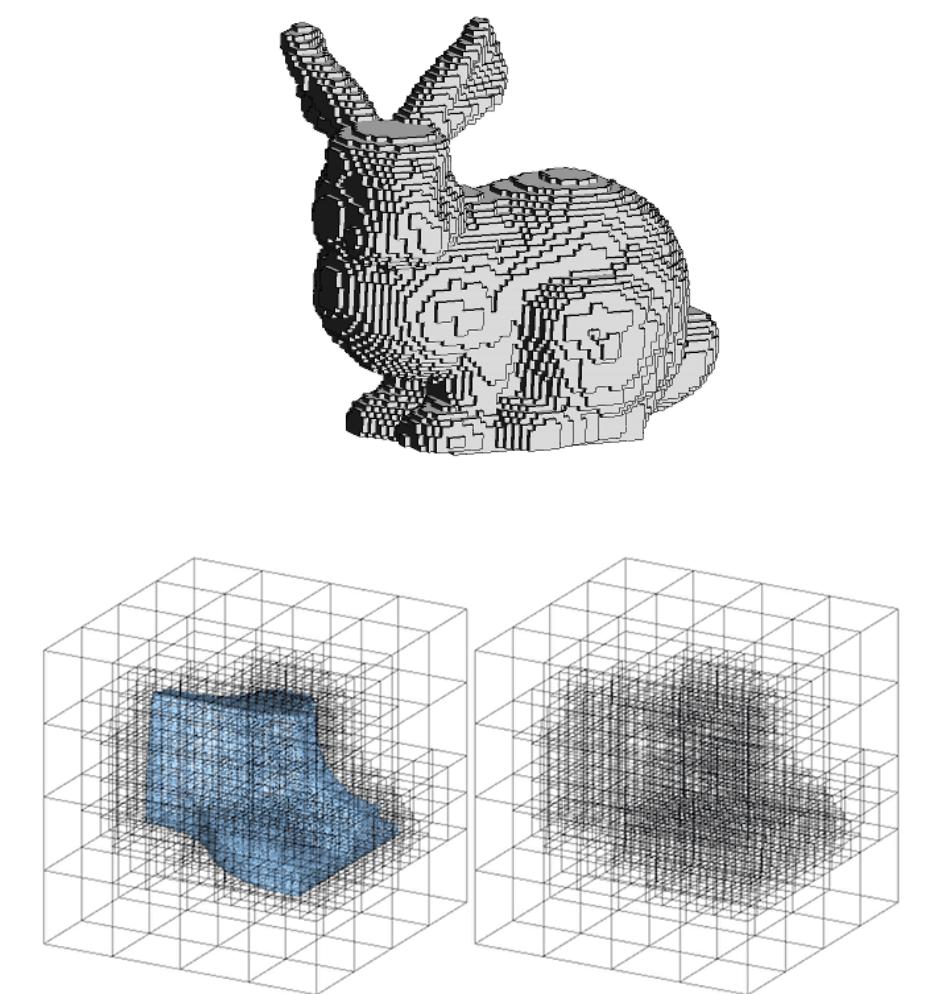
2.5D Representations



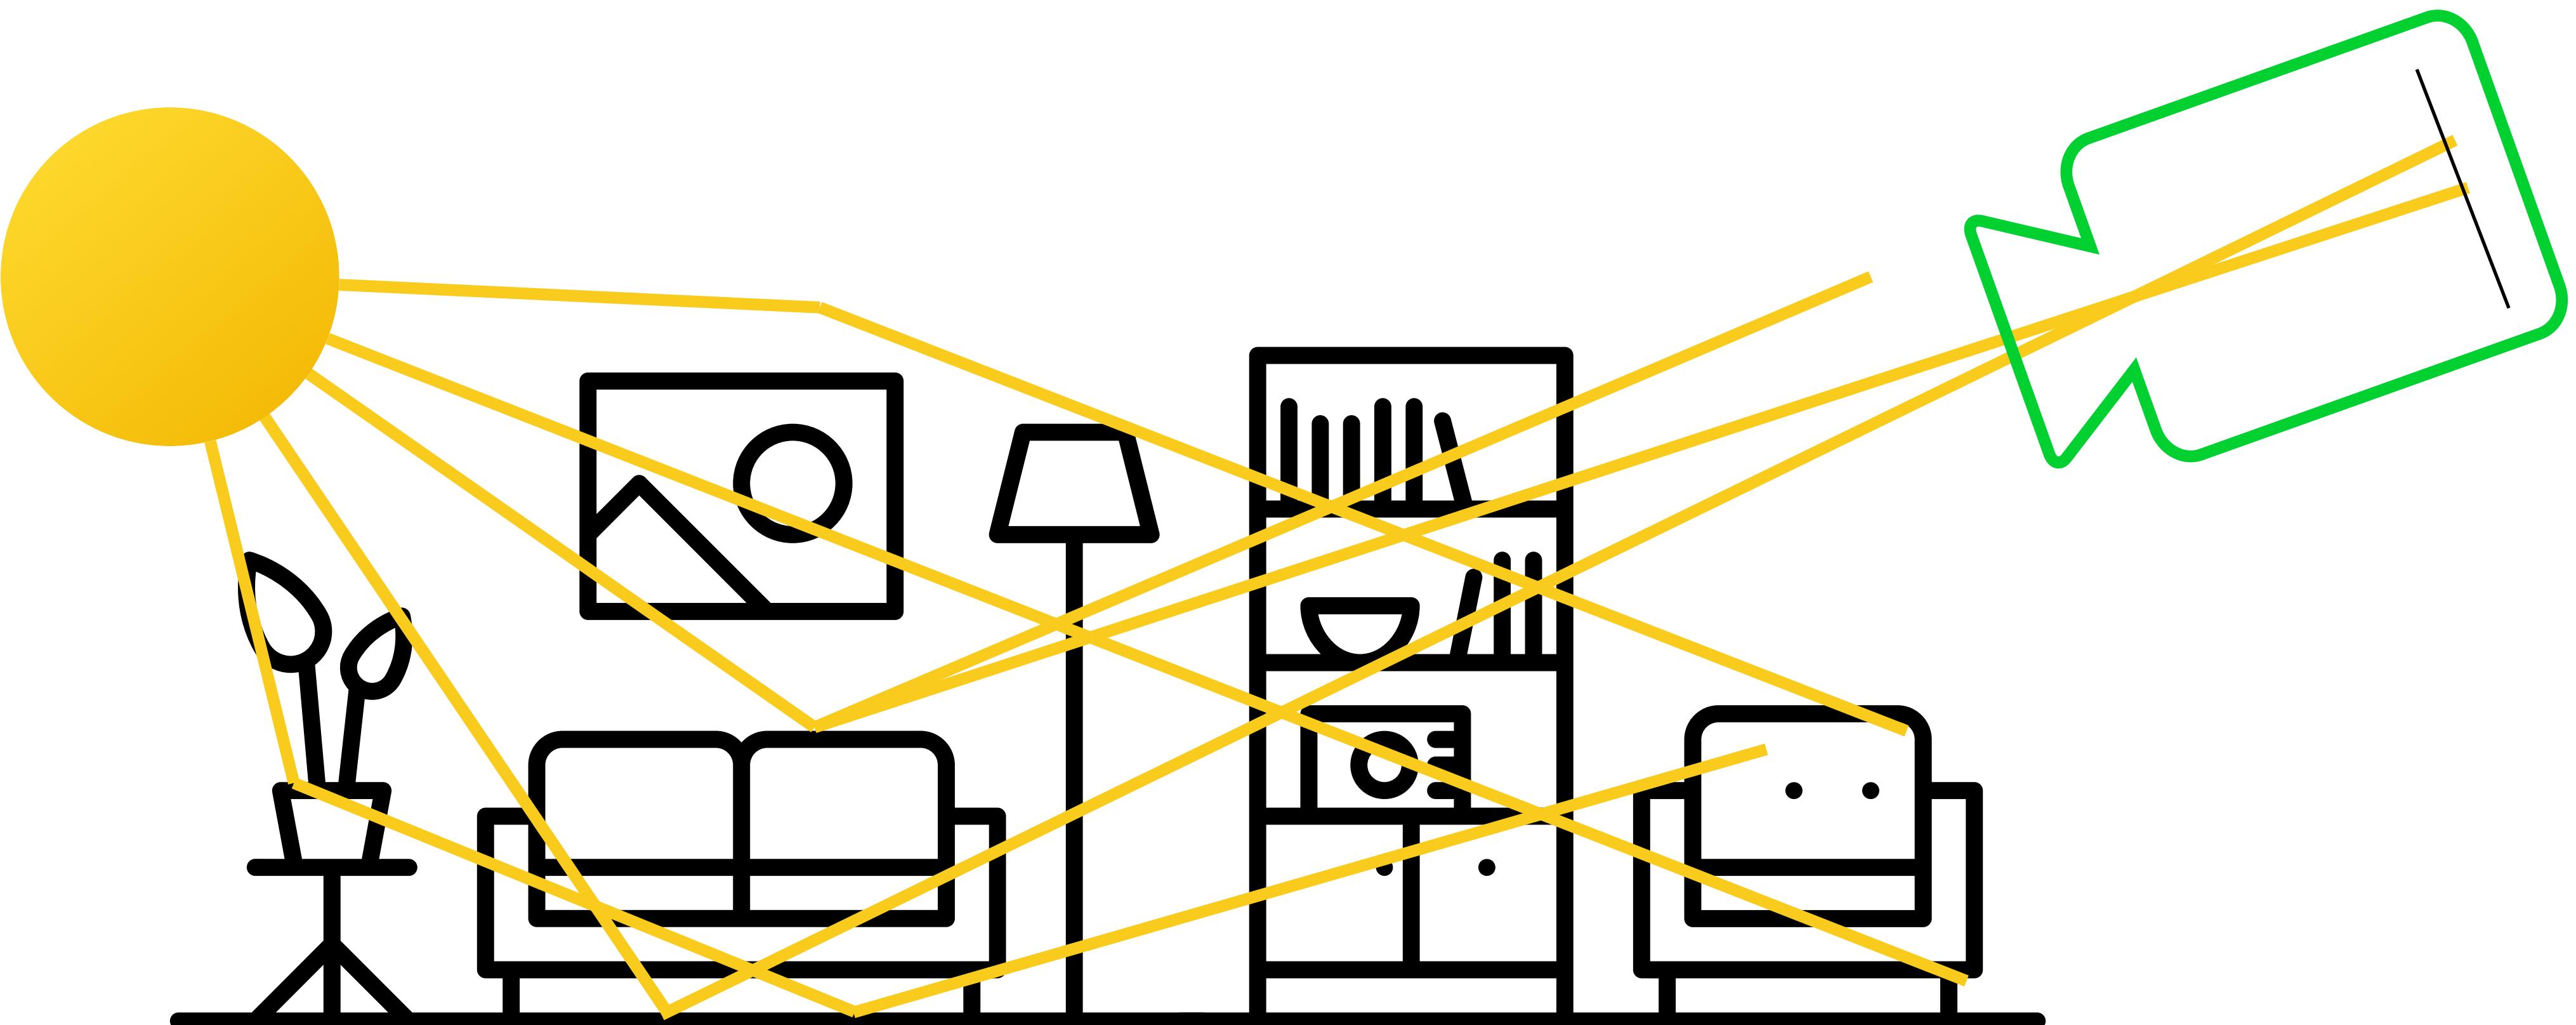
Surface Representations



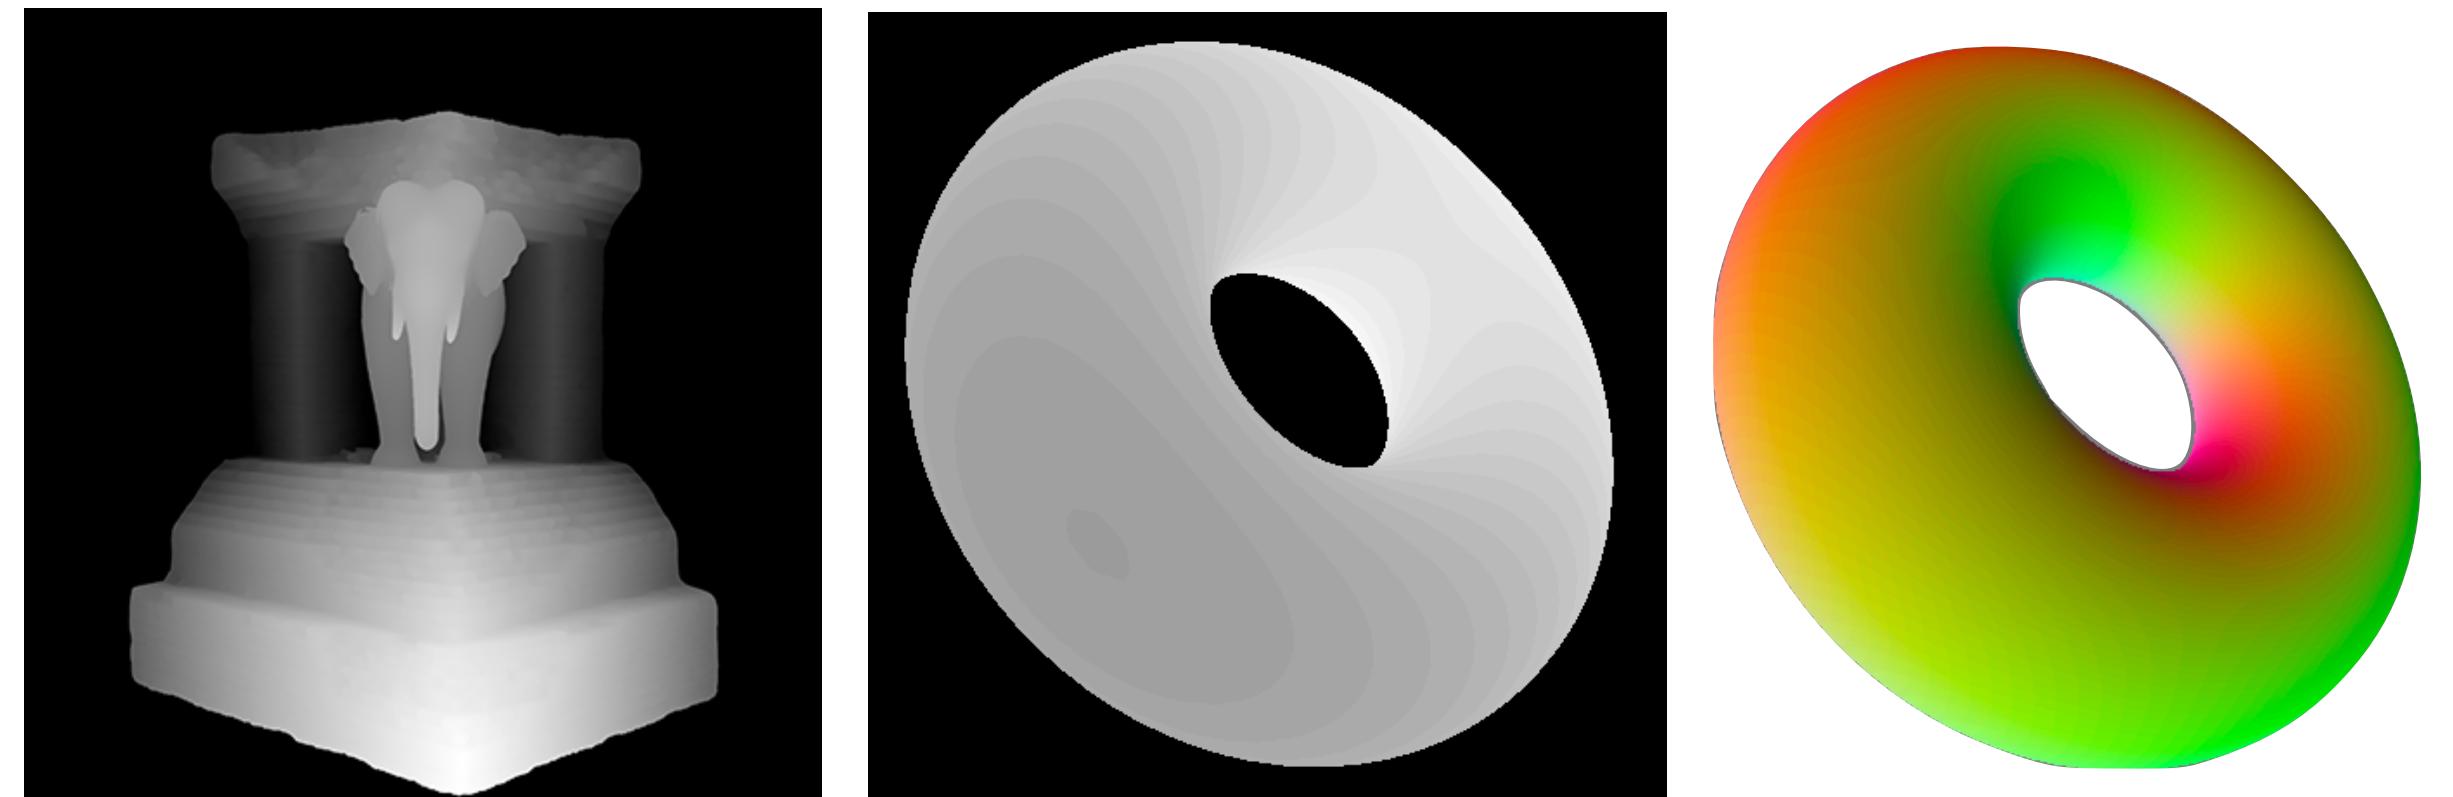
Field Parameterizations



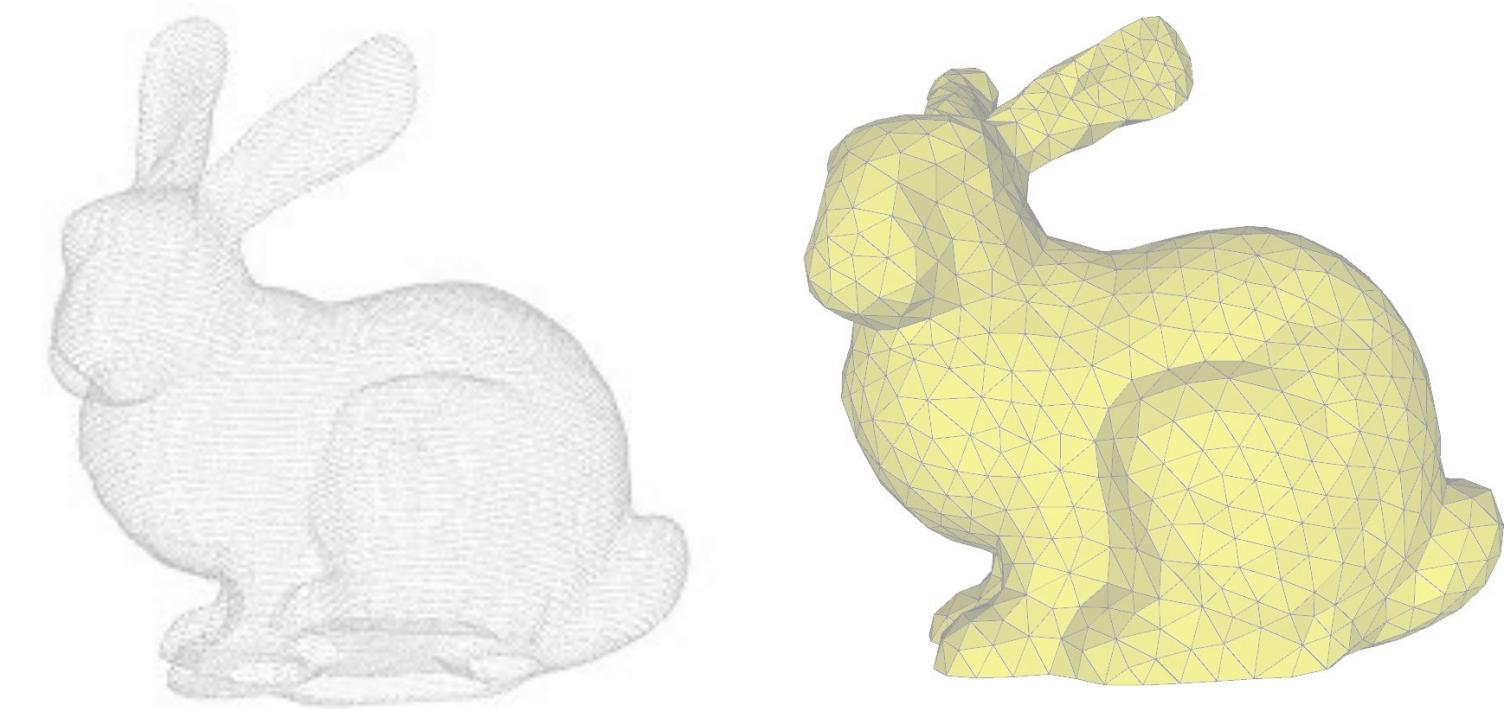
What does the camera see?



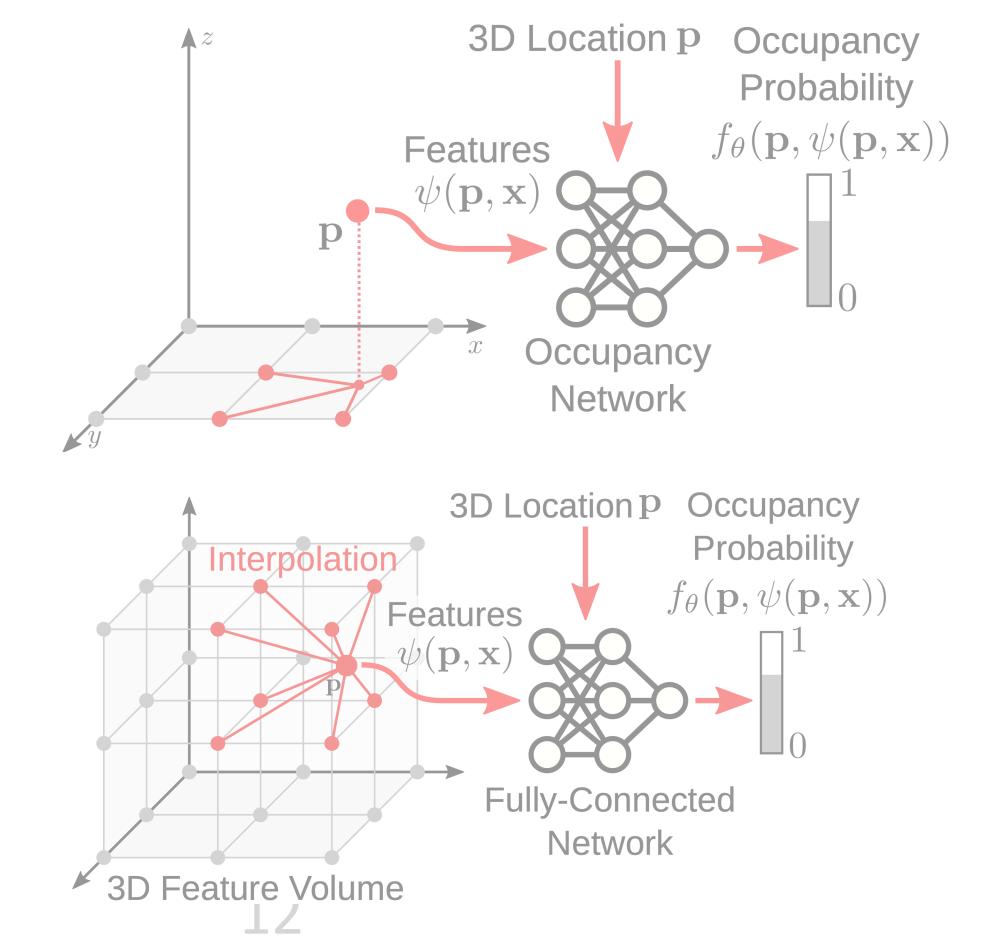
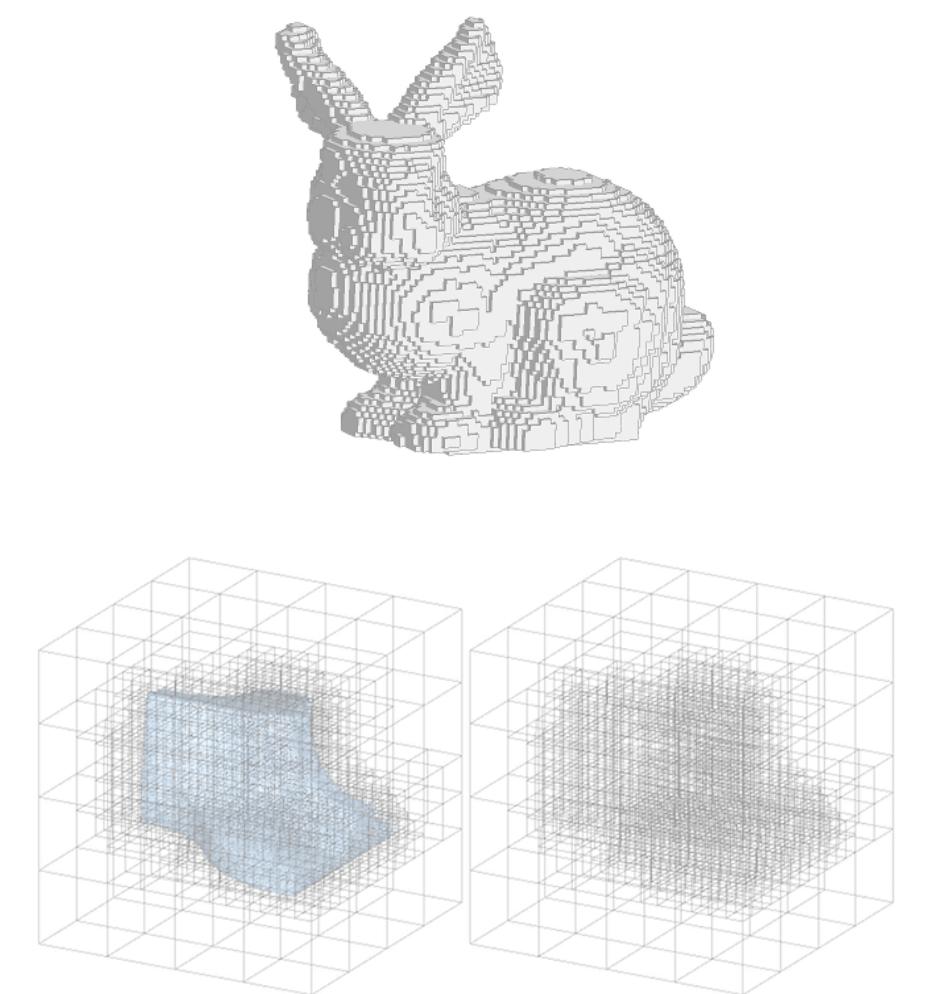
2.5D Representations



Surface Representations



Field Parameterizations

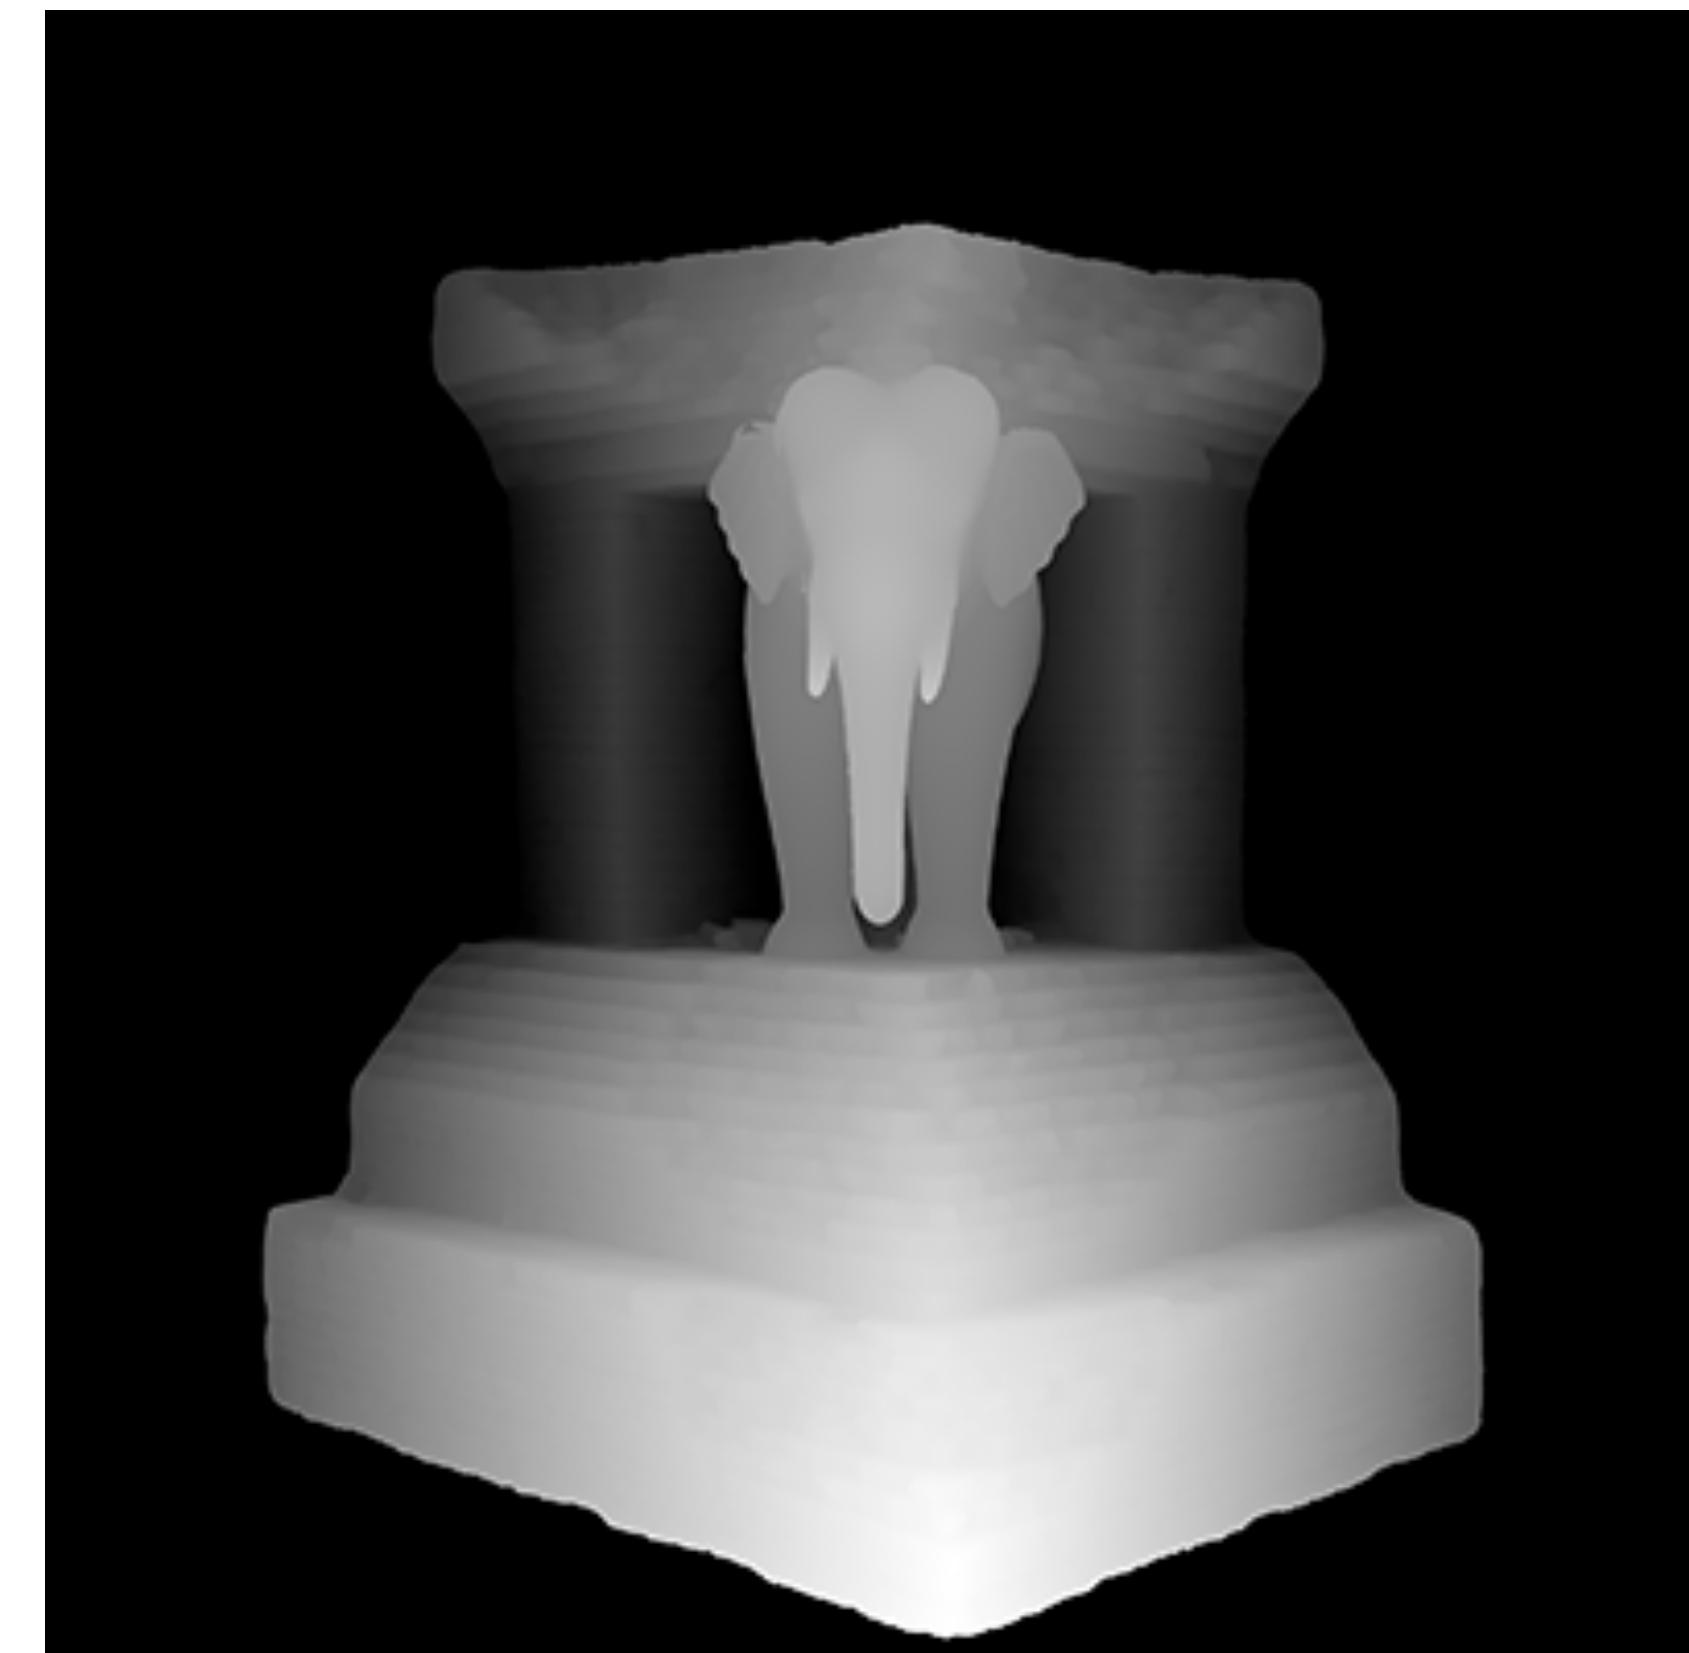


Depth Maps



Common modality, captured with Stereo, LIDAR,

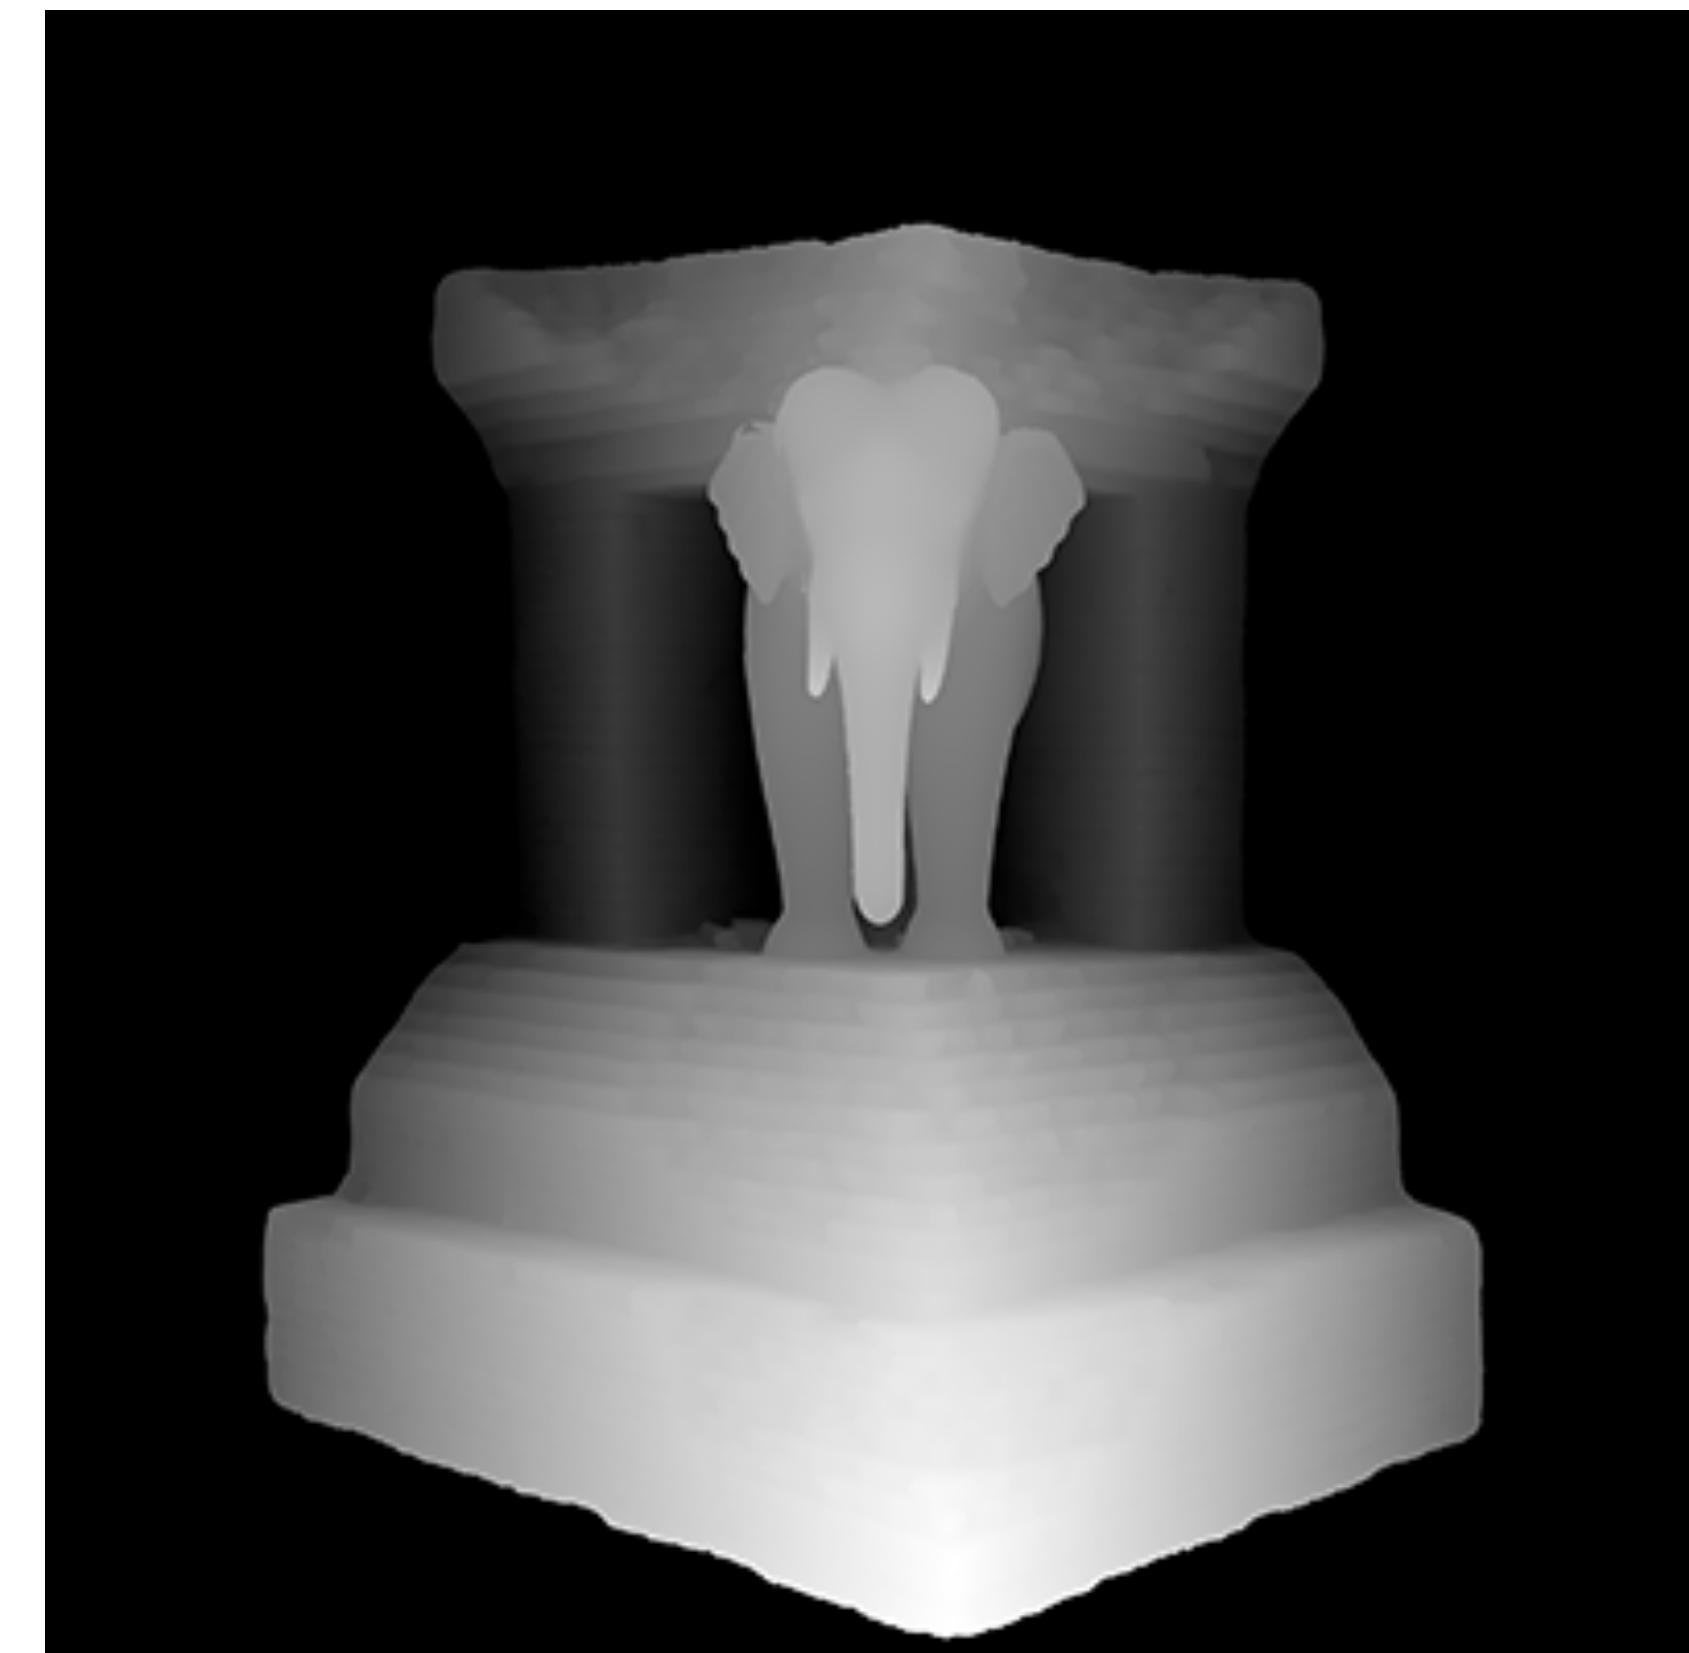
Depth Maps



An **image** that represents
how far each pixel \mathbf{p} is

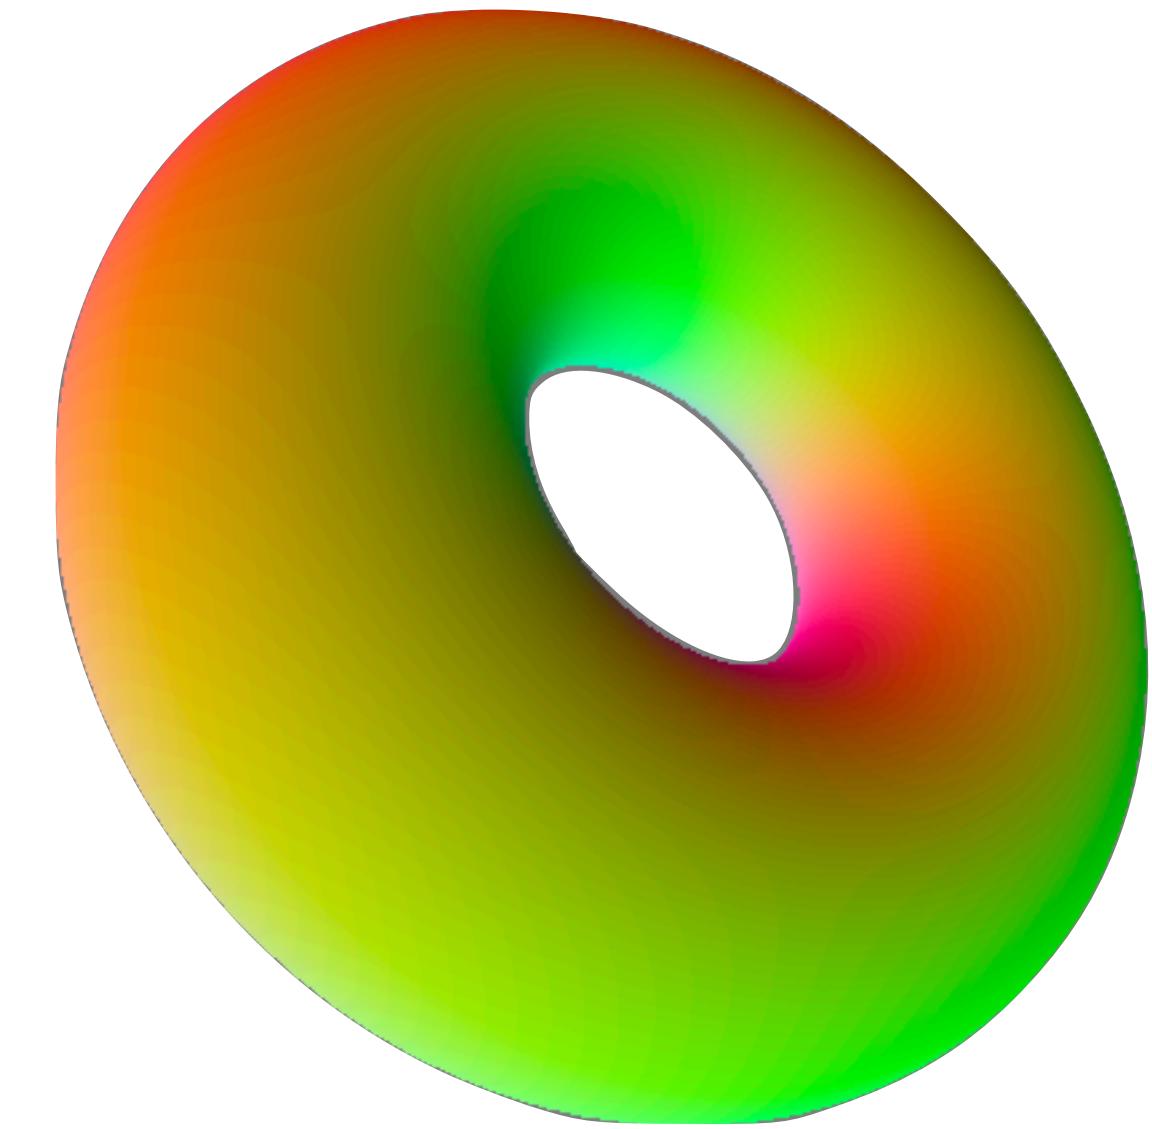
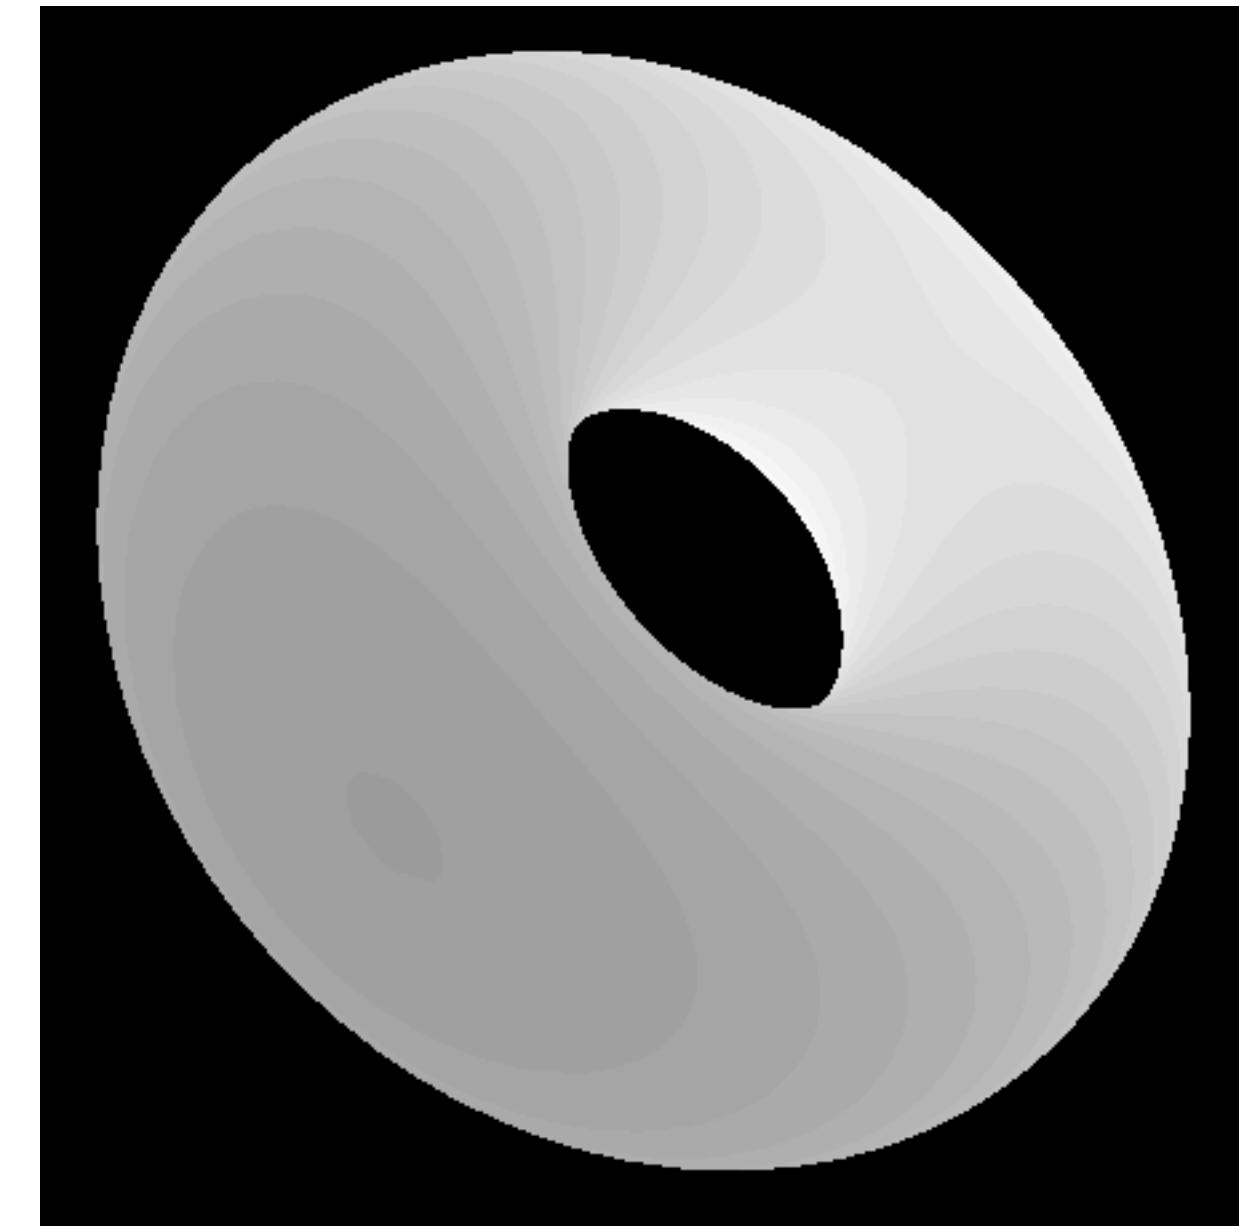
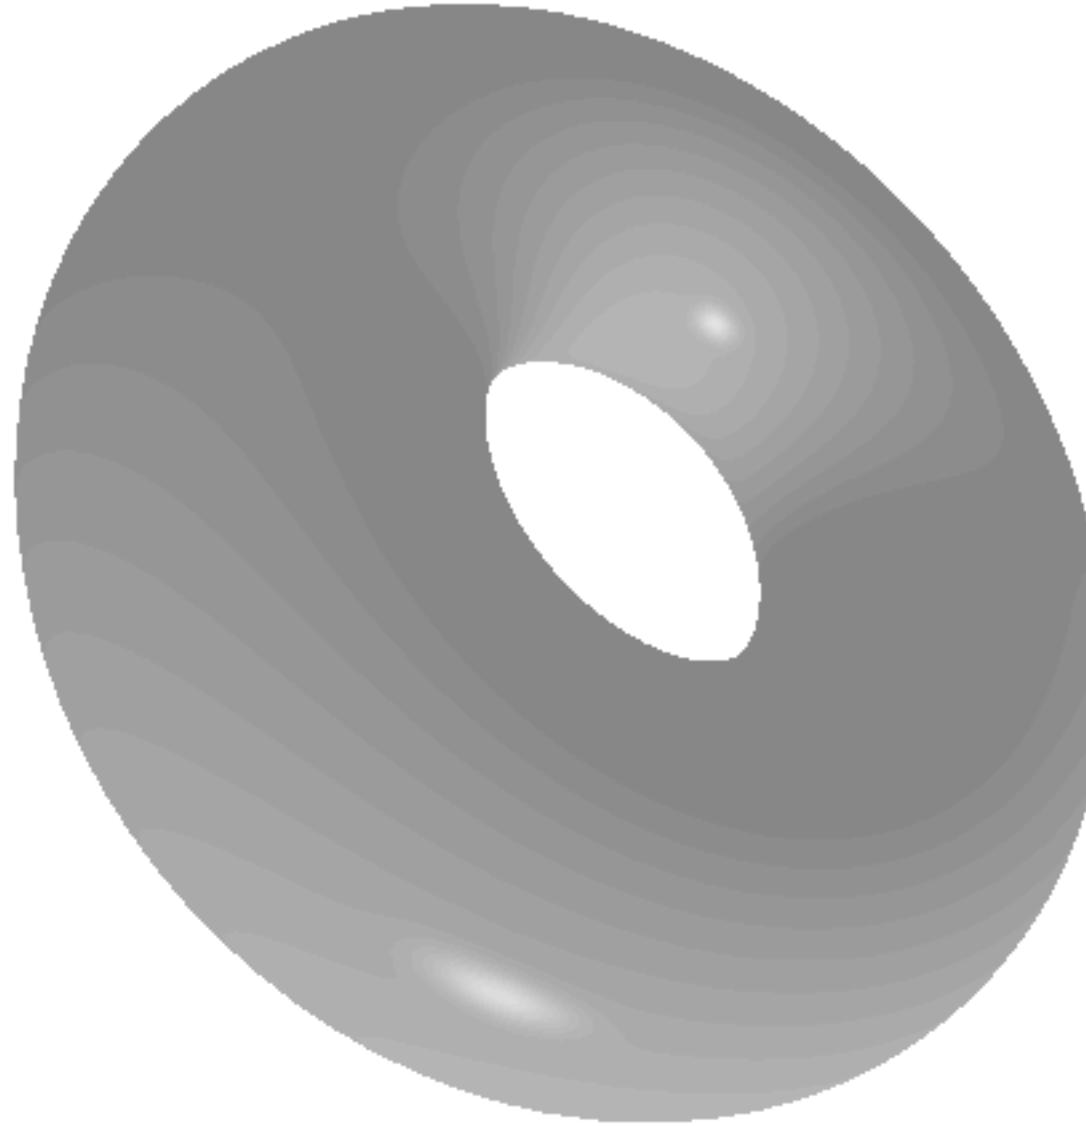
$$D[\mathbf{p}] \in \mathbb{R}^+$$

Depth Maps



The ‘standard’ image processing tools can be used

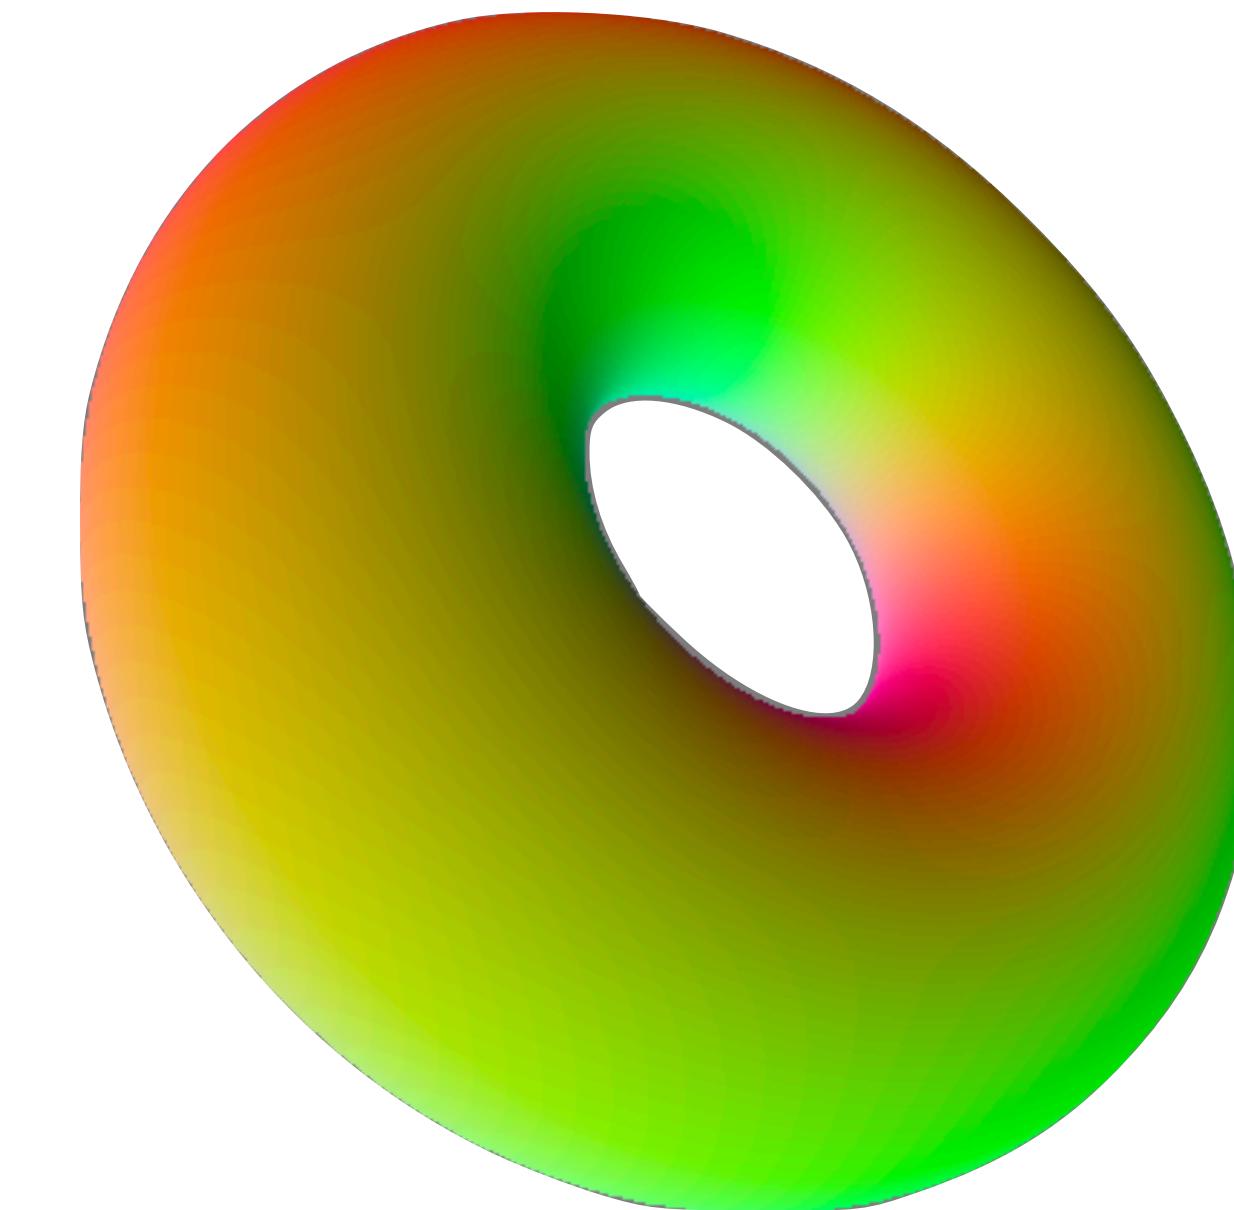
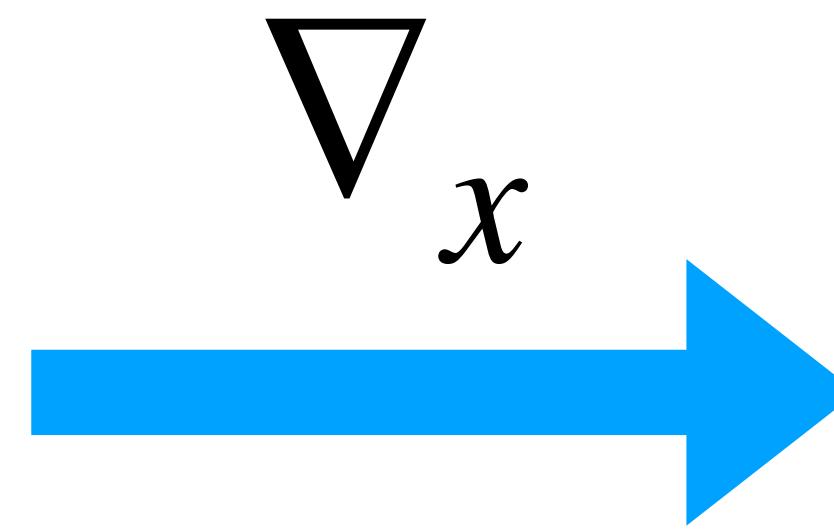
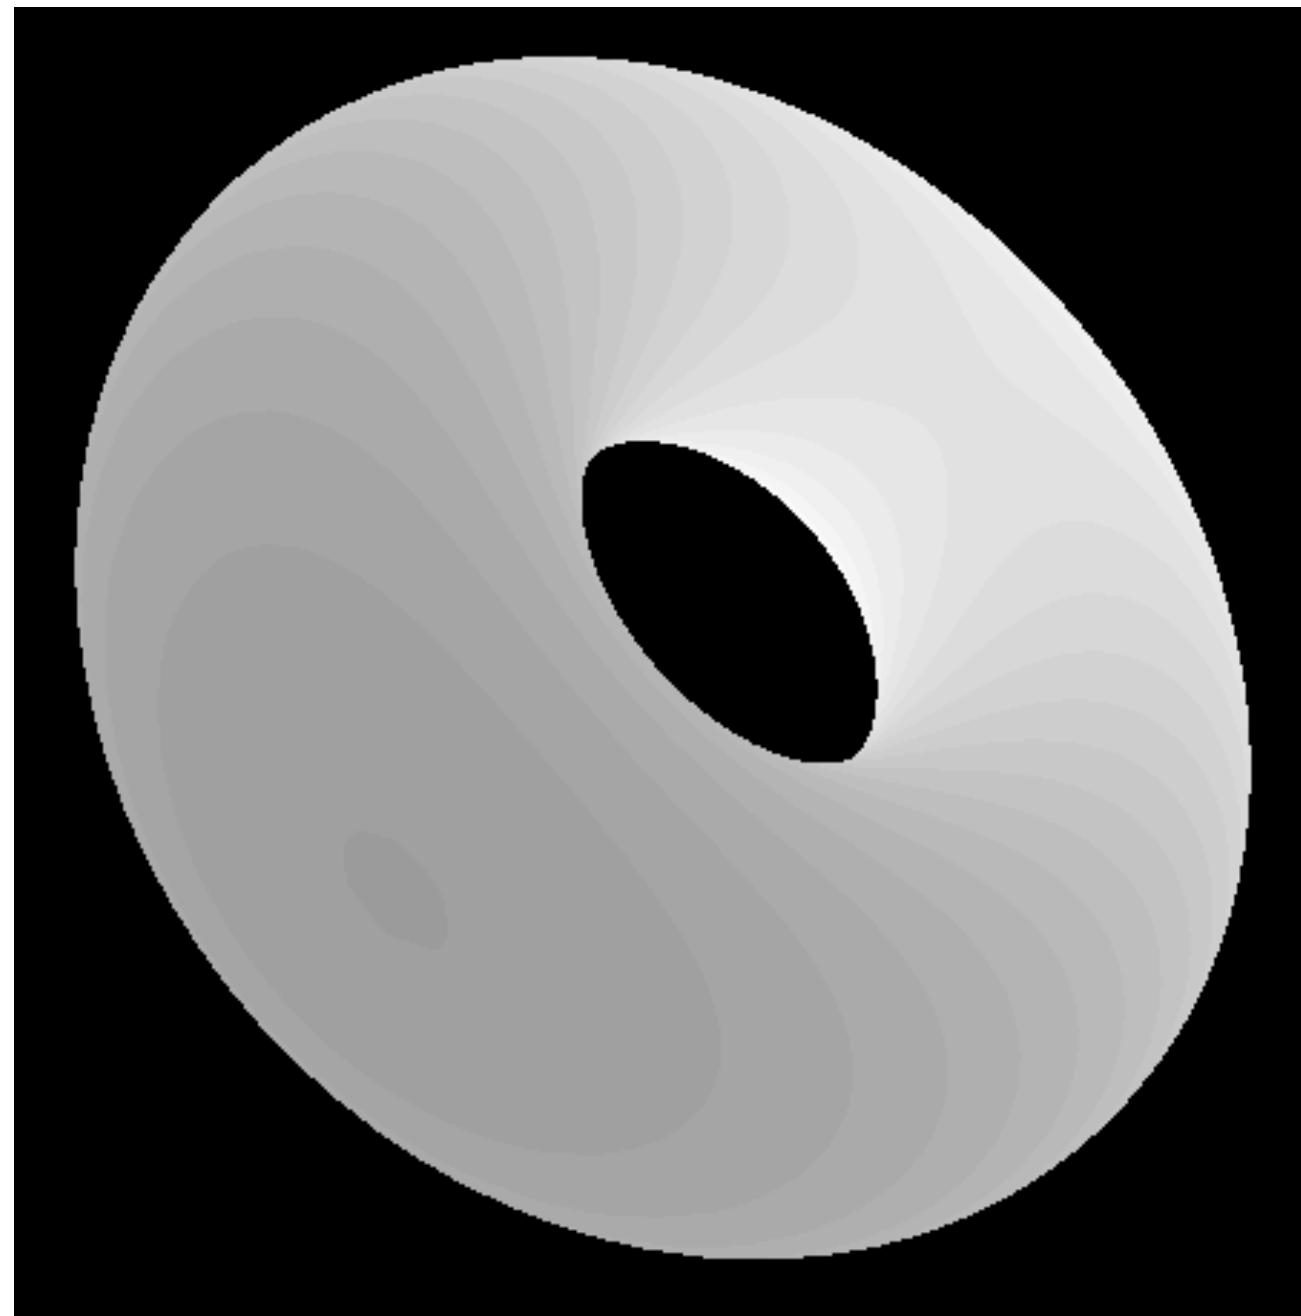
Depth Maps and Normal Maps



$$N[p] \in \mathbb{S}^2$$

Can extend to capture other surface properties e.g. surface normals

Normal map is spatial derivative of depth map



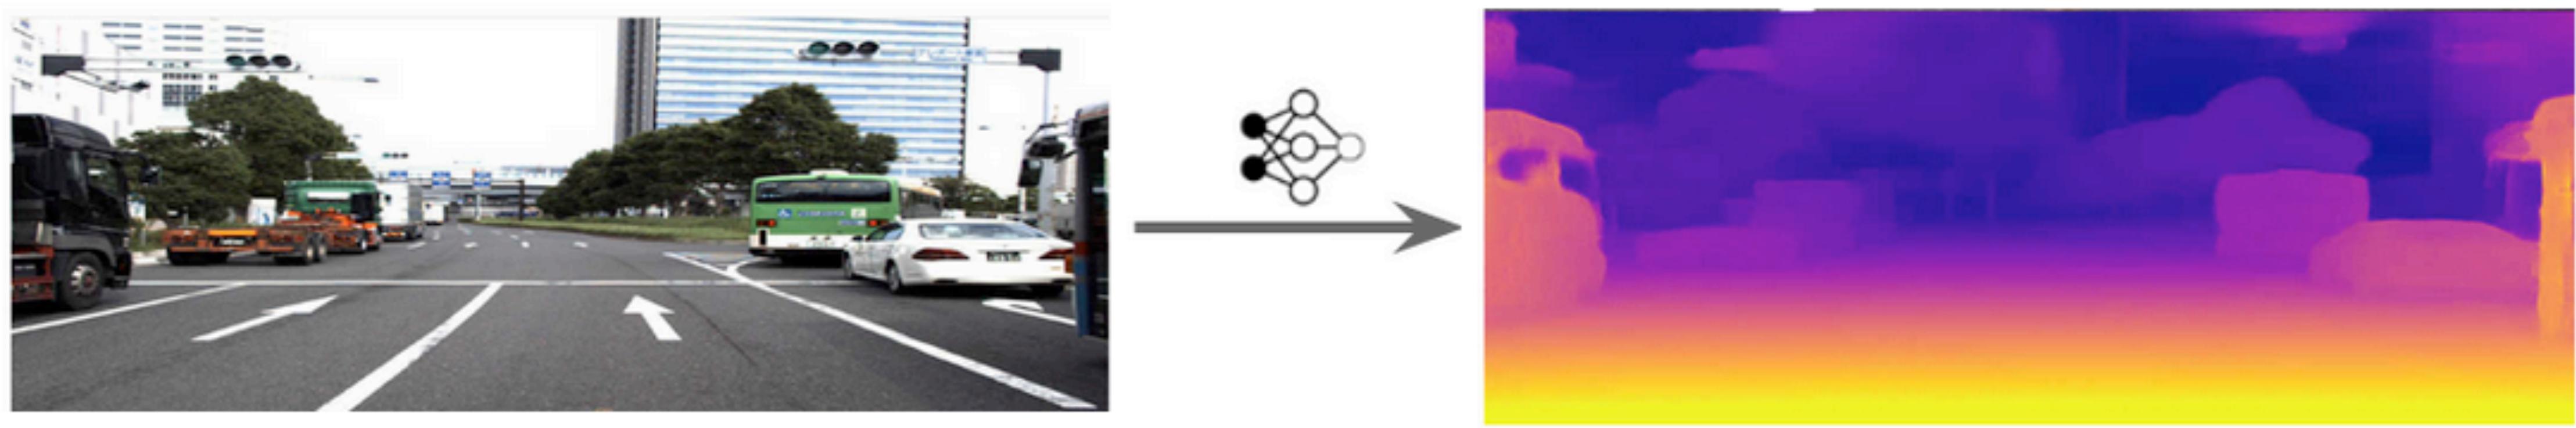
$$D : \mathbb{R}^2 \rightarrow \mathbb{R}^+$$

$$D(\mathbf{x}_{pix}) = d$$

$$N : \mathbb{R}^2 \rightarrow \mathbb{R}^3$$

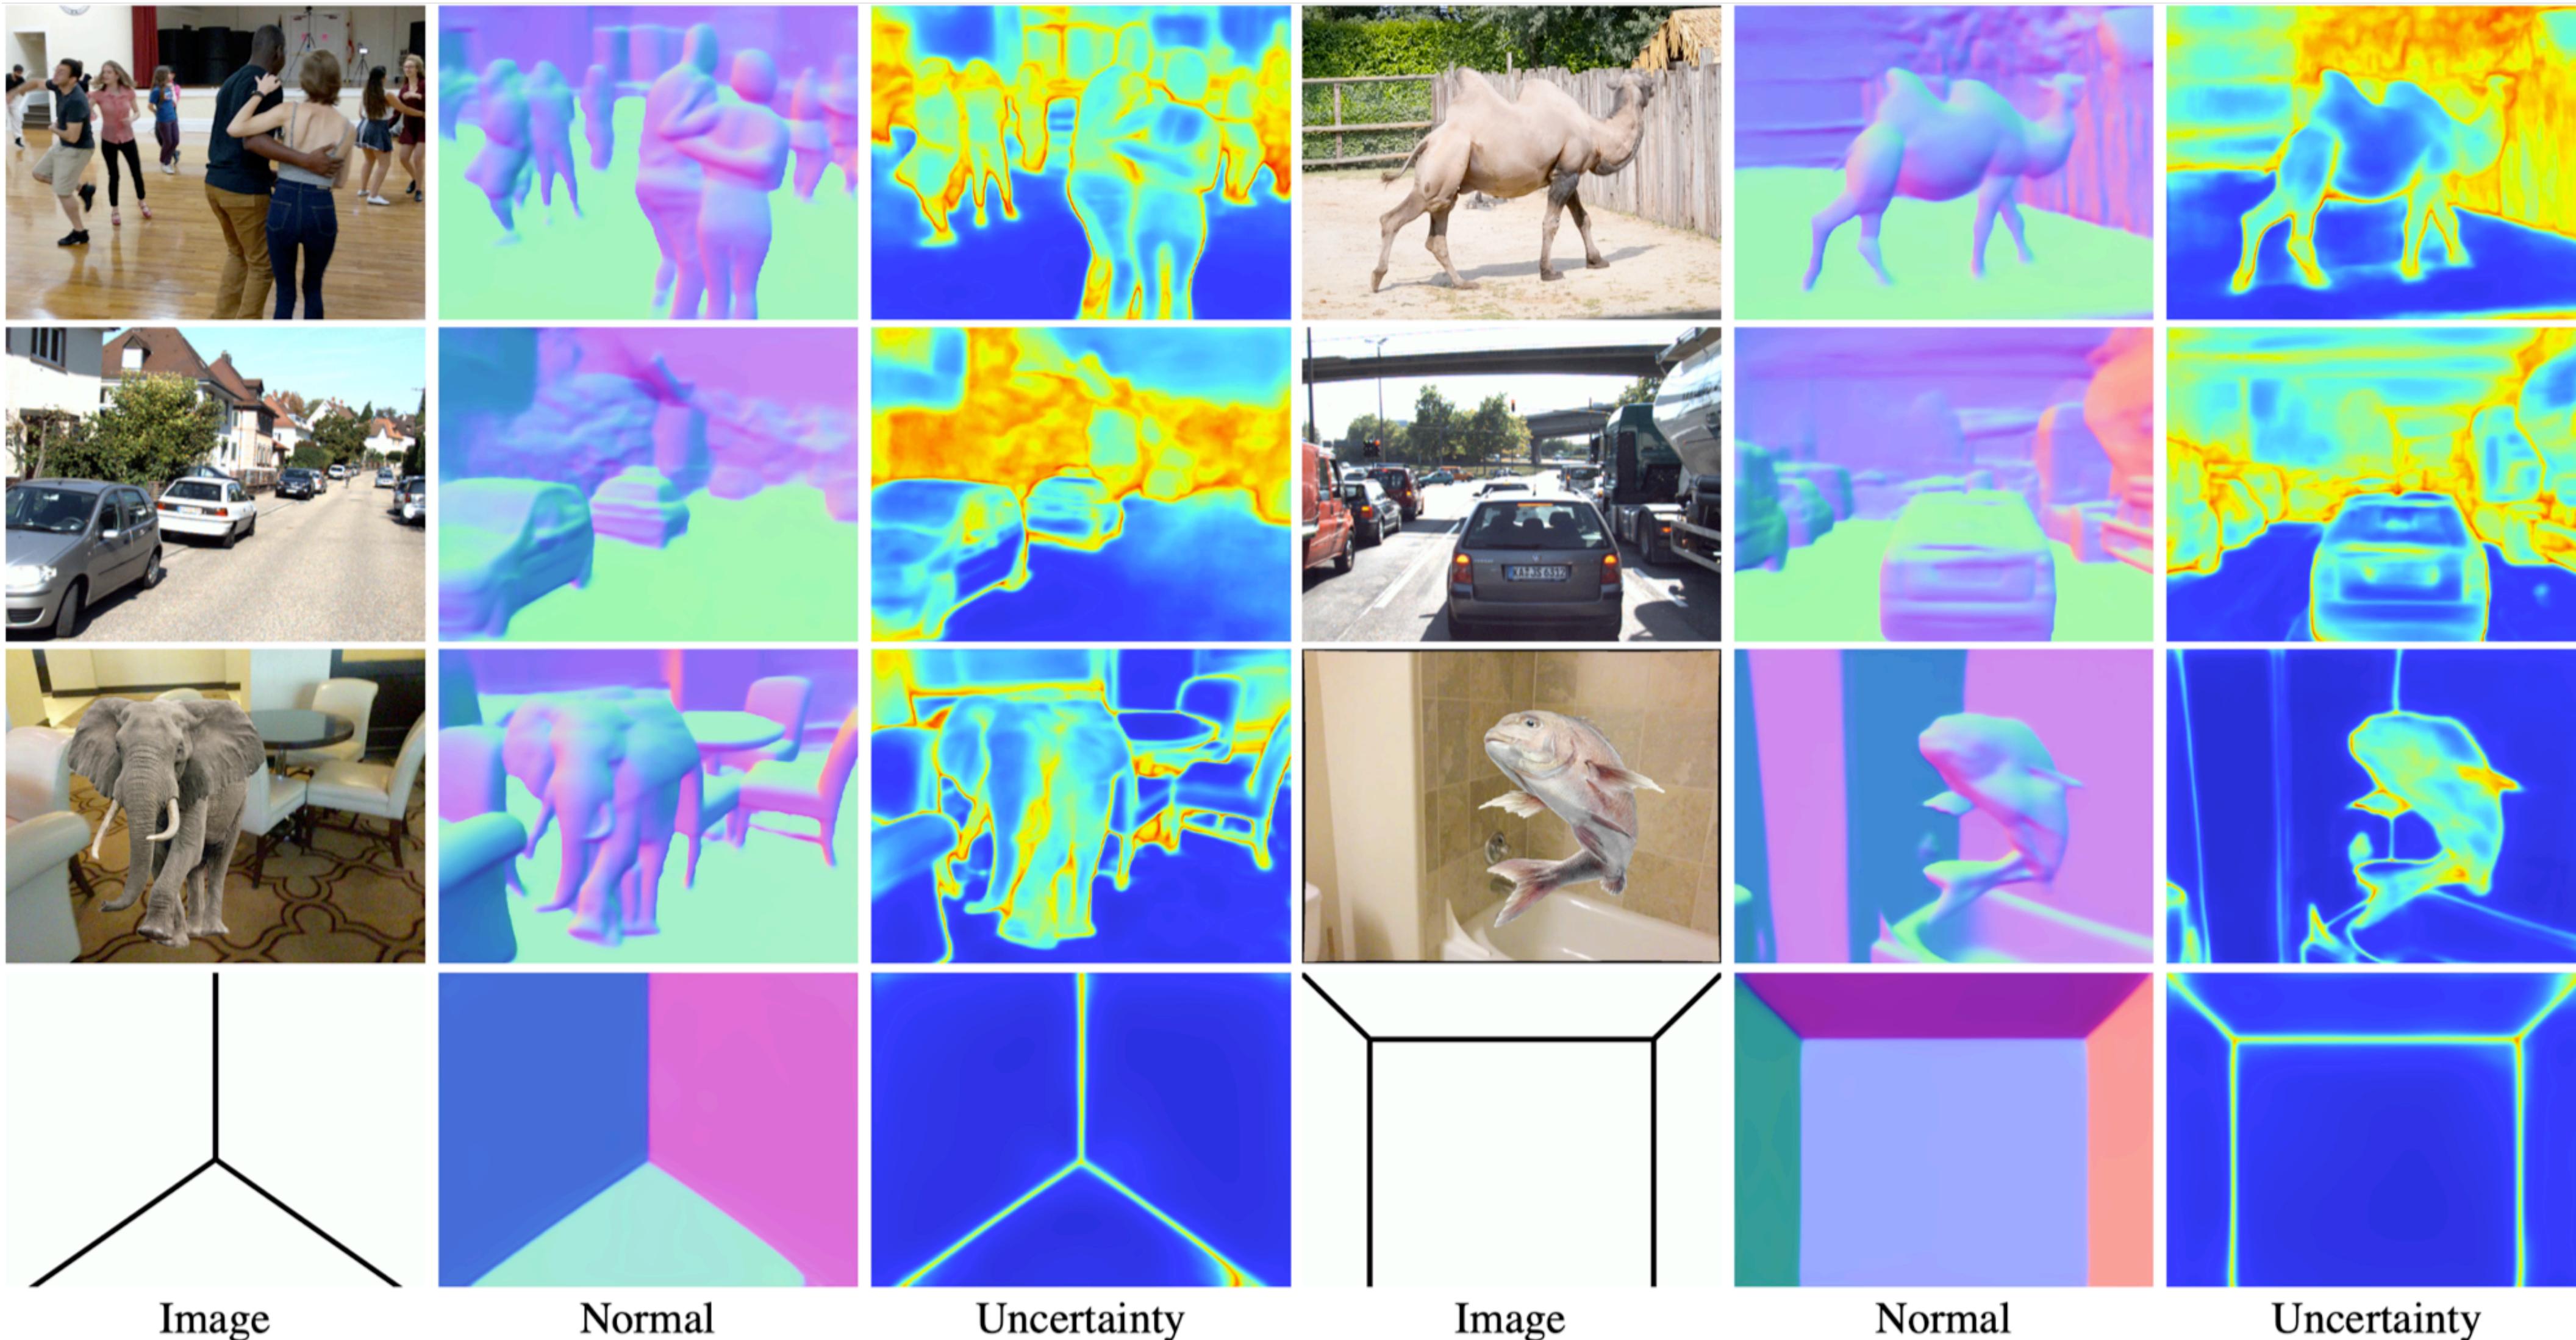
$$N(\mathbf{x}_{pix}) = \nabla_{\mathbf{x}} D(\mathbf{x}_{pix}) = \mathbf{n}$$

Supervised Monocular Depth Estimation

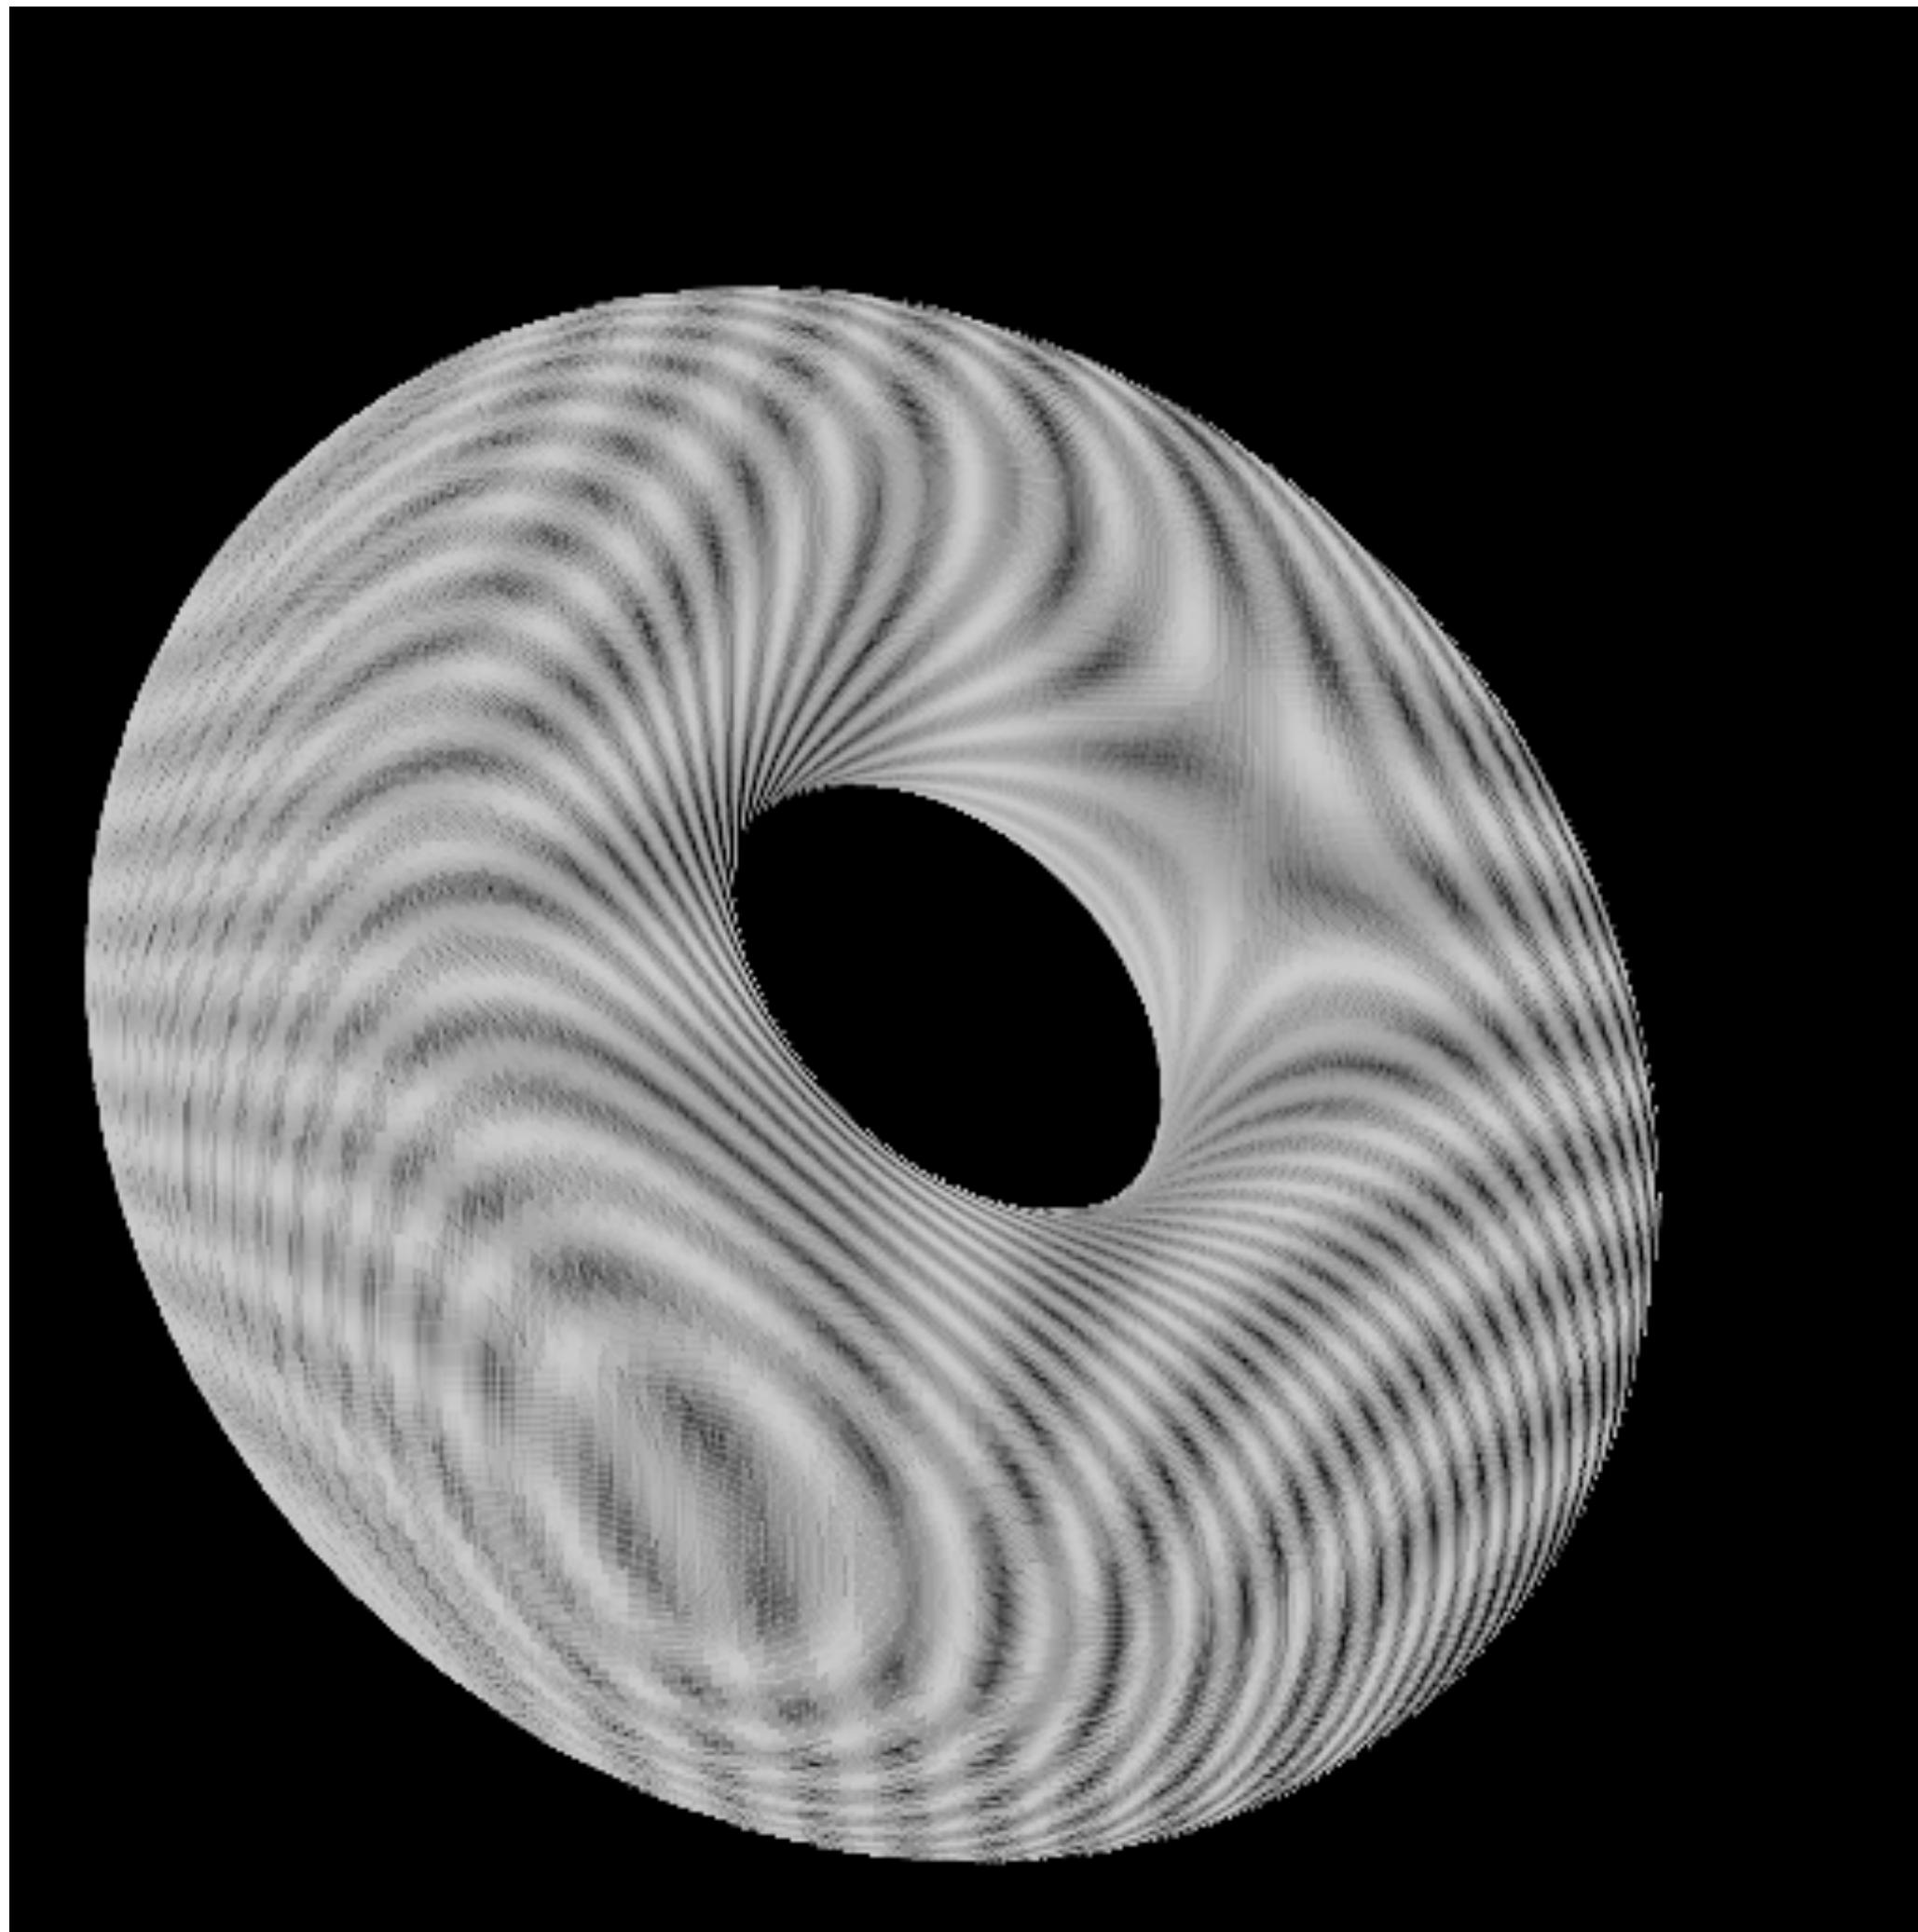


<https://medium.com/toyotaresearch/self-supervised-learning-in-depth-part-1-of-2-74825baaaa04>

It's easier to predict normals from pixels than depth



Depth Maps – Representing the Visible



‘2.5D’ representation –
properties associated
to image pixels

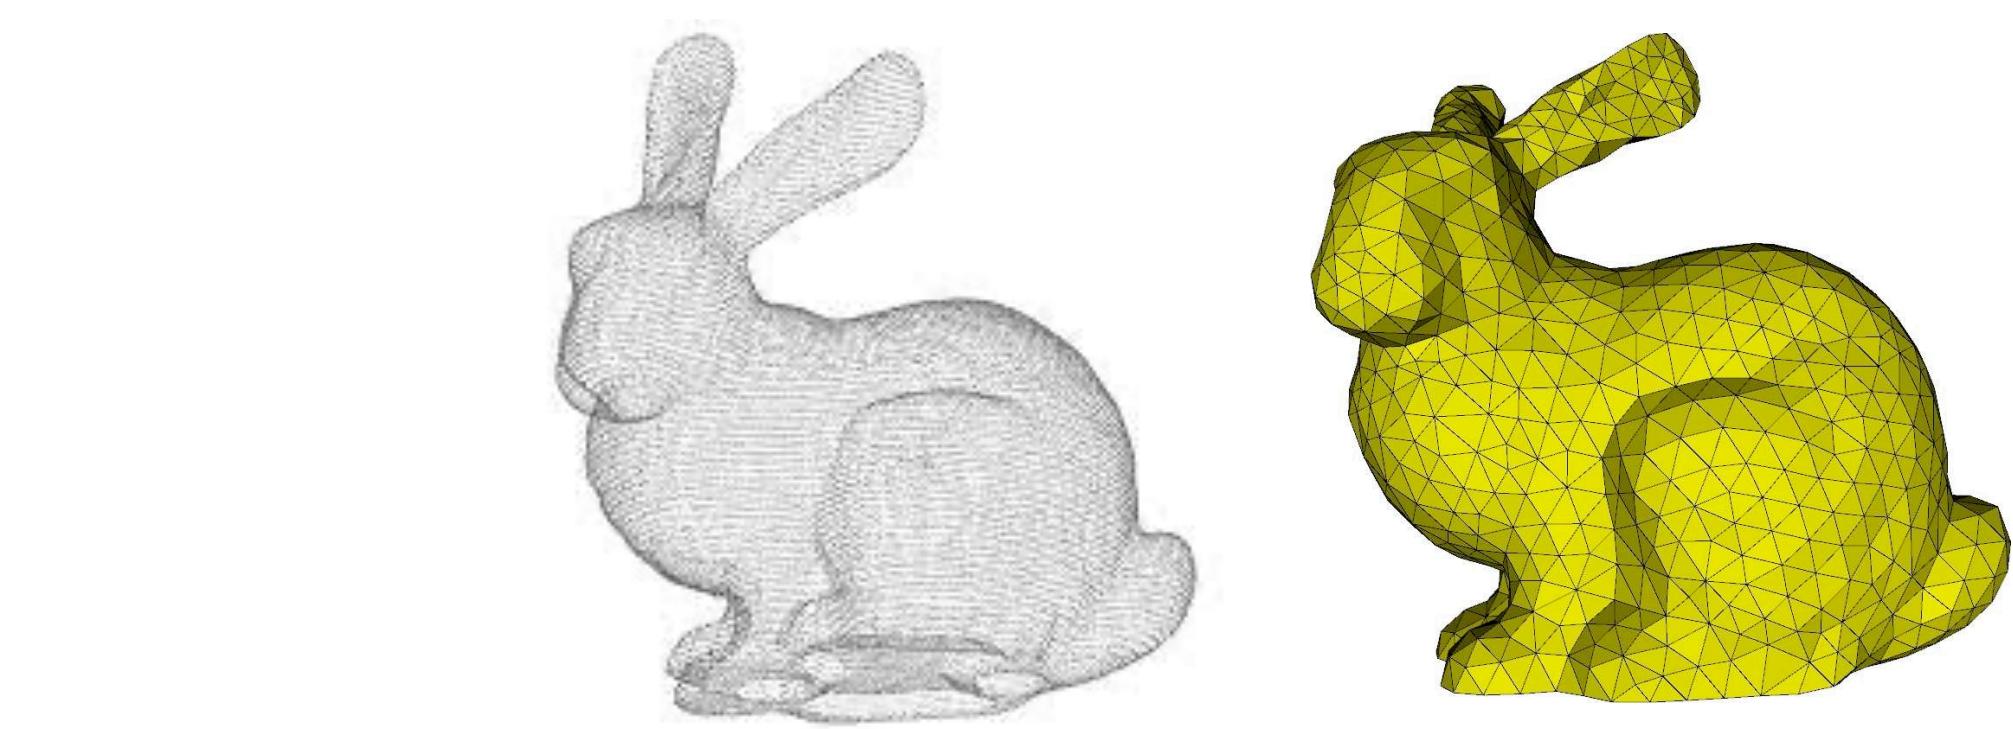
Does not capture the ‘full’
3D structure

can we relax this ‘image-
dependence’ constraint?

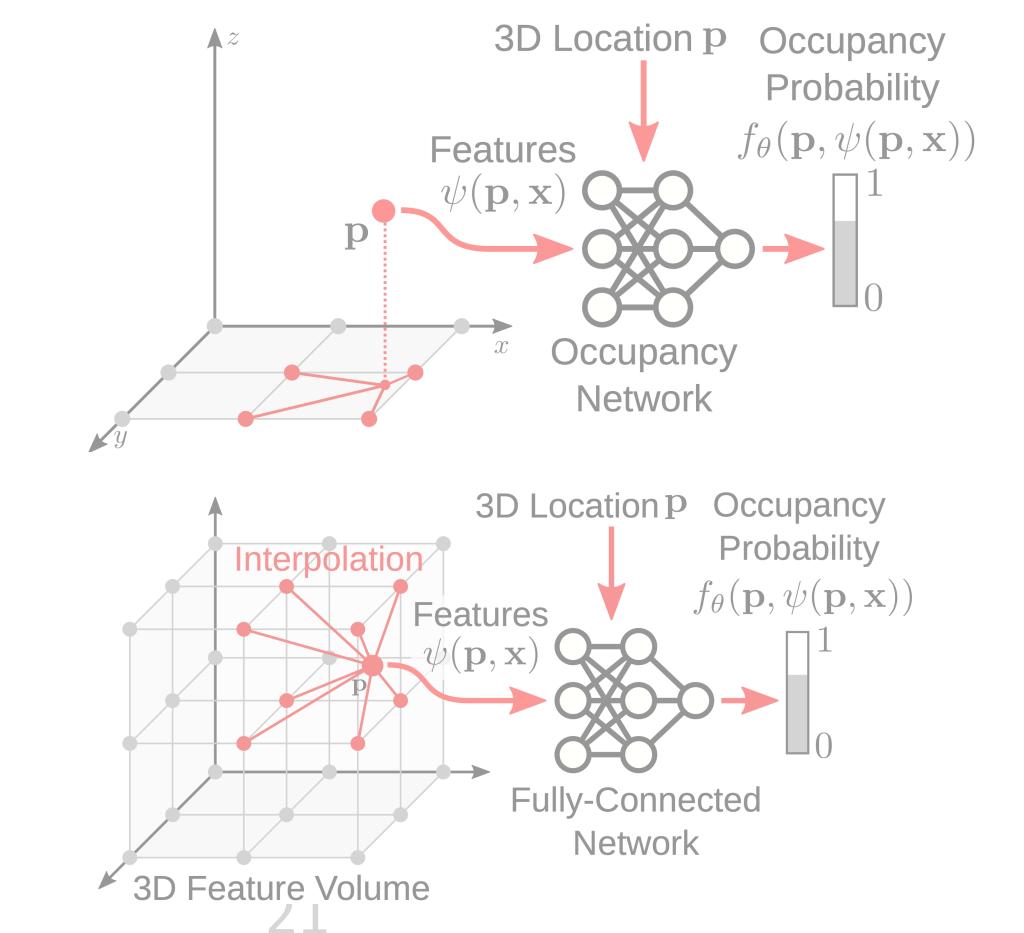
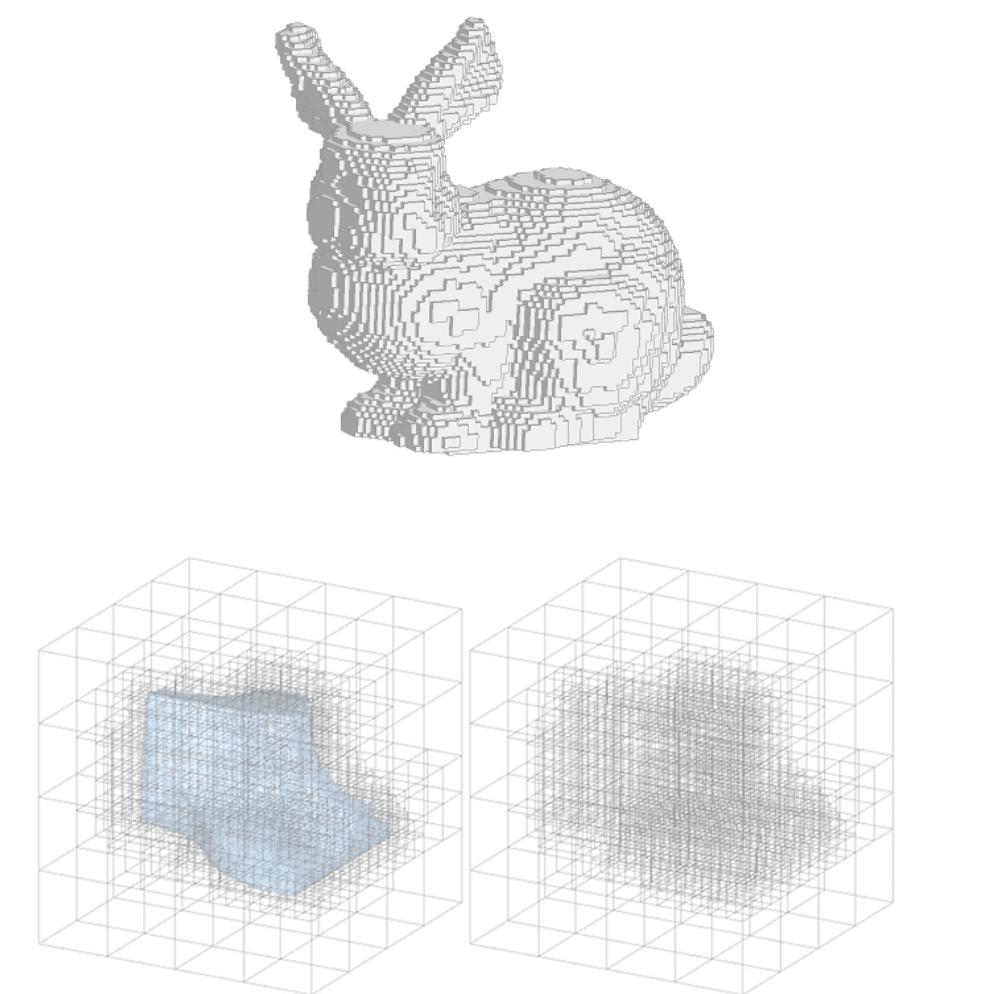
2.5D Representations



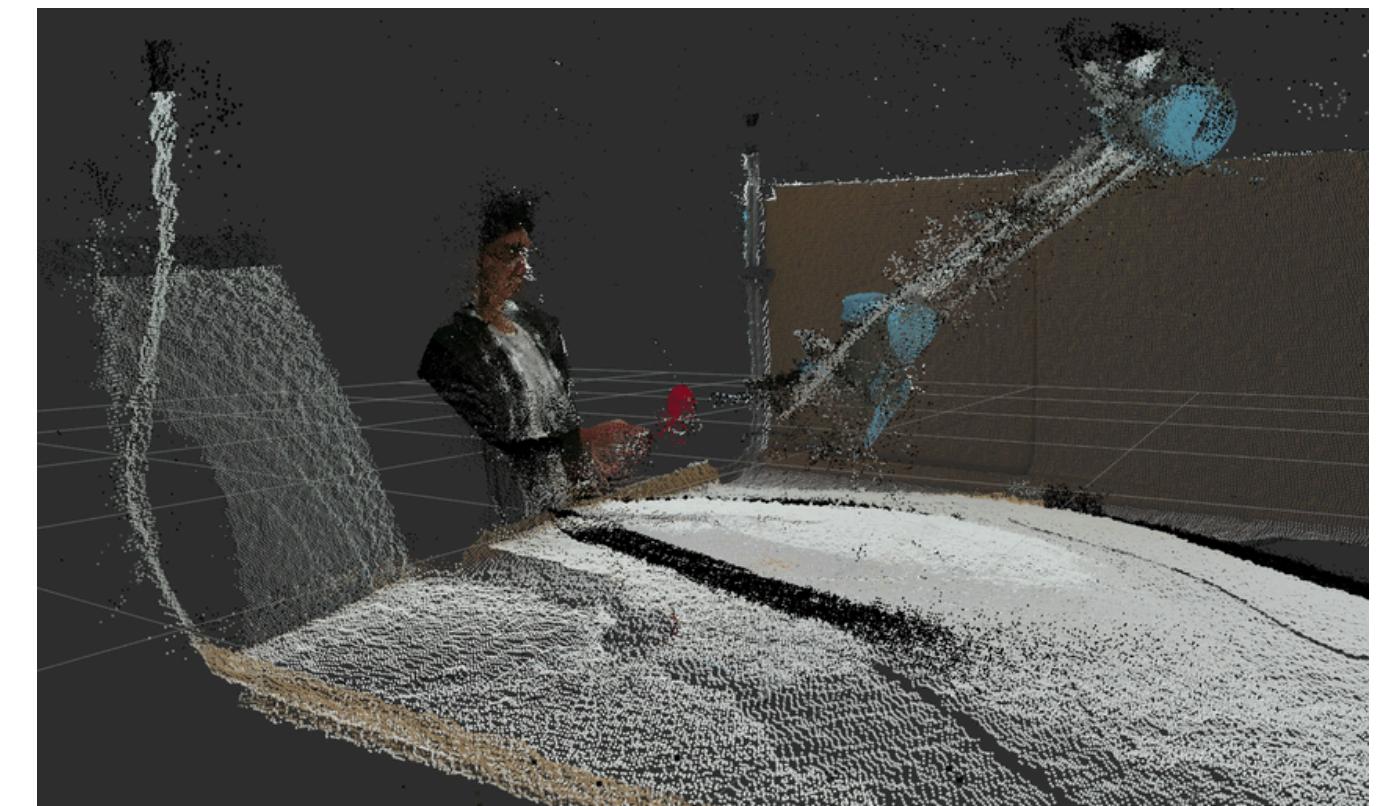
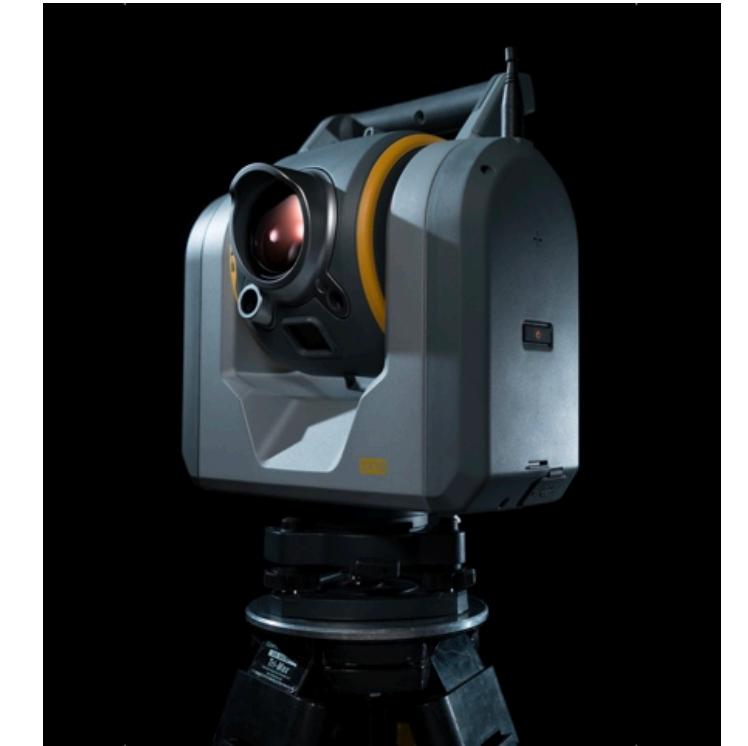
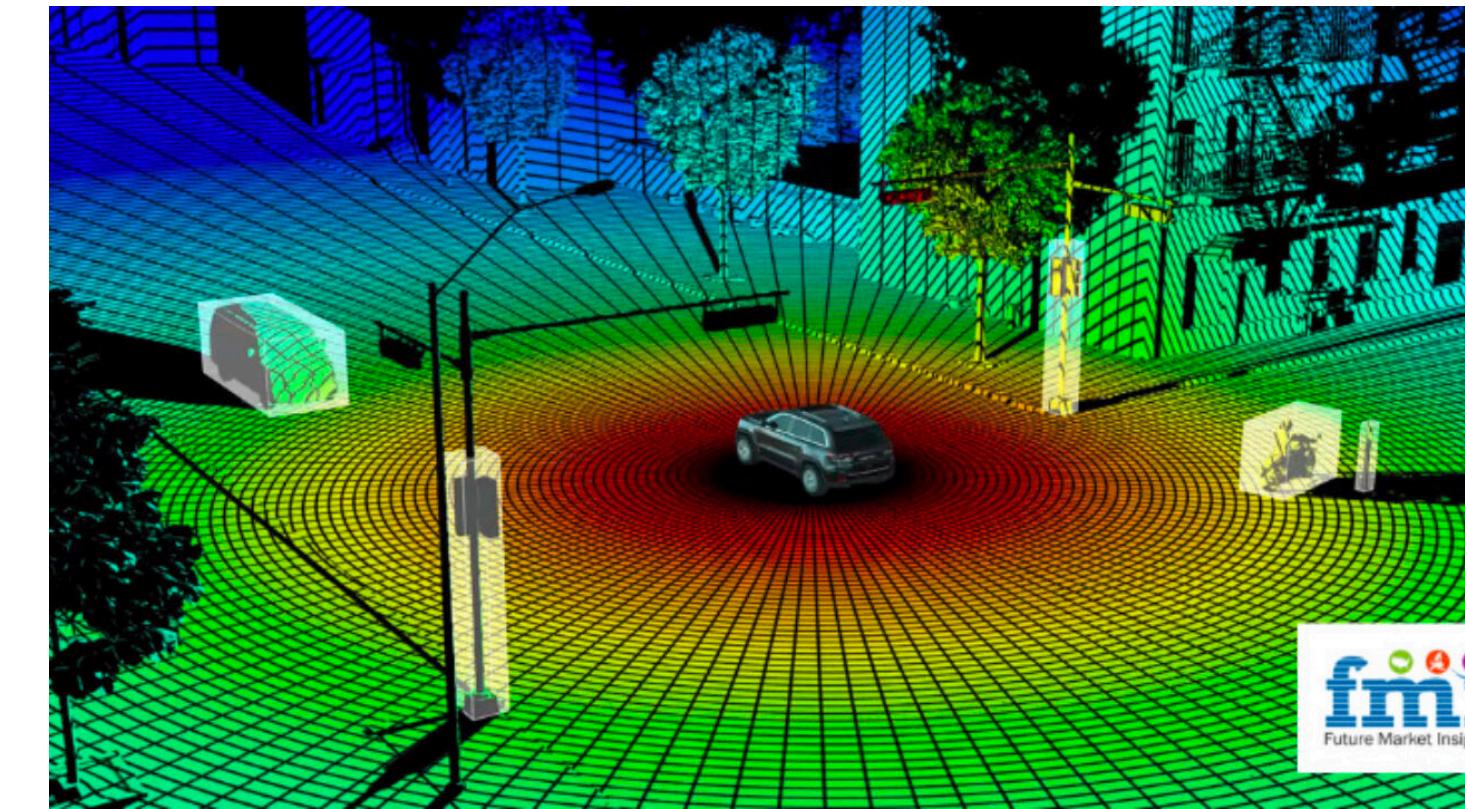
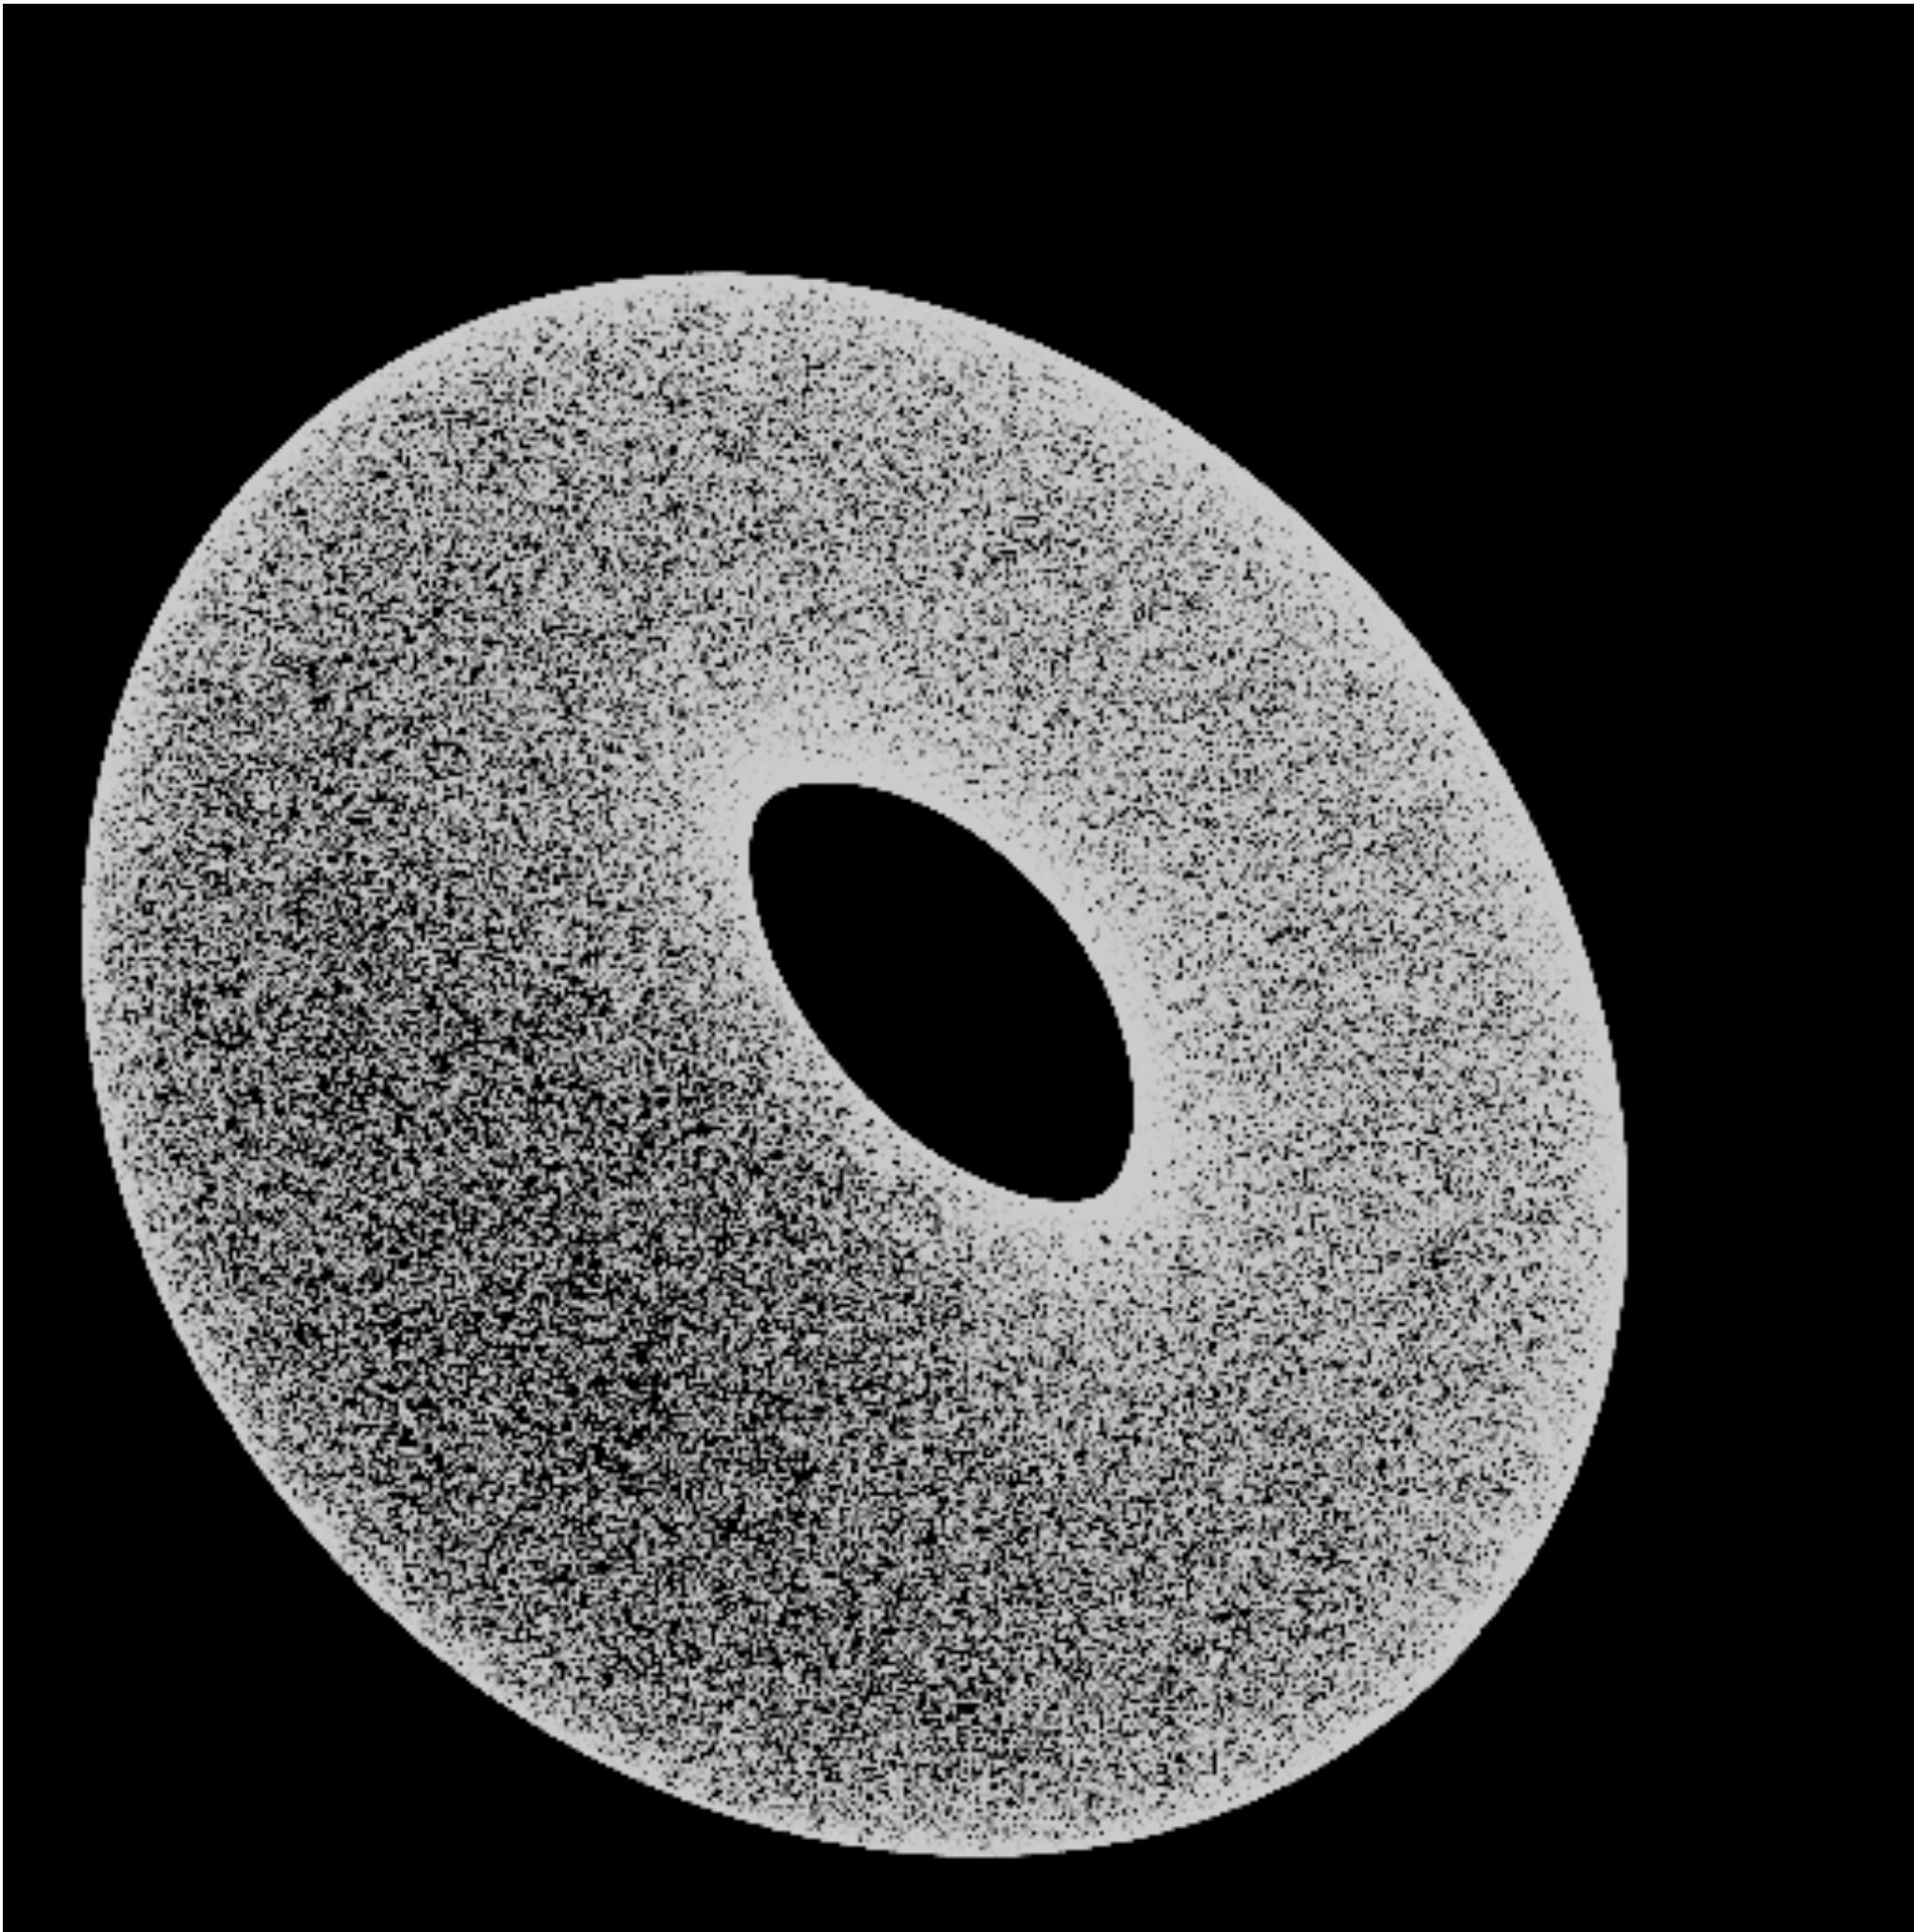
Surface Representations



Field Parameterizations



Point Clouds

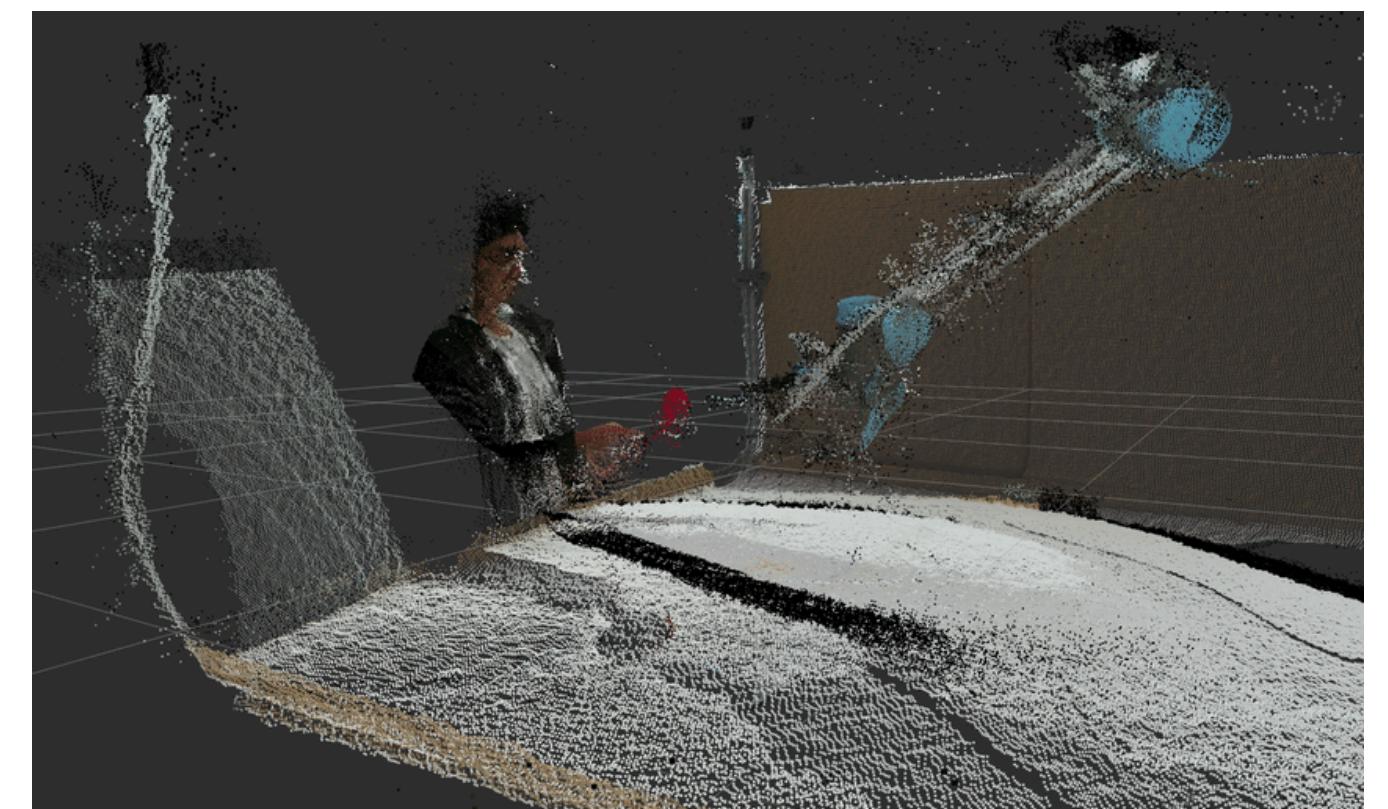
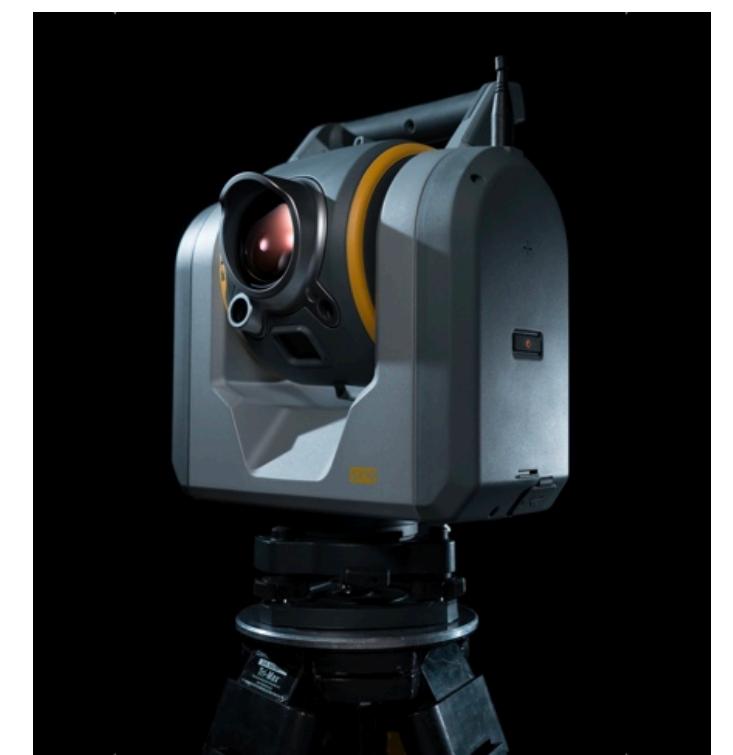
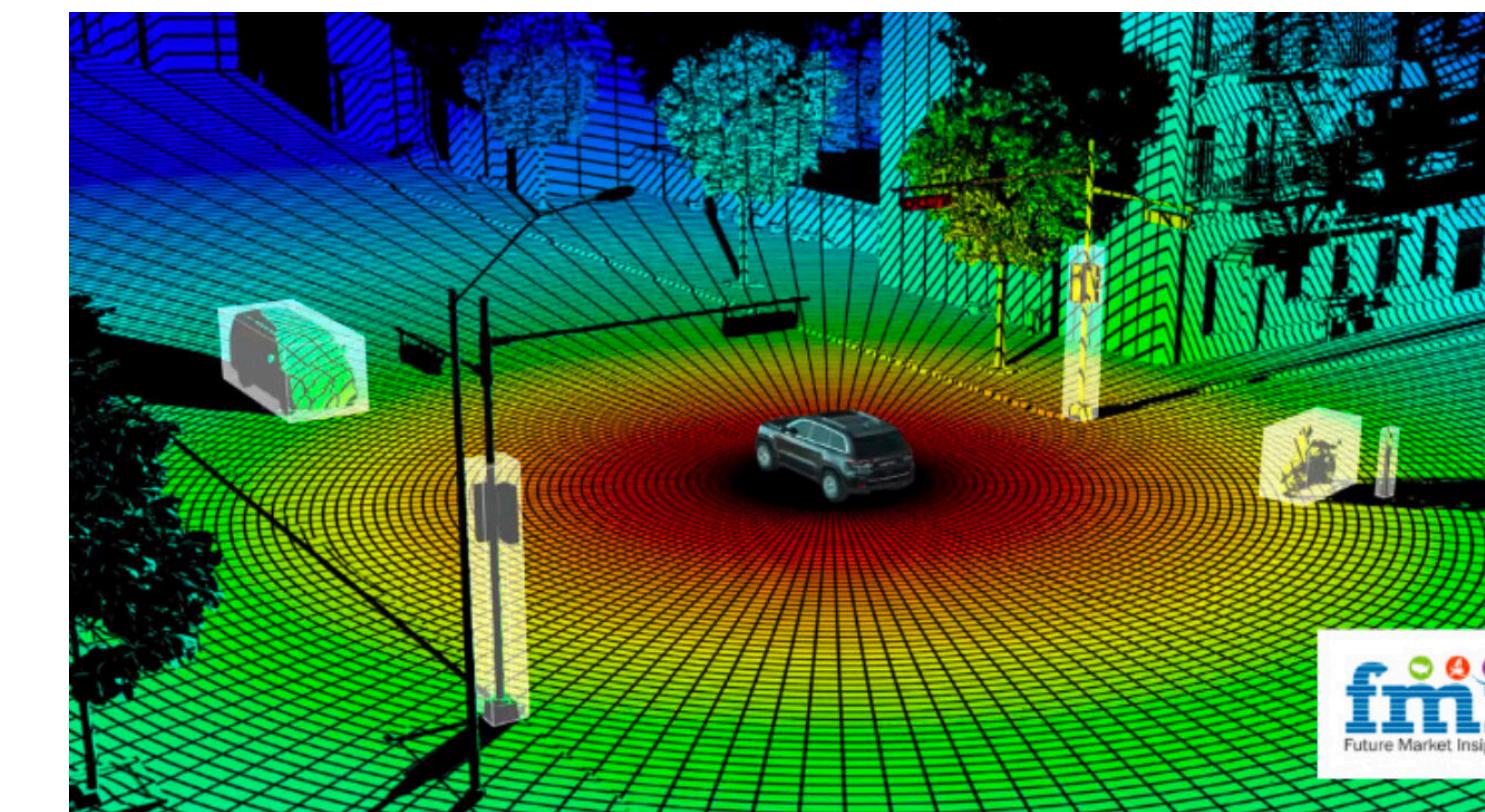


Point Clouds

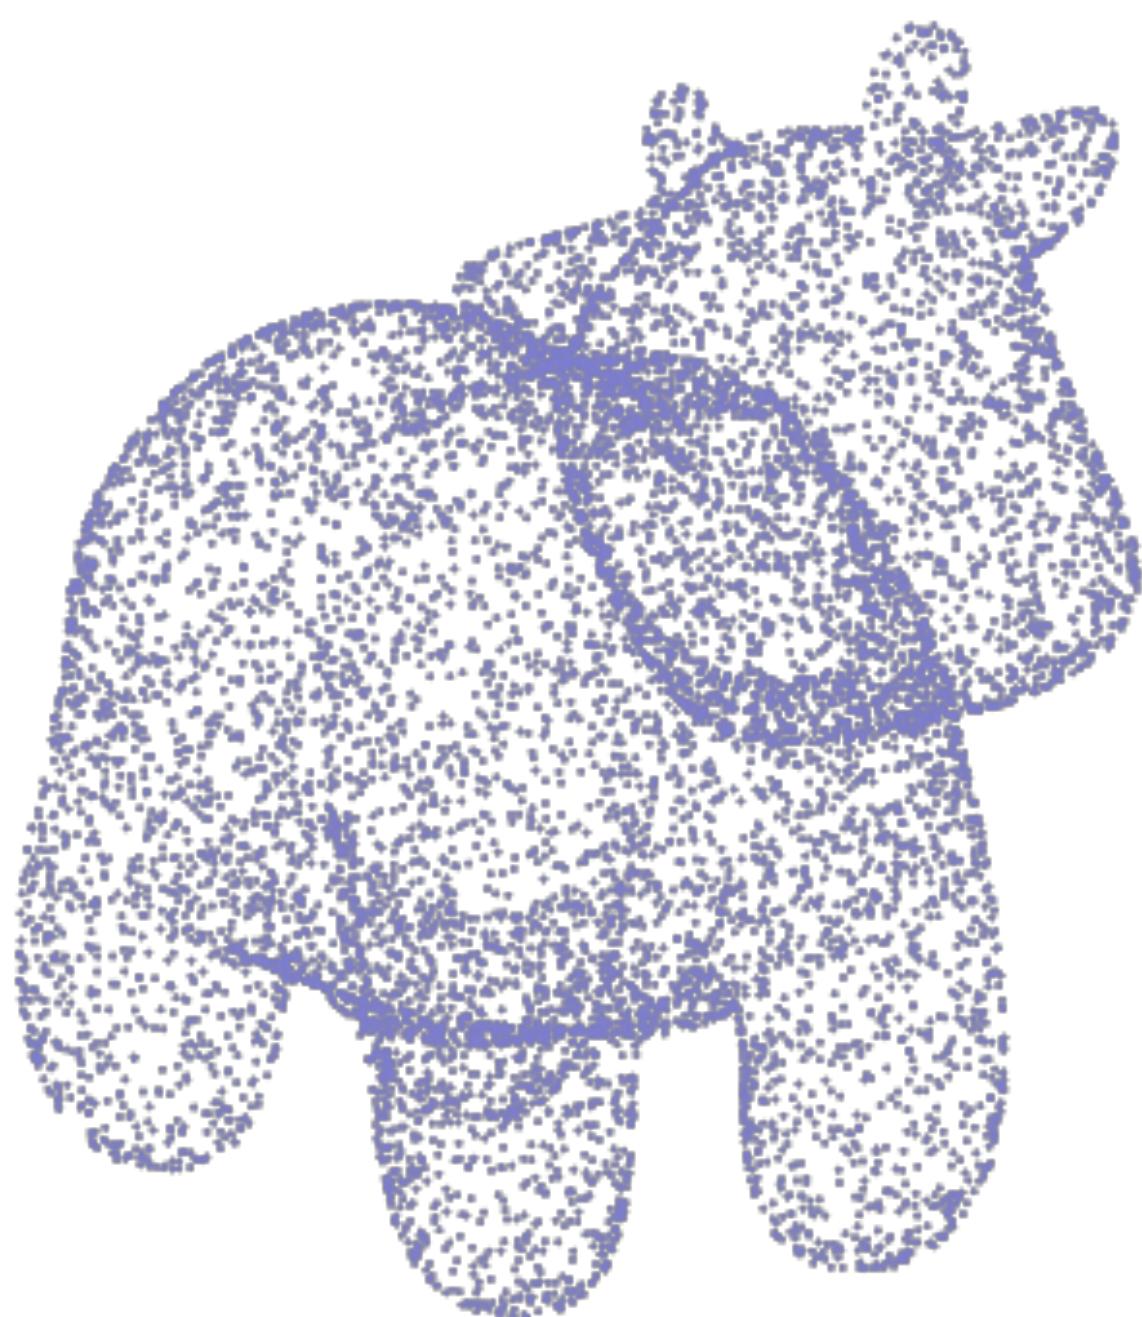
Often acquired in an image-like manner from sensors but can be useful to ‘forget’ this

e.g. point clouds from multiple views can make a single point cloud

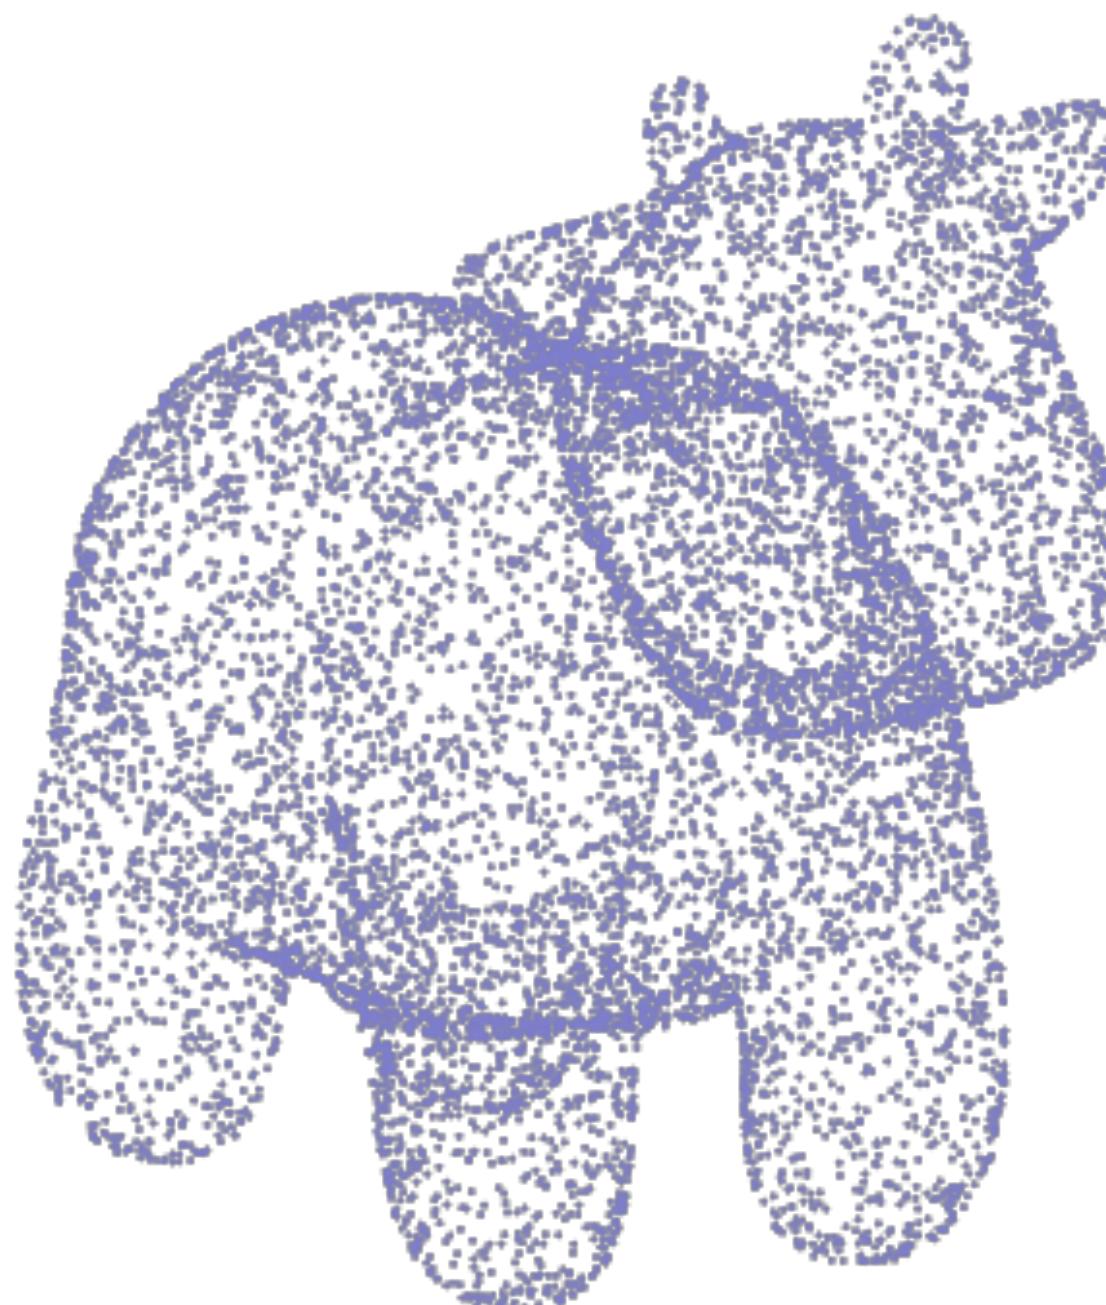
can help unify approaches independent of sensing modality



Point Clouds



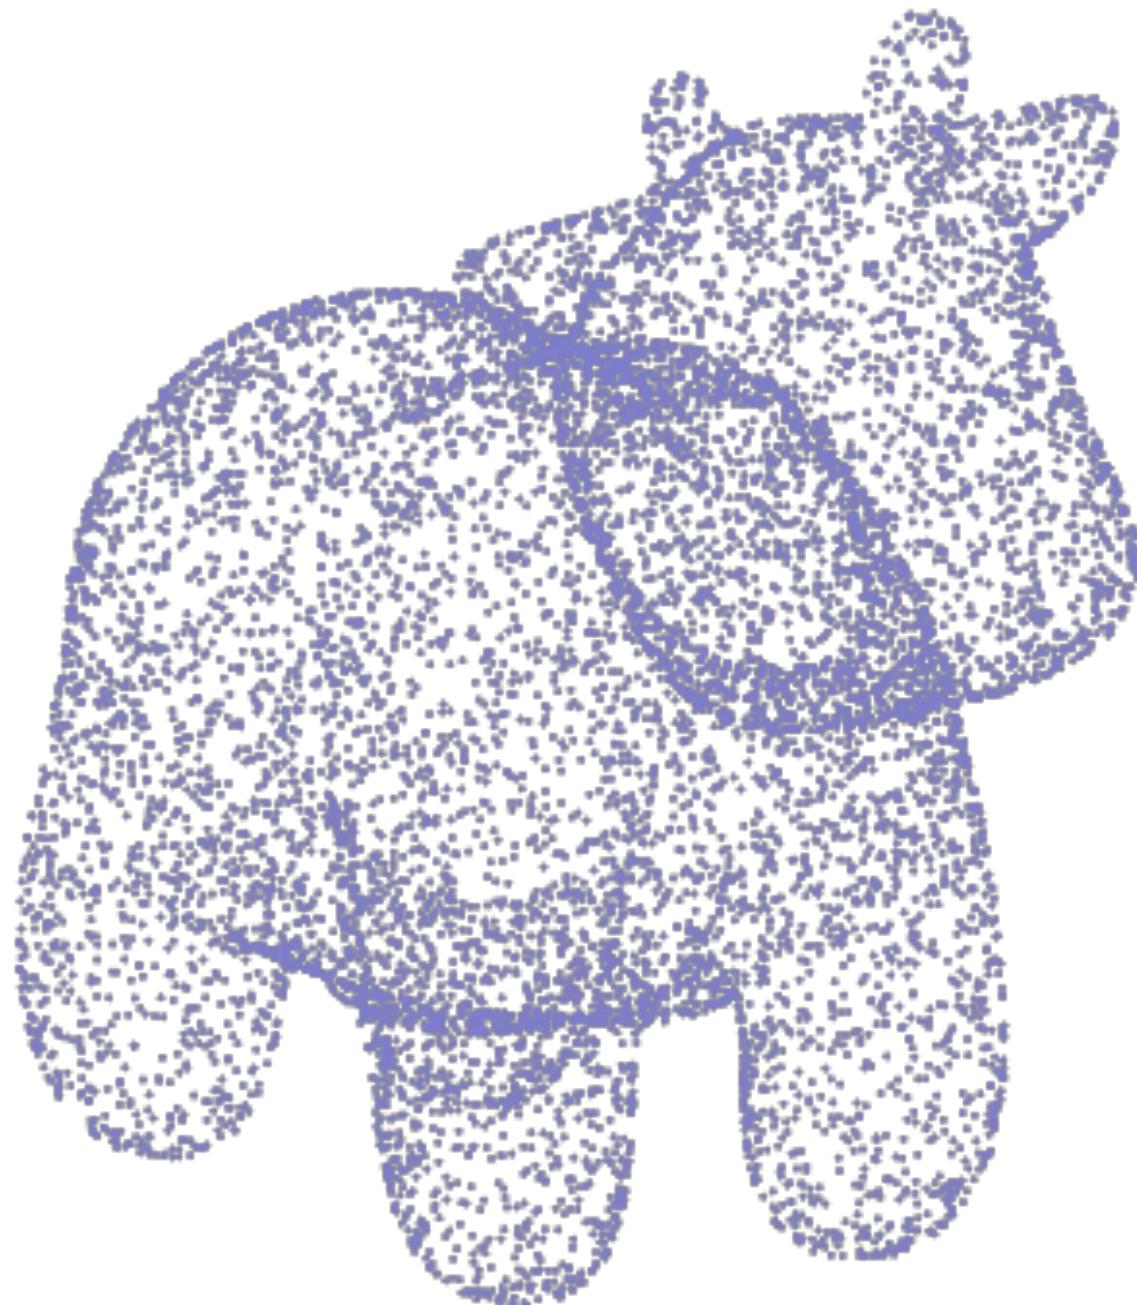
Permutation σ
→



$$\{\mathbf{p}_1, \mathbf{p}_2, \dots, \mathbf{p}_N\}$$

$$\{\mathbf{p}_{\sigma(1)}, \mathbf{p}_{\sigma(2)}, \dots, \mathbf{p}_{\sigma(N)}\}$$

Point Clouds



$$\{\mathbf{p}_1, \mathbf{p}_2, \dots, \mathbf{p}_N\}$$

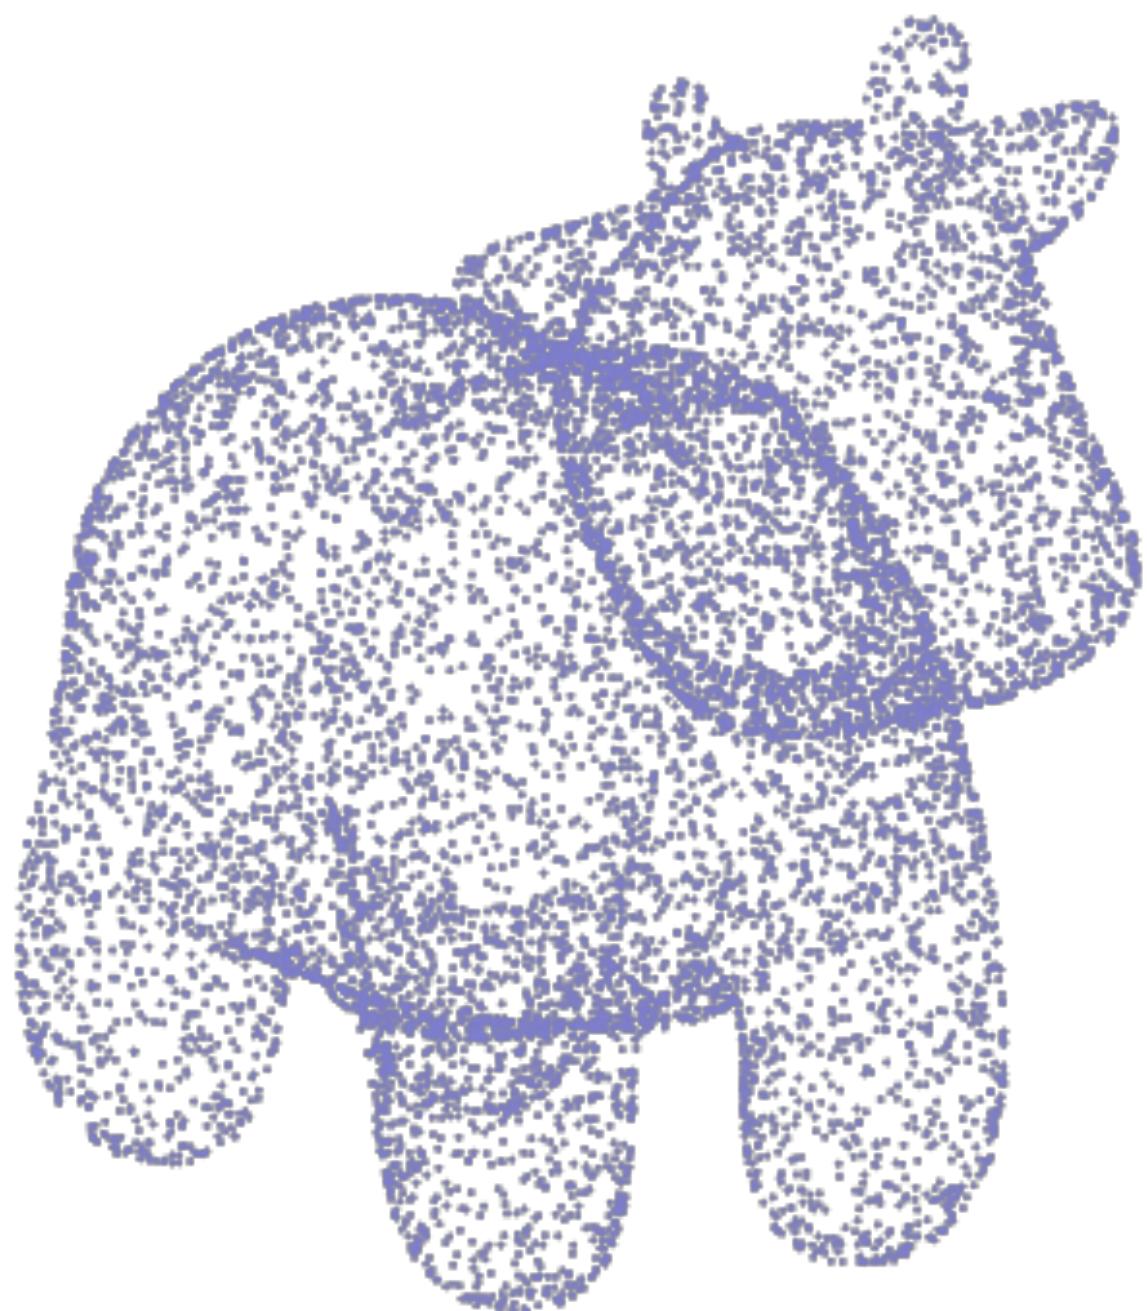
Unordered set of points

x1,y1,z1

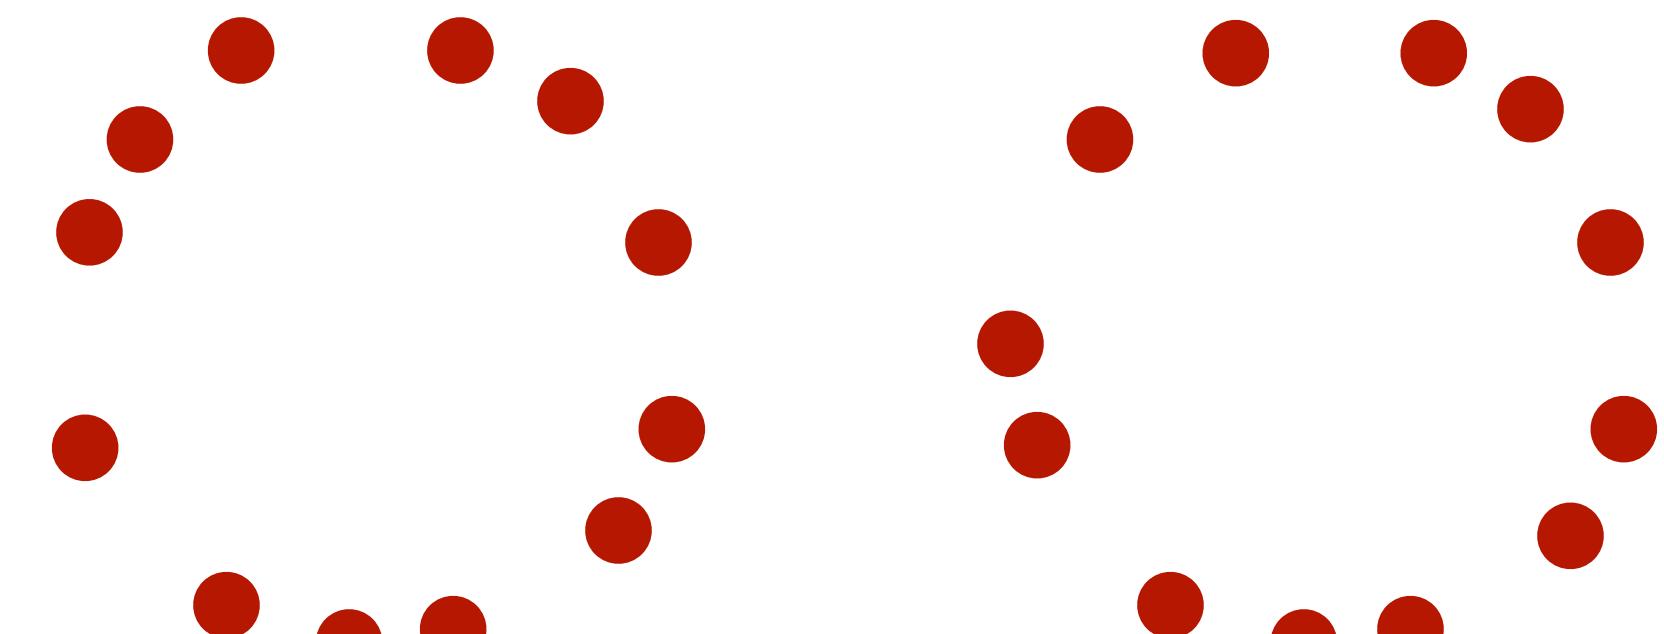
Often represented as a $N \times 3$ array, but ordering does not matter (unlike images)

Need processing/generation methods that are permutation invariant (e.g. fully connected layer will not work)

Point Clouds



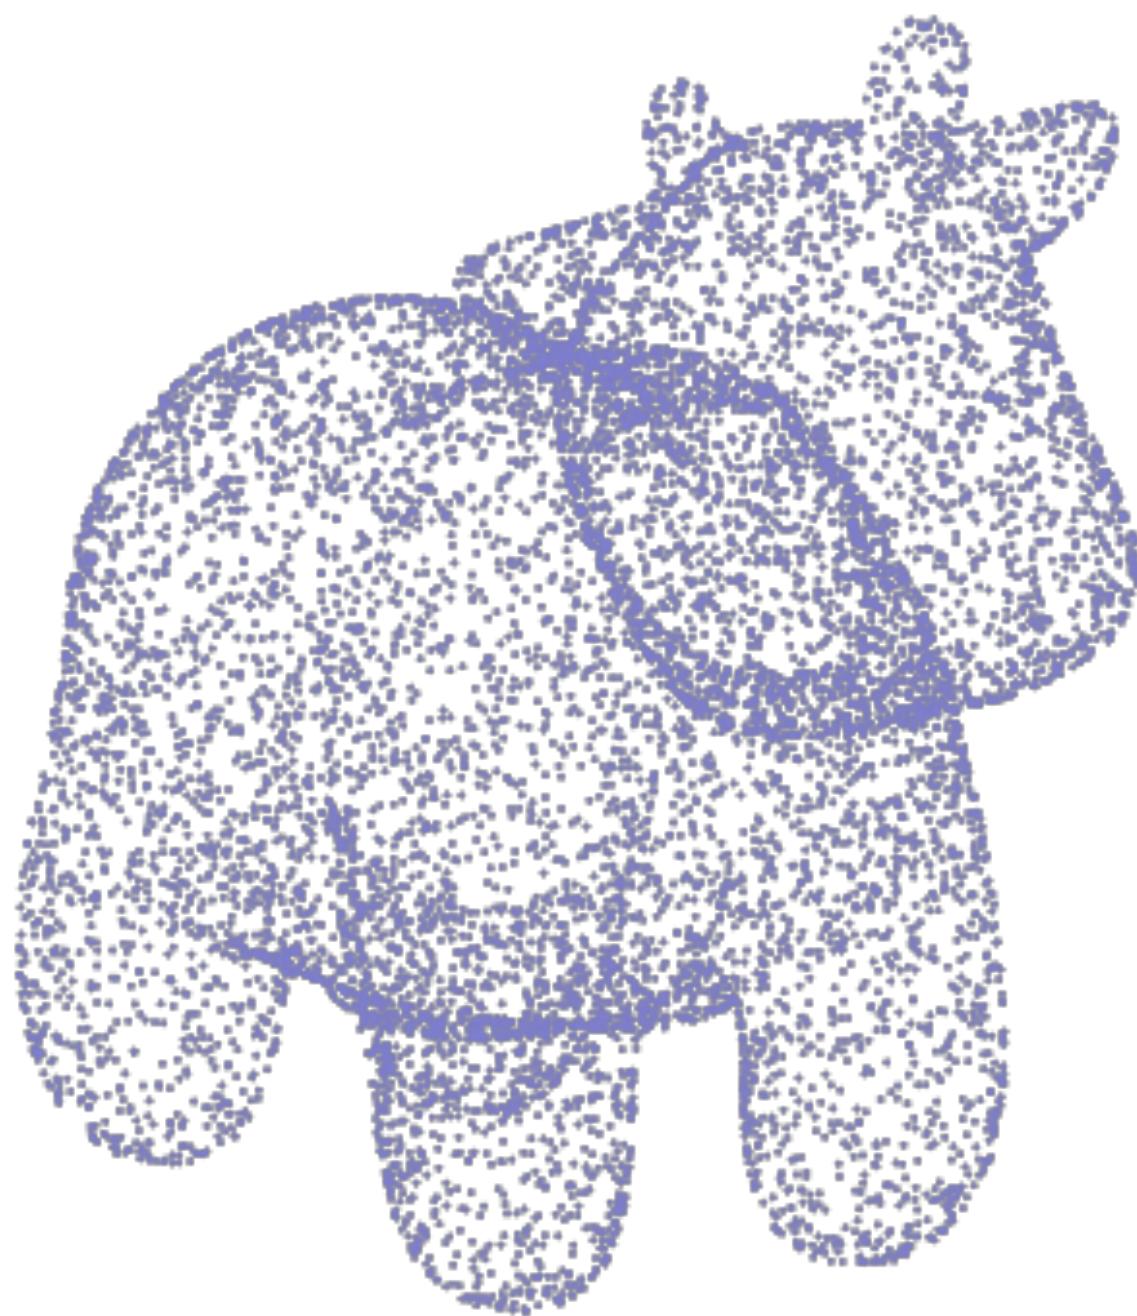
No explicit ‘connectivity’ information



$$\{\mathbf{p}_1, \mathbf{p}_2, \dots, \mathbf{p}_N\}$$

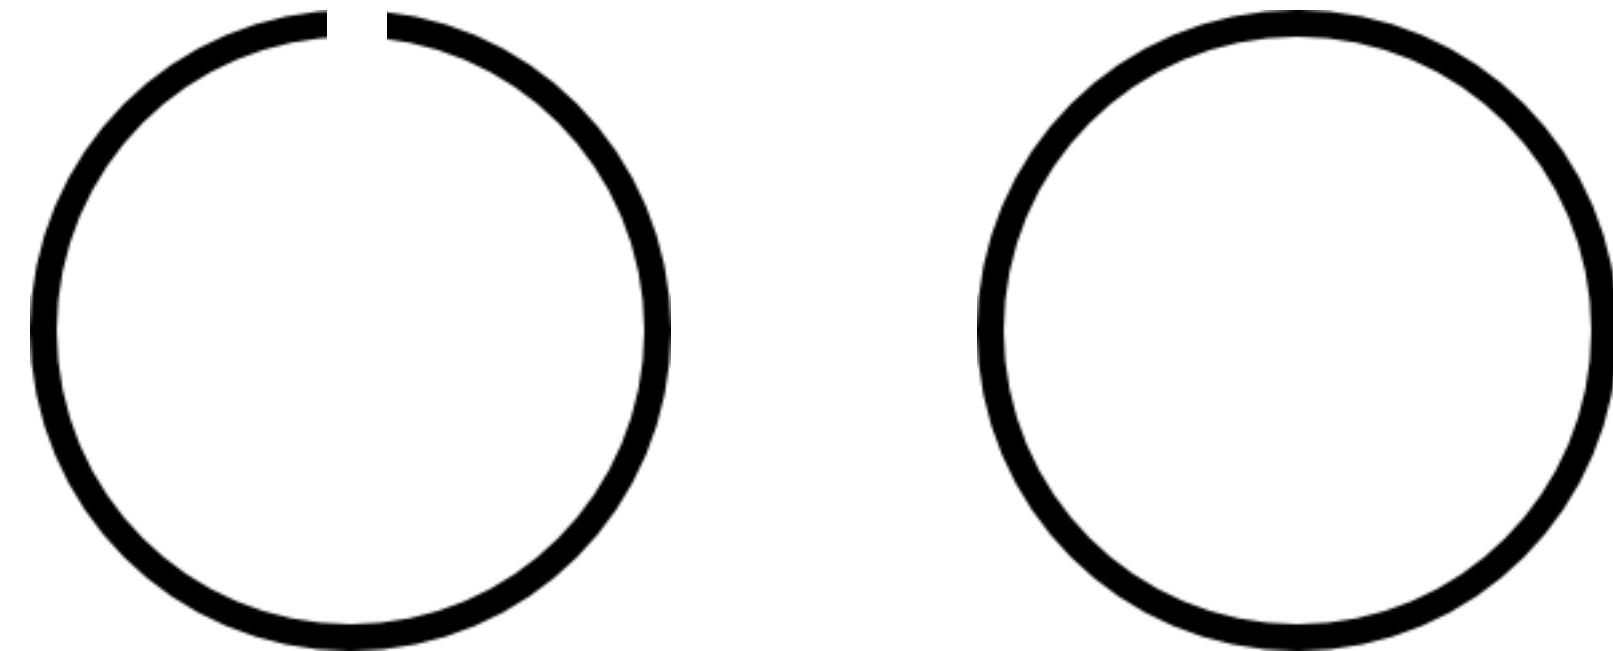
Unordered set of points

Point Clouds



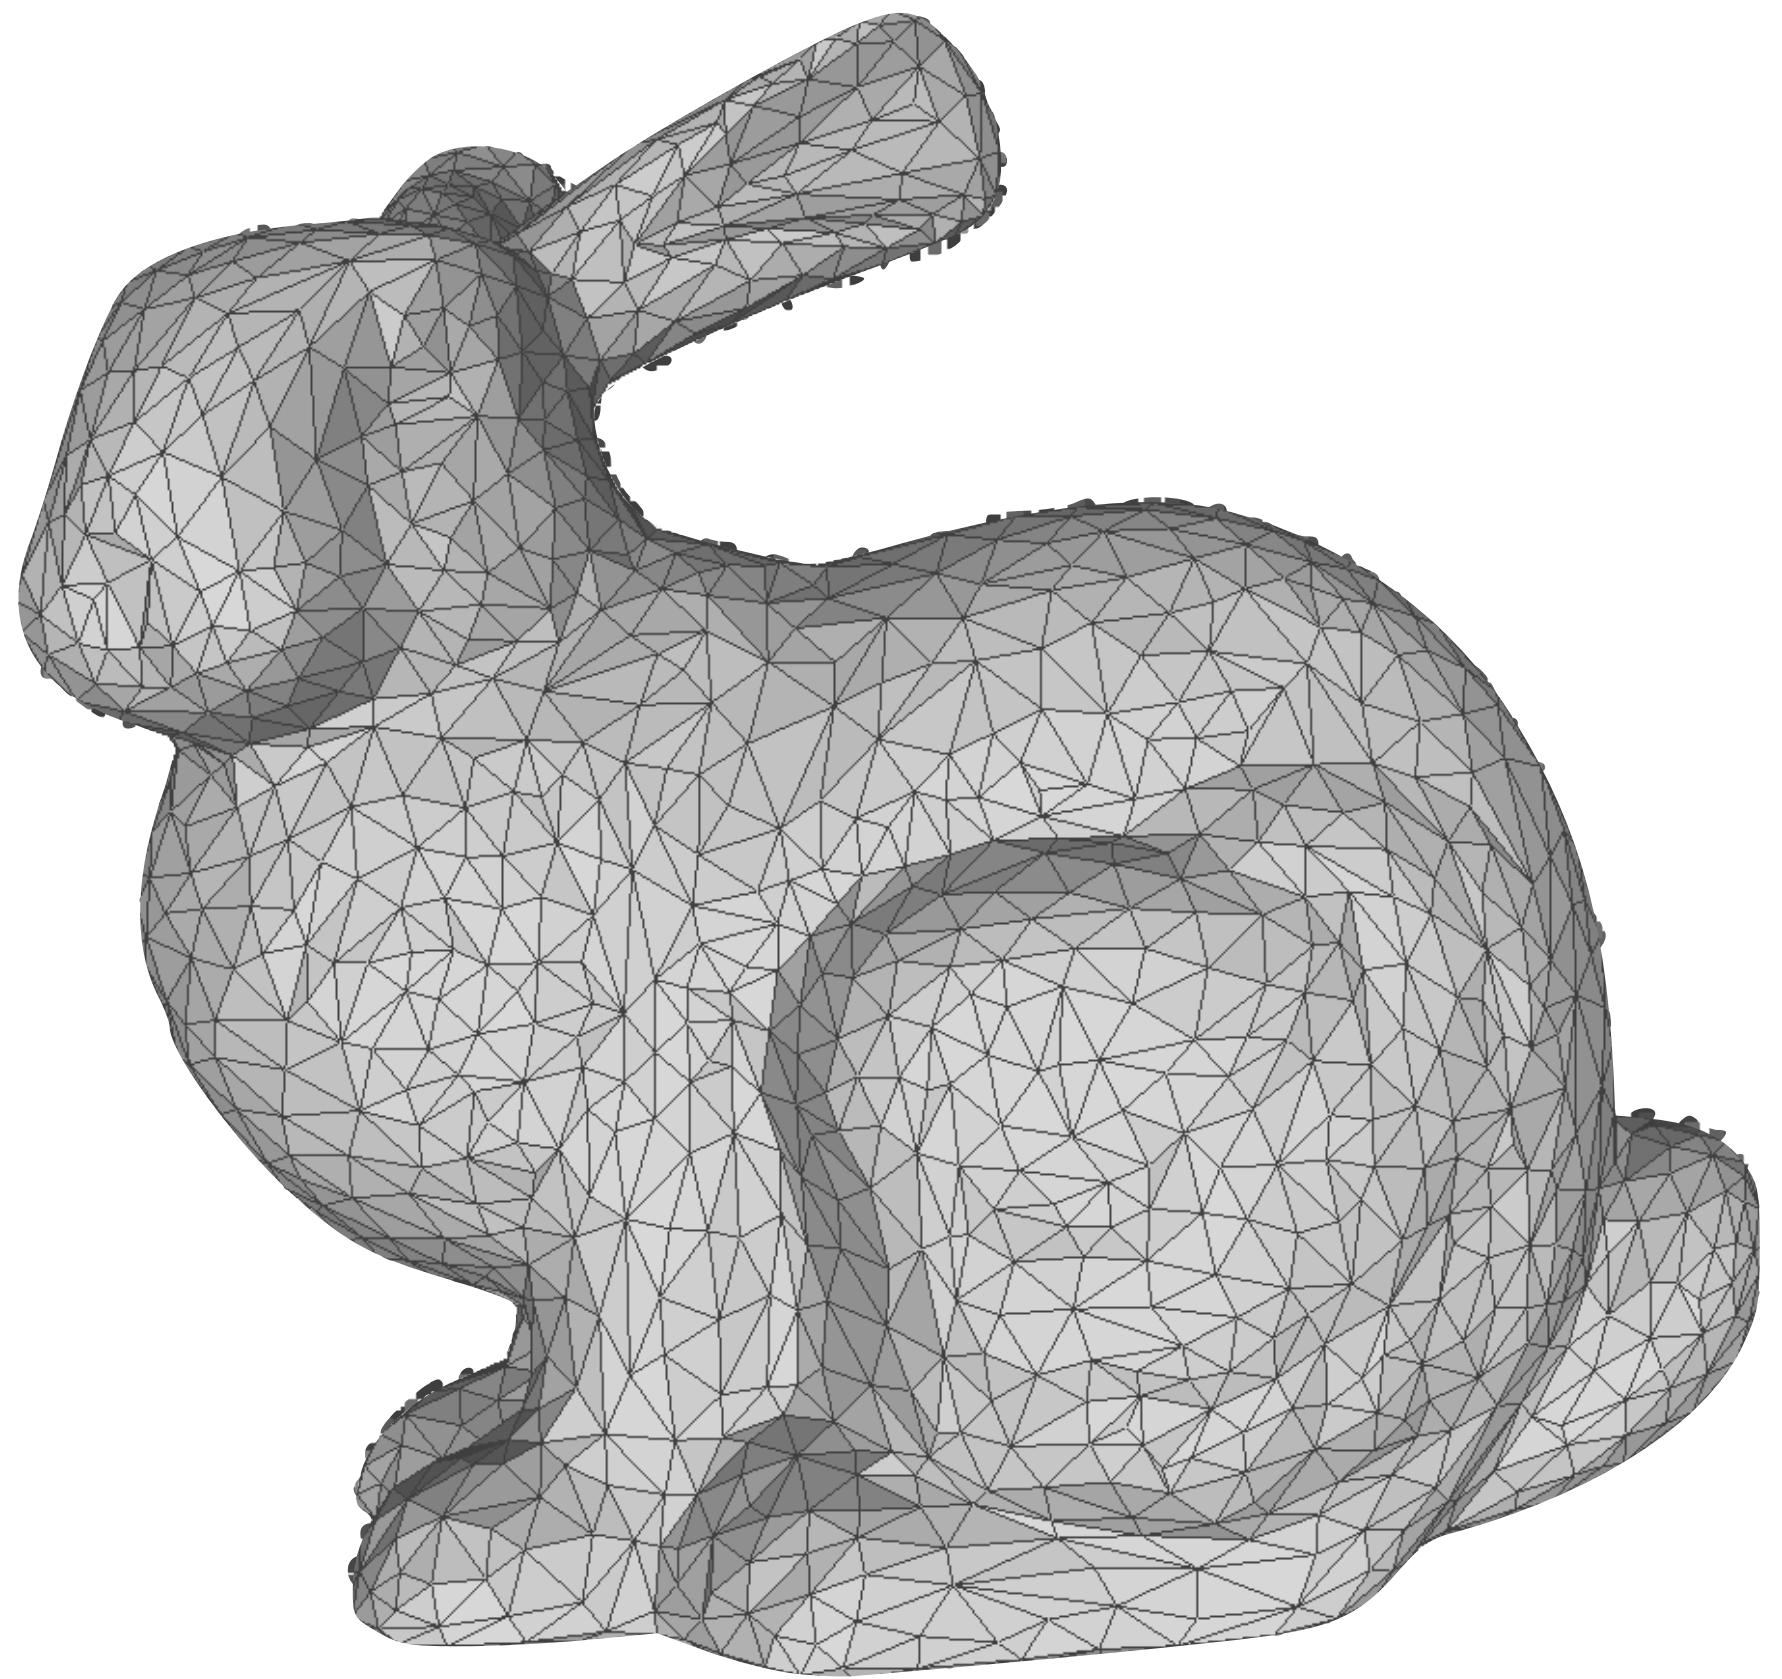
$\{p_1, p_2, \dots, p_N\}$
Unordered set of points

No explicit ‘connectivity’ information

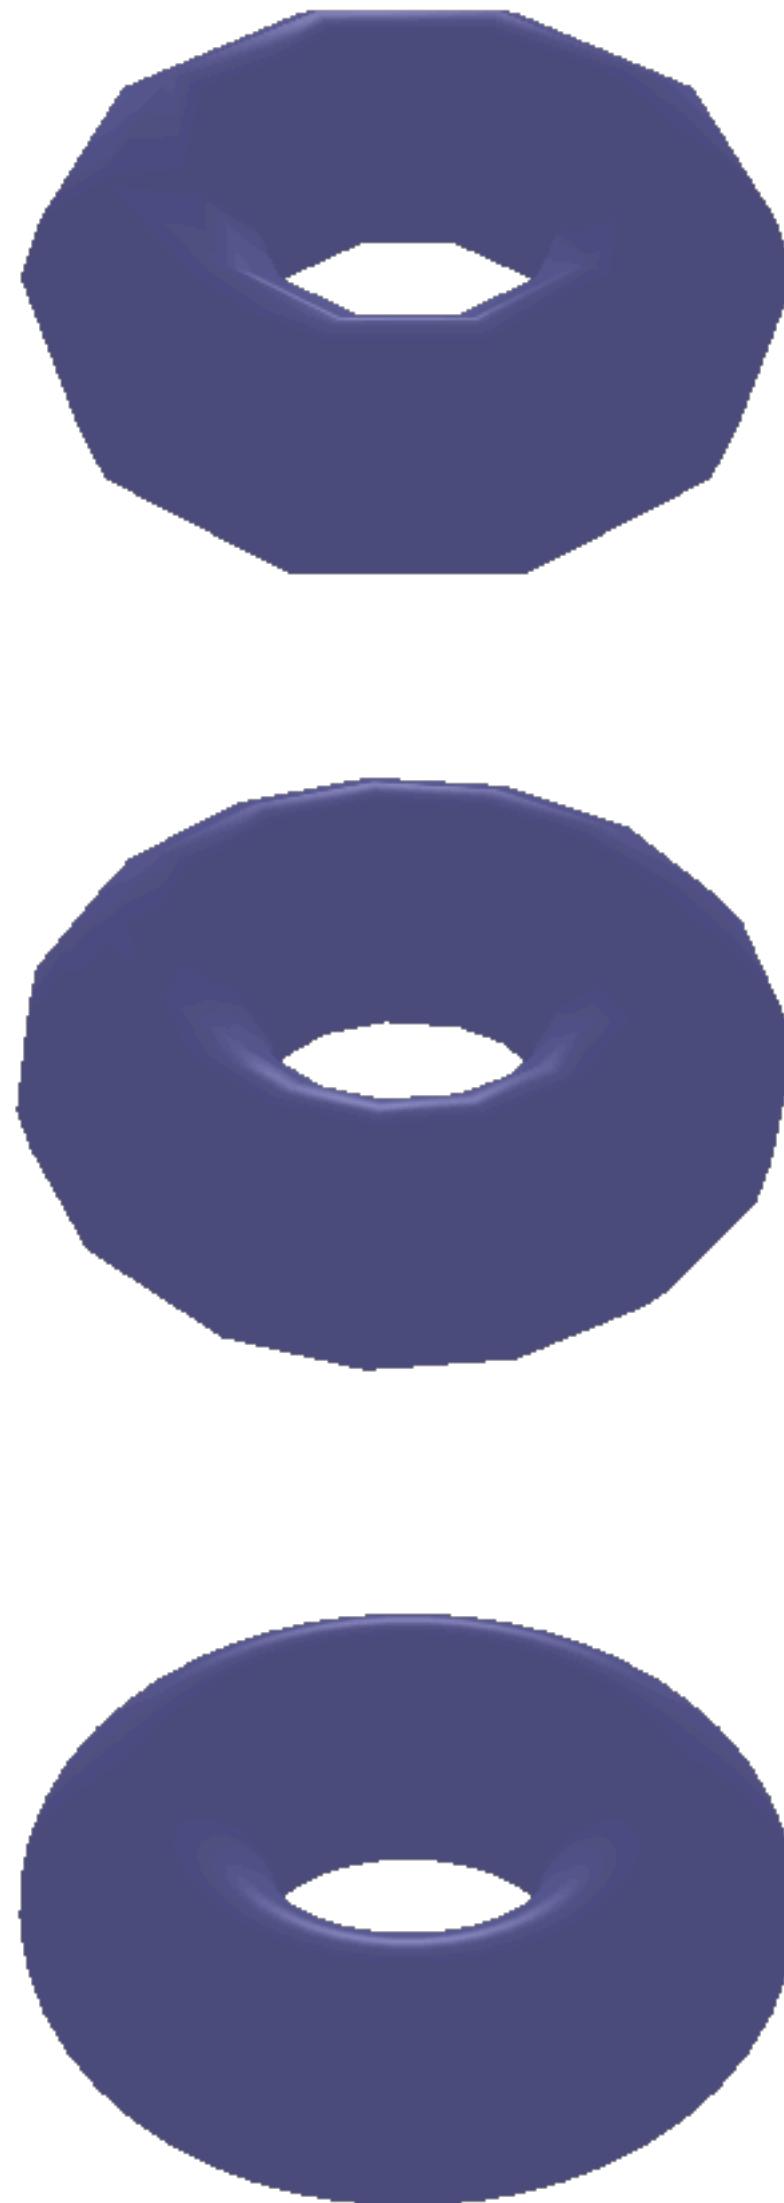


(yes, we can sample more - but
more efficient to add edges)

Meshes



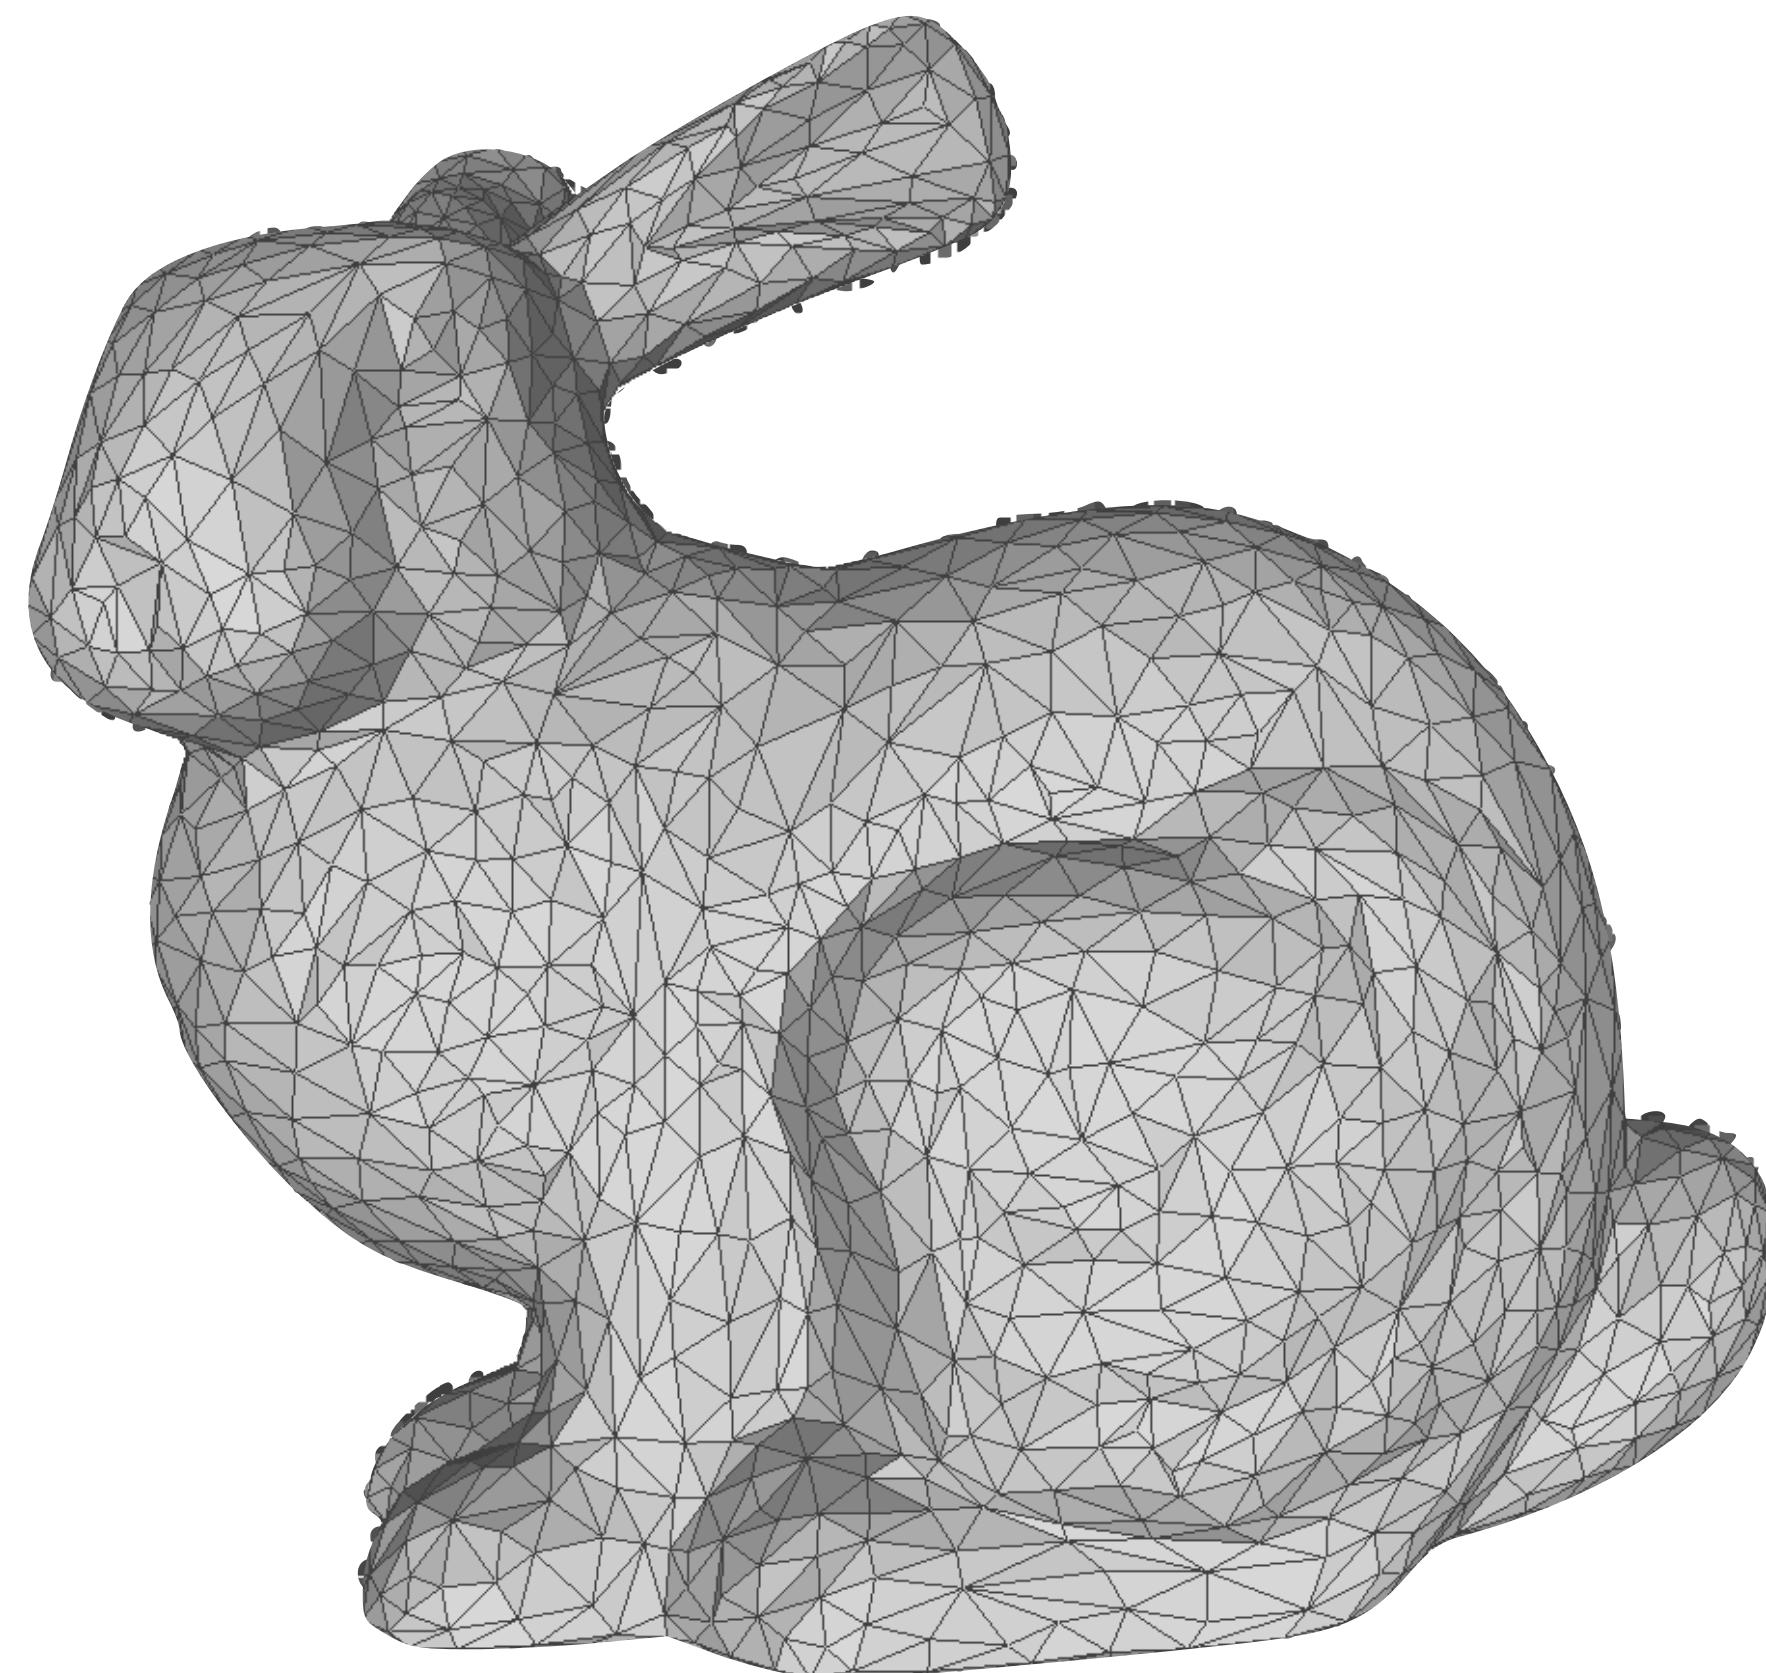
Vertices + Faces



Piecewise linear
approximations of
underlying surface

More elements allow
better approximation

Meshes



Vertices + Faces

Vertices

Positions of Vertices

Faces

Connectivity

(indices of vertices
that make a ‘face’)

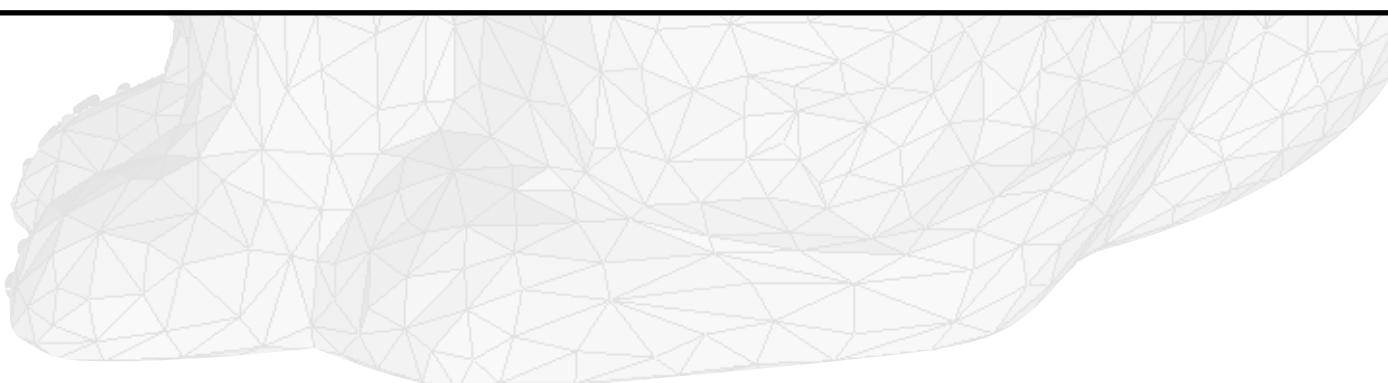
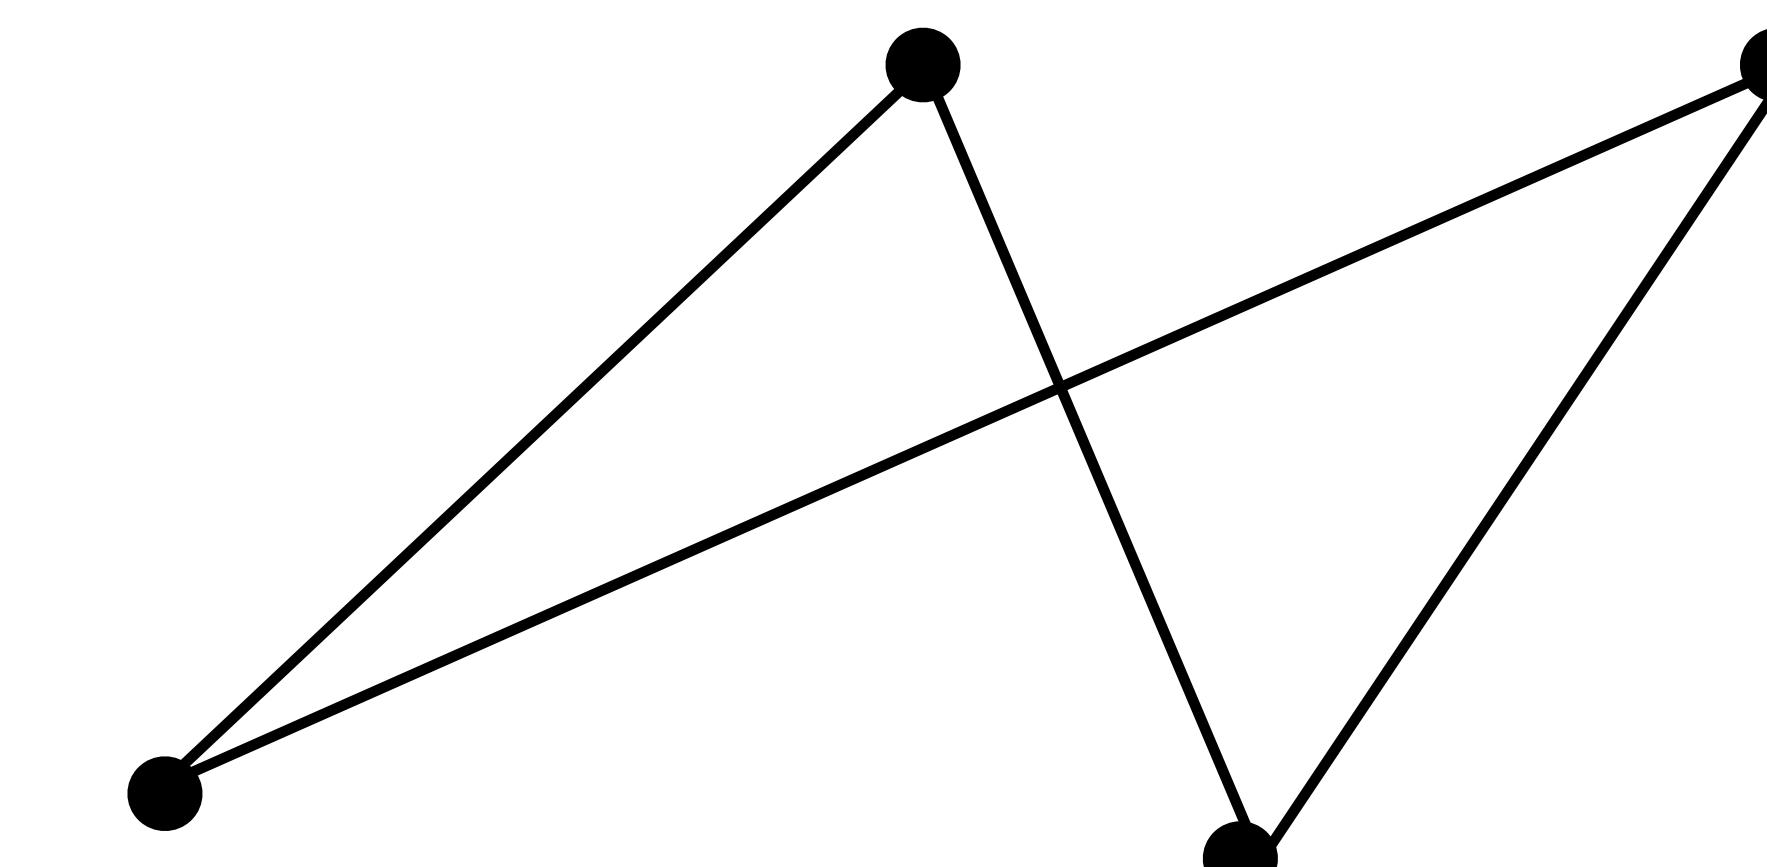
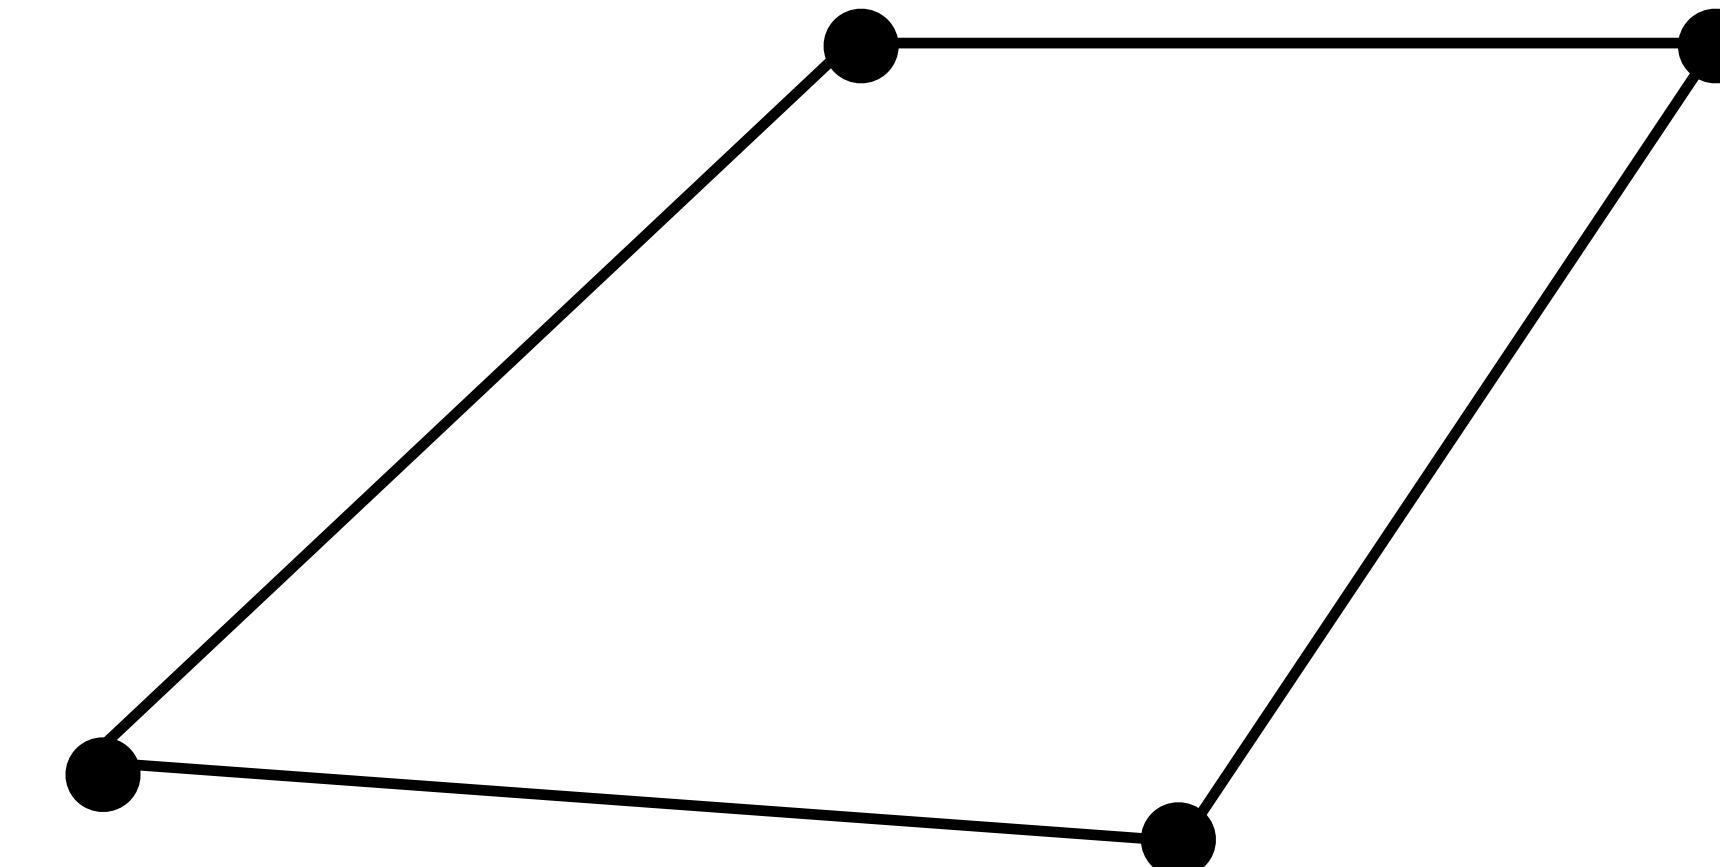
Meshes

Faces

Vertices

i,j,k

Ordering Matters in Face Indices



Vertices + Faces

Positions of Vertices

1
2
3

Connectivity
(indices of vertices
that make a ‘face’)

i,j,k
.

1
2
3

1
2
3
4

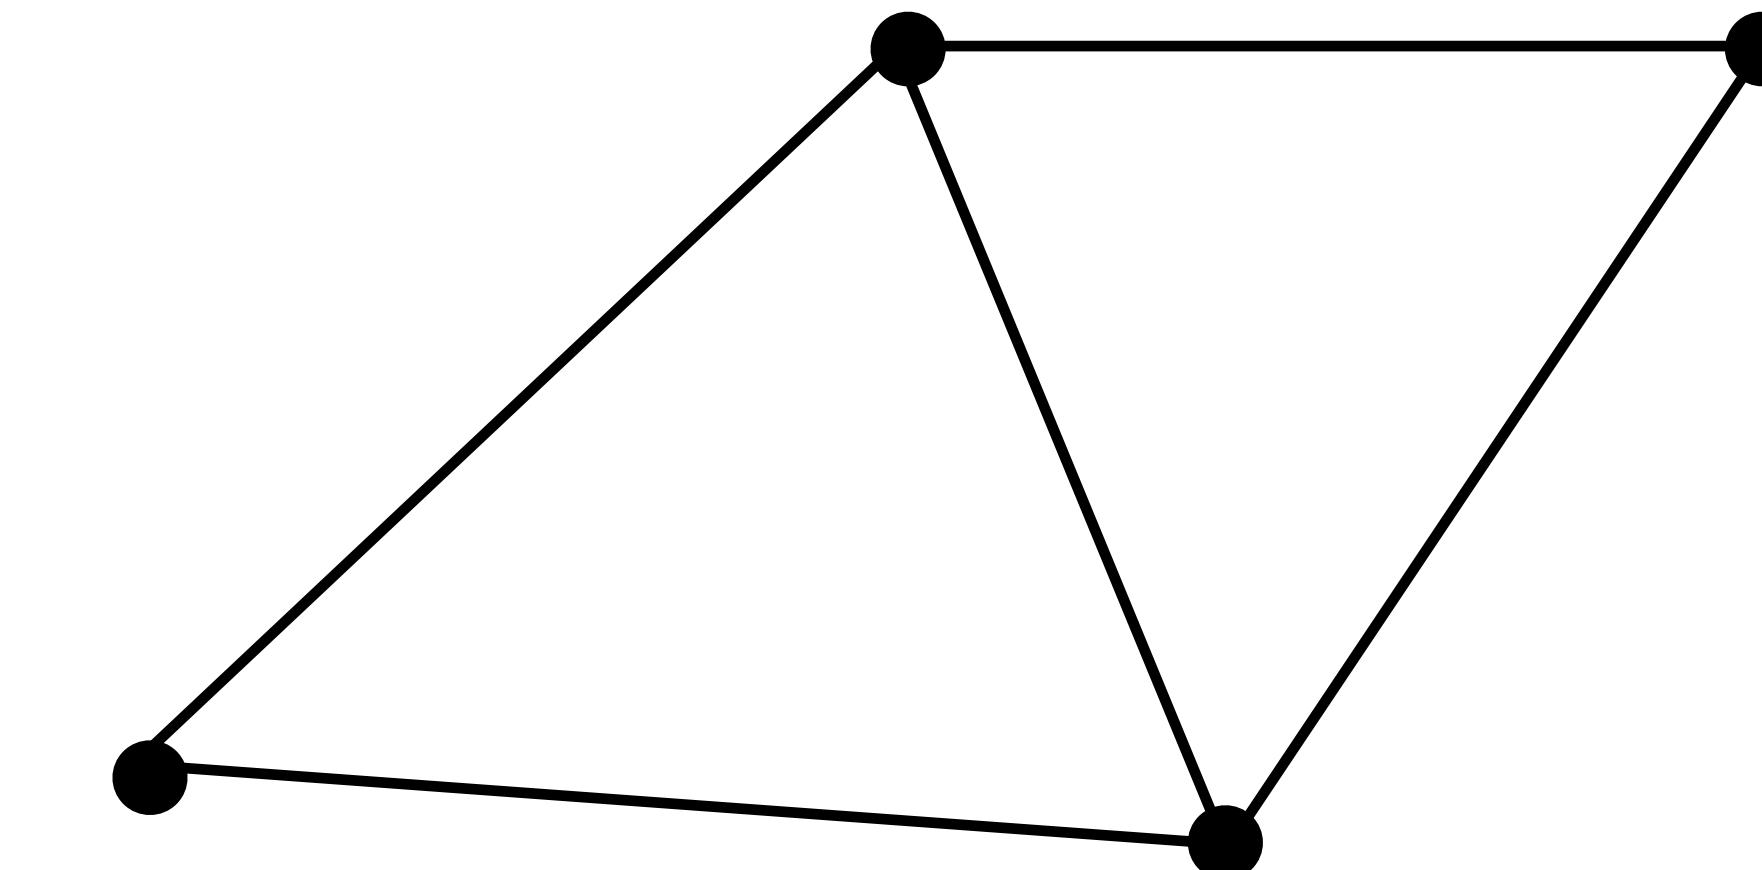
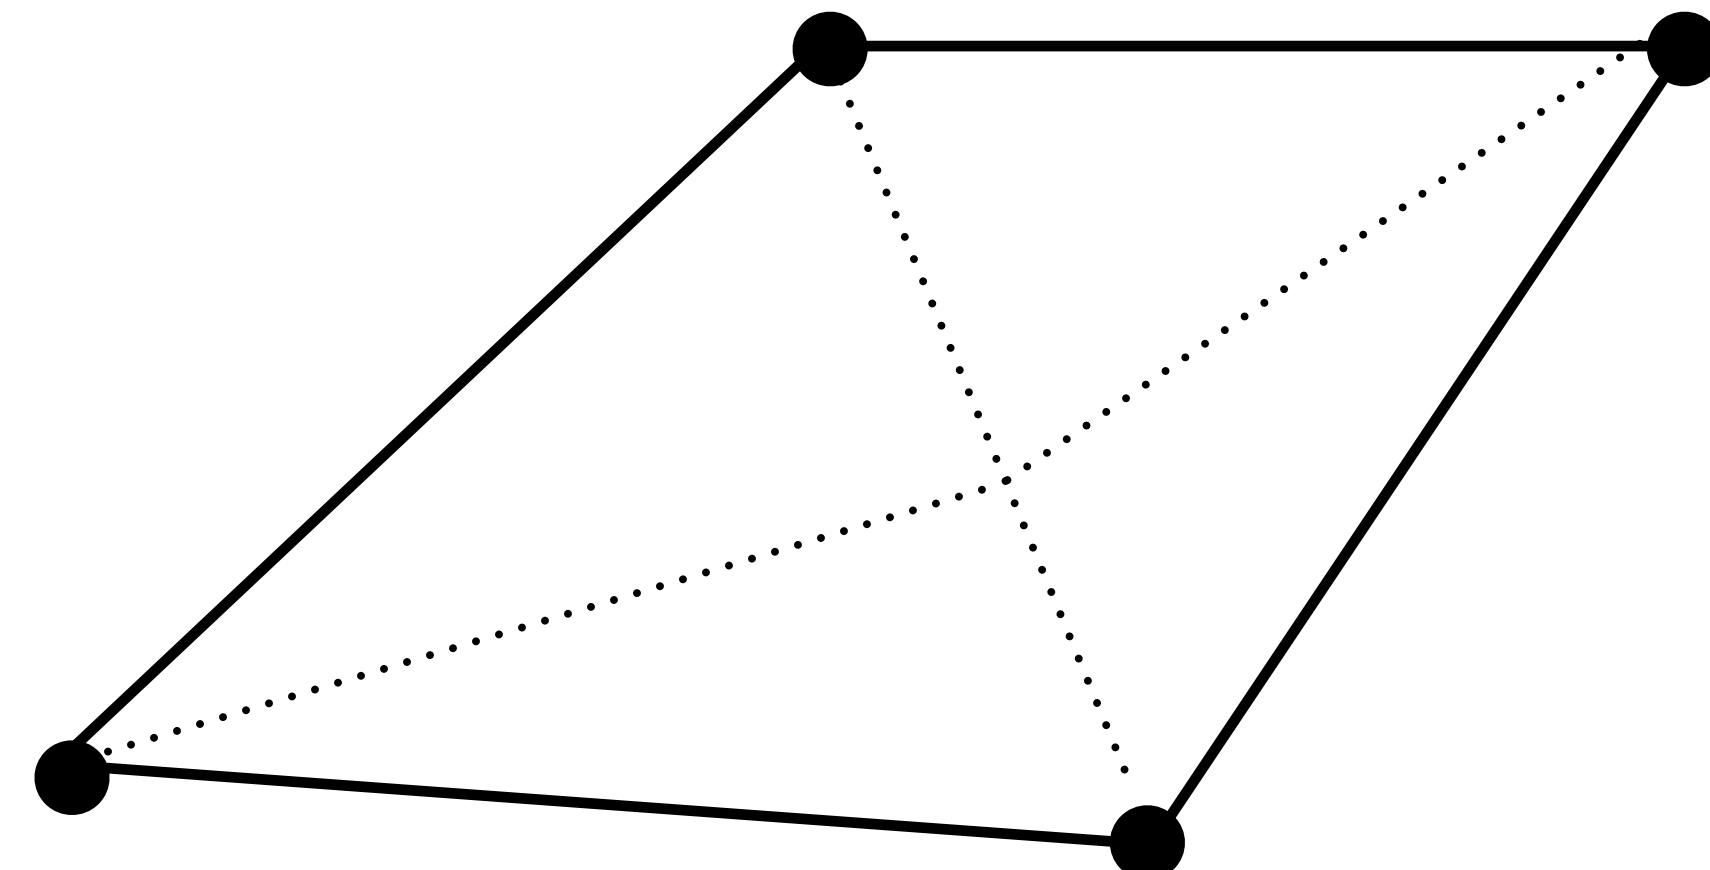
Meshes

Faces

Vertices

i,j,k

Arbitrary Polygons can be non-planar



restrict to *Triangular*
Meshes

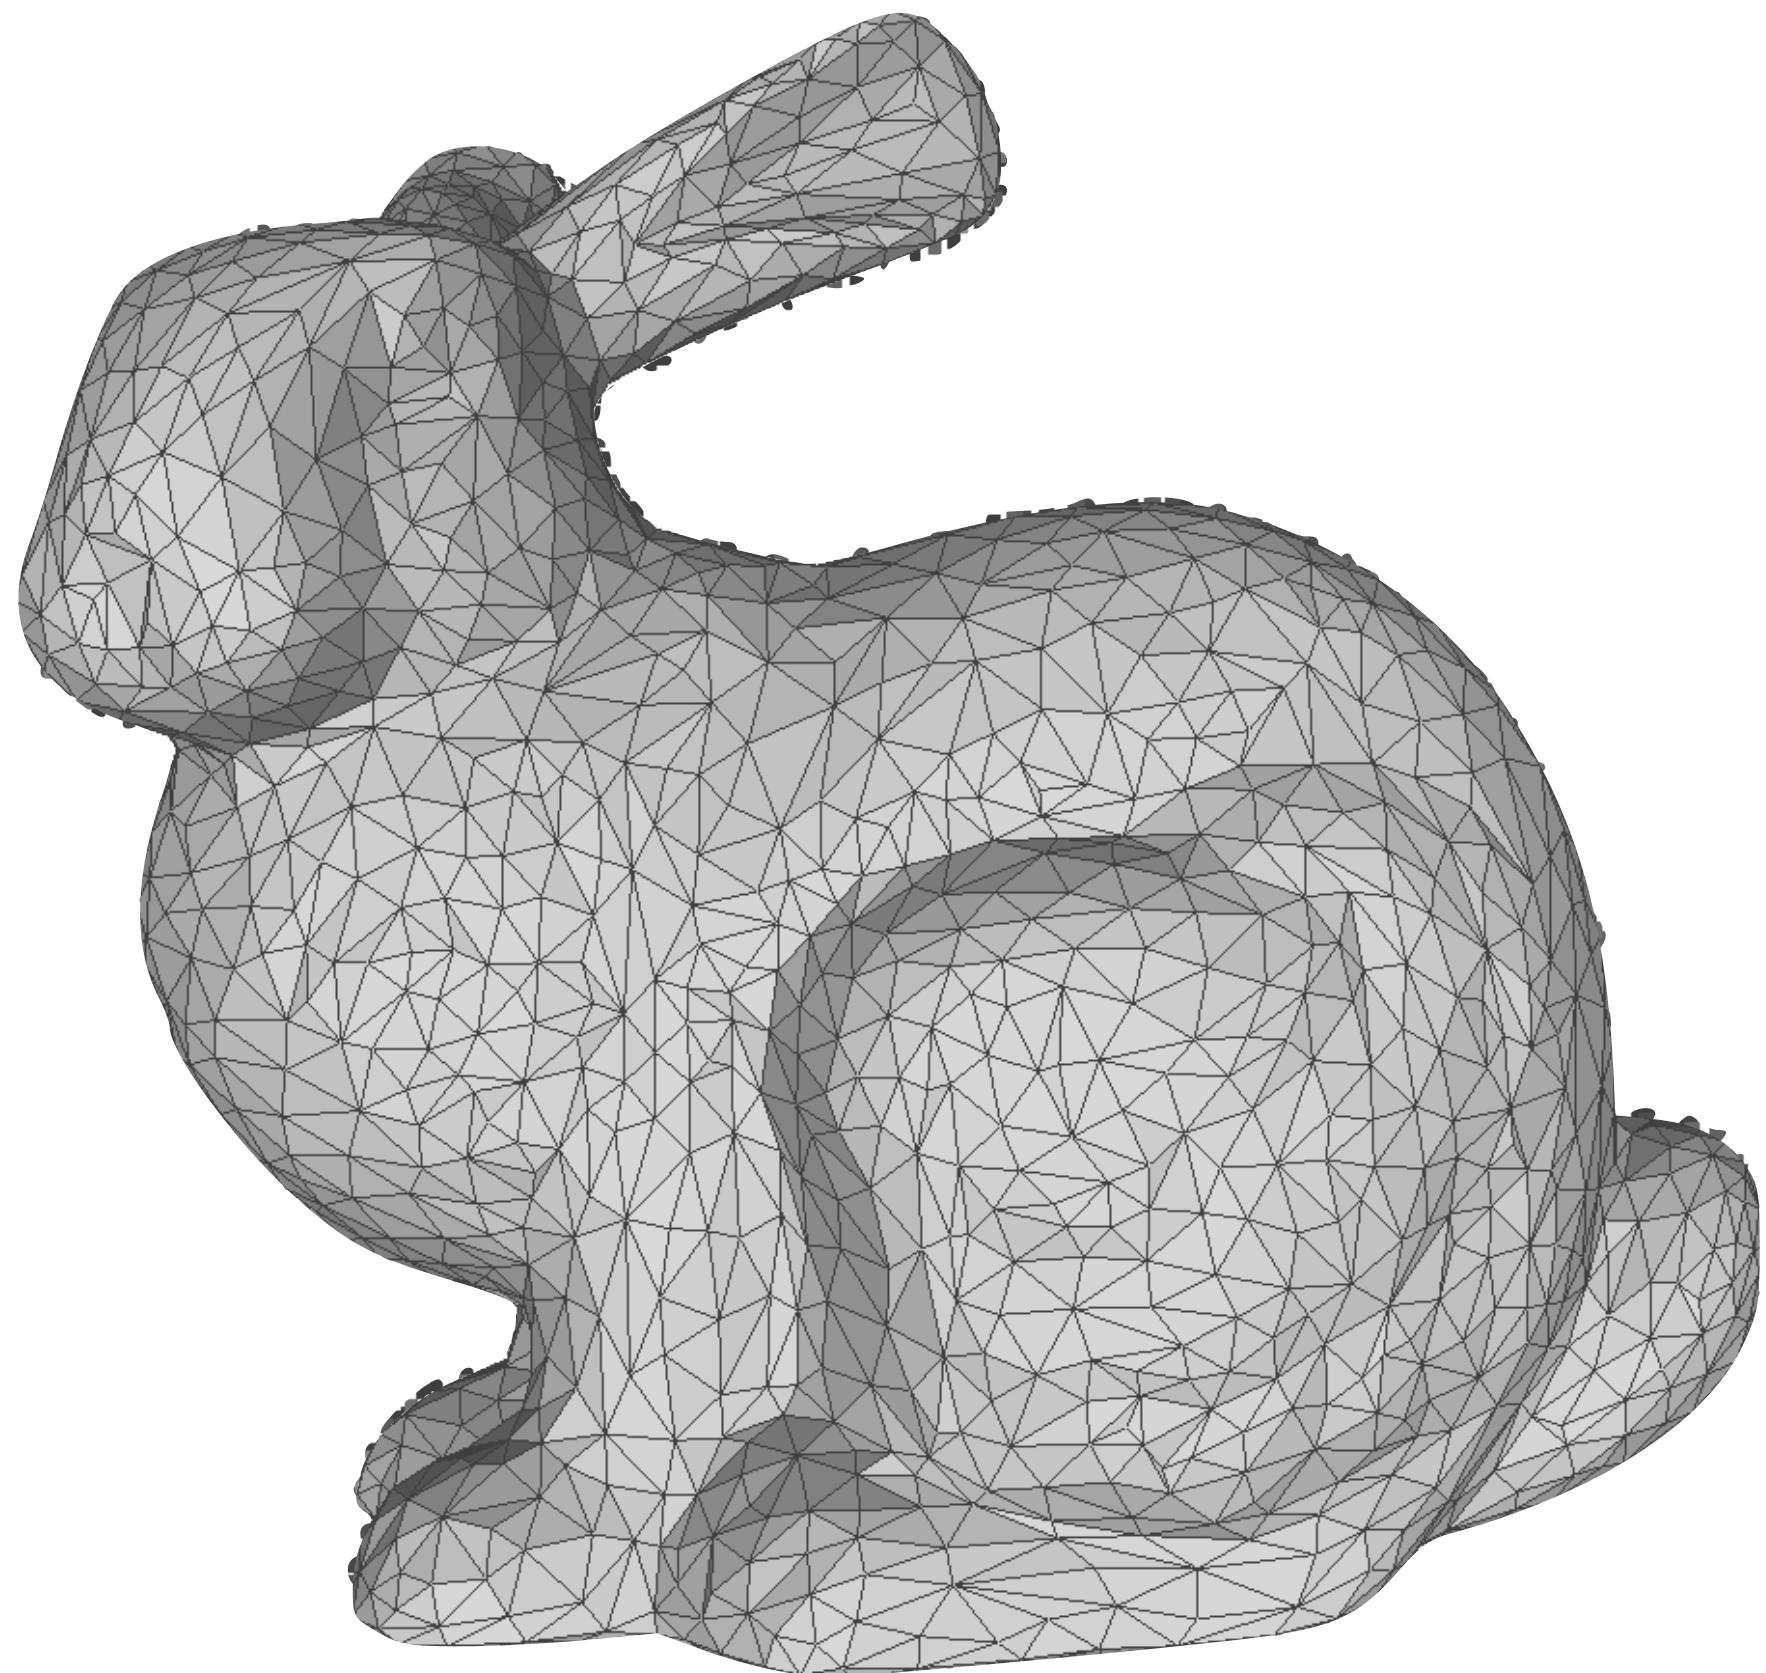
Vertices + Faces

Positions of Vertices

Connectivity
(indices of vertices
that make a ‘face’)

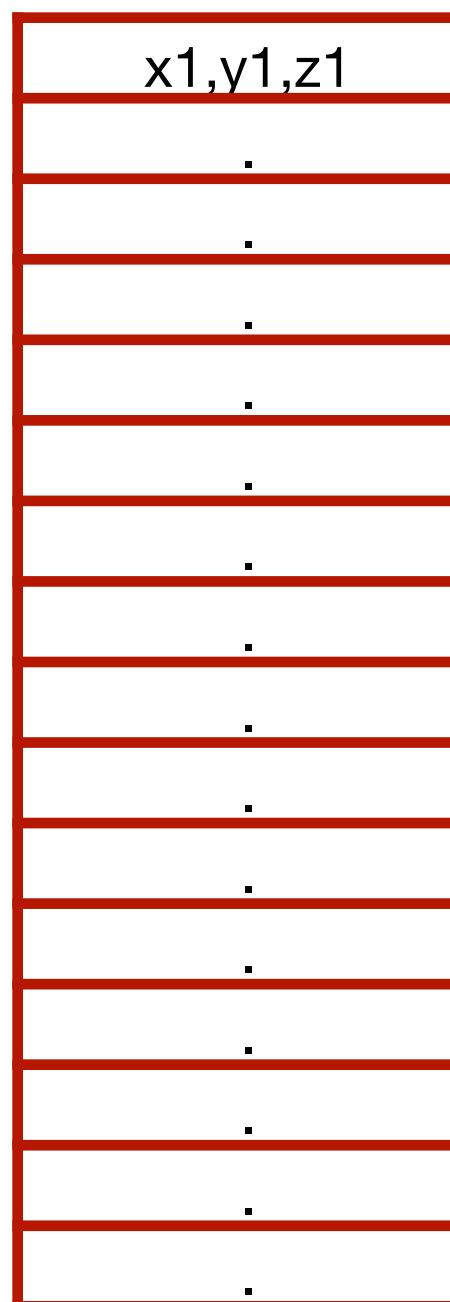
Triangular Meshes

Faces

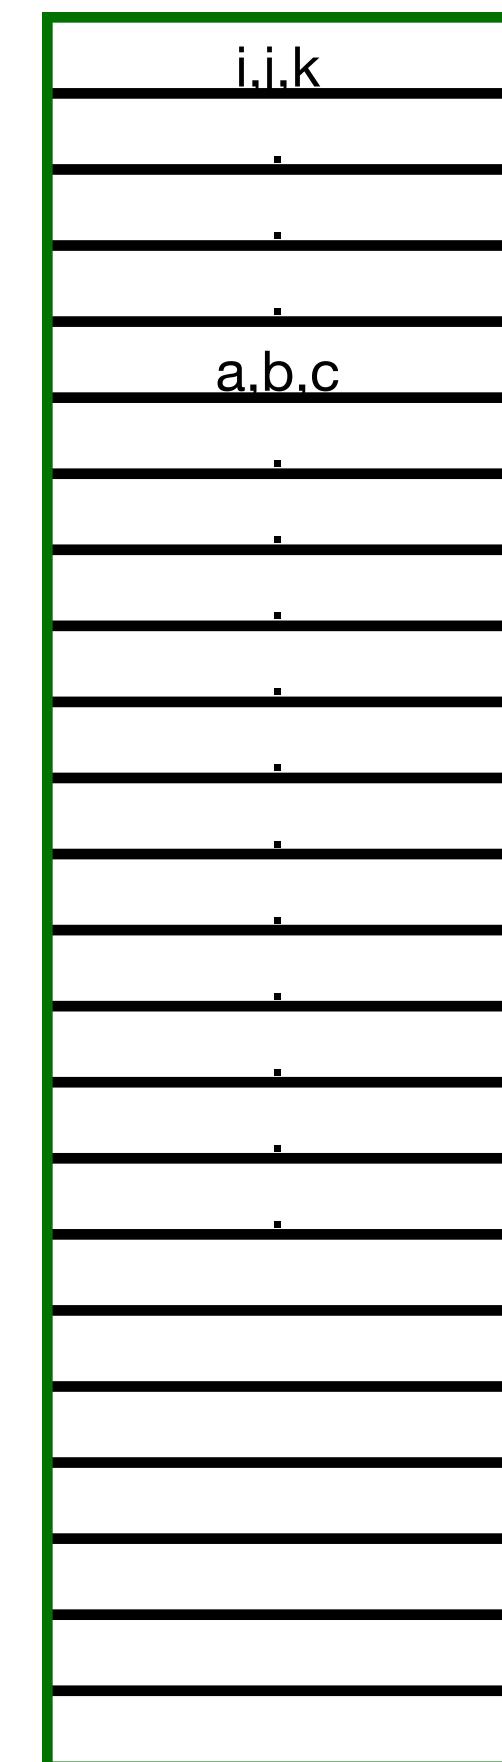


Vertices + Faces

Vertices

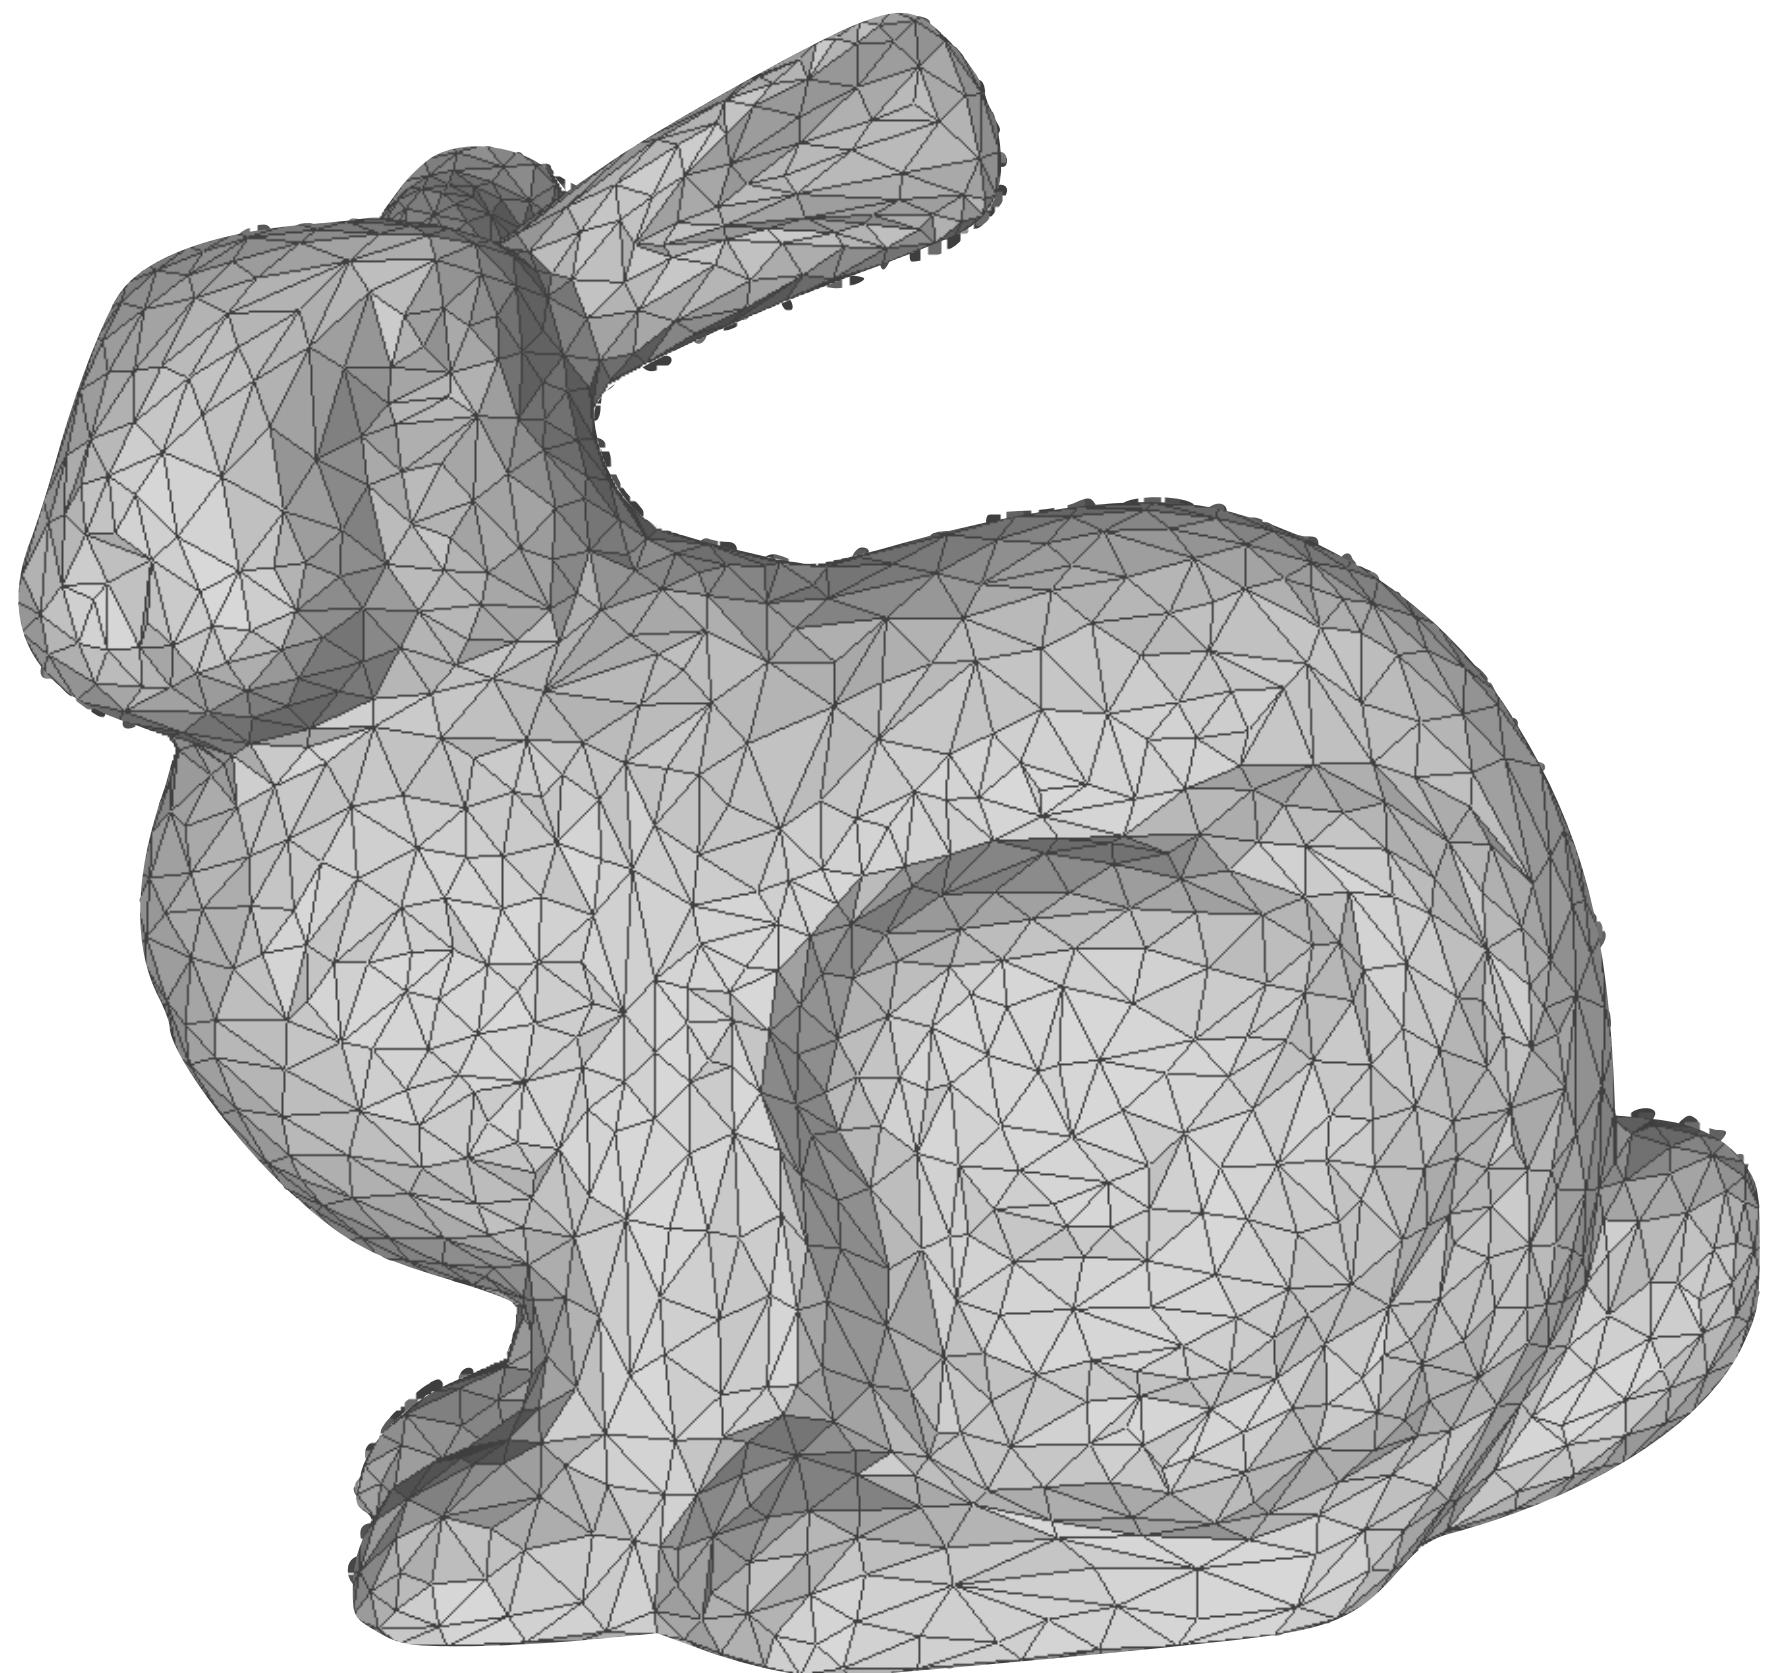


Positions of Vertices



Connectivity (indices of **three** vertices that make a ‘face’)

Is a point inside or outside the shape?



Vertices + Faces

Meshes can be “watertight”, i.e., if you filled the mesh with water, nothing would leak out

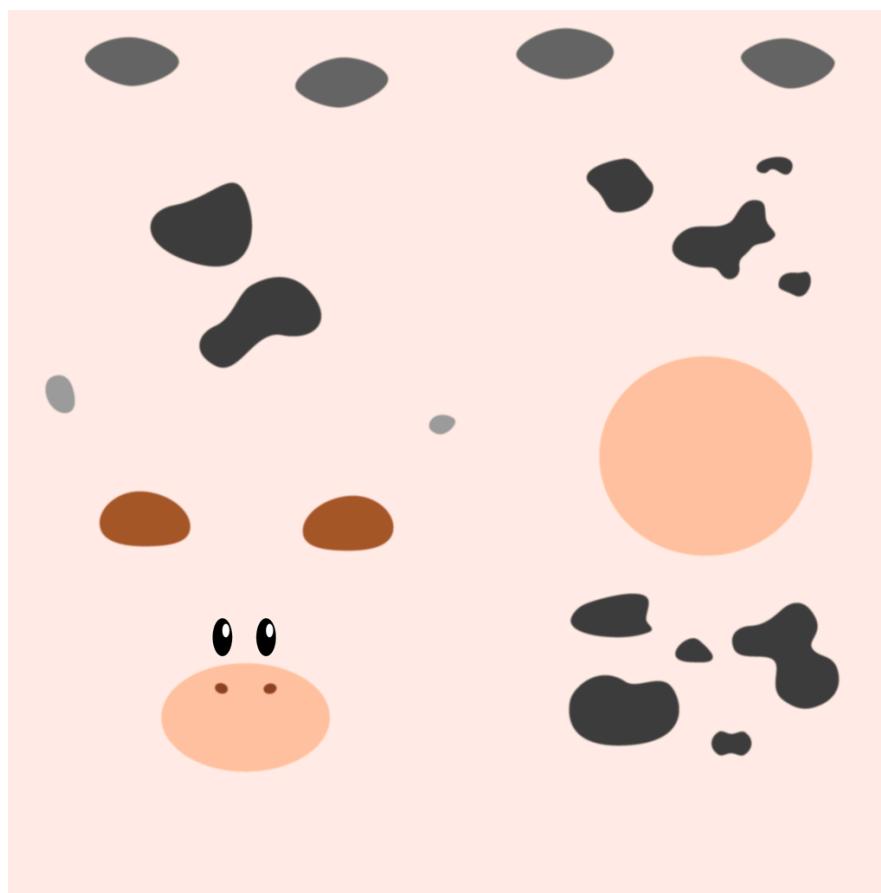
If mesh is watertight, can define “inside” and “outside”

If not, needs further processing, see “Generalized Winding Numbers”

Triangular Meshes

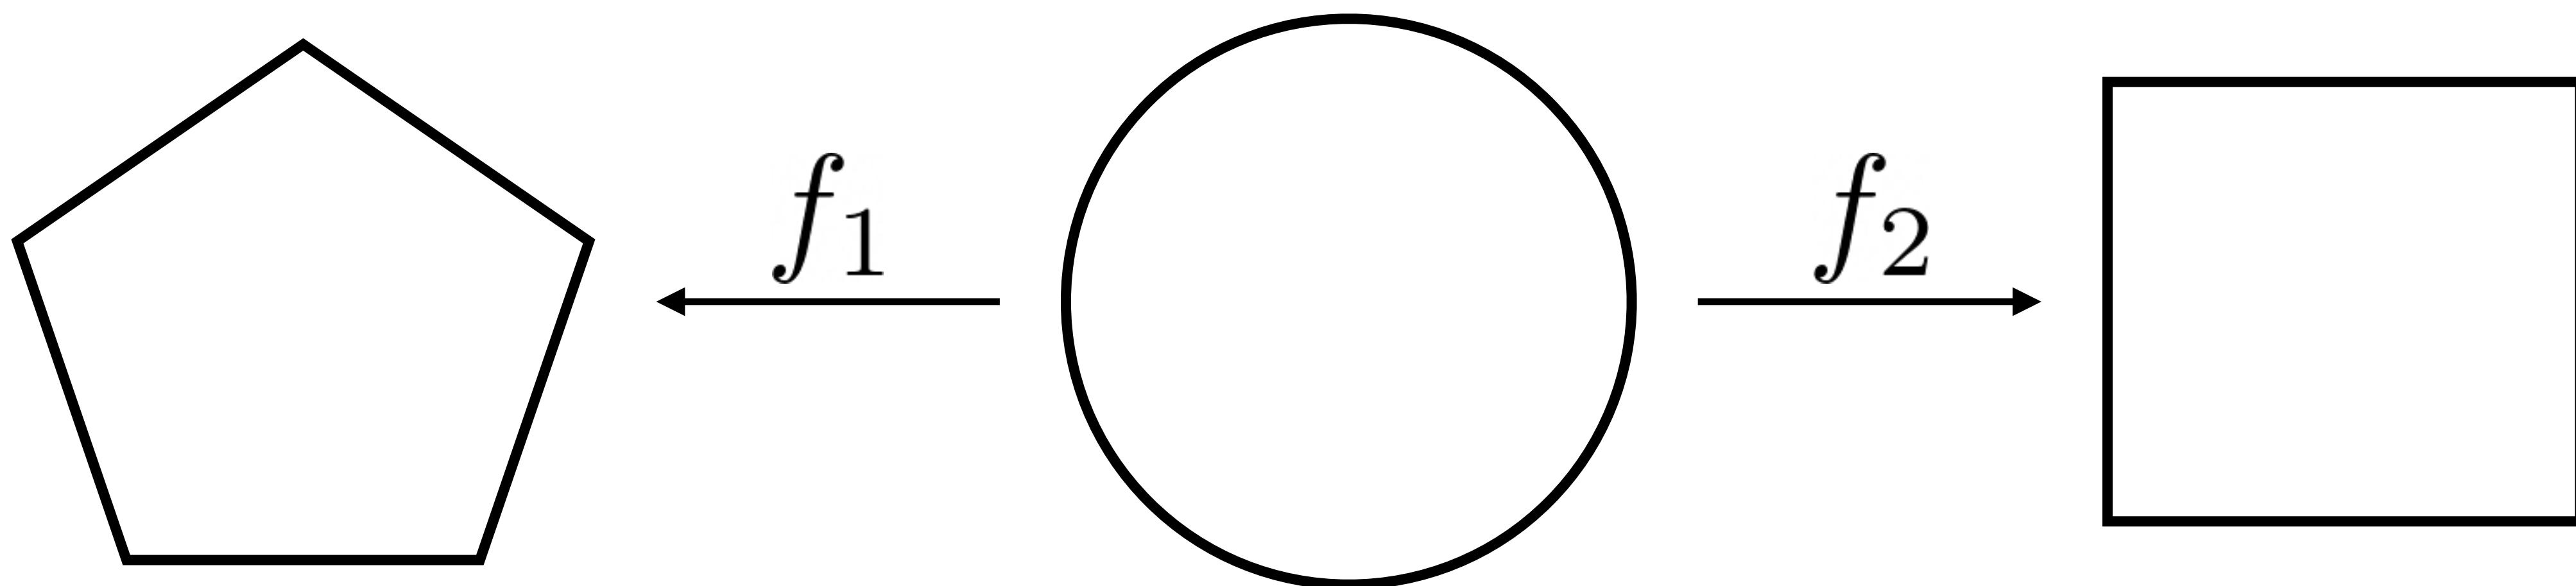


Efficient to compute
ray-triangle
intersections



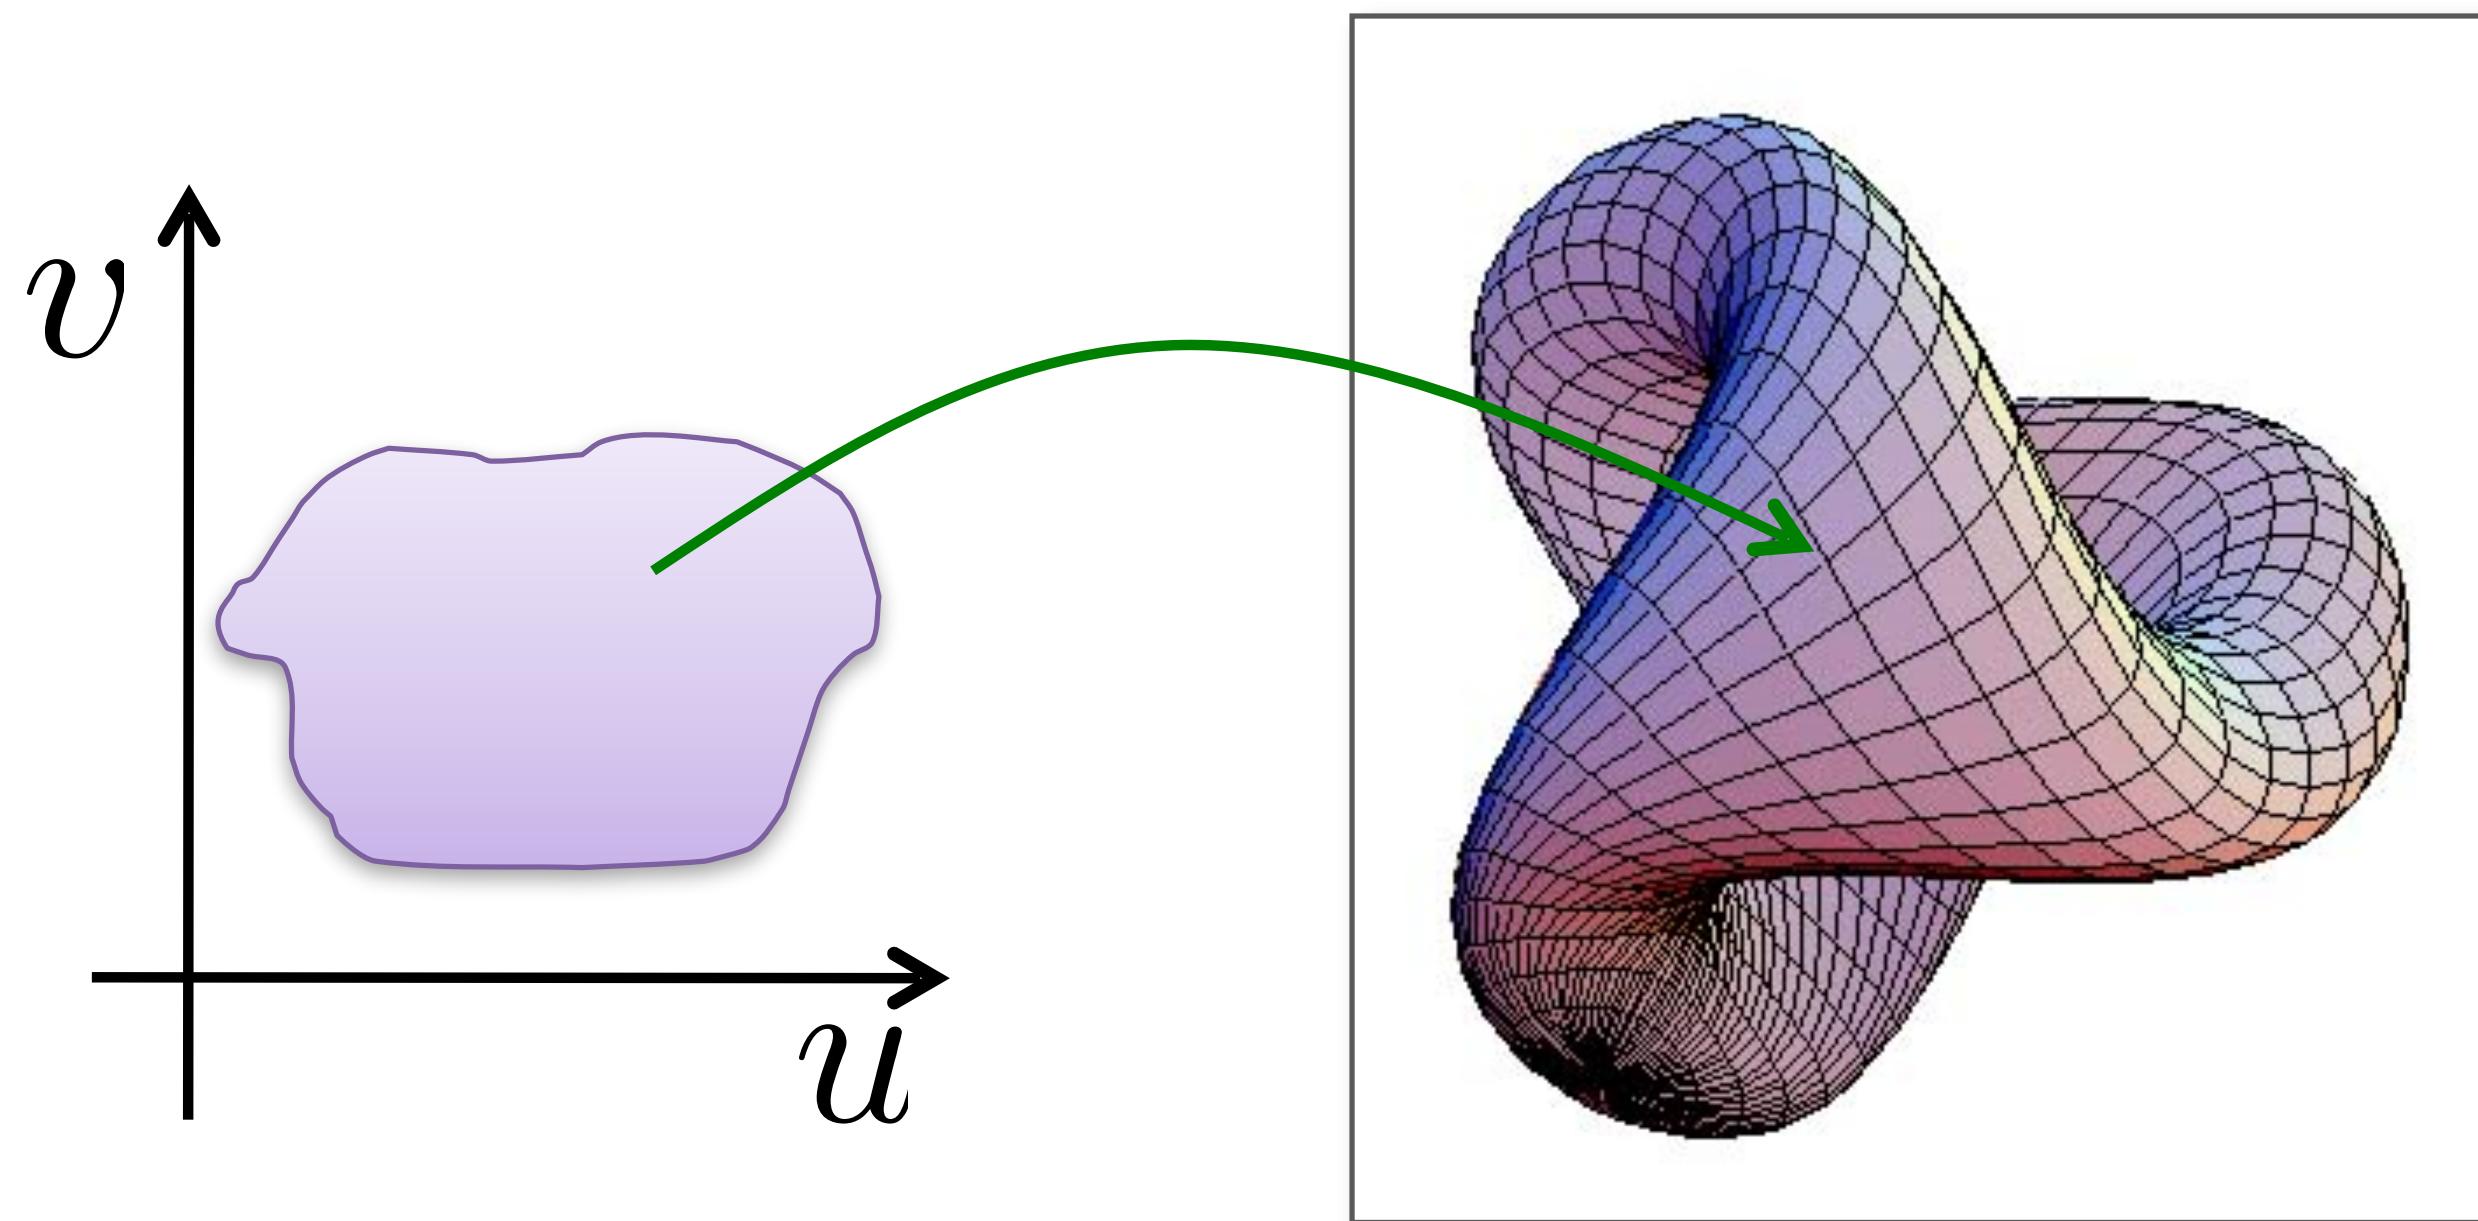
Easy to texture and render
(common representation
across graphics)

Parametric Surfaces (in 2D)



$$f(\mathbf{u}) = \mathbf{p} \in \mathbb{R}^2; \mathbf{u} \in \mathbb{S}^1$$

Parametric Surfaces



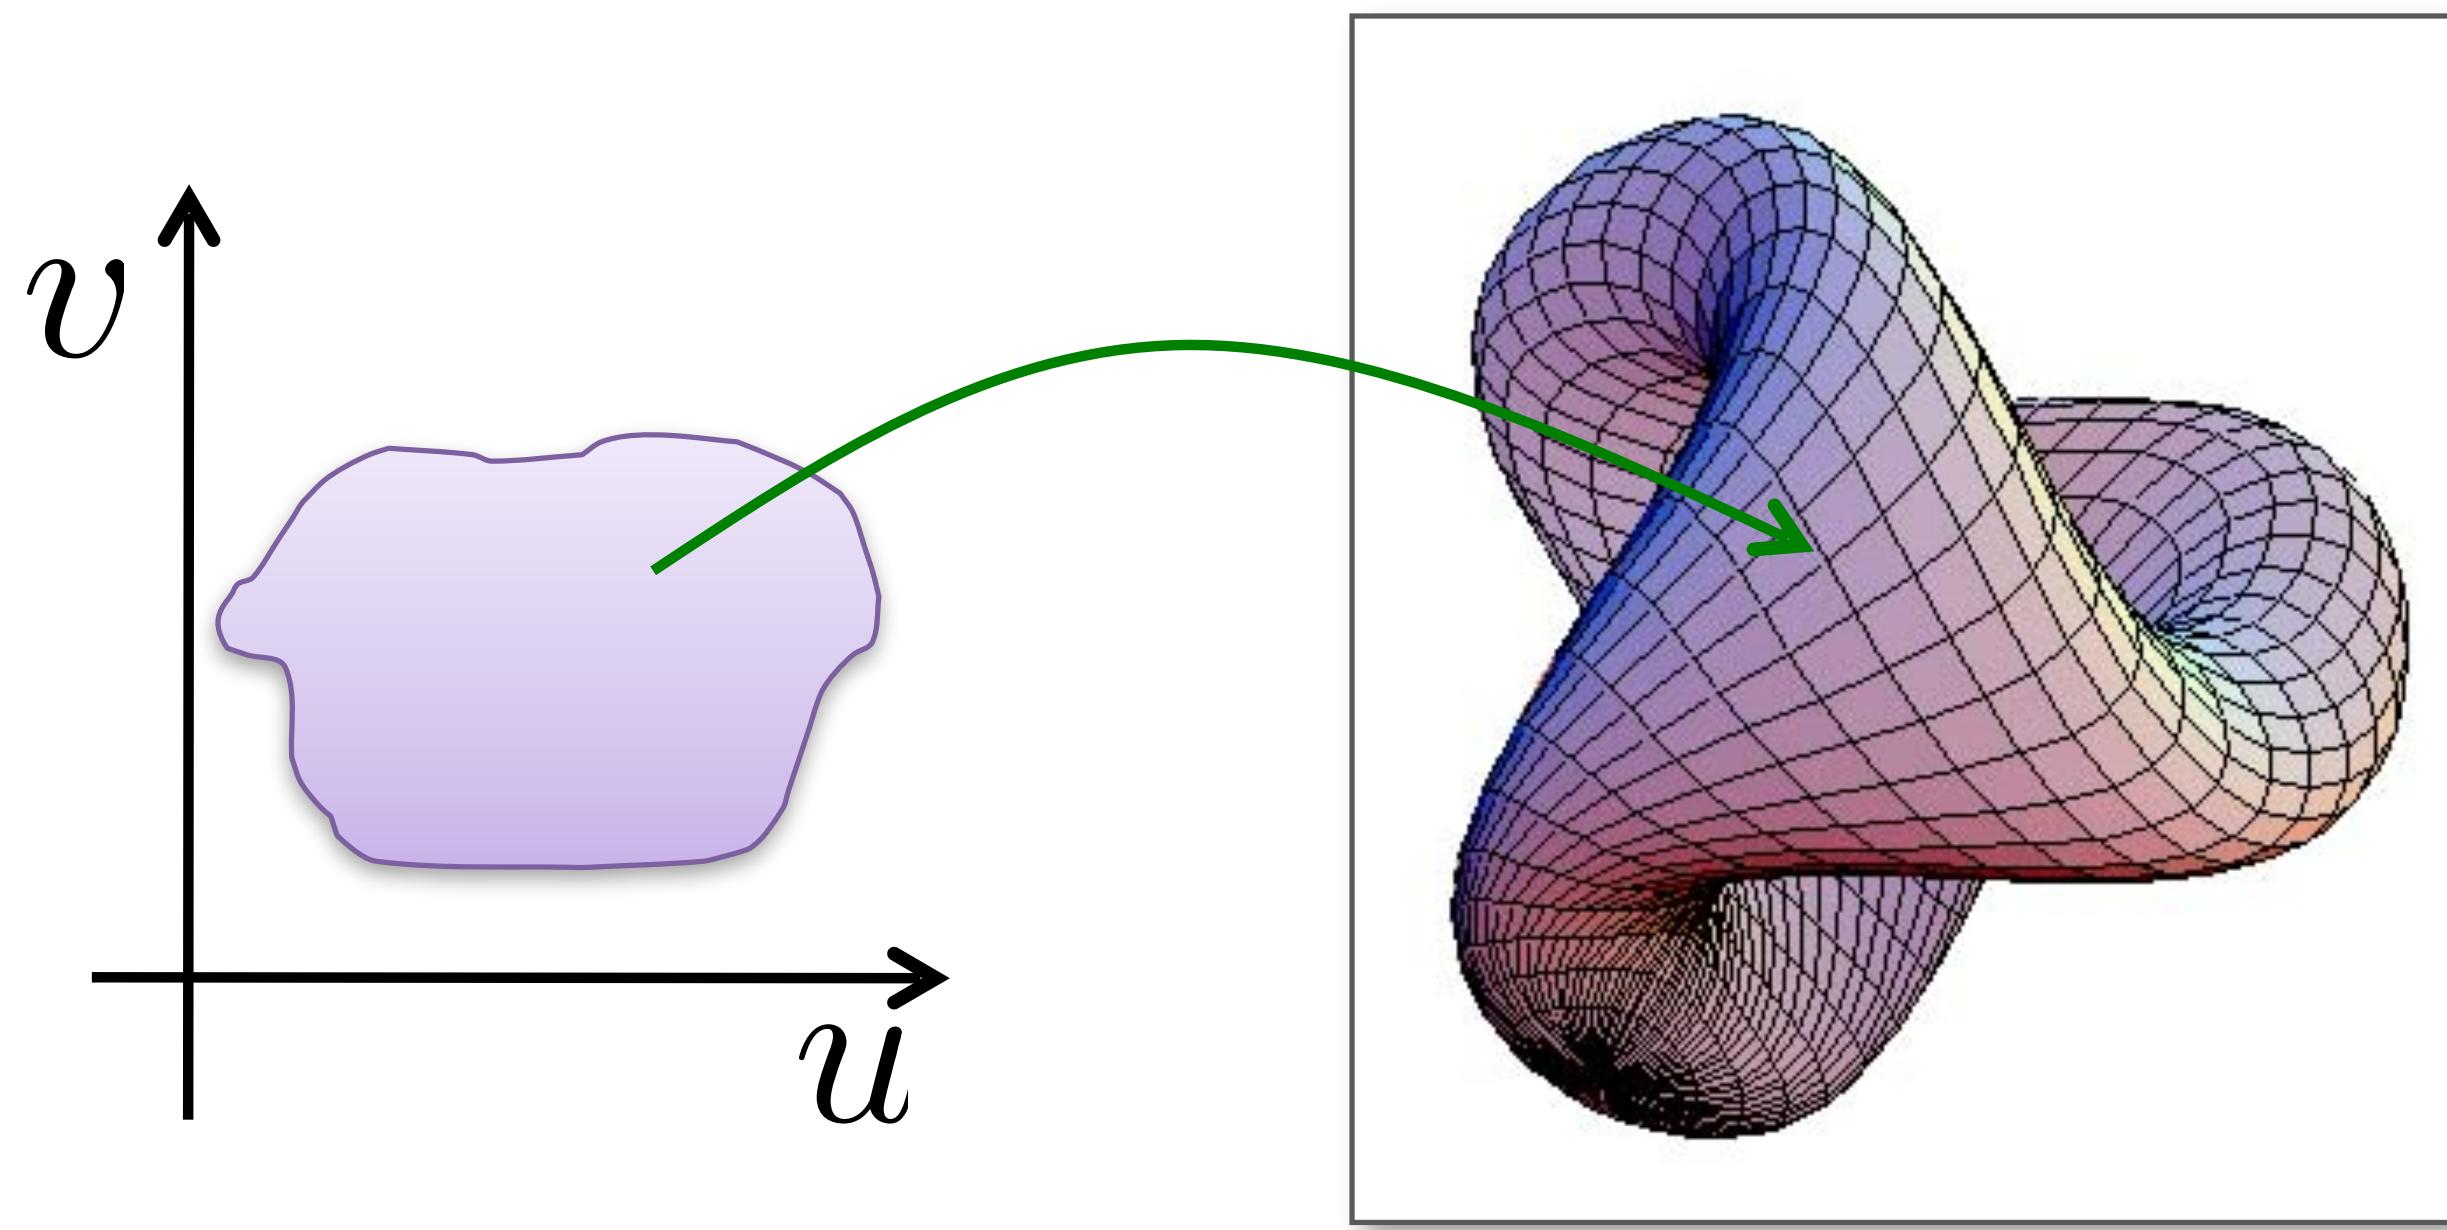
$$f(\mathbf{u}) = \mathbf{p} \in \mathbb{R}^3; \mathbf{u} \in \mathcal{M}$$

A continuous function f over a 2D manifold \mathcal{M} can parametrize a surface

The connectivity/topology is constrained to be same as \mathcal{M}

Disentangles discretization
(vertices/faces) from shape representation

Parametric Surfaces



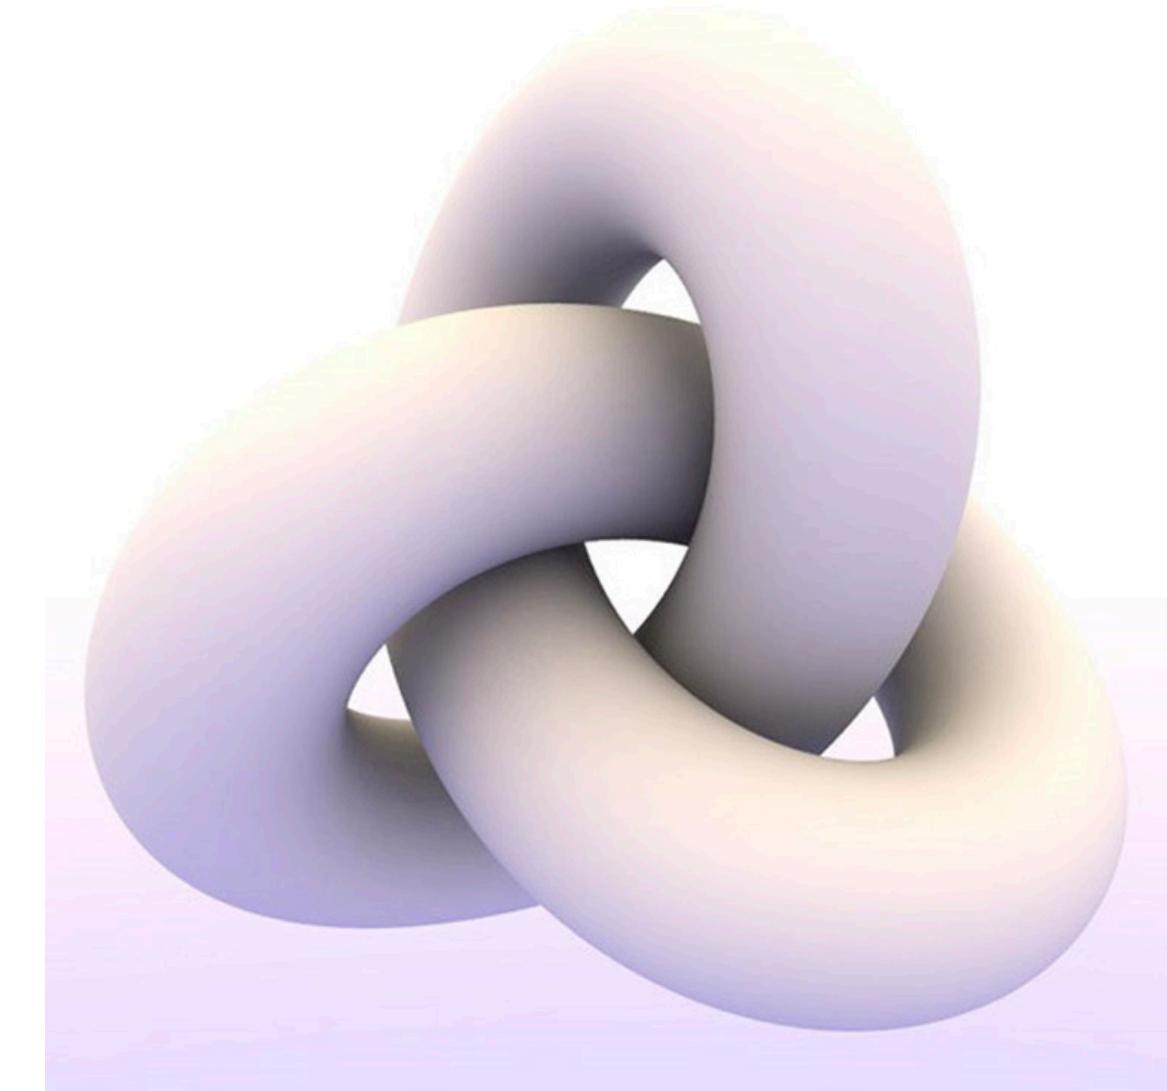
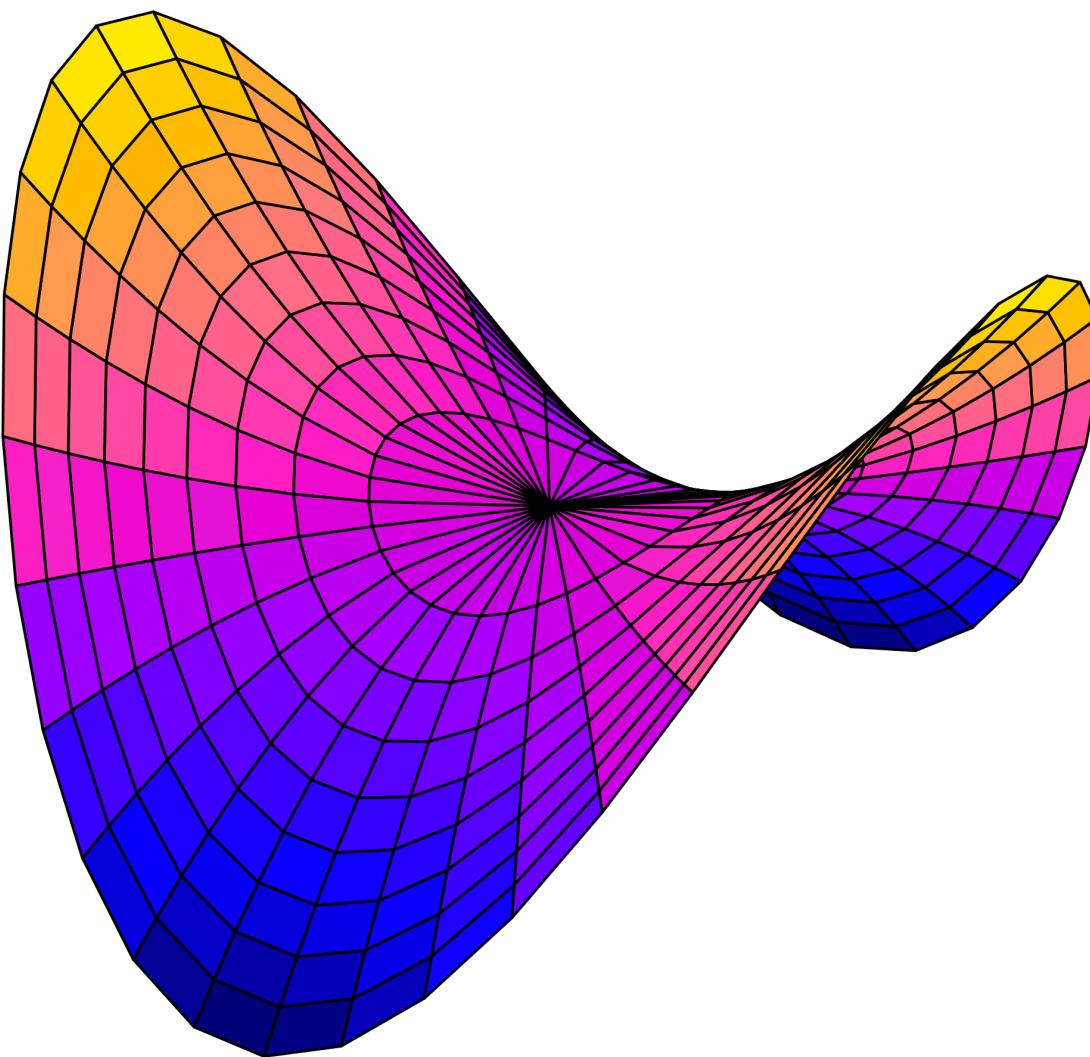
A continuous function f over a 2D manifold \mathcal{M} can parametrize a surface

```
def f(u):  
    ...  
Input:  
    u: N-dimensional vector from a 2D manifold  
Output:  
    p: 3D point  
    ...
```

$$f(\mathbf{u}) = \mathbf{p} \in \mathbb{R}^3; \mathbf{u} \in \mathcal{M}$$

Can be a fixed analytic function, parametrized family, or even a neural network!

Parametric Surfaces

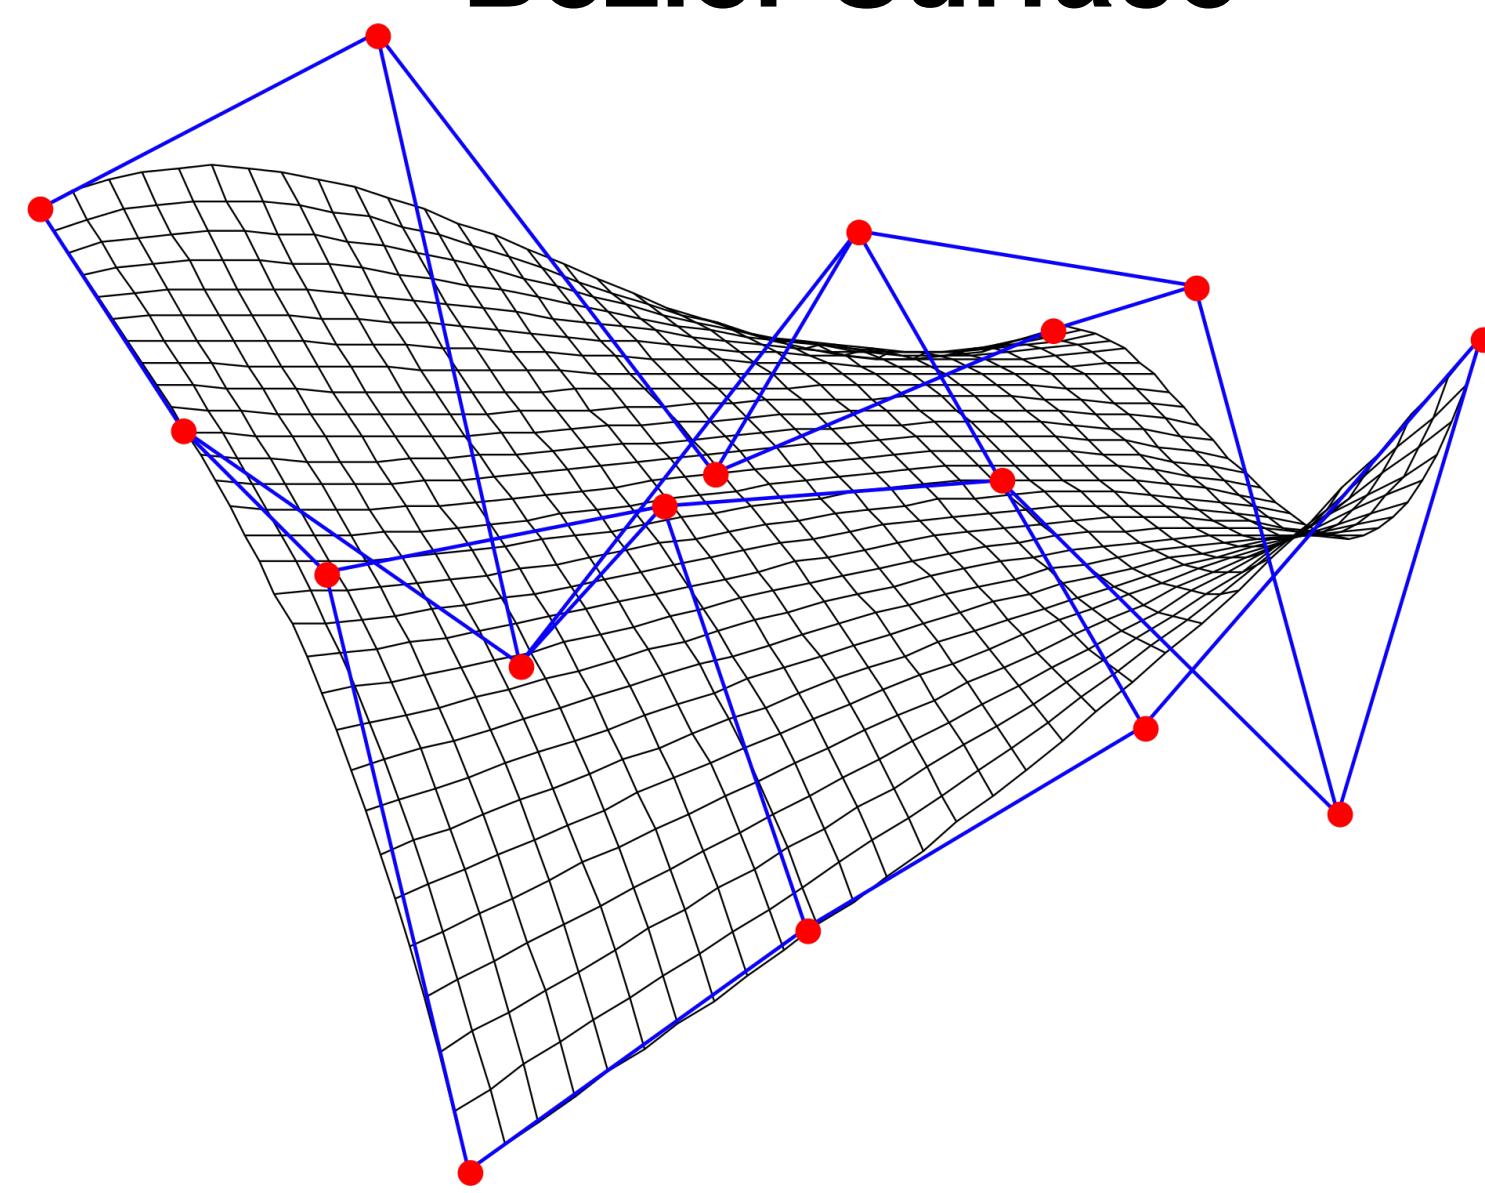


$$f((u, v)) = (u, v, v^2 - u^2)$$

```
def trefoil(u, v):
    x = r * np.sin(3 * u) / (2 + np.cos(v))
    y = r * (np.sin(u) + 2 * np.sin(2 * u)) / (2 + np.cos(v + np.pi * 2 / 3))
    z = r / 2 * (np.cos(u) - 2 * np.cos(2 * u)) * (2 + np.cos(v)) * (2 + np.cos(v + np.pi * 2 / 3)) / 4
```

Parametric Surfaces

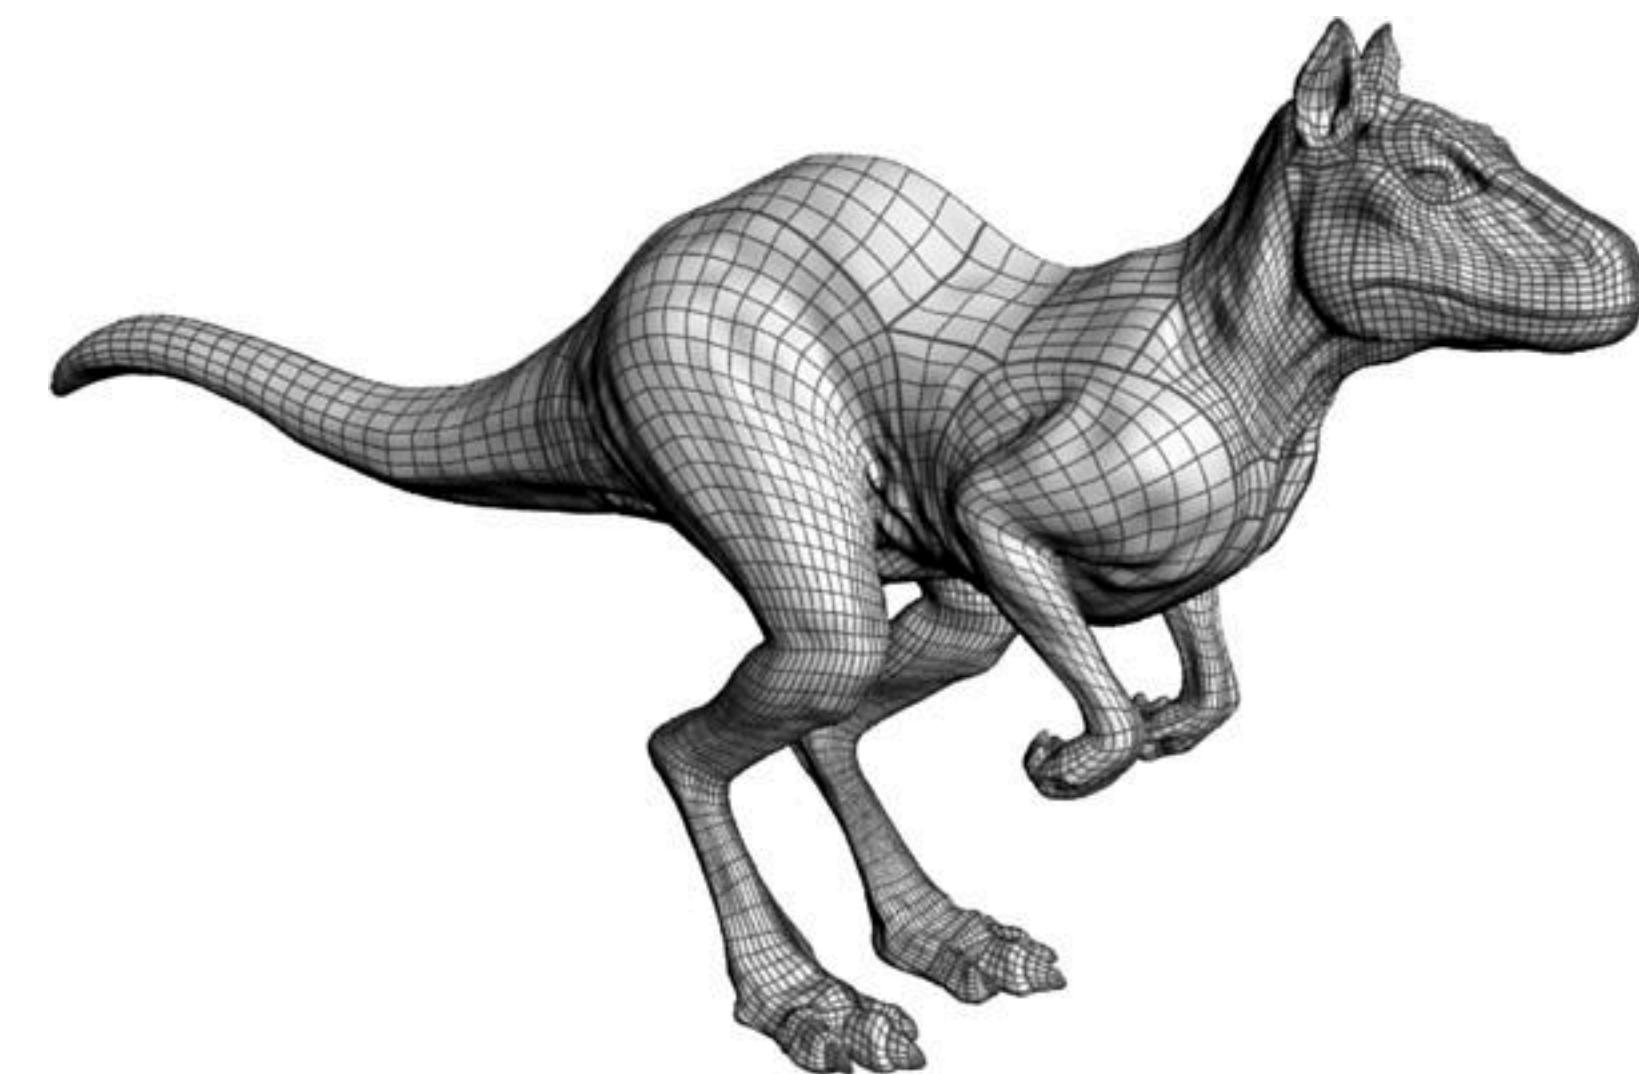
Bezier Surface



$$f((u, v)) = \sum_i^N \sum_j^M B_i^N(u) B_j^M(v) k_{i,j}$$

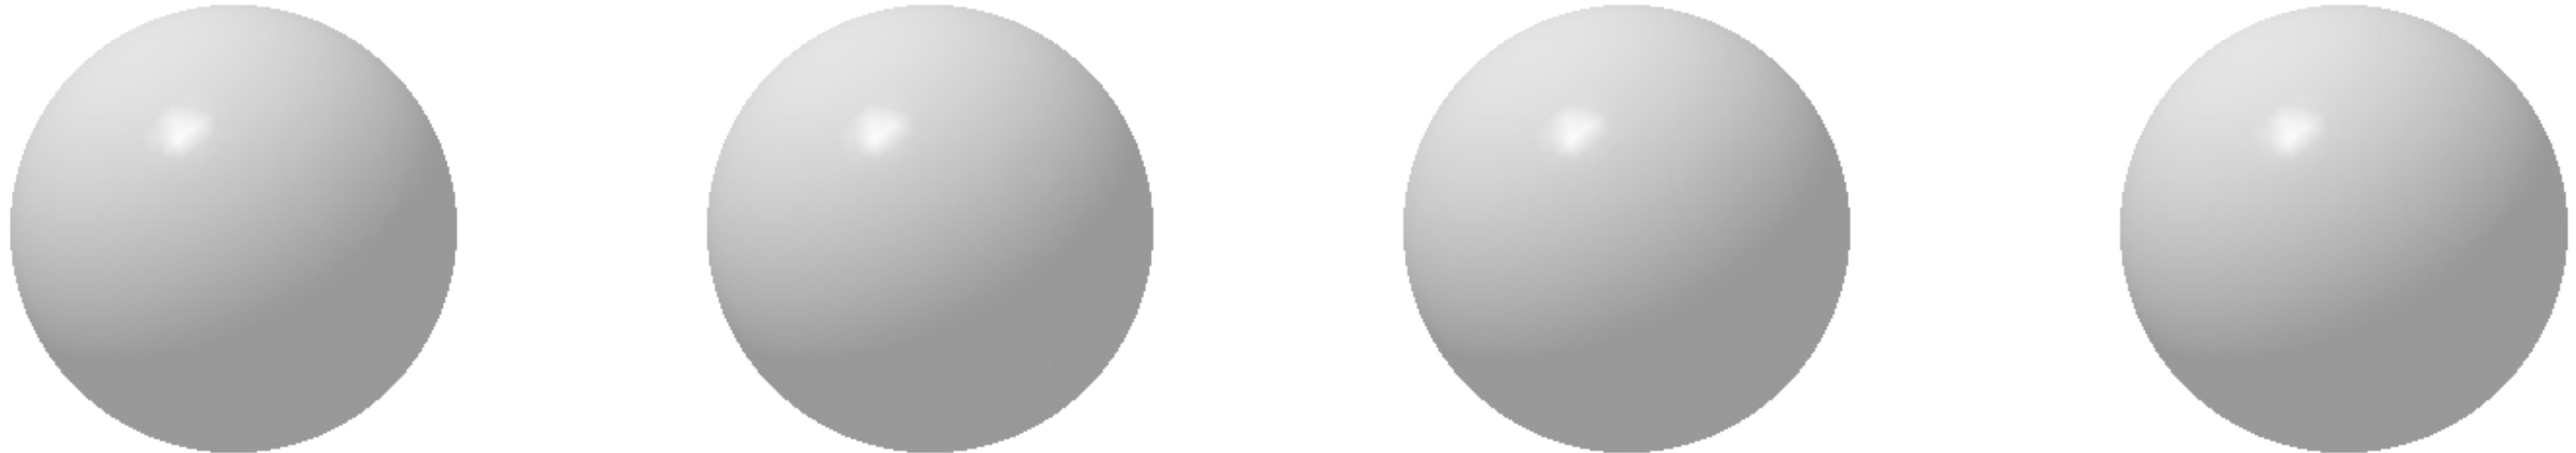
$$B_i^N(u) = \binom{N}{i} u^i (1-u)^{N-i}$$

M^*N control points can determine a smooth surface in 3D



Commonly used for designing meshes

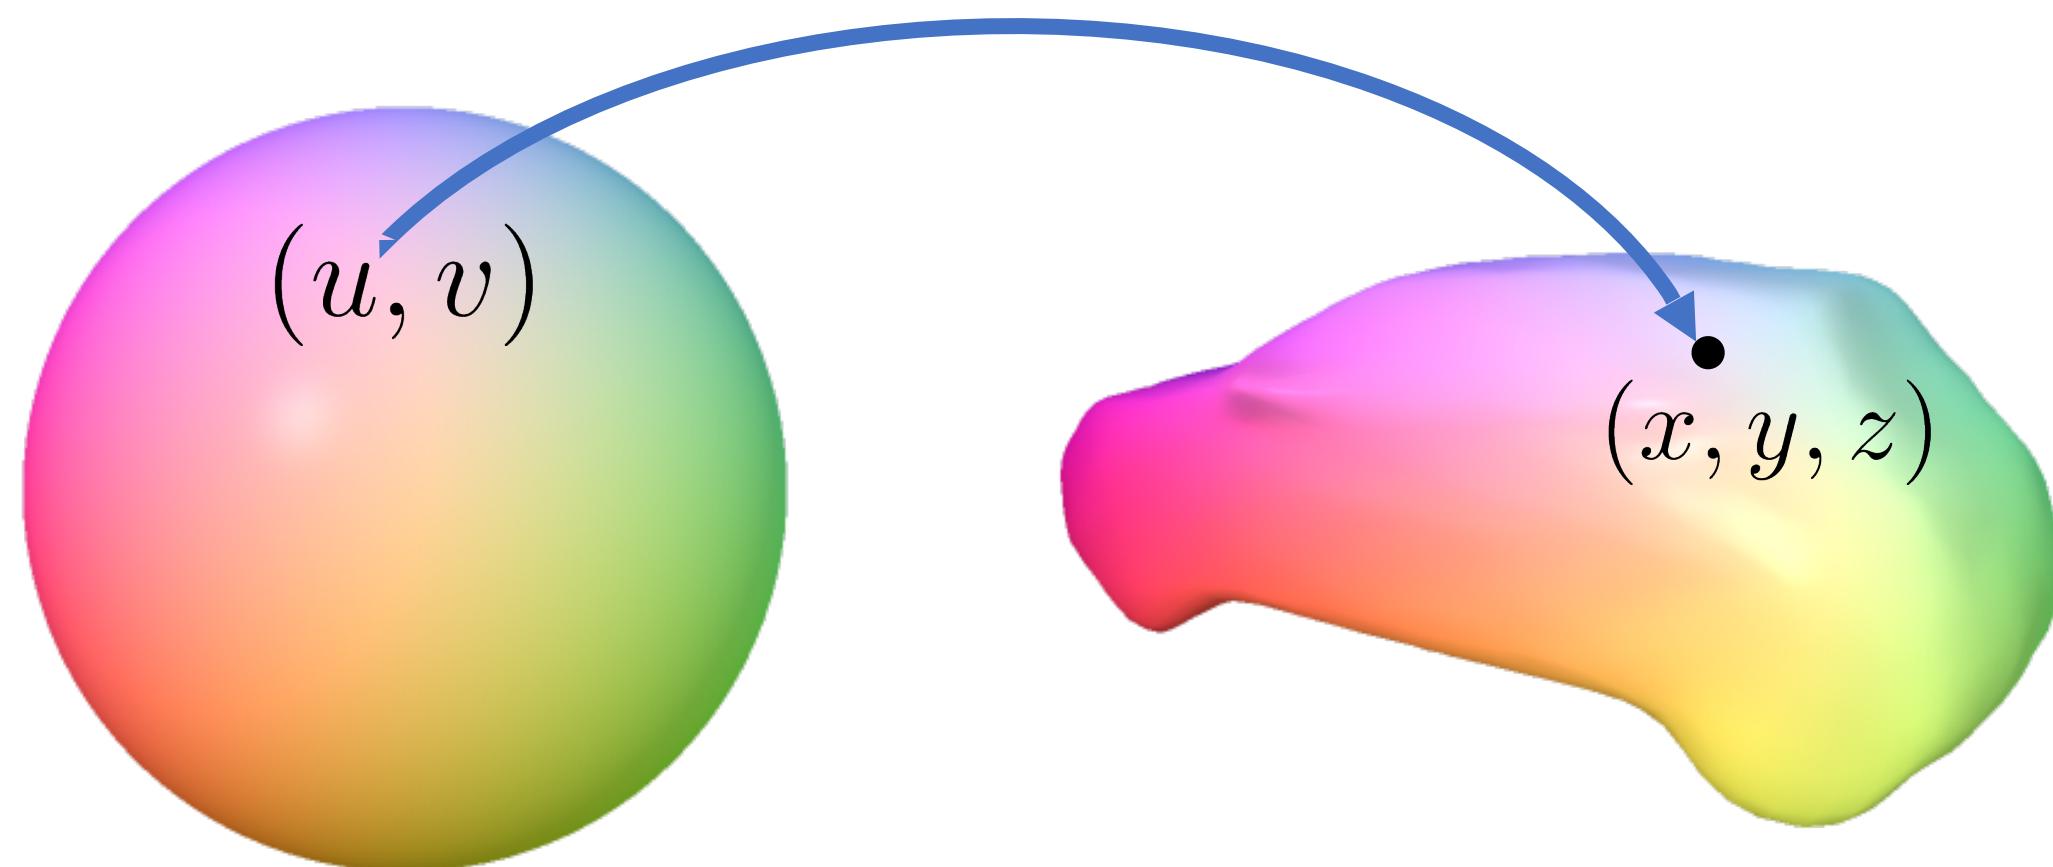
Parametric Surfaces



$$f_{\theta}(\mathbf{u}) = \mathbf{p} \in \mathbb{R}^3; \quad \mathbf{u} \in \mathbb{S}^2$$

Neural network with parameters θ

Parametric Surfaces



A continuous function f over a 2D manifold \mathcal{M} can parametrize a surface

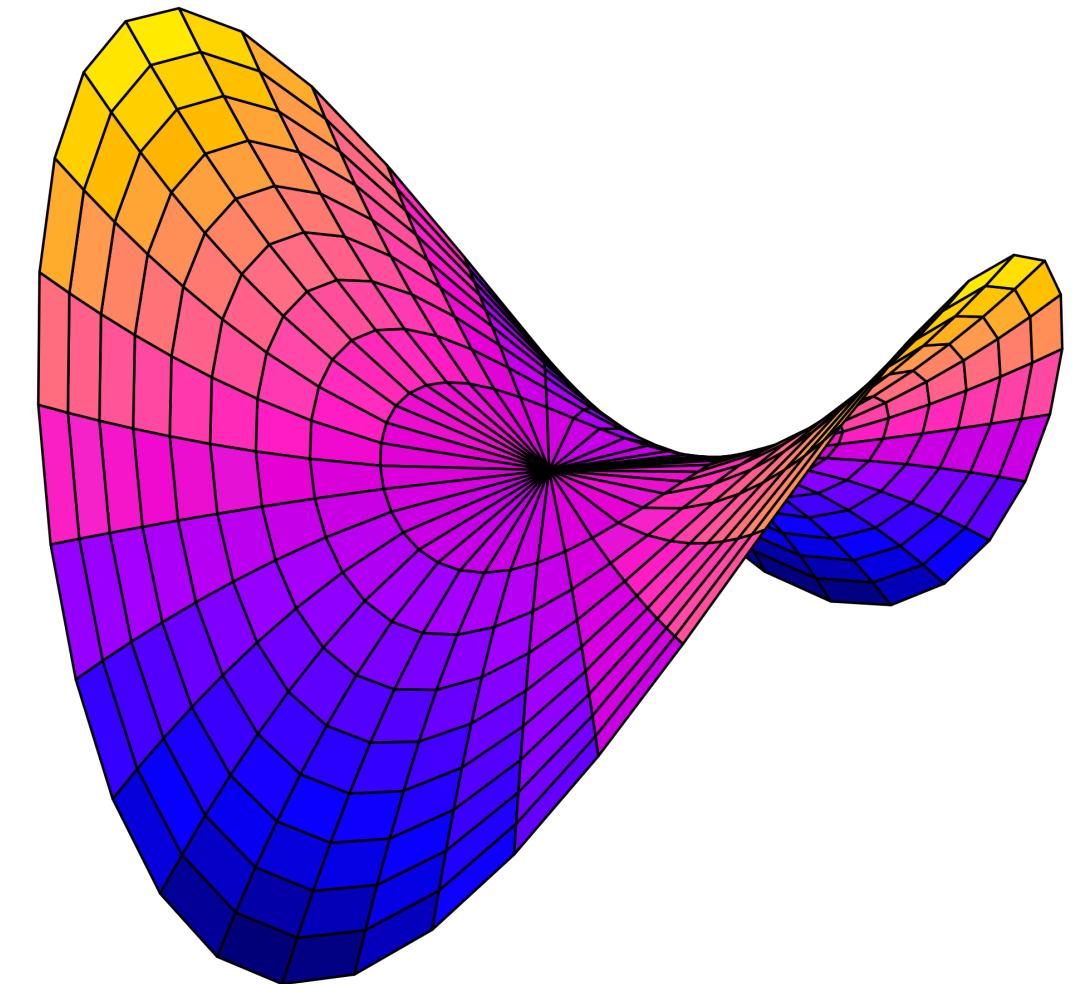
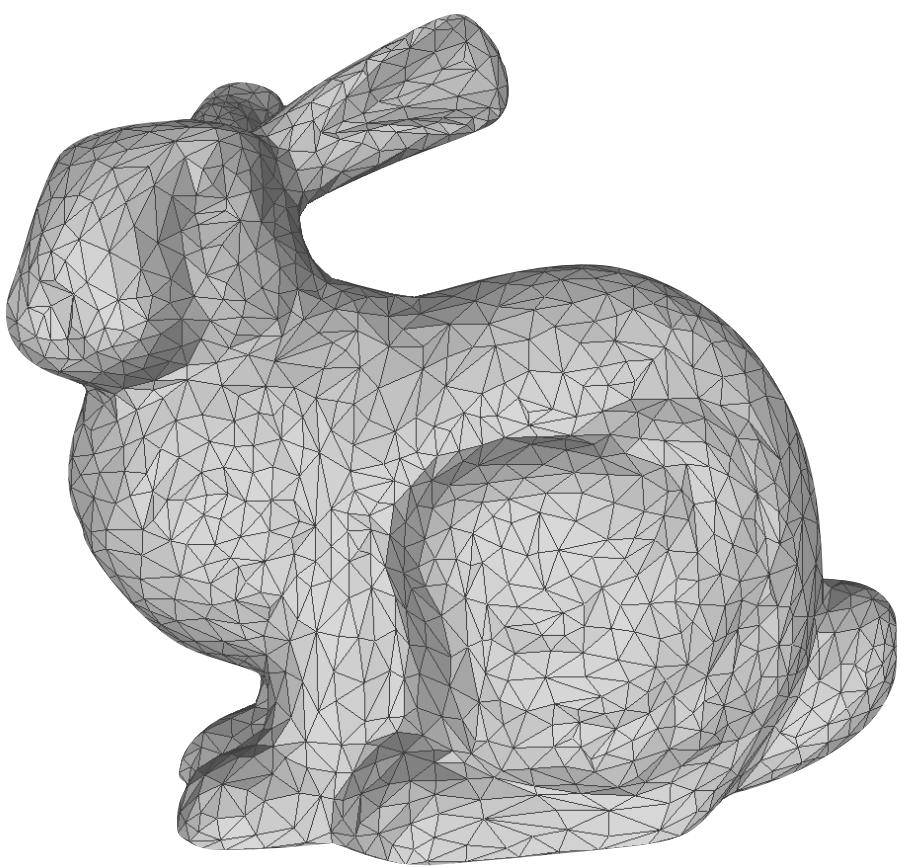
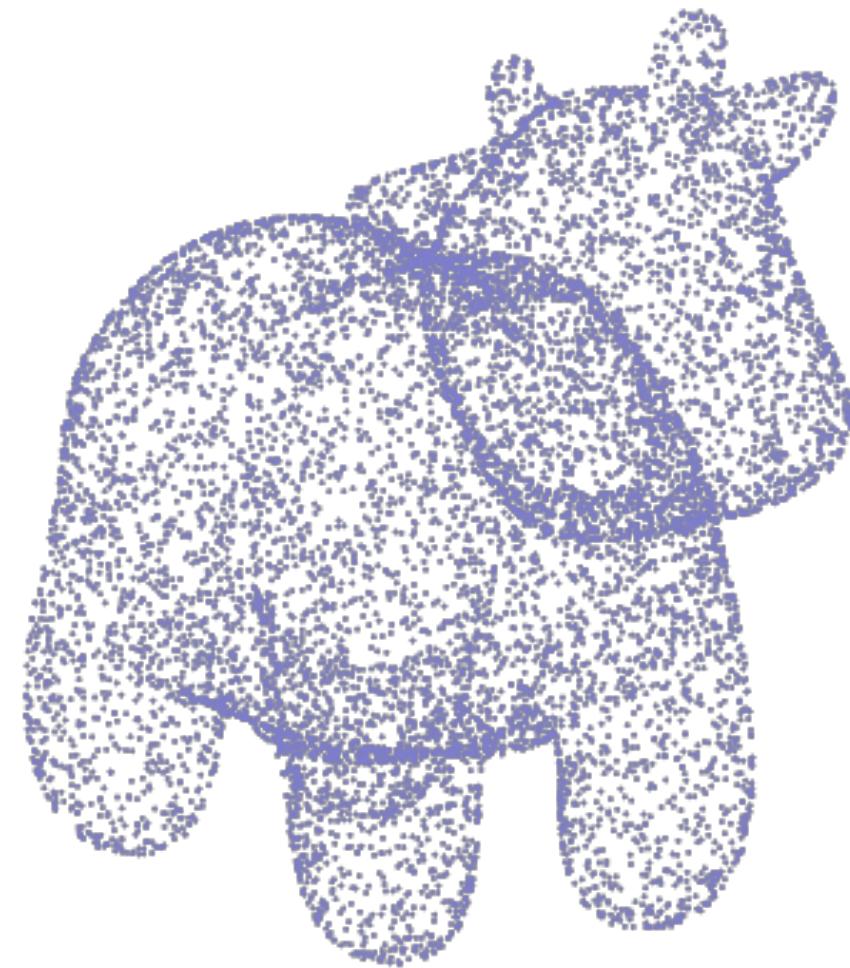
Easy to sample points on surface

Can do computation in the domain
 \mathcal{M}

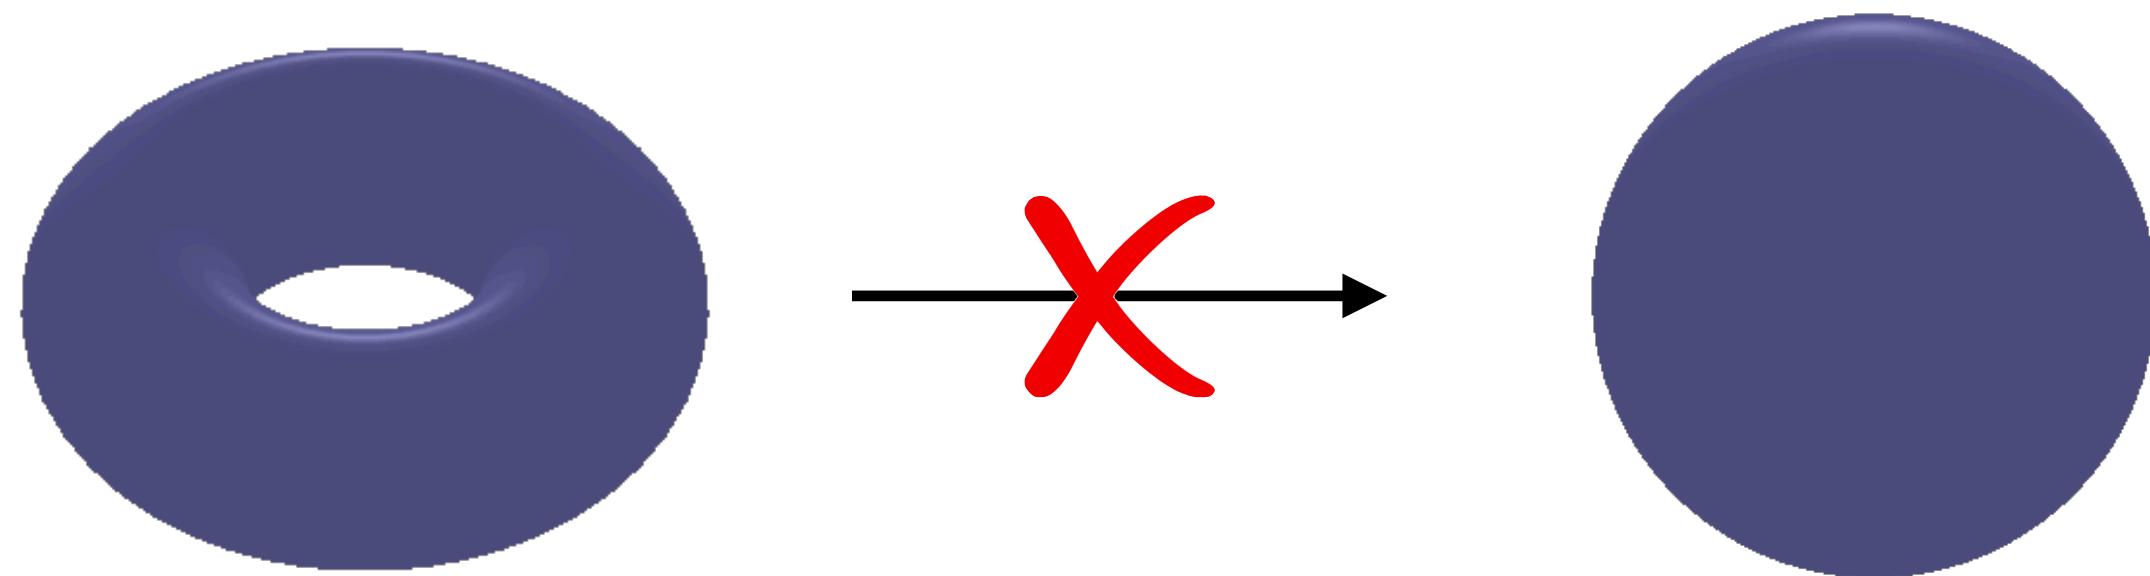
Difficult to reason about global structure (e.g. is a query point \mathbf{q} outside the shape?)

$$f(\mathbf{u}) = \mathbf{p} \in \mathbb{R}^3; \mathbf{u} \in \mathcal{M}$$

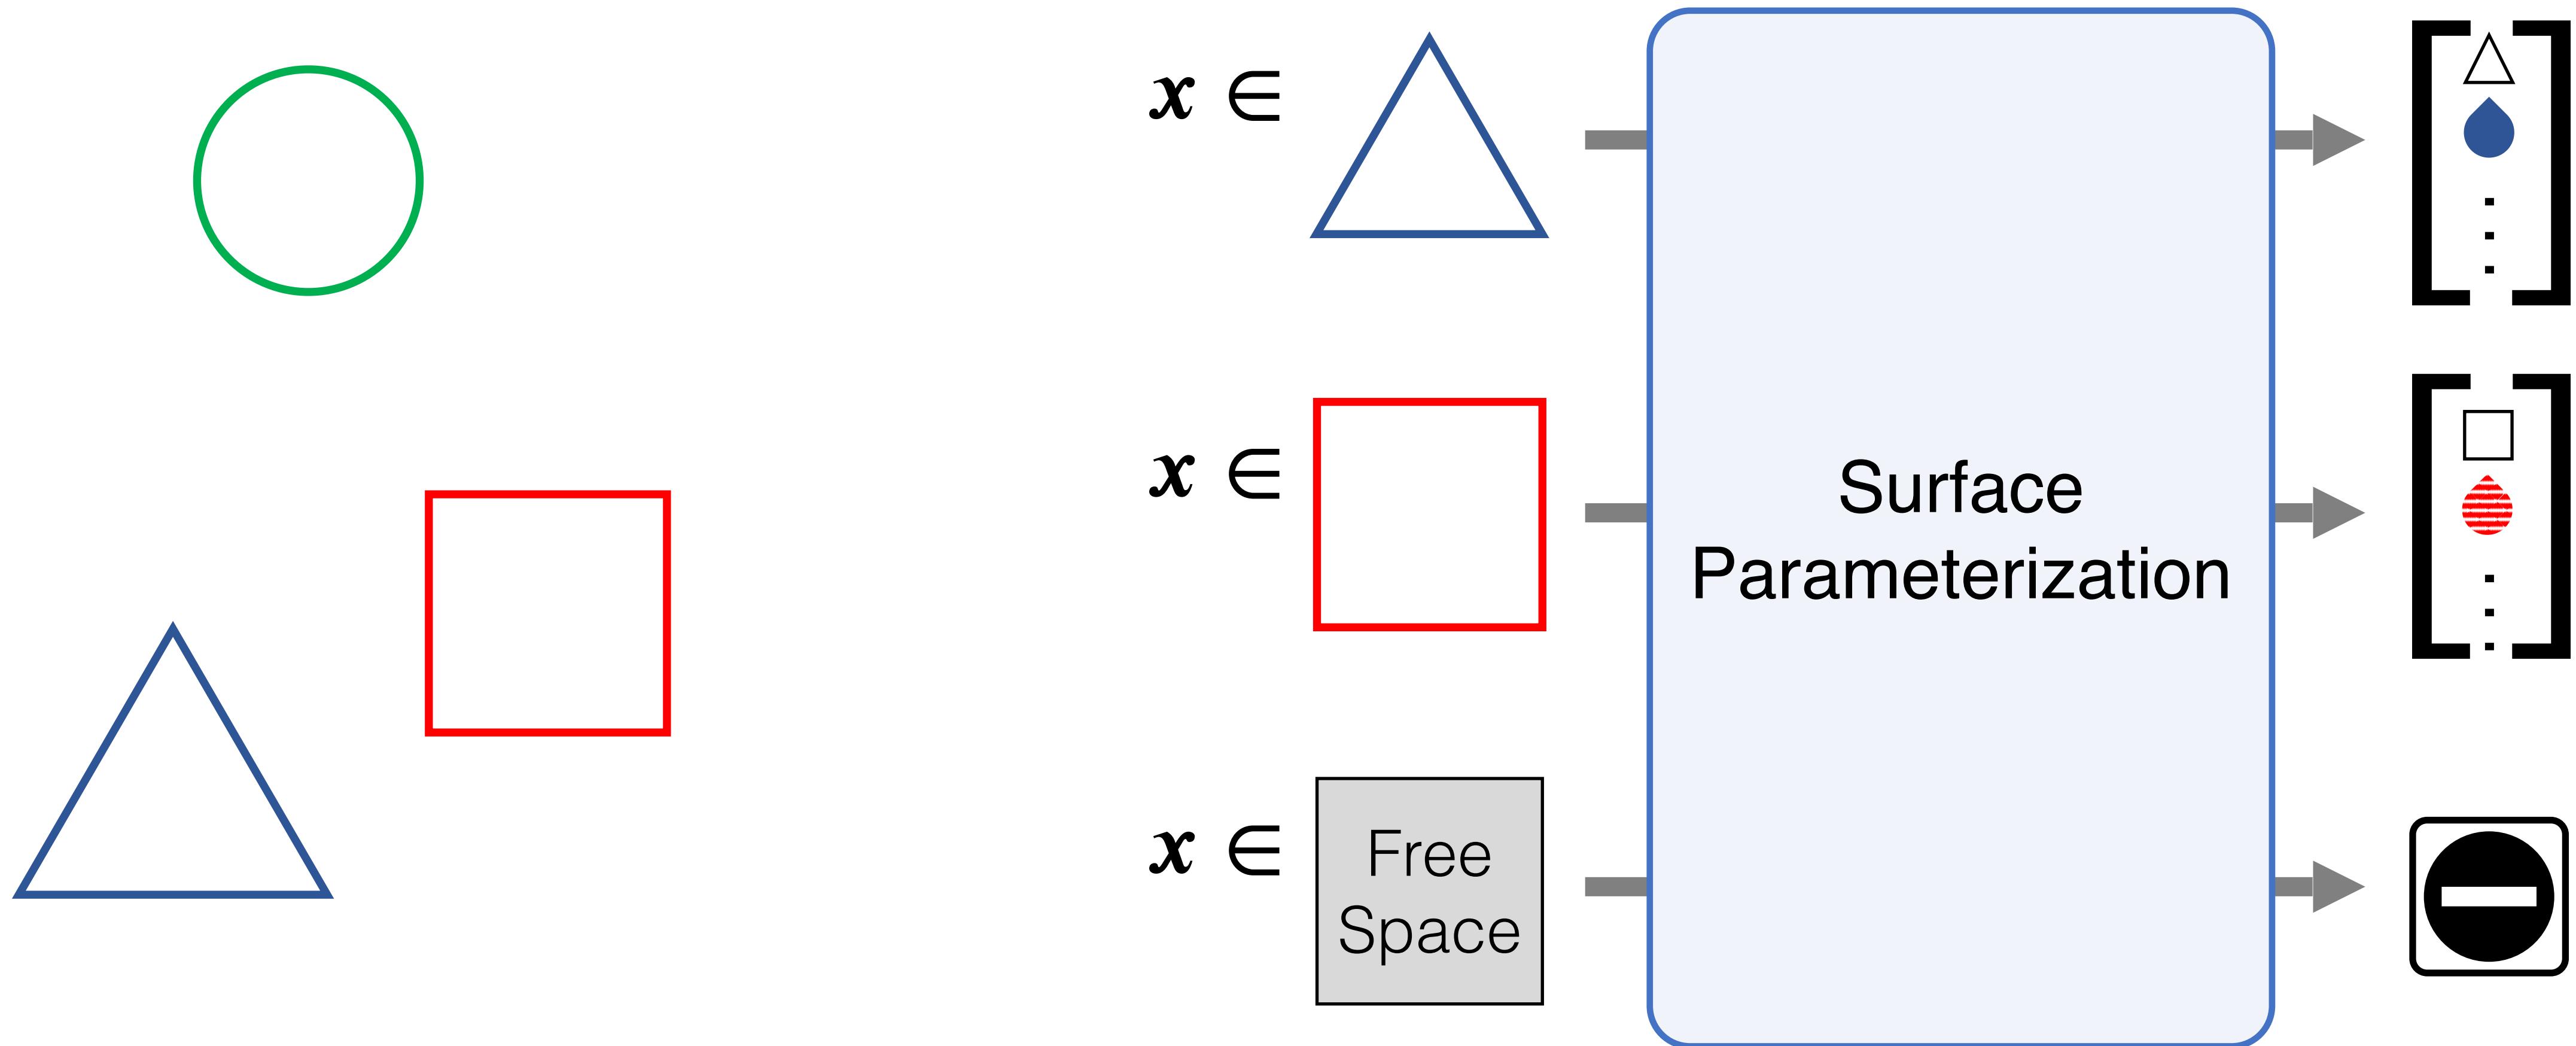
Difficult to analytically render (ray-surface intersection is hard for arbitrary f)



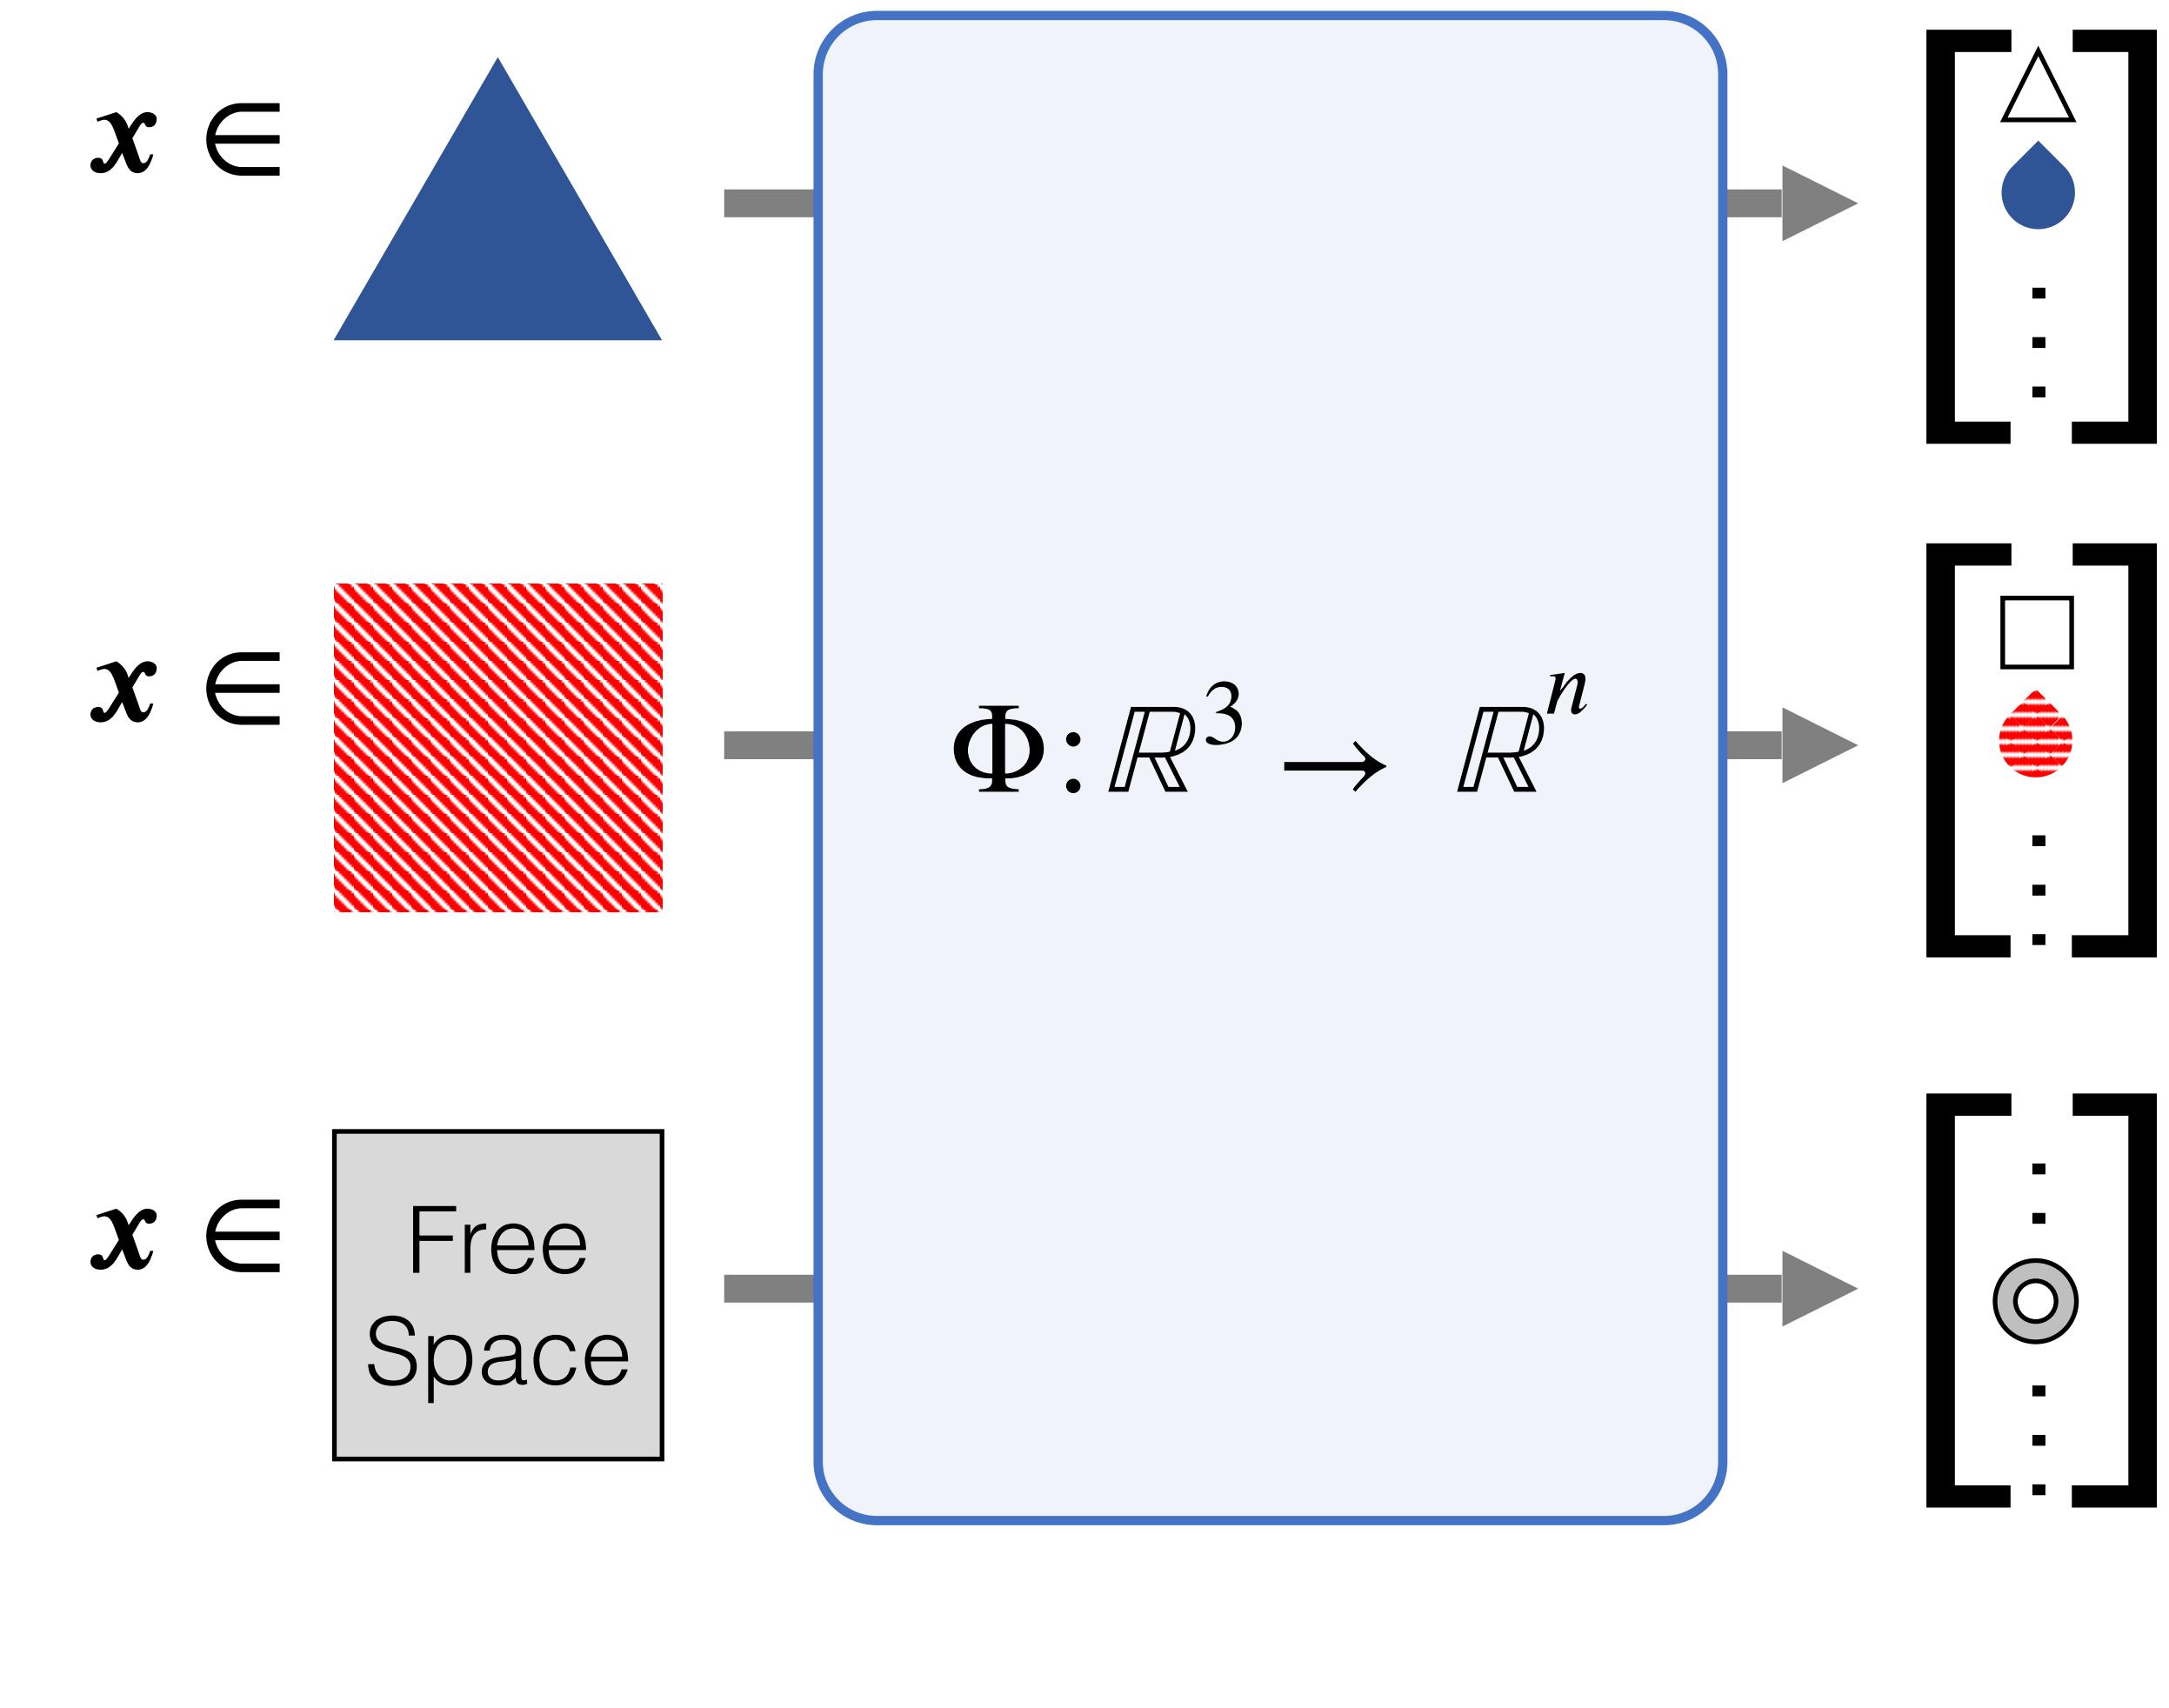
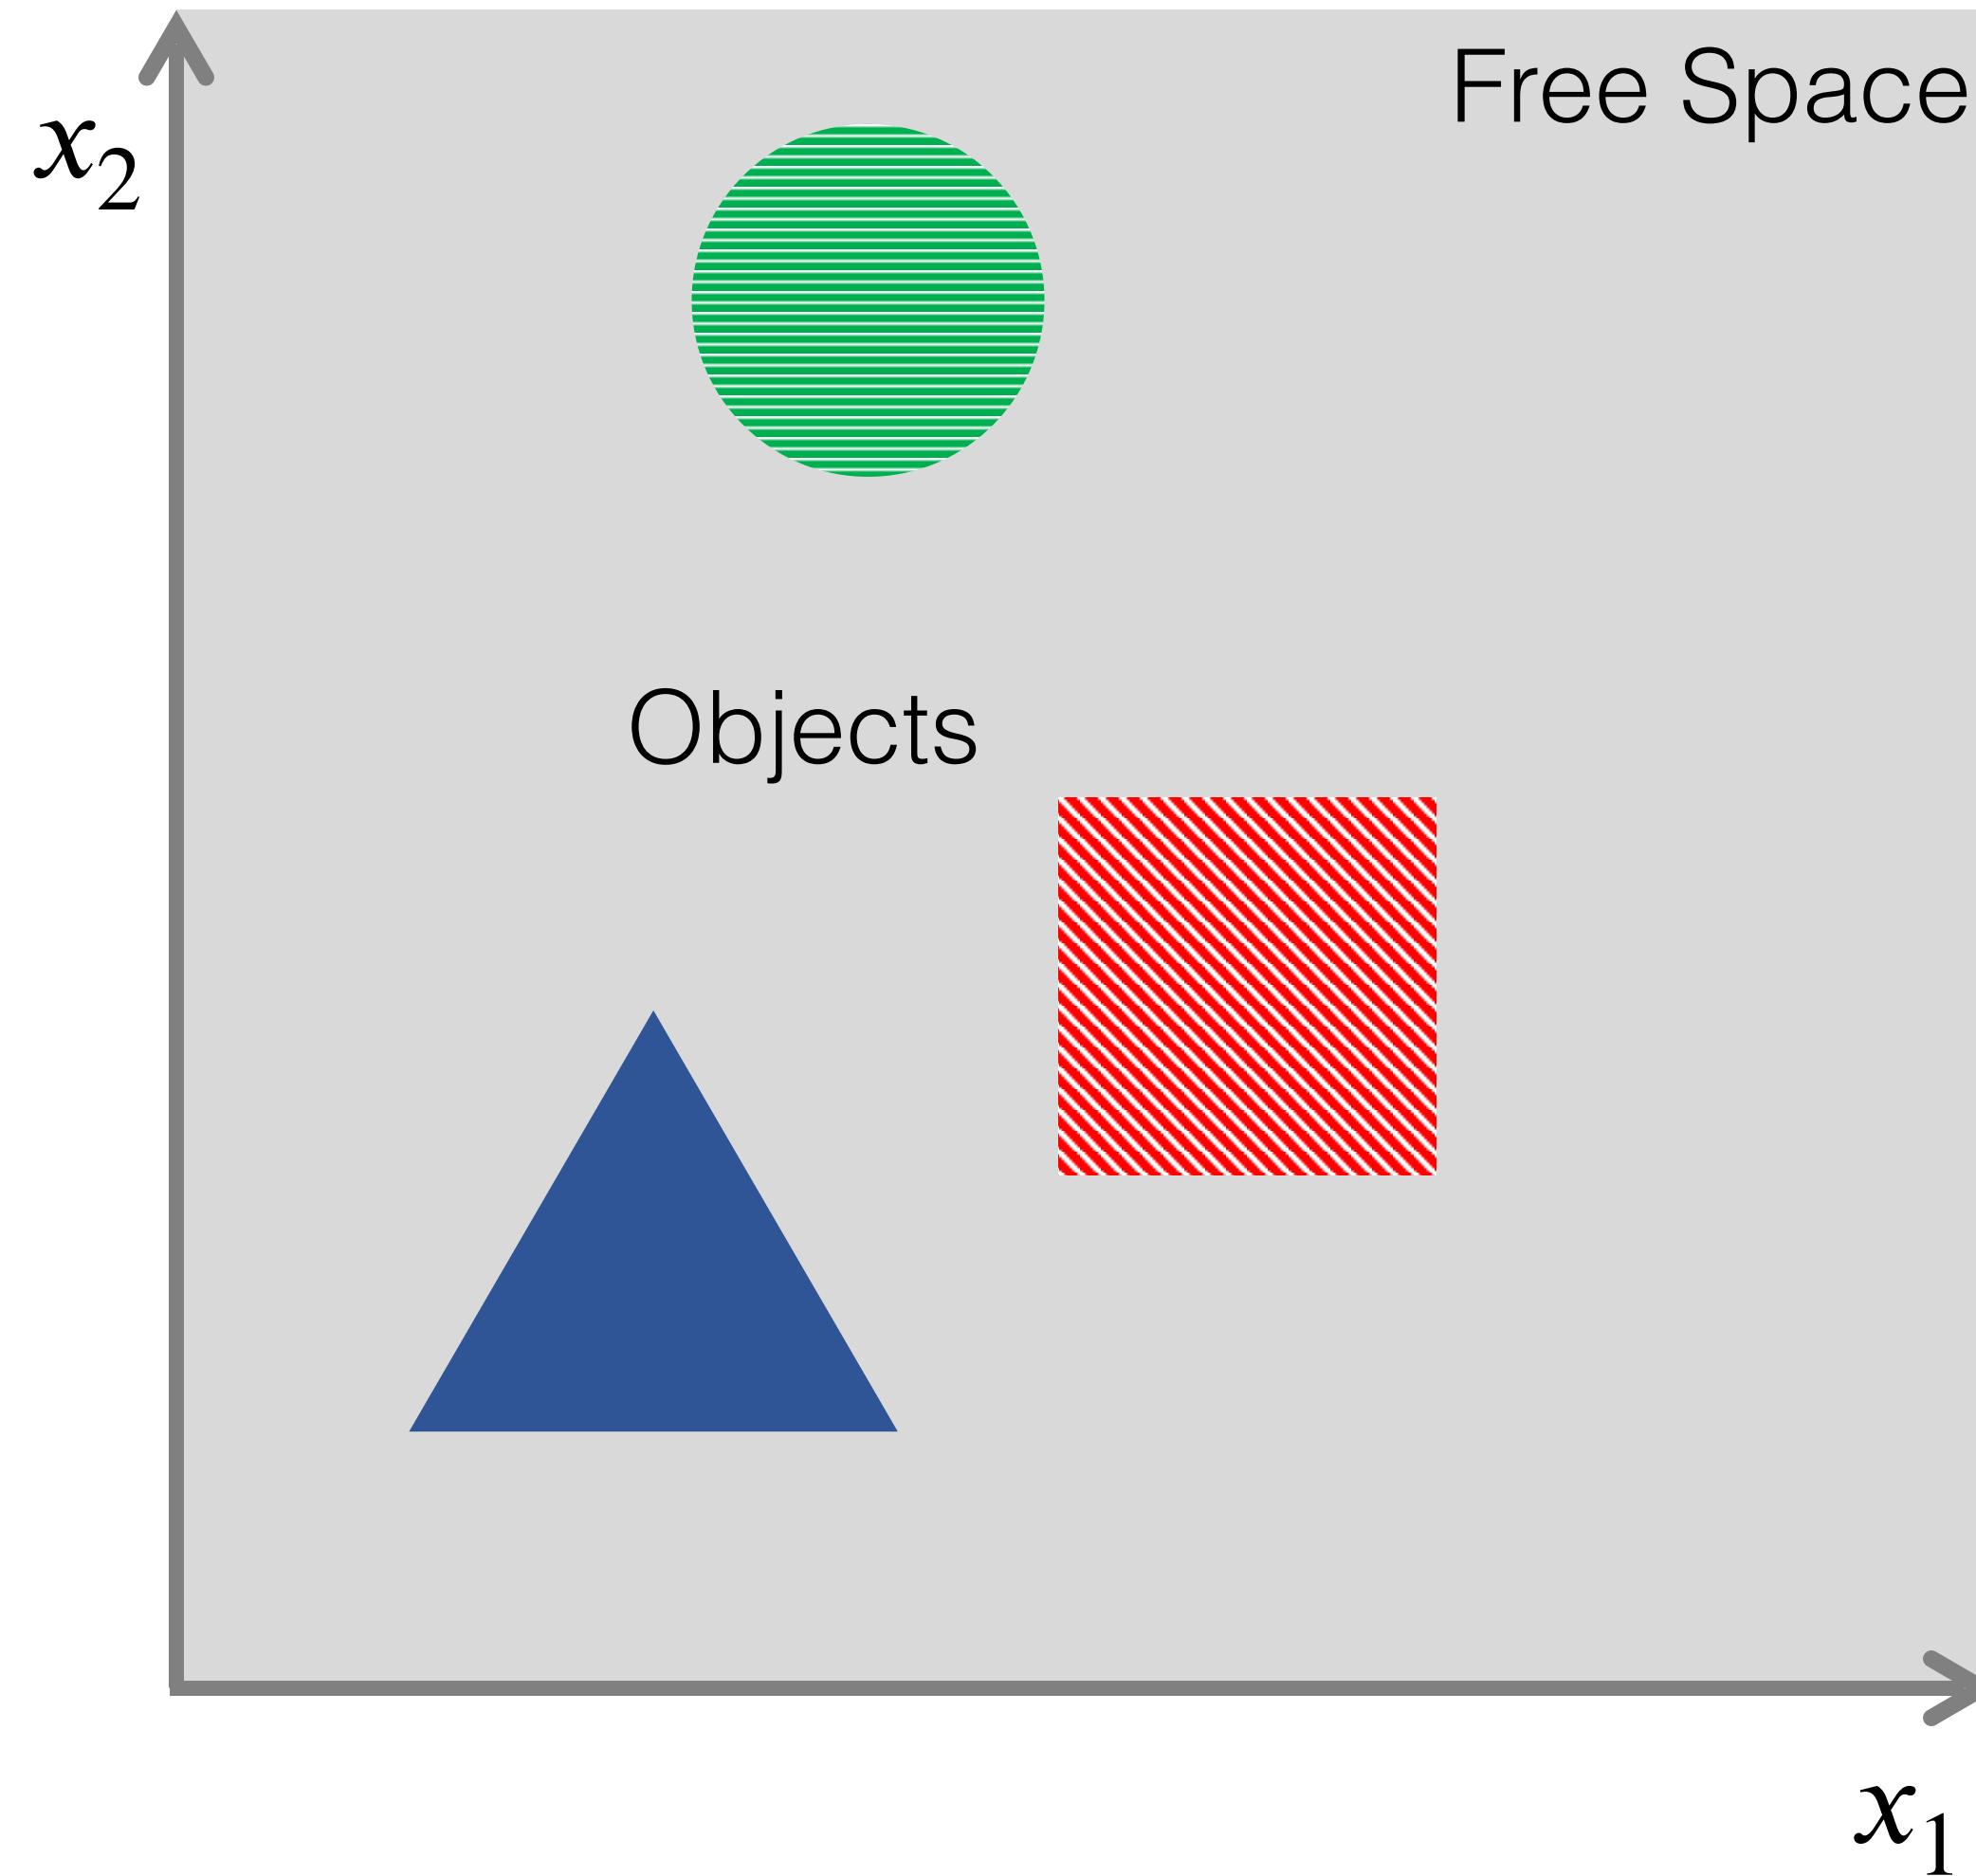
Surface Representations



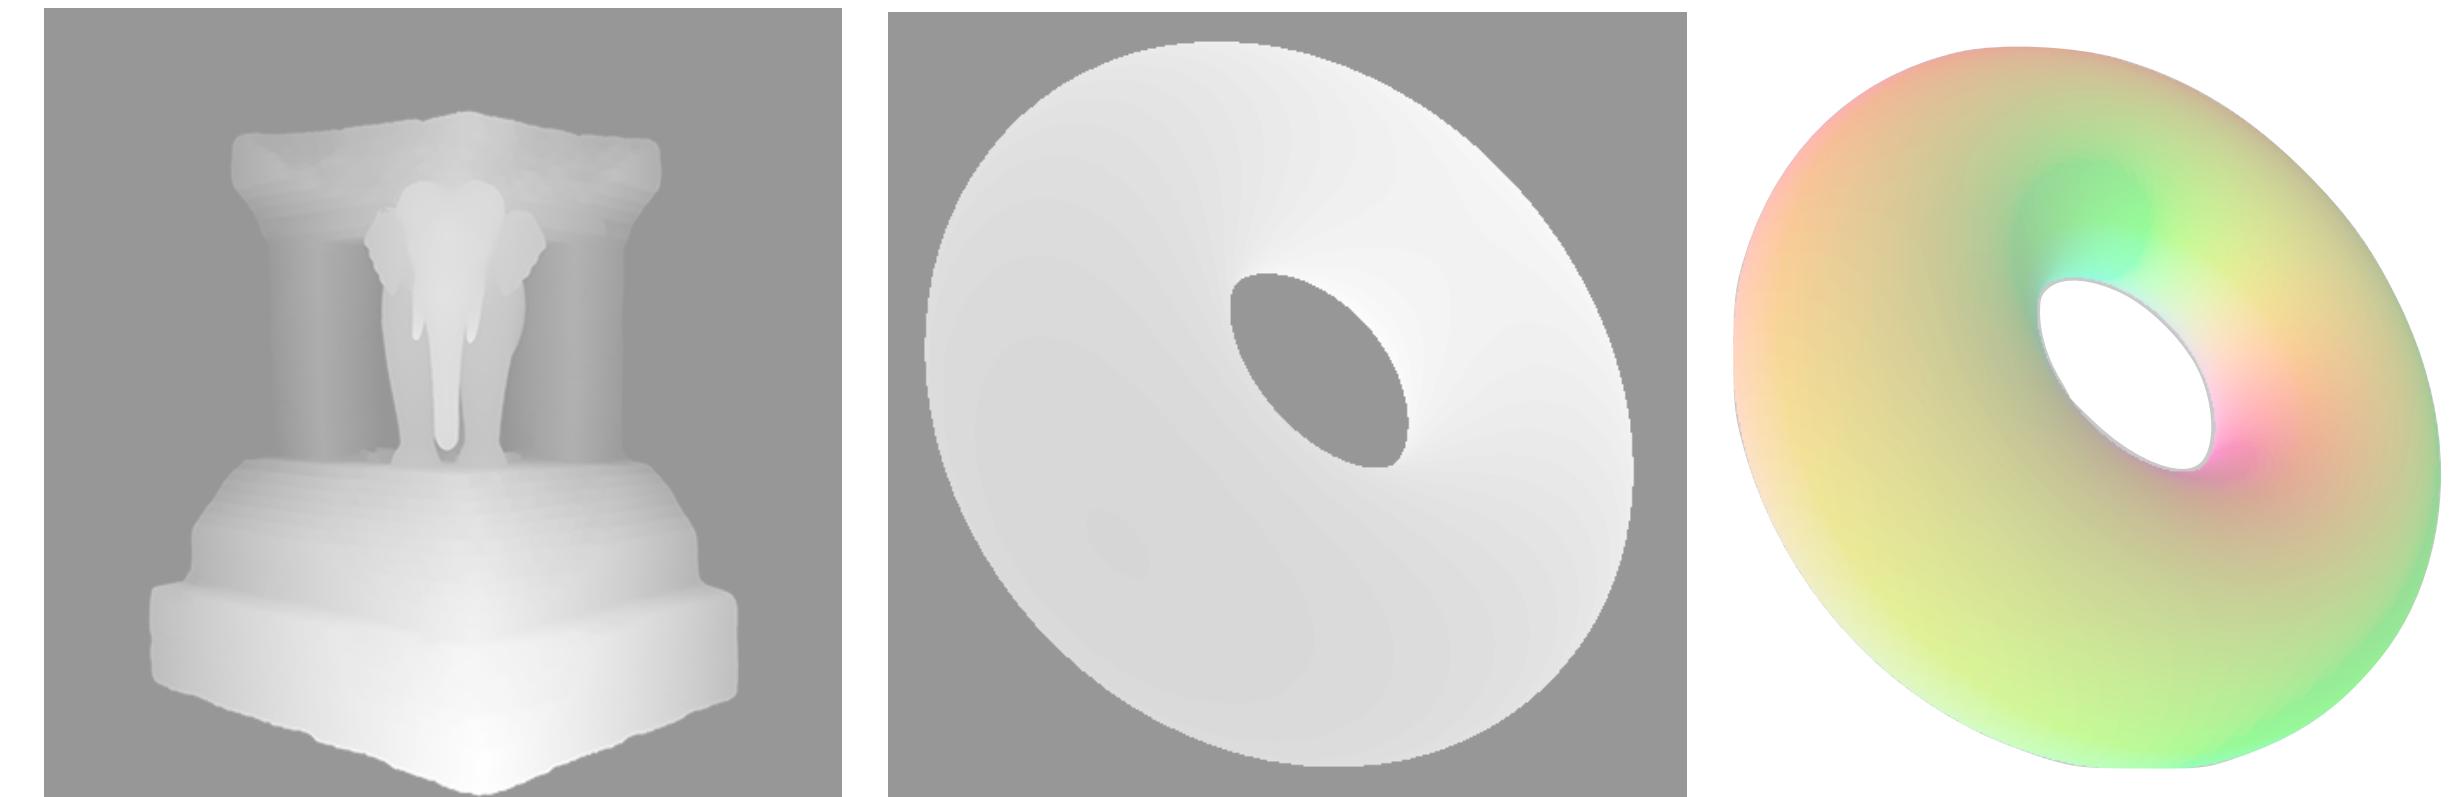
Surface representations only parameterize stuff on surfaces of things.



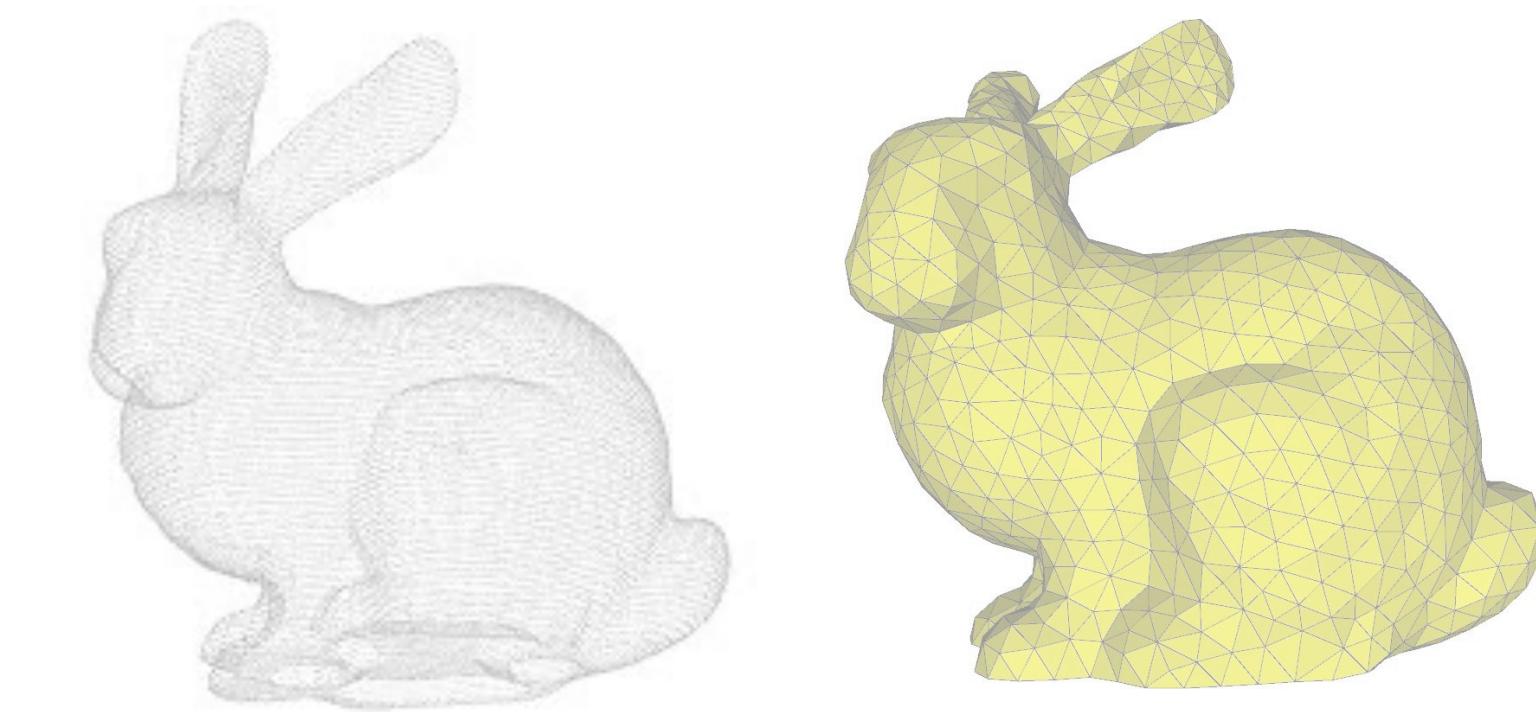
What if we wanted to parameterize stuff **everywhere**?



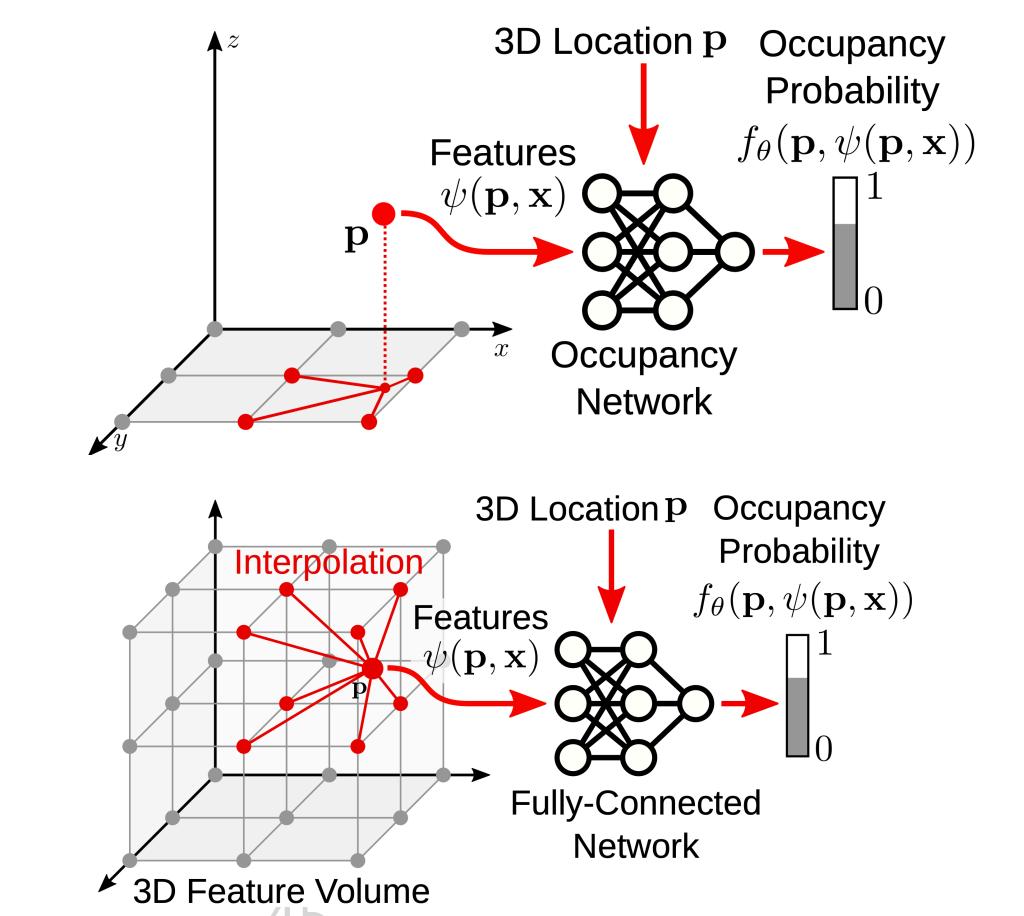
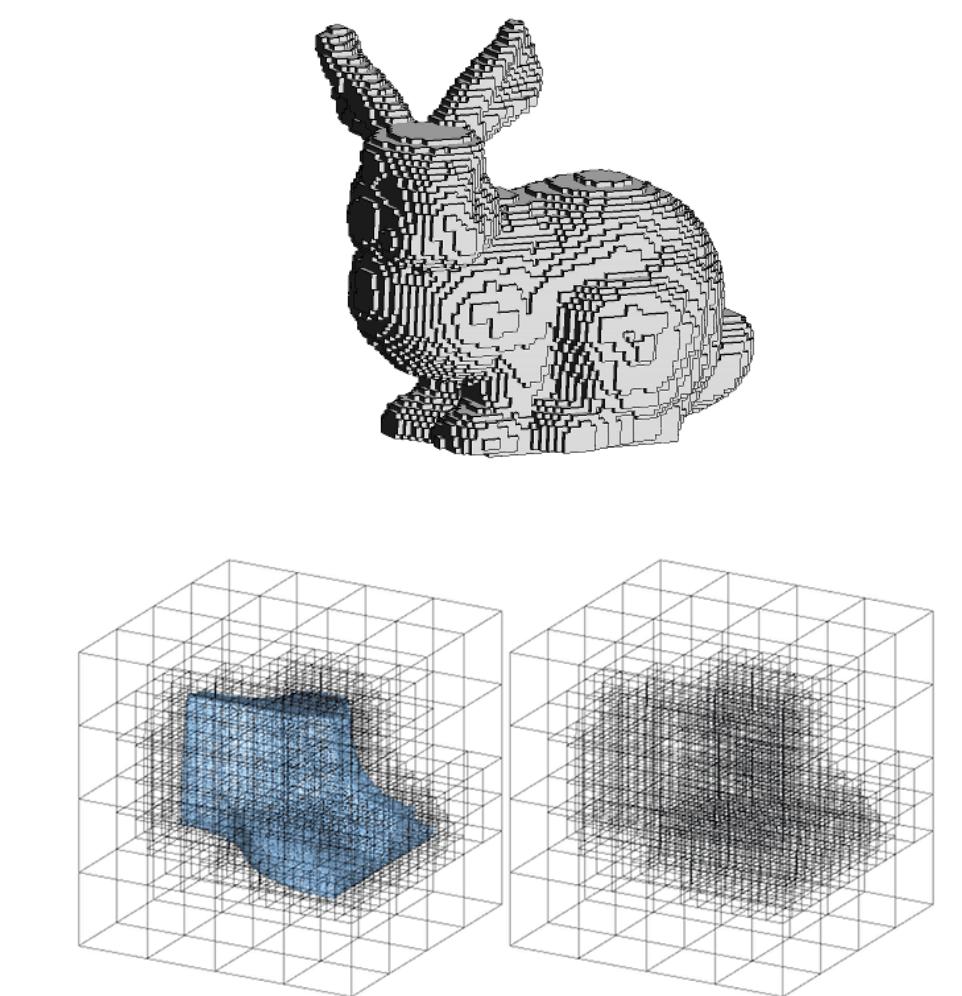
2.5D Representations



Surface Representations



Field Parameterizations



Definition: Field (physics)

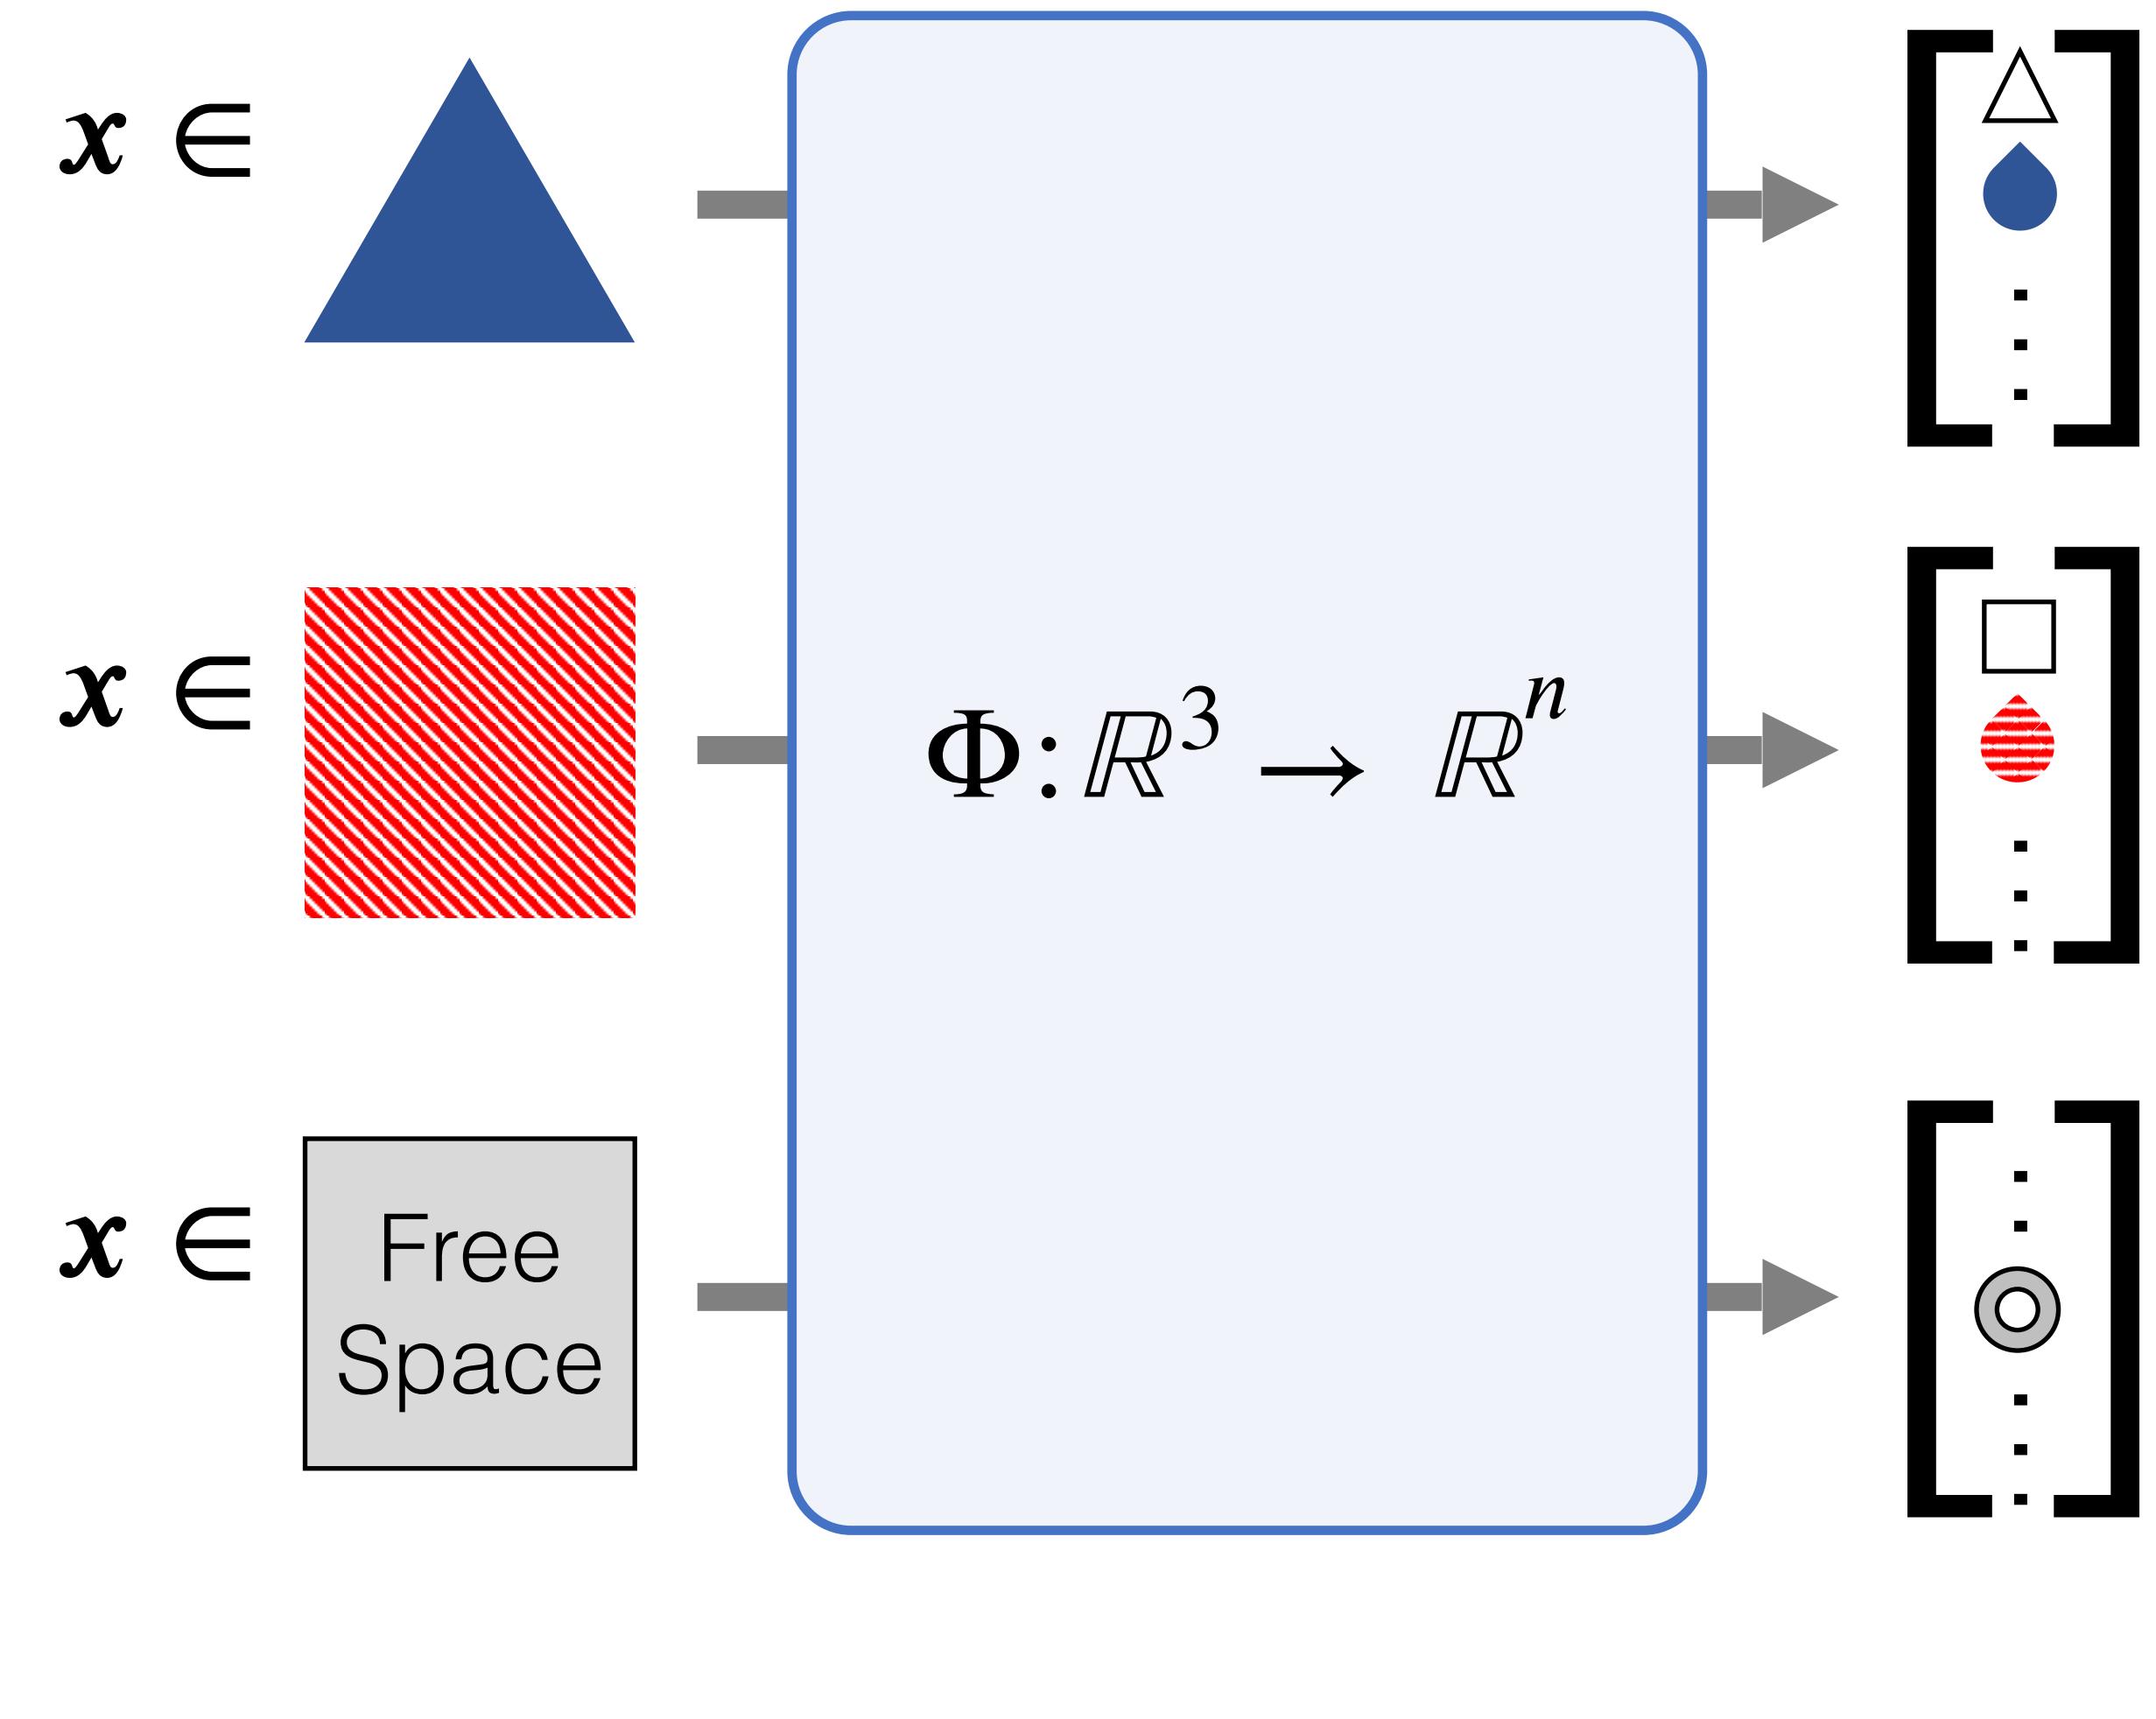
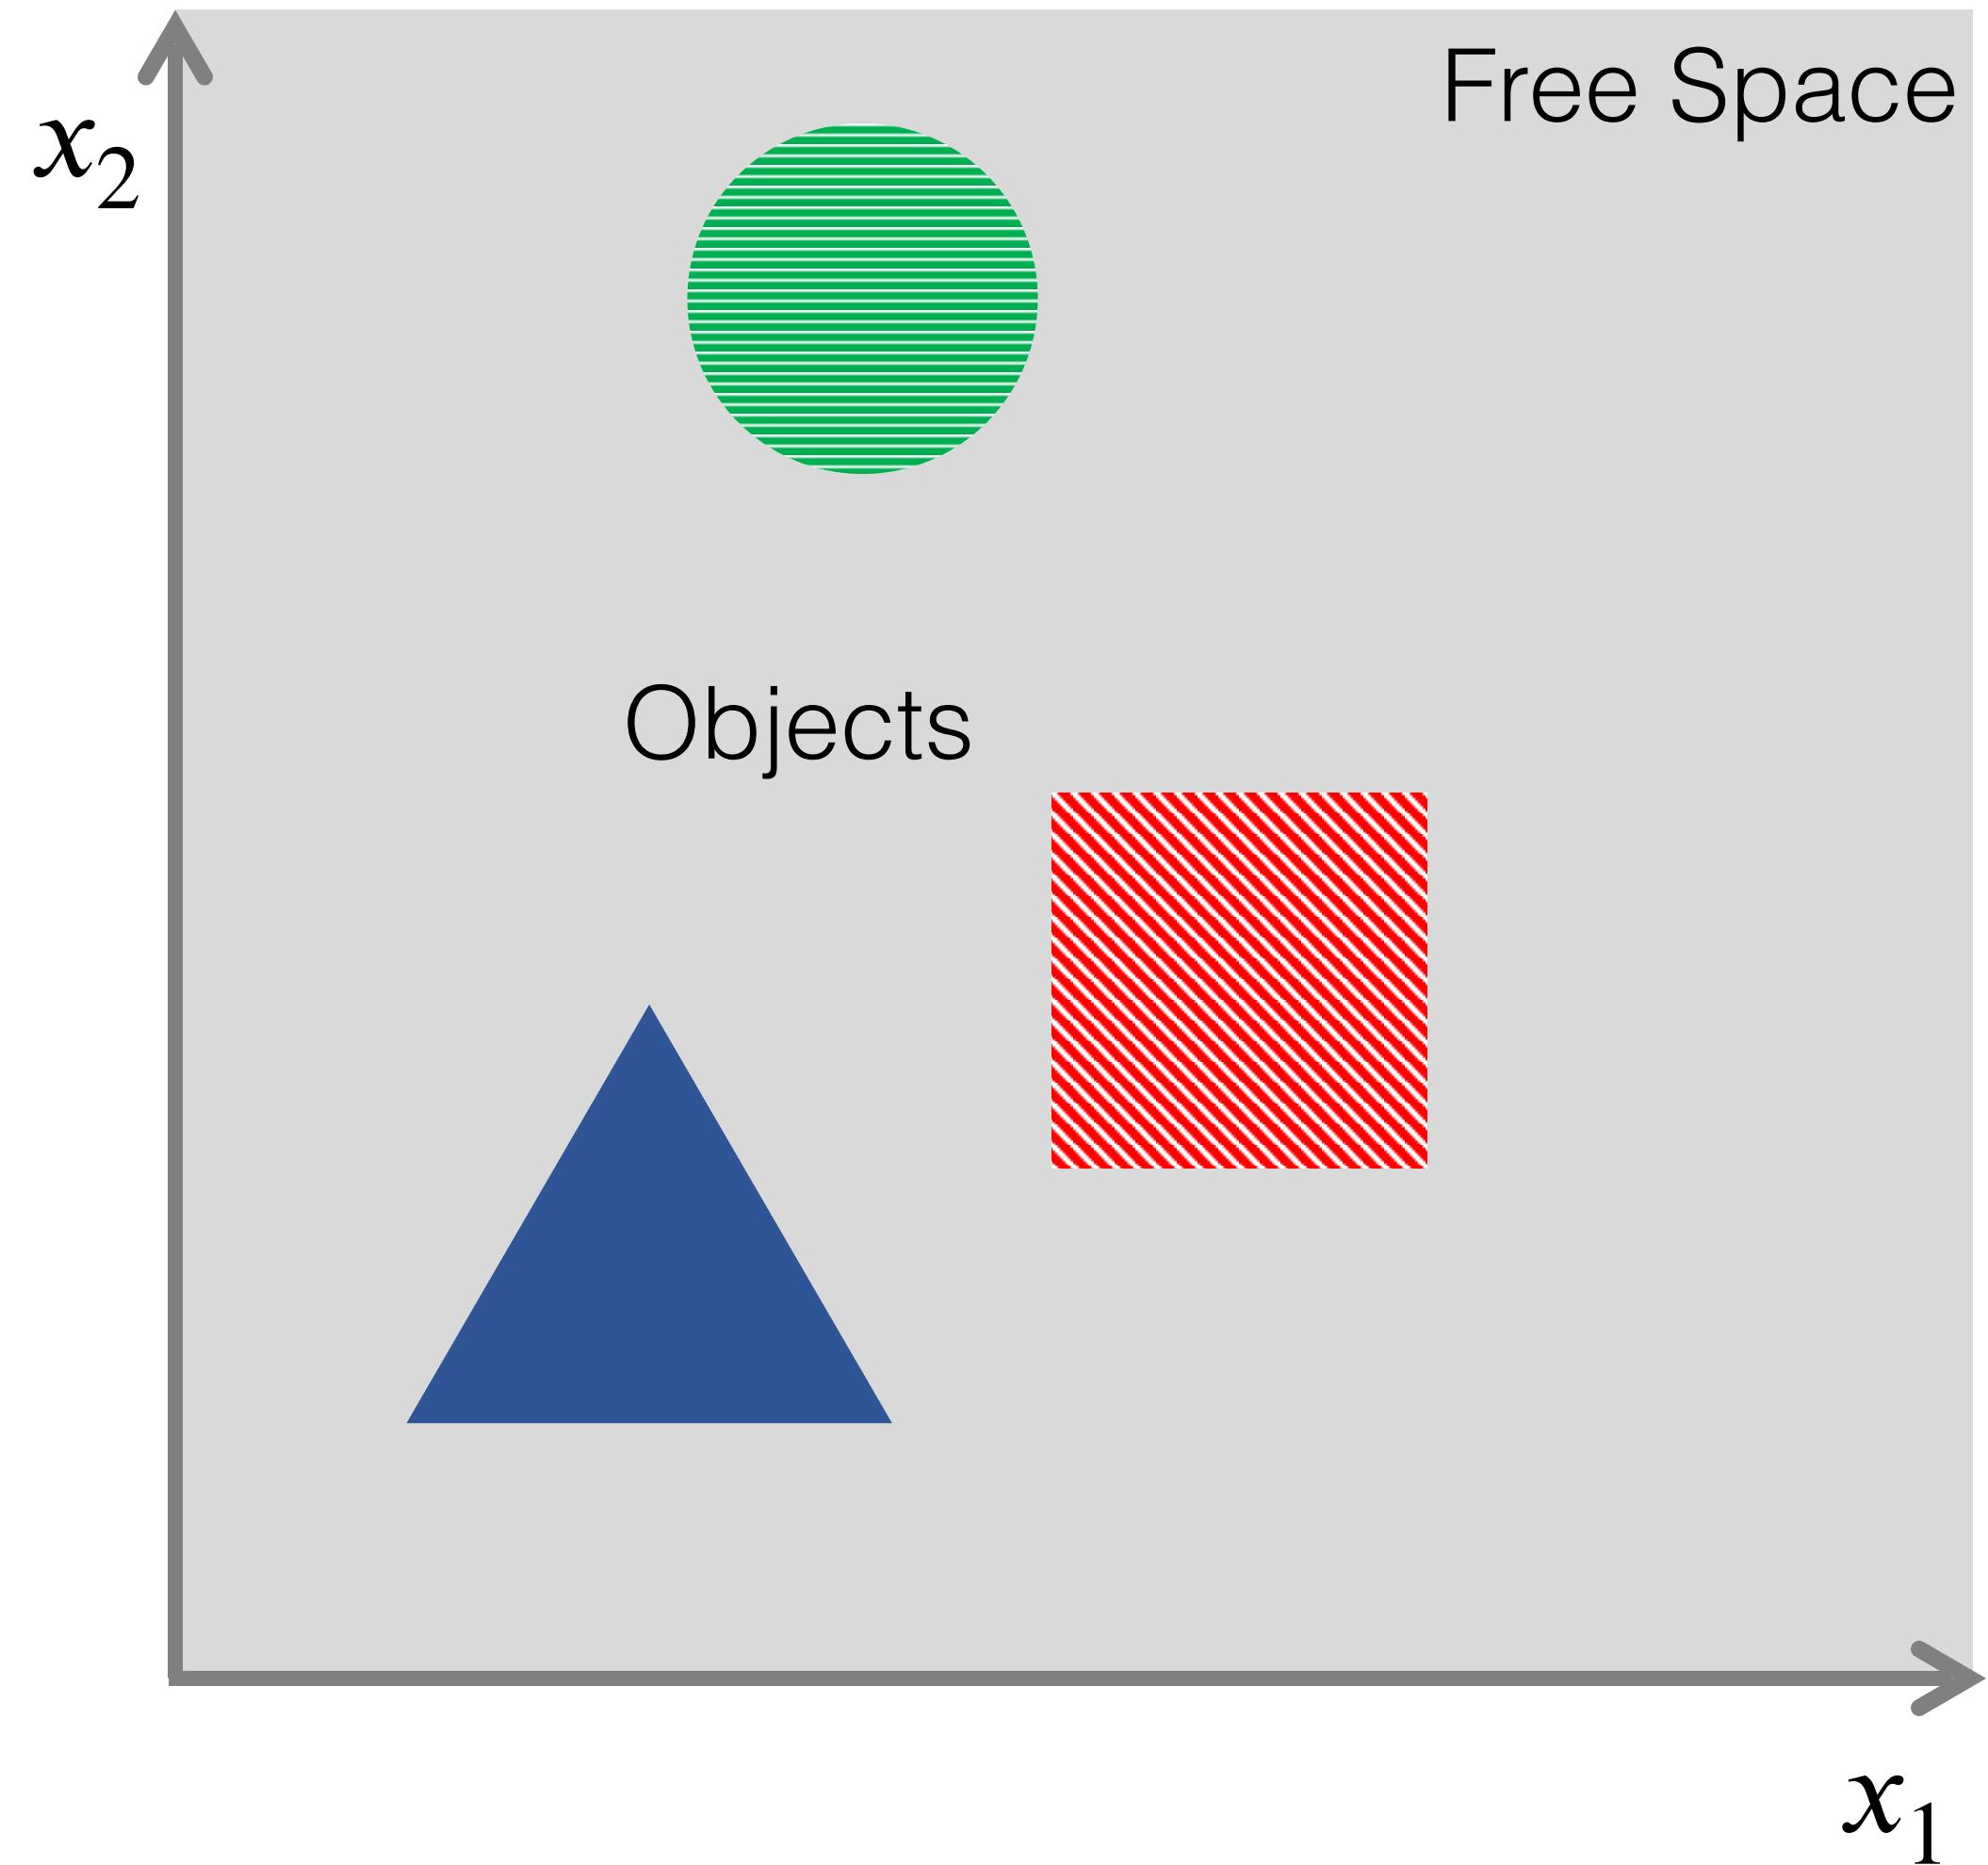
“

A physical quantity represented by a scalar, vector, or tensor, that has a value for each point in space and time.

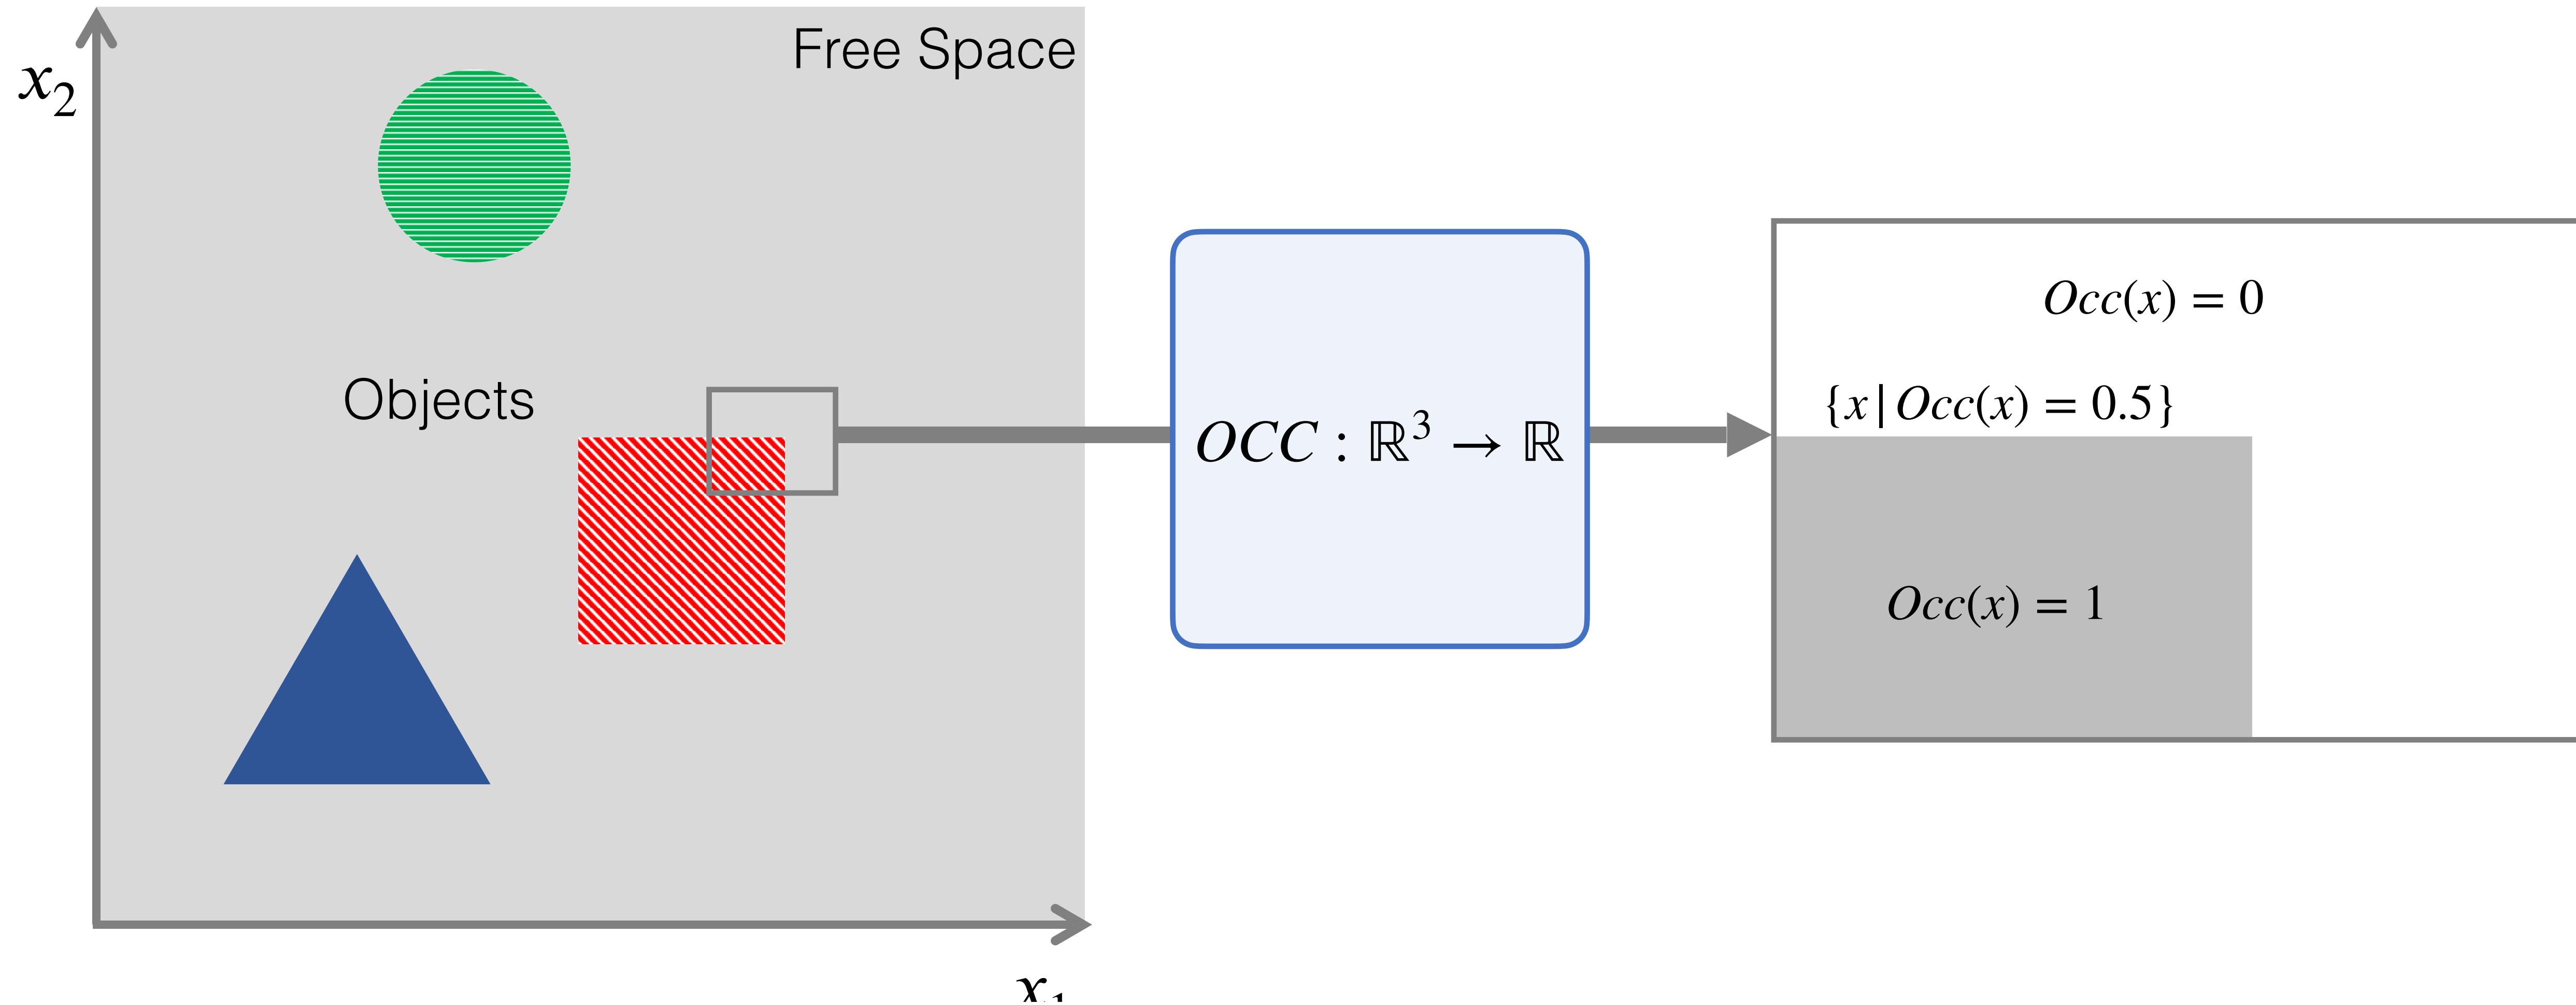
Will use term “field” for any function that takes a space, time, or space-time coordinate as input to map it to some quantity at that coordinate.

$$\Phi: \mathbb{R}^3 \rightarrow \mathbb{R}^n$$

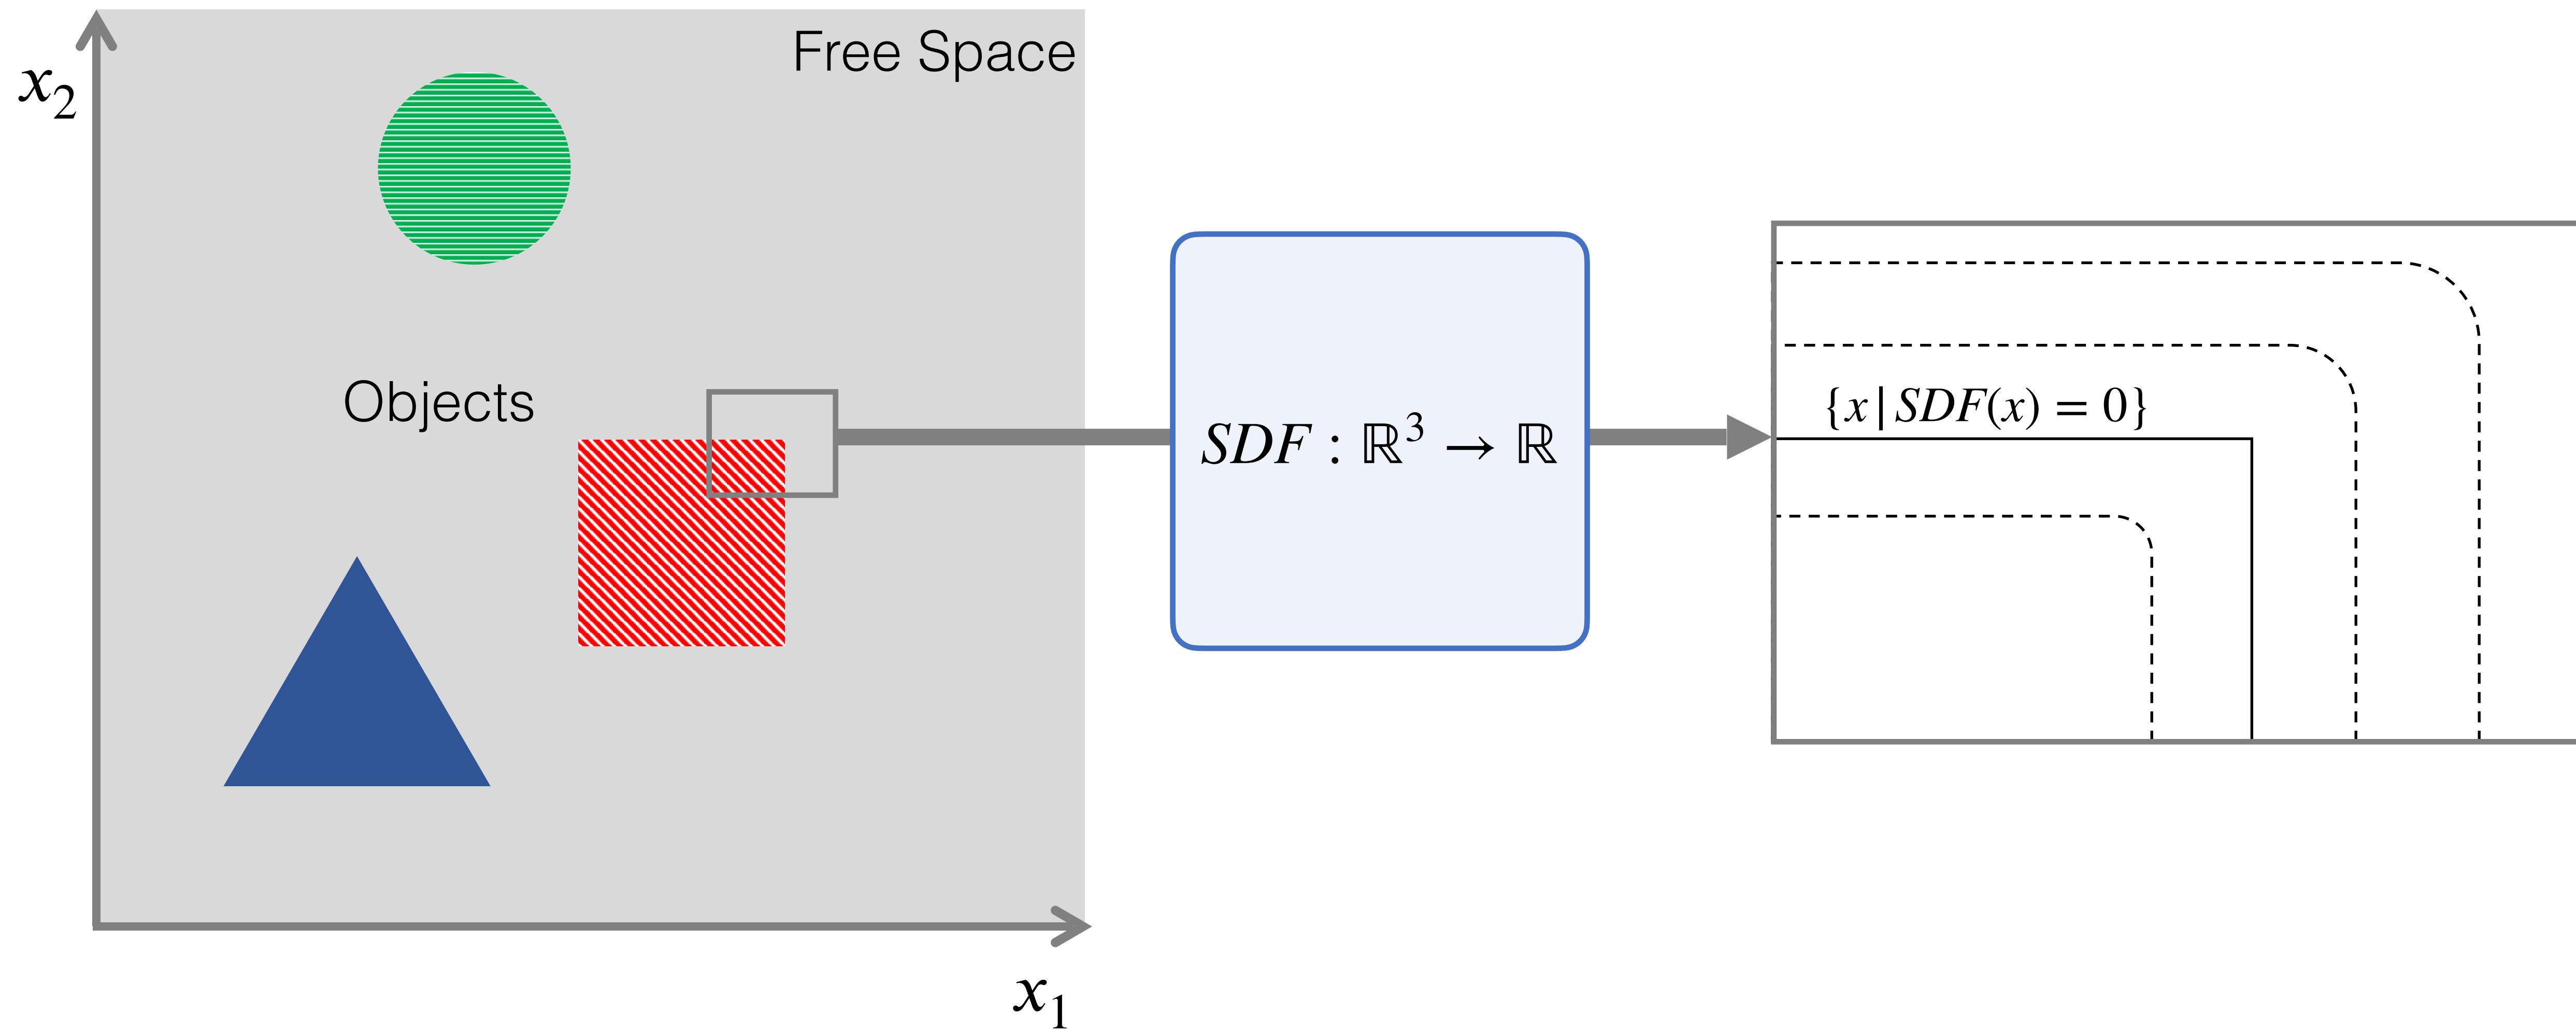
How to parameterize surfaces?



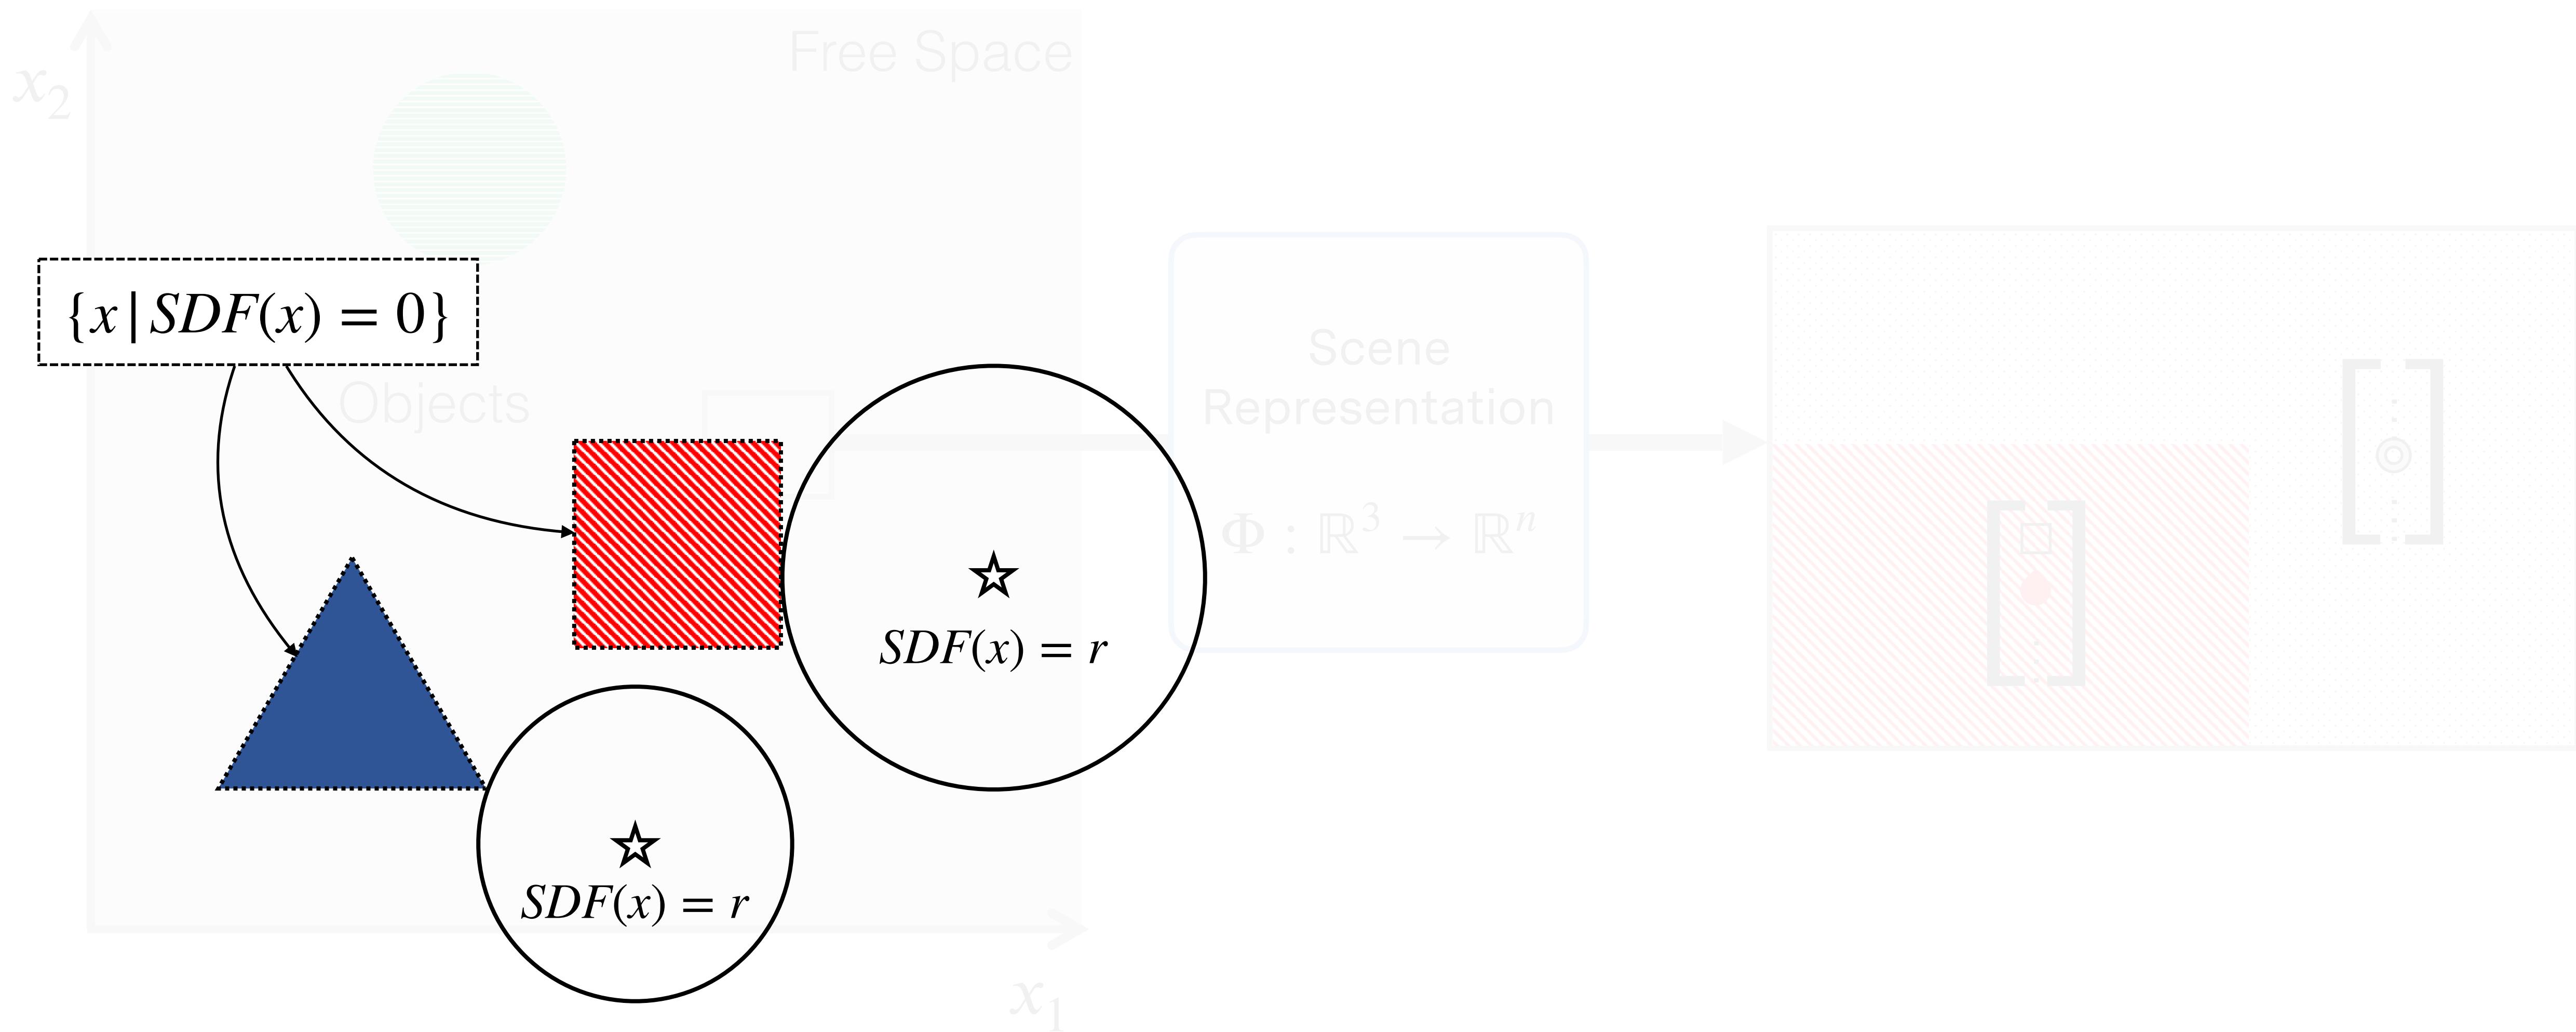
The Occupancy Field



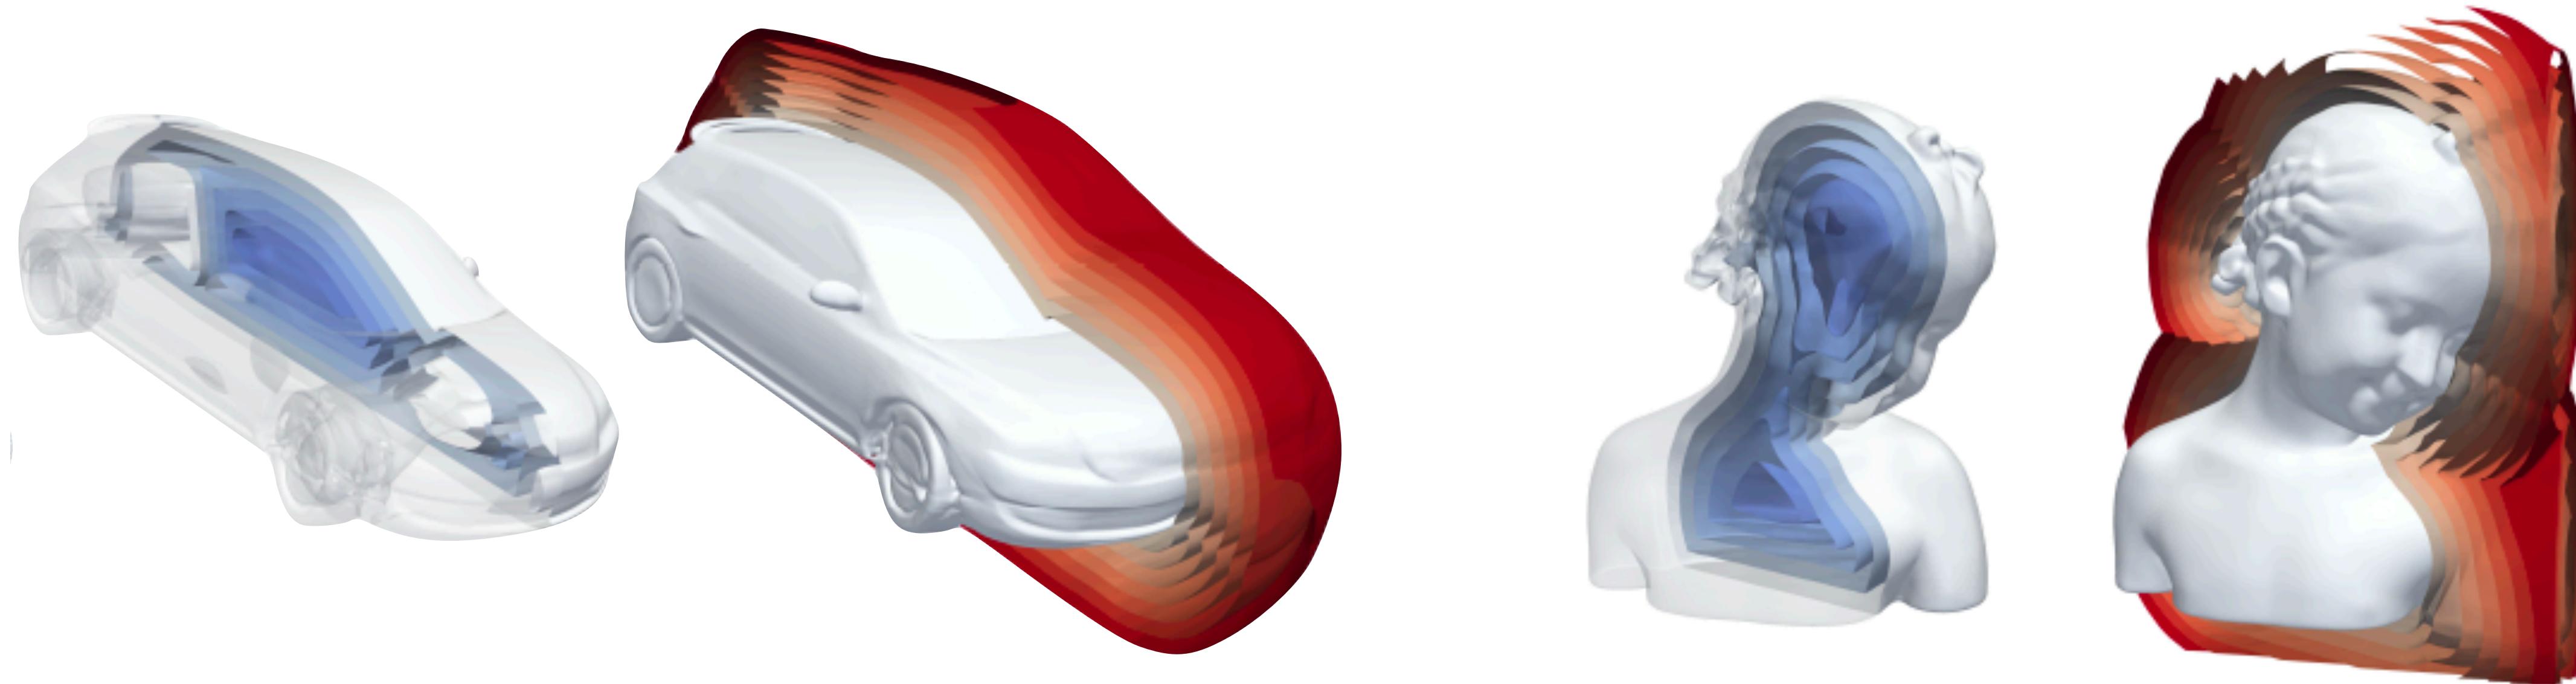
The Signed Distance Field



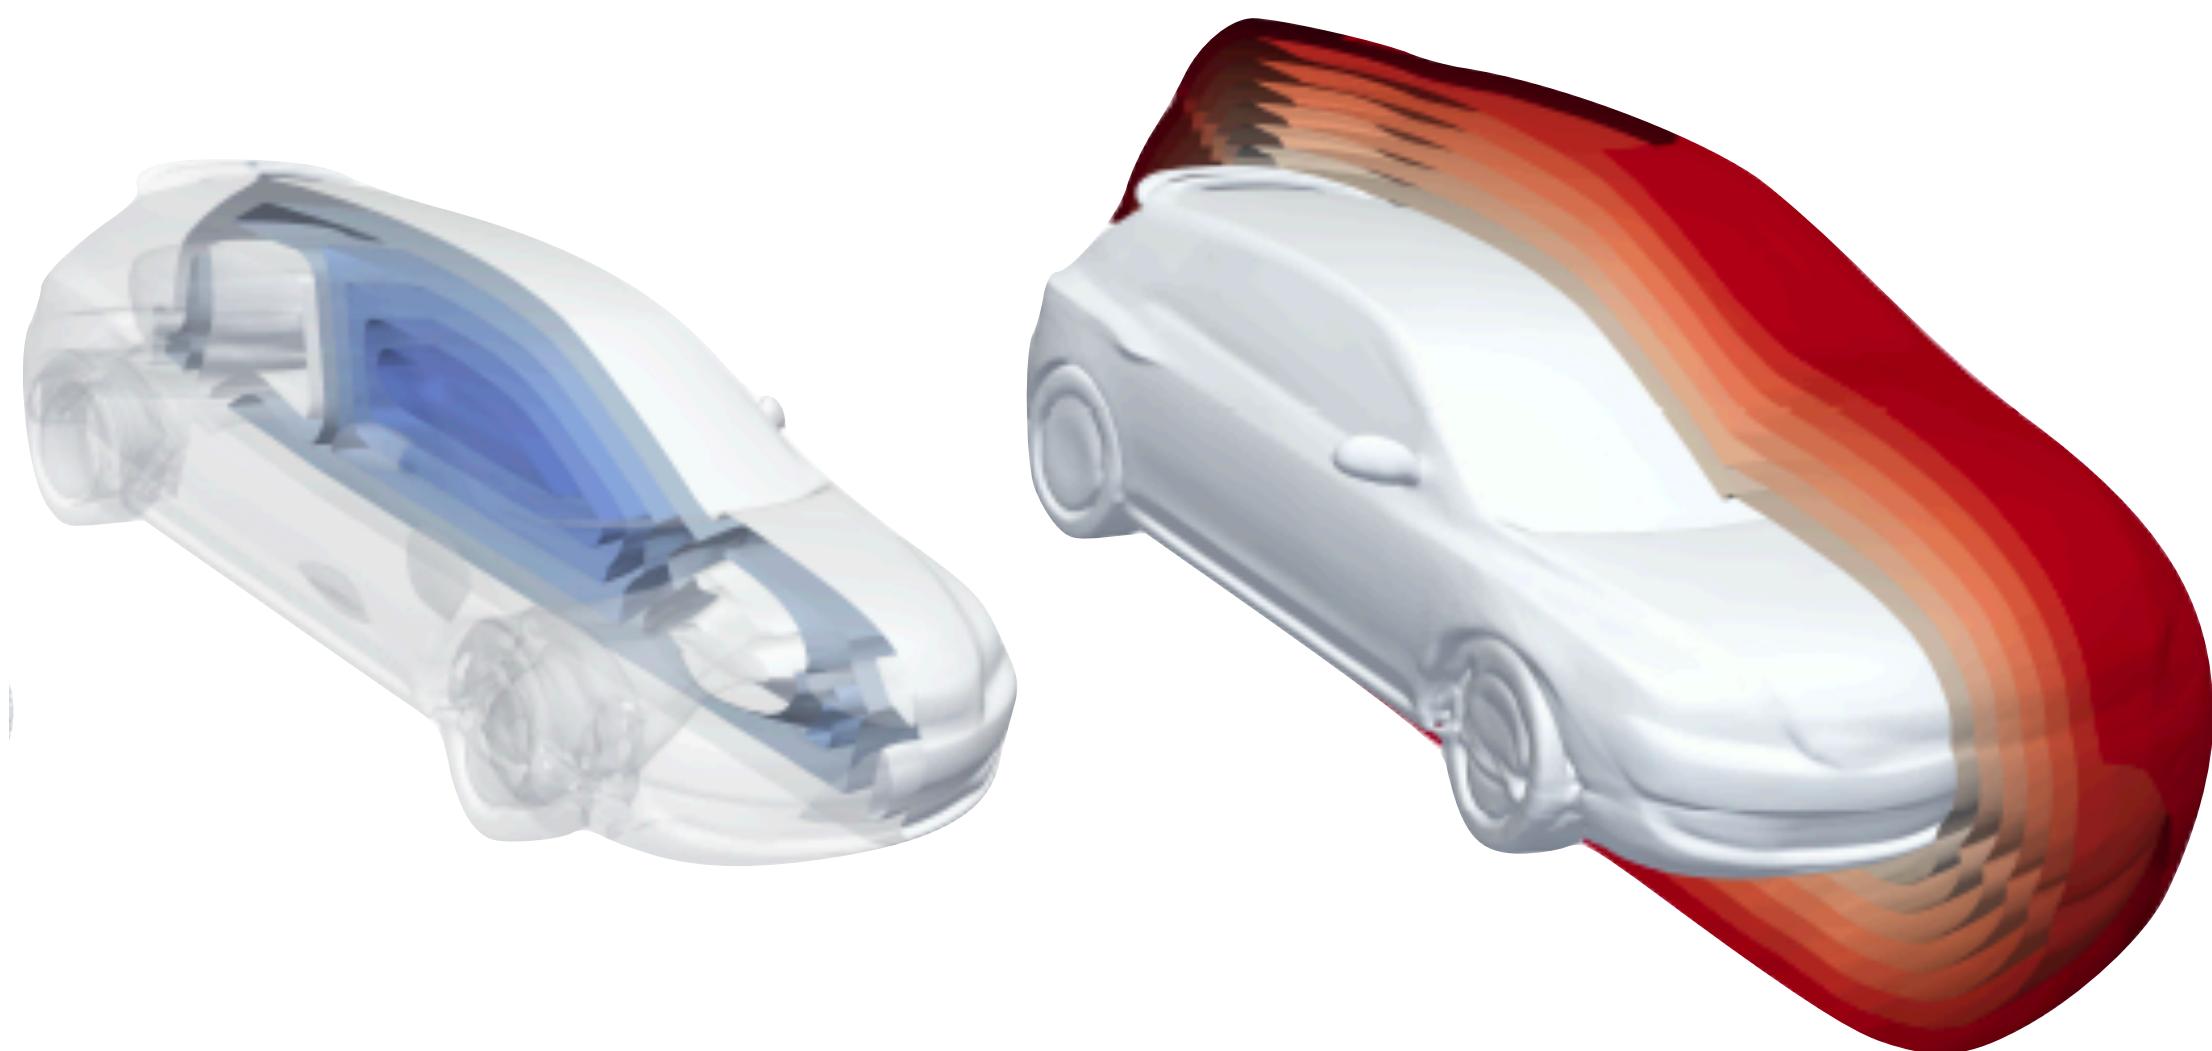
The Signed Distance Field



Implicit Surfaces



Implicit Surfaces

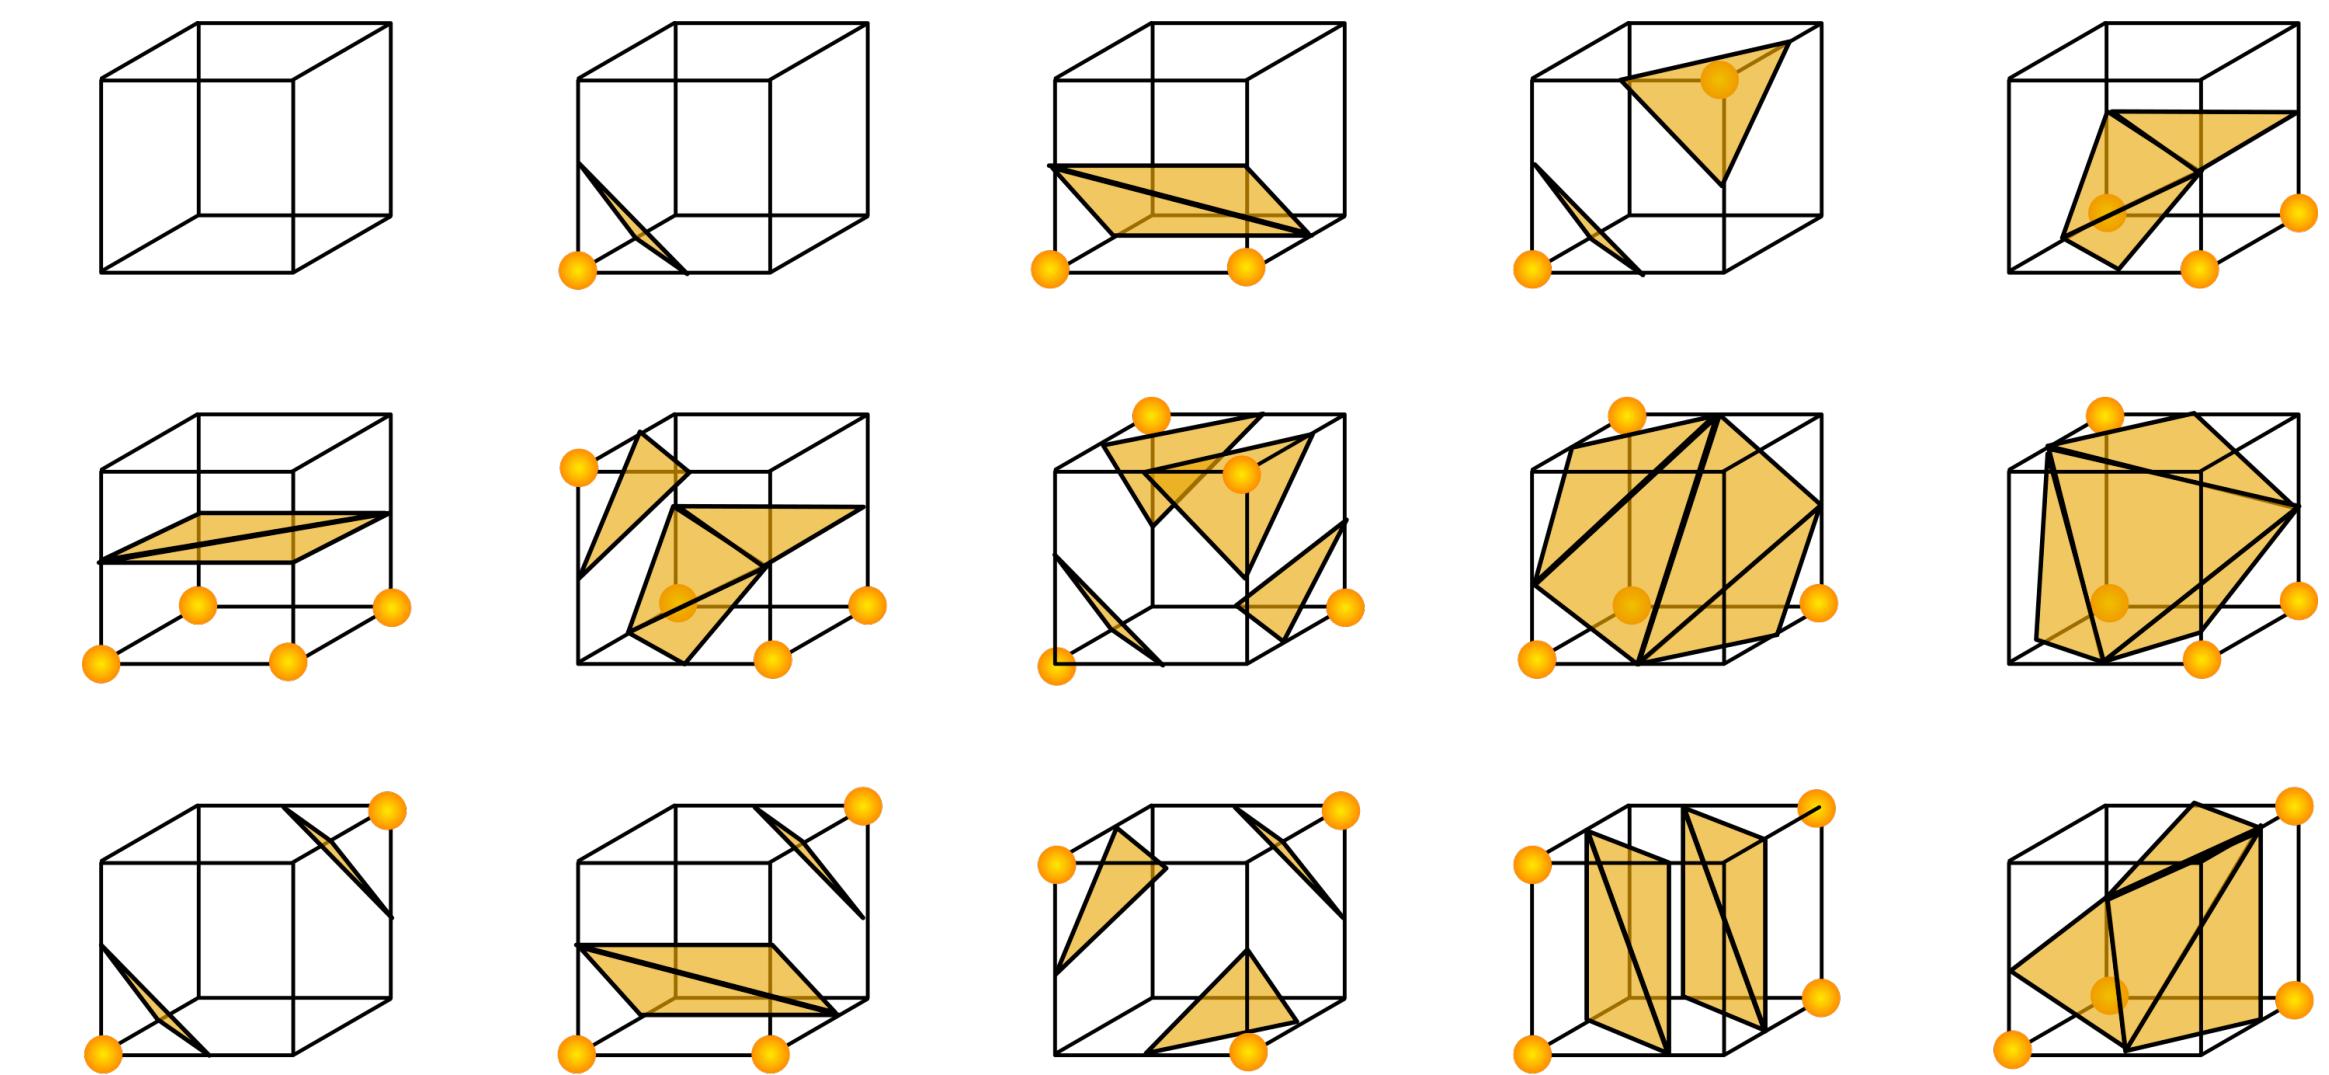
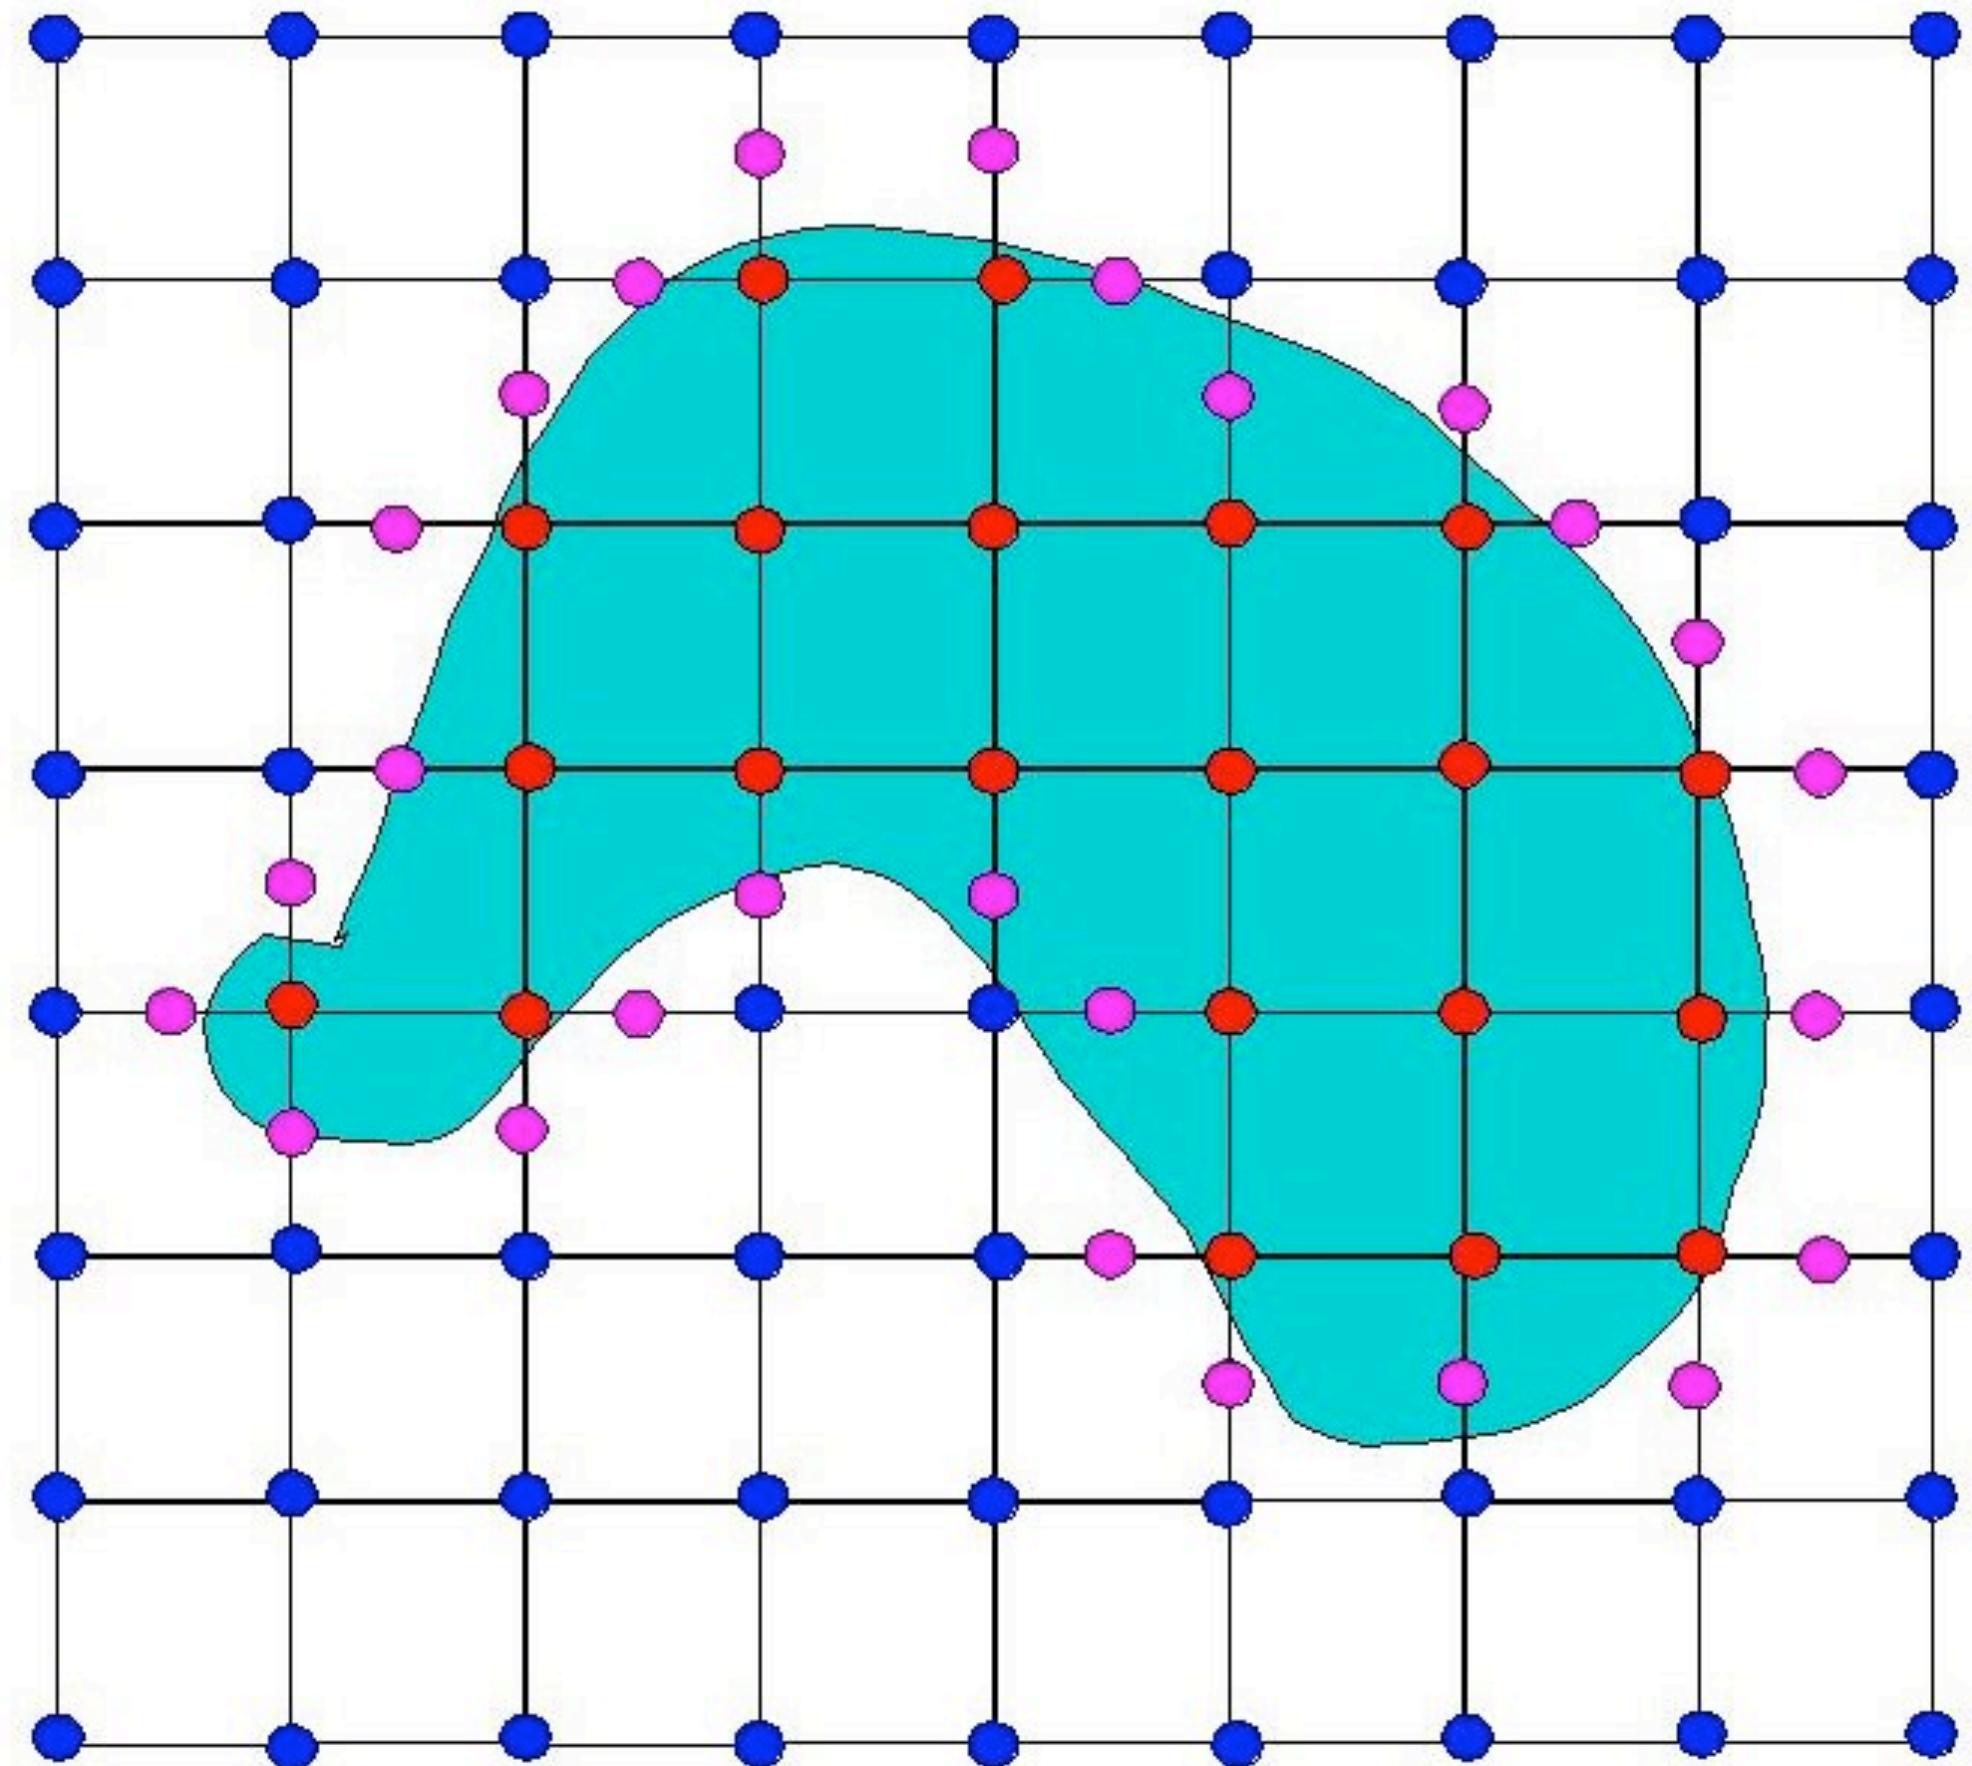


Level sets of Occupancy, SDF
parameterize surface

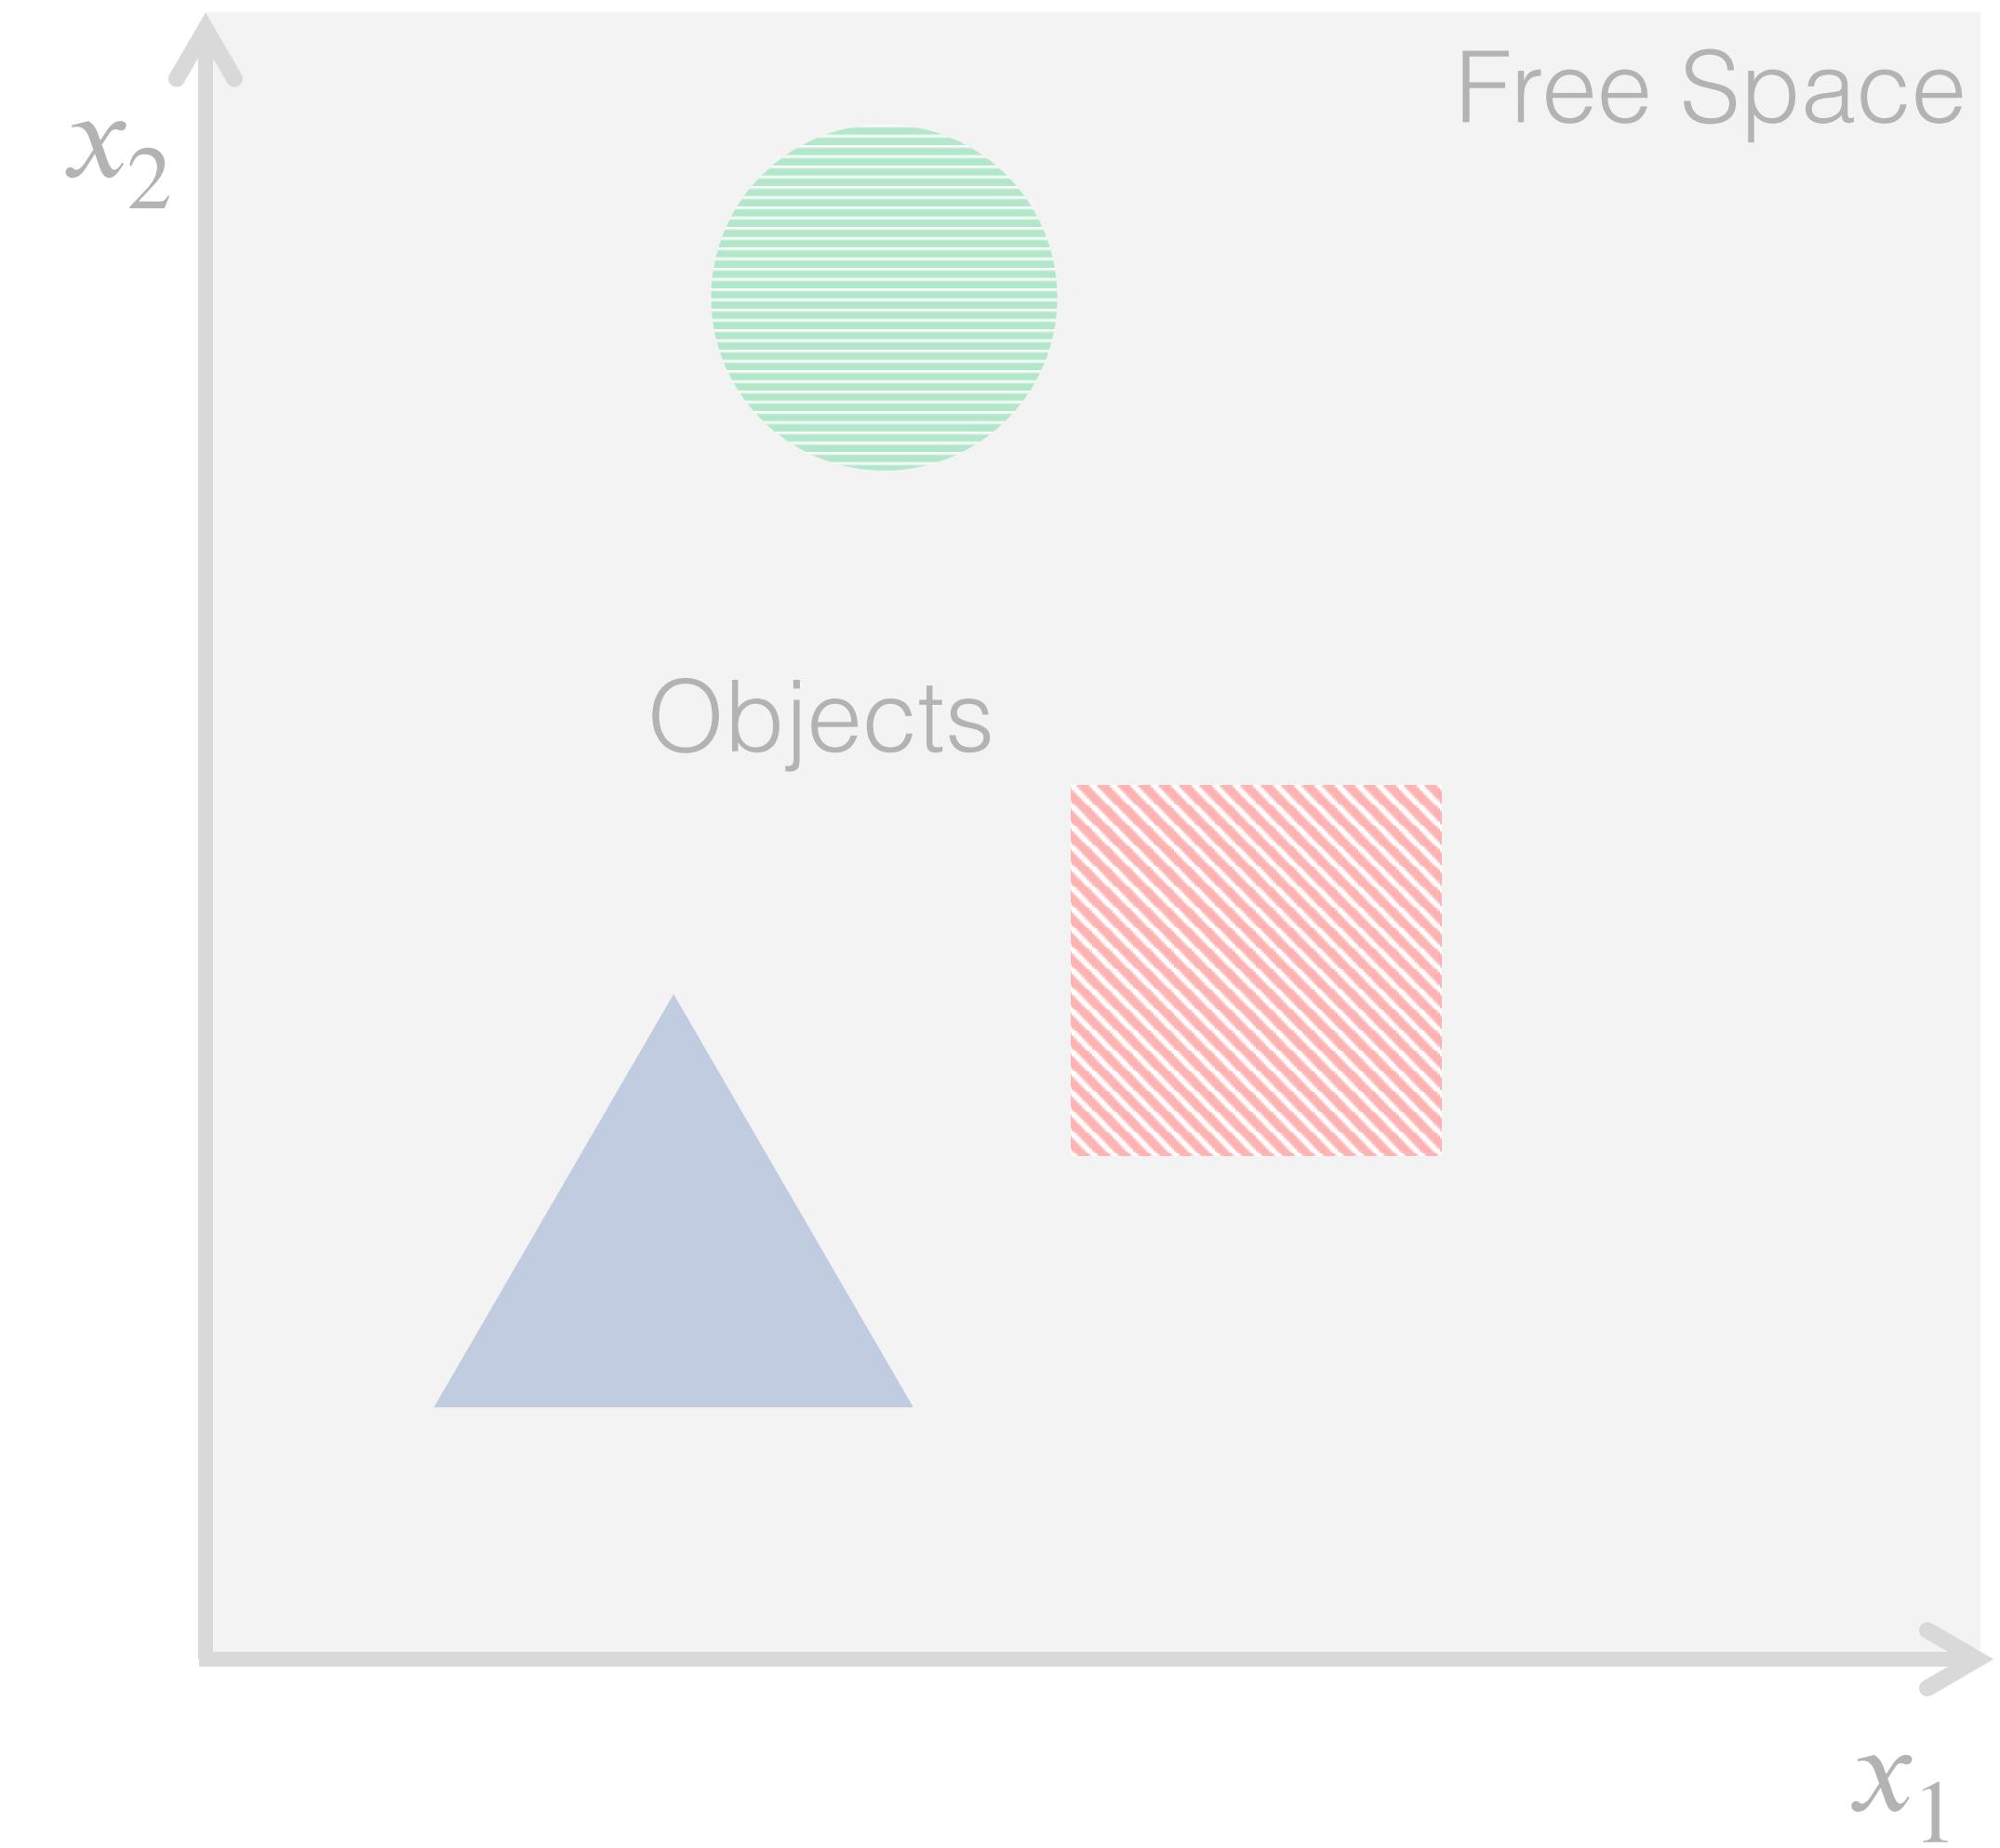
Easy to answer questions about
occupancy

Difficult to reason about the surface:
have to first **find** zero-level-set!

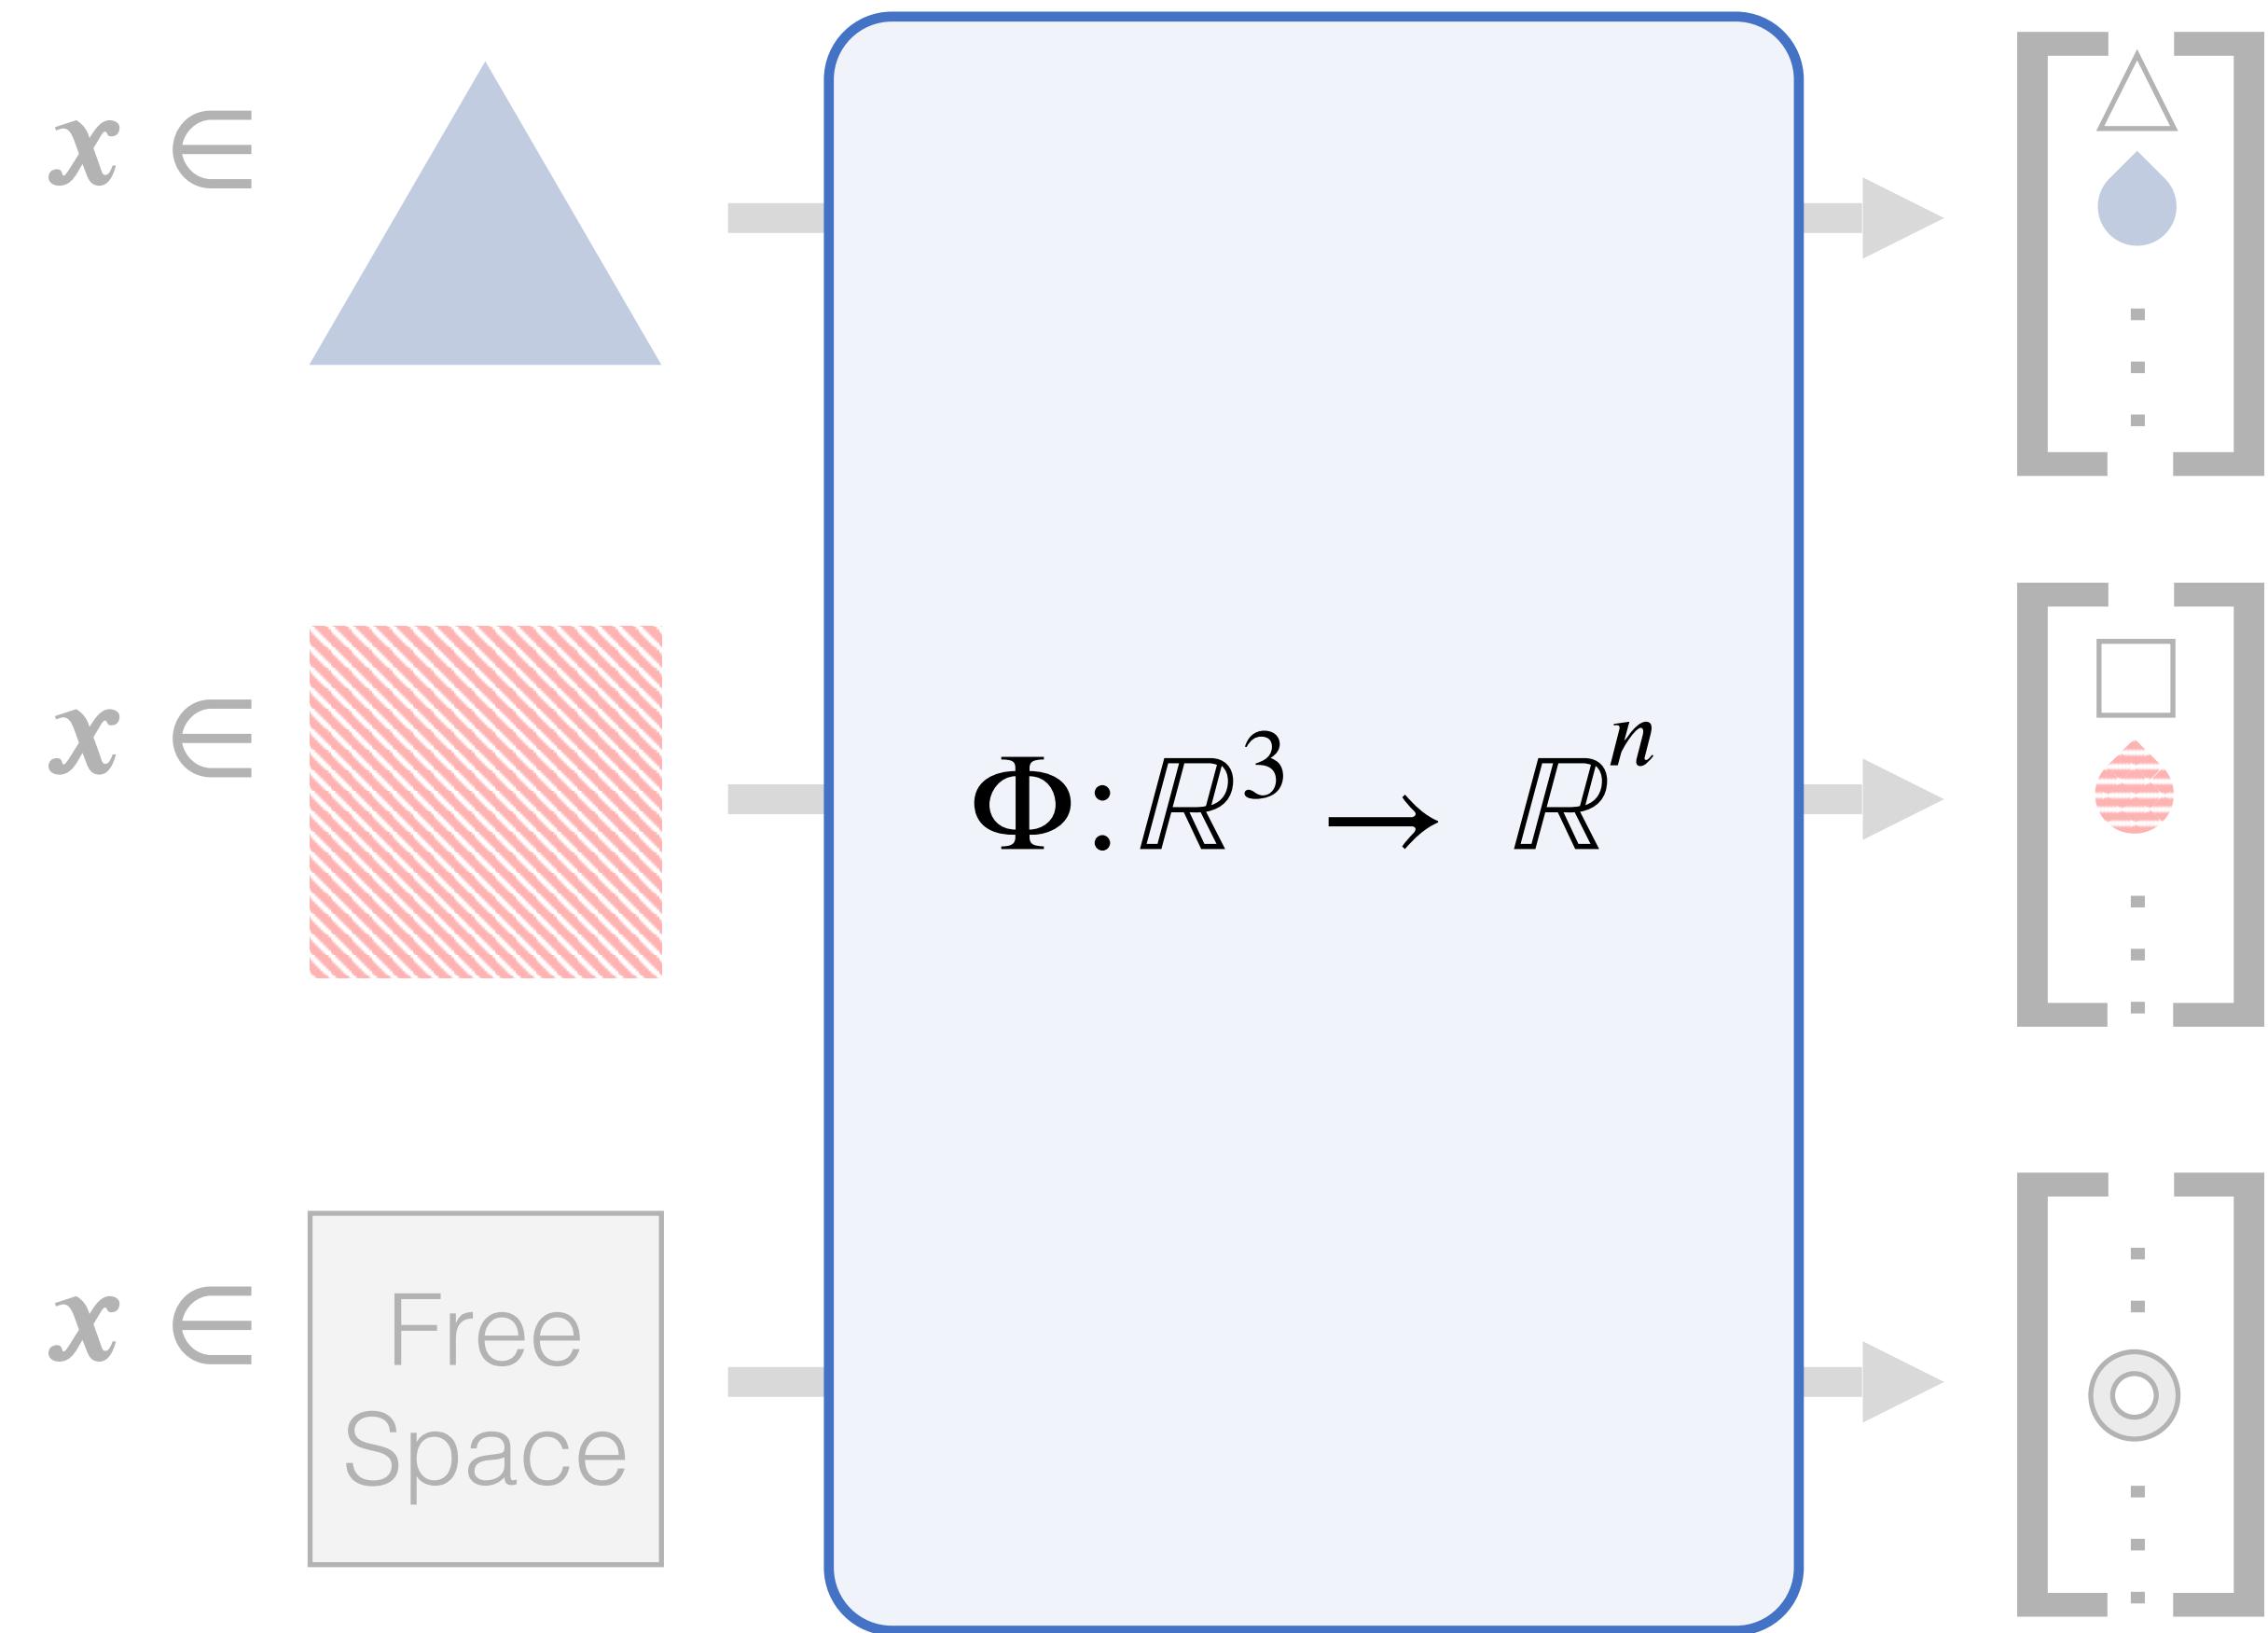
Marching Cubes: Extracting 2D Mesh from 3D Implicit Representation



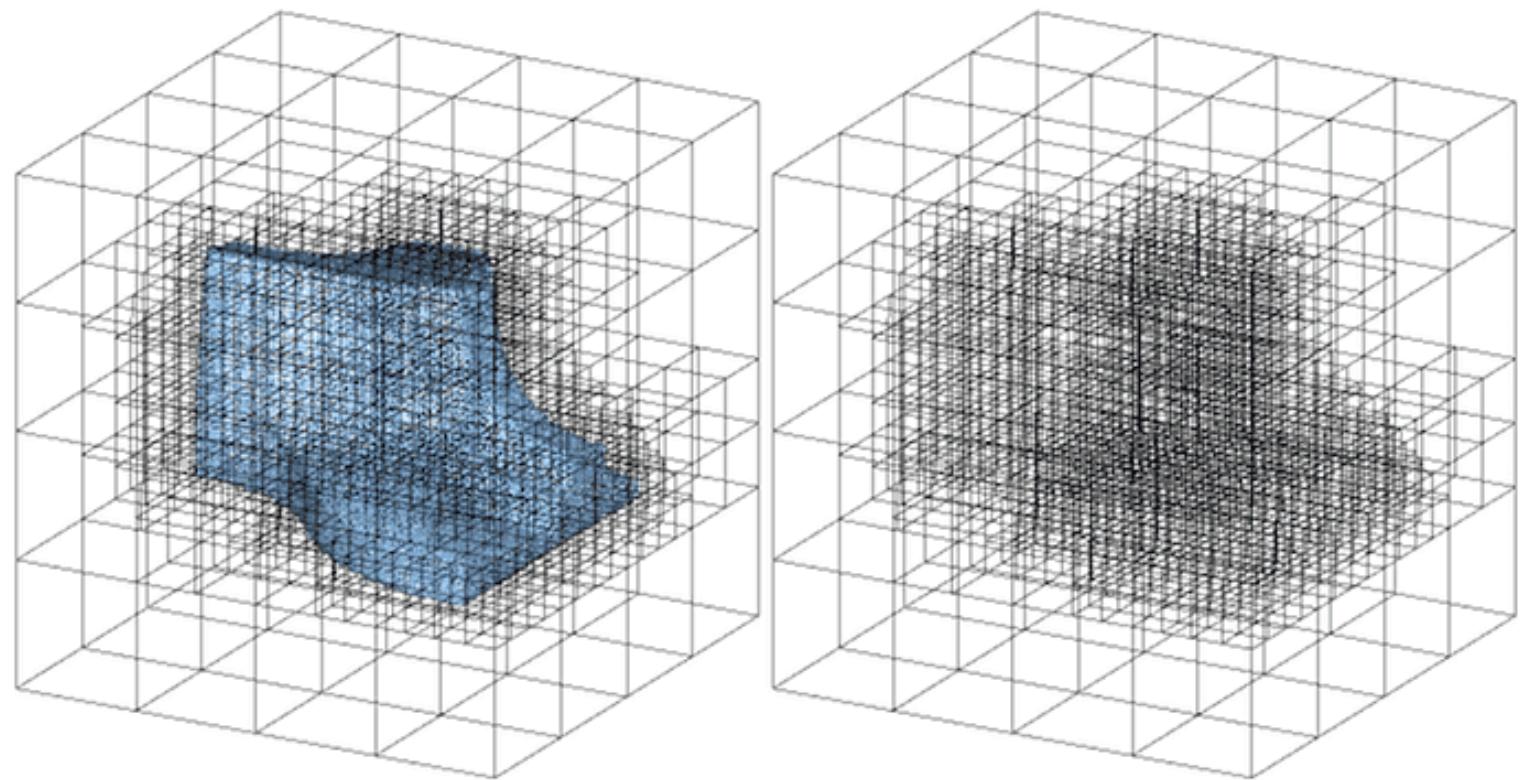
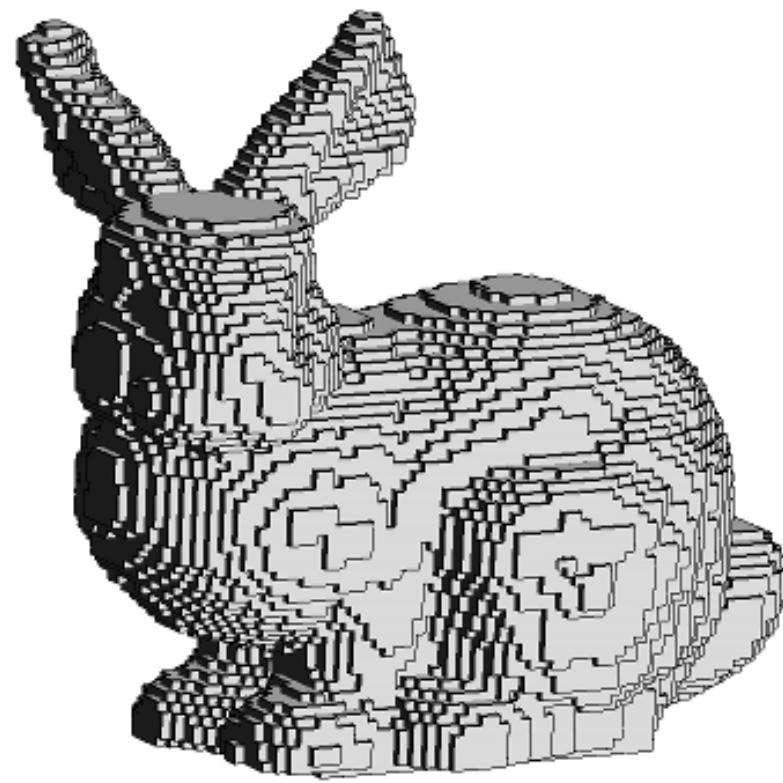
How to parameterize surfaces?



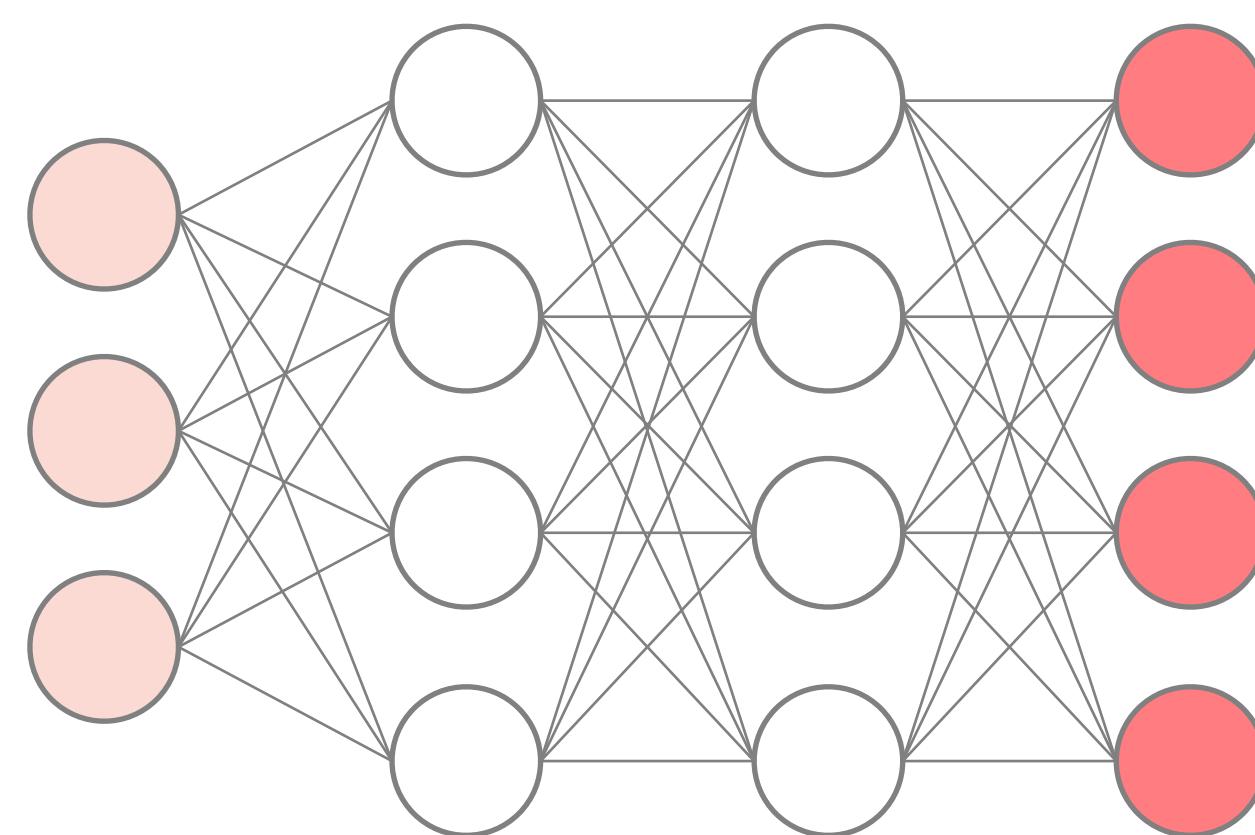
How to parameterize 3D function?



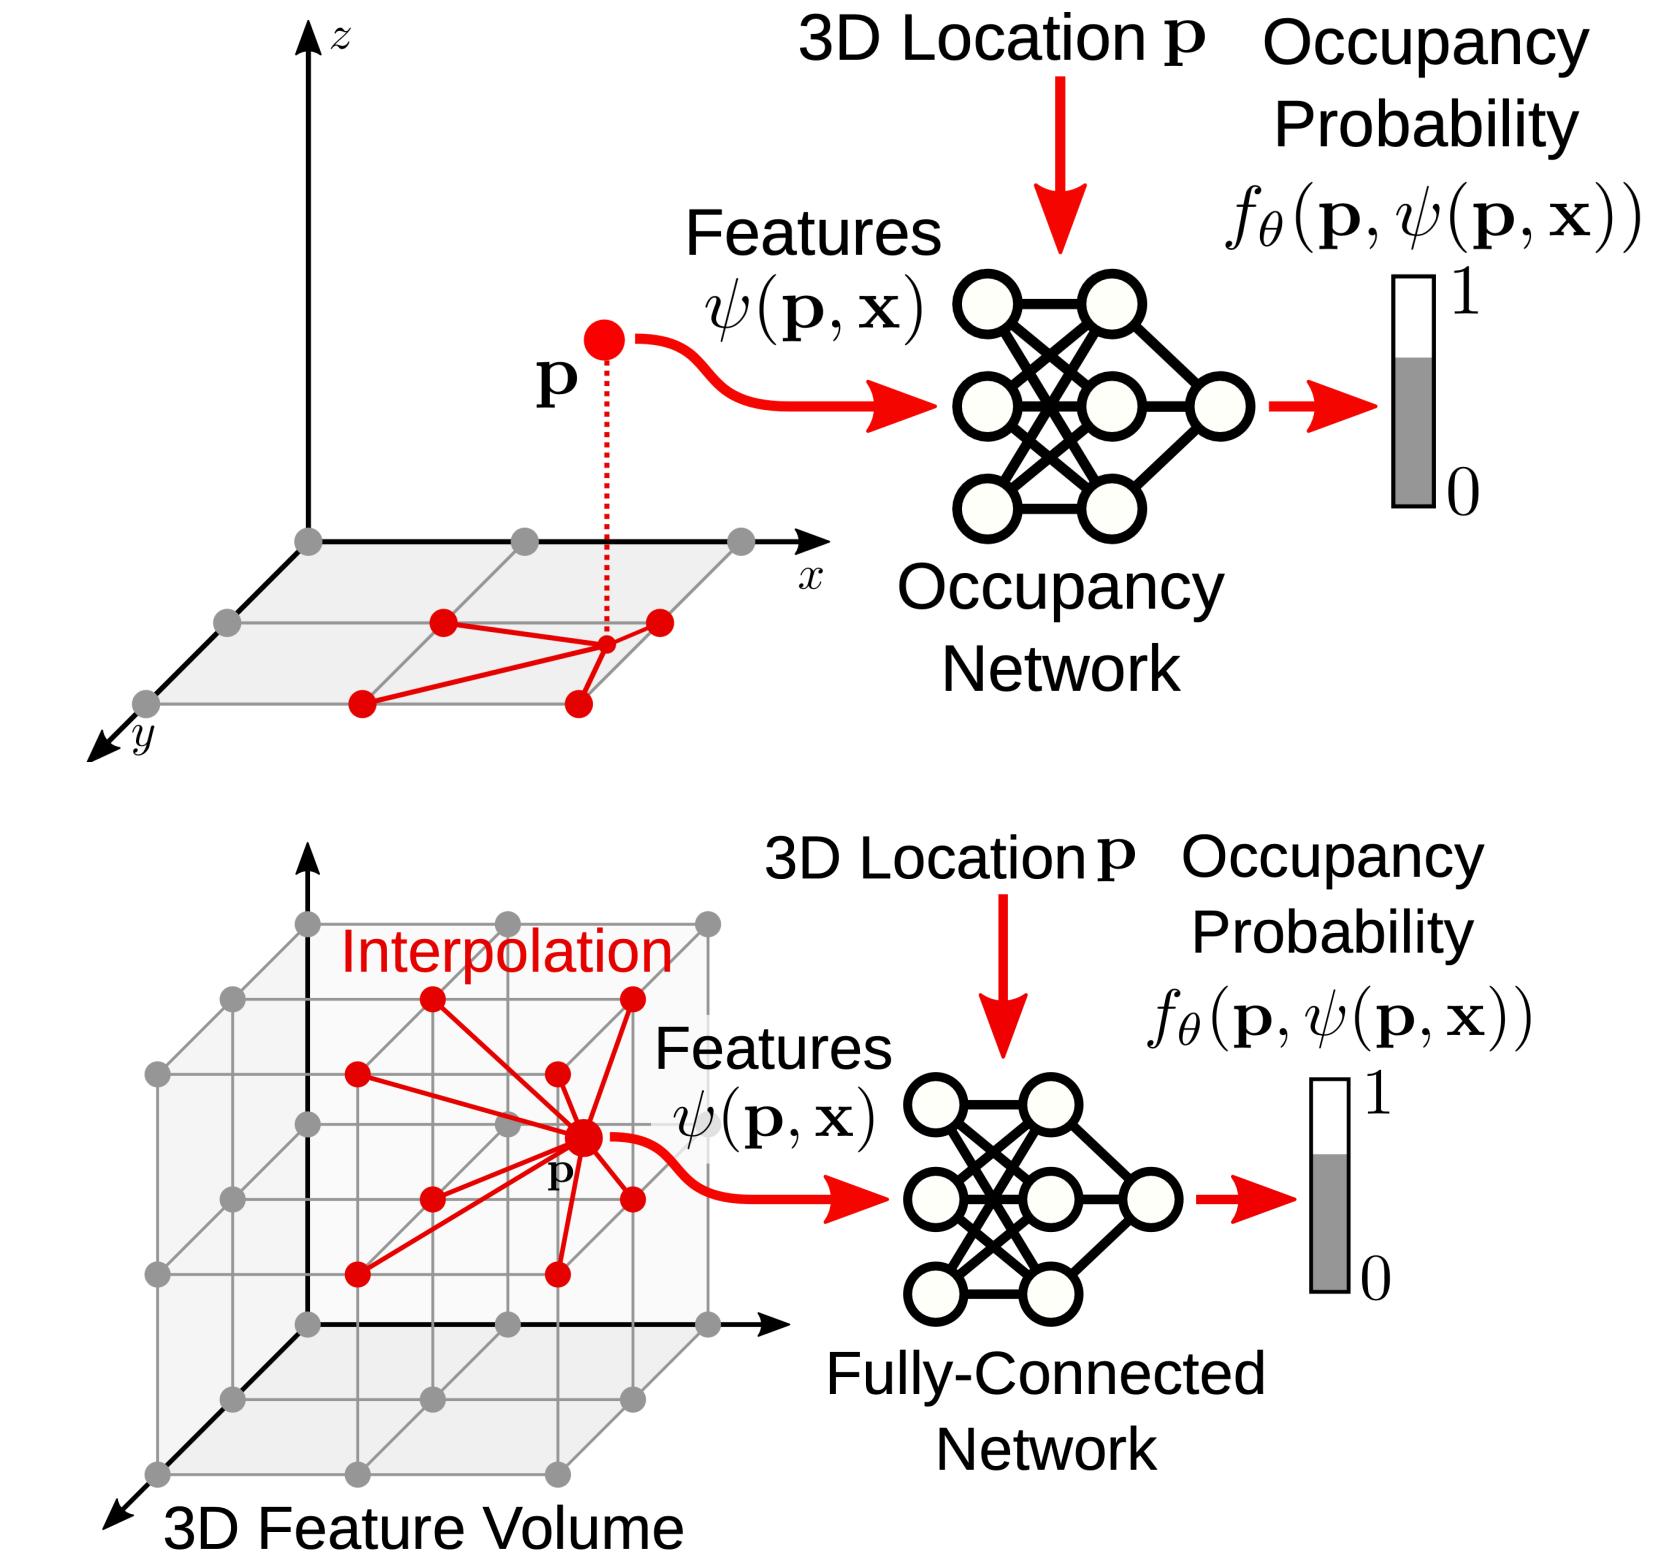
Parameterizing Fields



Discrete Parameterizations

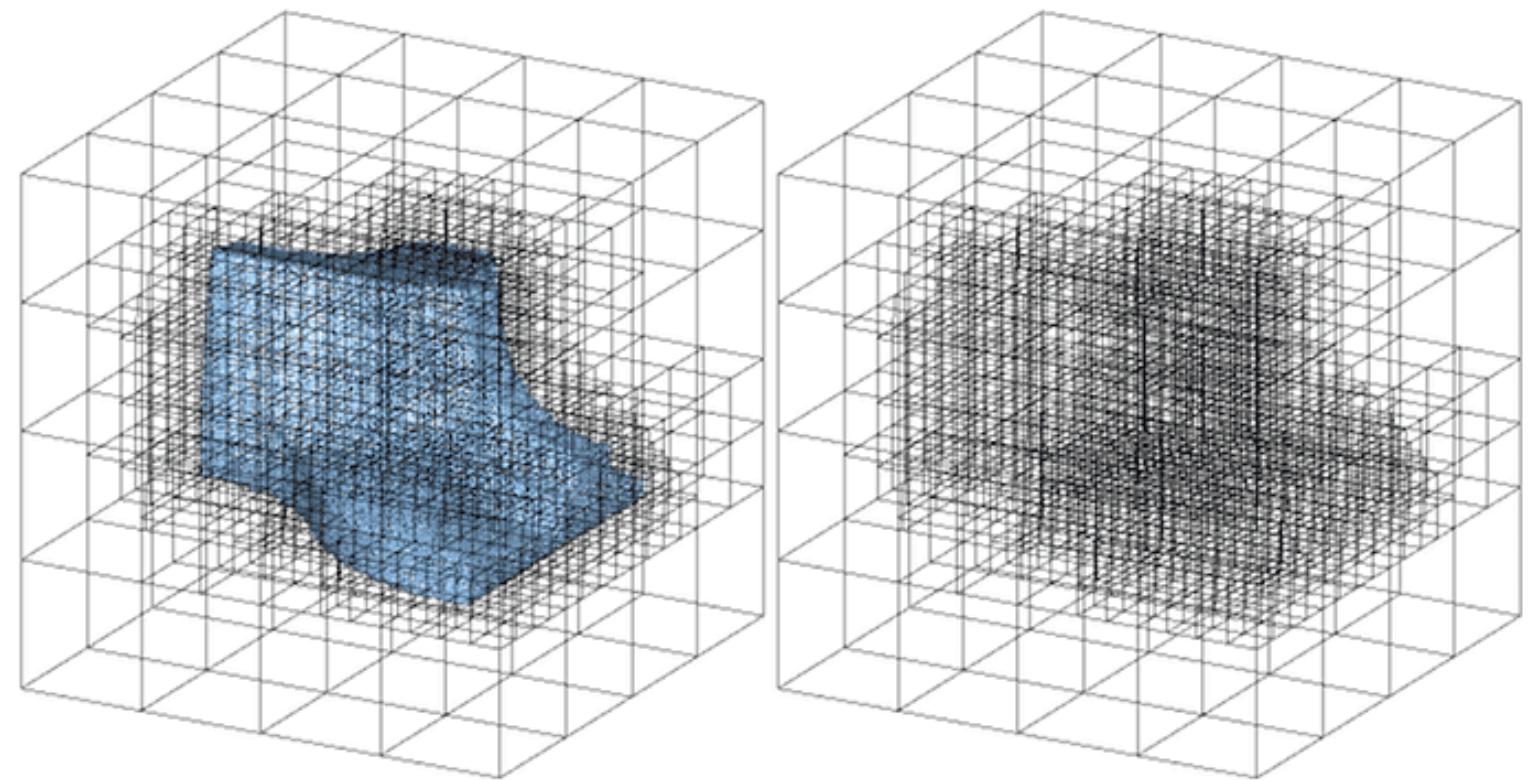
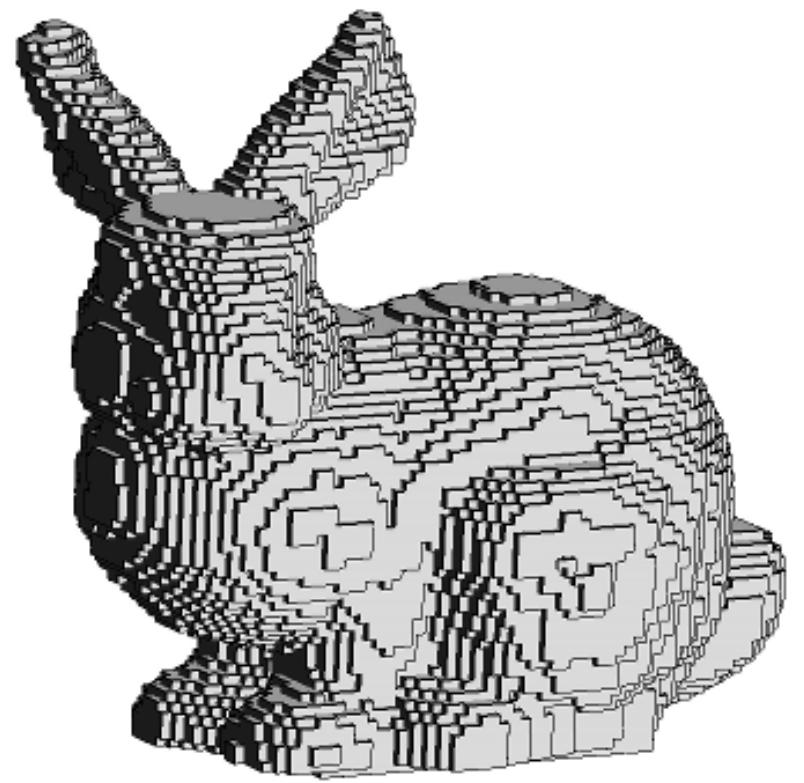


Continuous Parameterizations

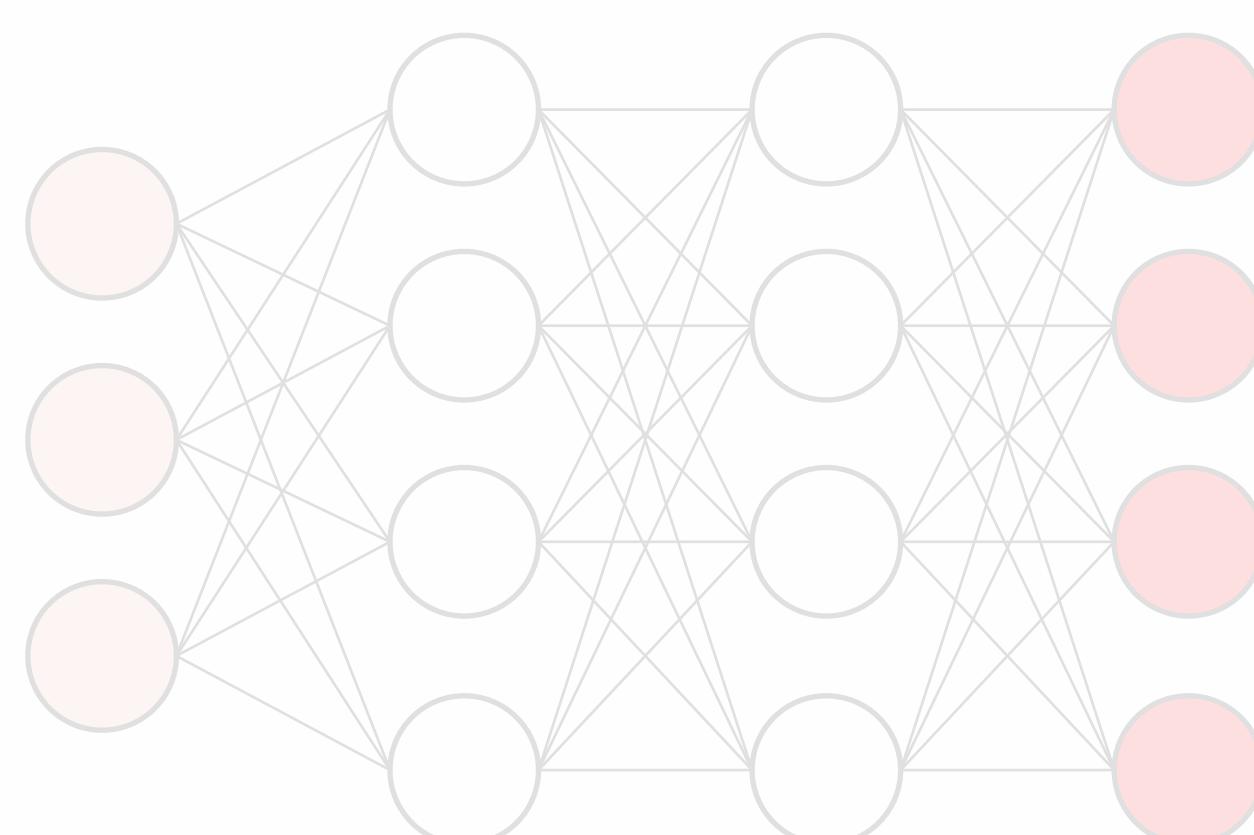


Hybrid Parameterizations

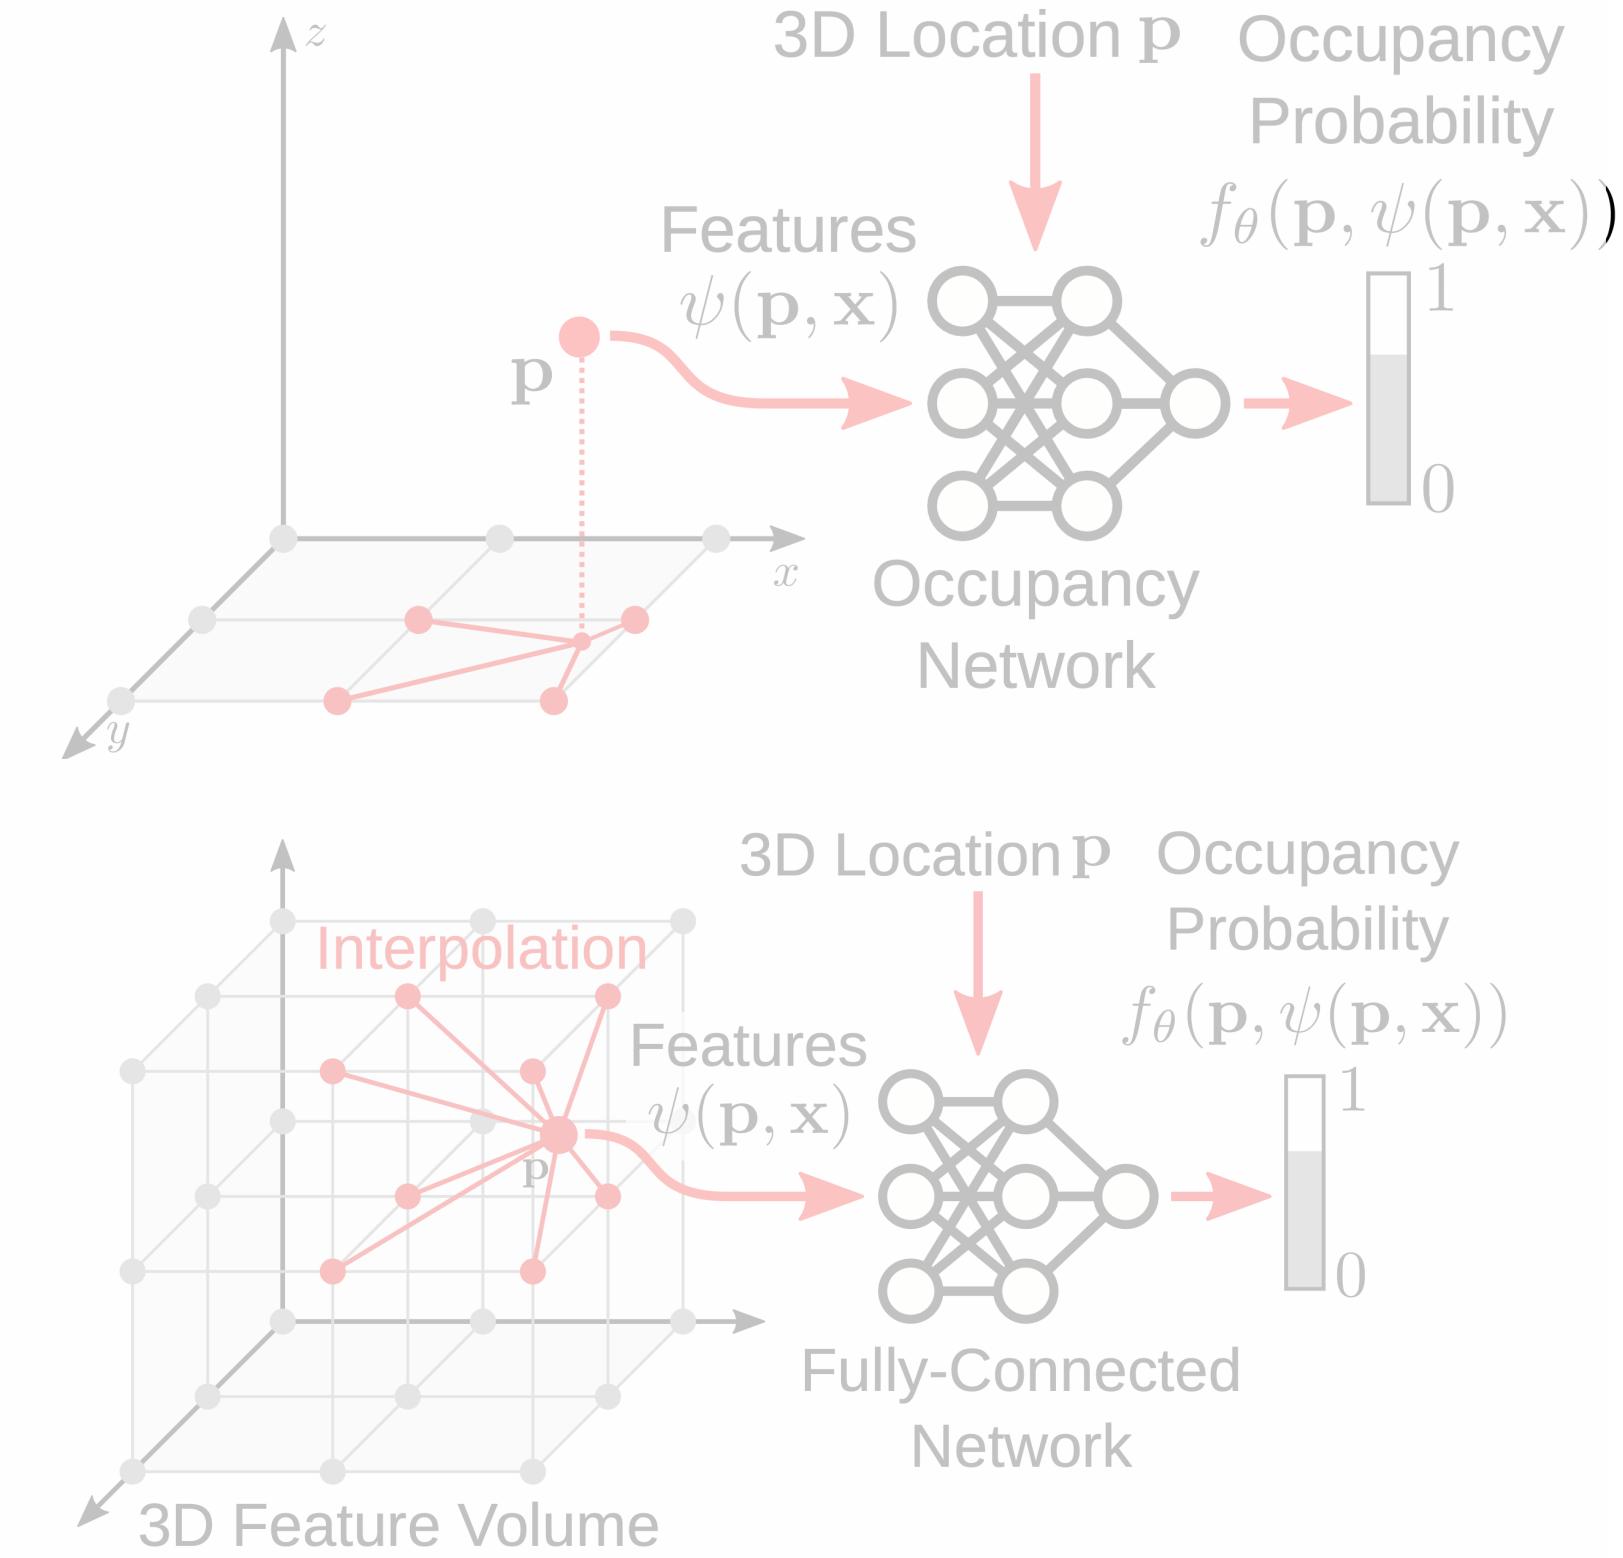
Voxel Grids



Discrete Parameterizations

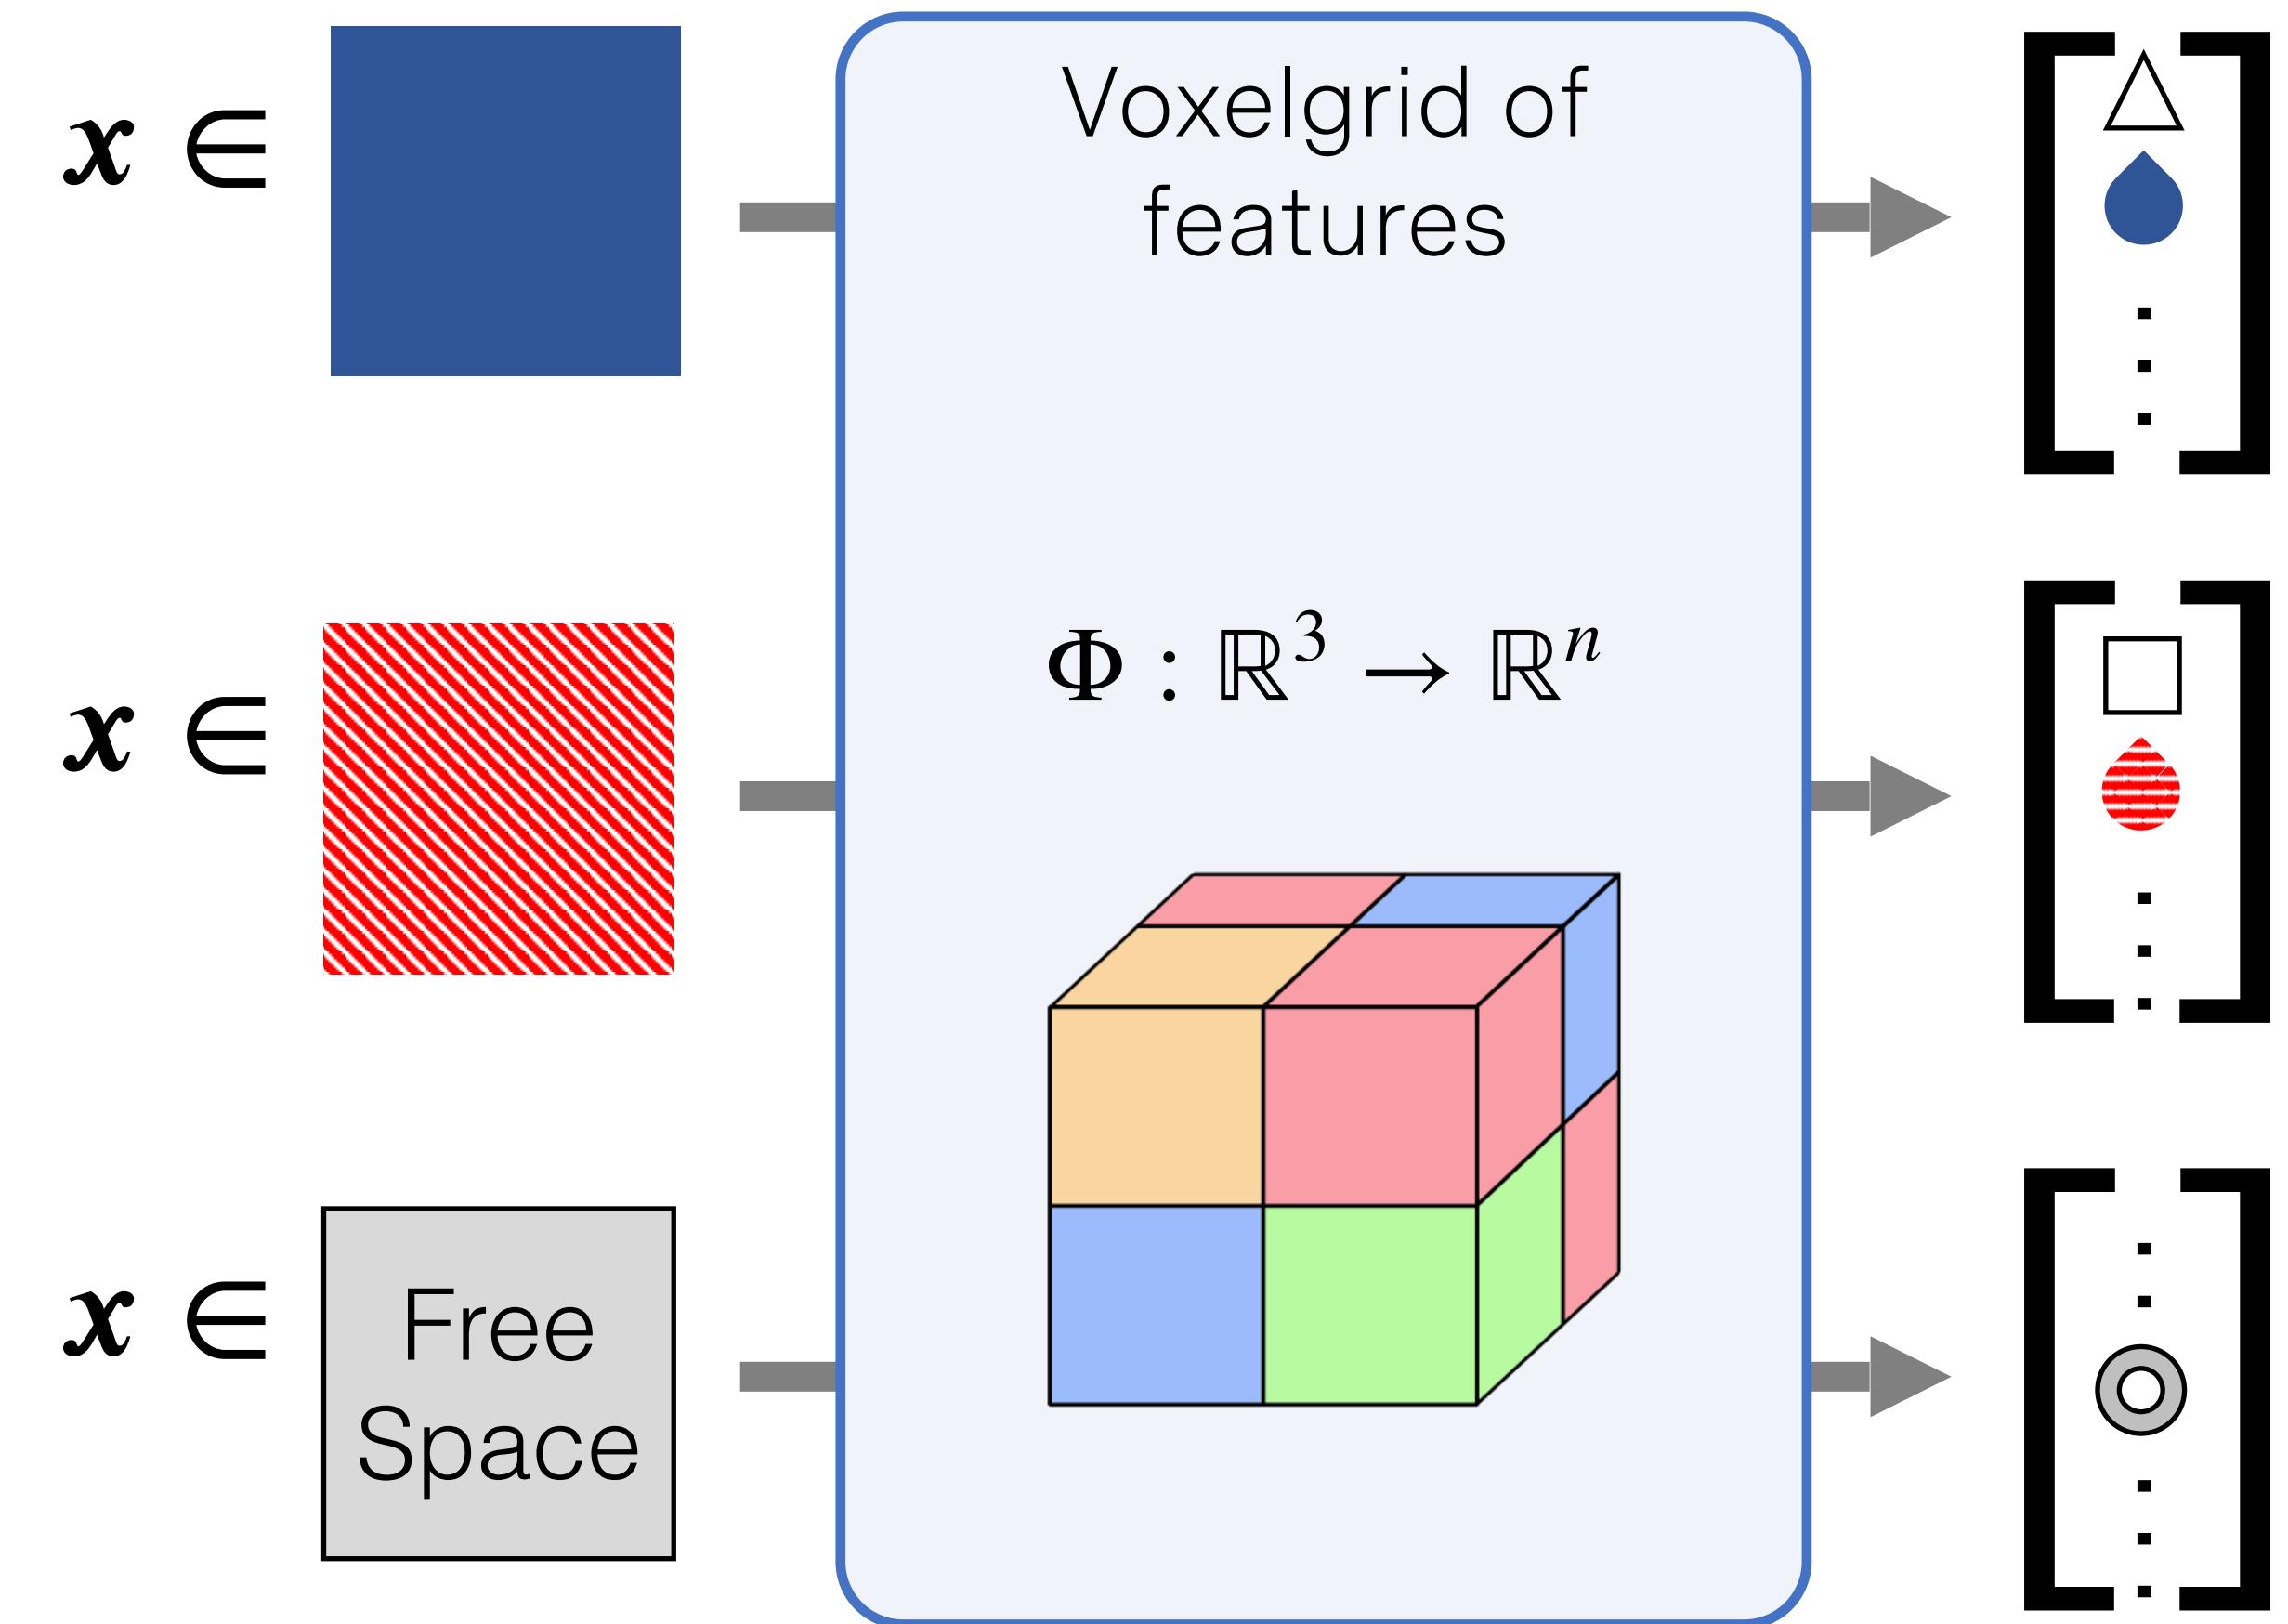
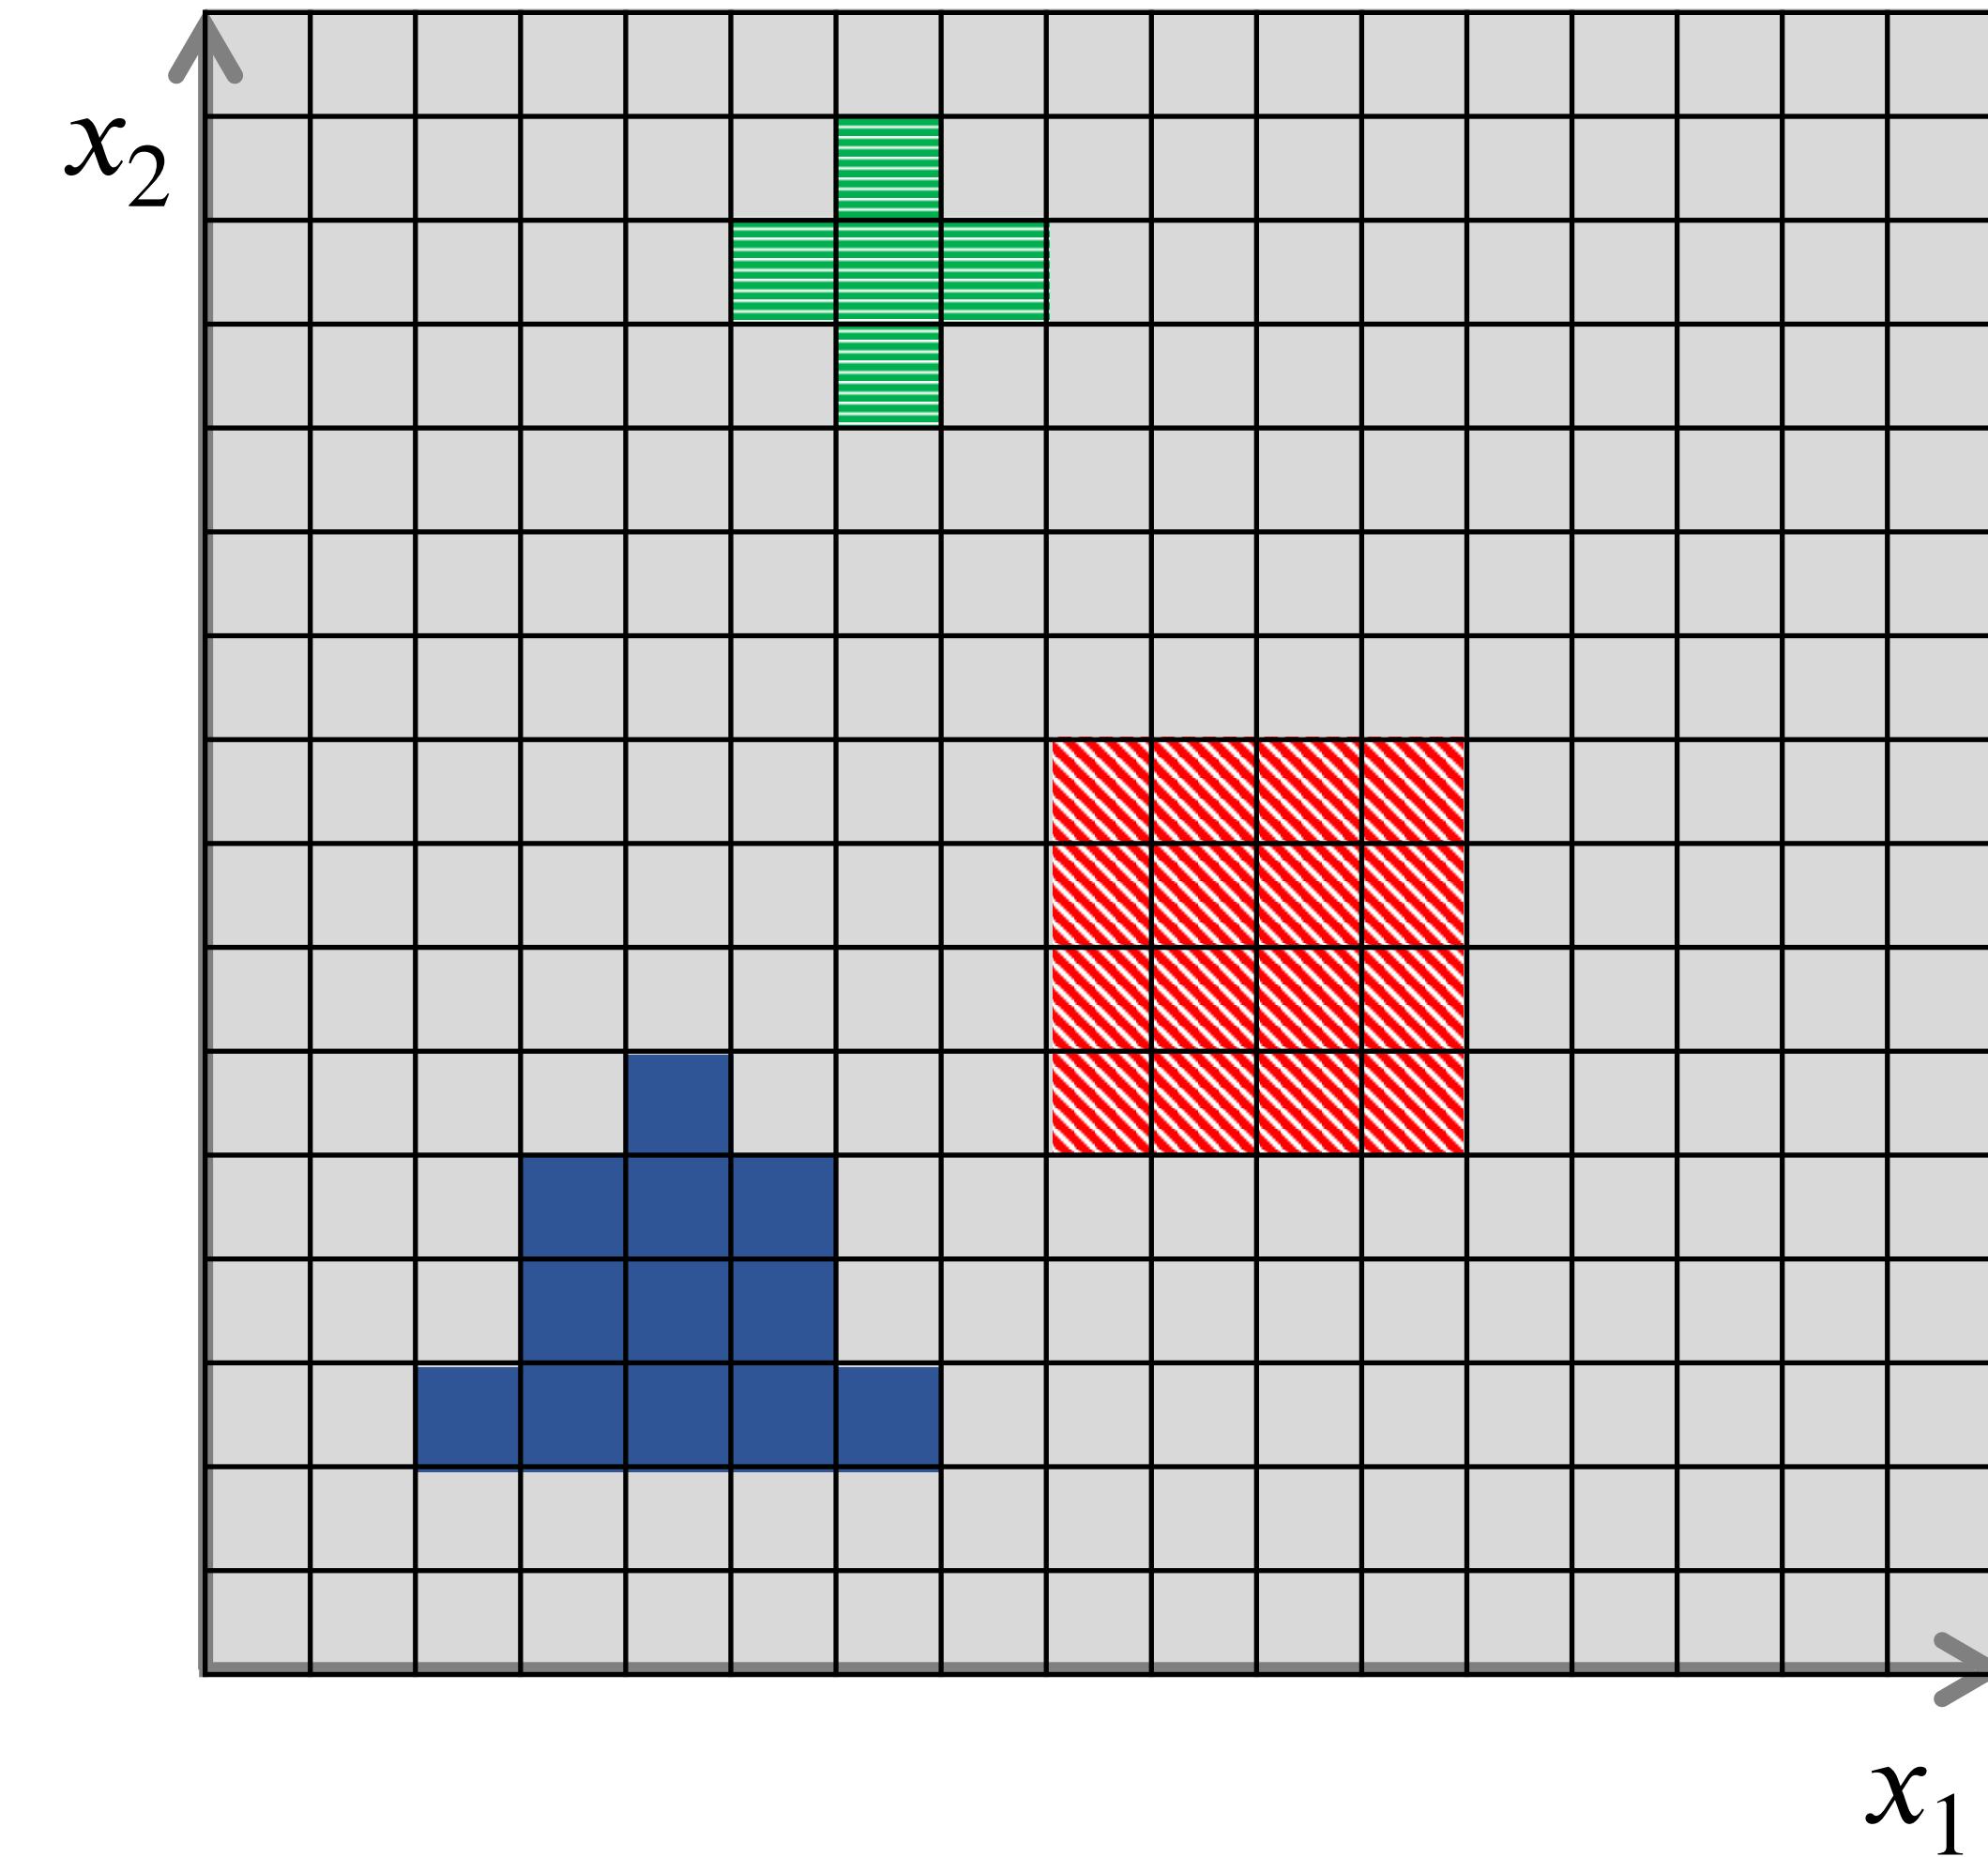


Continuous Parameterizations

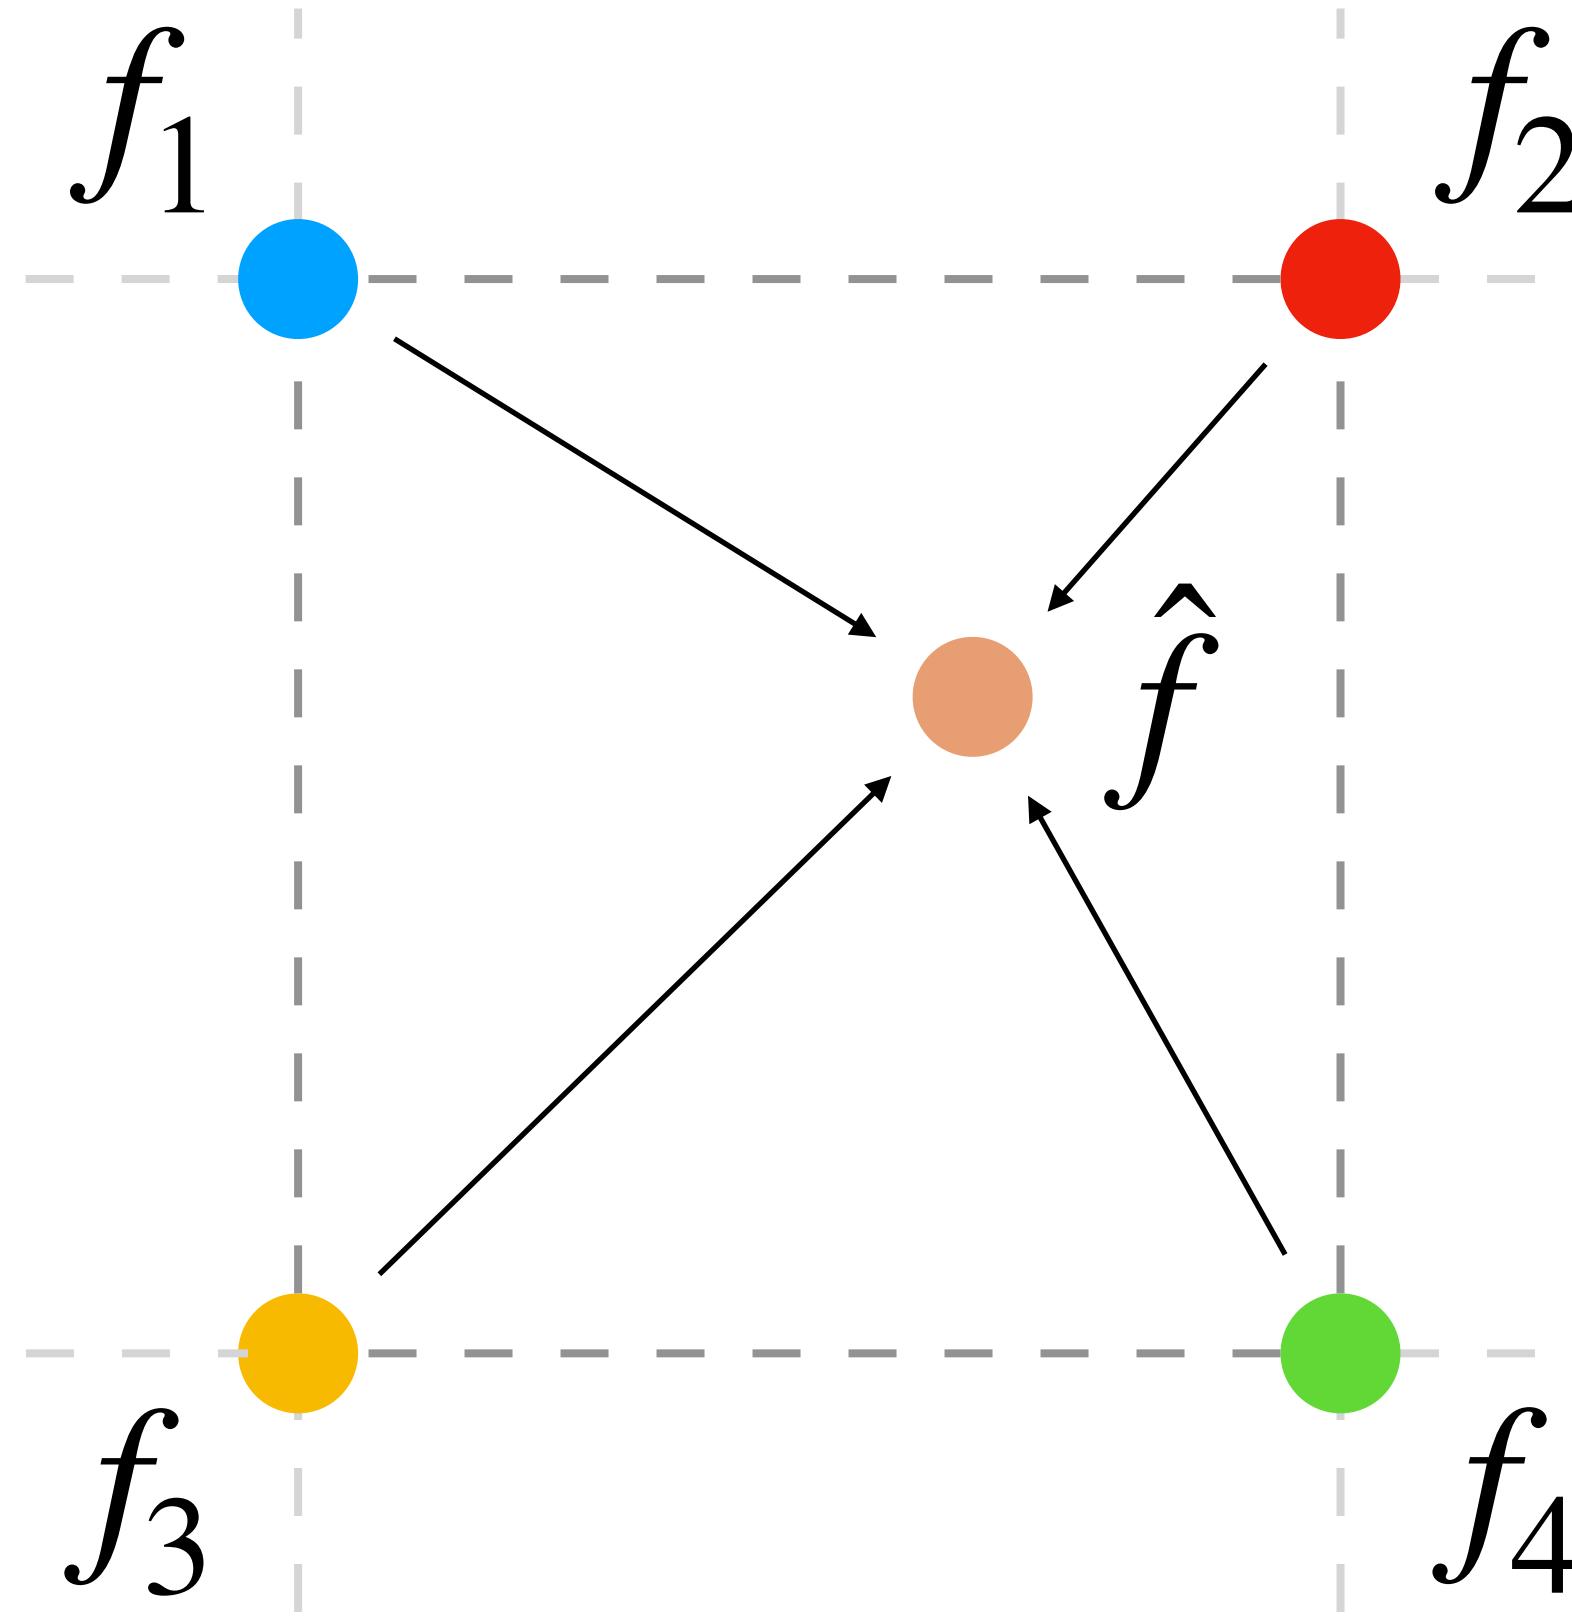


Hybrid Parameterizations

Voxel Grids



Interpolation: Querying *continuous* values



- Only stores values at vertex locations, i.e. corners
- Values at intermediate coordinates are defined via interpolation w/ some kernel k

$$\hat{f}(\mathbf{x}) = \sum_{i=1}^N f_i k(\mathbf{x}, \mathbf{x}_i)$$

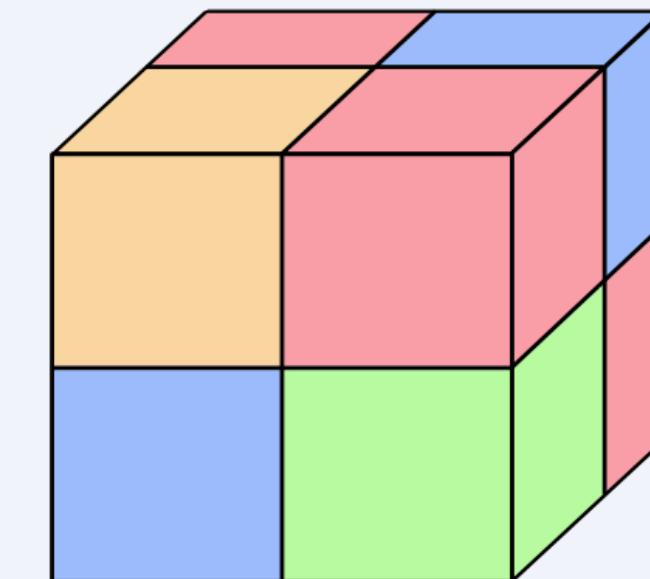
Interpolation in a 2D voxel grid.

Voxel grids

- For d spatial dimensions and resolution n ,
memory grows $O(n^d)$
- Fast sampling: d -linear interpolation
= index into array 2^d times & weighted sum.
- Convenient processing as it exposes *locality*. Can
easily run convolutions, nearest-neighbor lookups,
etc.
- Intractable in higher dimensions :/

Voxelgrid of
features

$$\Phi : \mathbb{R}^3 \rightarrow \mathbb{R}^n$$



Multi-Resolution Representations

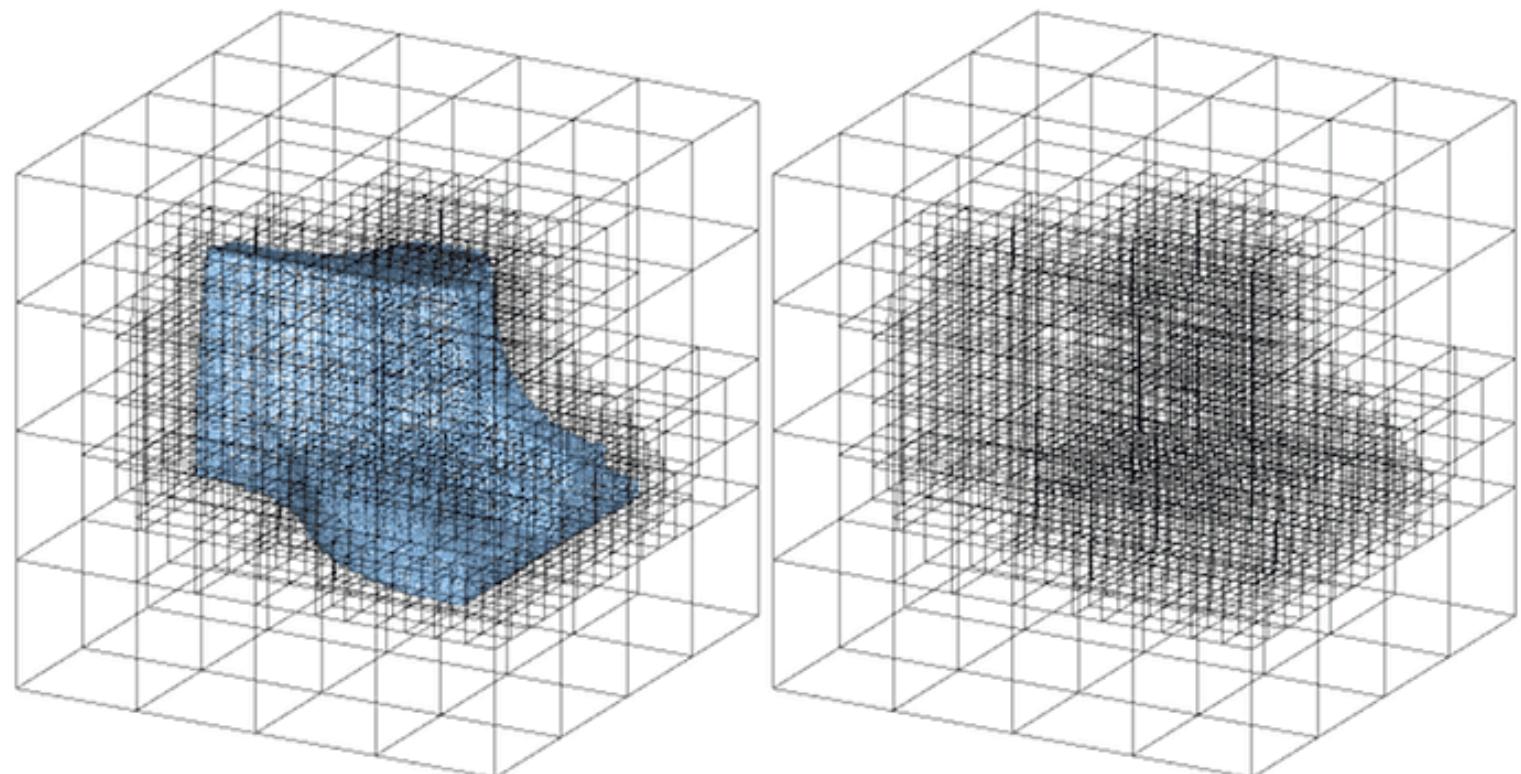
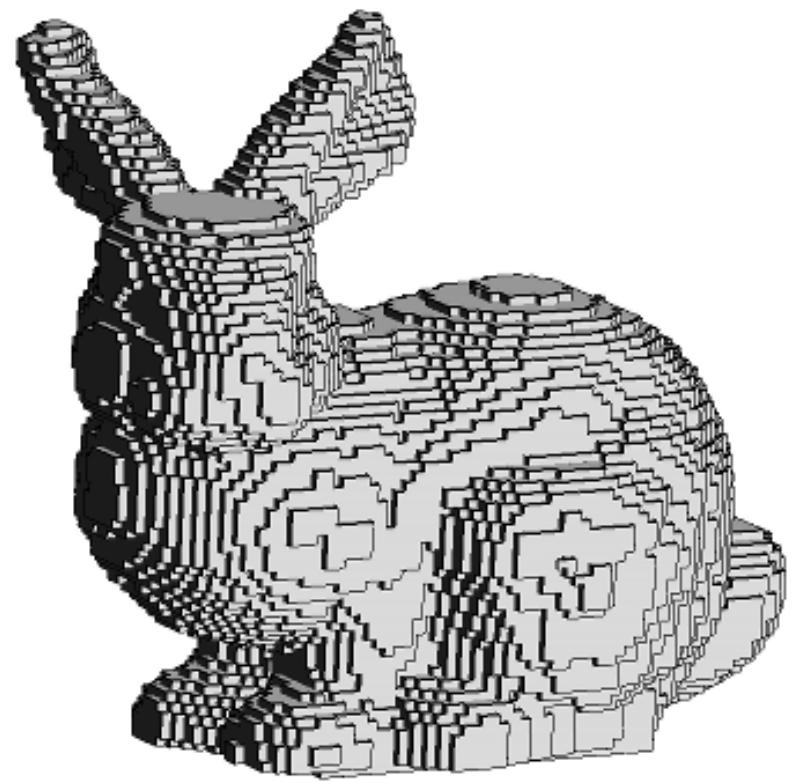
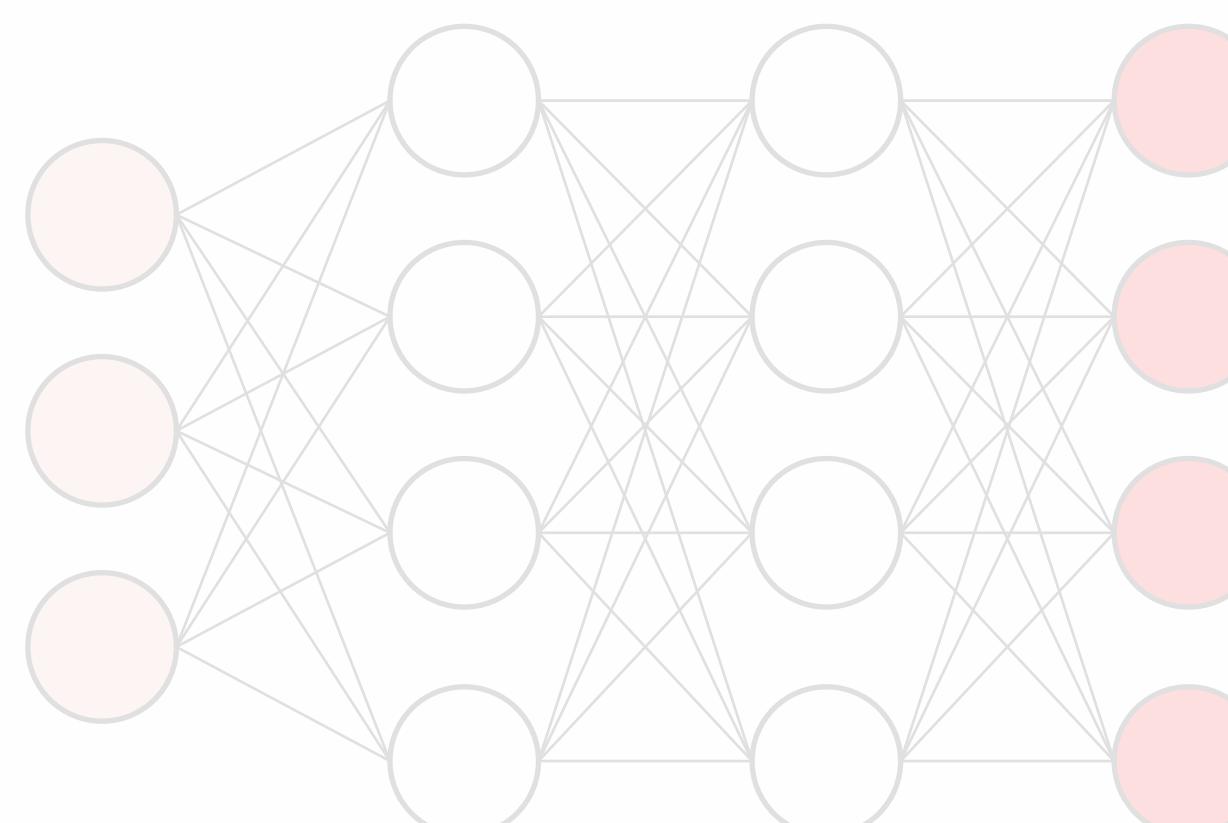
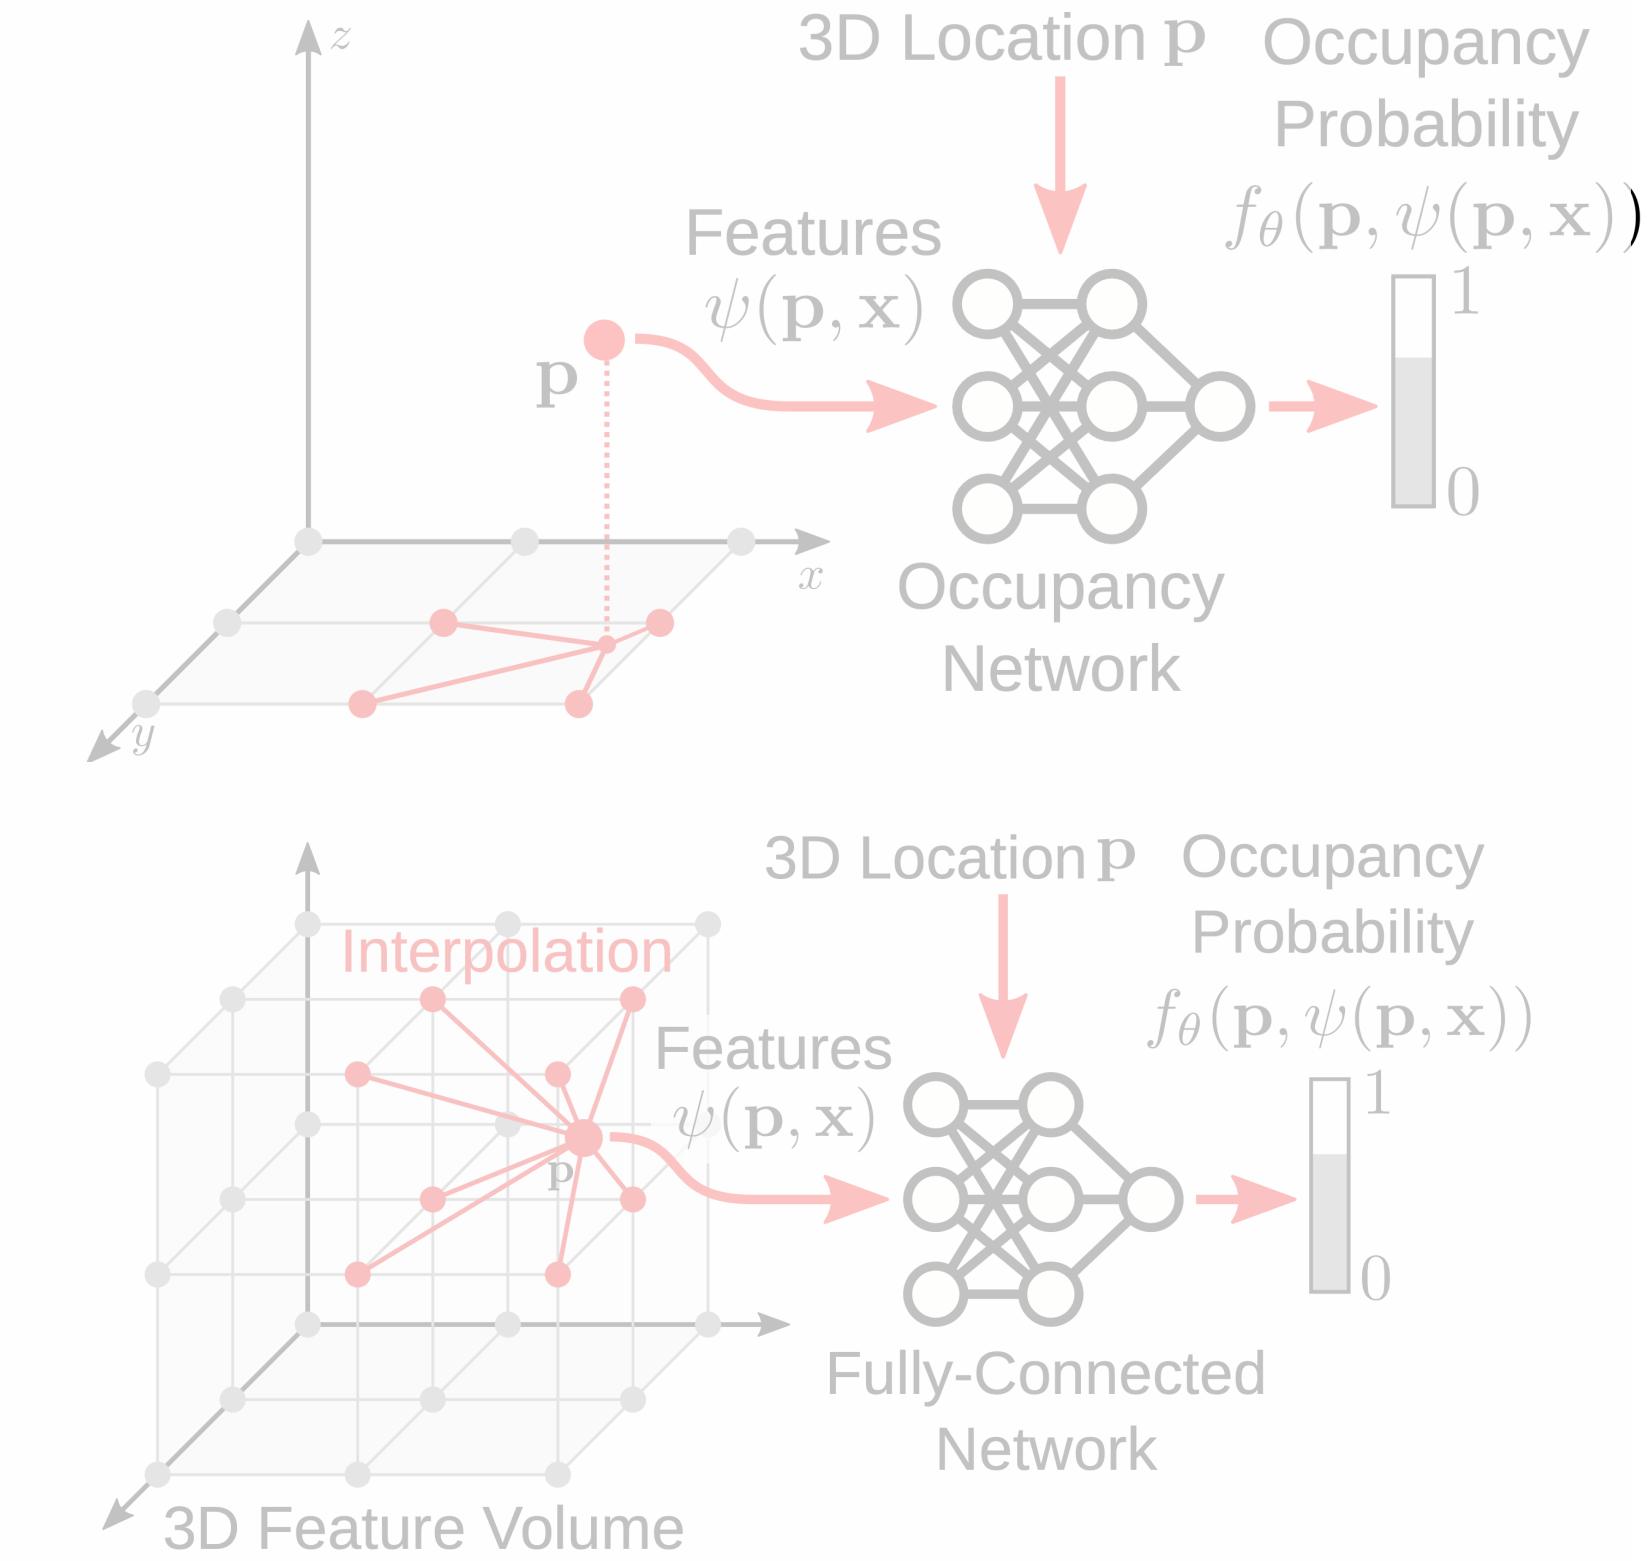


Image Source: <https://doc.cgal.org/latest/Orthree/index.html>

Discrete Parameterizations

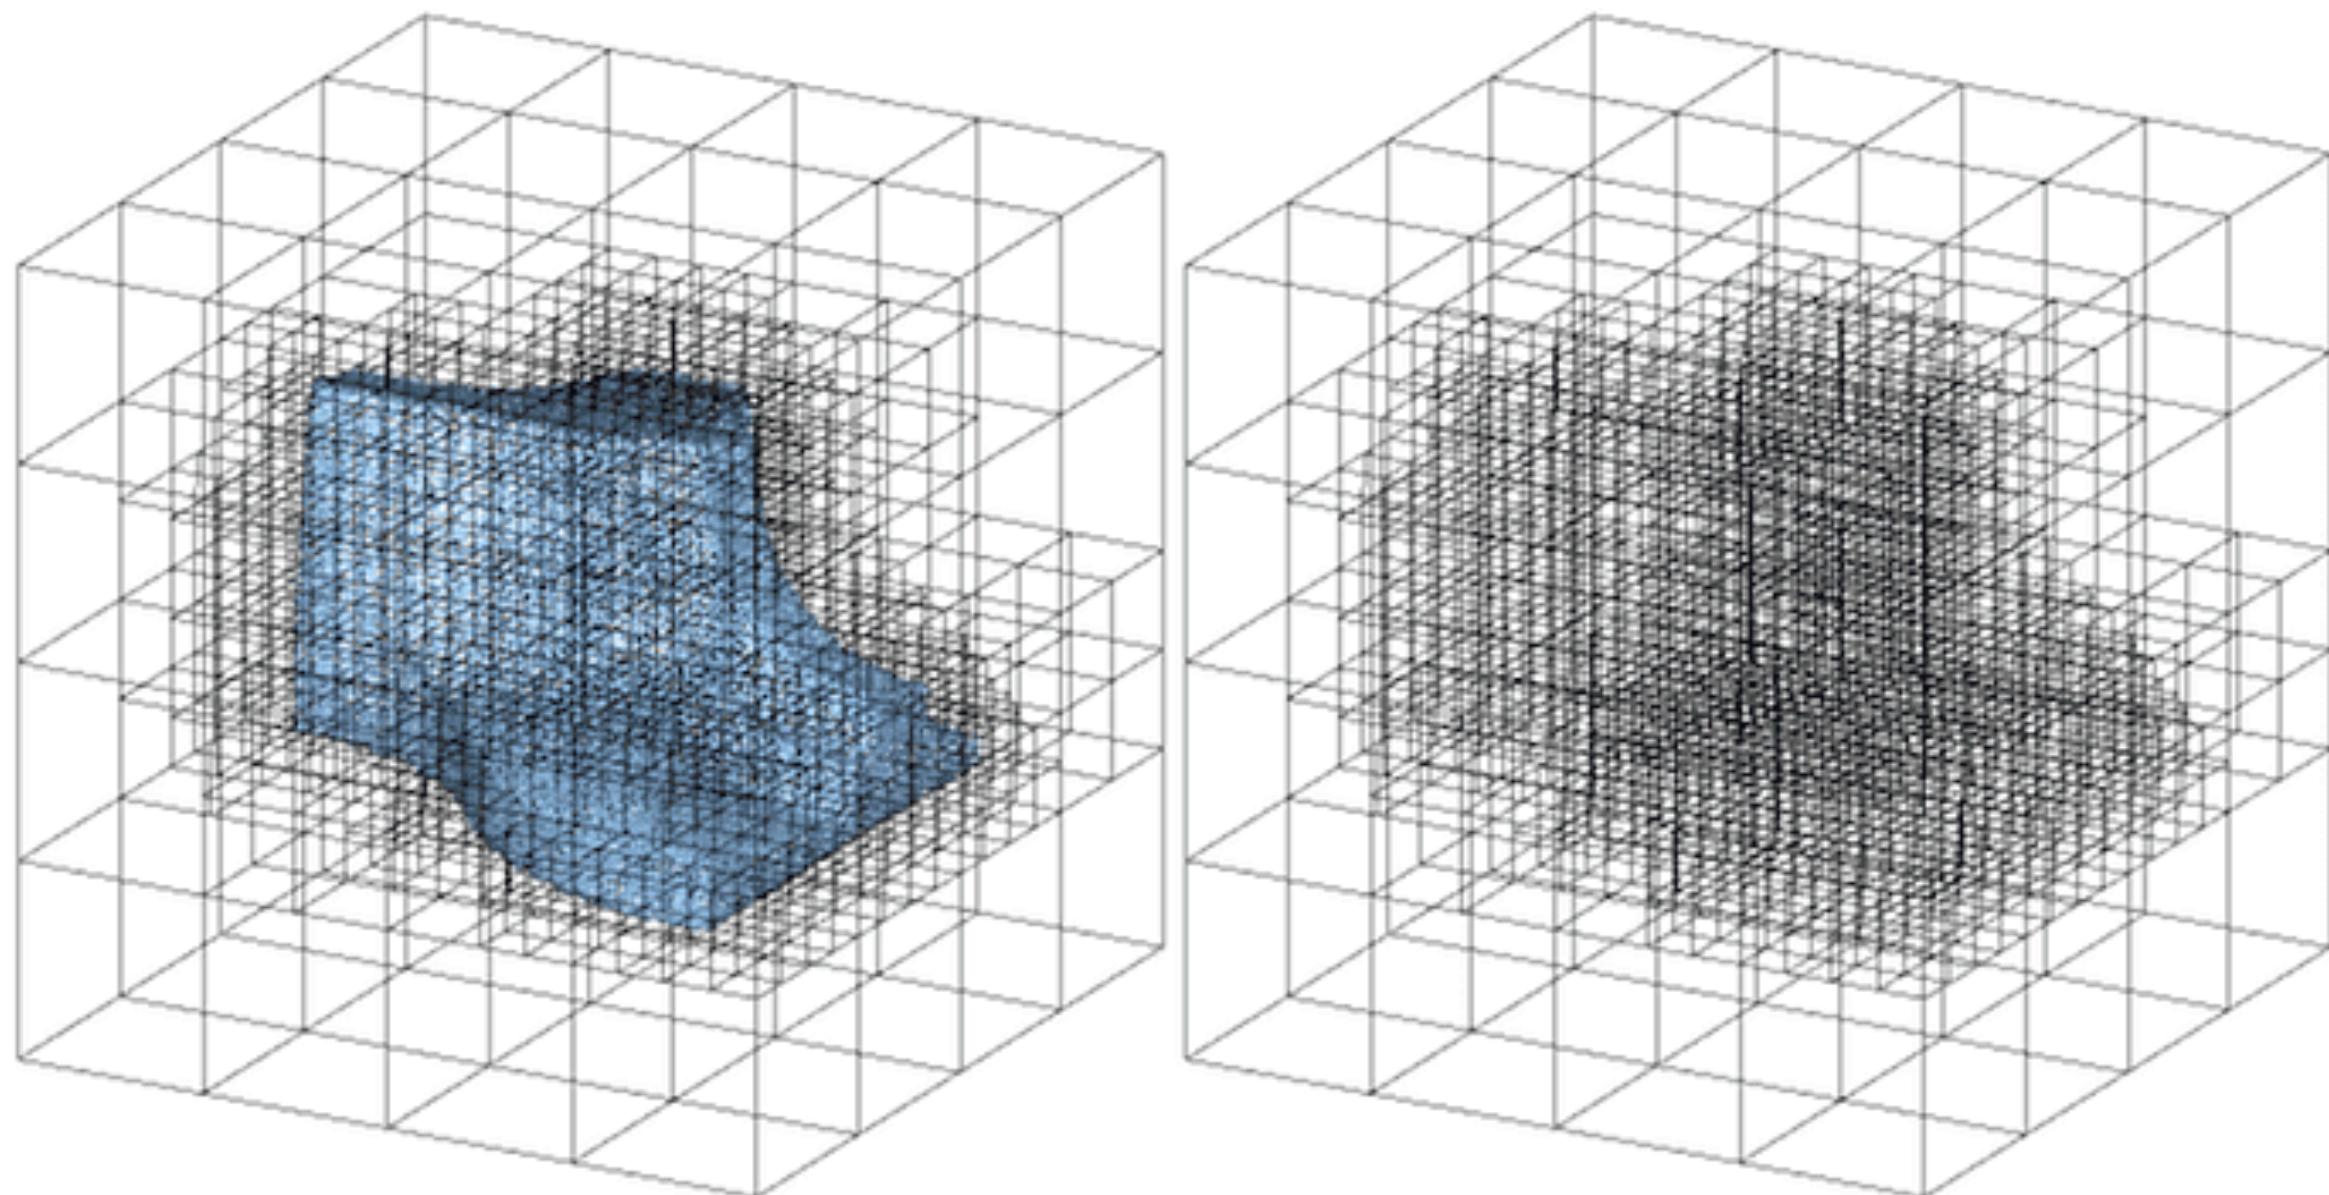


Continuous Parameterizations



Hybrid Parameterizations

Octrees



Divide volume into 8 voxels

If any box has more than N points in it, divide it into 8 voxels

Do not divide voxels with zero points.

Image Source: <https://doc.cgal.org/latest/Orytree/index.html>

Octrees

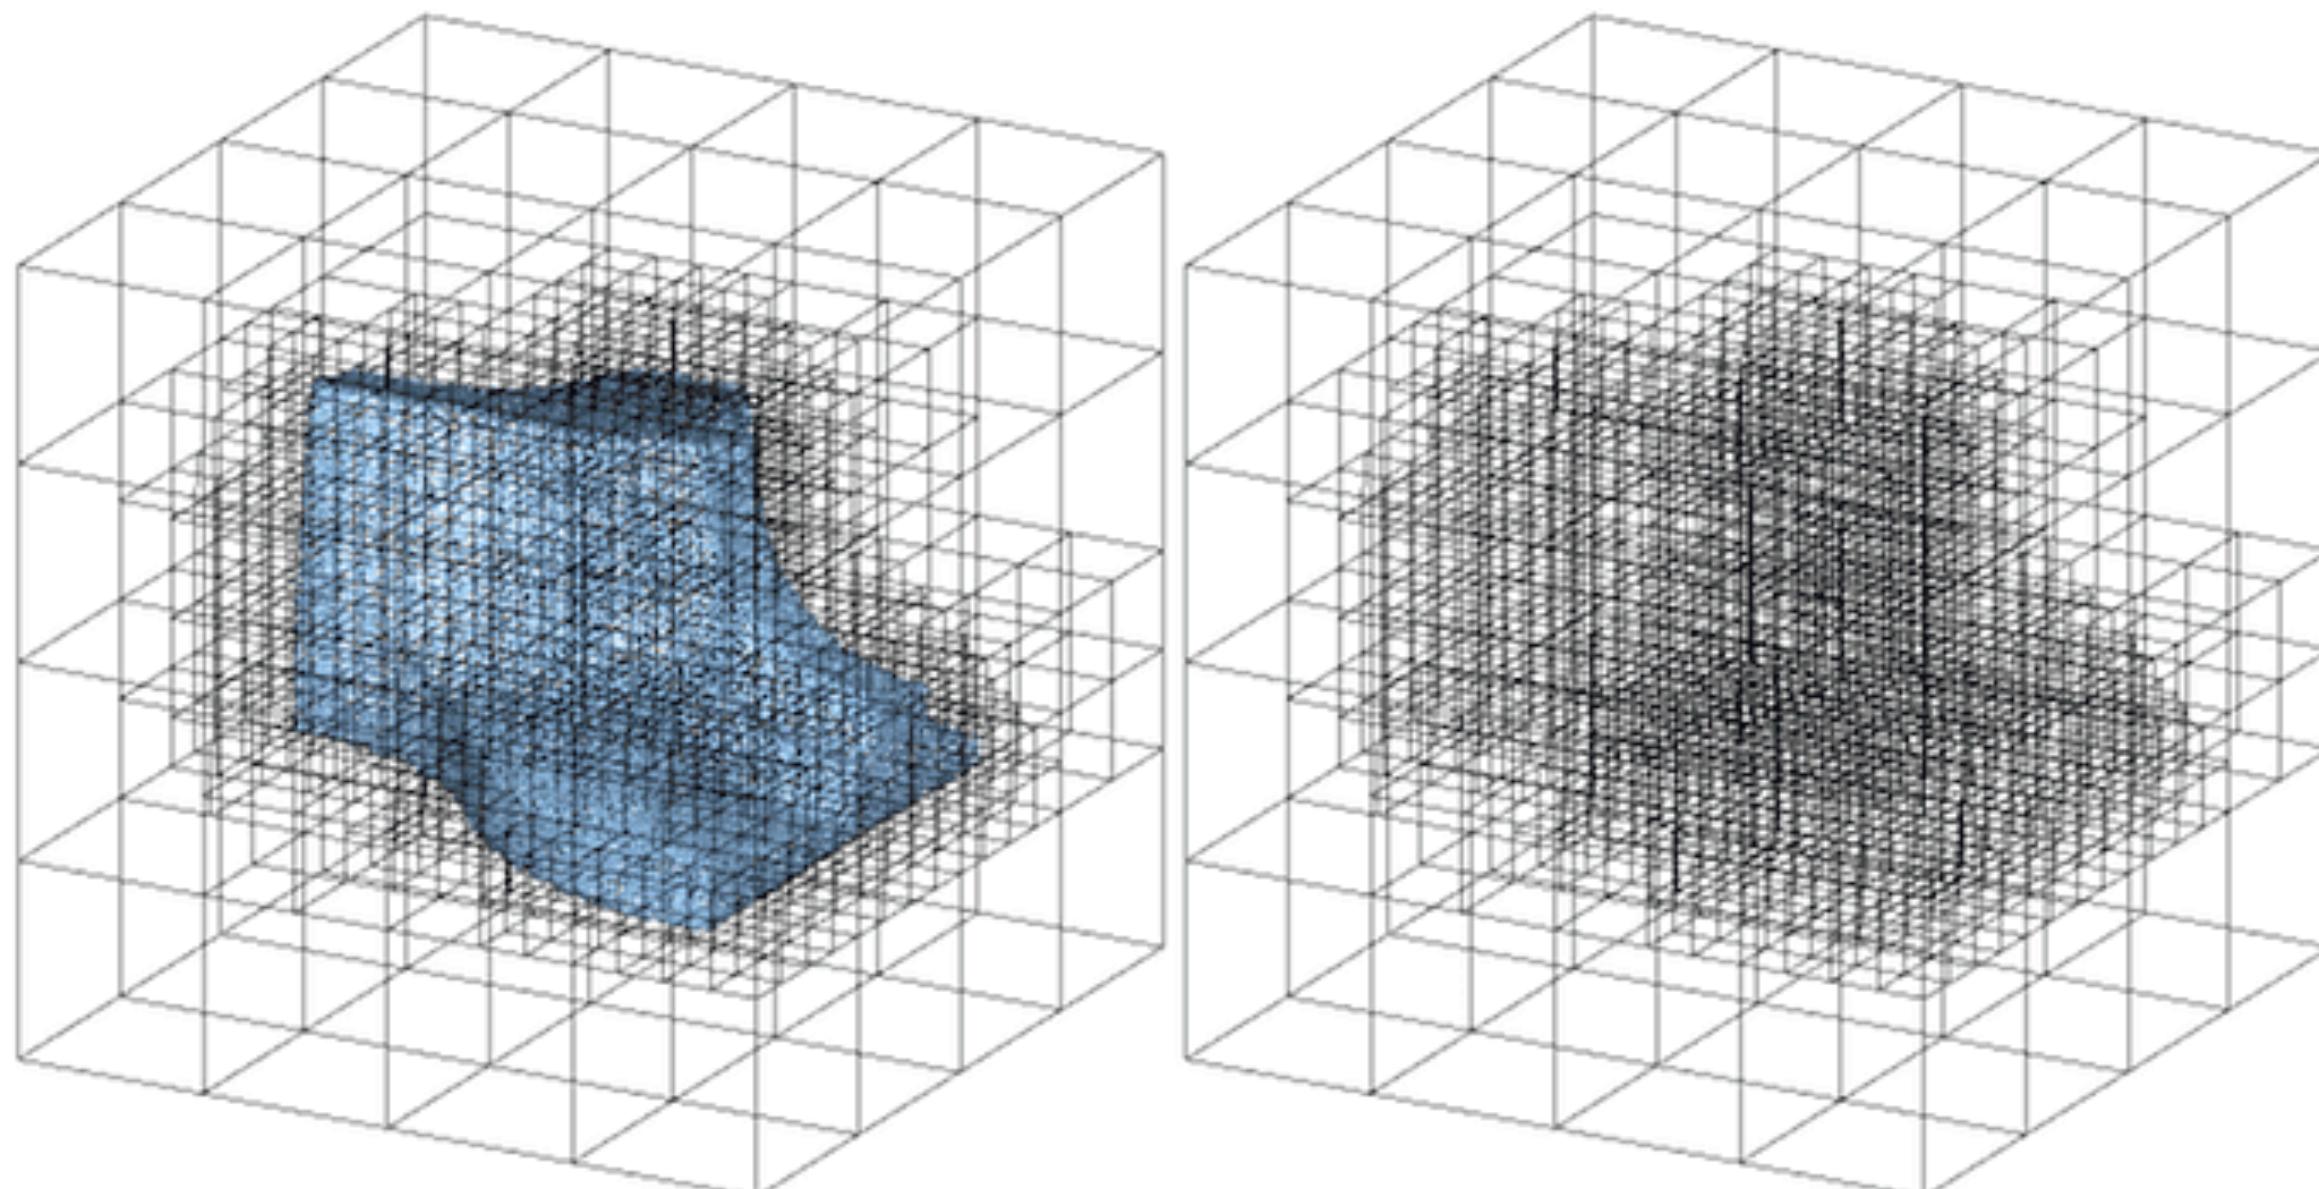


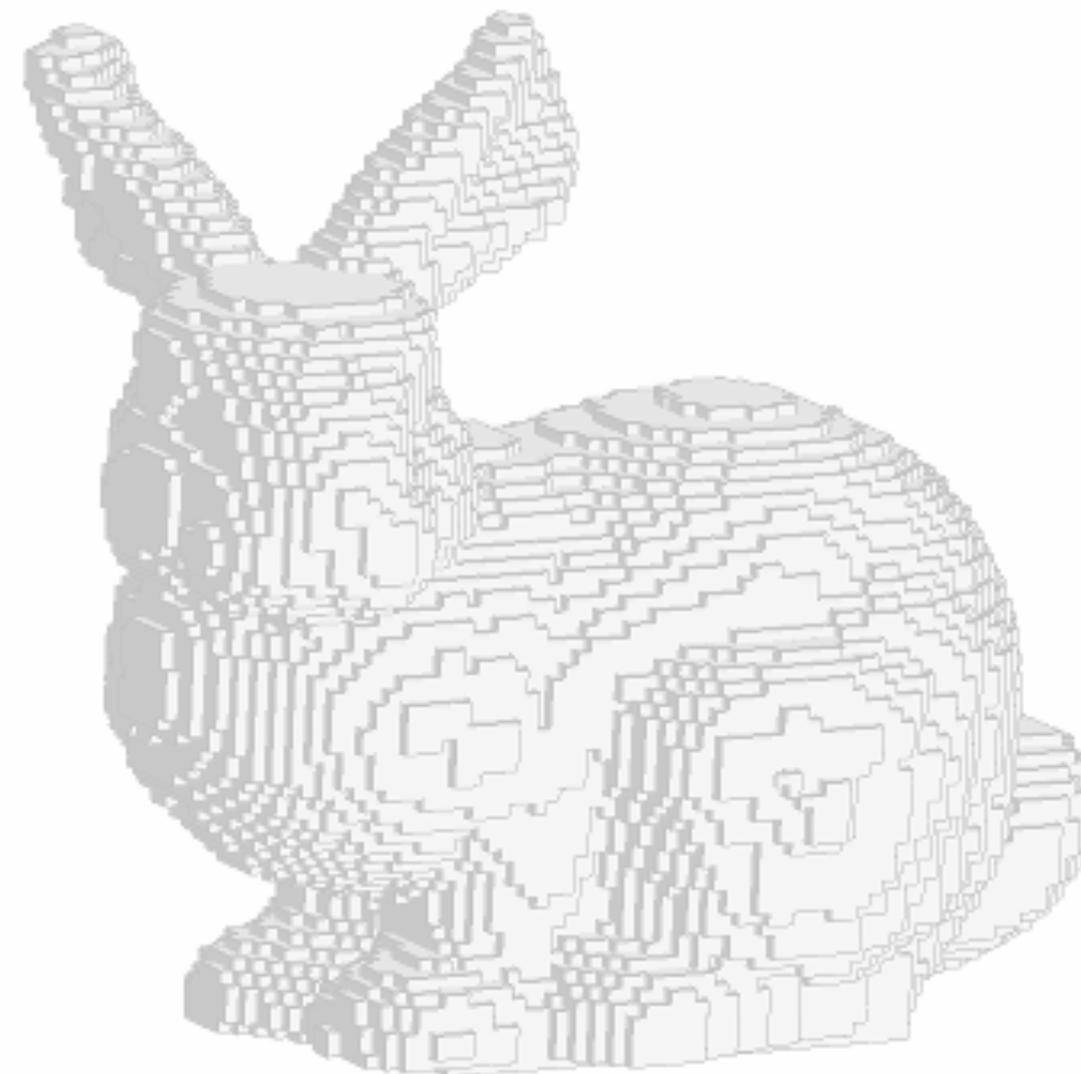
Image Source: <https://doc.cgal.org/latest/Orytree/index.html>

Saves memory (less resolution in empty parts)

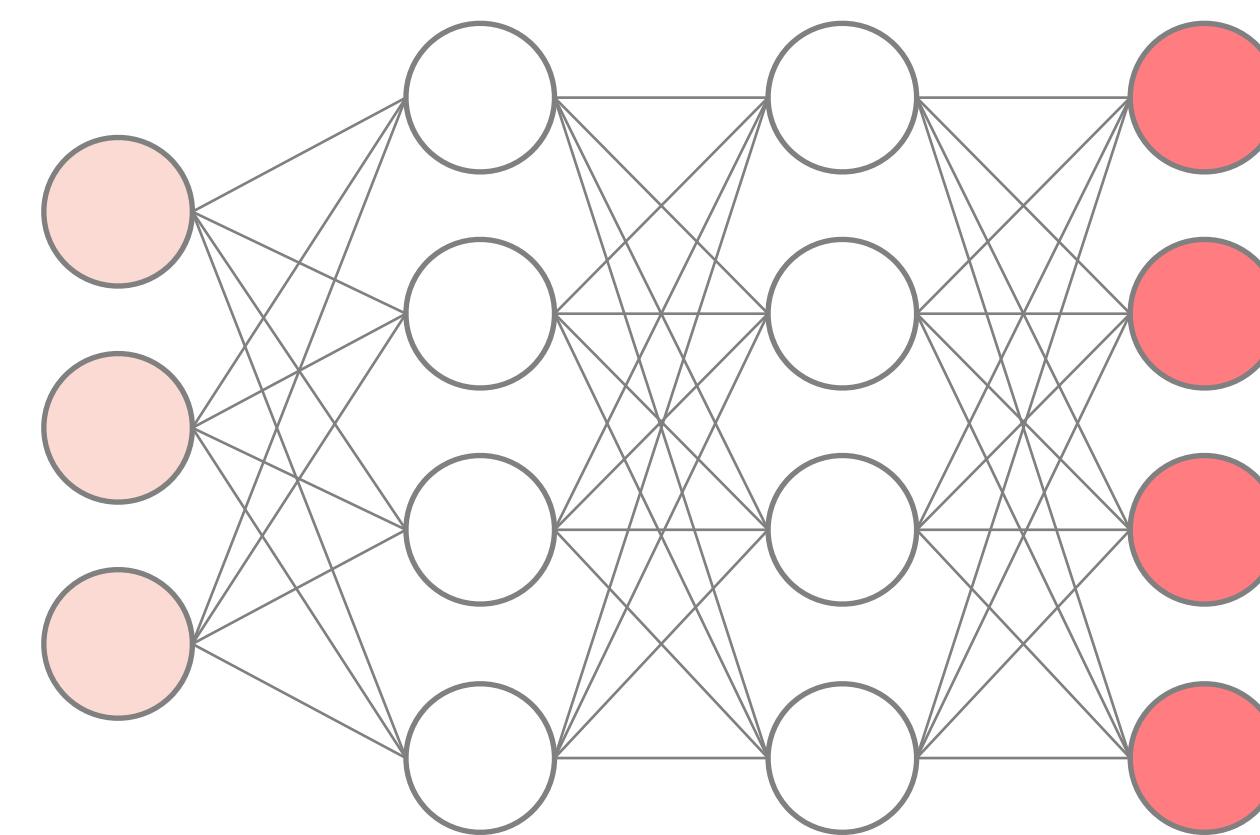
Sub-division is not differentiable:
can't easily build neural network
that outputs Octree

To sample point, have to traverse
multi-resolution hierarchy. Still
fast, though.

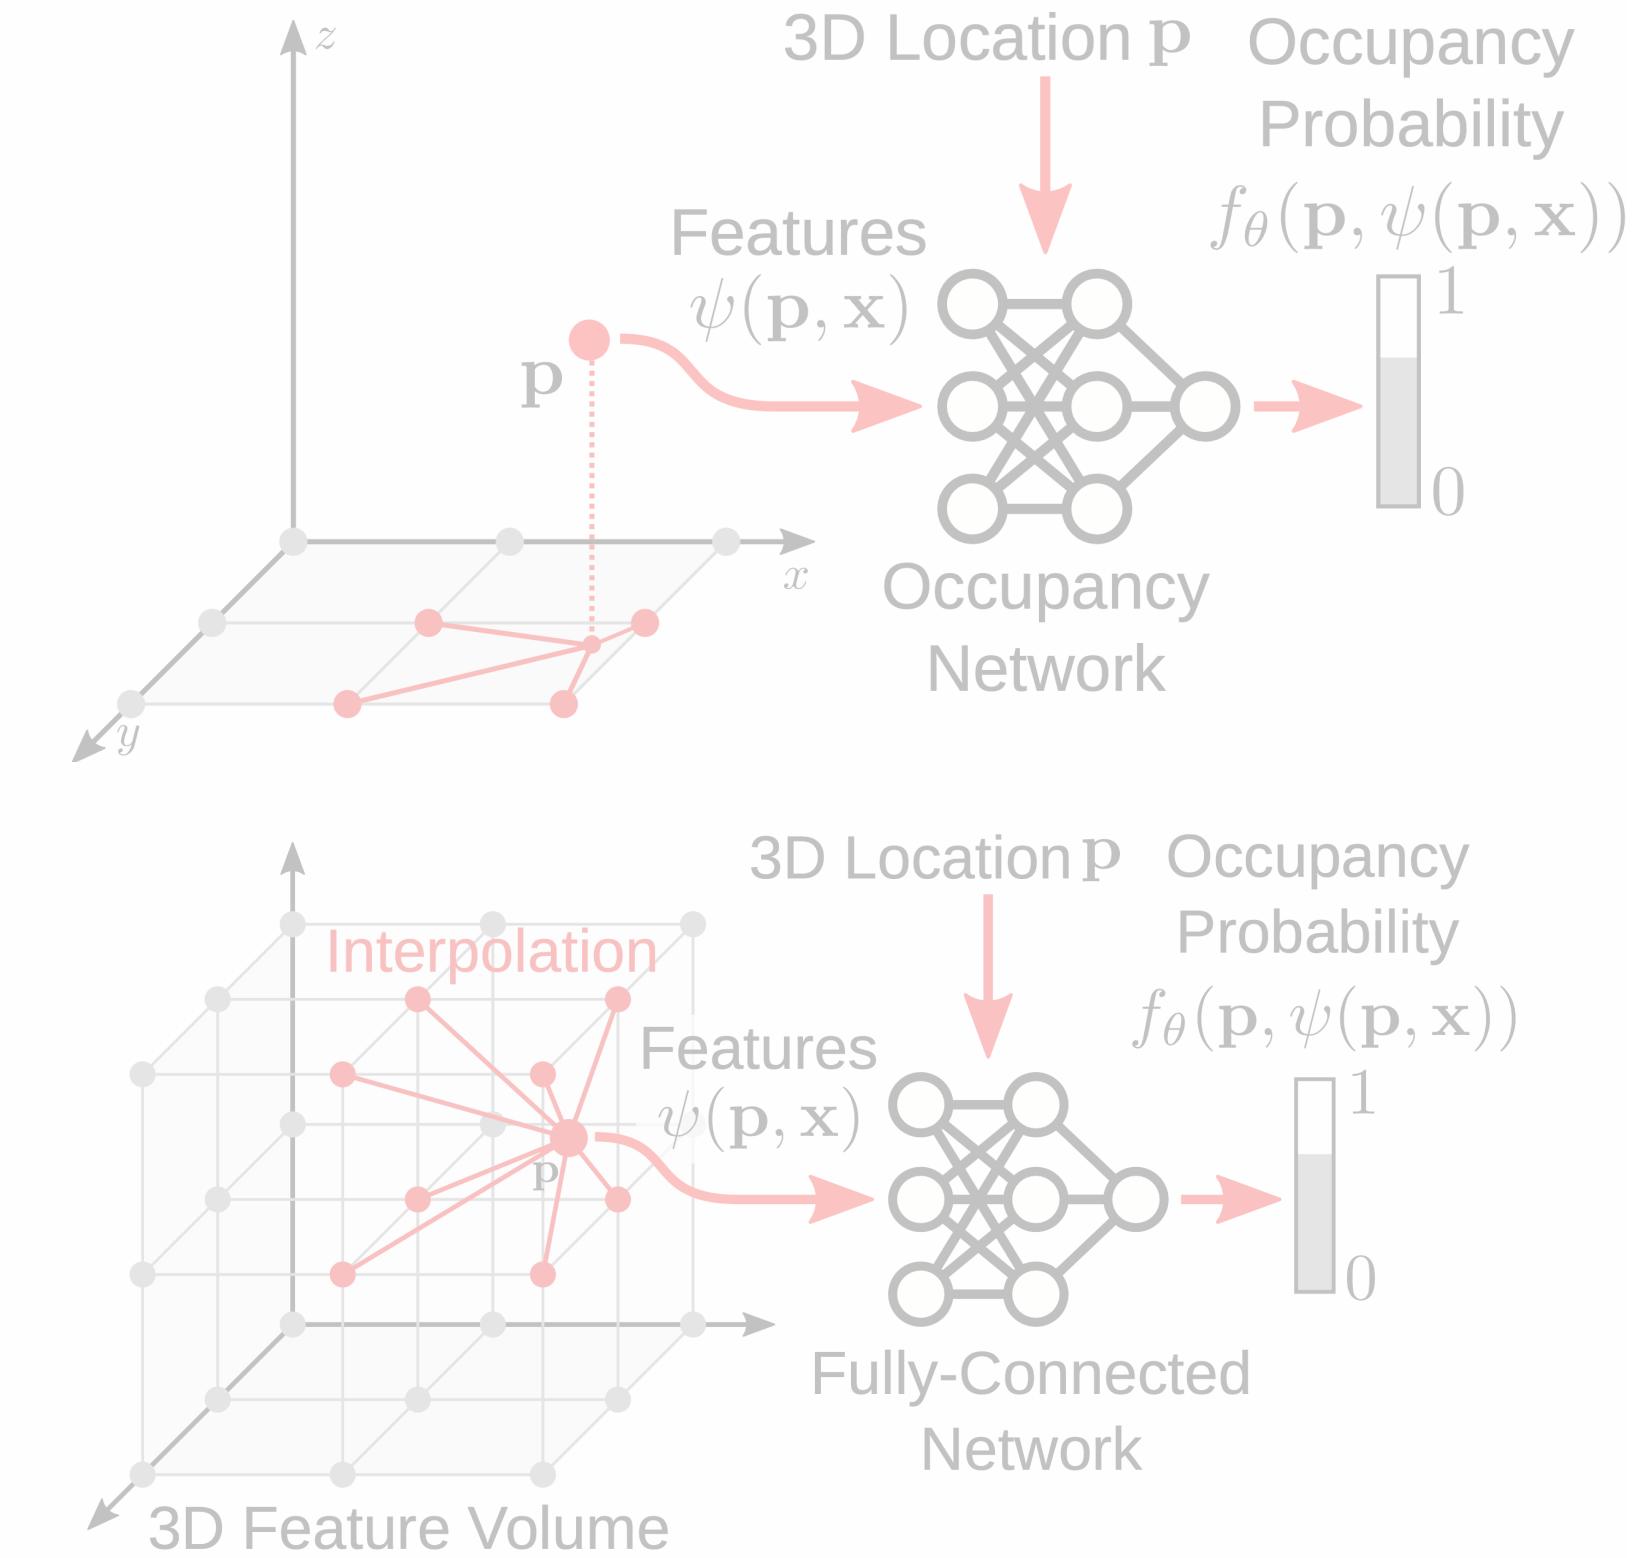
Neural Fields



Discrete Parameterizations

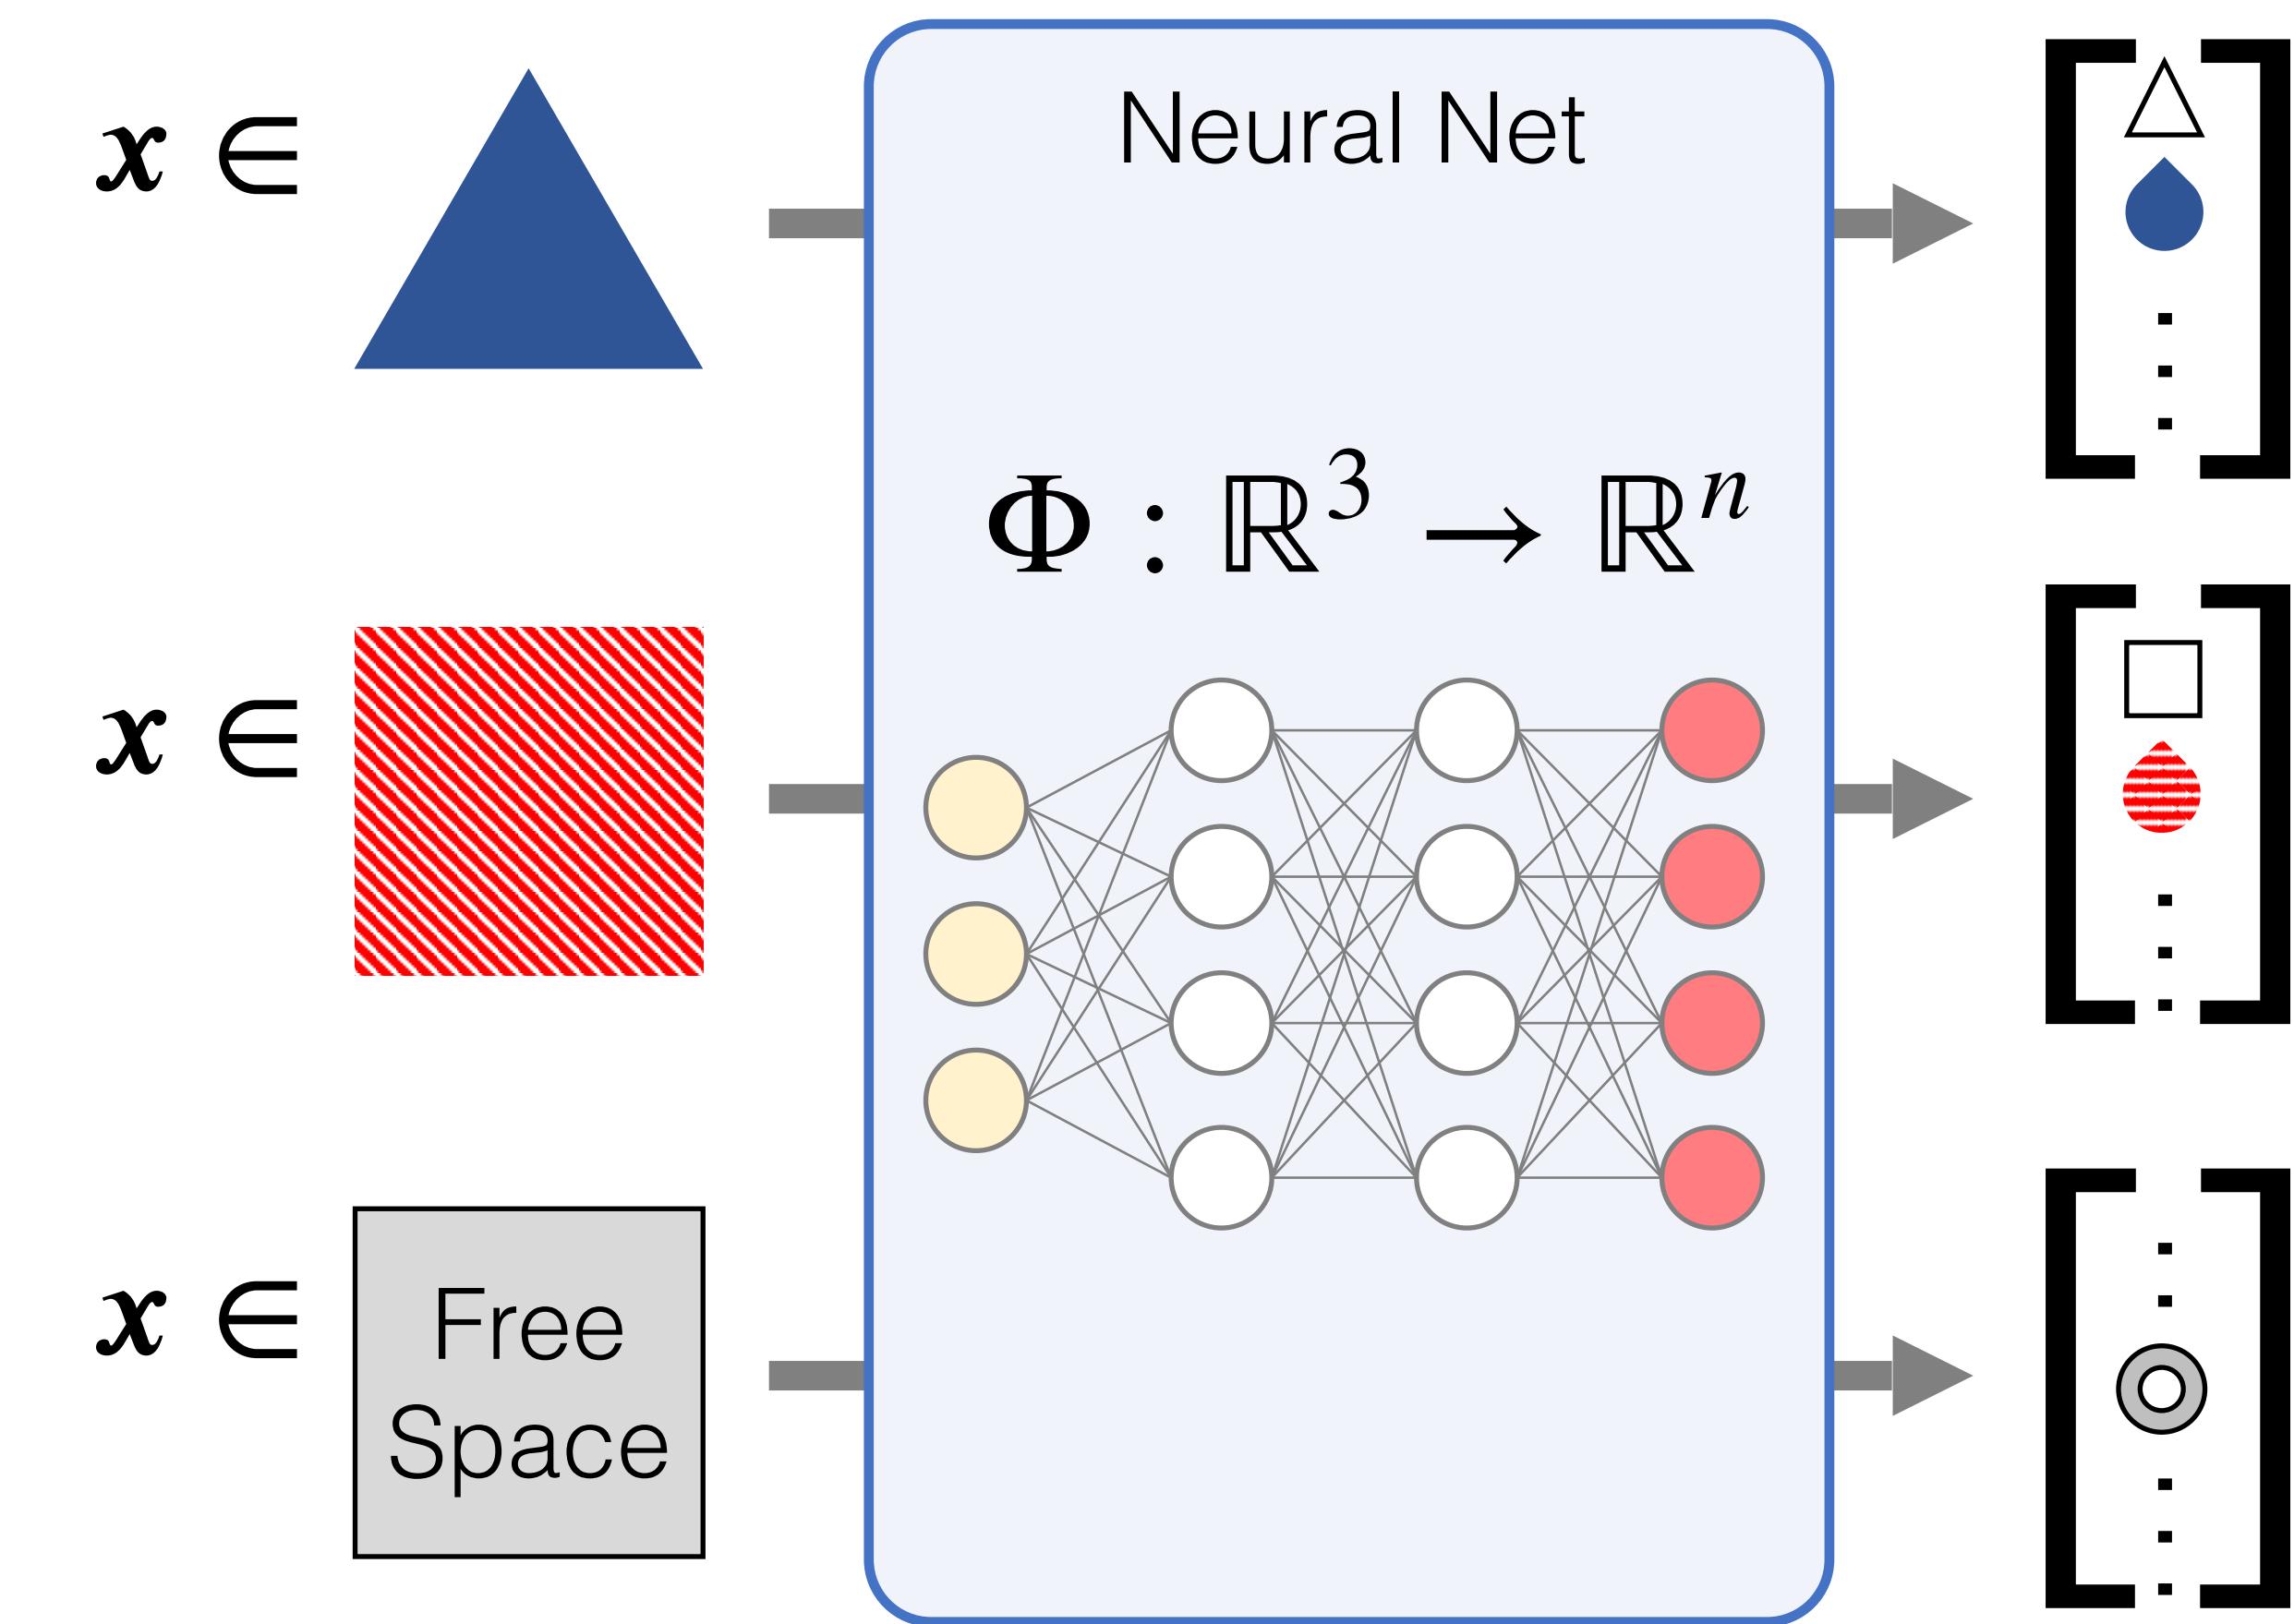
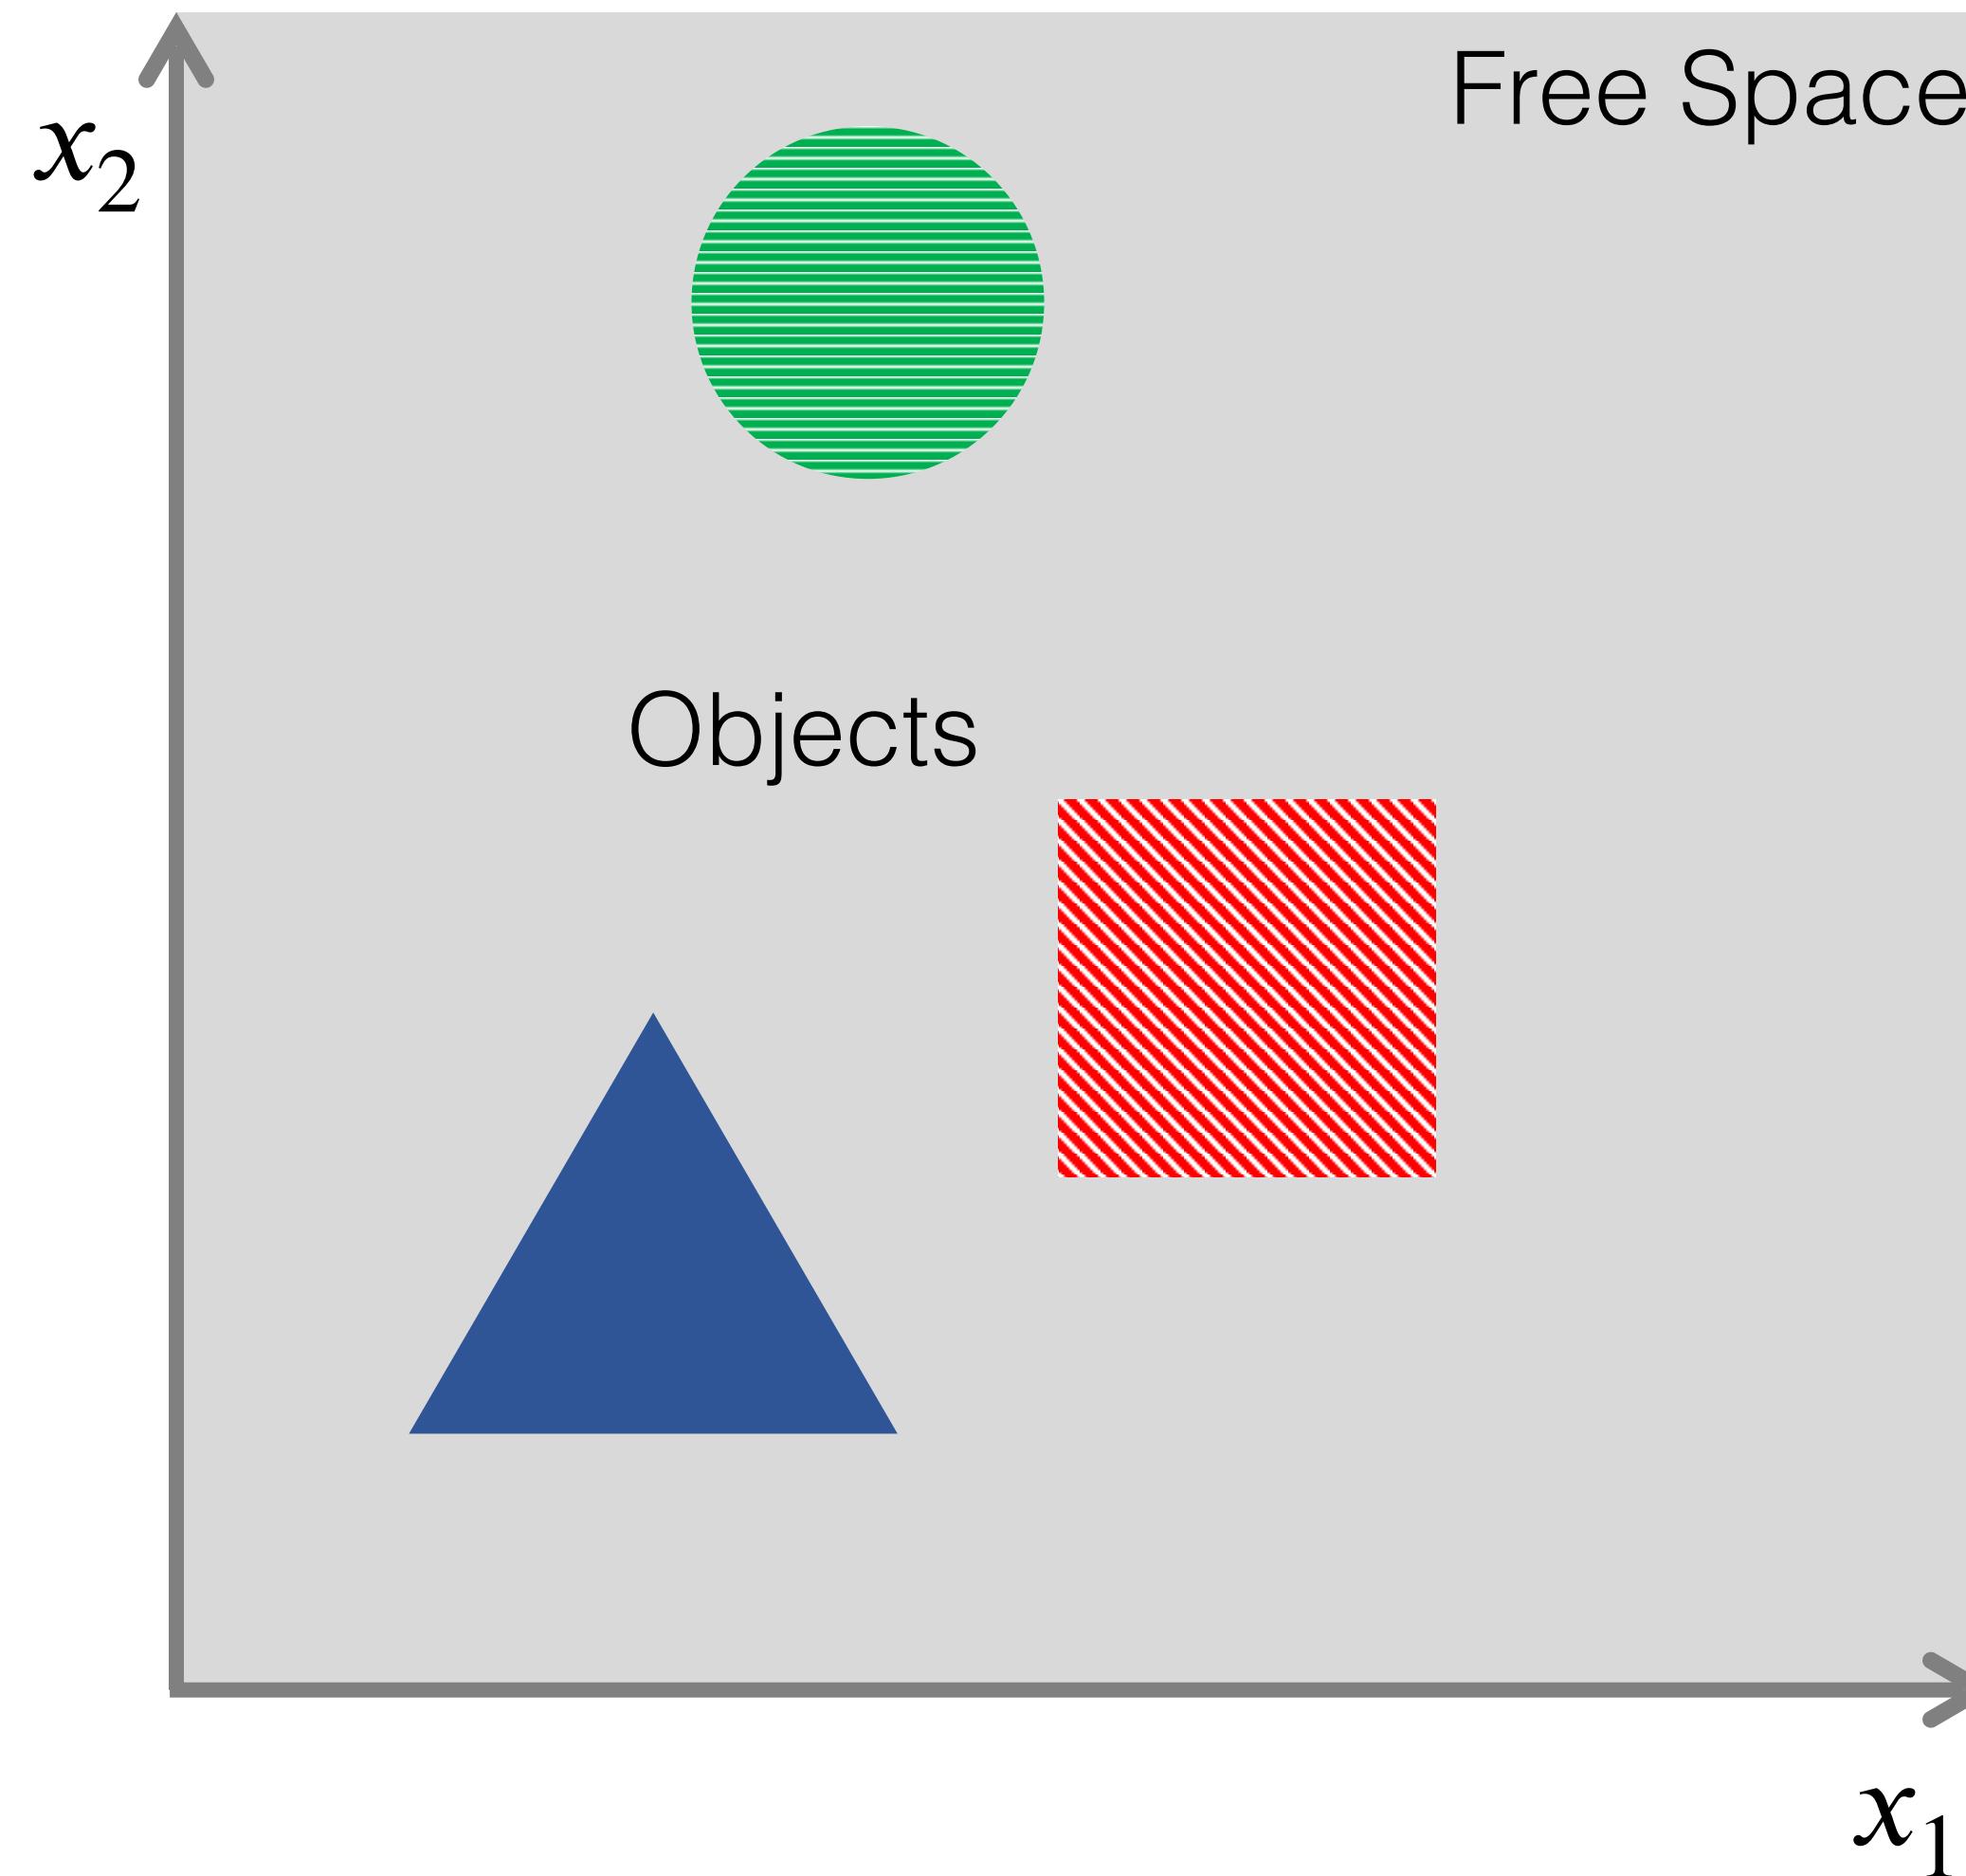


Neural Fields



Hybrid Parameterizations

Neural fields: Parameterize Field as Neural Network!



Occupancy Networks: Learning 3D Reconstruction in Function Space, Mescheder et al.

IM-Net: Learning Implicit Fields for Generative Shape Modeling, Chen et al.

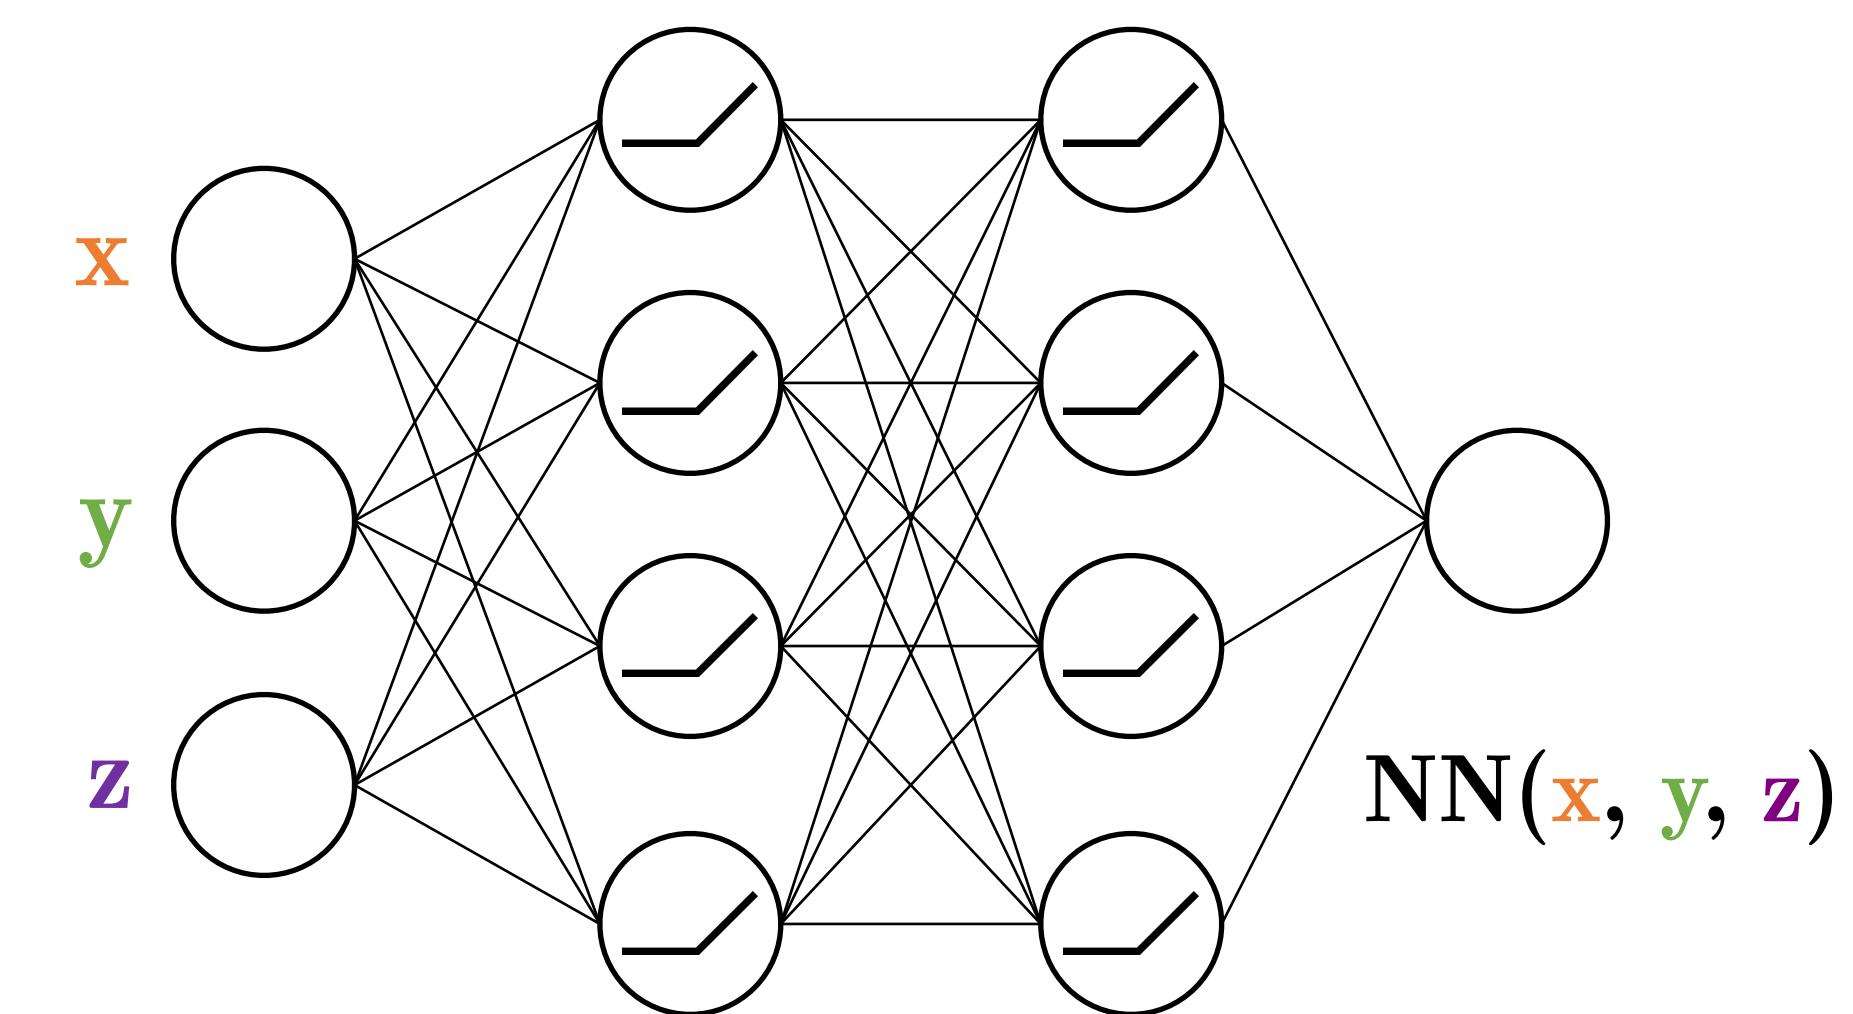
DeepSDF: Learning Continuous Signed Distance Functions for Shape Representation, Park et al. 2019

Sitzmann et al: Scene Representation Networks: Continuous 3D-Structure-Aware Neural Scene Representations, NeurIPS 2019.

Shapes



ReLU MLP

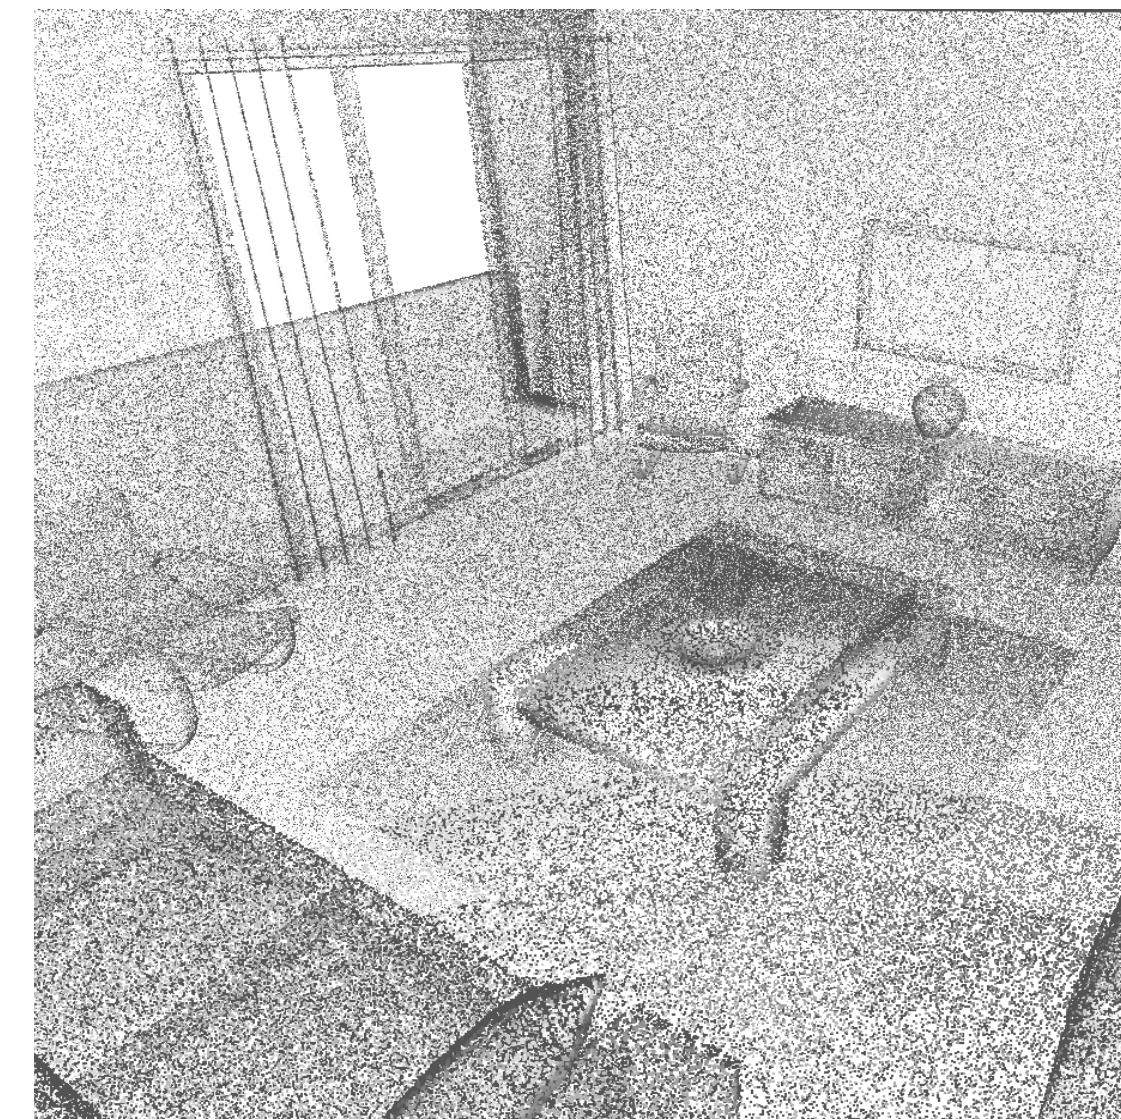


Mescheder et al. 2018
Park et al. 2018
Chen et al. 2018

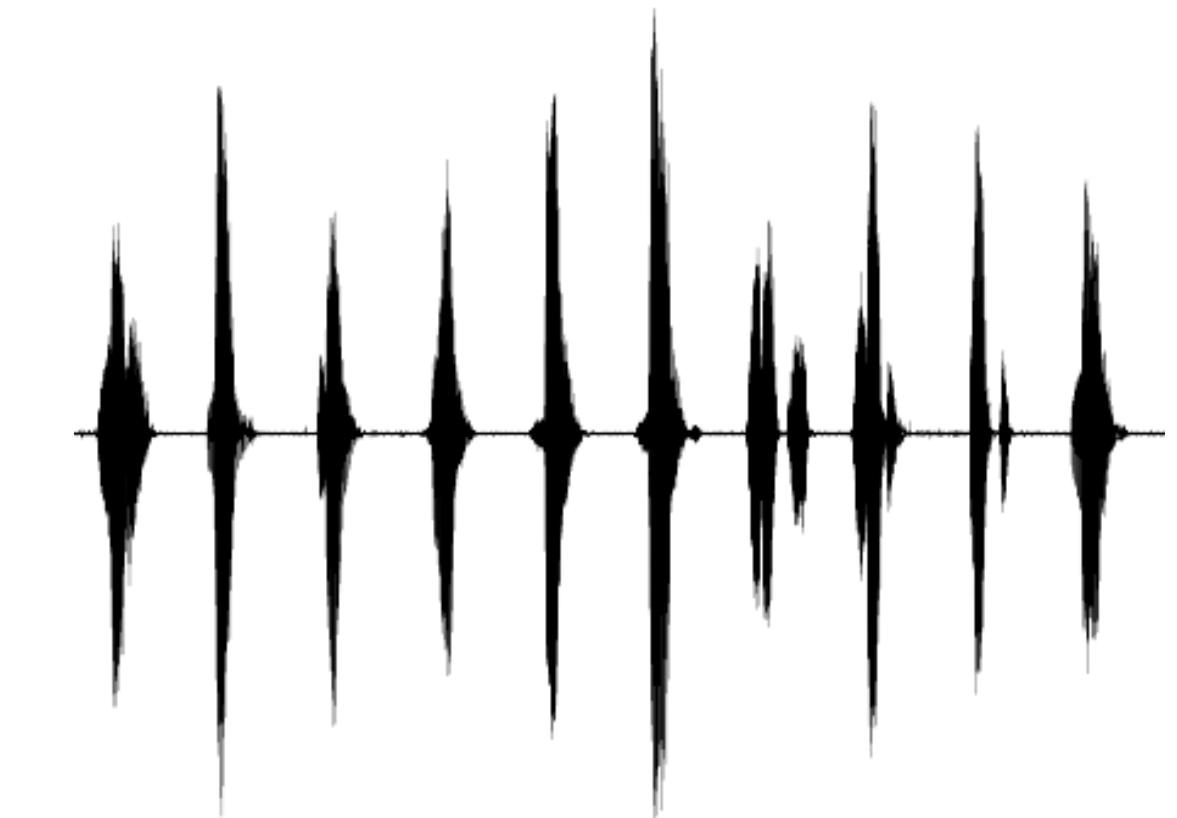
Images



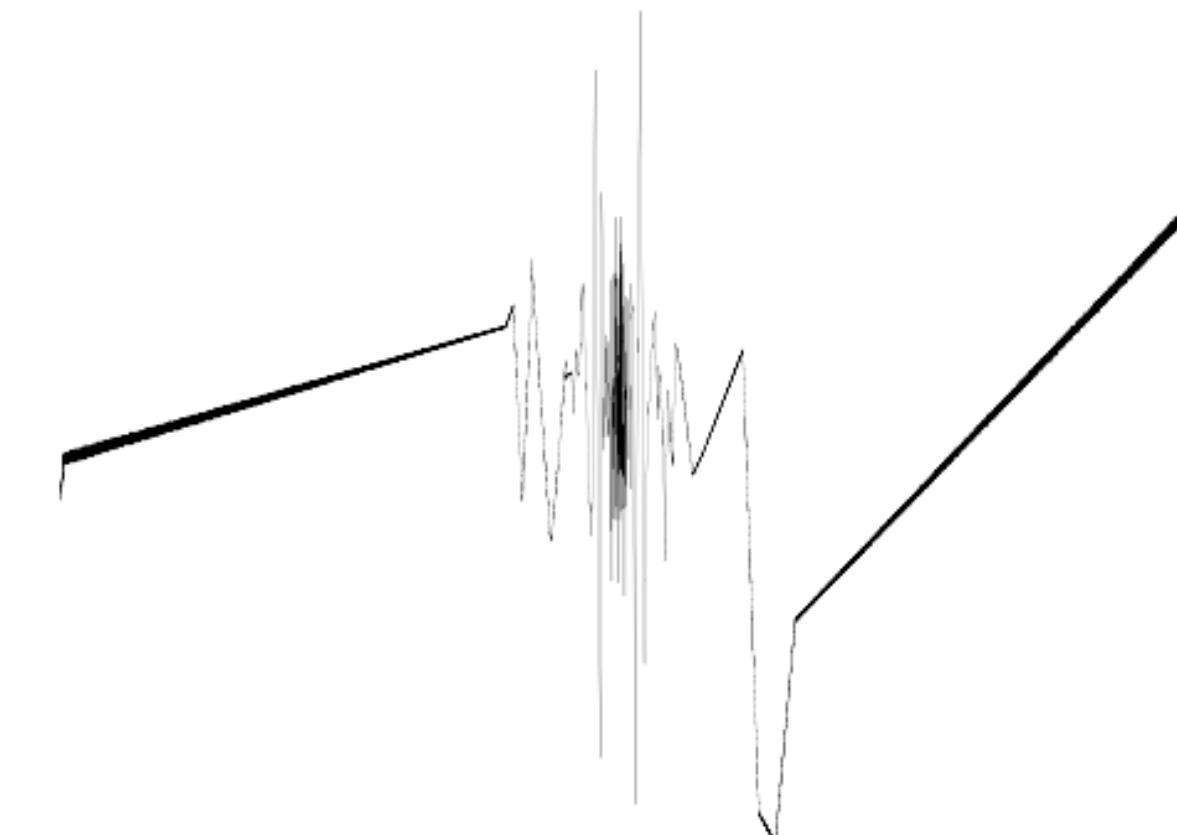
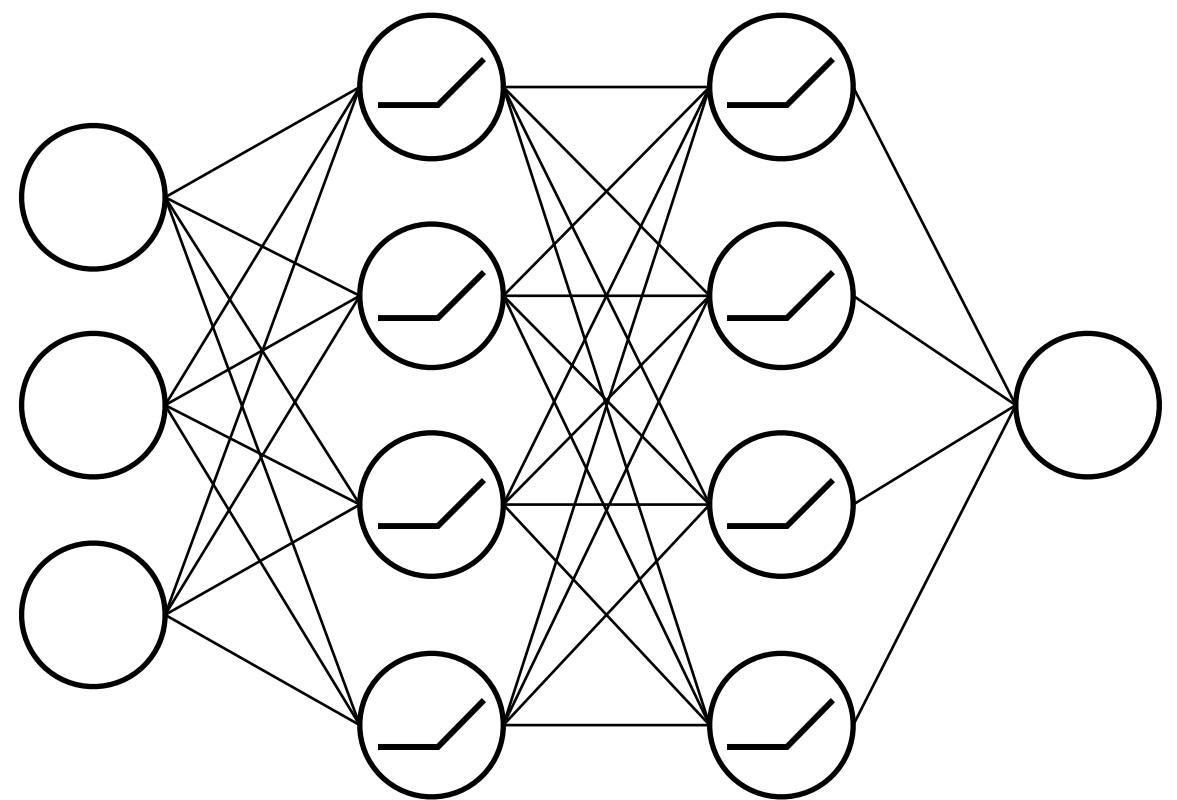
Shapes



Audio



ReLU MLP



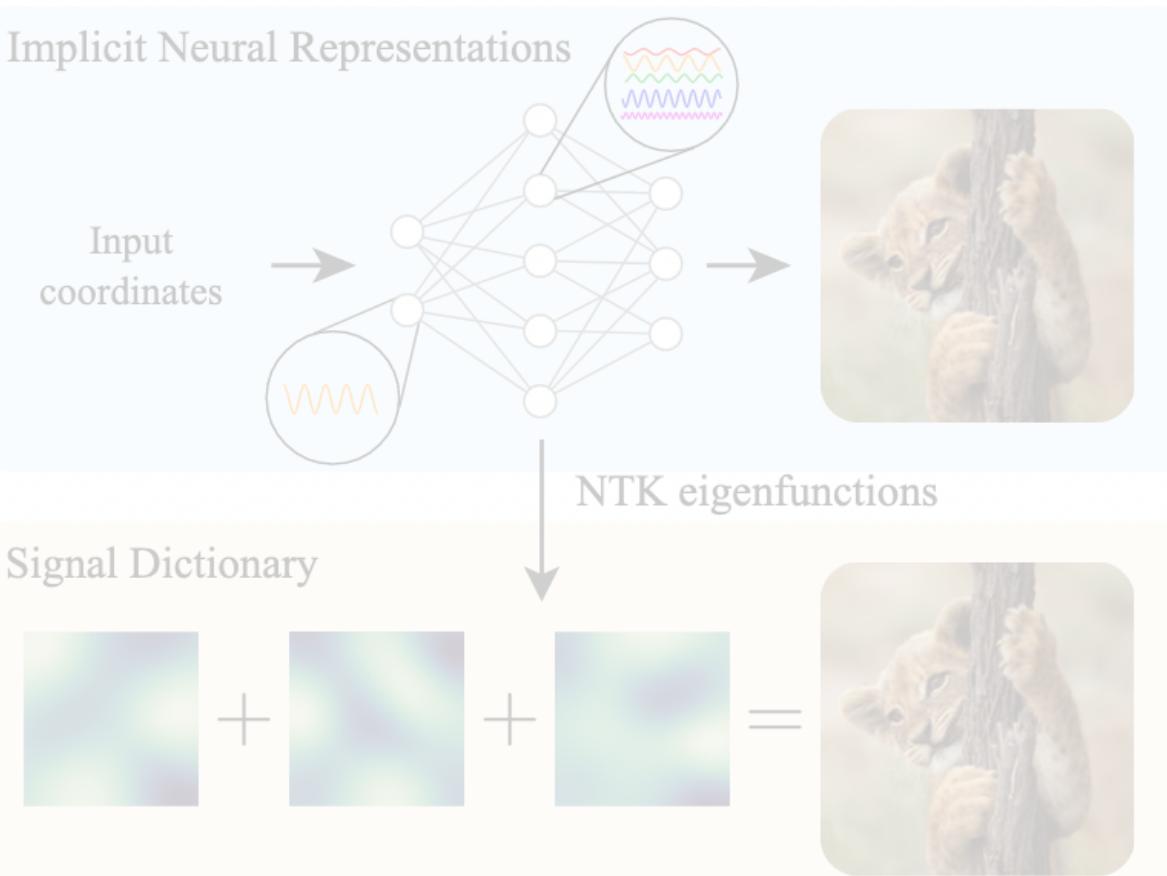
SIREN: Neural Implicit Representations With Periodic Activation Functions

Sitzmann & Martel et al. NeurIPS 2020



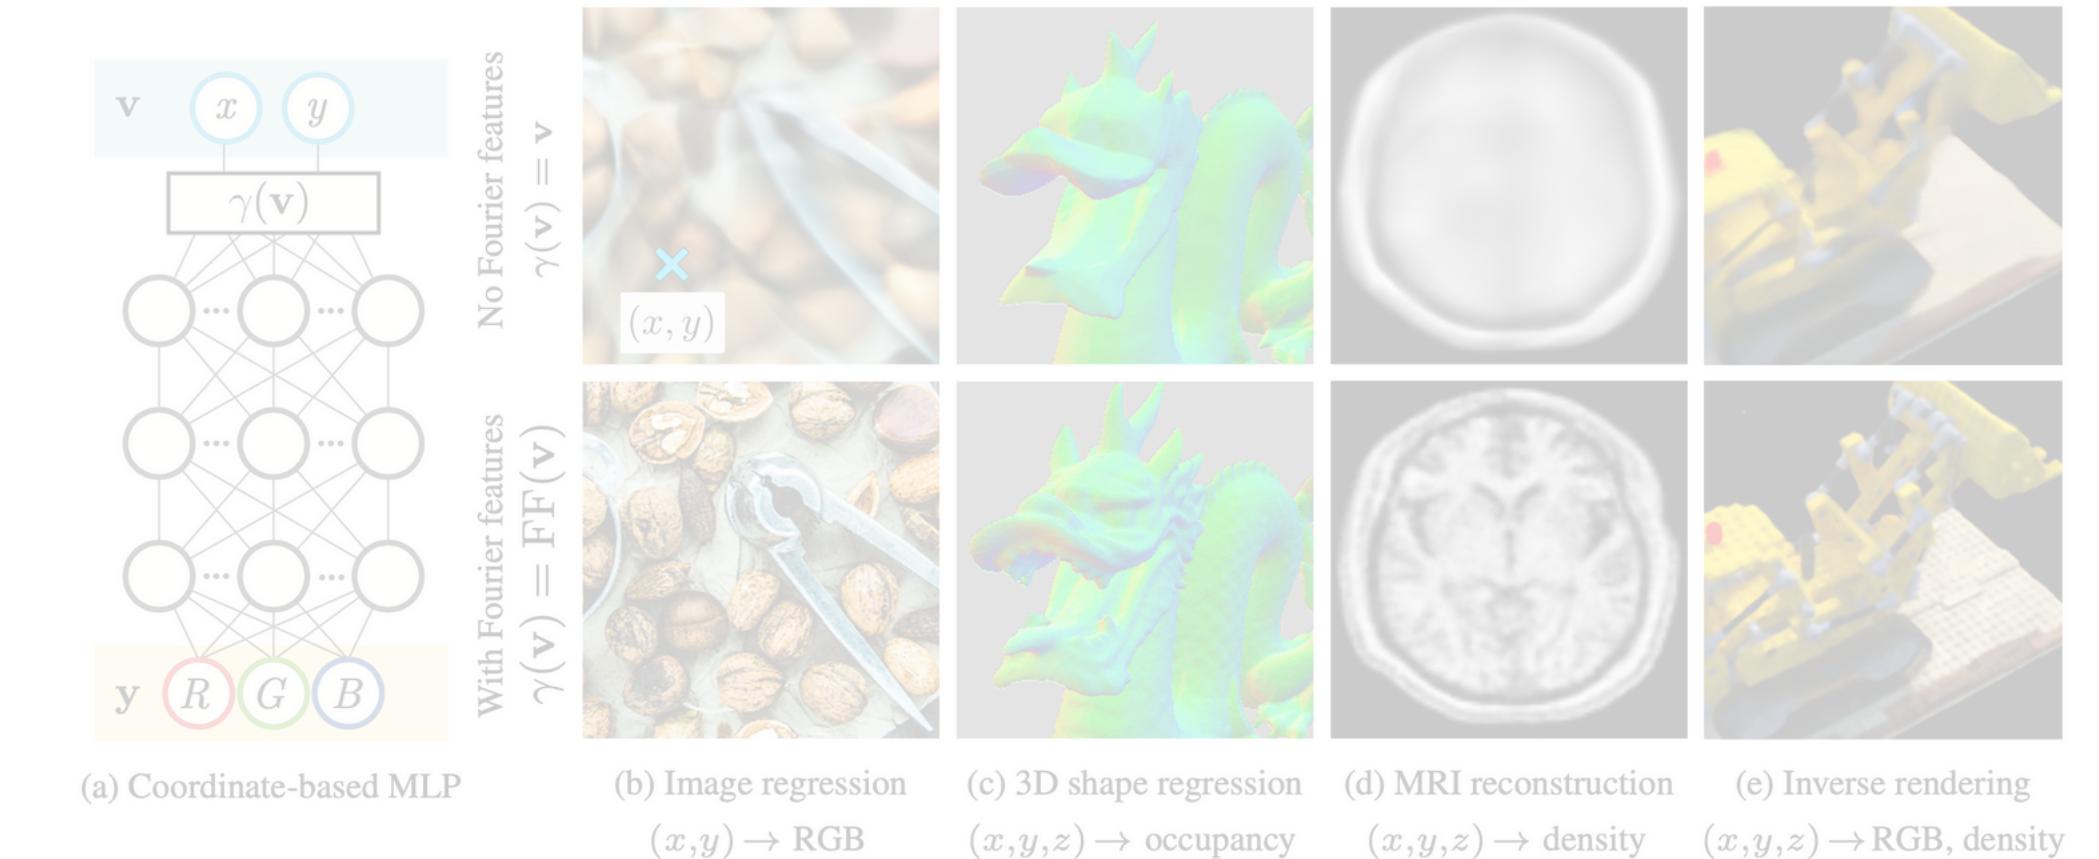
A Structured Dictionary Perspective on Implicit Neural Representations

Yüce & Ortiz-Jiménez et al. CVPR 2022



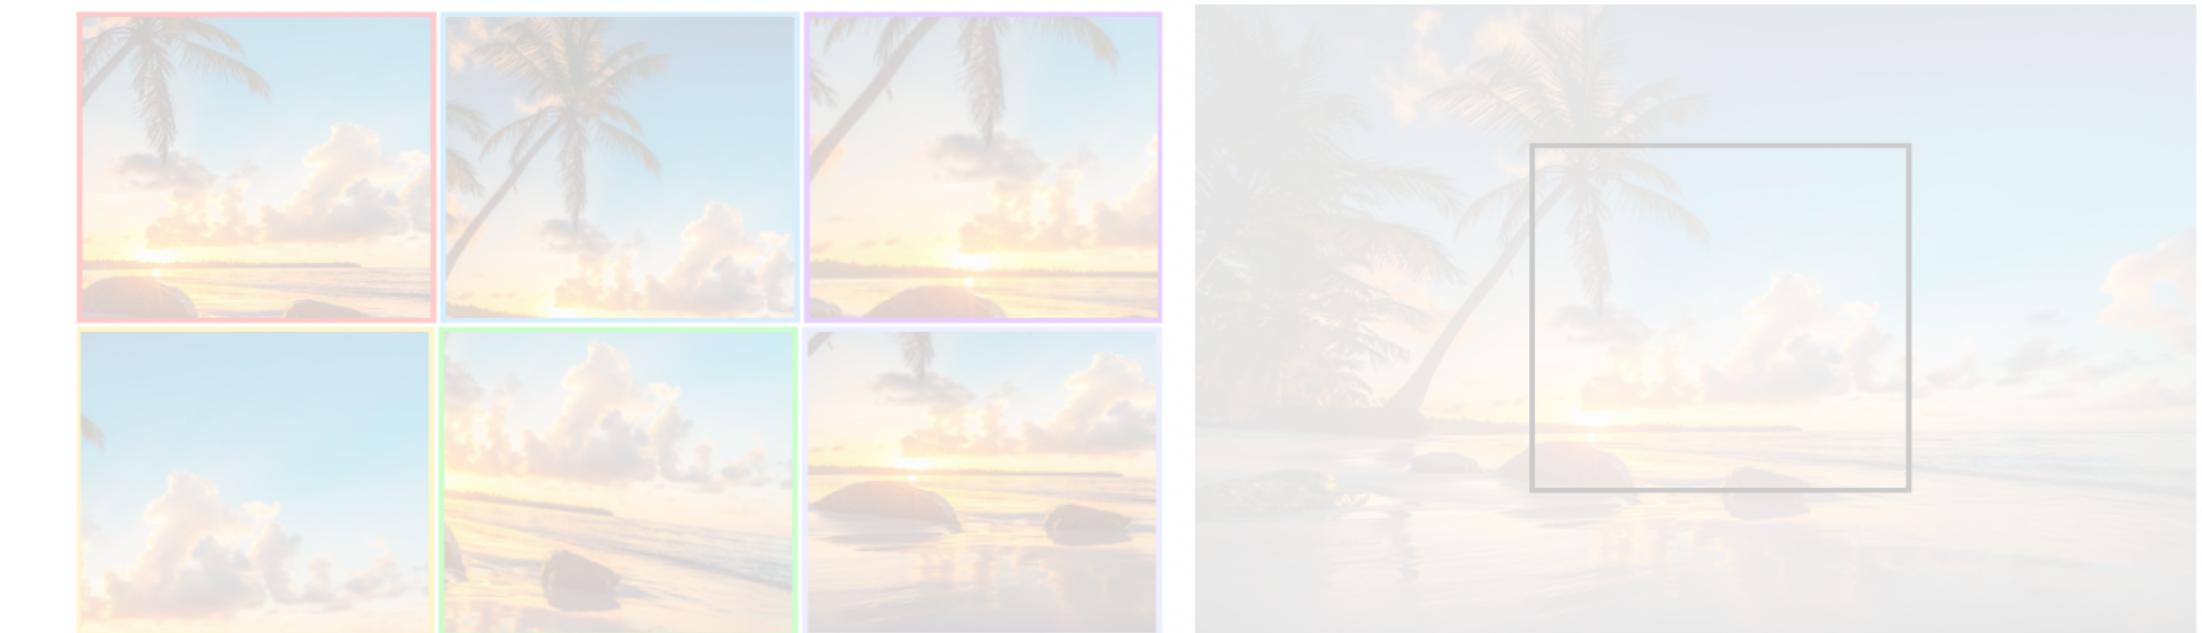
Fourier Features Let Networks Learn High Frequency Functions in Low Dimensional Domains

Tancik, Srinivasan, Mildenhall NeurIPS 2020

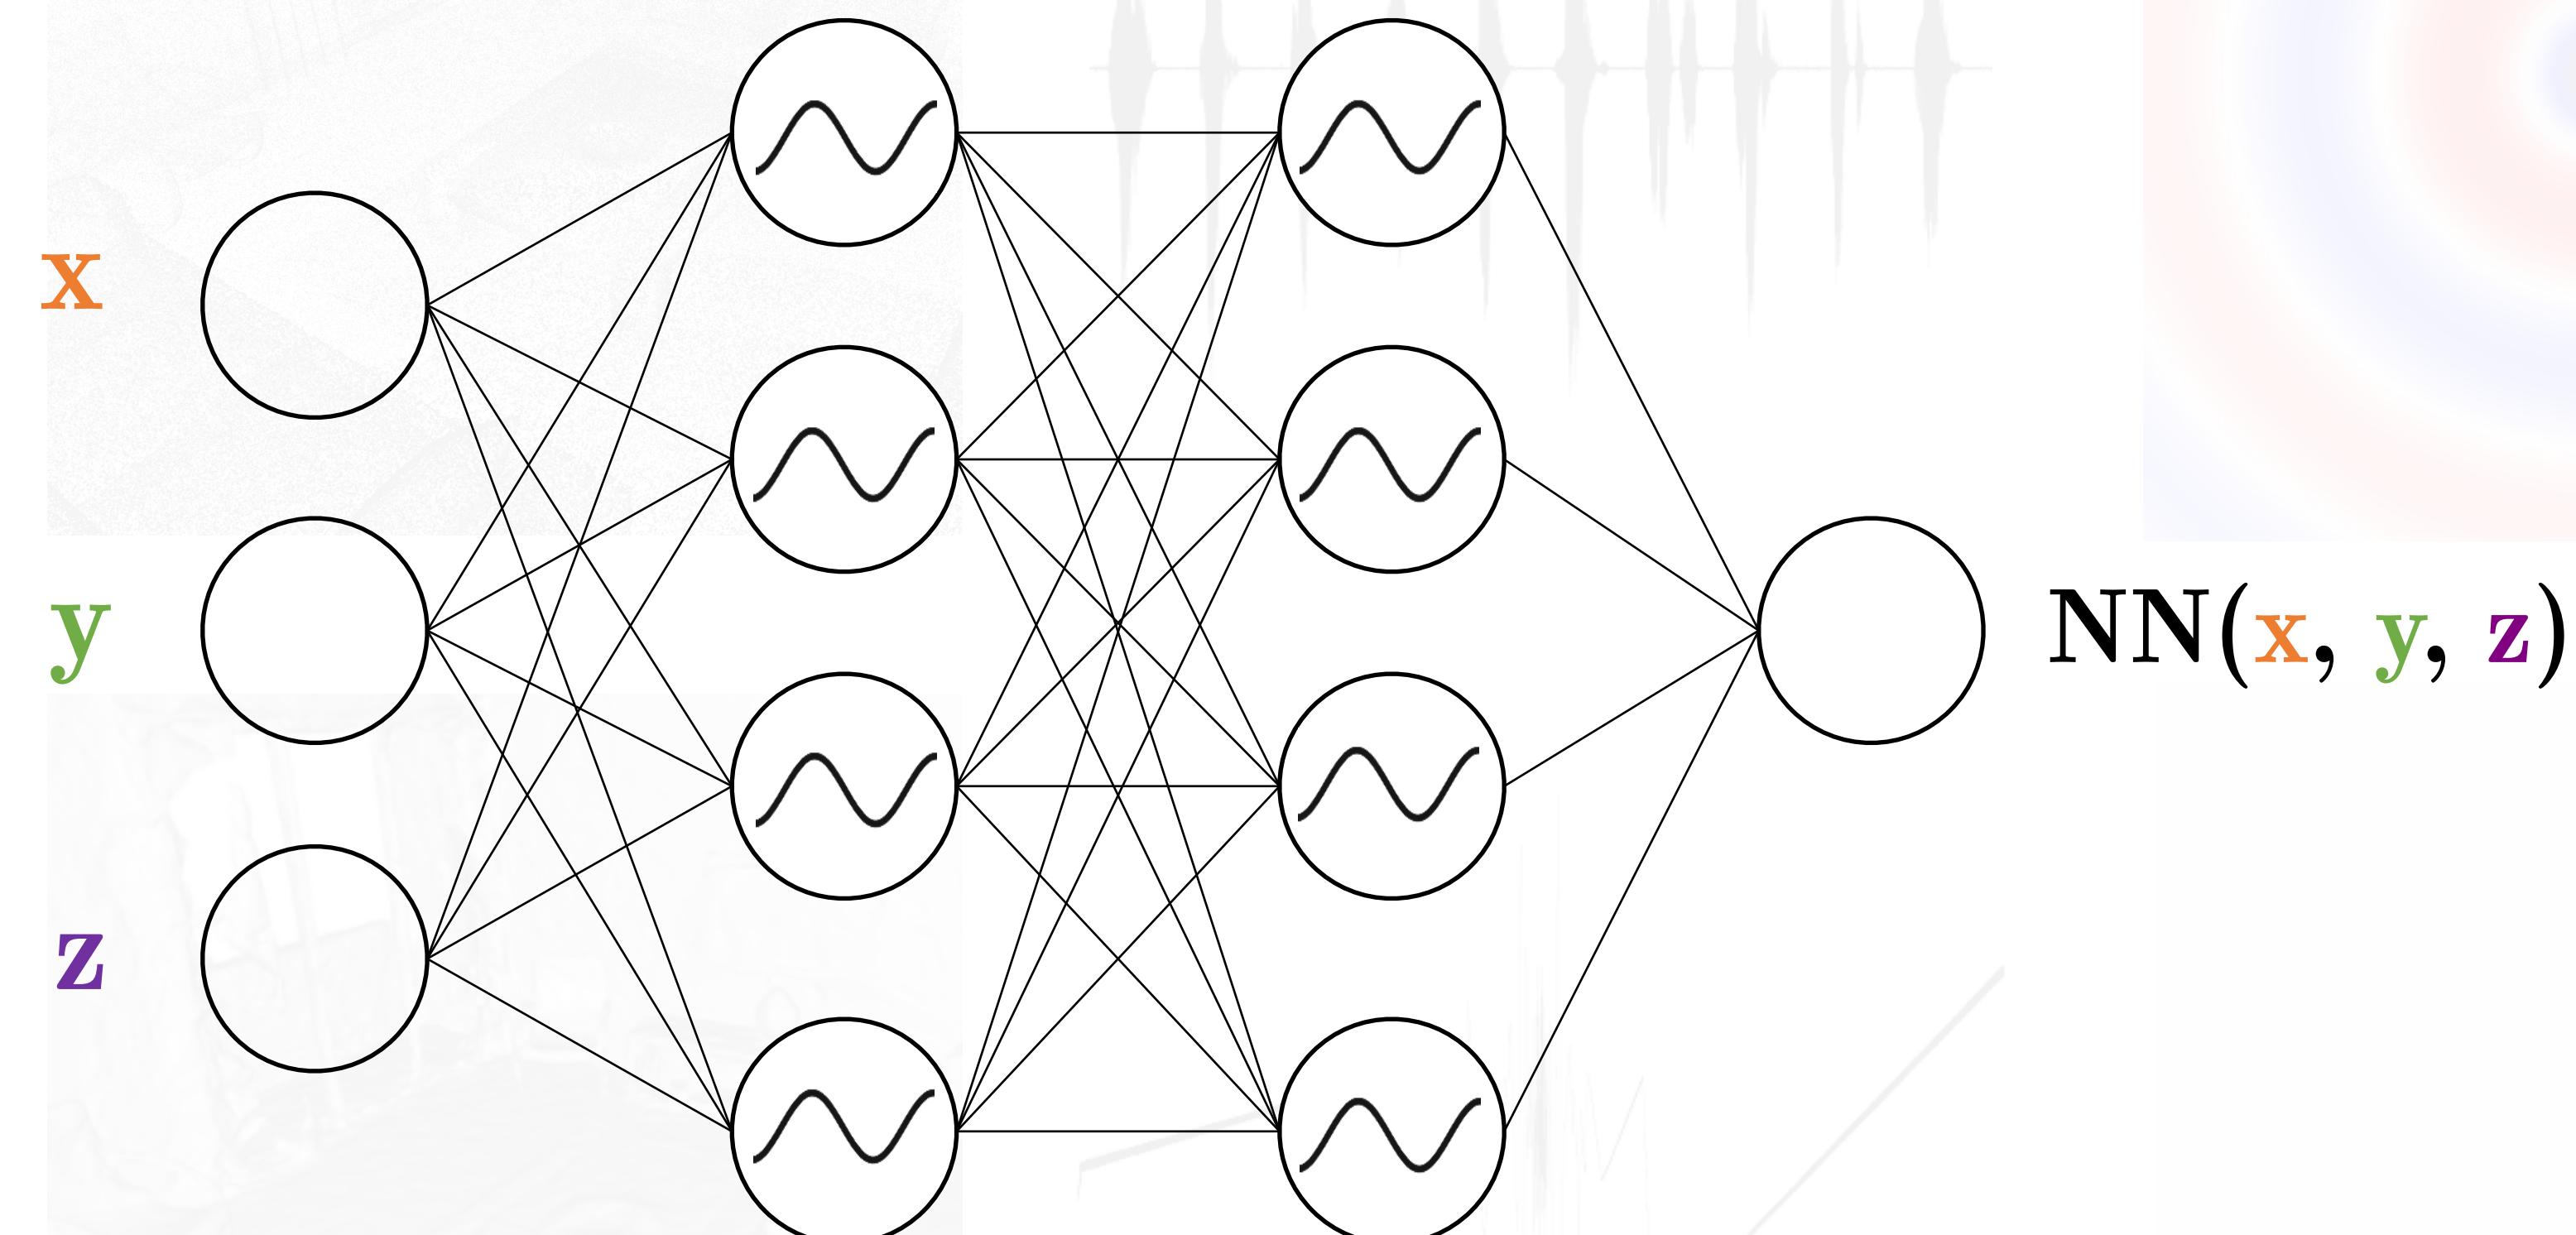


Gaussian Activated Neural Radiance Fields for High Fidelity Reconstruction & Pose Estimation

Chng et al. ECCV 2022



SIREN: Sinusoidal Representation Networks



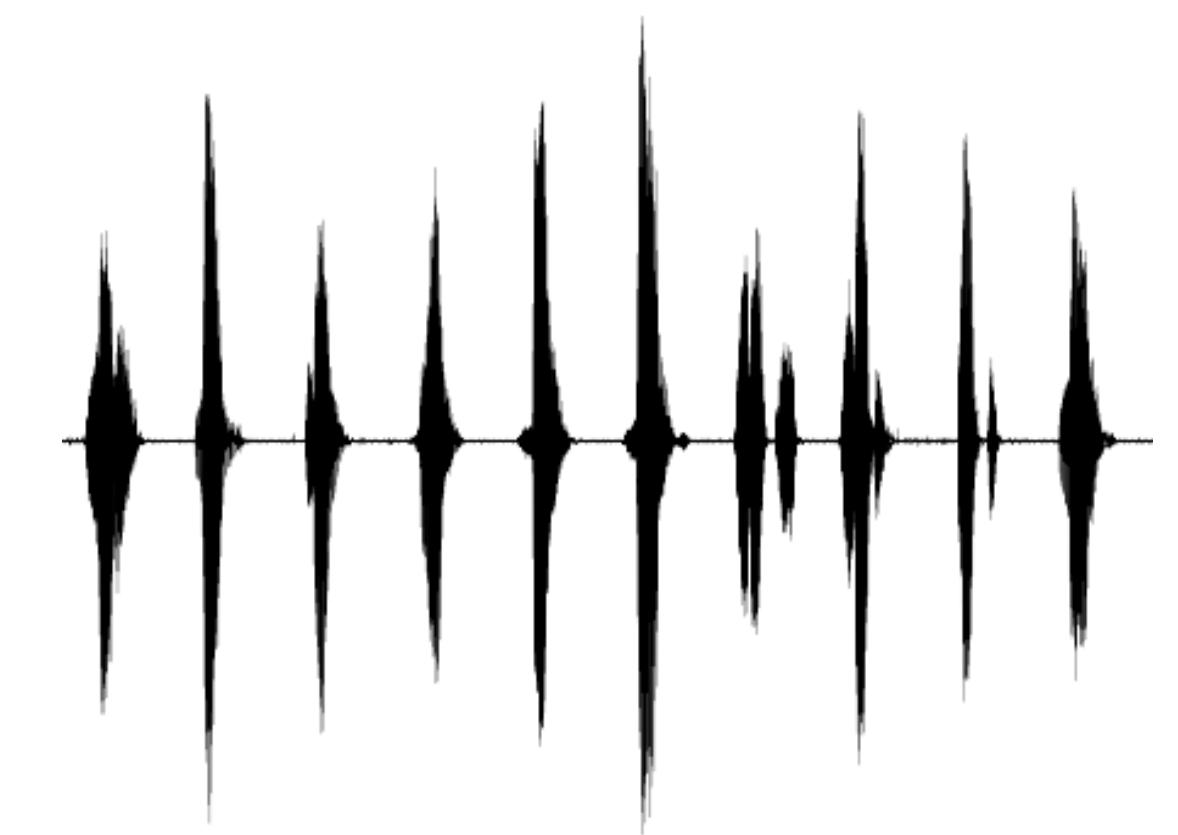
Images



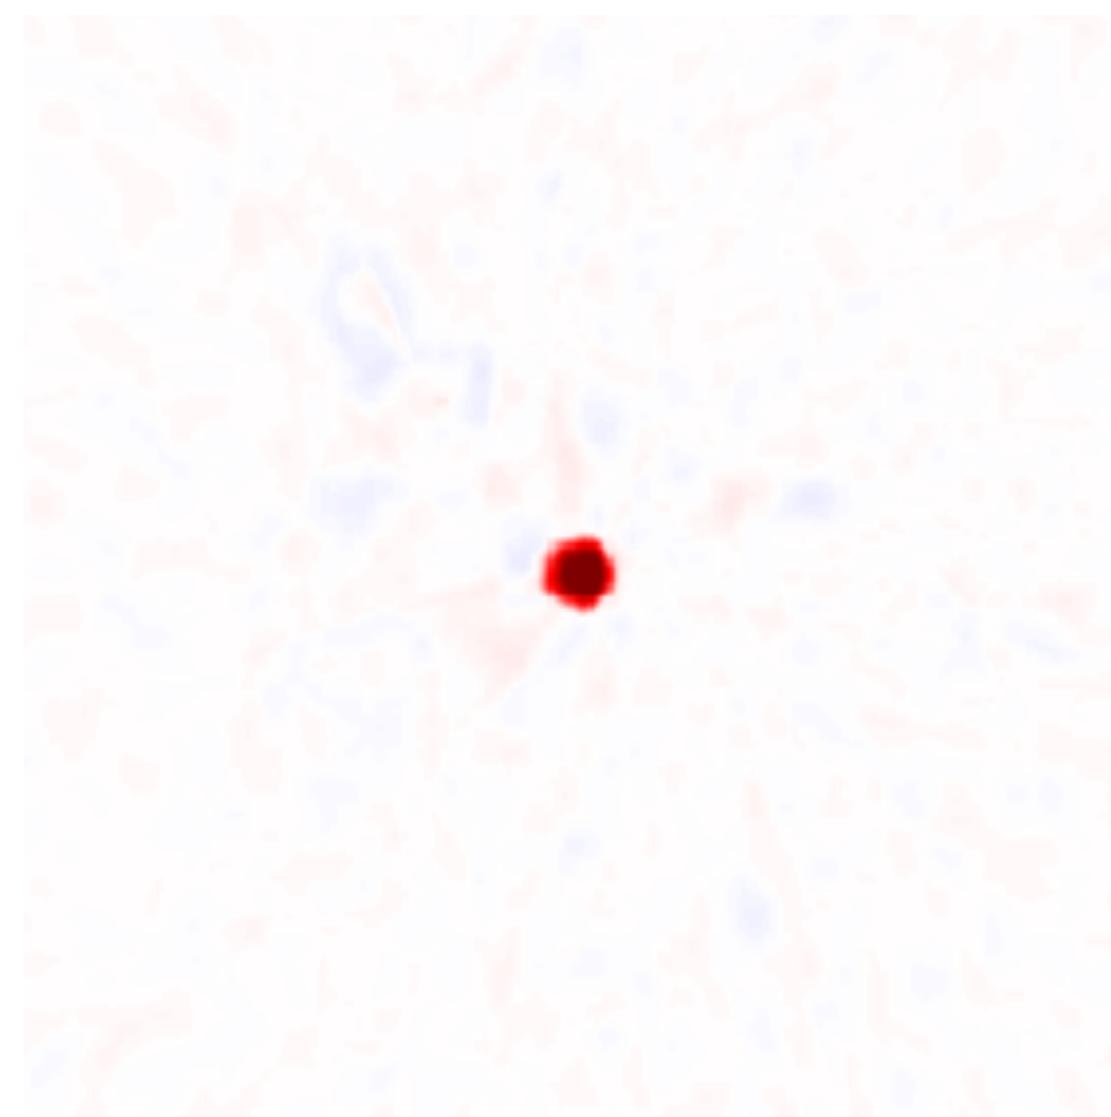
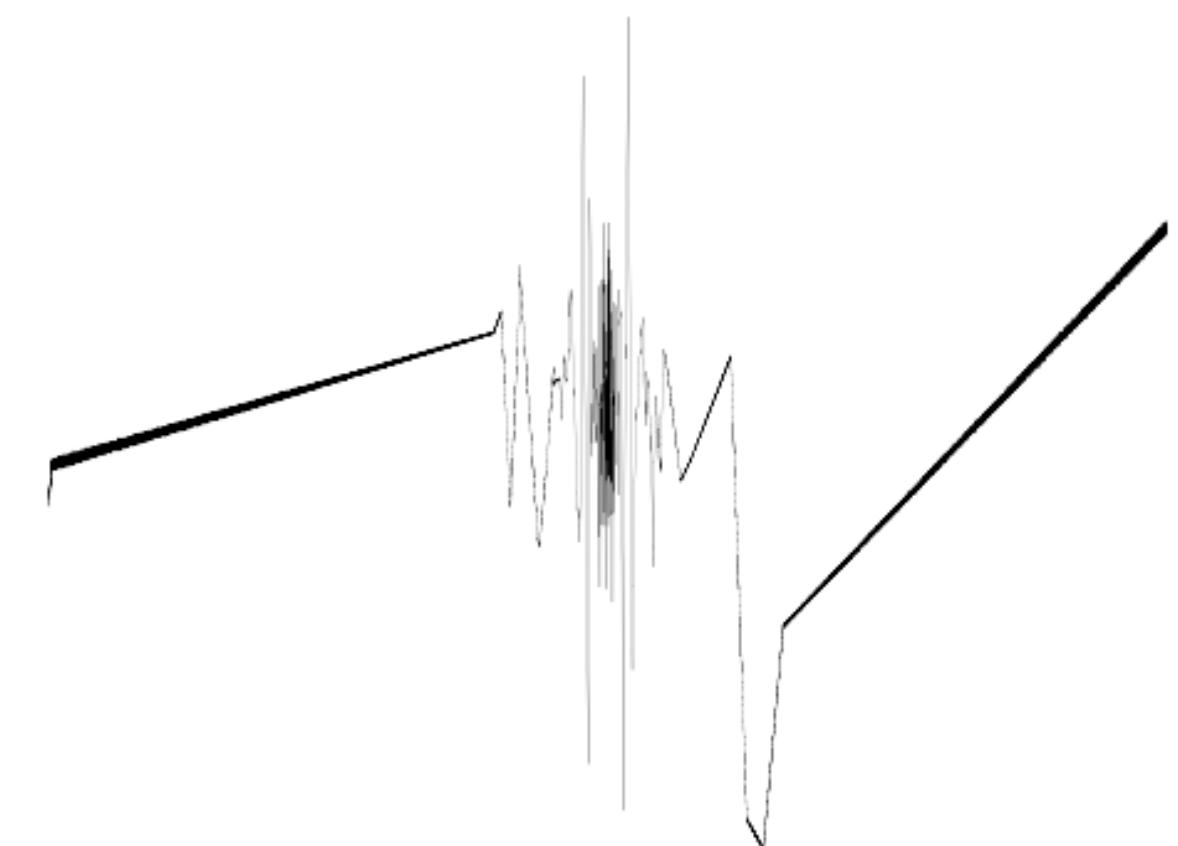
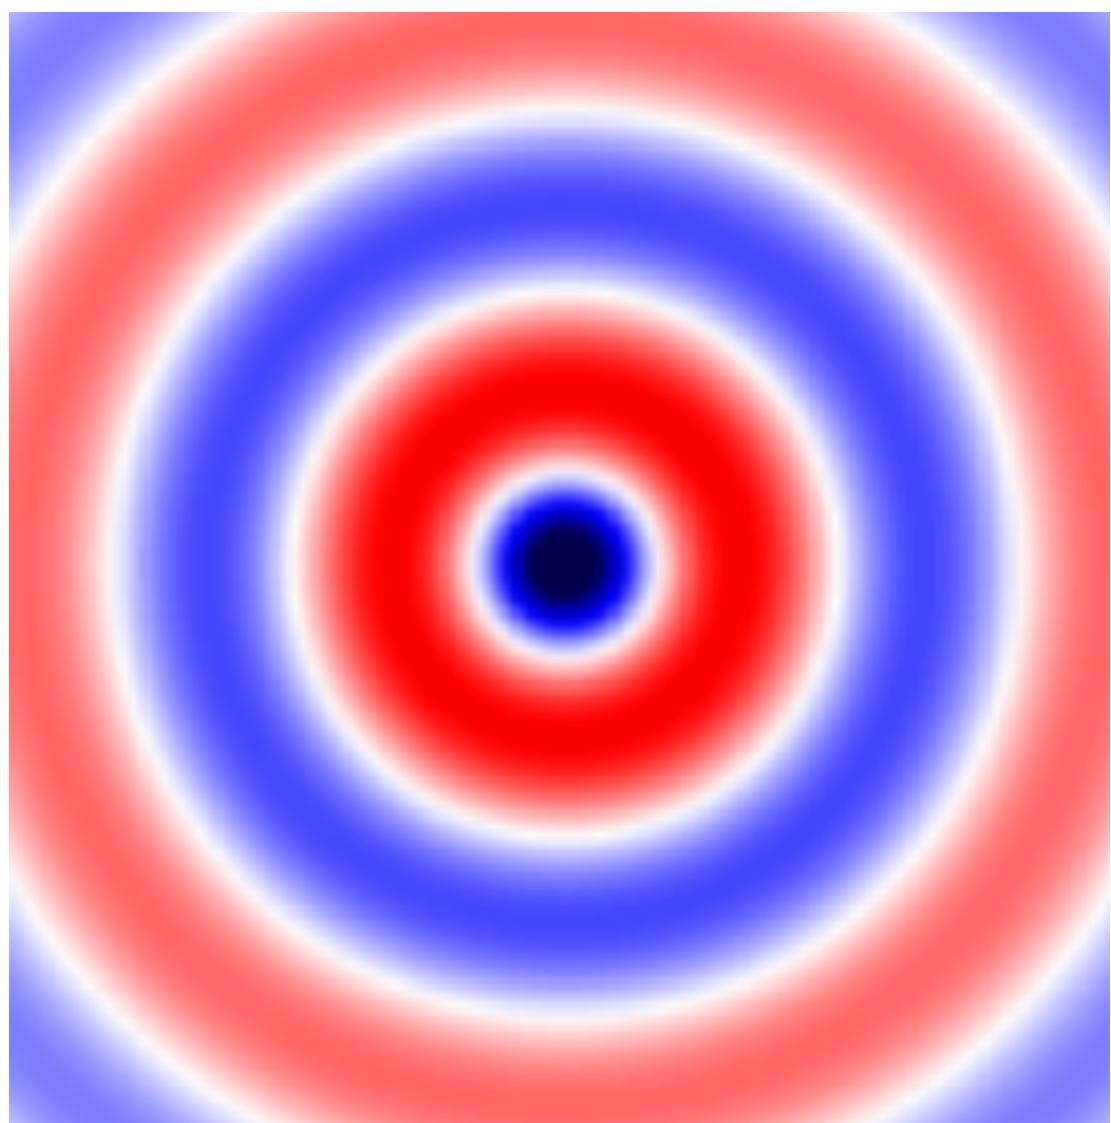
Shapes



Audio



Quantities defined by a differential equation



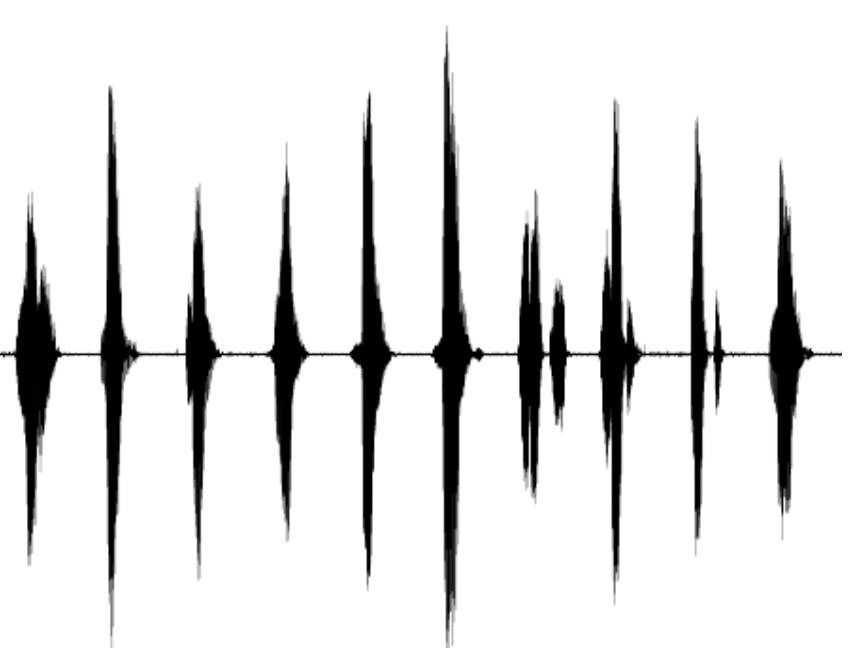
Images



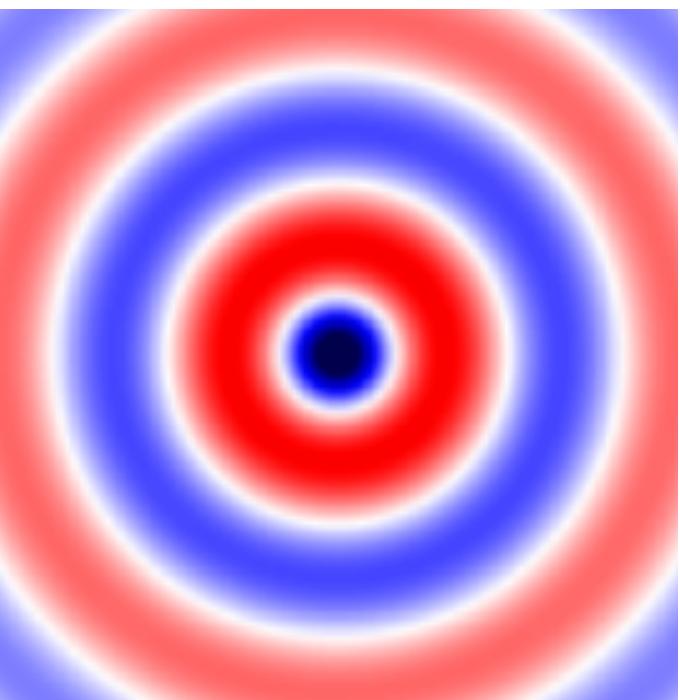
Shapes



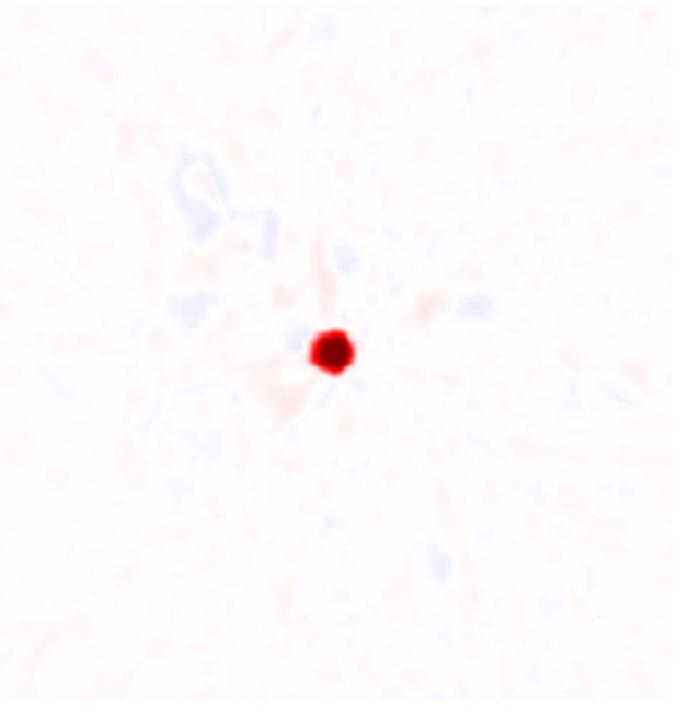
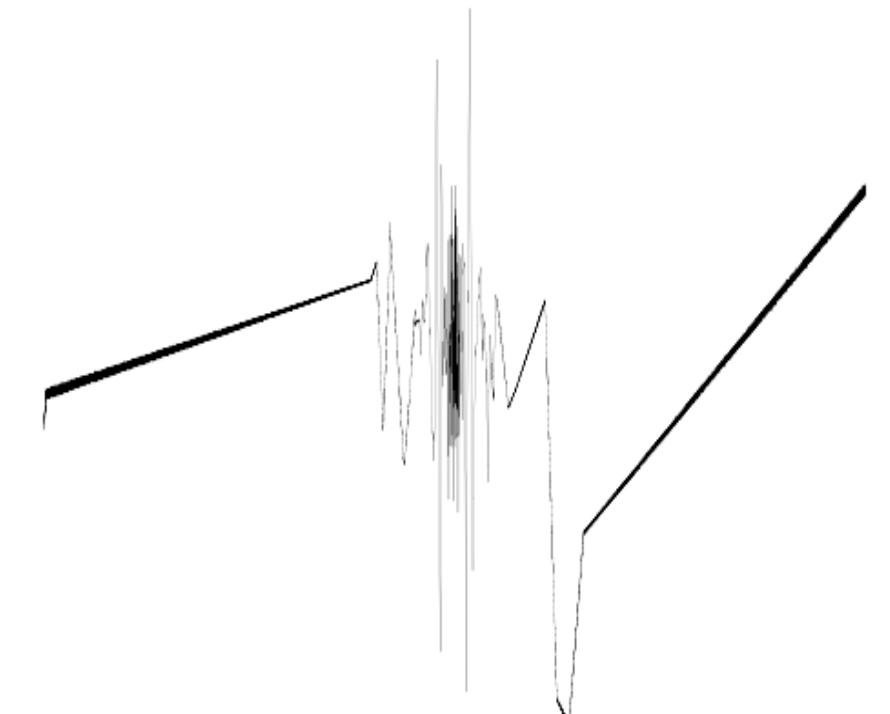
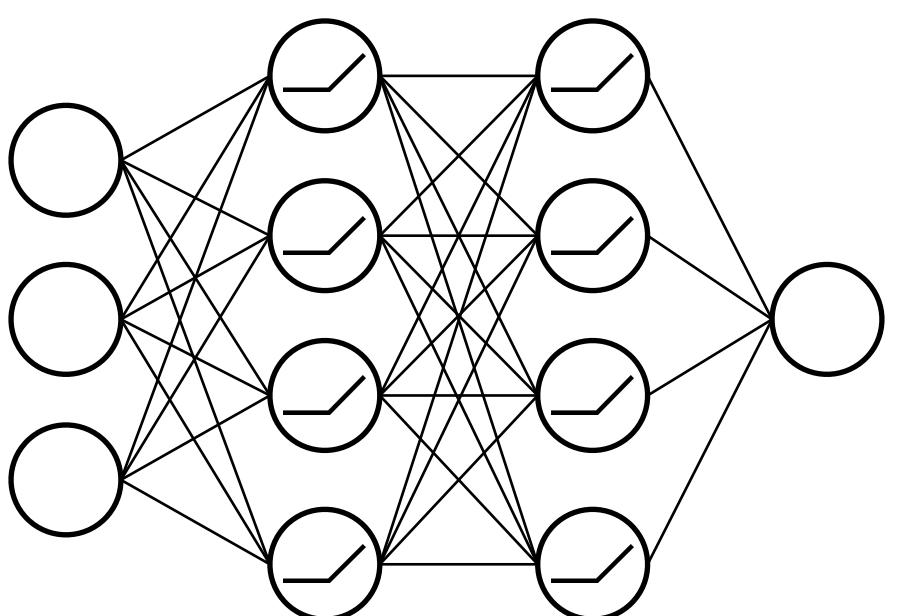
Audio



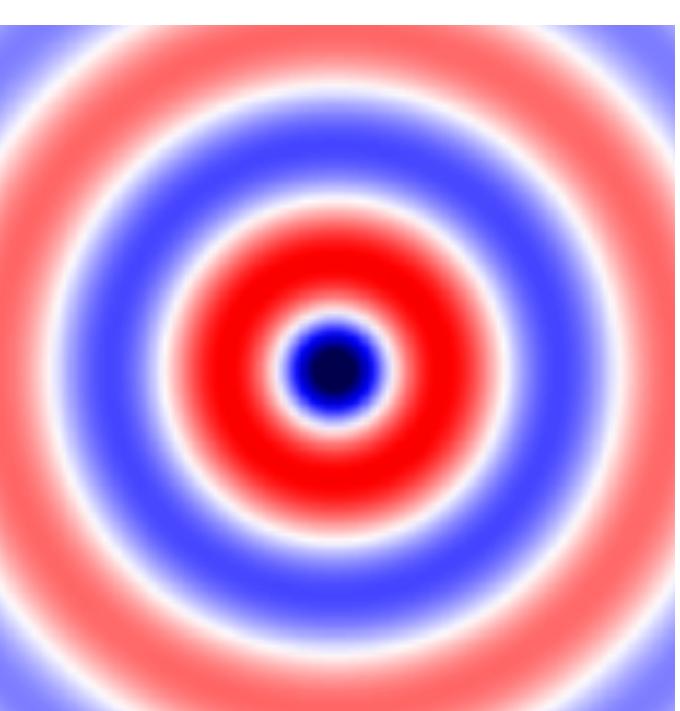
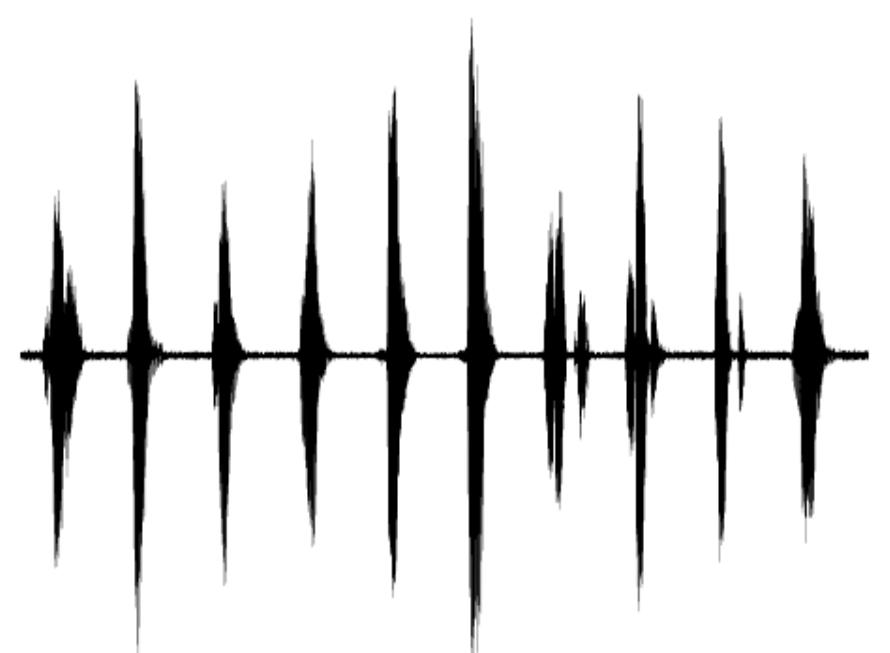
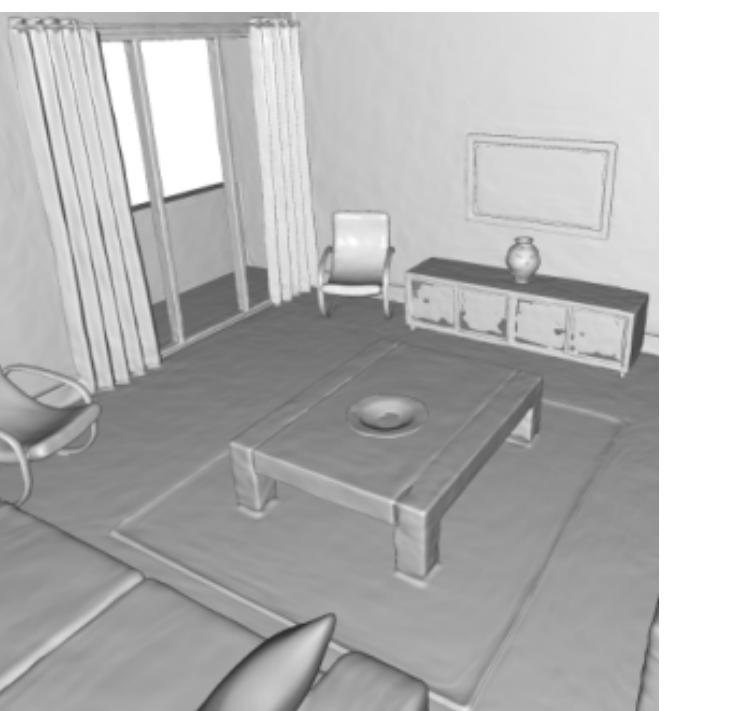
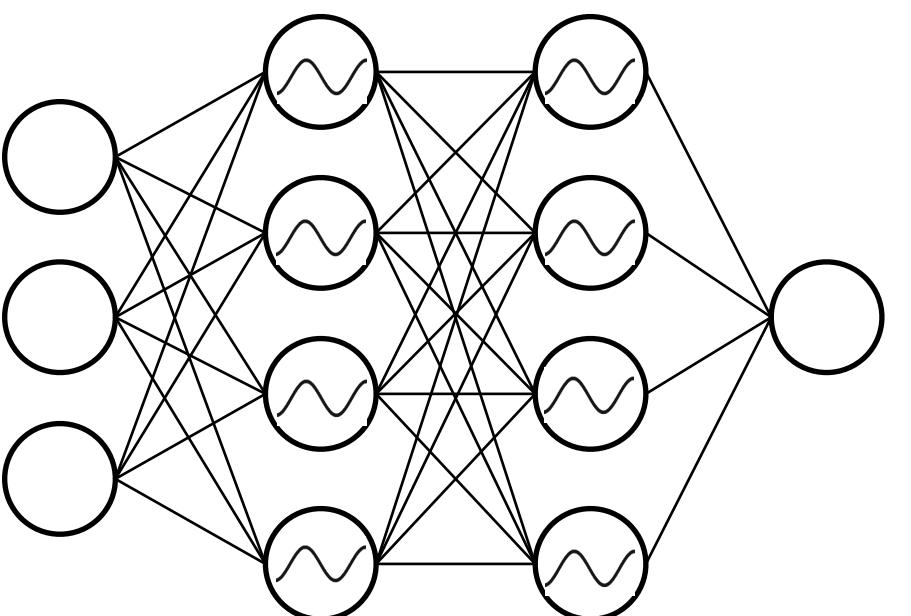
Quantities defined by a differential equation



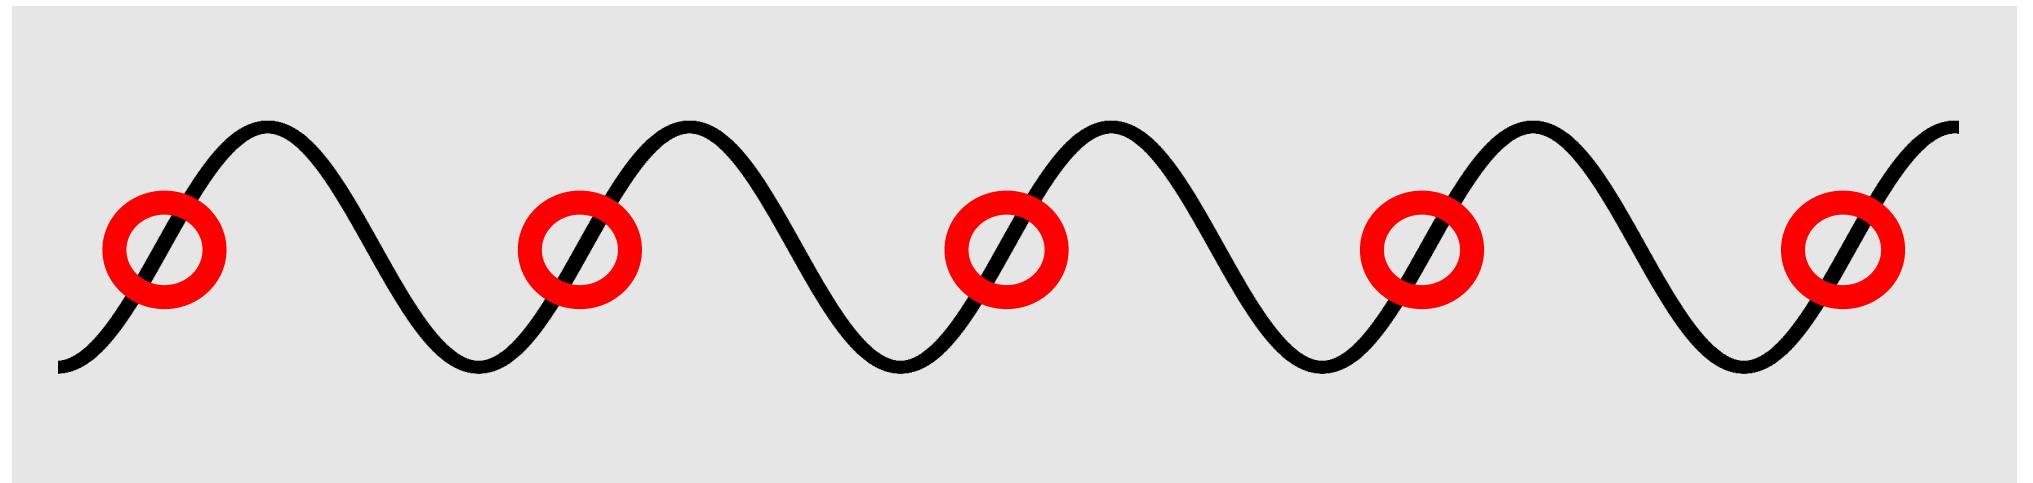
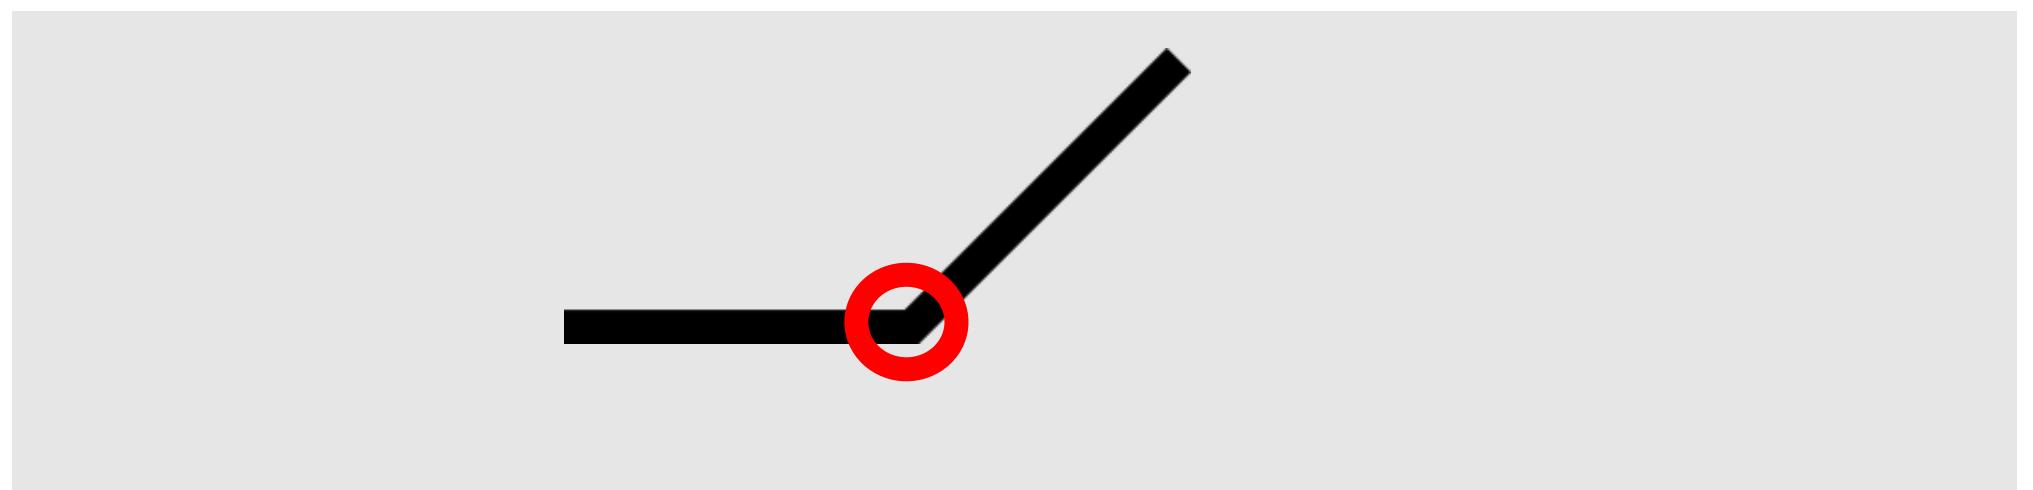
ReLU MLP



SIREN



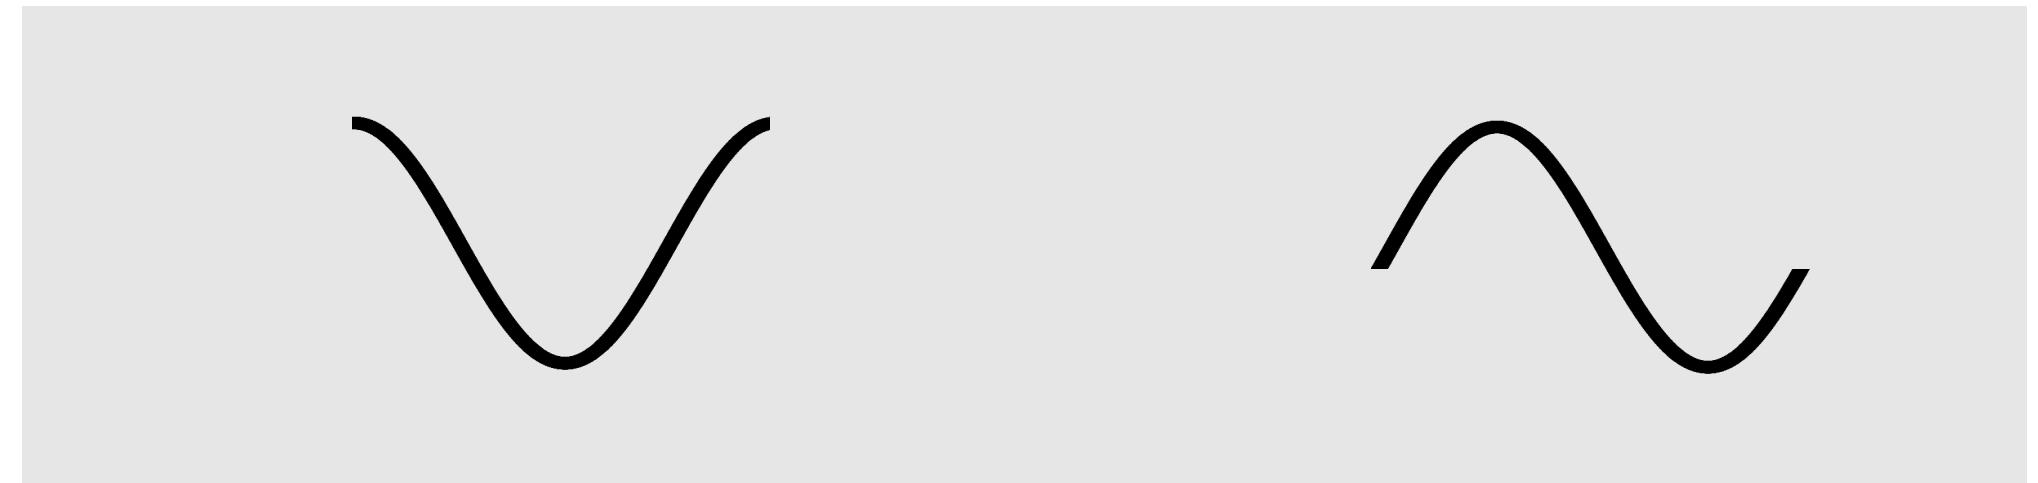
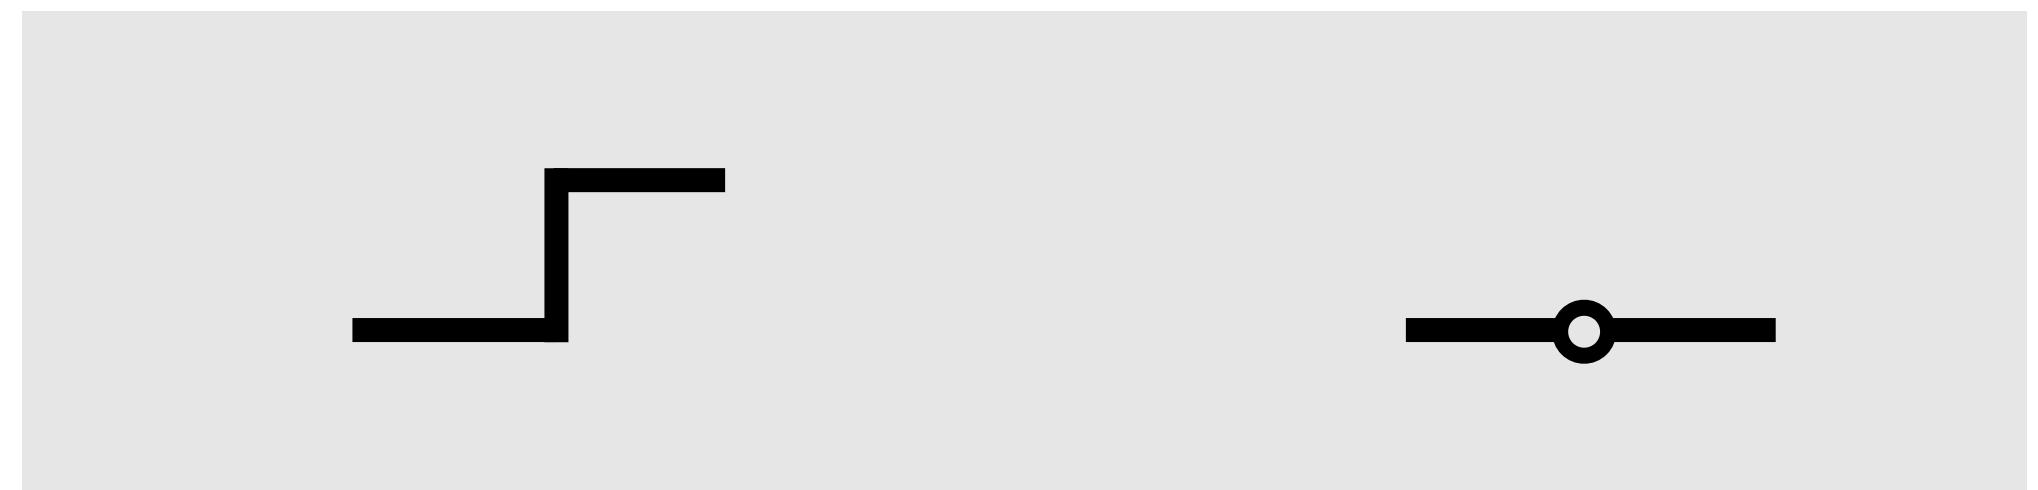
Locality



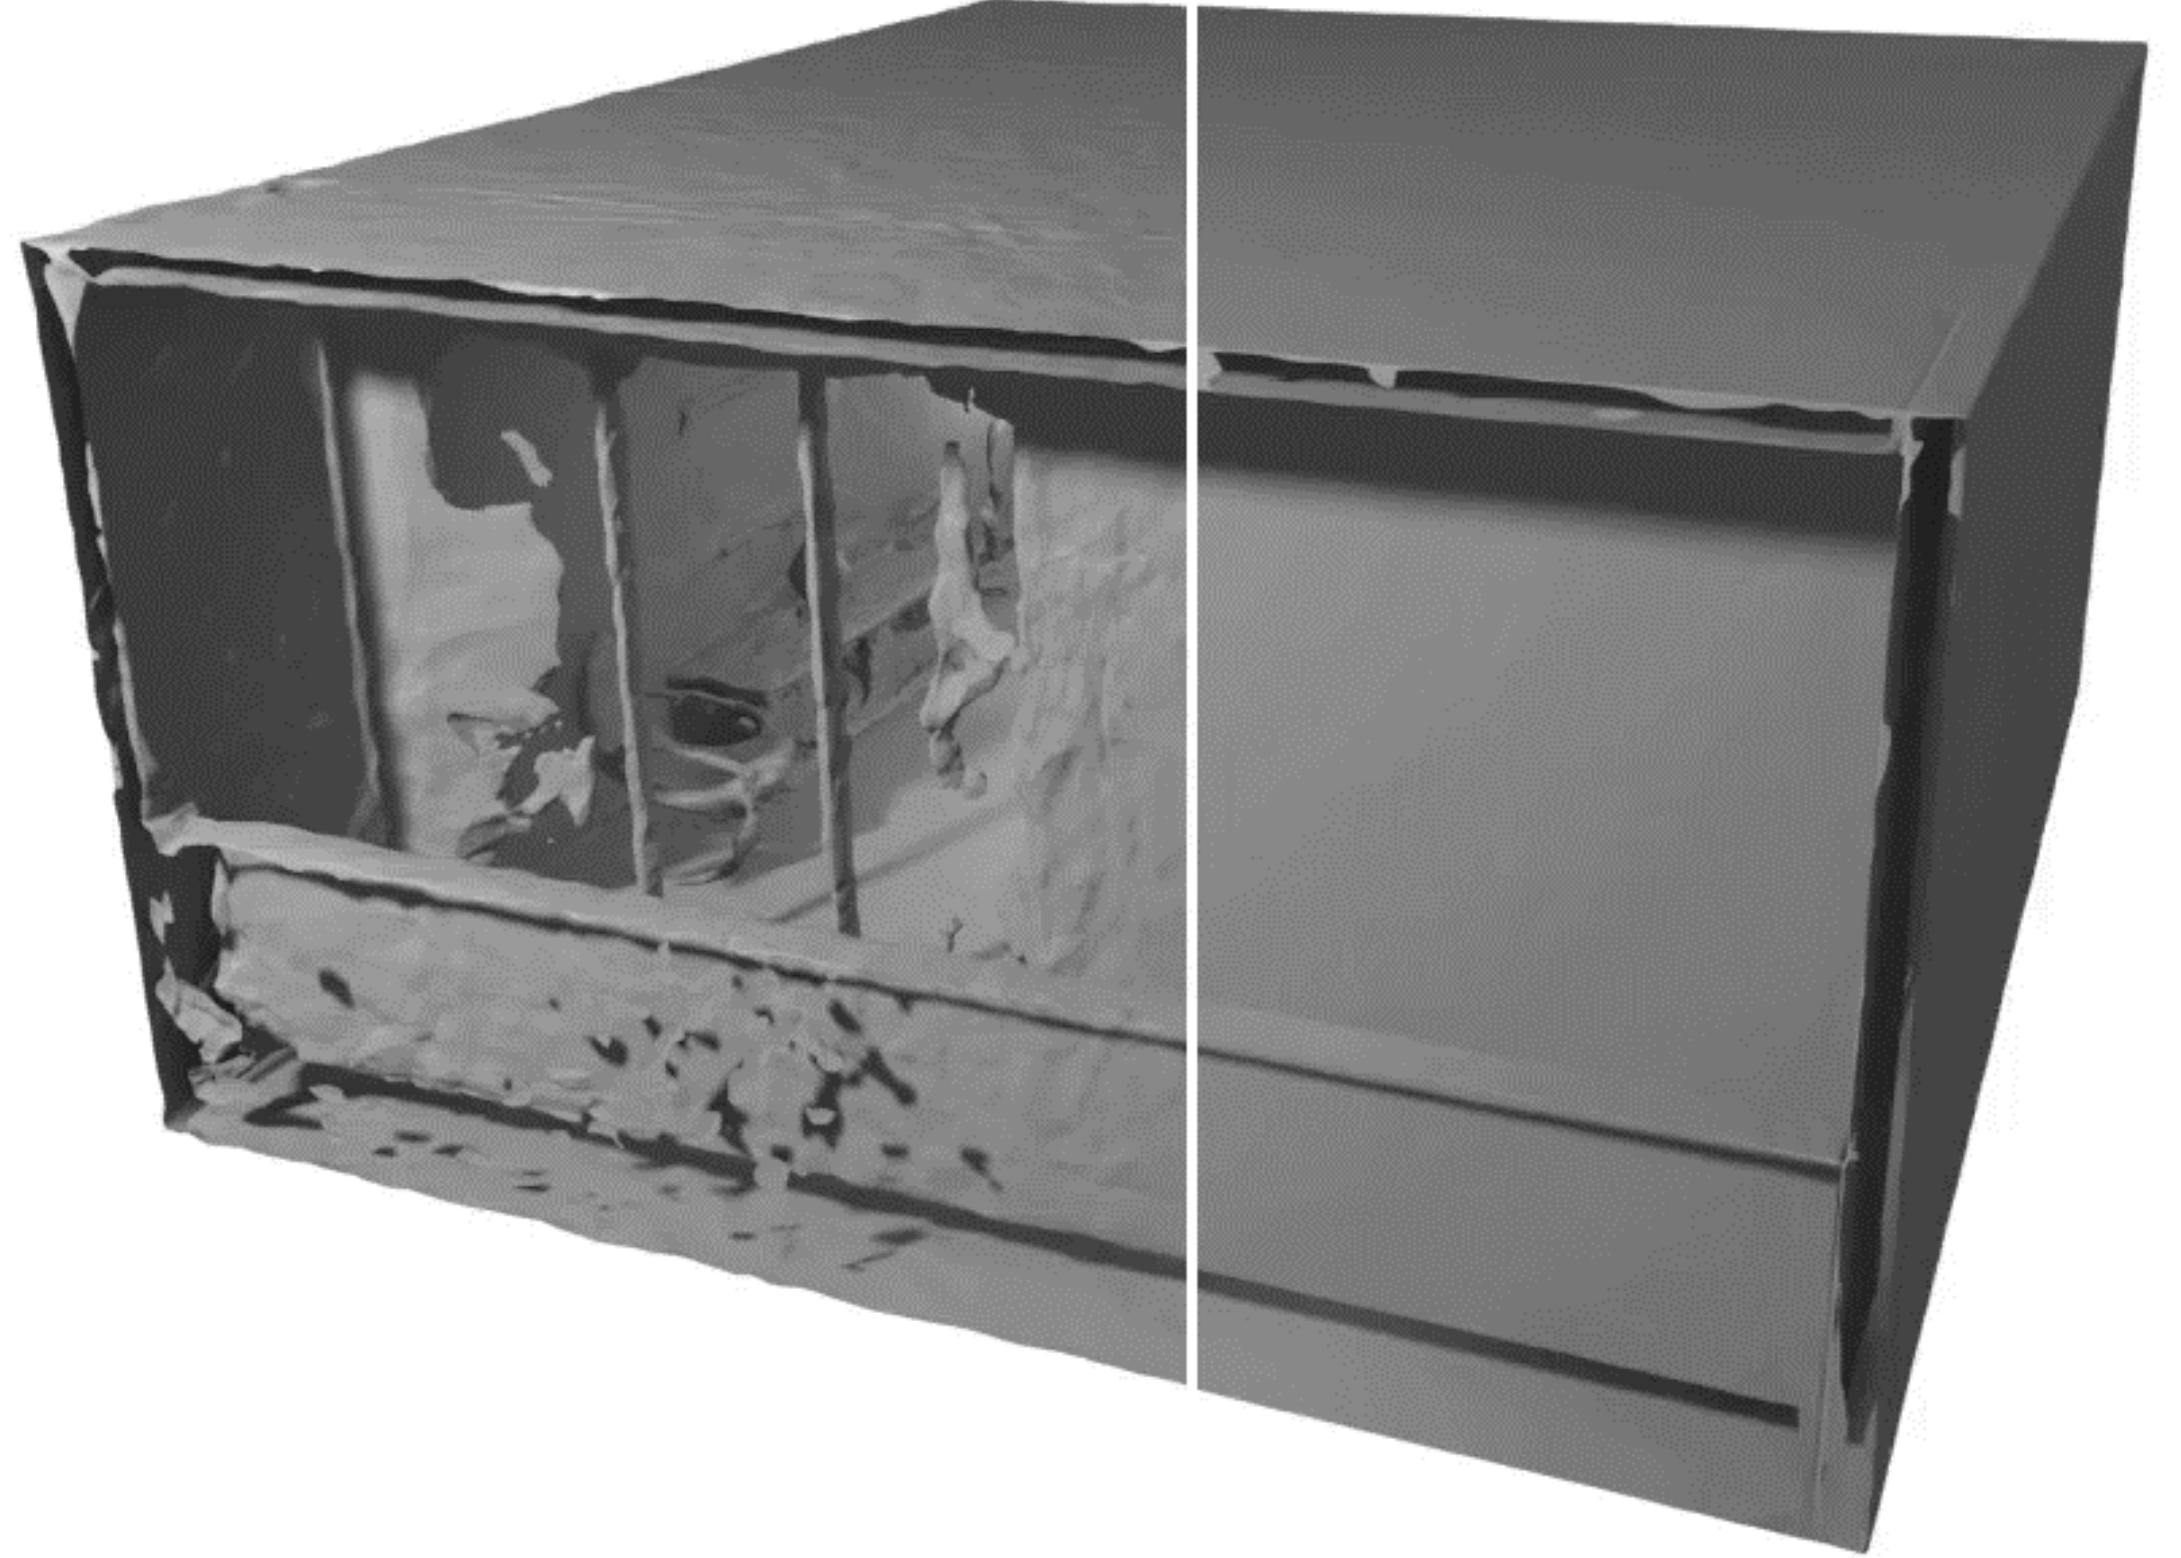
Periodicity allows SIREN to replicate activations across input domain

Derivatives

$$\delta/\delta x$$



All derivatives exist, are nonzero and bounded by 1



~1MB compared to
110MB of full mesh!

Background: Kernel Regression

Given “training” set $\{(x_i, y_i)\}_i$, a kernel function makes predictions on a point \mathbf{x} by interpolating labels y_i in the training set according to pairwise weights between x_i to \mathbf{x} as measured by a kernel function:

$$f(\mathbf{x}) = \sum_{i=1}^n (\mathbf{K}^{-1} \mathbf{y})_i k(\mathbf{x}_i, \mathbf{x})$$

Where $\mathbf{K}_{ij} = k(\mathbf{x}_i, \mathbf{x}_j)$, k is kernel function with $k \geq 0$.

The Neural Tangent Kernel

$$f(\mathbf{x}) = \sum_{i=1}^n (\mathbf{K}^{-1}\mathbf{y})_i k(\mathbf{x}_i, \mathbf{x})$$

It turns out (Jacot et al. 2018) that in the limit of an infinitely wide MLP $f(x)$ with initialization θ_0 , test-time predictions are made according to a kernel:

$$k_{NTK}(\mathbf{x}_1, \mathbf{x}_2) = \langle \nabla_{\theta} f_{\theta_0}(\mathbf{x}_1), \nabla_{\theta} f_{\theta_0}(\mathbf{x}_2) \rangle$$

The Neural Tangent Kernel

$$f(\mathbf{x}) = \sum_{i=1}^n (\mathbf{K}^{-1}\mathbf{y})_i k(\mathbf{x}_i, \mathbf{x})$$

It turns out (Jacot et al. 2018) that in the limit of an infinitely wide MLP $f(x)$ with initialization θ_0 , test-time predictions are made according to a kernel:

$$k_{NTK}(\mathbf{x}_1, \mathbf{x}_2) = \langle \nabla_{\theta} f_{\theta_0}(\mathbf{x}_1), \nabla_{\theta} f_{\theta_0}(\mathbf{x}_2) \rangle$$

In other words: “Neural” implicit representations interpolate the training set according to a *relatively* simple rule. No “intelligence”, no AI, no magic.

The Neural Tangent Kernel

$$f(\mathbf{x}) = \sum_{i=1}^n (\mathbf{K}^{-1}\mathbf{y})_i k(\mathbf{x}_i, \mathbf{x})$$

It turns out (Jacot et al. 2018) that in the limit of an infinitely wide MLP $f(x)$ with initialization θ_0 , test-time predictions are made according to a kernel:

$$k_{NTK}(\mathbf{x}_1, \mathbf{x}_2) = \langle \nabla_{\theta} f_{\theta_0}(\mathbf{x}_1), \nabla_{\theta} f_{\theta_0}(\mathbf{x}_2) \rangle$$

What we can do is design the kernel and pick the space of (\mathbf{x}, \mathbf{y}) , the space in which our neural network will interpolate.

The Neural Tangent Kernel

$$f(\mathbf{x}) = \sum_{i=1}^n (\mathbf{K}^{-1}\mathbf{y})_i k(\mathbf{x}_i, \mathbf{x})$$

It turns out (Jacot et al. 2018) that in the limit of an infinitely wide MLP $f(x)$ with initialization θ_0 , test-time predictions are made according to a kernel:

$$k_{NTK}(\mathbf{x}_1, \mathbf{x}_2) = \langle \nabla_{\theta} f_{\theta_0}(\mathbf{x}_1), \nabla_{\theta} f_{\theta_0}(\mathbf{x}_2) \rangle$$

We do not know of an equivalent rule for transformers to date - work in progress!

The Neural Tangent Kernel

$$f(\mathbf{x}) = \sum_{i=1}^n (\mathbf{K}^{-1}\mathbf{y})_i k(\mathbf{x}_i, \mathbf{x})$$

It turns out (Jacot et al. 2018) that in the limit of an infinitely wide MLP $f(x)$ with initialization θ_0 , test-time predictions are made according to a kernel:

$$k_{NTK}(\mathbf{x}_1, \mathbf{x}_2) = \langle \nabla_{\theta} f_{\theta_0}(\mathbf{x}_1), \nabla_{\theta} f_{\theta_0}(\mathbf{x}_2) \rangle$$

What does this mean for neural fields?

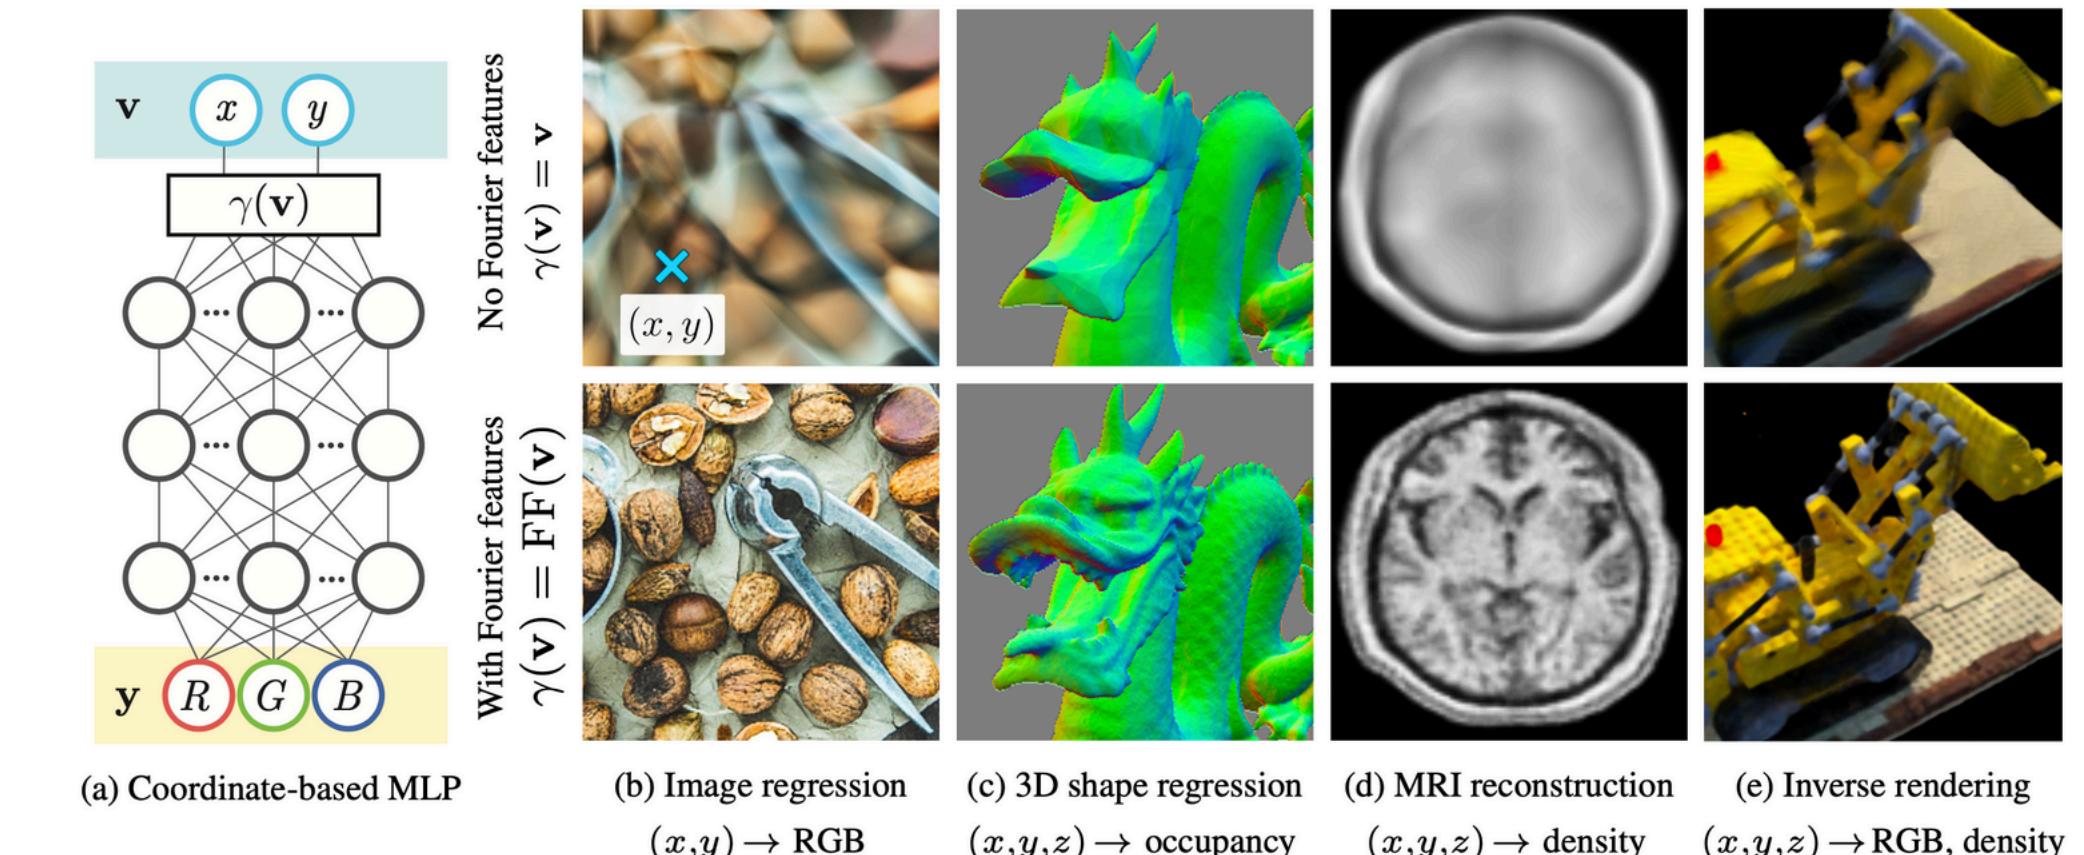
SIREN: Neural Implicit Representations With Periodic Activation Functions

Sitzmann & Martel et al. NeurIPS 2020



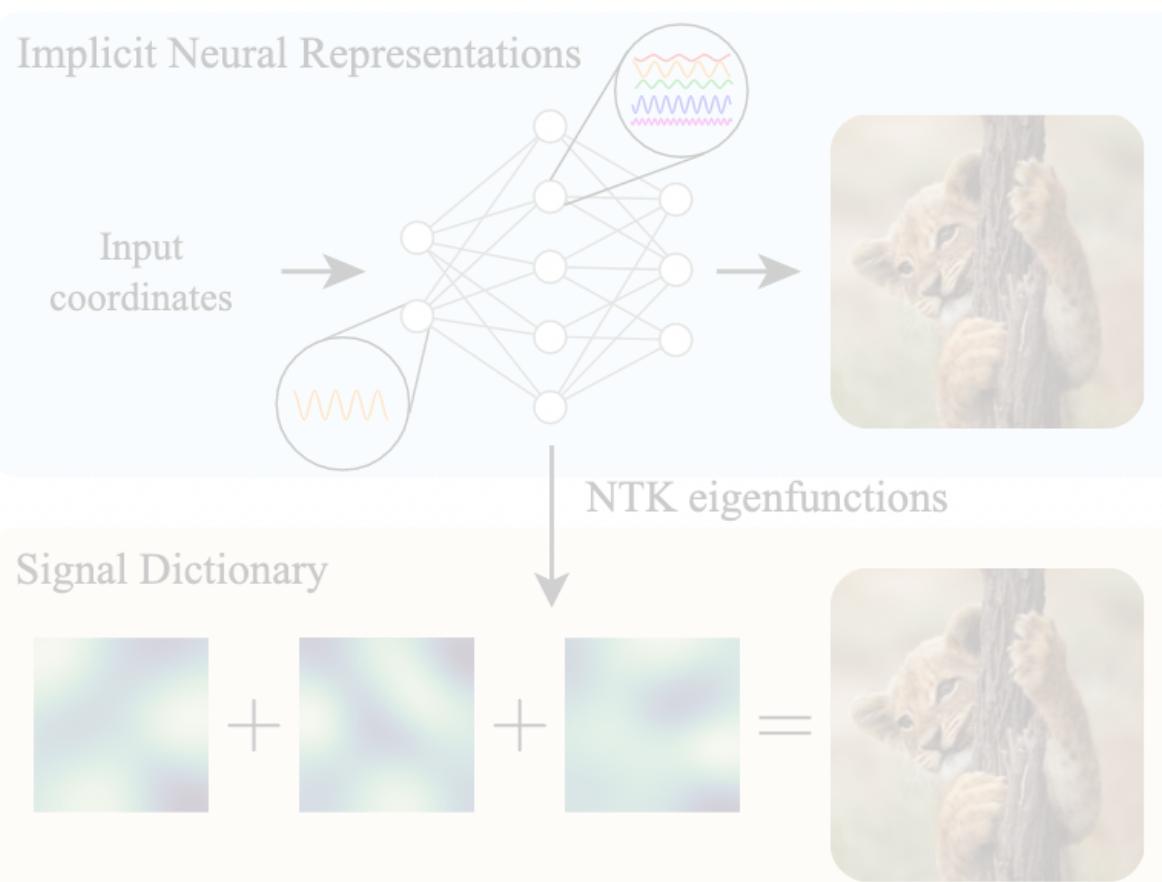
Fourier Features Let Networks Learn High Frequency Functions in Low Dimensional Domains

Tancik, Srinivasan, Mildenhall et al, NeurIPS 2020



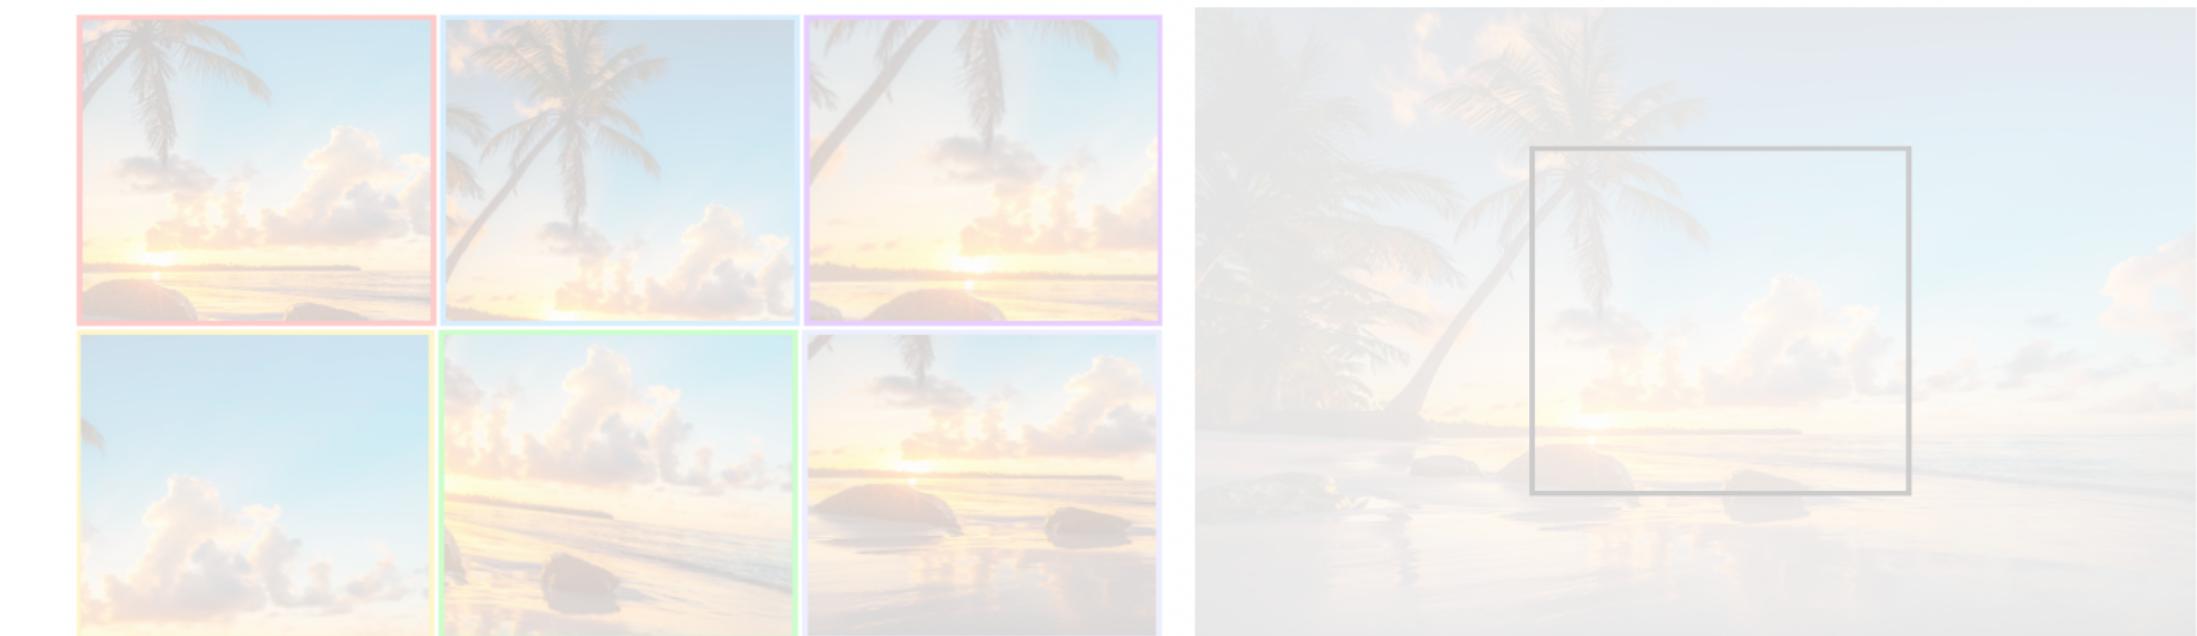
A Structured Dictionary Perspective on Implicit Neural Representations

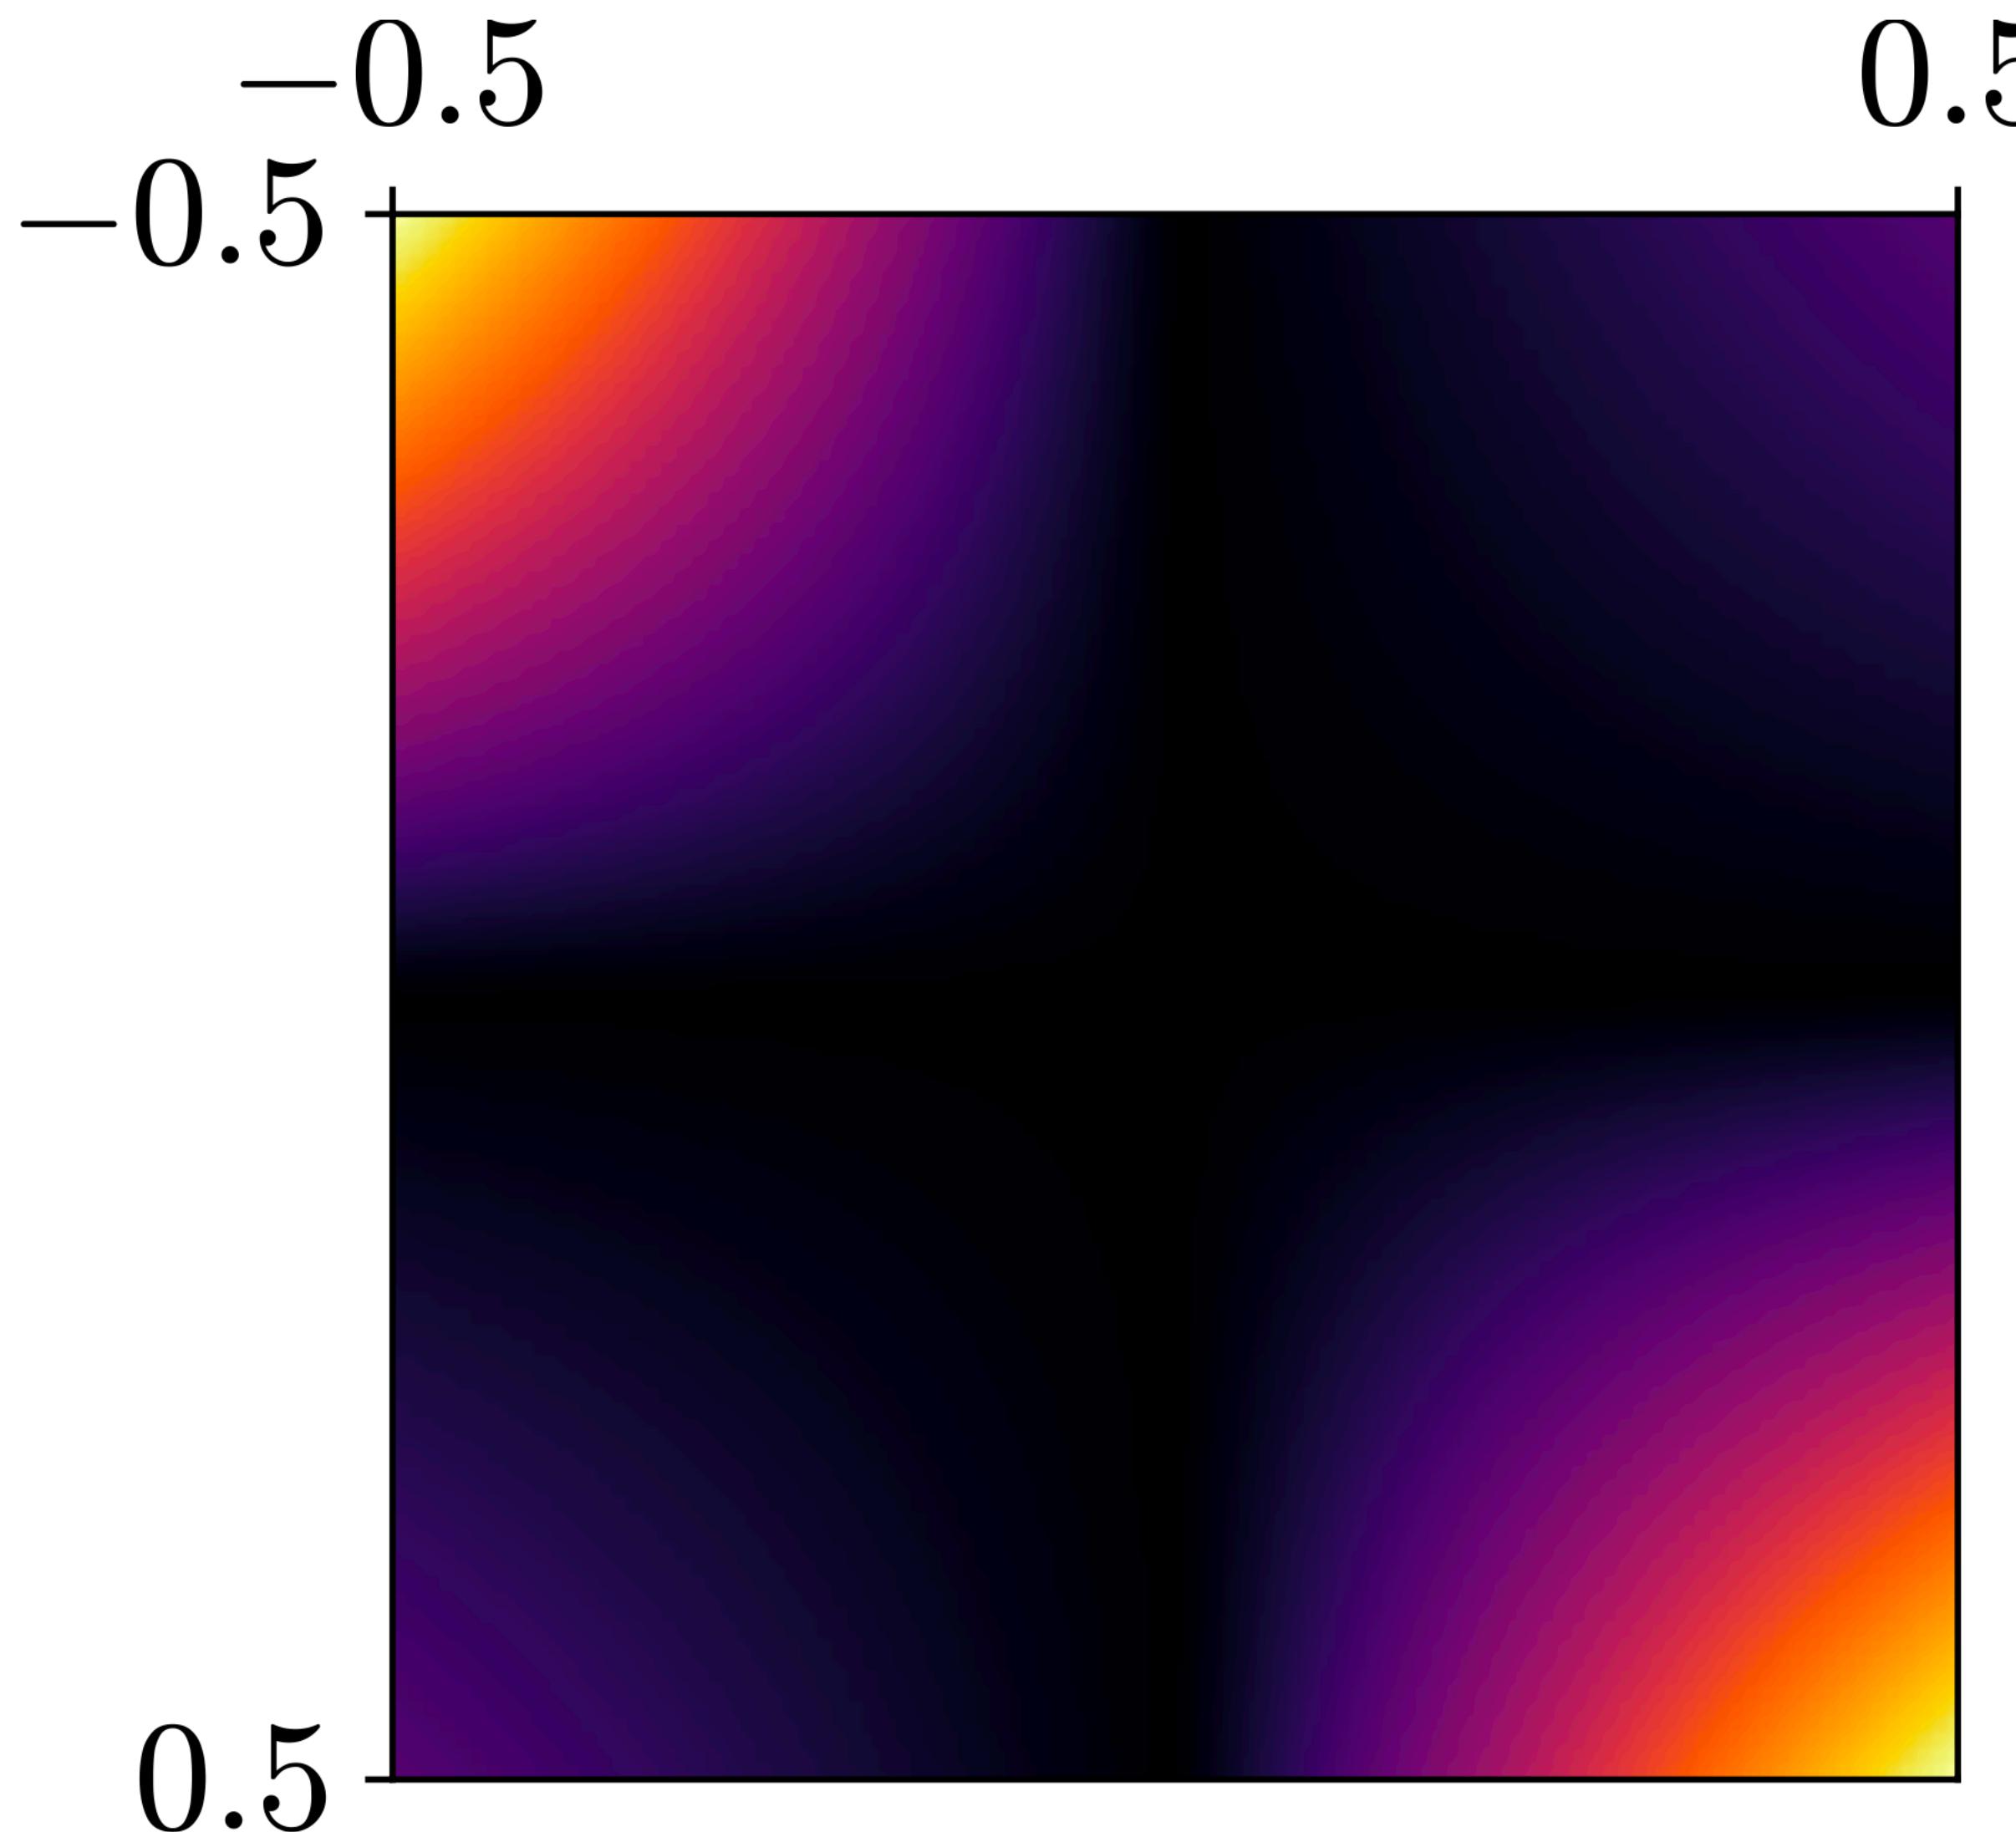
Yüce & Ortiz-Jiménez et al. CVPR 2022



Gaussian Activated Neural Radiance Fields for High Fidelity Reconstruction & Pose Estimation

Chng et al. ECCV 2022



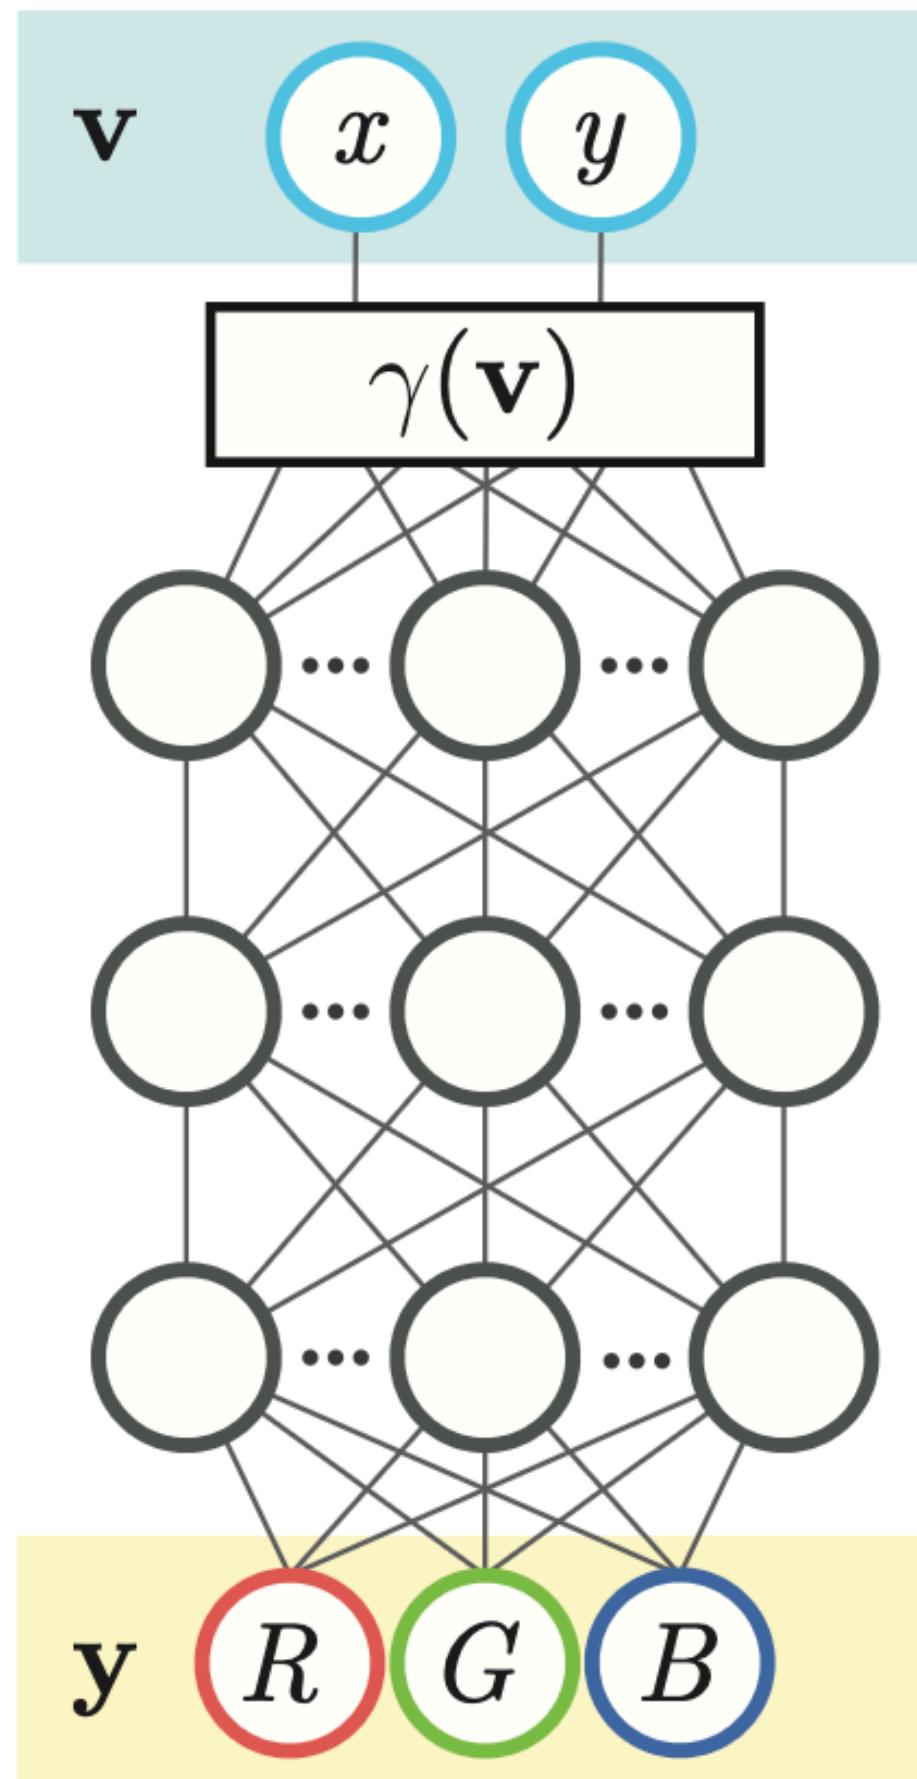


NTK kernel $k_{NTK}(x_i, x_j)$ for a 4-layer
ReLU MLP with one scalar input.

From “Fourier Features Let Networks Learn High Frequency Functions in Low Dimensional Domains”, Tancik, Srinivasan, Mildenhall et al.

NTK of ReLU neural net is quite non-local.

Feature Embeddings



$$\gamma(\mathbf{x}) = [\gamma_1(\mathbf{x}), \dots, \gamma_n(\mathbf{x})]$$

$$\gamma_{2i}(\mathbf{x}) = \sin(2^{i-1}\pi x_i)$$

$$\gamma_{(2i+1)}(\mathbf{x}) = \cos(2^{i-1}\pi x_i)$$

Sinusoidal Embeddings

Zhong et al. ICLR 2020

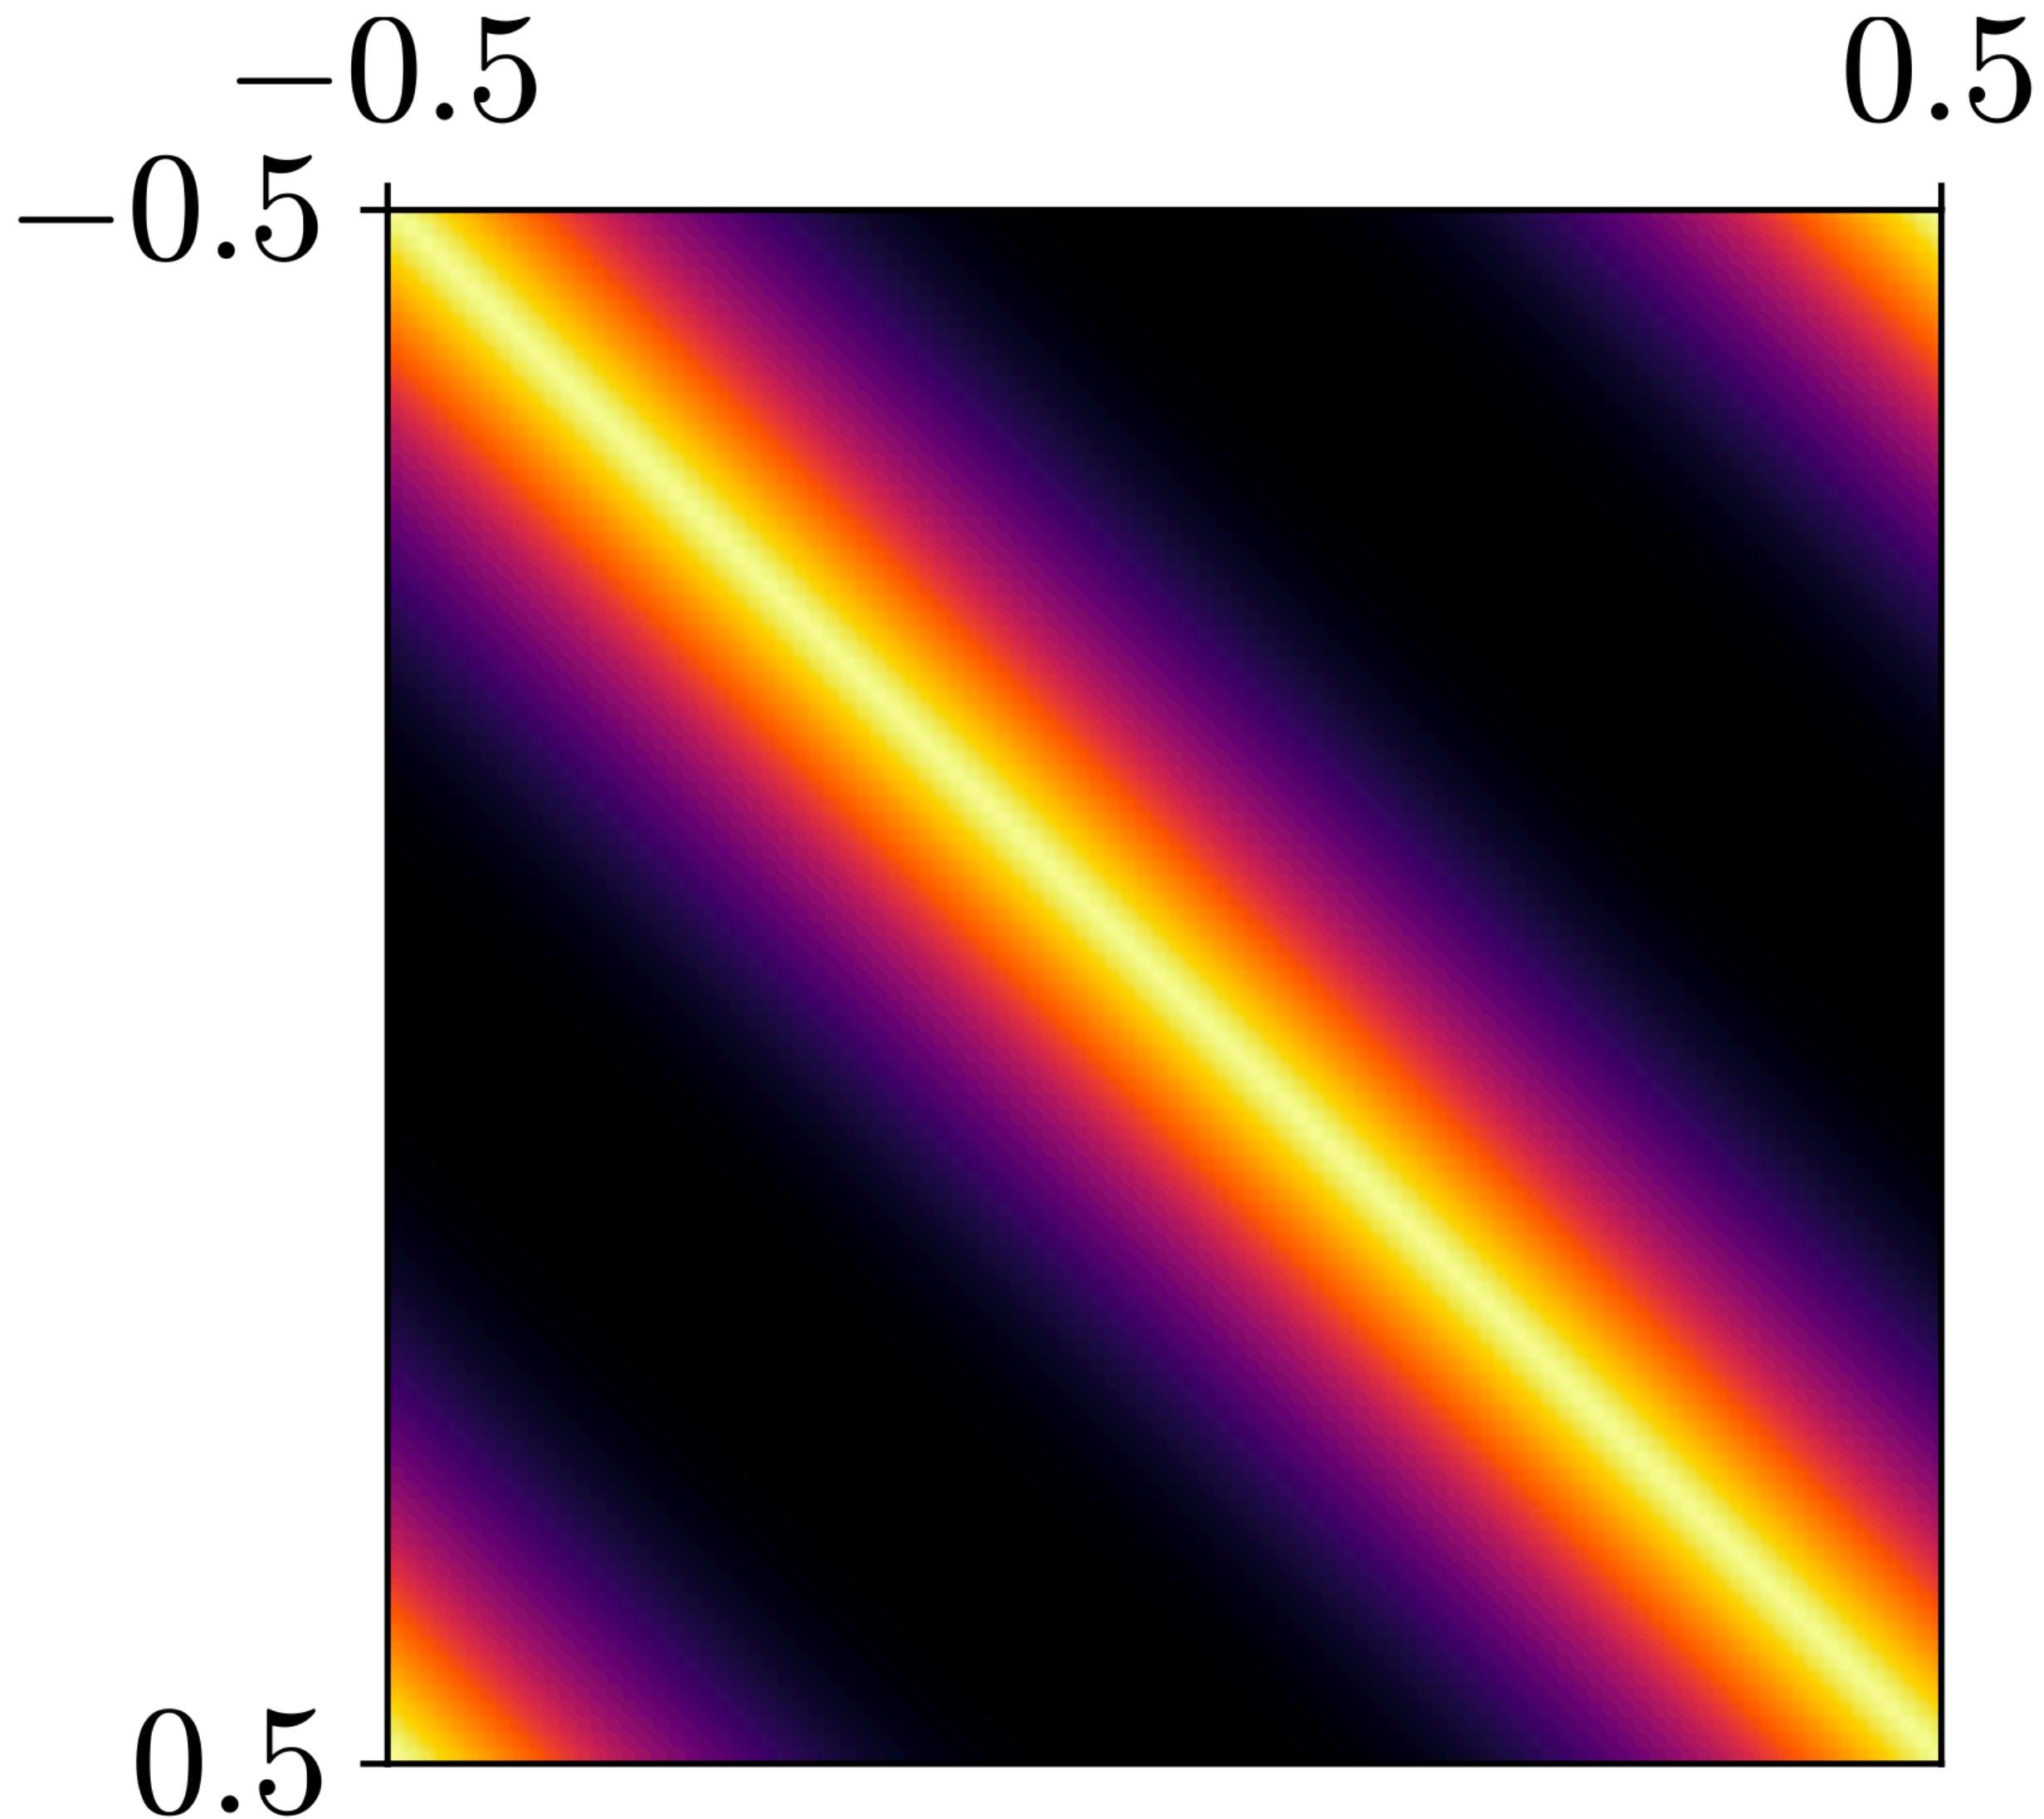
Mildenhall et al., ECCV 2020

$$\gamma(\mathbf{x}) = \exp\left(-\frac{\|t - x\|^2}{2\sigma^2}\right)$$

Gaussian Embeddings

Zheng et al., arXiv 2021

• • •



NTK kernel $k_{NTK}(\gamma(x_i), \gamma(x_j))$ for a 4-layer ReLU MLP with one scalar input and **Sinusoidal Feature Mapping**

$$\gamma(x) = [\cos(2\pi x), \sin(2\pi x)].$$

From “Fourier Features Let Networks Learn High Frequency Functions in Low Dimensional Domains”, Tancik, Srinivasan, Mildenhall et al.

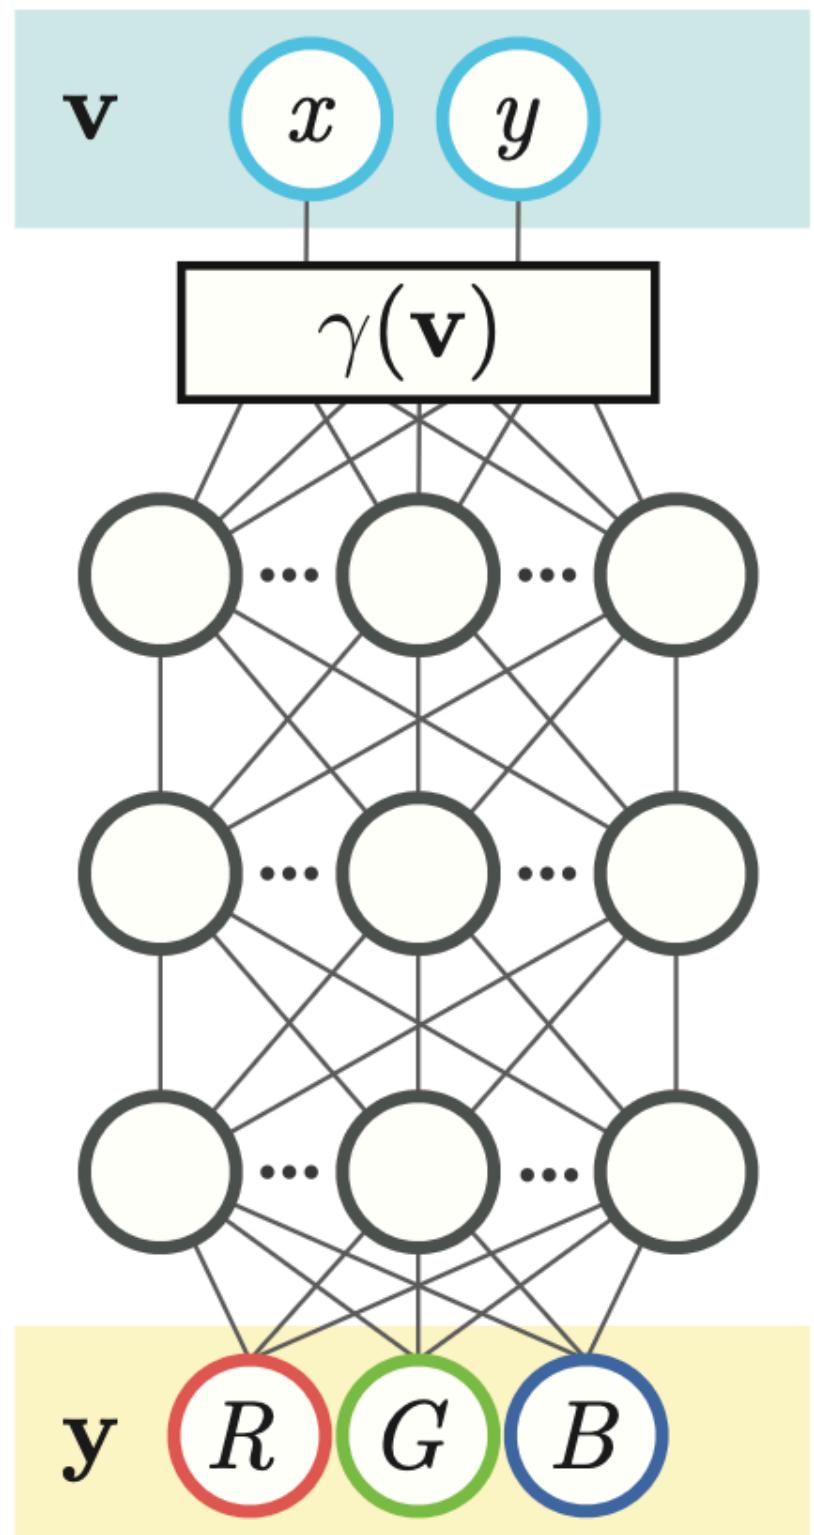
NTK is local (note diagonal).

SIREN



~1MB compared to
110MB of full mesh!

Fourier Features



(a) Coordinate-based MLP

No Fourier features
 $\gamma(\mathbf{v}) = \mathbf{v}$

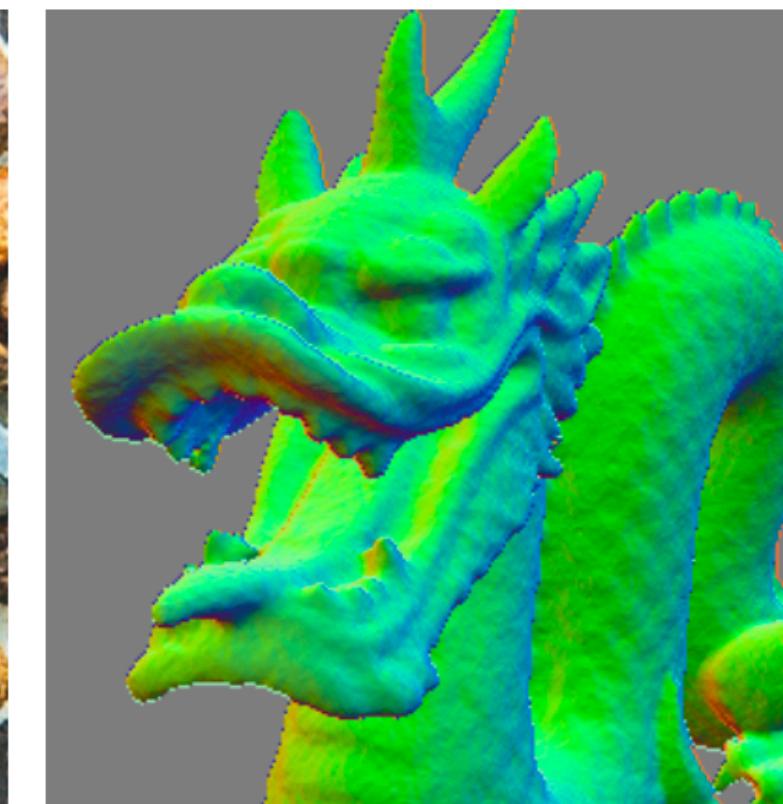
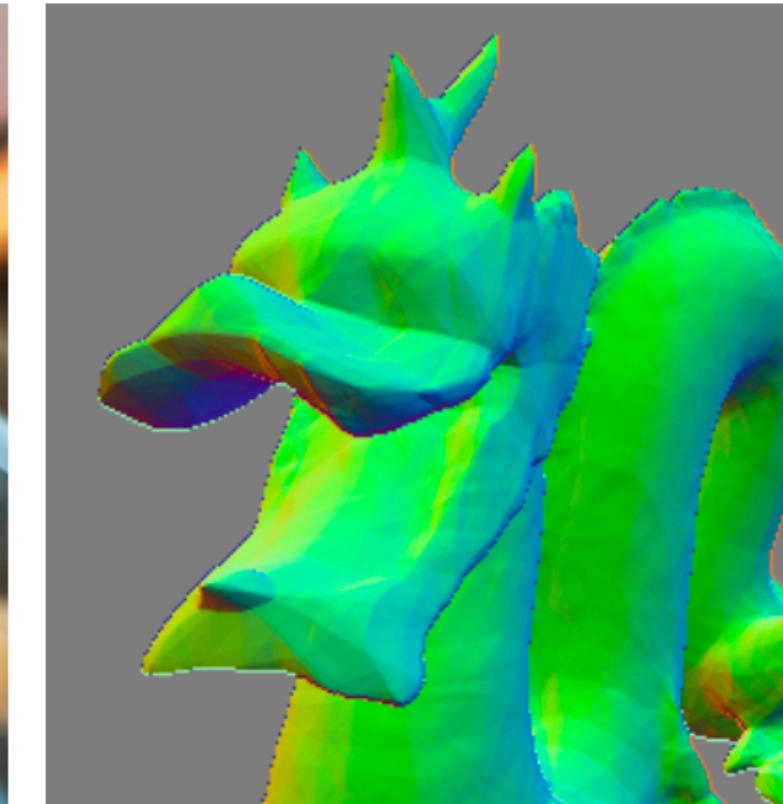


With Fourier features
 $\gamma(\mathbf{v}) = \text{FF}(\mathbf{v})$



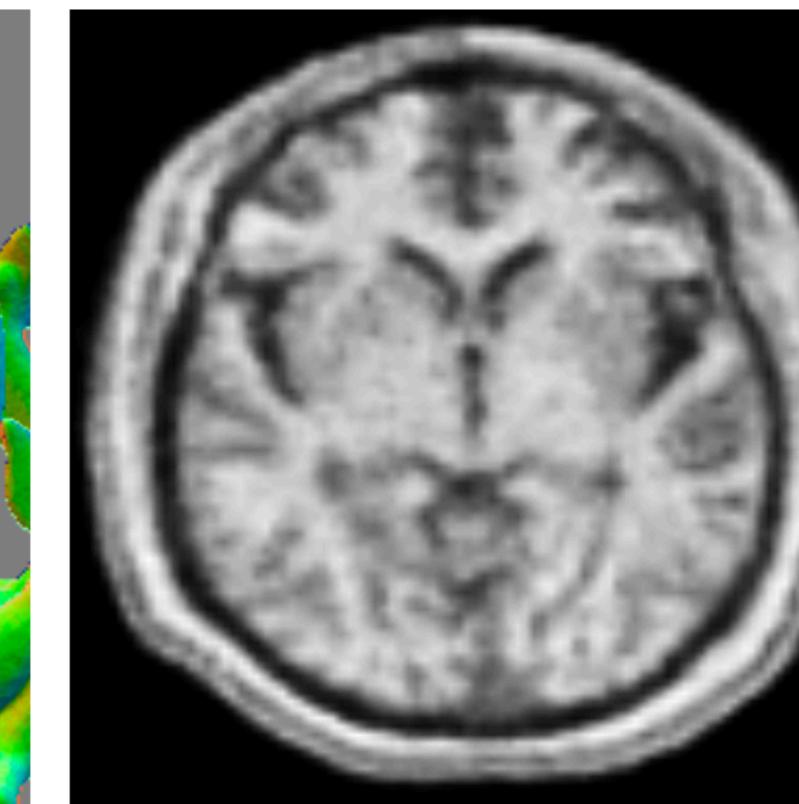
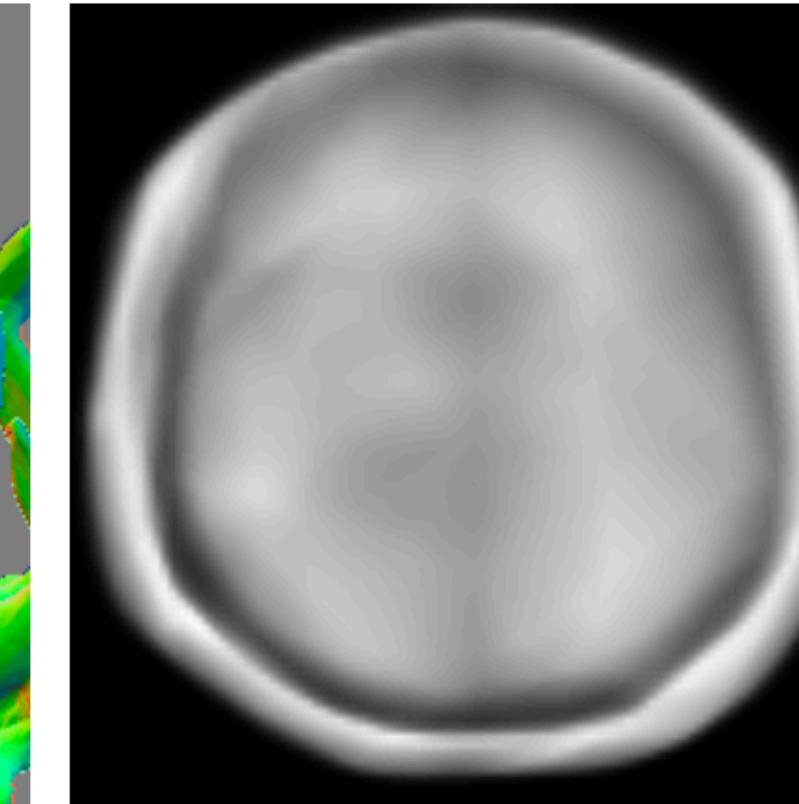
(b) Image regression

$$(x, y) \rightarrow \text{RGB}$$



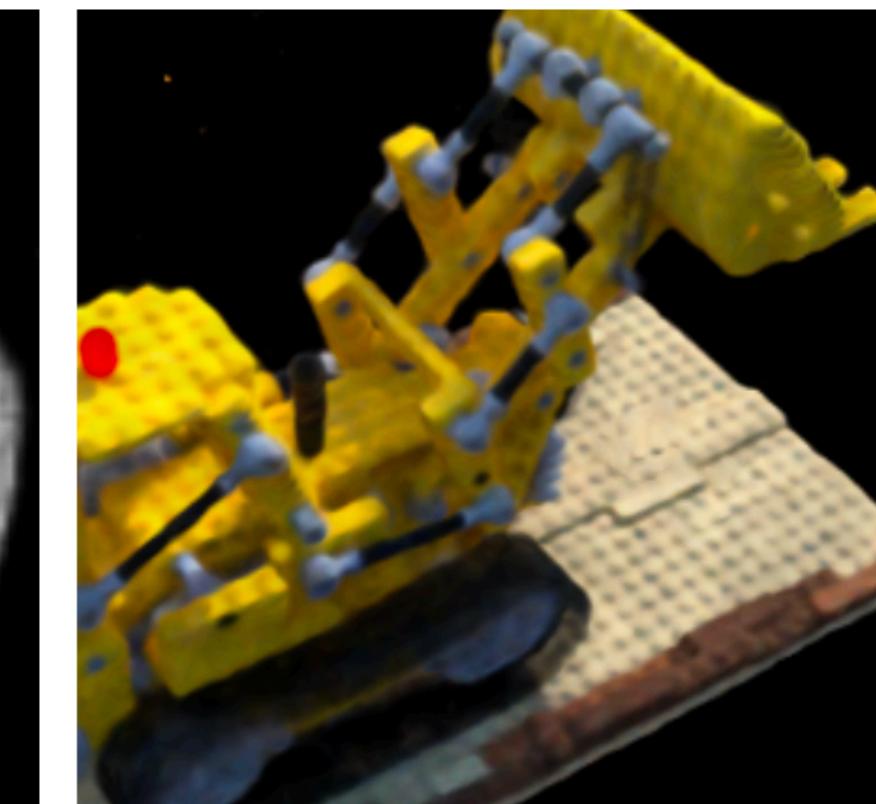
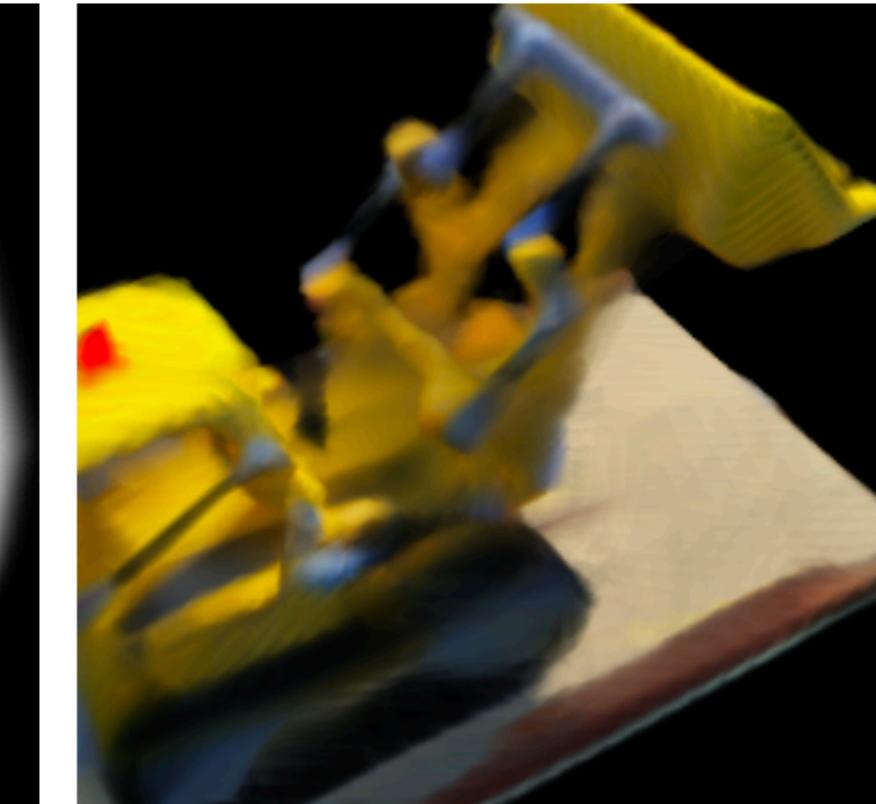
(c) 3D shape regression

$$(x, y, z) \rightarrow \text{occupancy}$$



(d) MRI reconstruction

$$(x, y, z) \rightarrow \text{density}$$



(e) Inverse rendering

$$(x, y, z) \rightarrow \text{RGB, density}$$

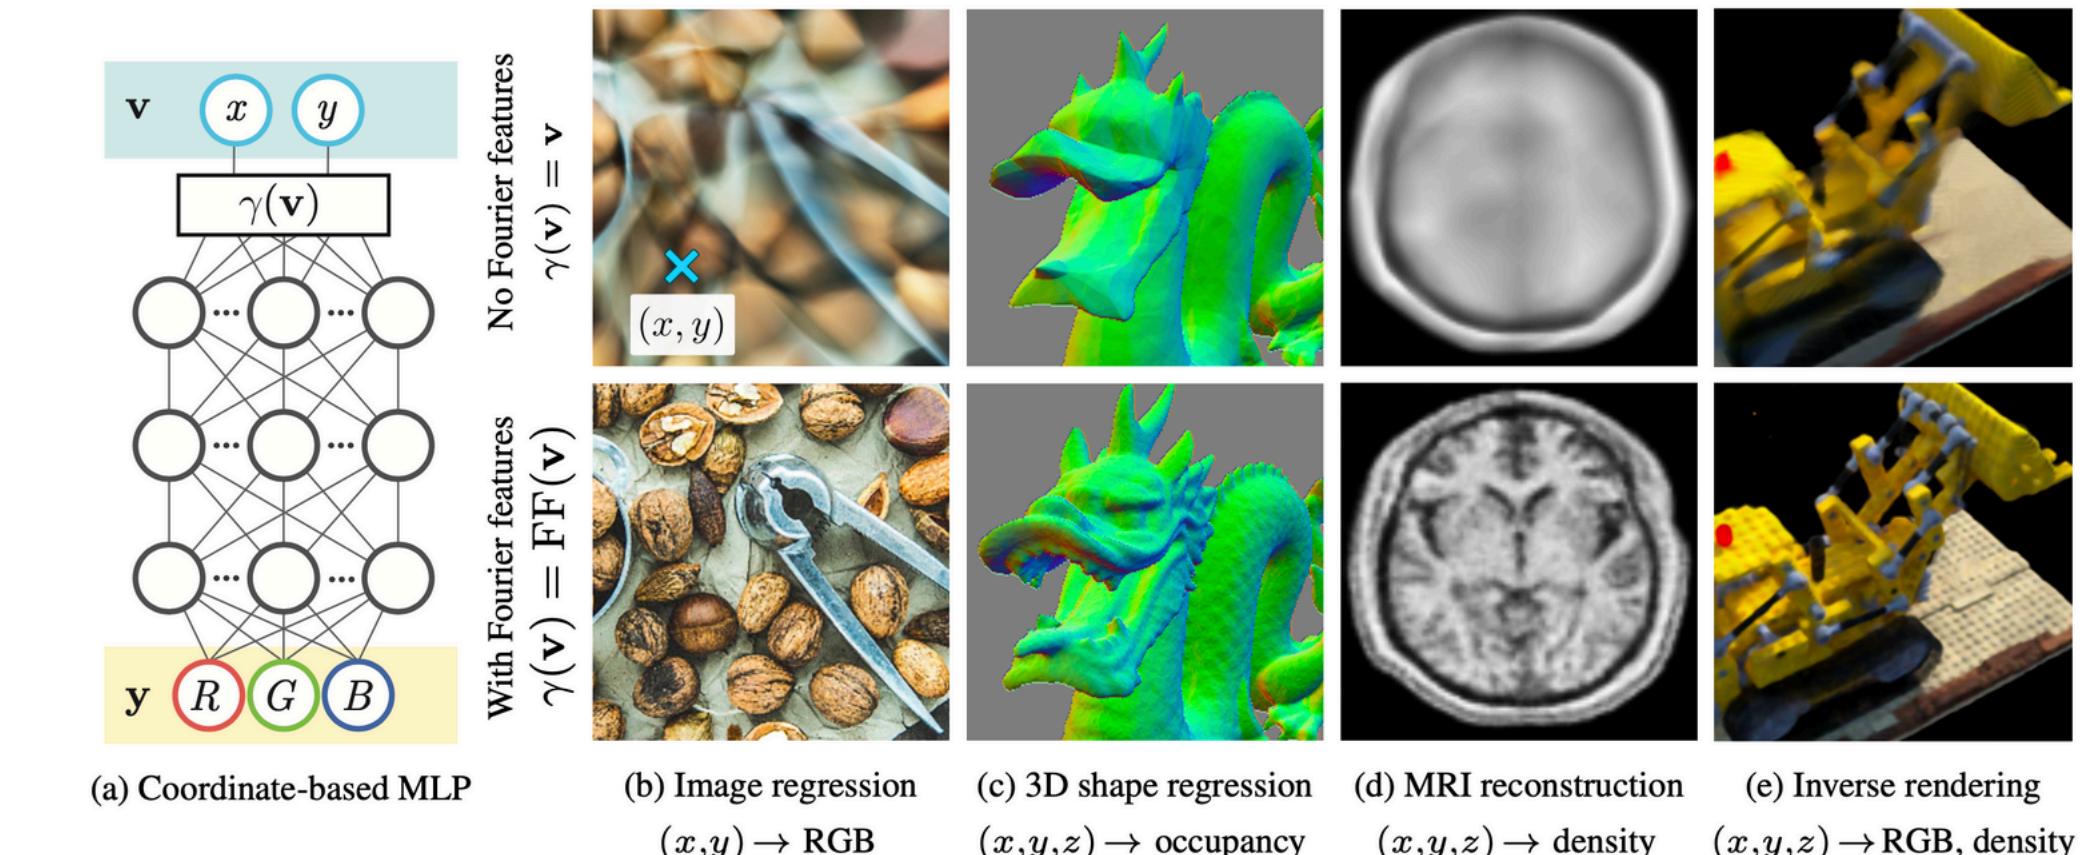
SIREN: Neural Implicit Representations With Periodic Activation Functions

Sitzmann & Martel et al. NeurIPS 2020



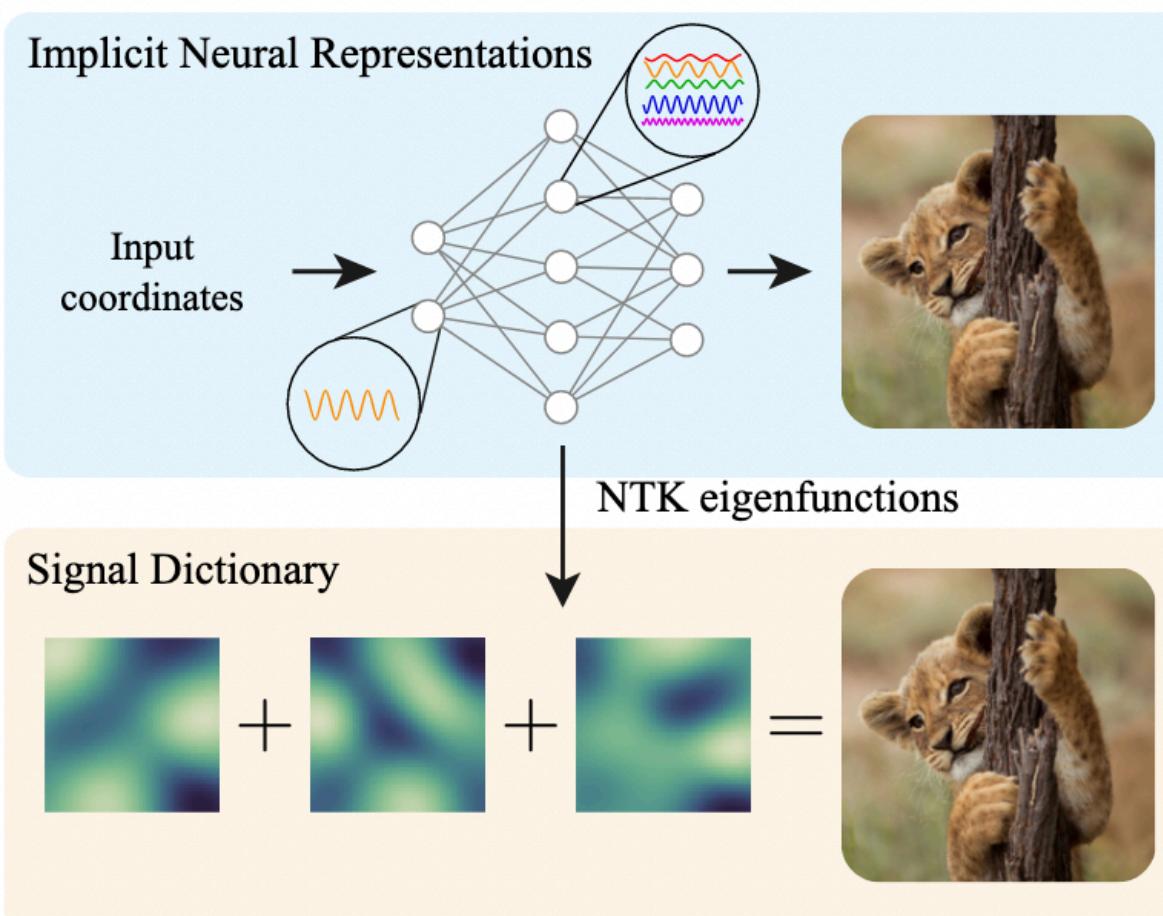
Fourier Features Let Networks Learn High Frequency Functions in Low Dimensional Domains

Tancik, Srinivasan, Mildenhall NeurIPS 2020



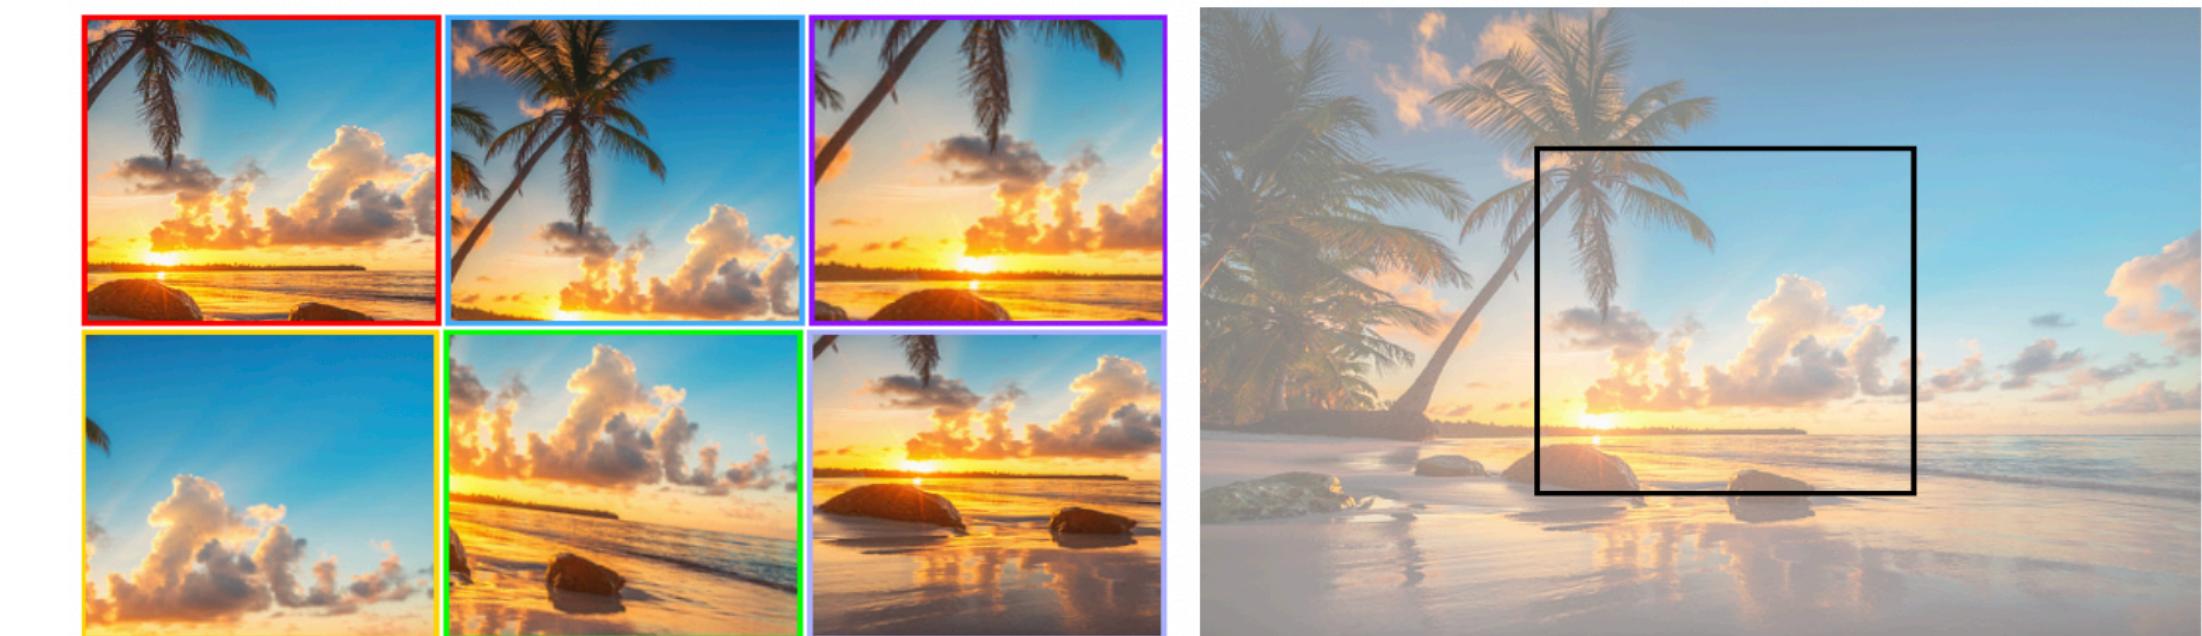
A Structured Dictionary Perspective on Implicit Neural Representations

Yüce & Ortiz-Jiménez et al. CVPR 2022



Gaussian Activated Neural Radiance Fields for High Fidelity Reconstruction & Pose Estimation

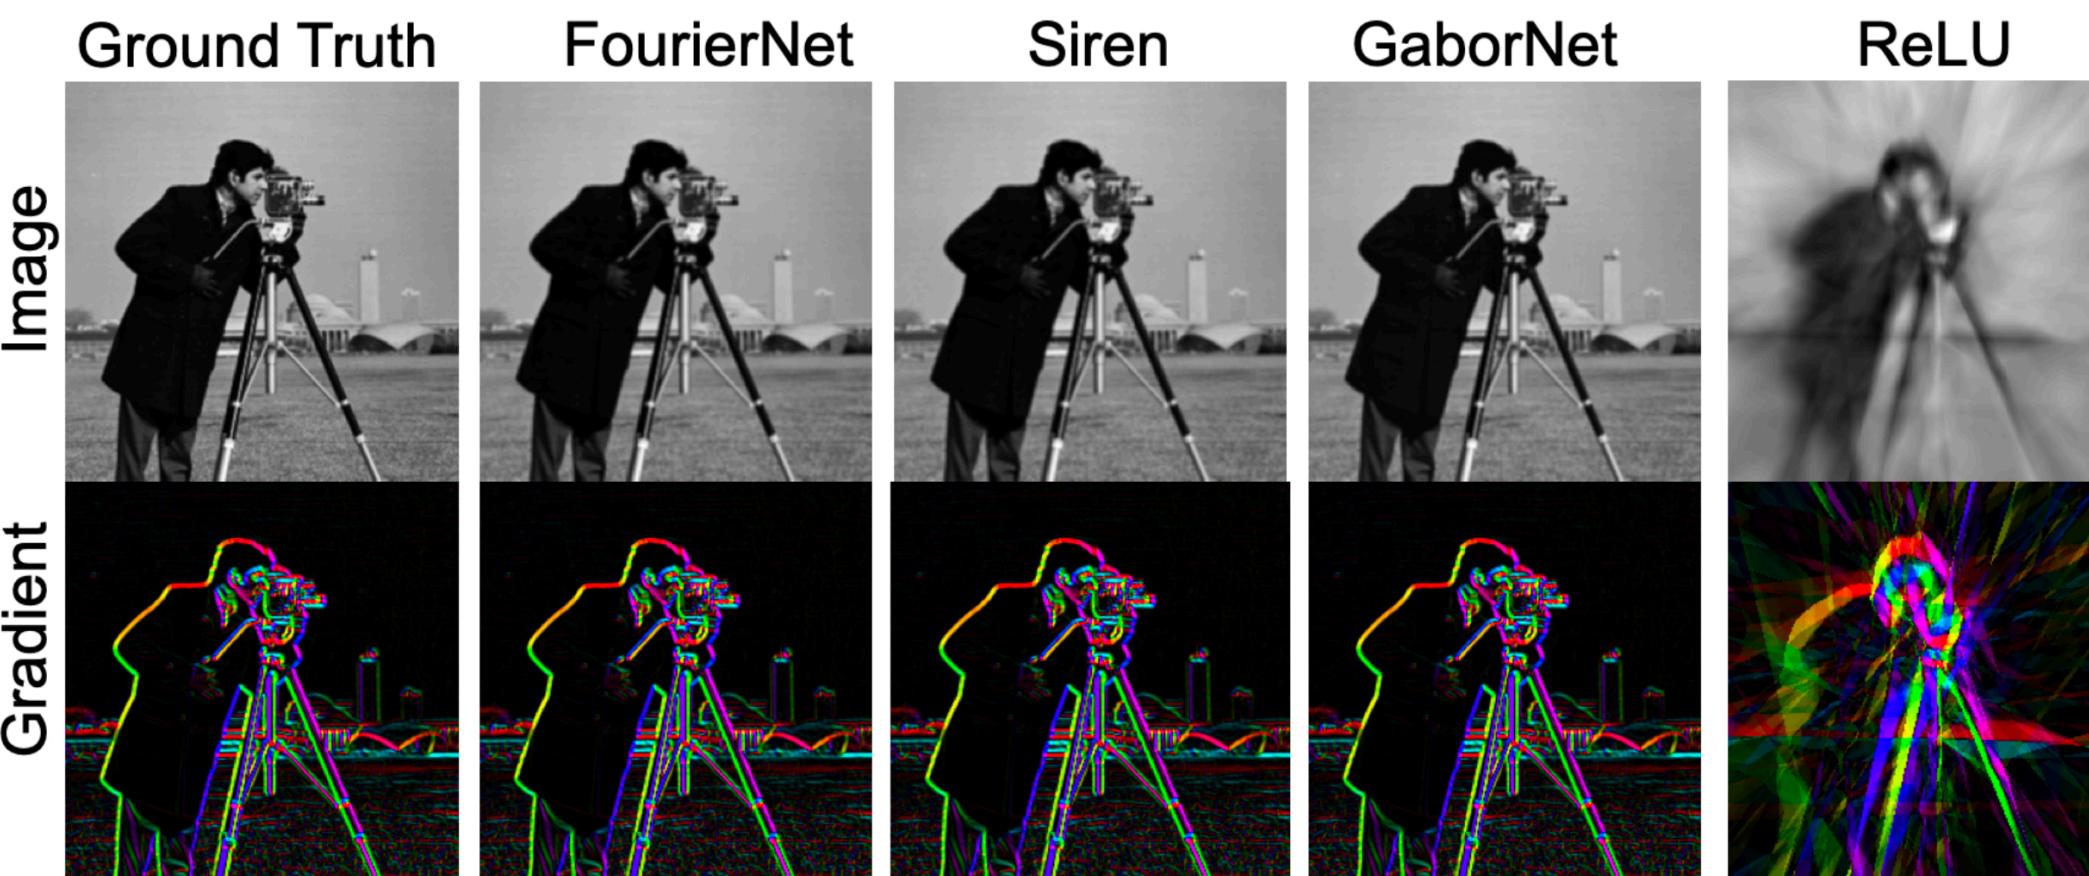
Chng et al. ECCV 2022



Neural (?) Fields with analytical (but not tractably computable) Fourier spectrum!

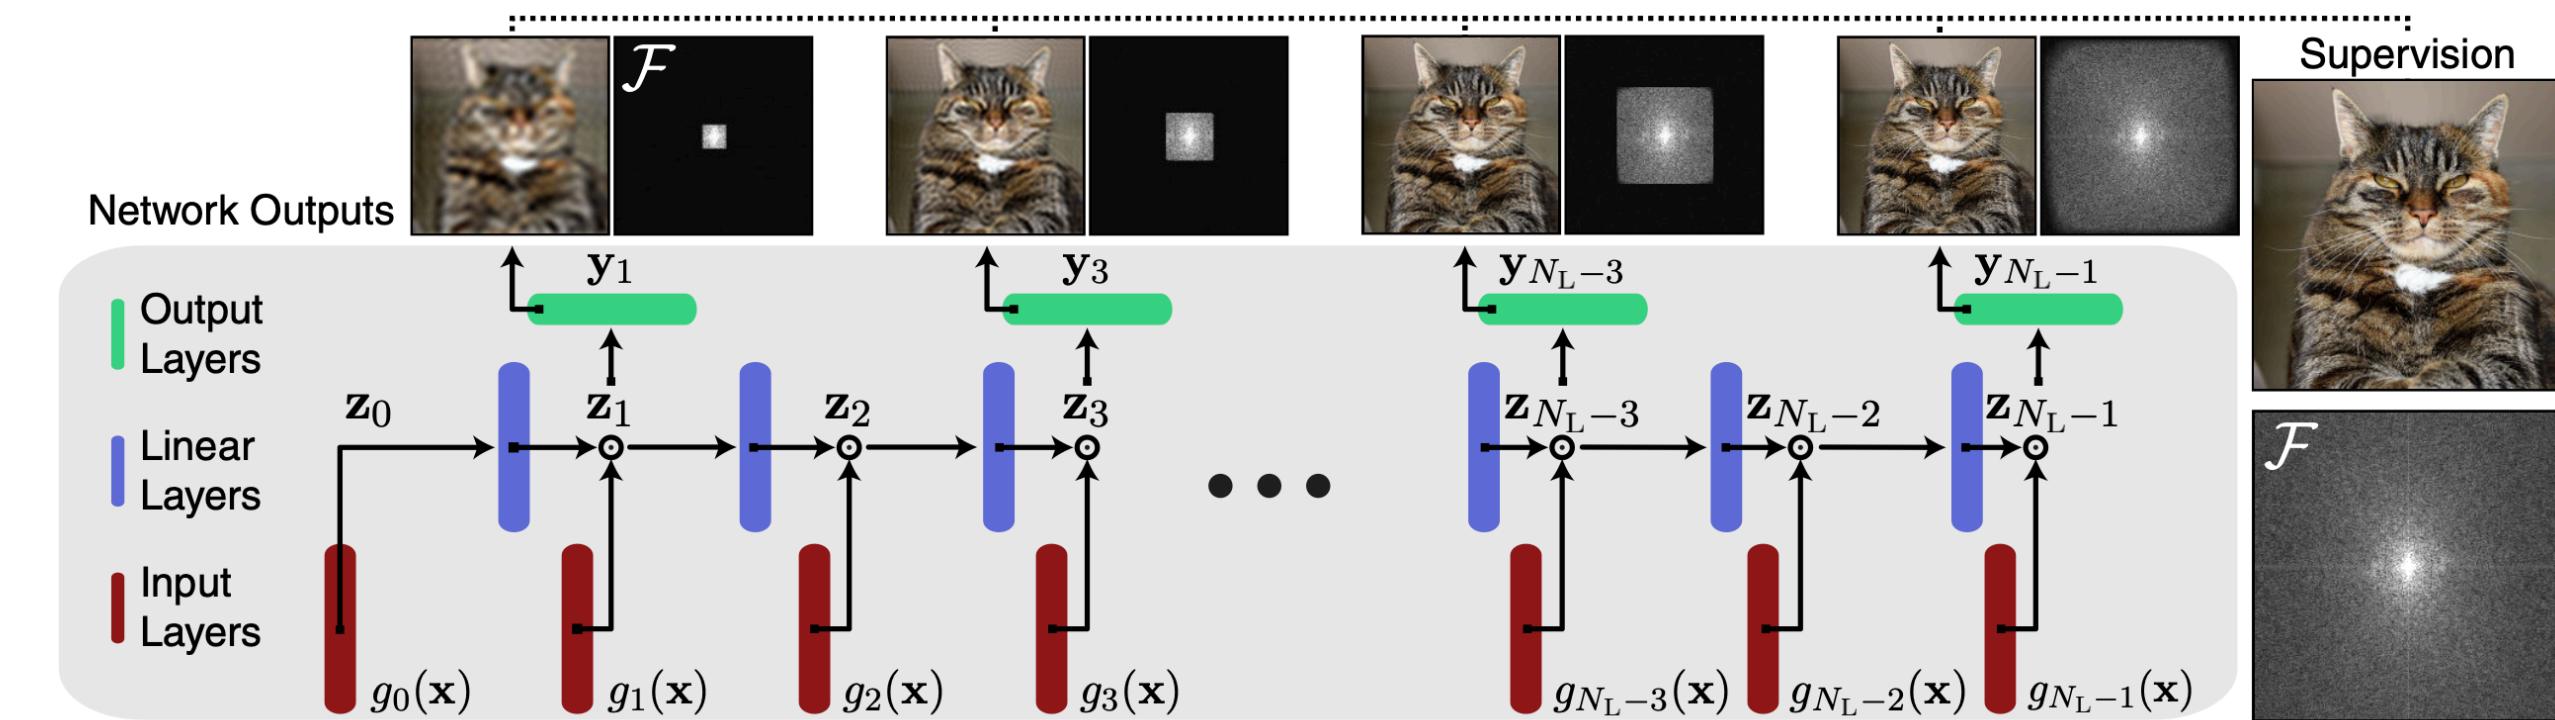
Multiplicative Filter Networks

Fathony et al. ICLR 2021



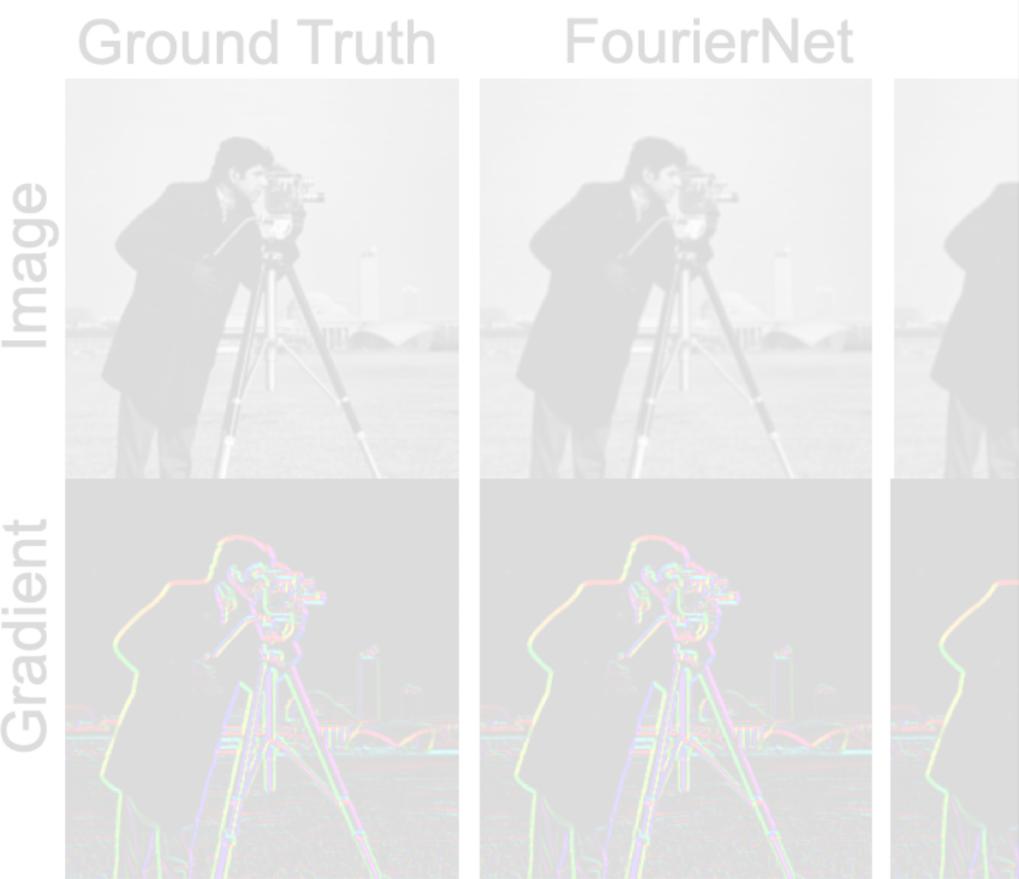
BACON: Band-limited coordinate networks for multiscale scene representation

Lindell et al. CVPR 2022



(but

Multiplicative Fathy et al.



arXiv:2111.11426v4 [cs.CV] 5 Apr 2022

EUROGRAPHICS 2022
D. Meneveaux and G. Patanè
(Guest Editors)

Volume 41 (2022), Number 2
STAR – State of The Art Report

rum!

Neural Fields in Visual Computing and Beyond

Yiheng Xie^{1,2} Towaki Takikawa^{3,4} Shunsuke Saito⁵ Or Litany⁴ Shiqin Yan¹ Numair Khan¹ Federico Tombari^{6,7}
James Tompkin¹ Vincent Sitzmann^{8†} Srinath Sridhar^{1†}

¹Brown University ²Unity Technologies ³University of Toronto ⁴NVIDIA ⁵Meta Reality Labs Research ⁶Google ⁷Technical University of Munich
⁸Massachusetts Institute of Technology [†]Equal advising

<https://neuralfields.cs.brown.edu/>

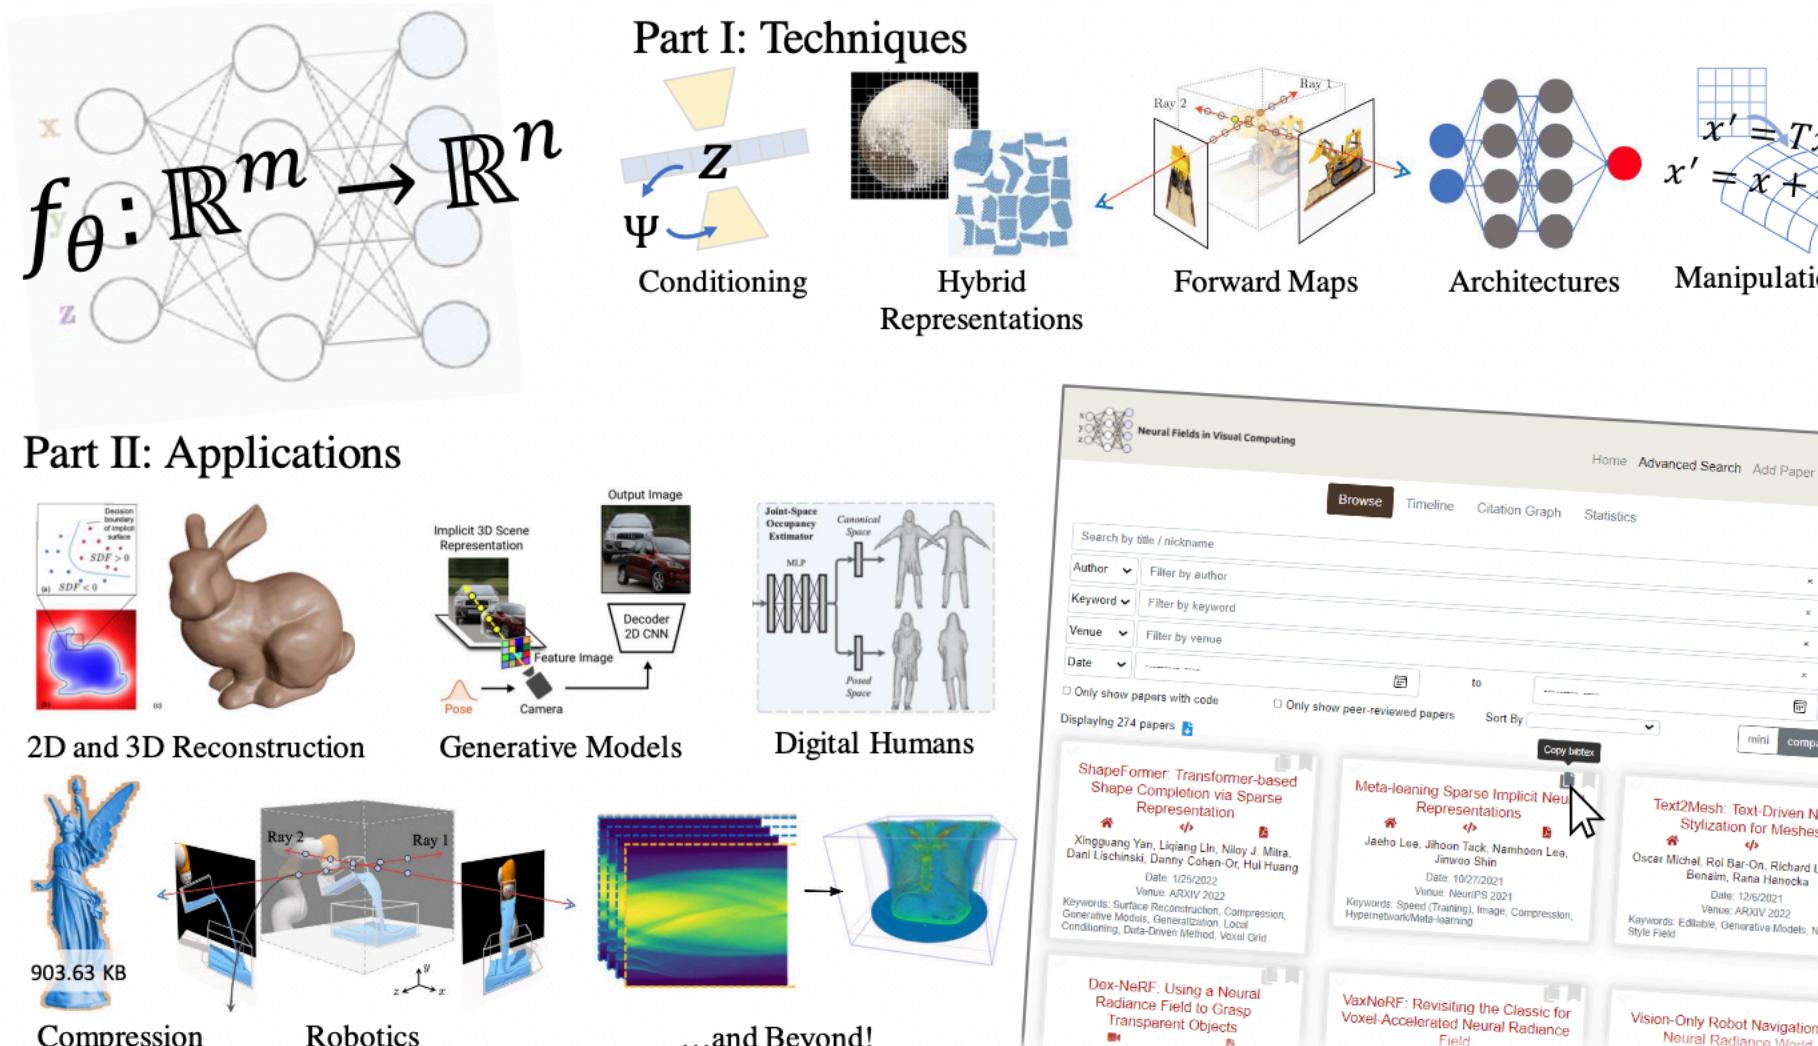
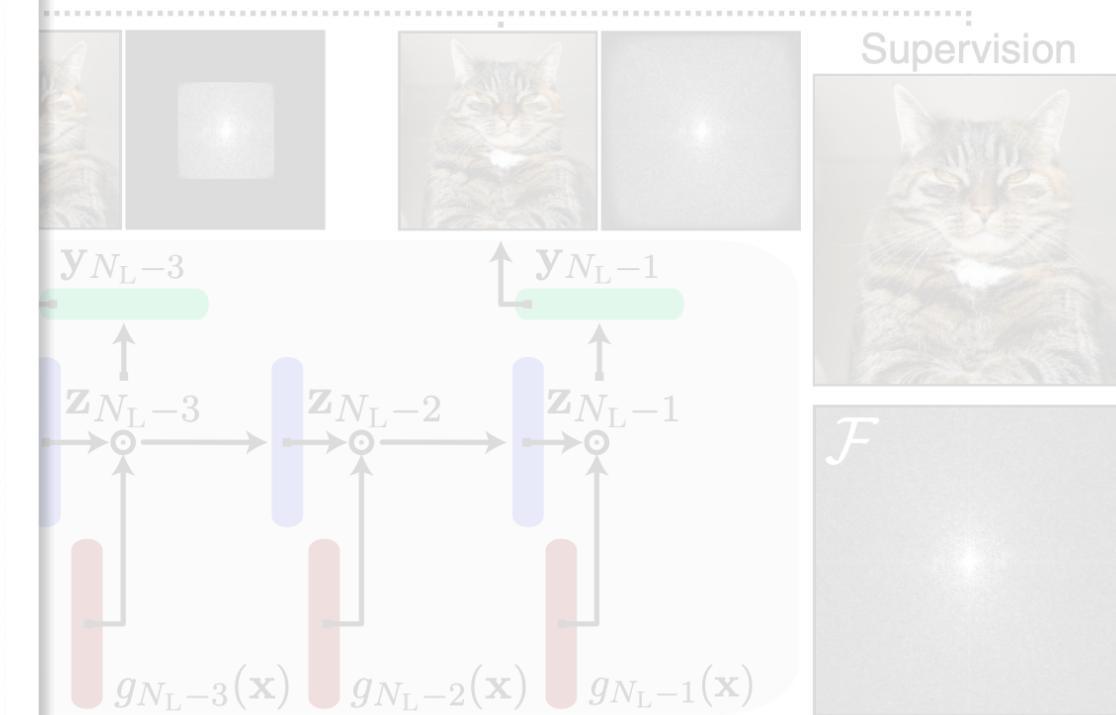


Figure 1: Contribution of this report. Following a survey of over 250 papers, we provide a review of (**Part I**) techniques in neural fields such as prior learning and conditioning, representations, forward maps, architectures, and manipulation, and of (**Part II**) applications in visual computing including 2D image processing, 3D scene reconstruction, generative modeling, digital humans, compression, robotics, and beyond. This report is complemented by a [community-driven website](https://neuralfields.cs.brown.edu/) with search, filtering, bibliographic, and visualization features.

ordinate networks for
representation
CVPR 2022

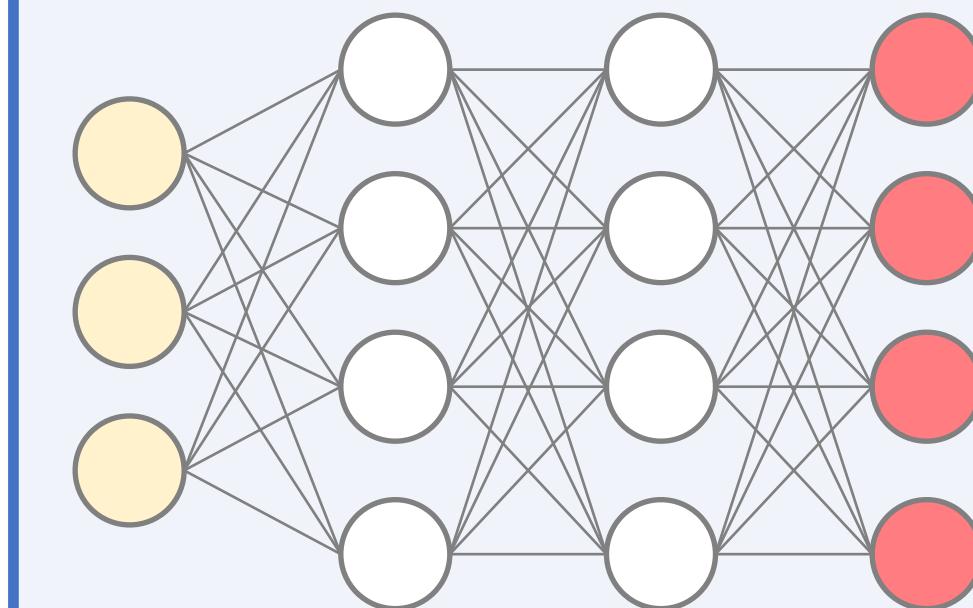


Neural Fields

- Storage memory **does not grow with spatial resolution or number of spatial dimensions.**
- **Slow** sampling: Each query takes a forward pass through the neural net. Means GPU-memory intensive forward passes.
- Does **not** expose locality: Can't identify set of parameters / direction in parameter space that encodes particular spatial location.
- **Inconvenient processing:** cannot run convolutions.
- But: automatic adaptive resolution. Over optimization, will converge to assign more compute to higher-frequency areas.

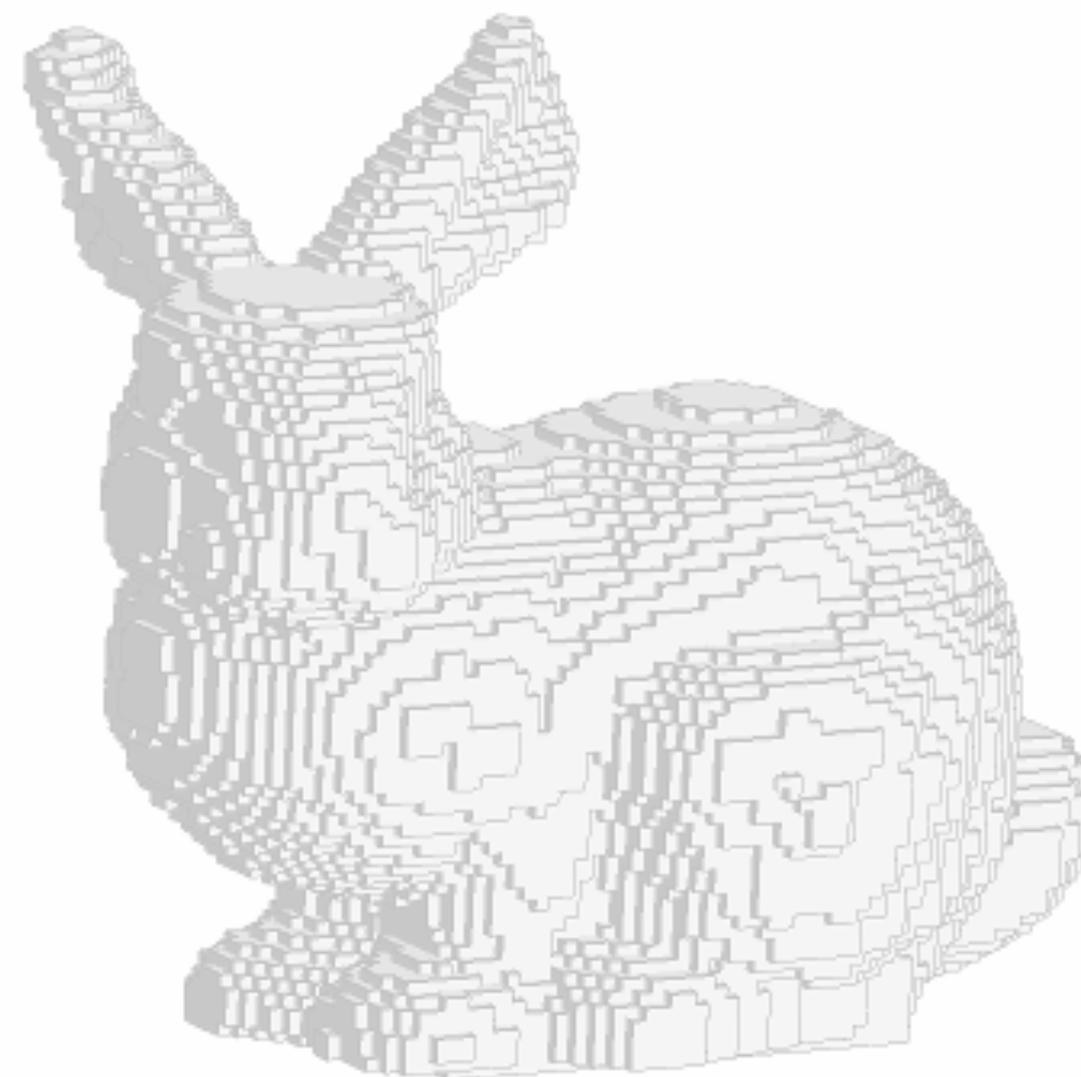
Neural Fields

$$\Phi : \mathbb{R}^3 \rightarrow \mathbb{R}^n$$

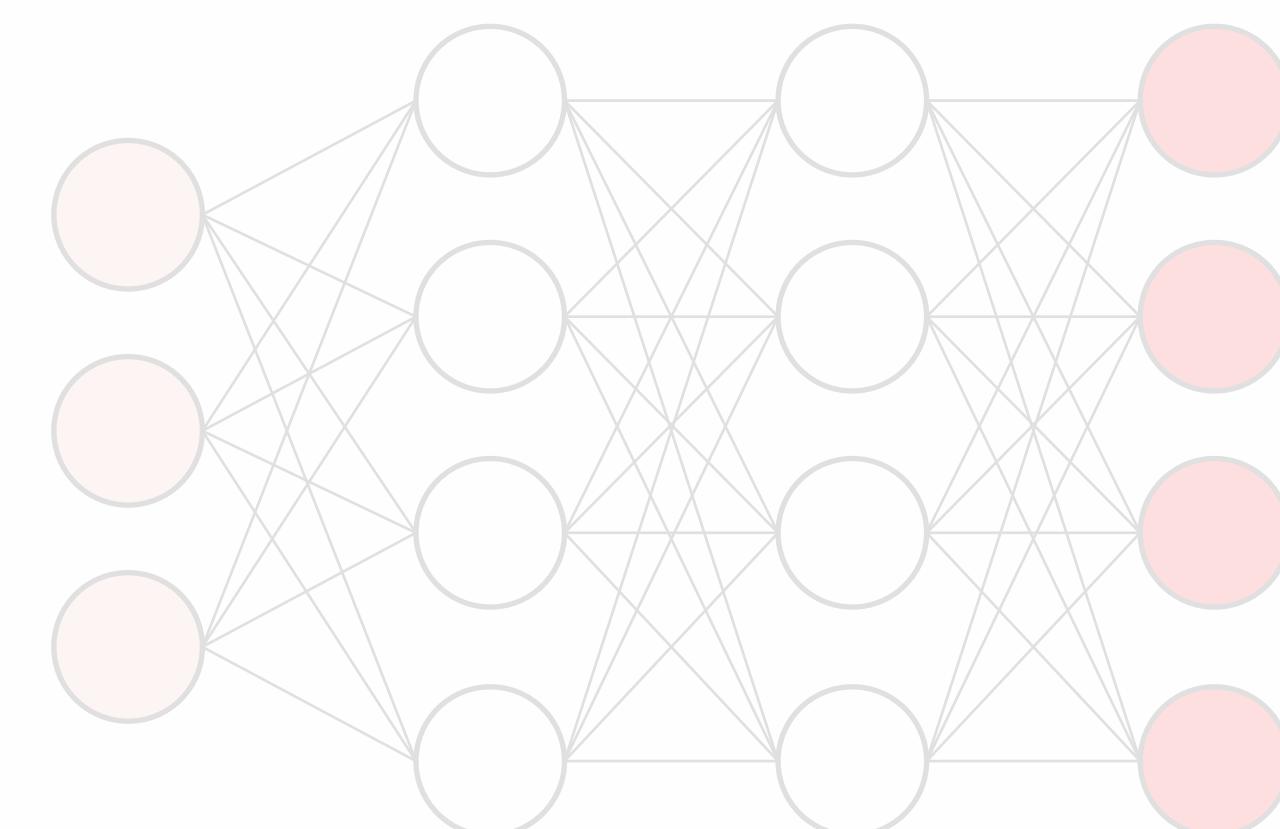


Note: Despite their name, has nothing to do with Machine Learning itself.
It's as “smart” as a voxelgrid.

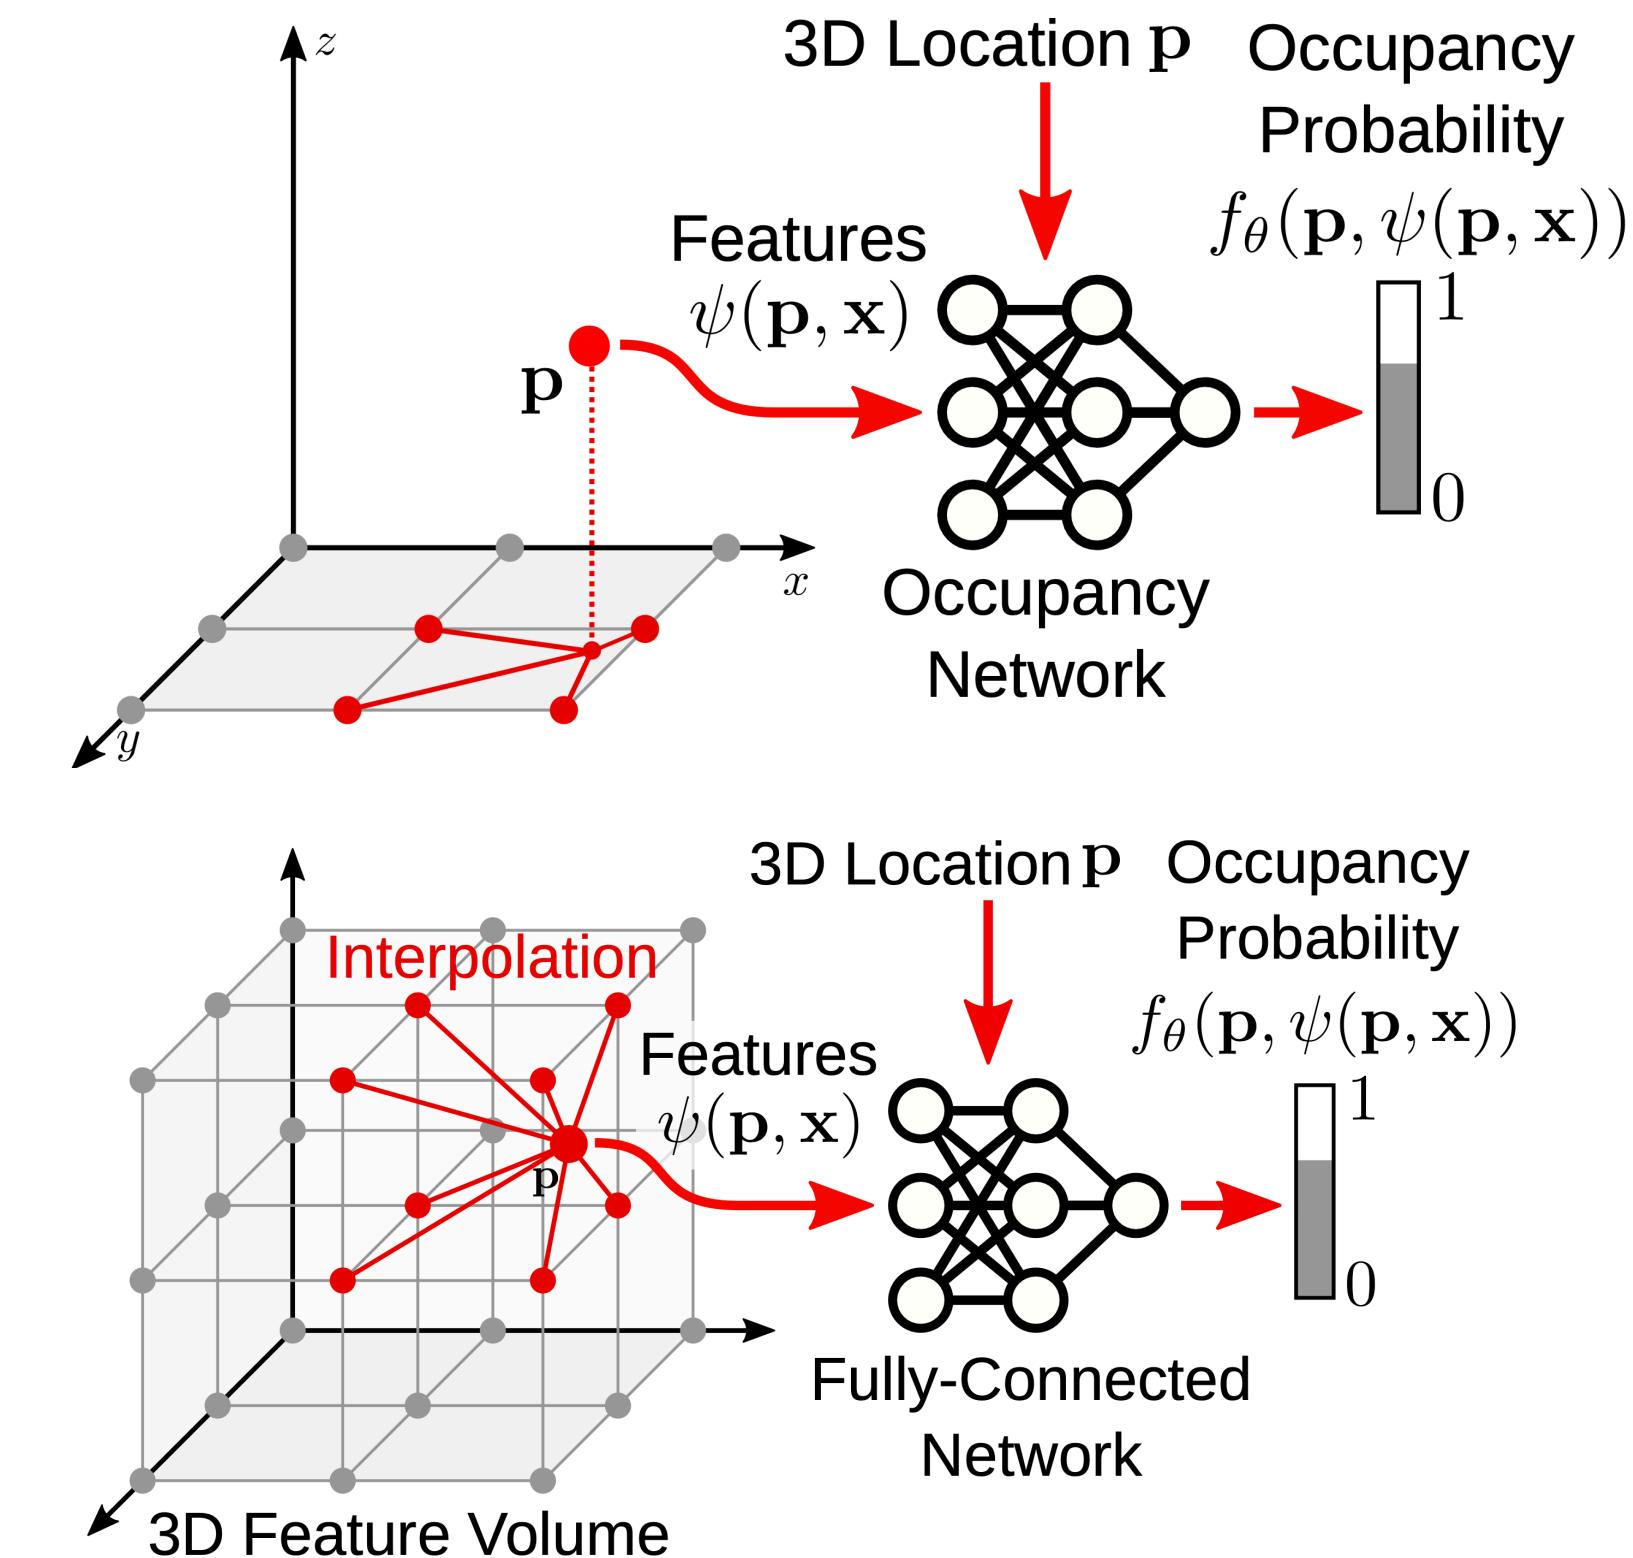
Hybrid discrete / continuous reps



Discrete Parameterizations

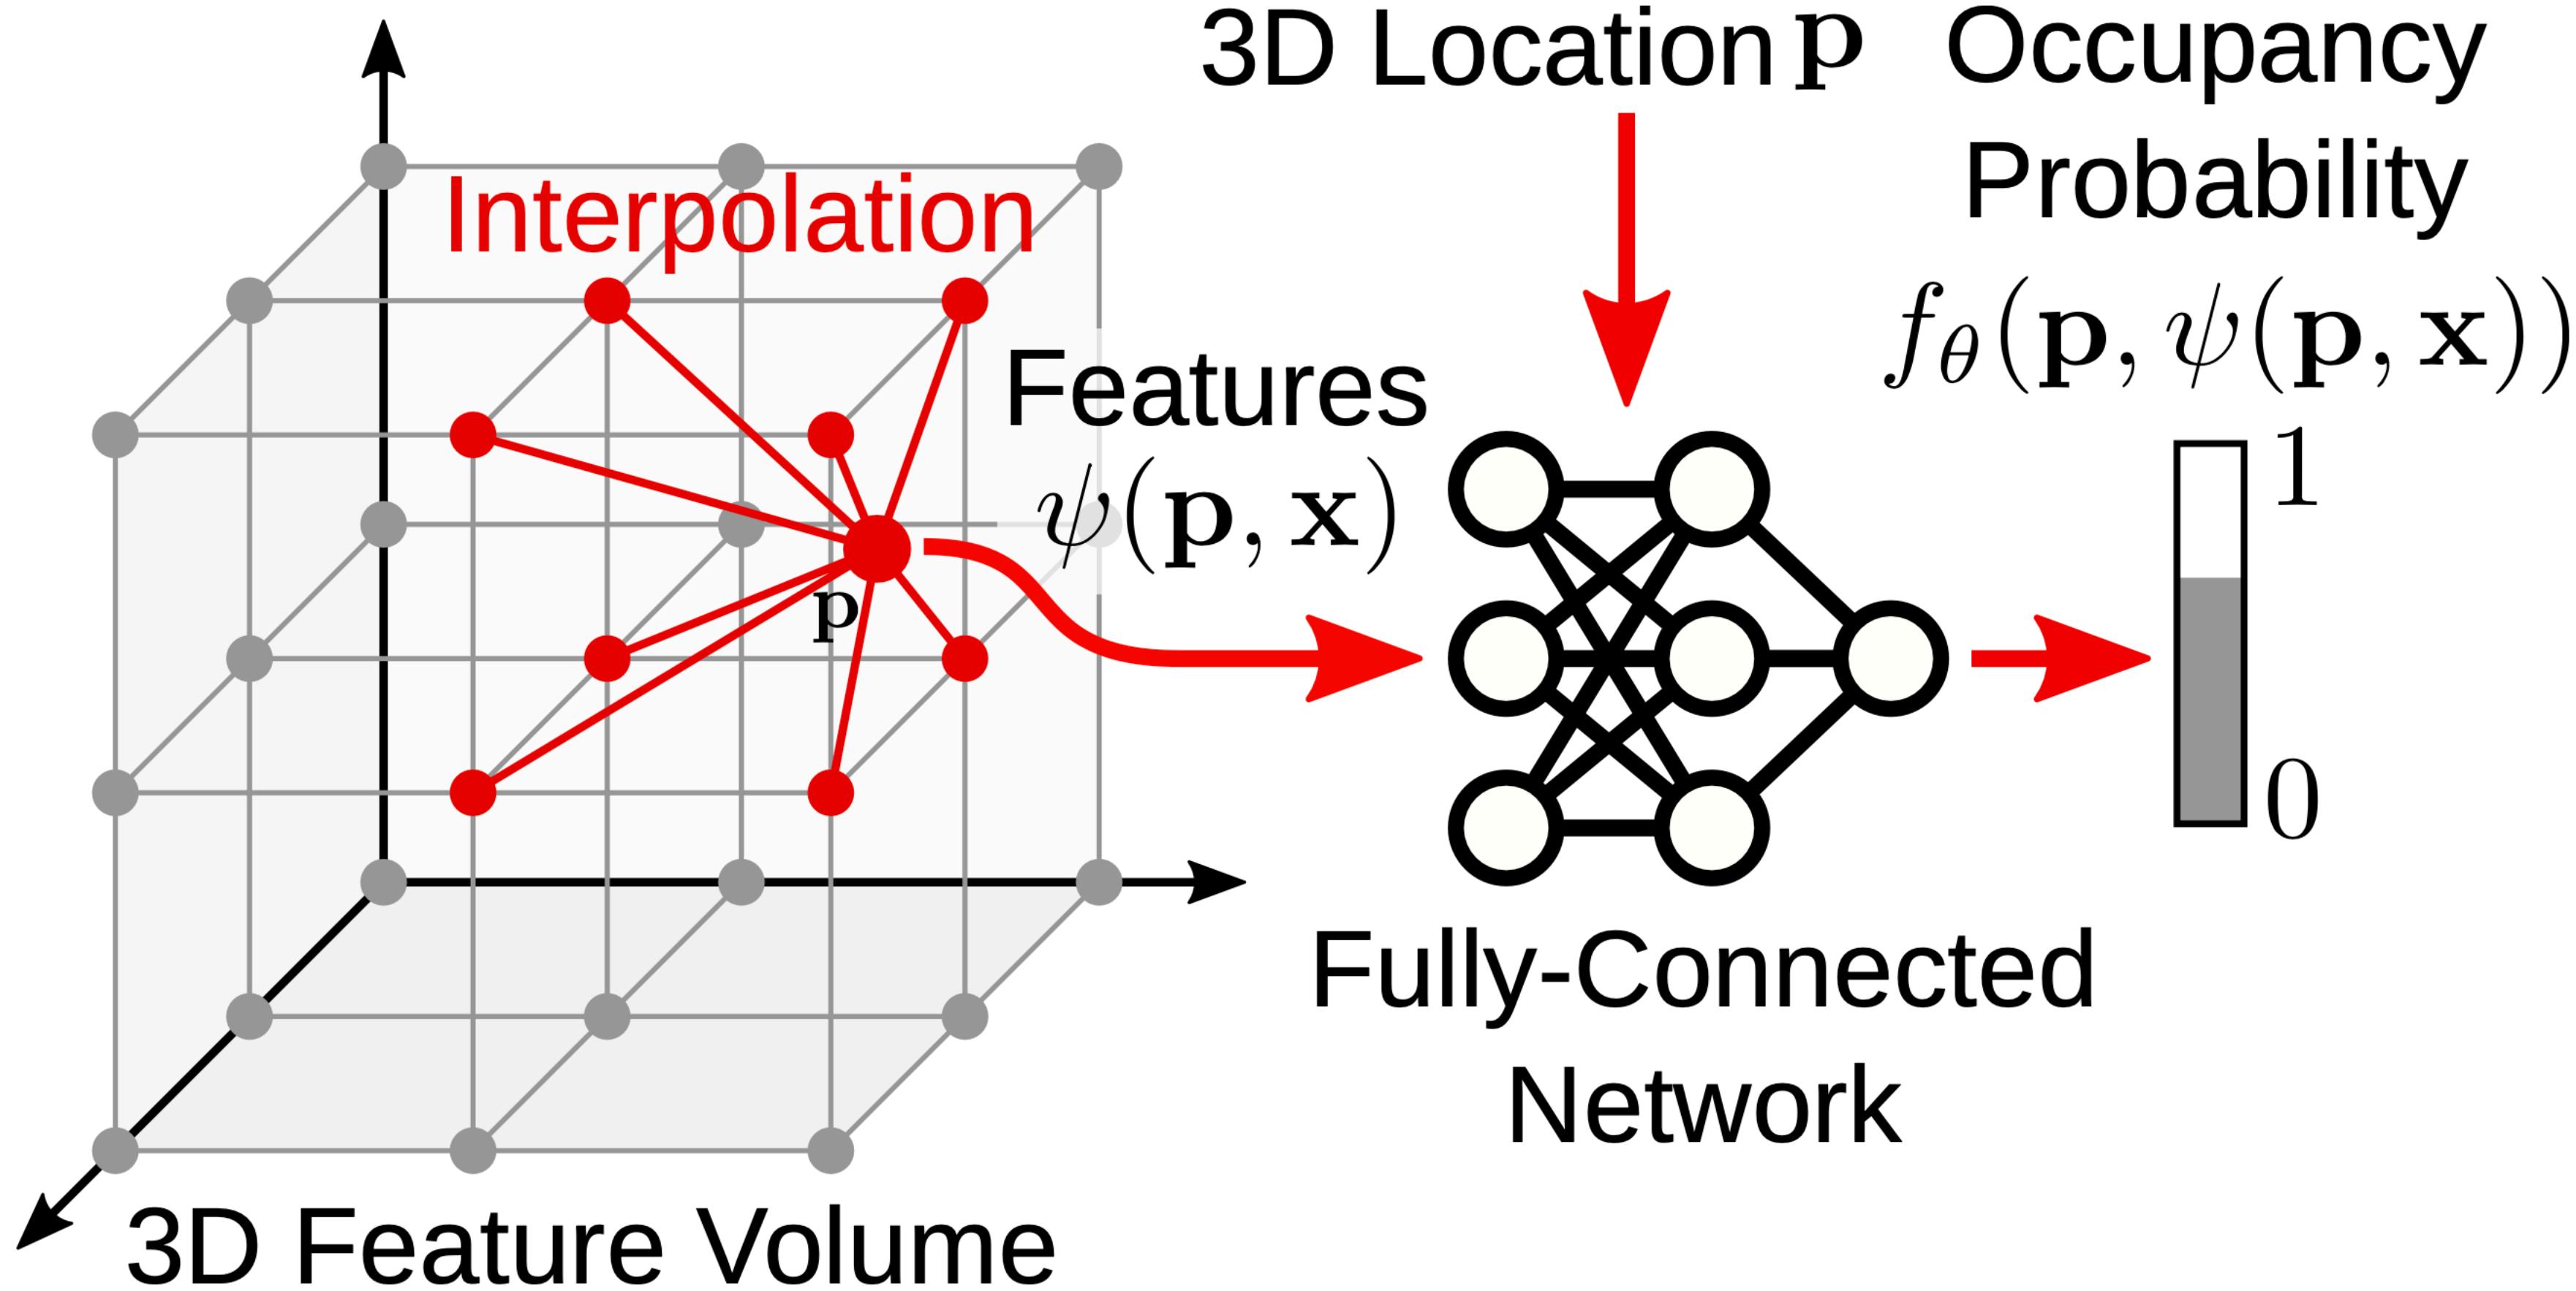


Continuous Parameterizations



Hybrid Parameterizations

Basic idea: Use Neural Net as Interpolation Kernel



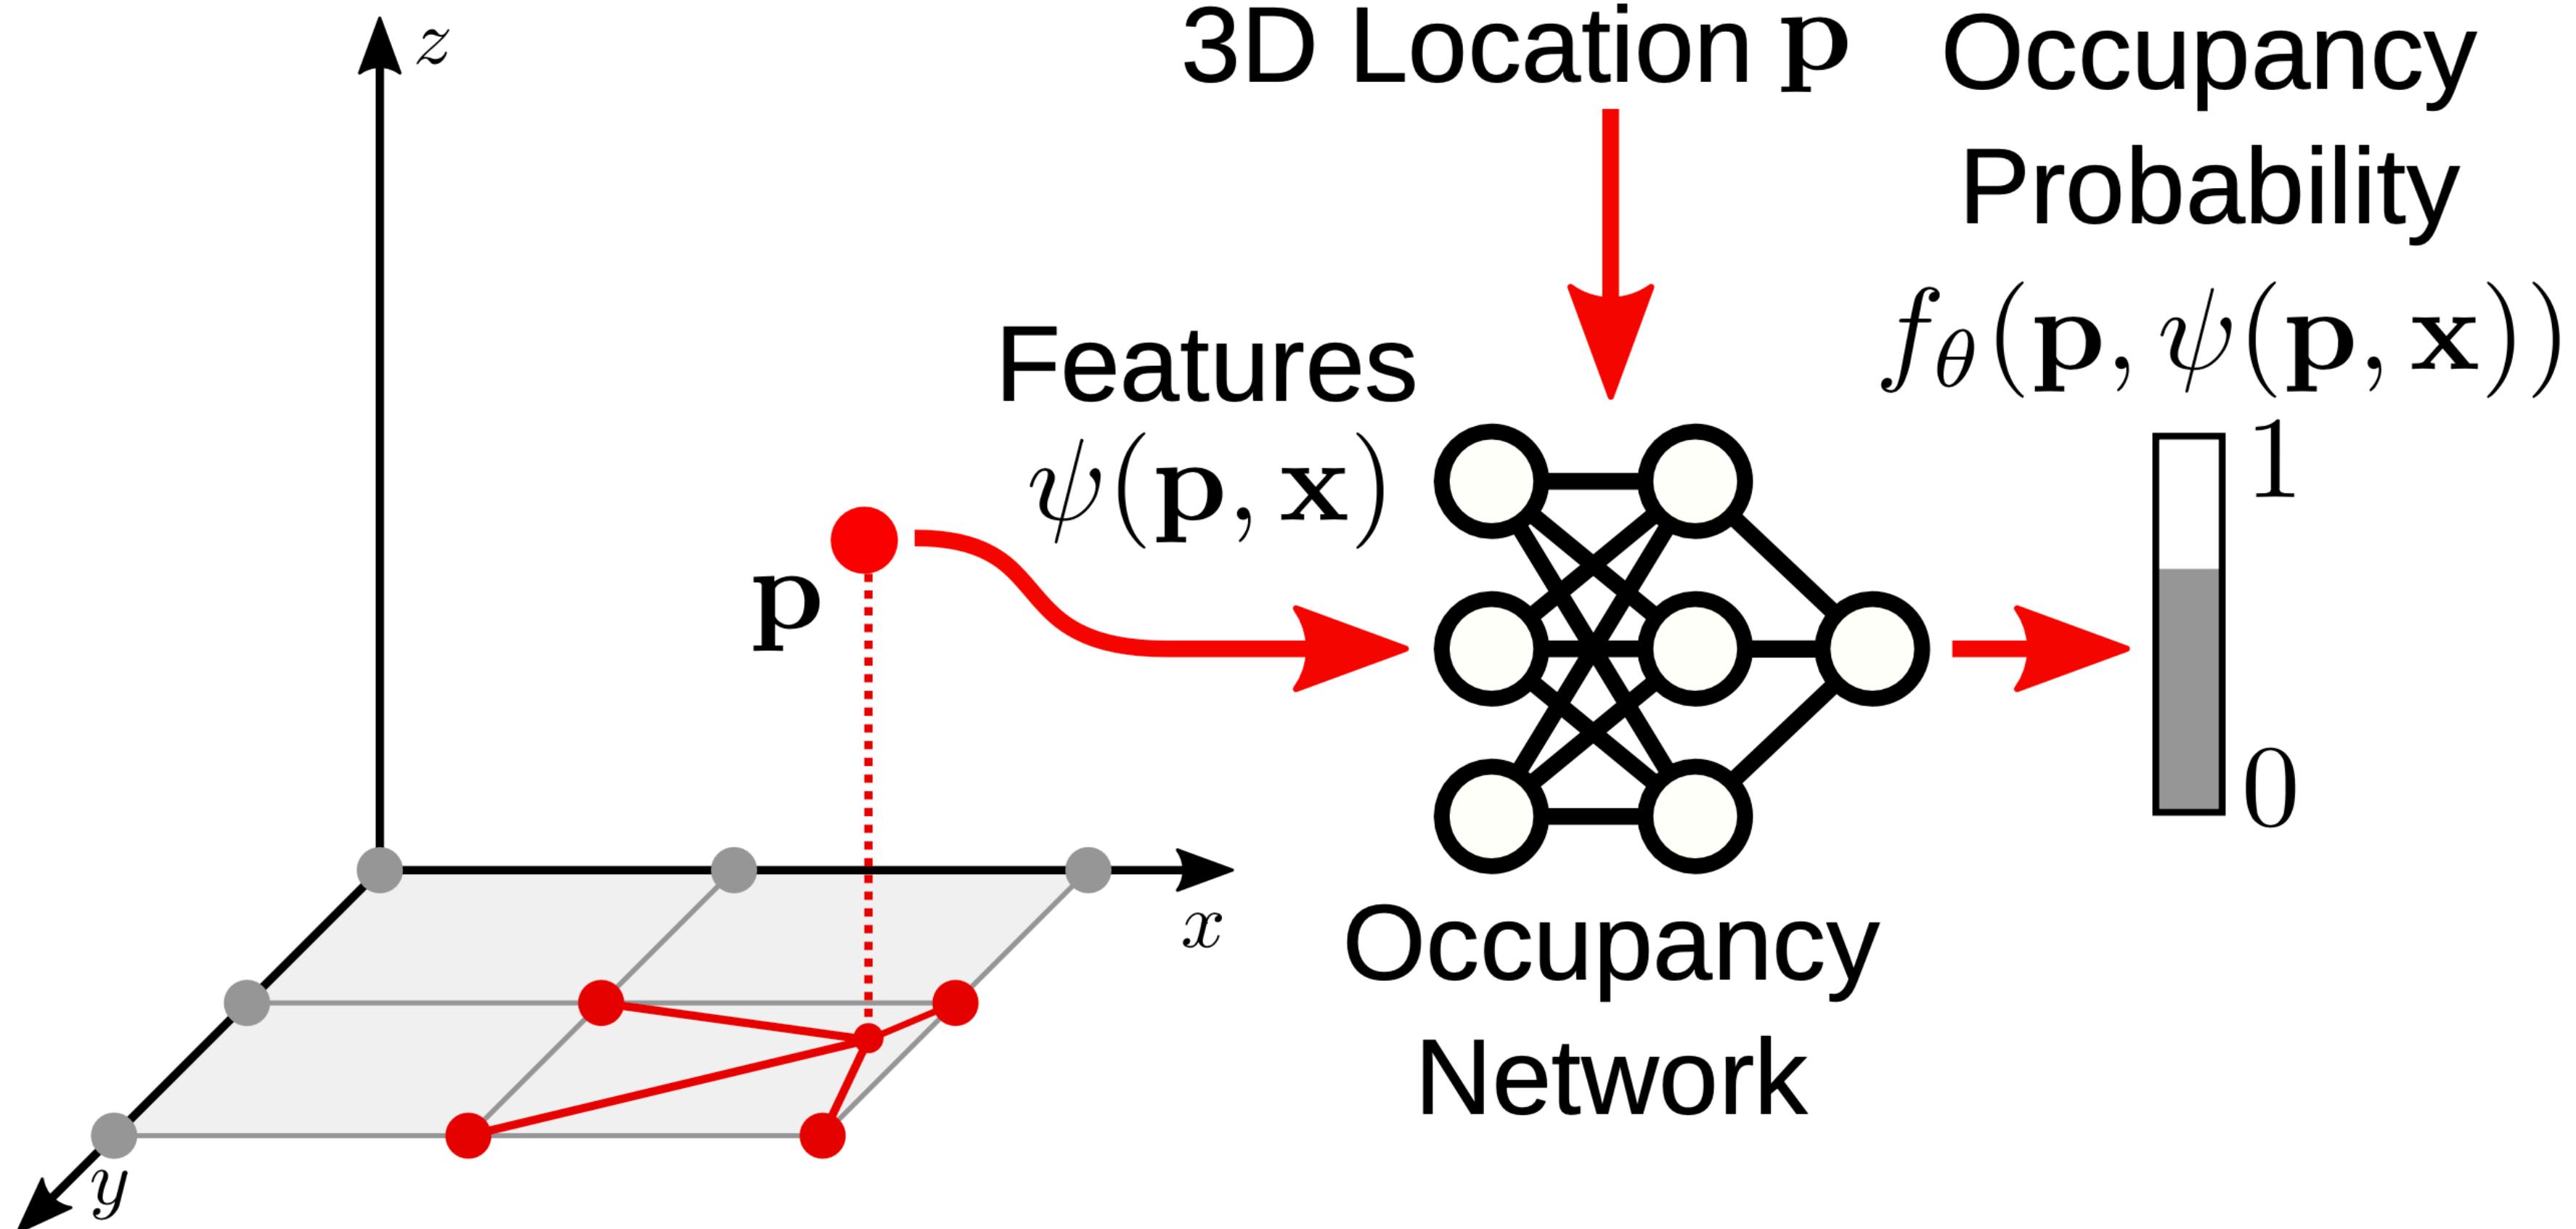
Convolutional Occupancy Networks [Peng et al. 2020]

Local Implicit Grid Representations for 3D Scenes [Jiang et al. 2020]

Implicit Functions in Feature Space for 3D Shape Reconstruction and Completion [Chibane et al. 2020]

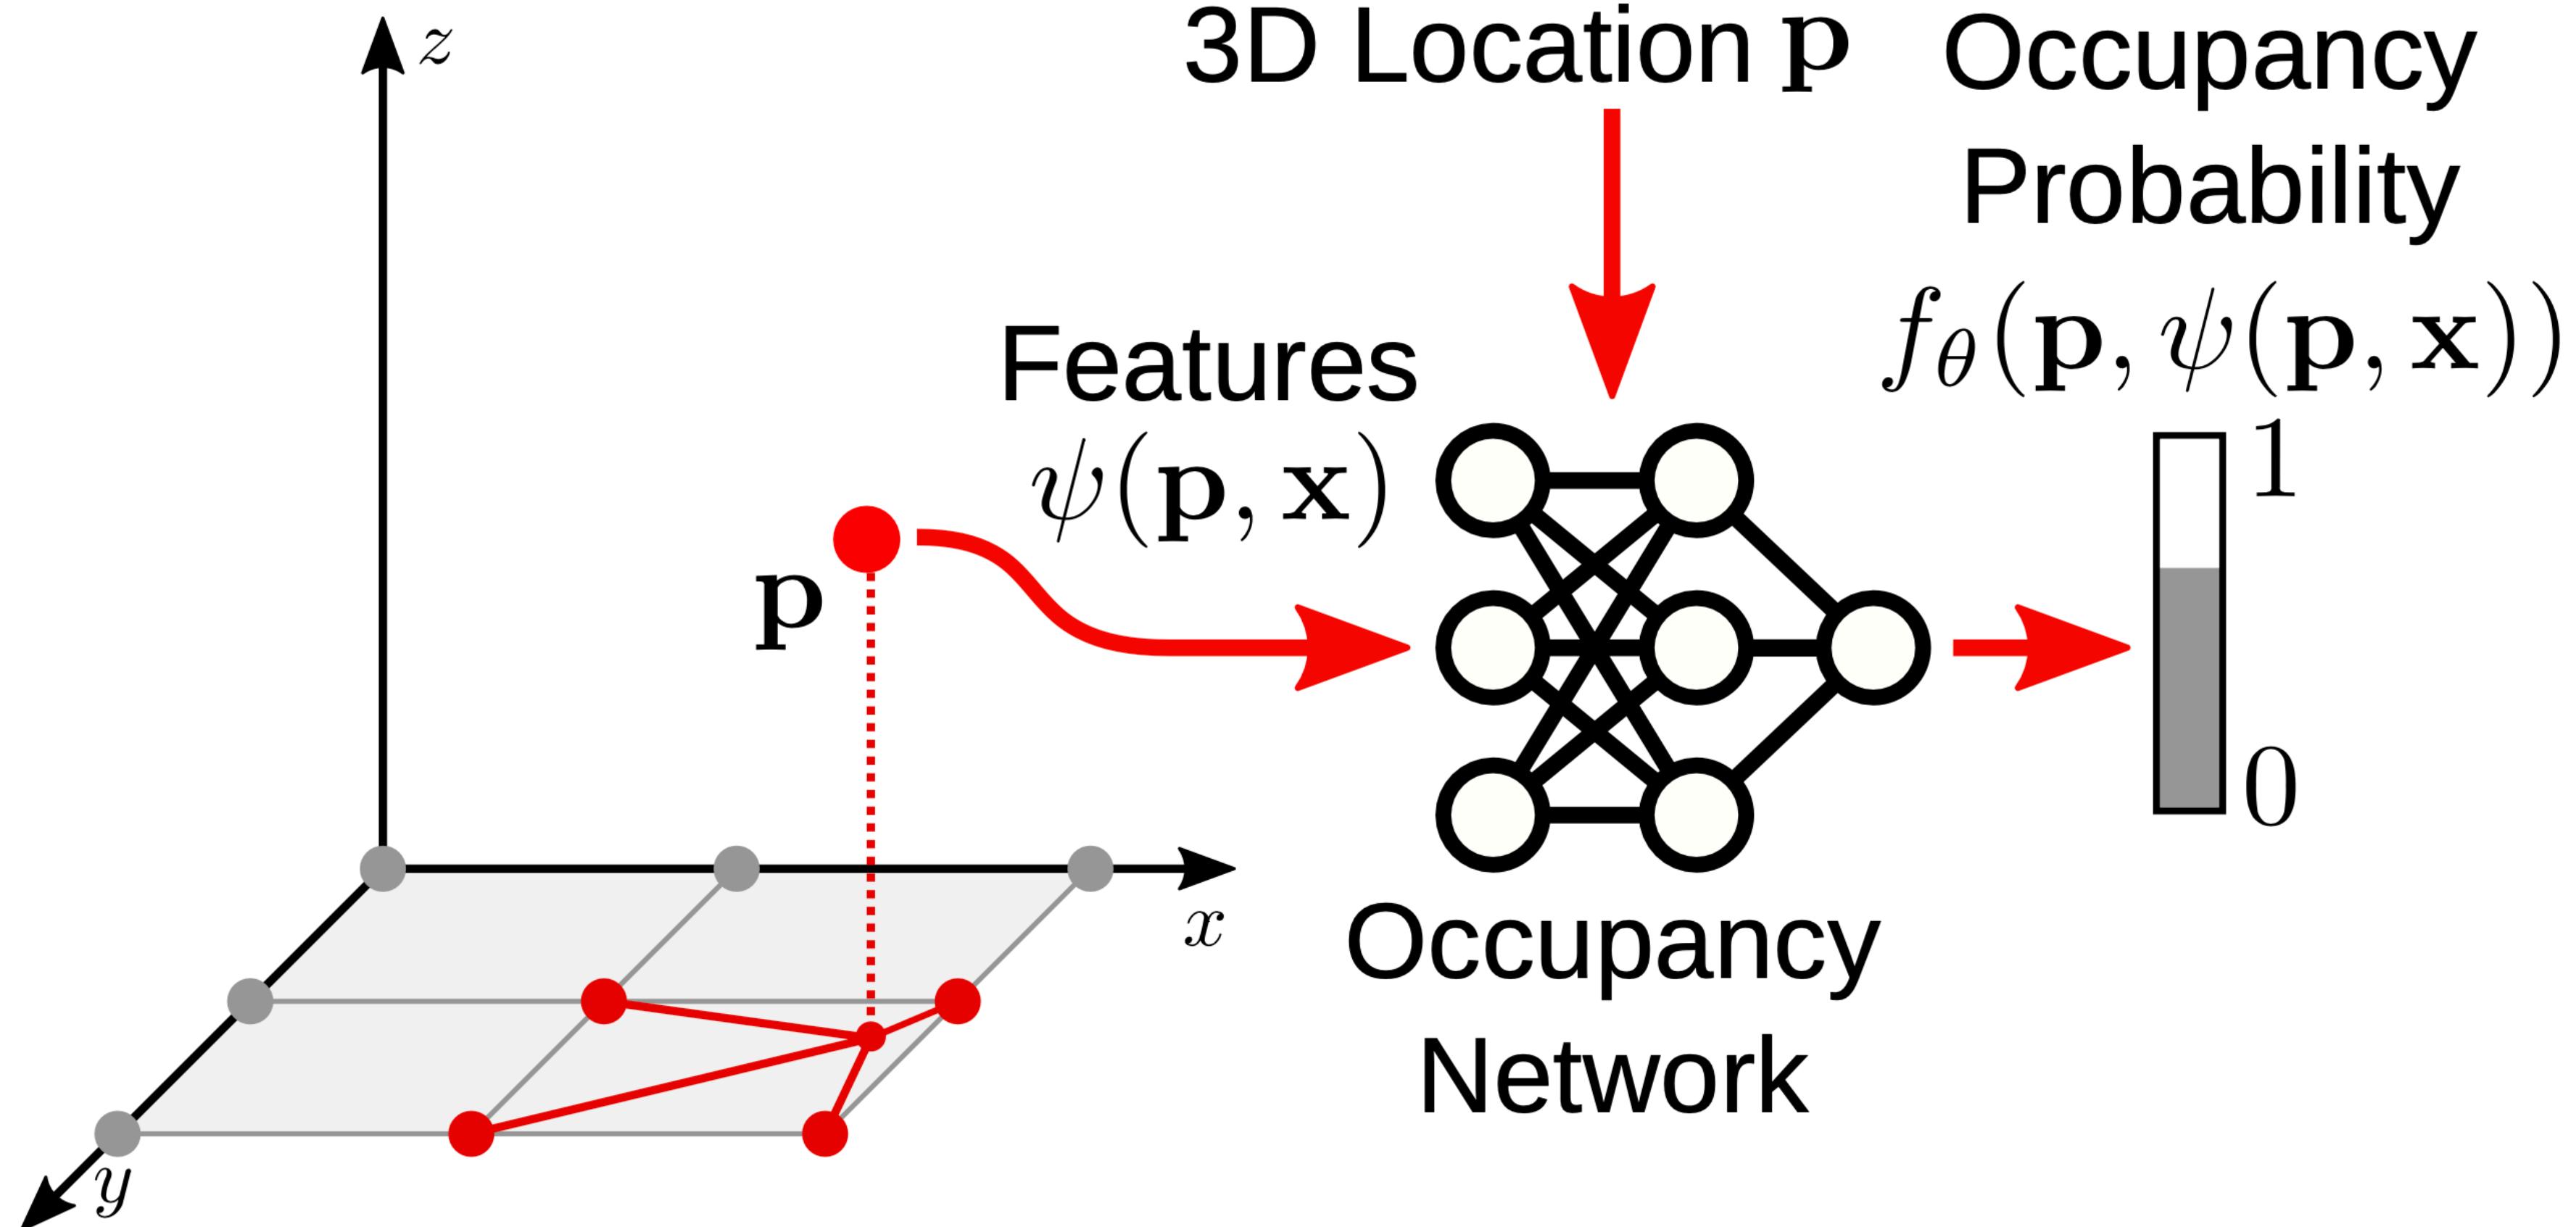
Deep Local Shapes: Learning Local SDF Priors for Detailed 3D Reconstruction [Chabra et al. 2020]

Ground plan & Orthographic Projection



Any assumptions made on scene? Any limitations?

Ground plan & Orthographic Projection



No. You can still represent *any* scene - some info stored in NN, some in grid.

Tri-Planes and Orthographic Projection

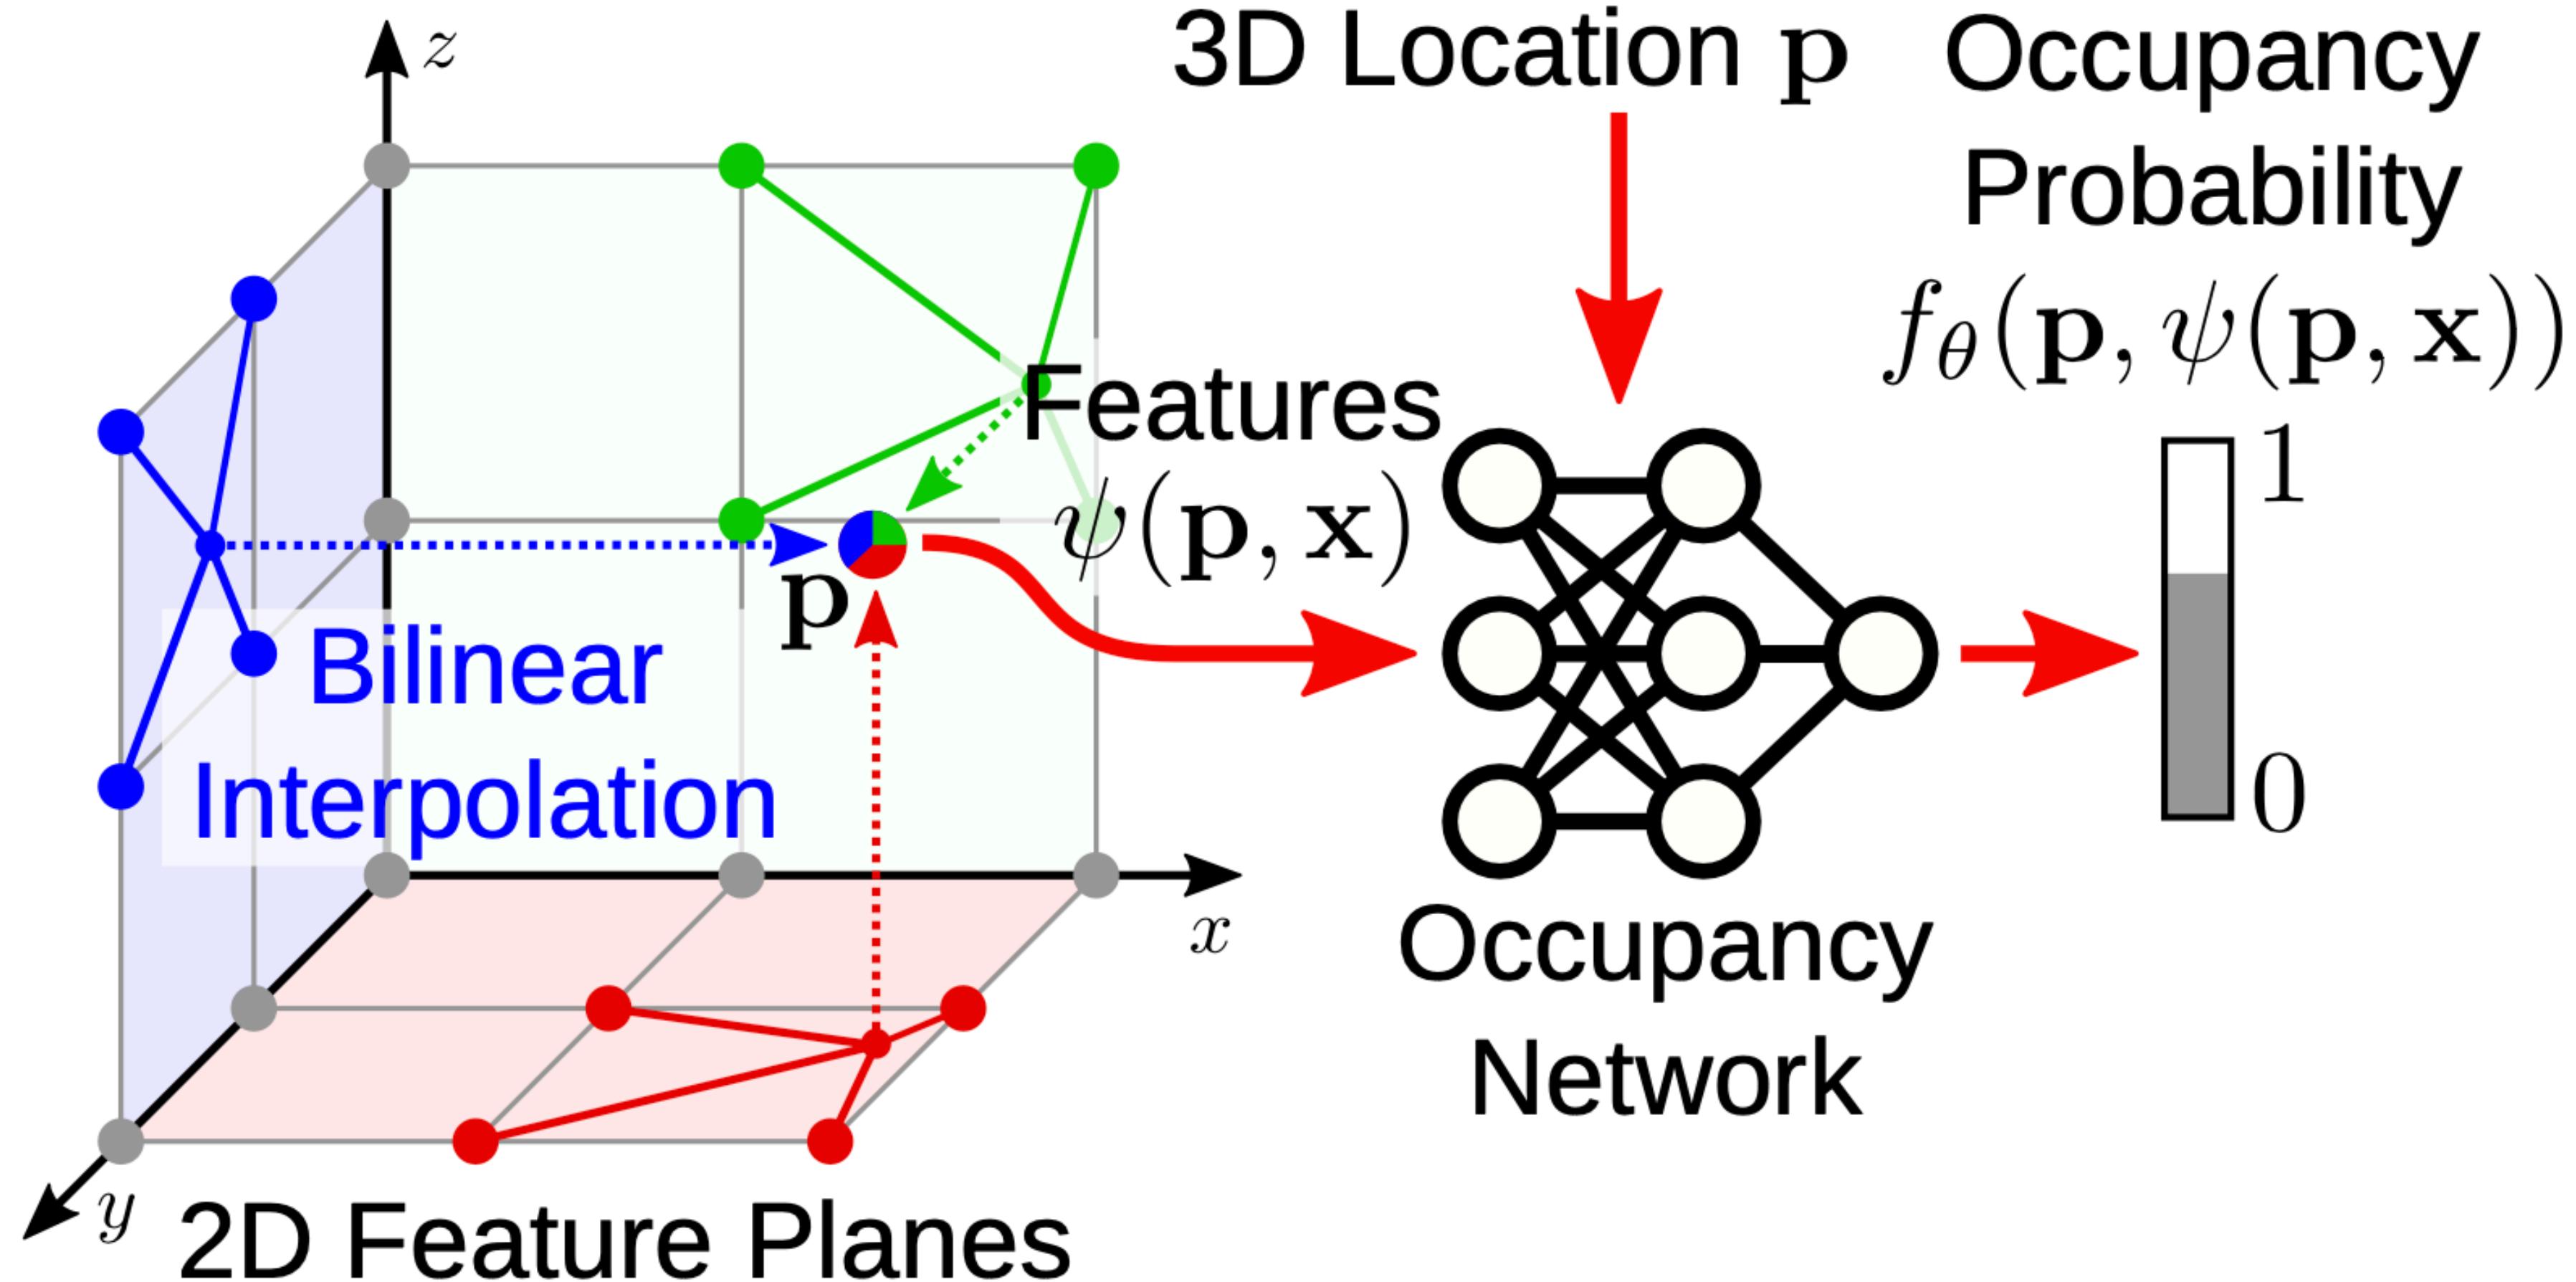
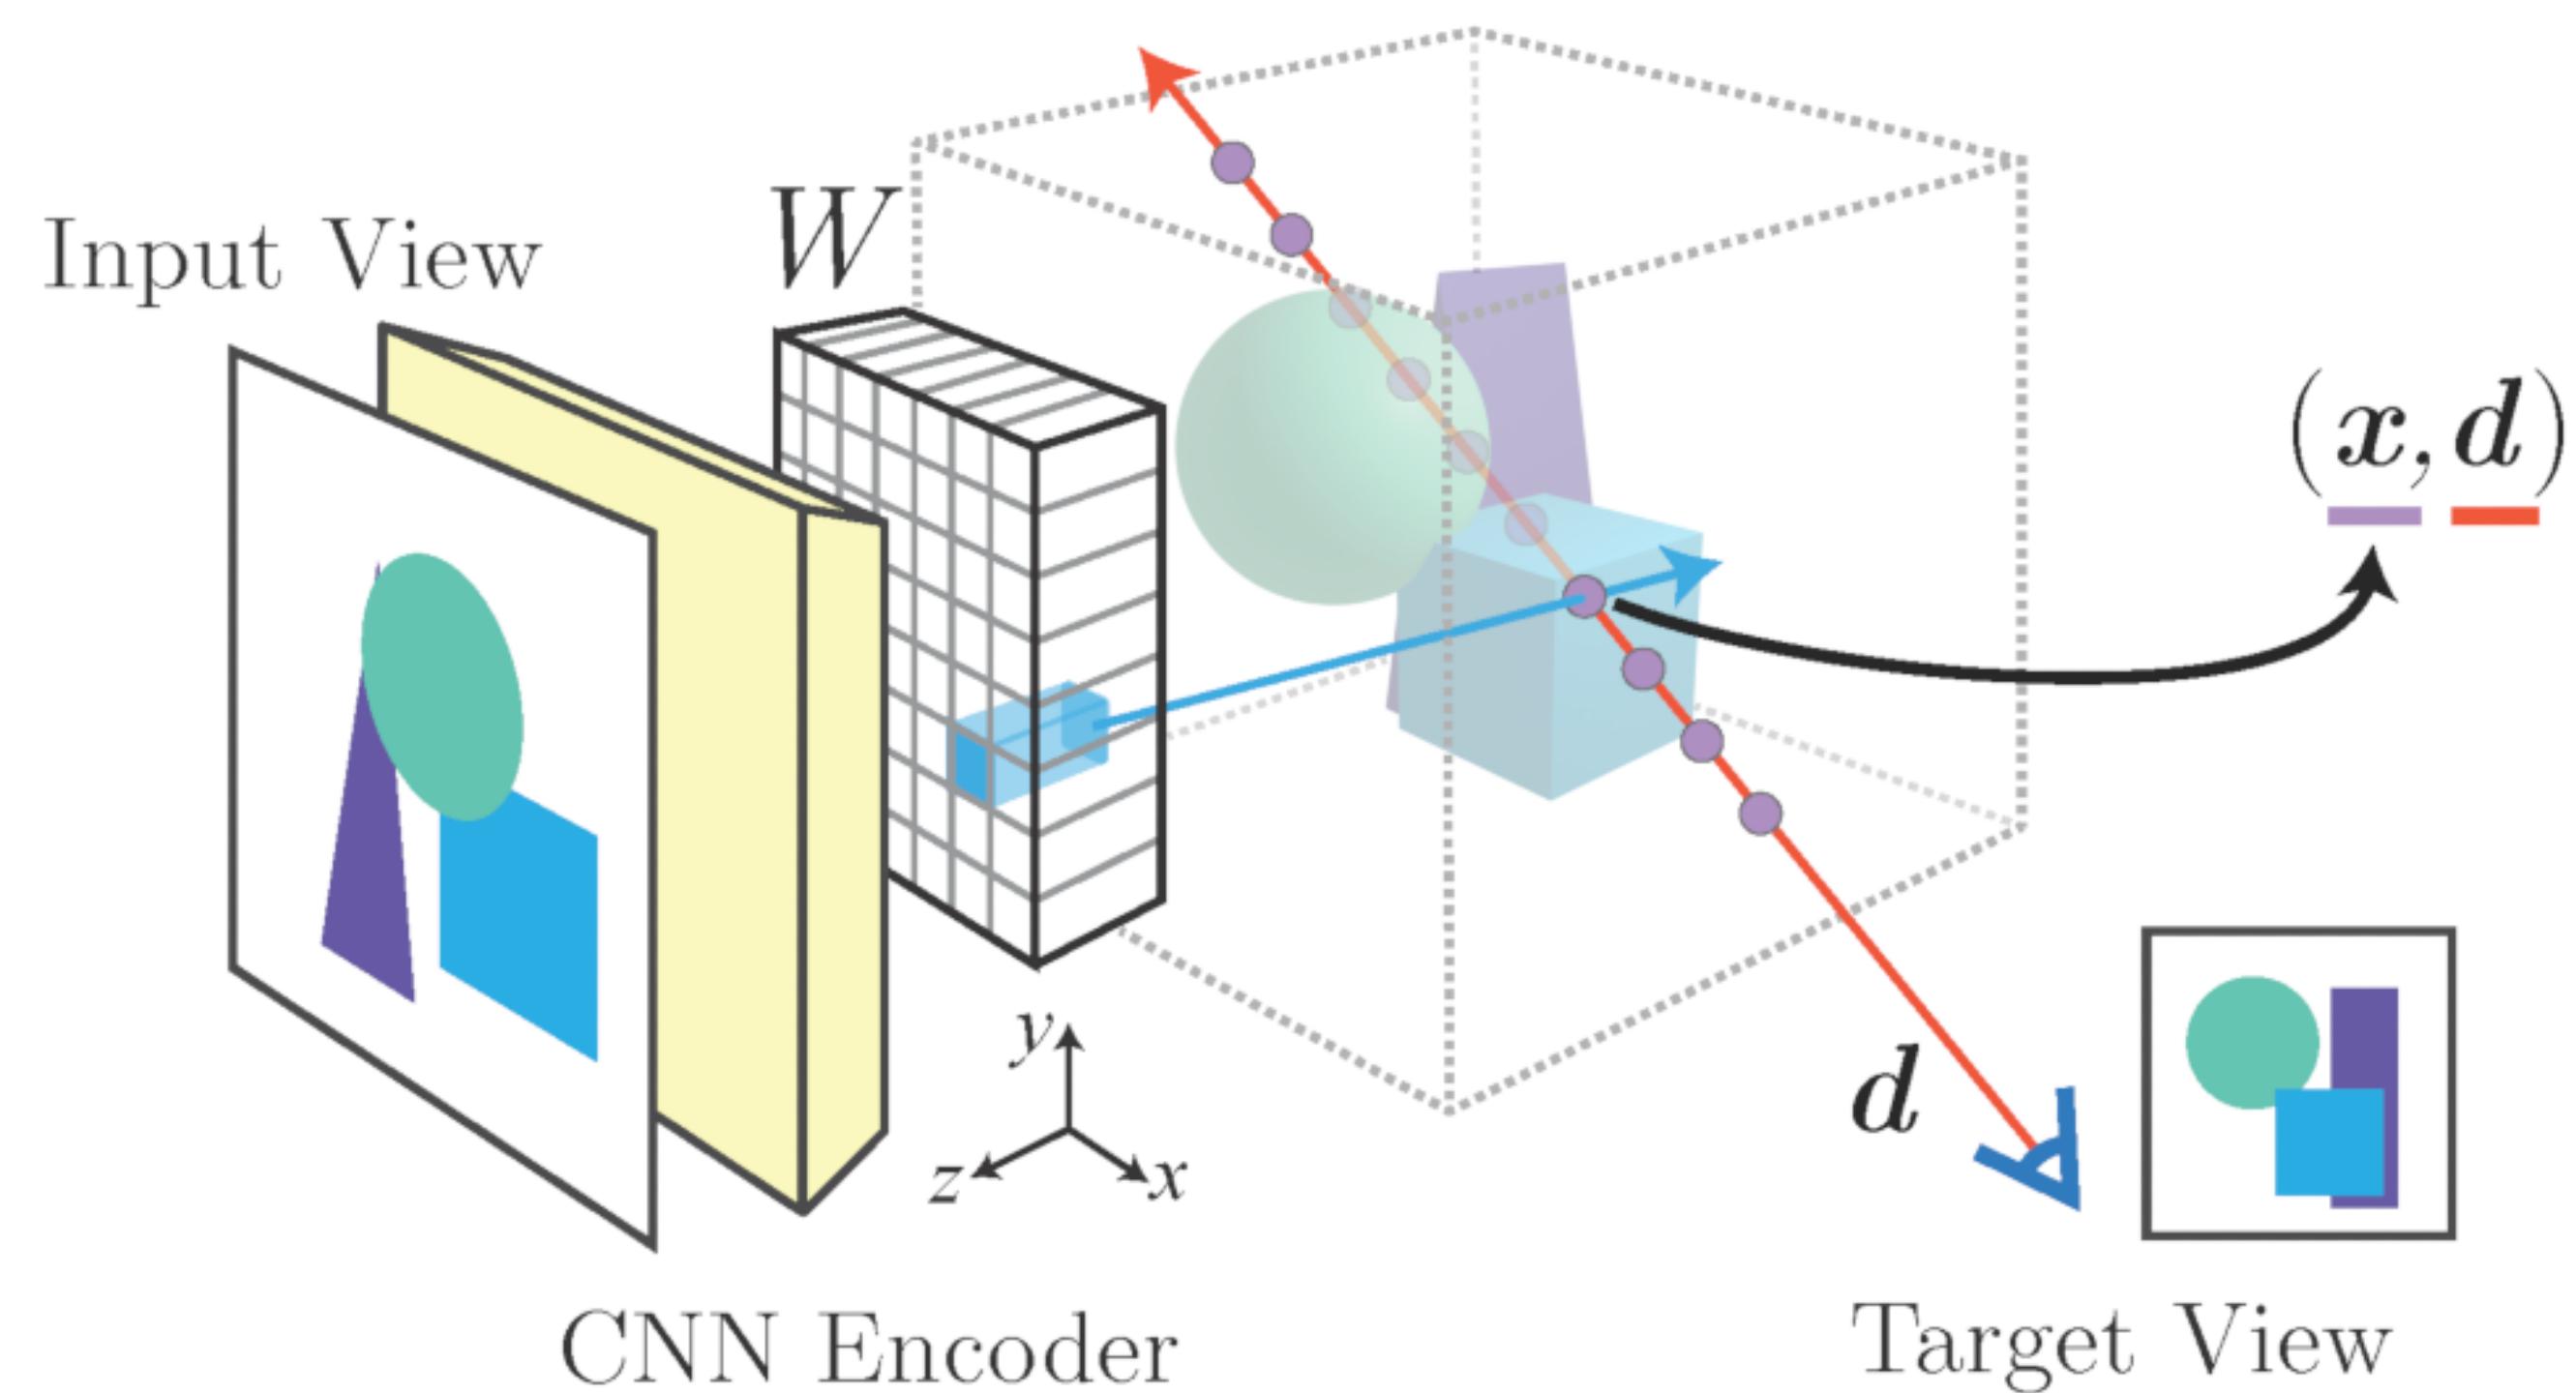
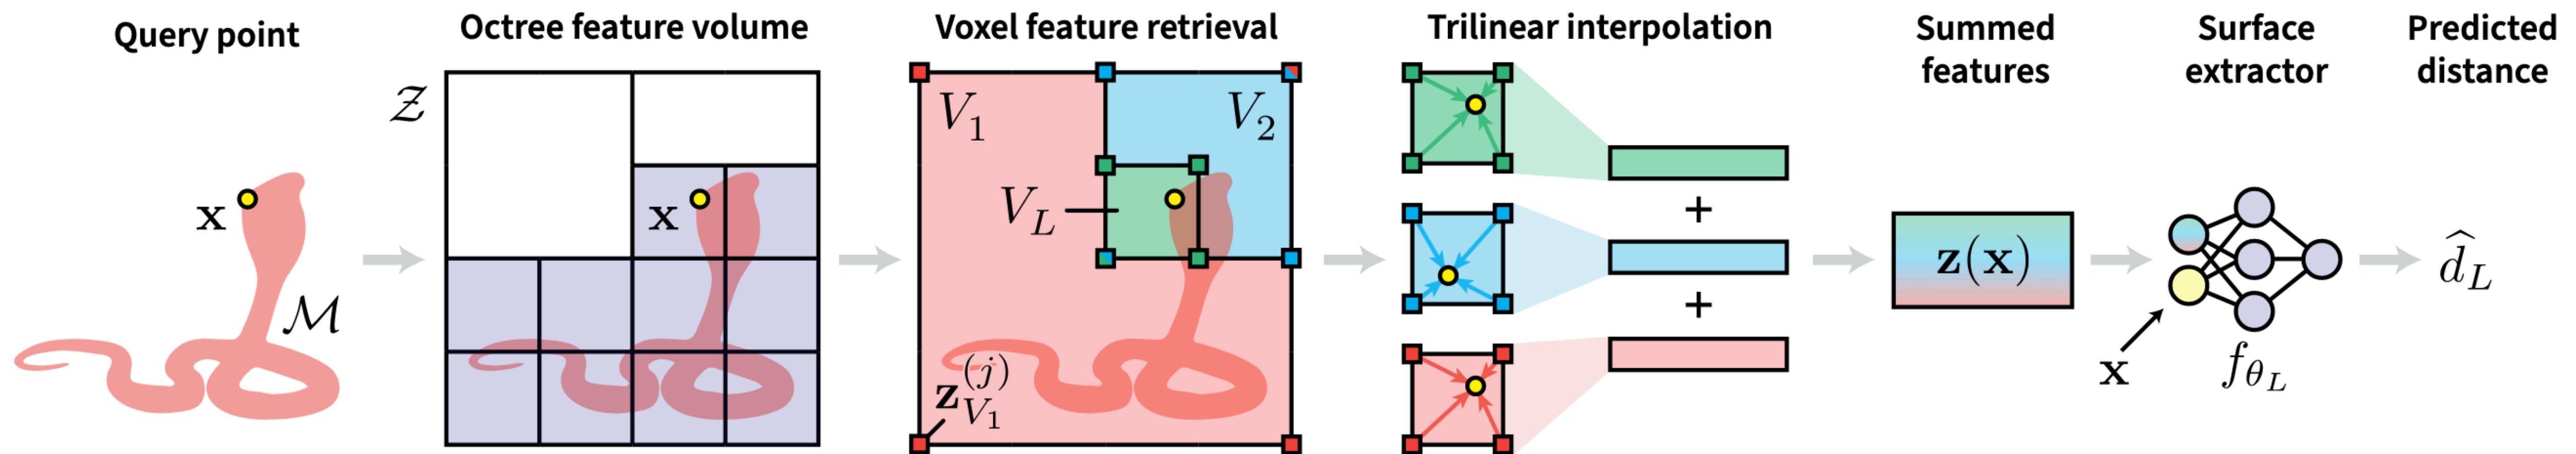


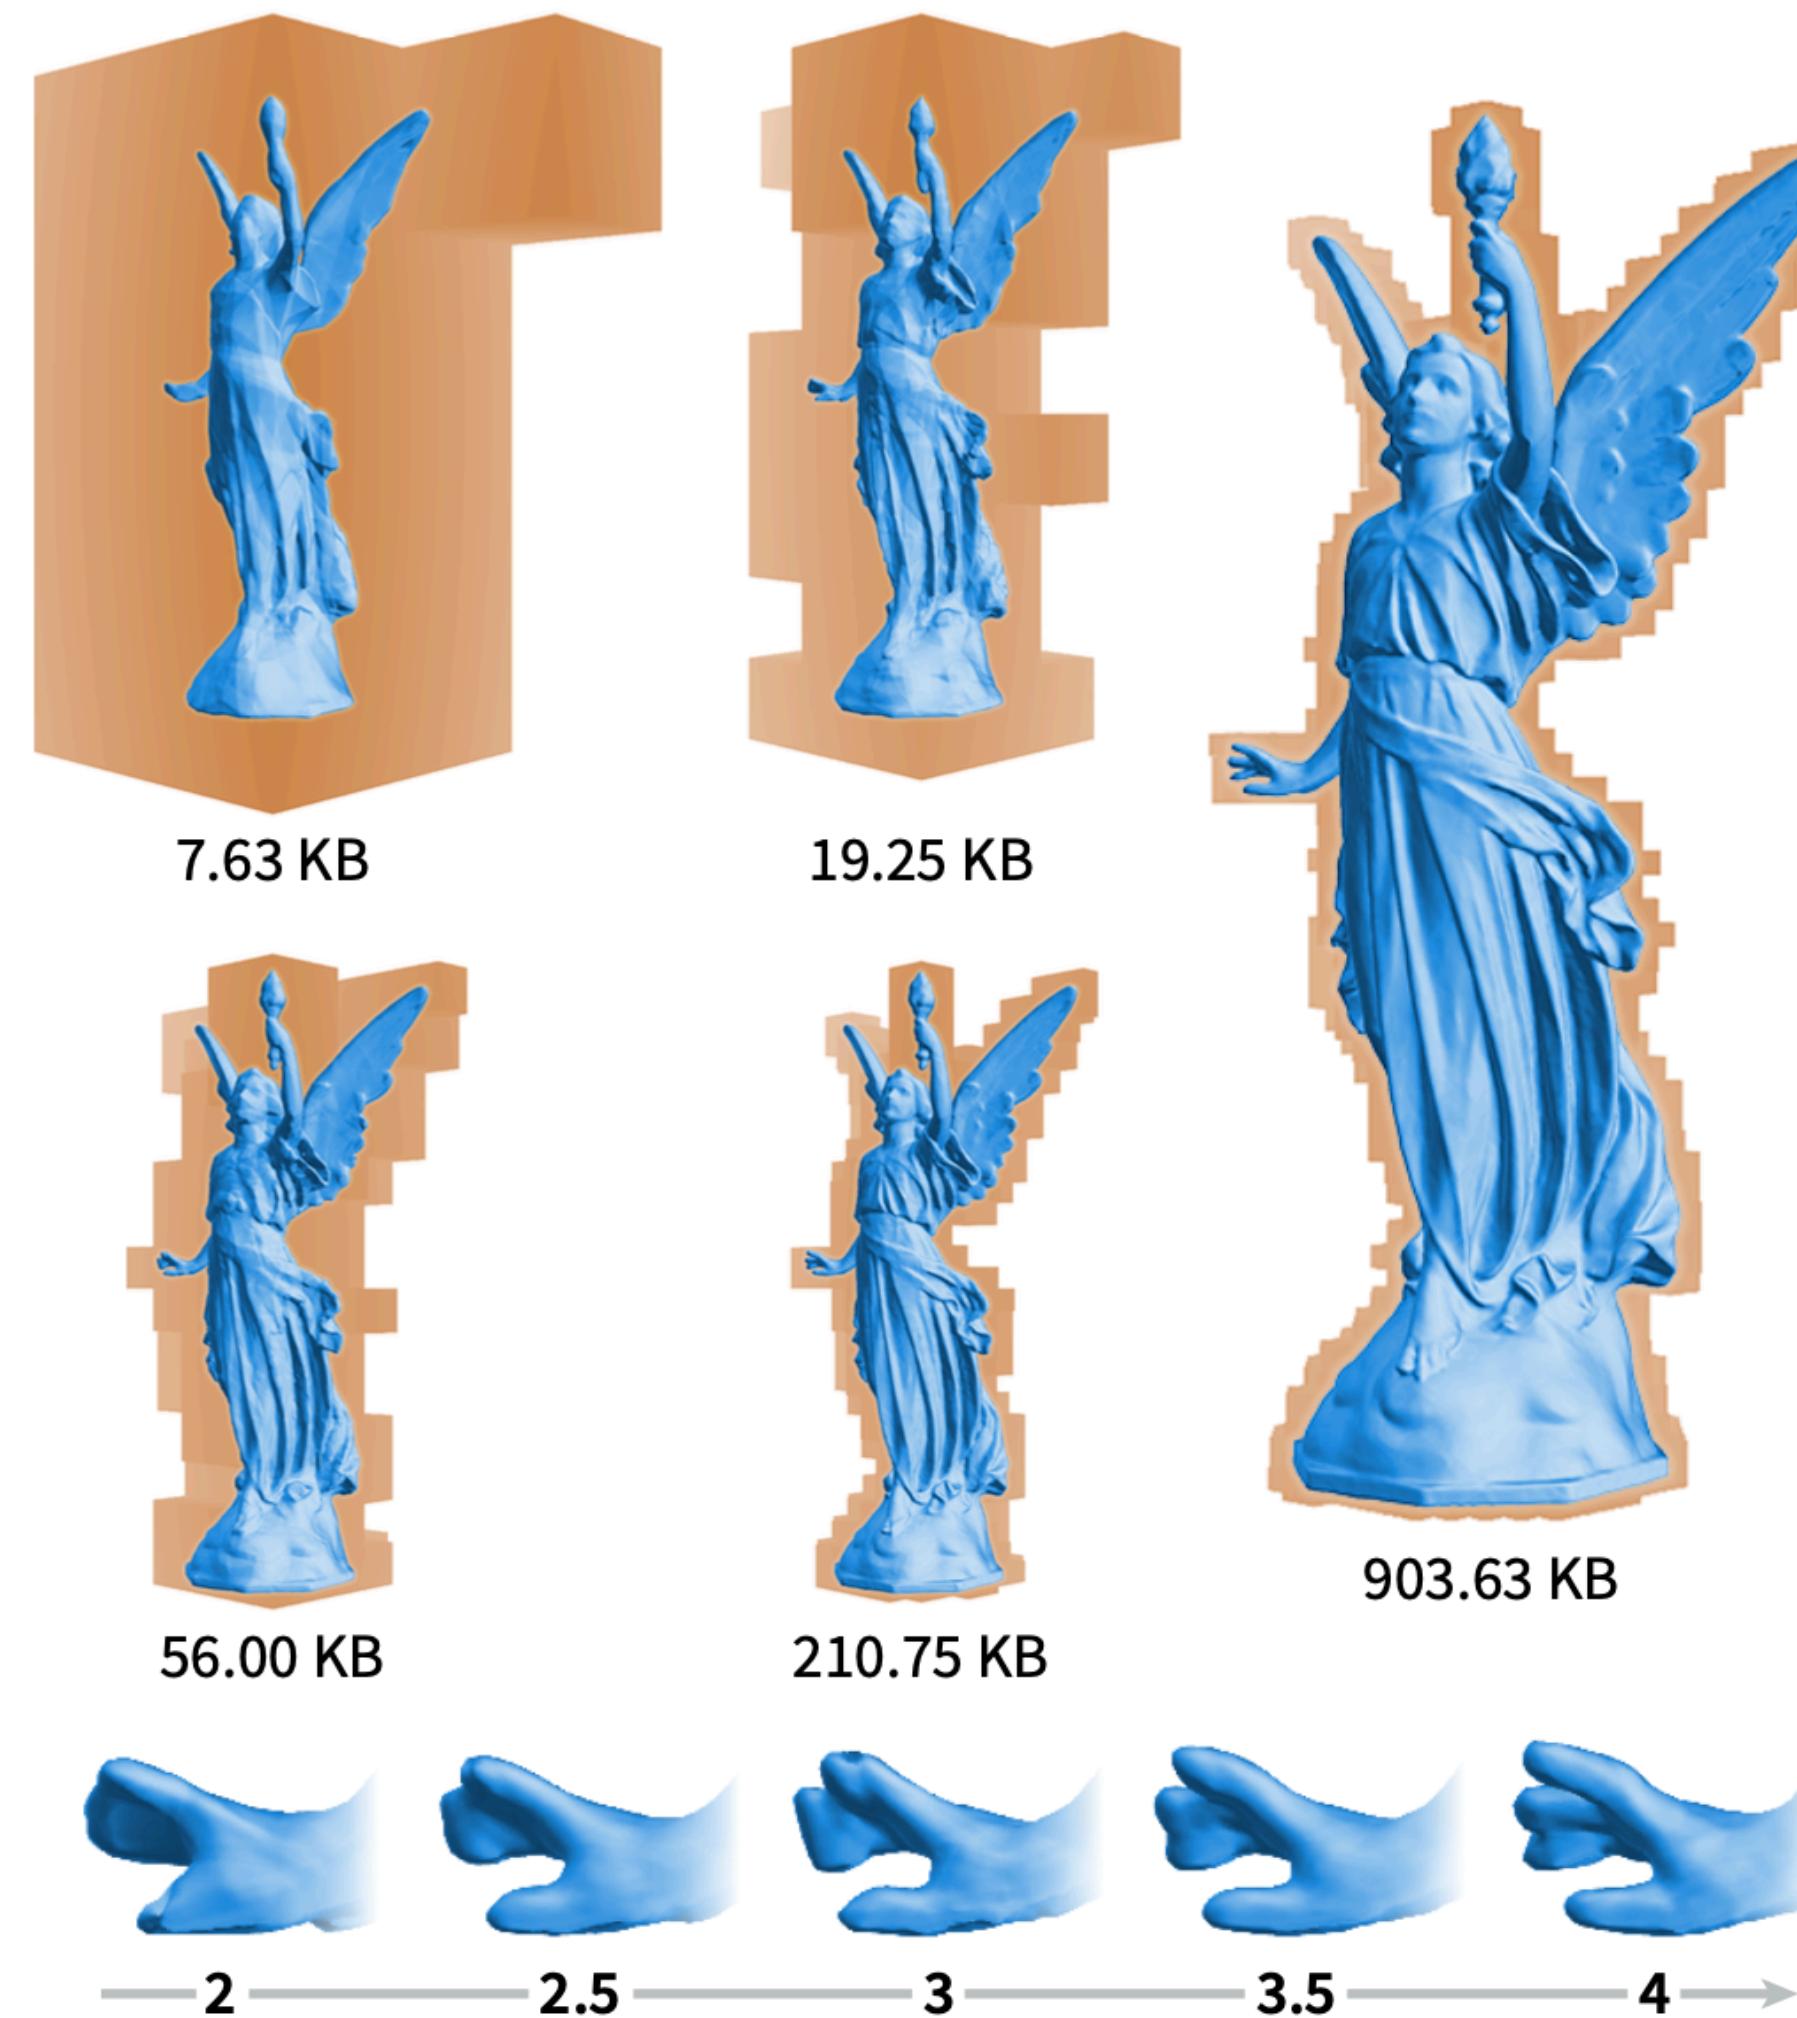
Image Grids & Perspective Projection



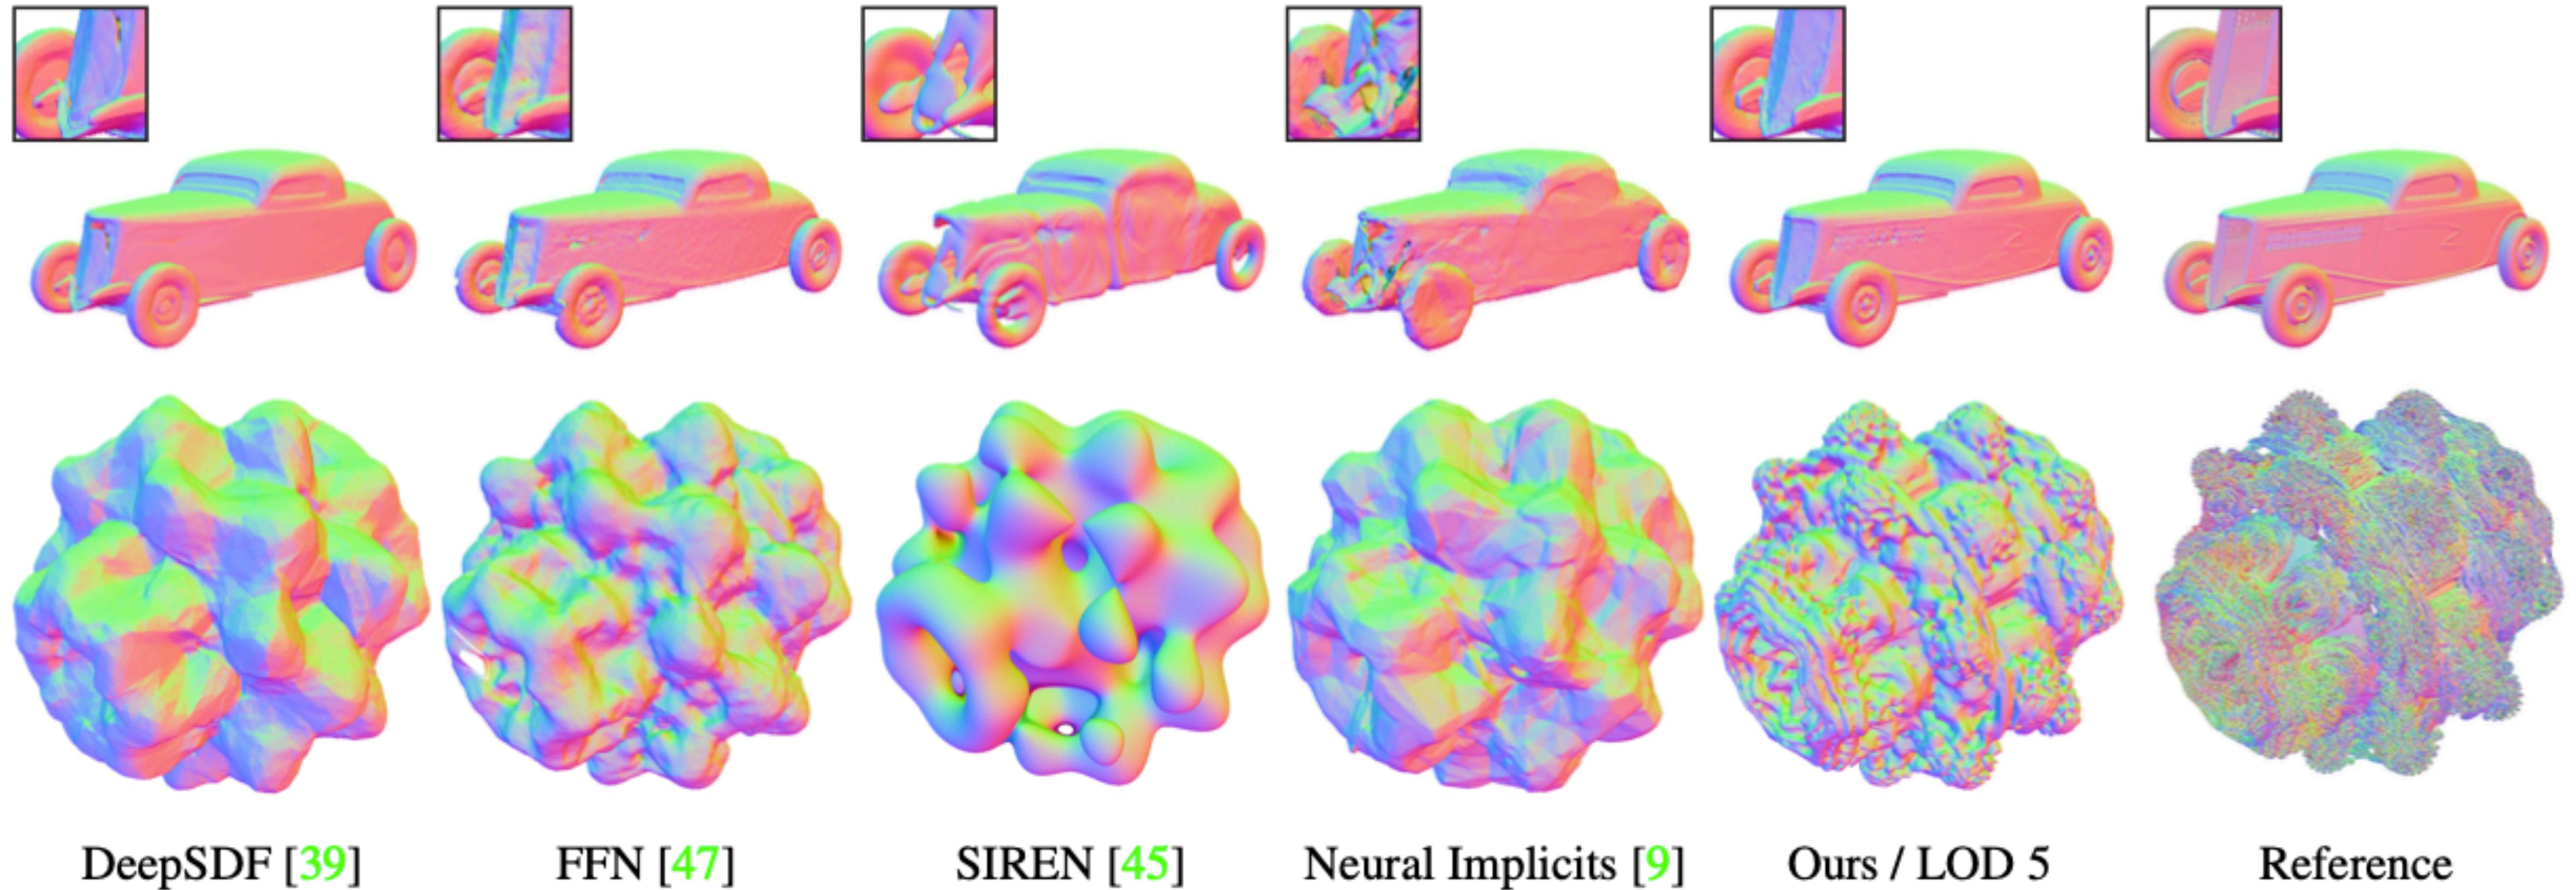
Multi-scale Voxel Grids



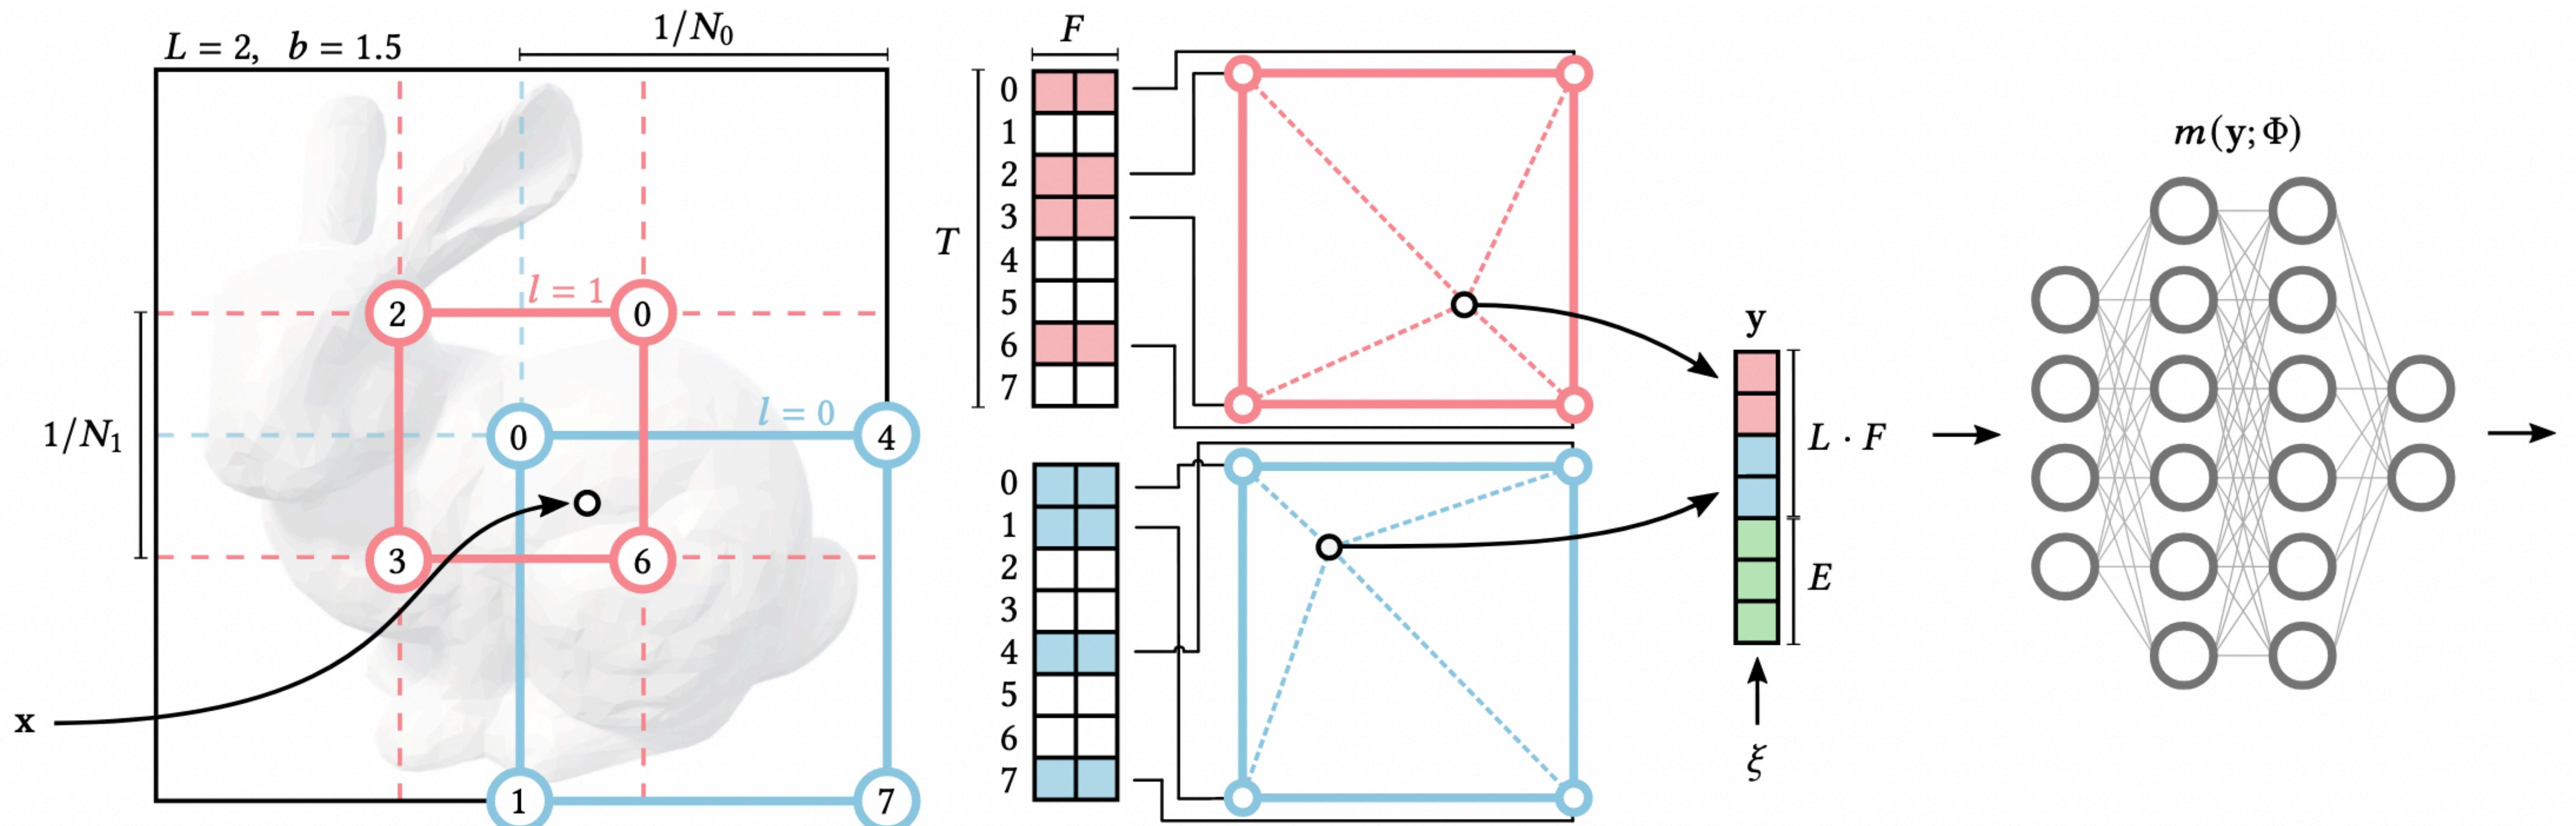
Multi-scale Voxel Grids



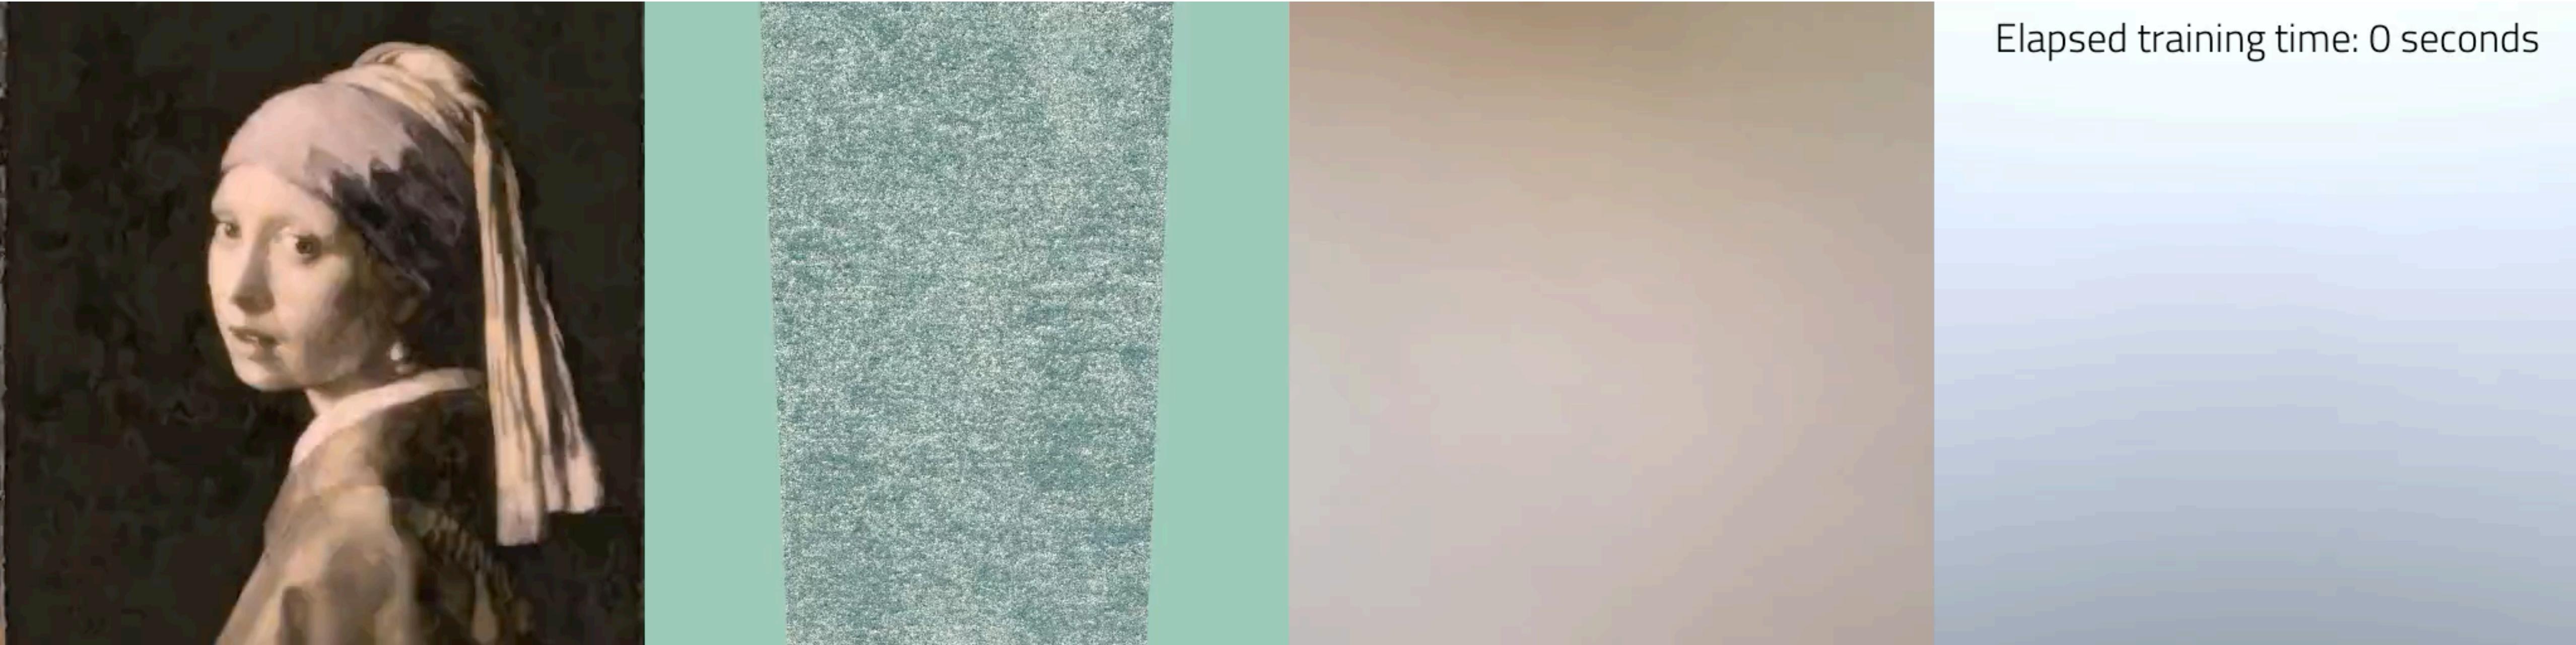
Multi-scale Voxel Grids



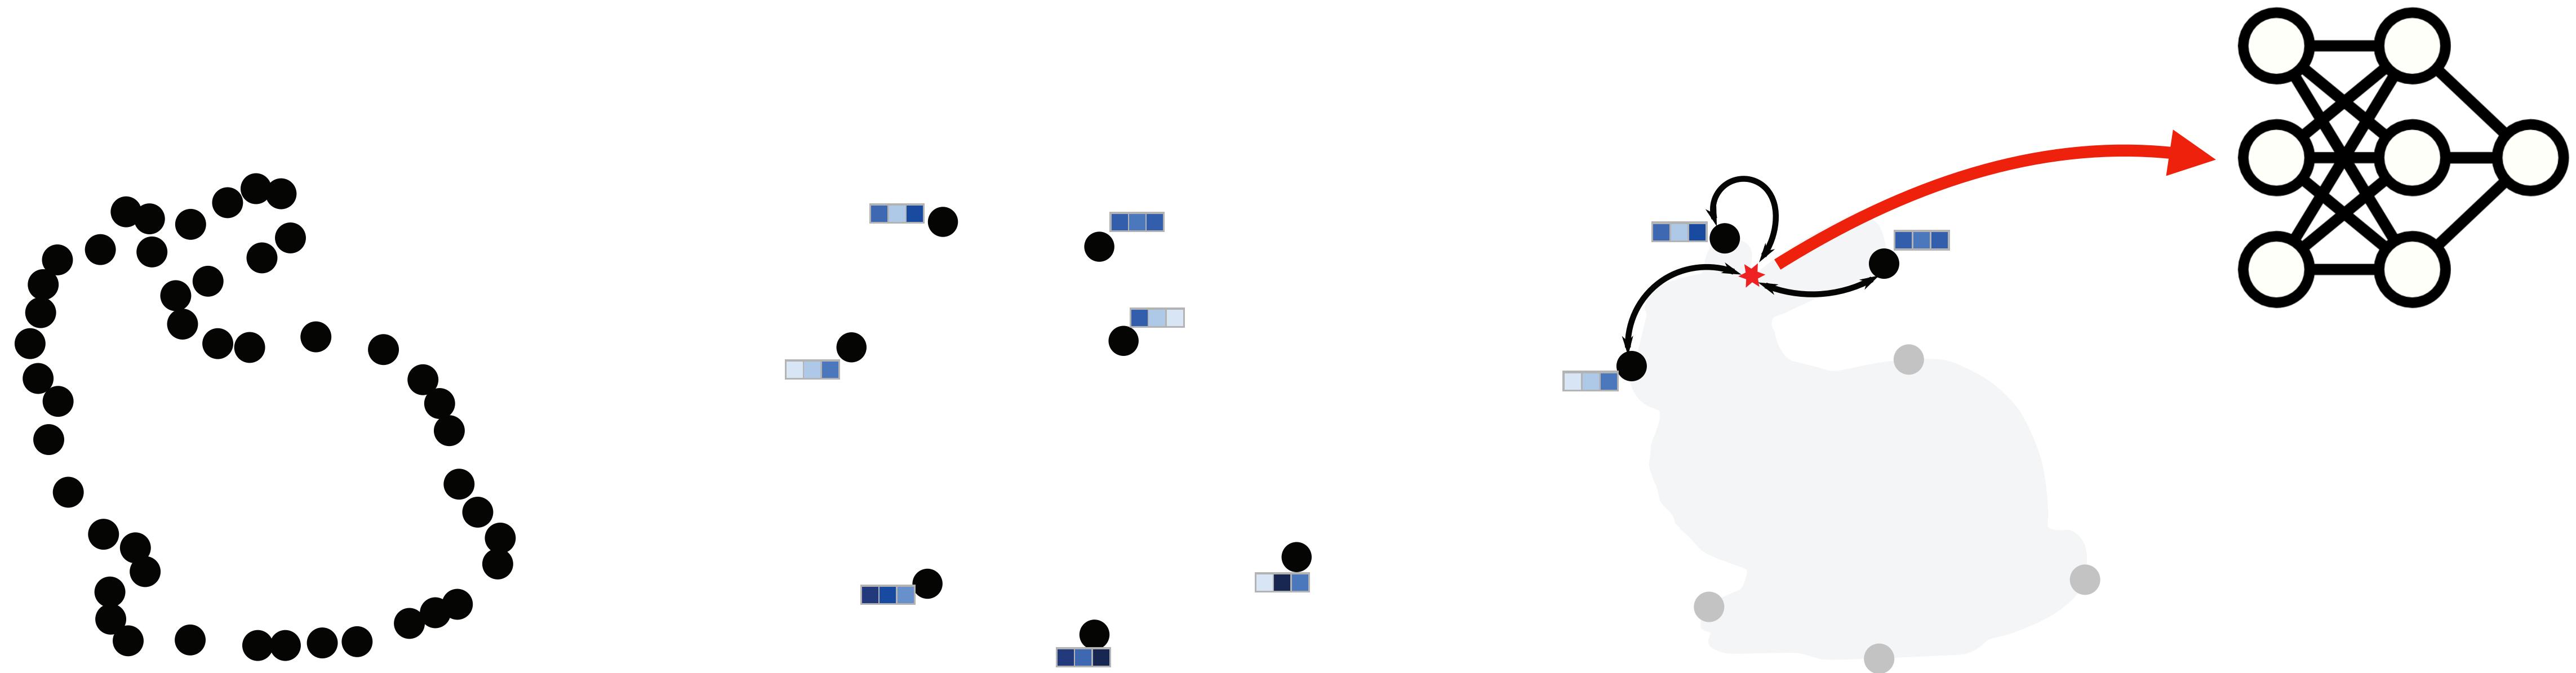
Multi-scale Voxel Grids + Hash Table



Multi-scale Voxel Grids + Hash Table

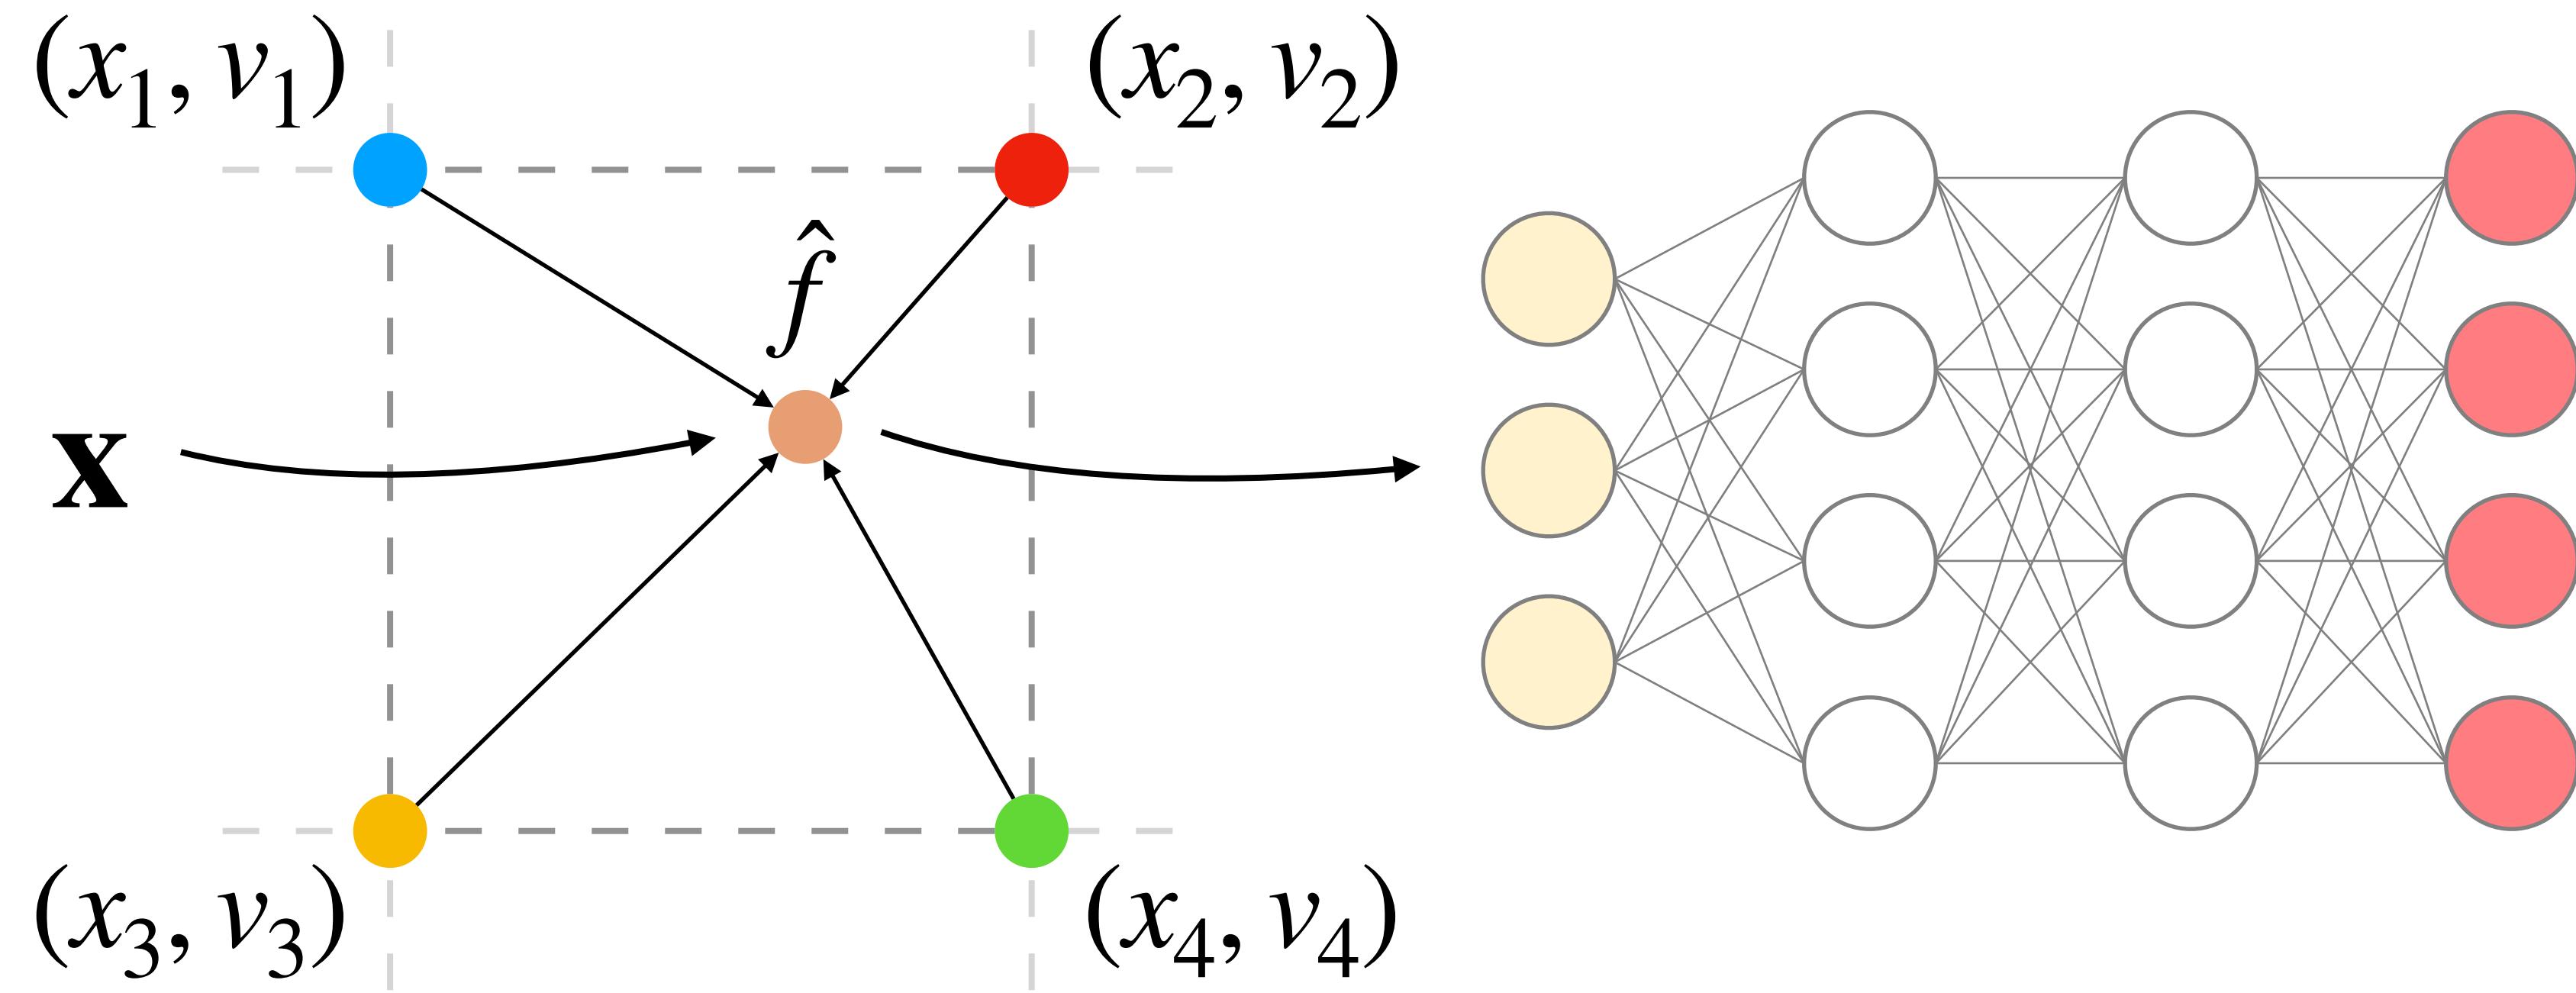


Point Cloud & Nearest Neighbor

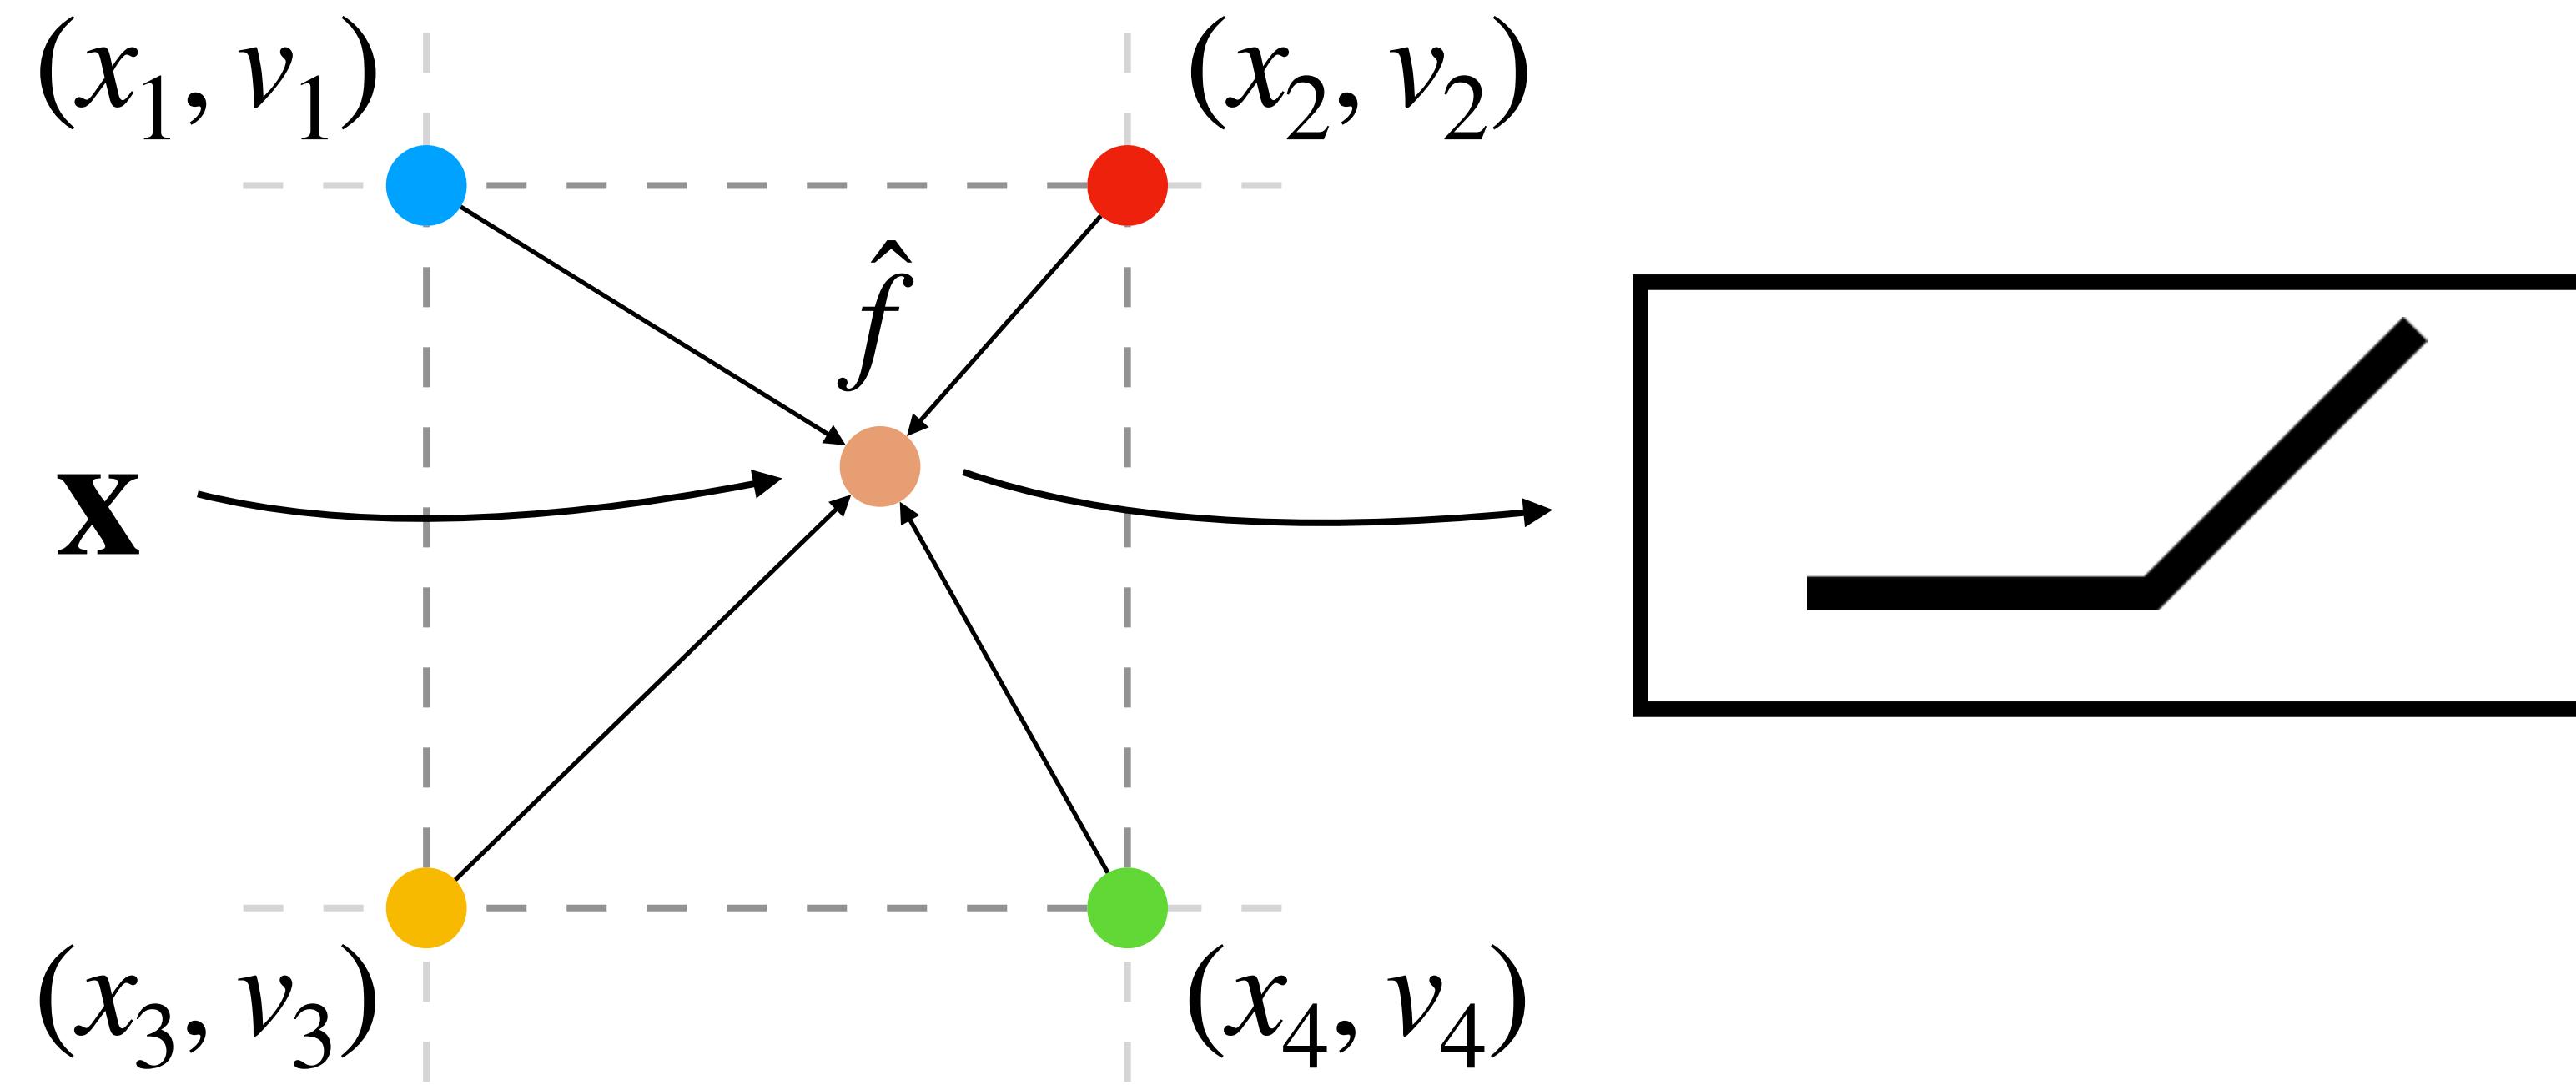


Hybrid representations can be used to decode a surface representation into a volume representation!

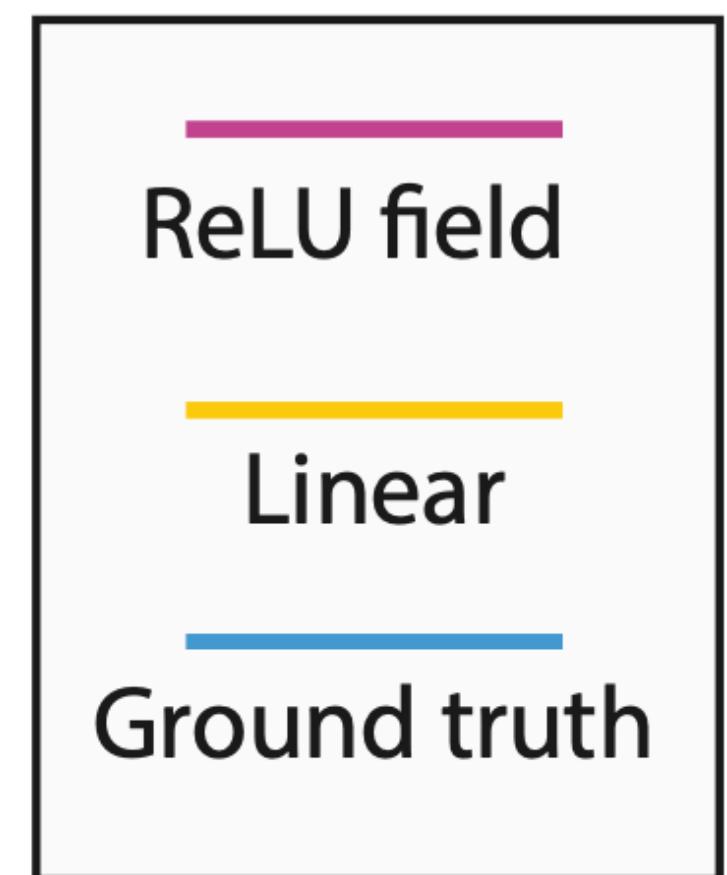
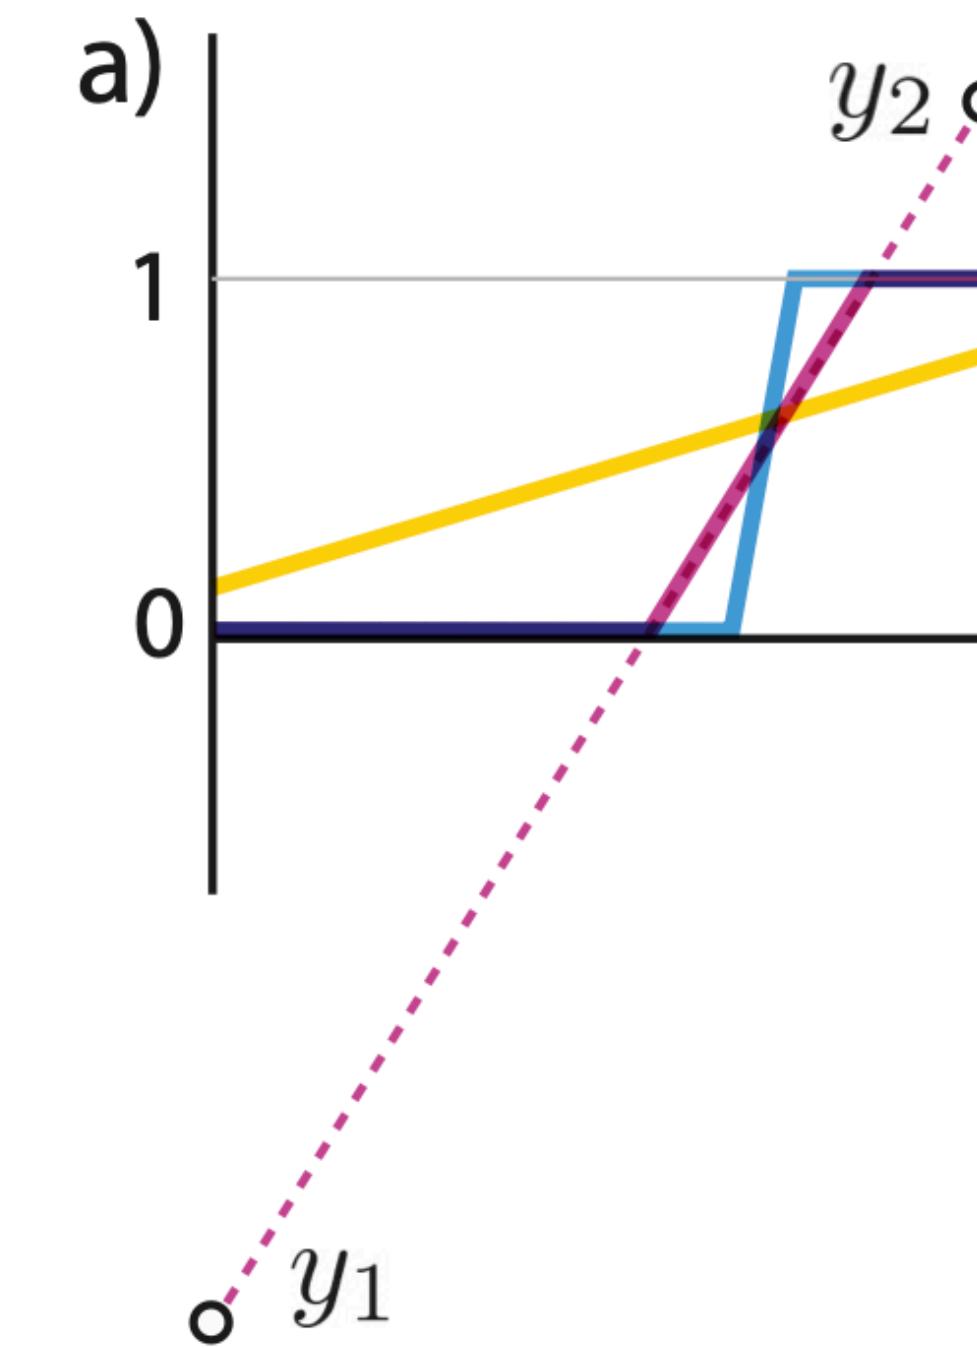
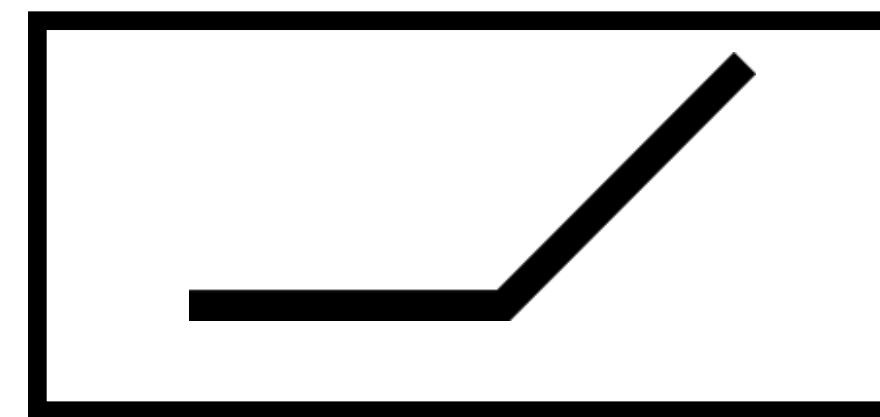
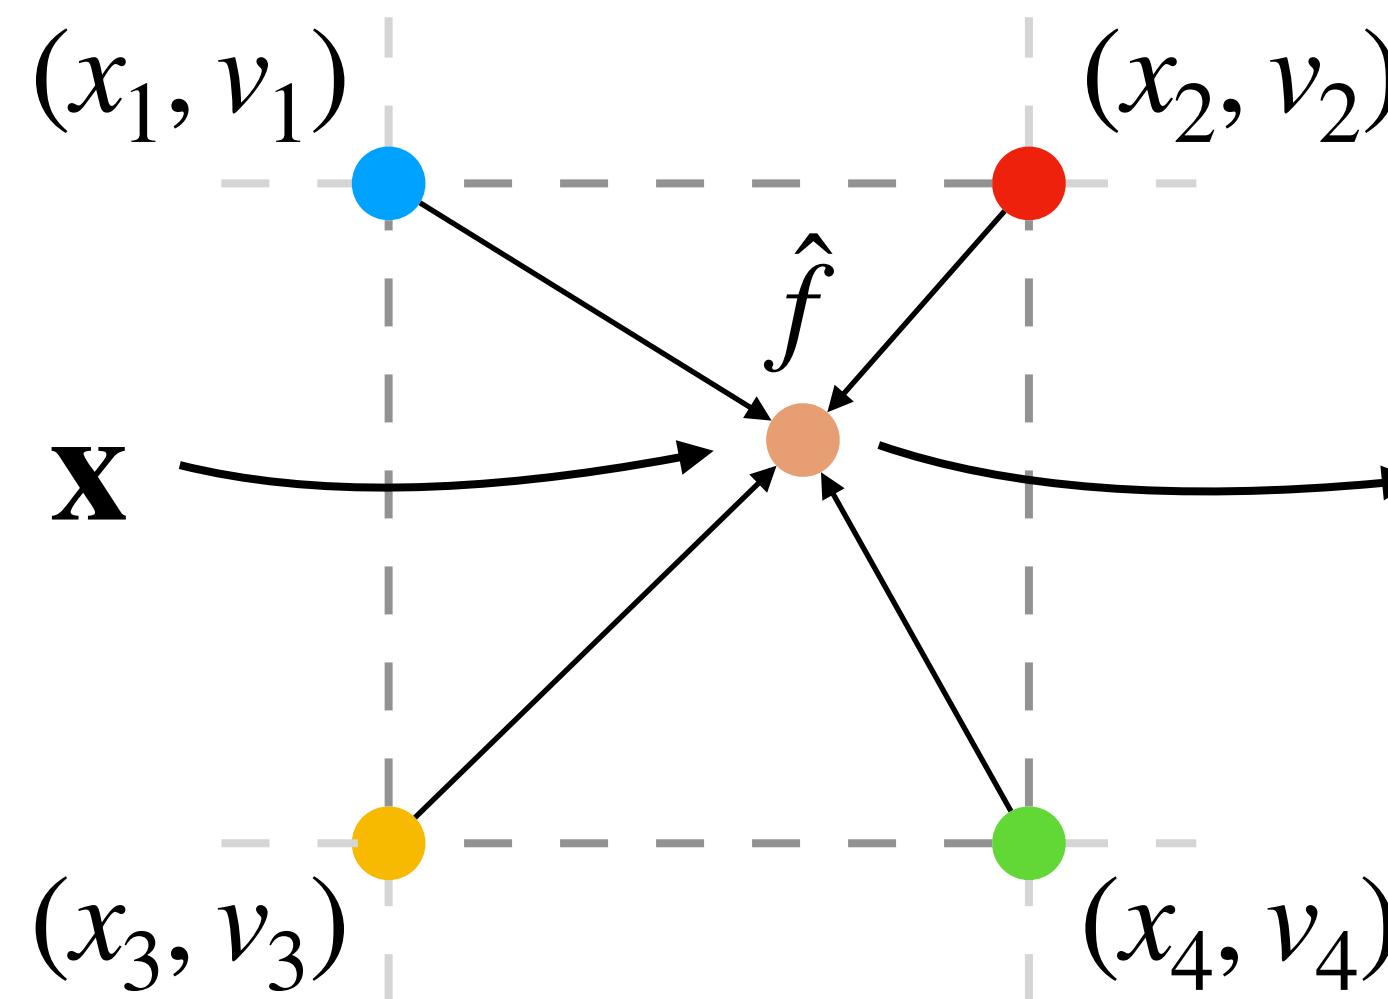
ReLU Fields: How much MLP do you need?



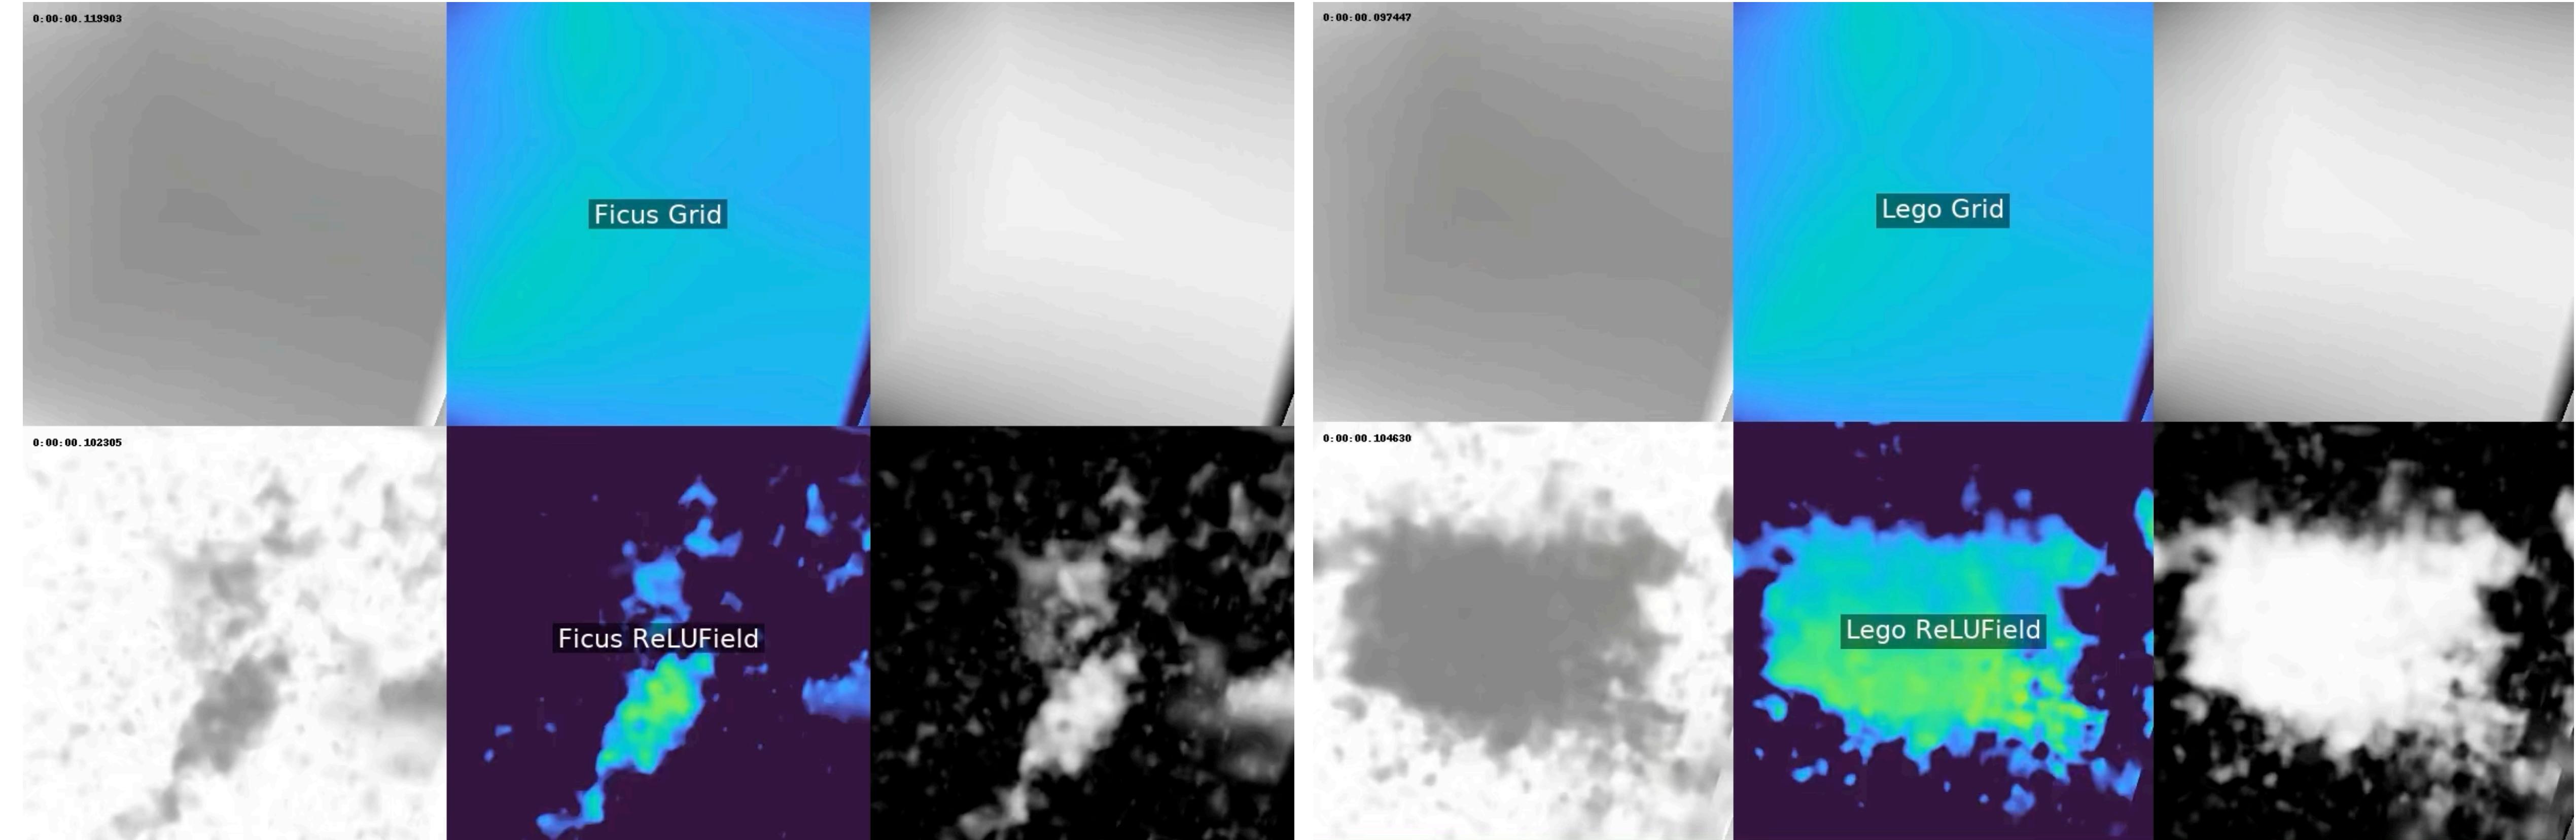
ReLU Fields: How much MLP do you need?



ReLU Fields: How much MLP do you need?

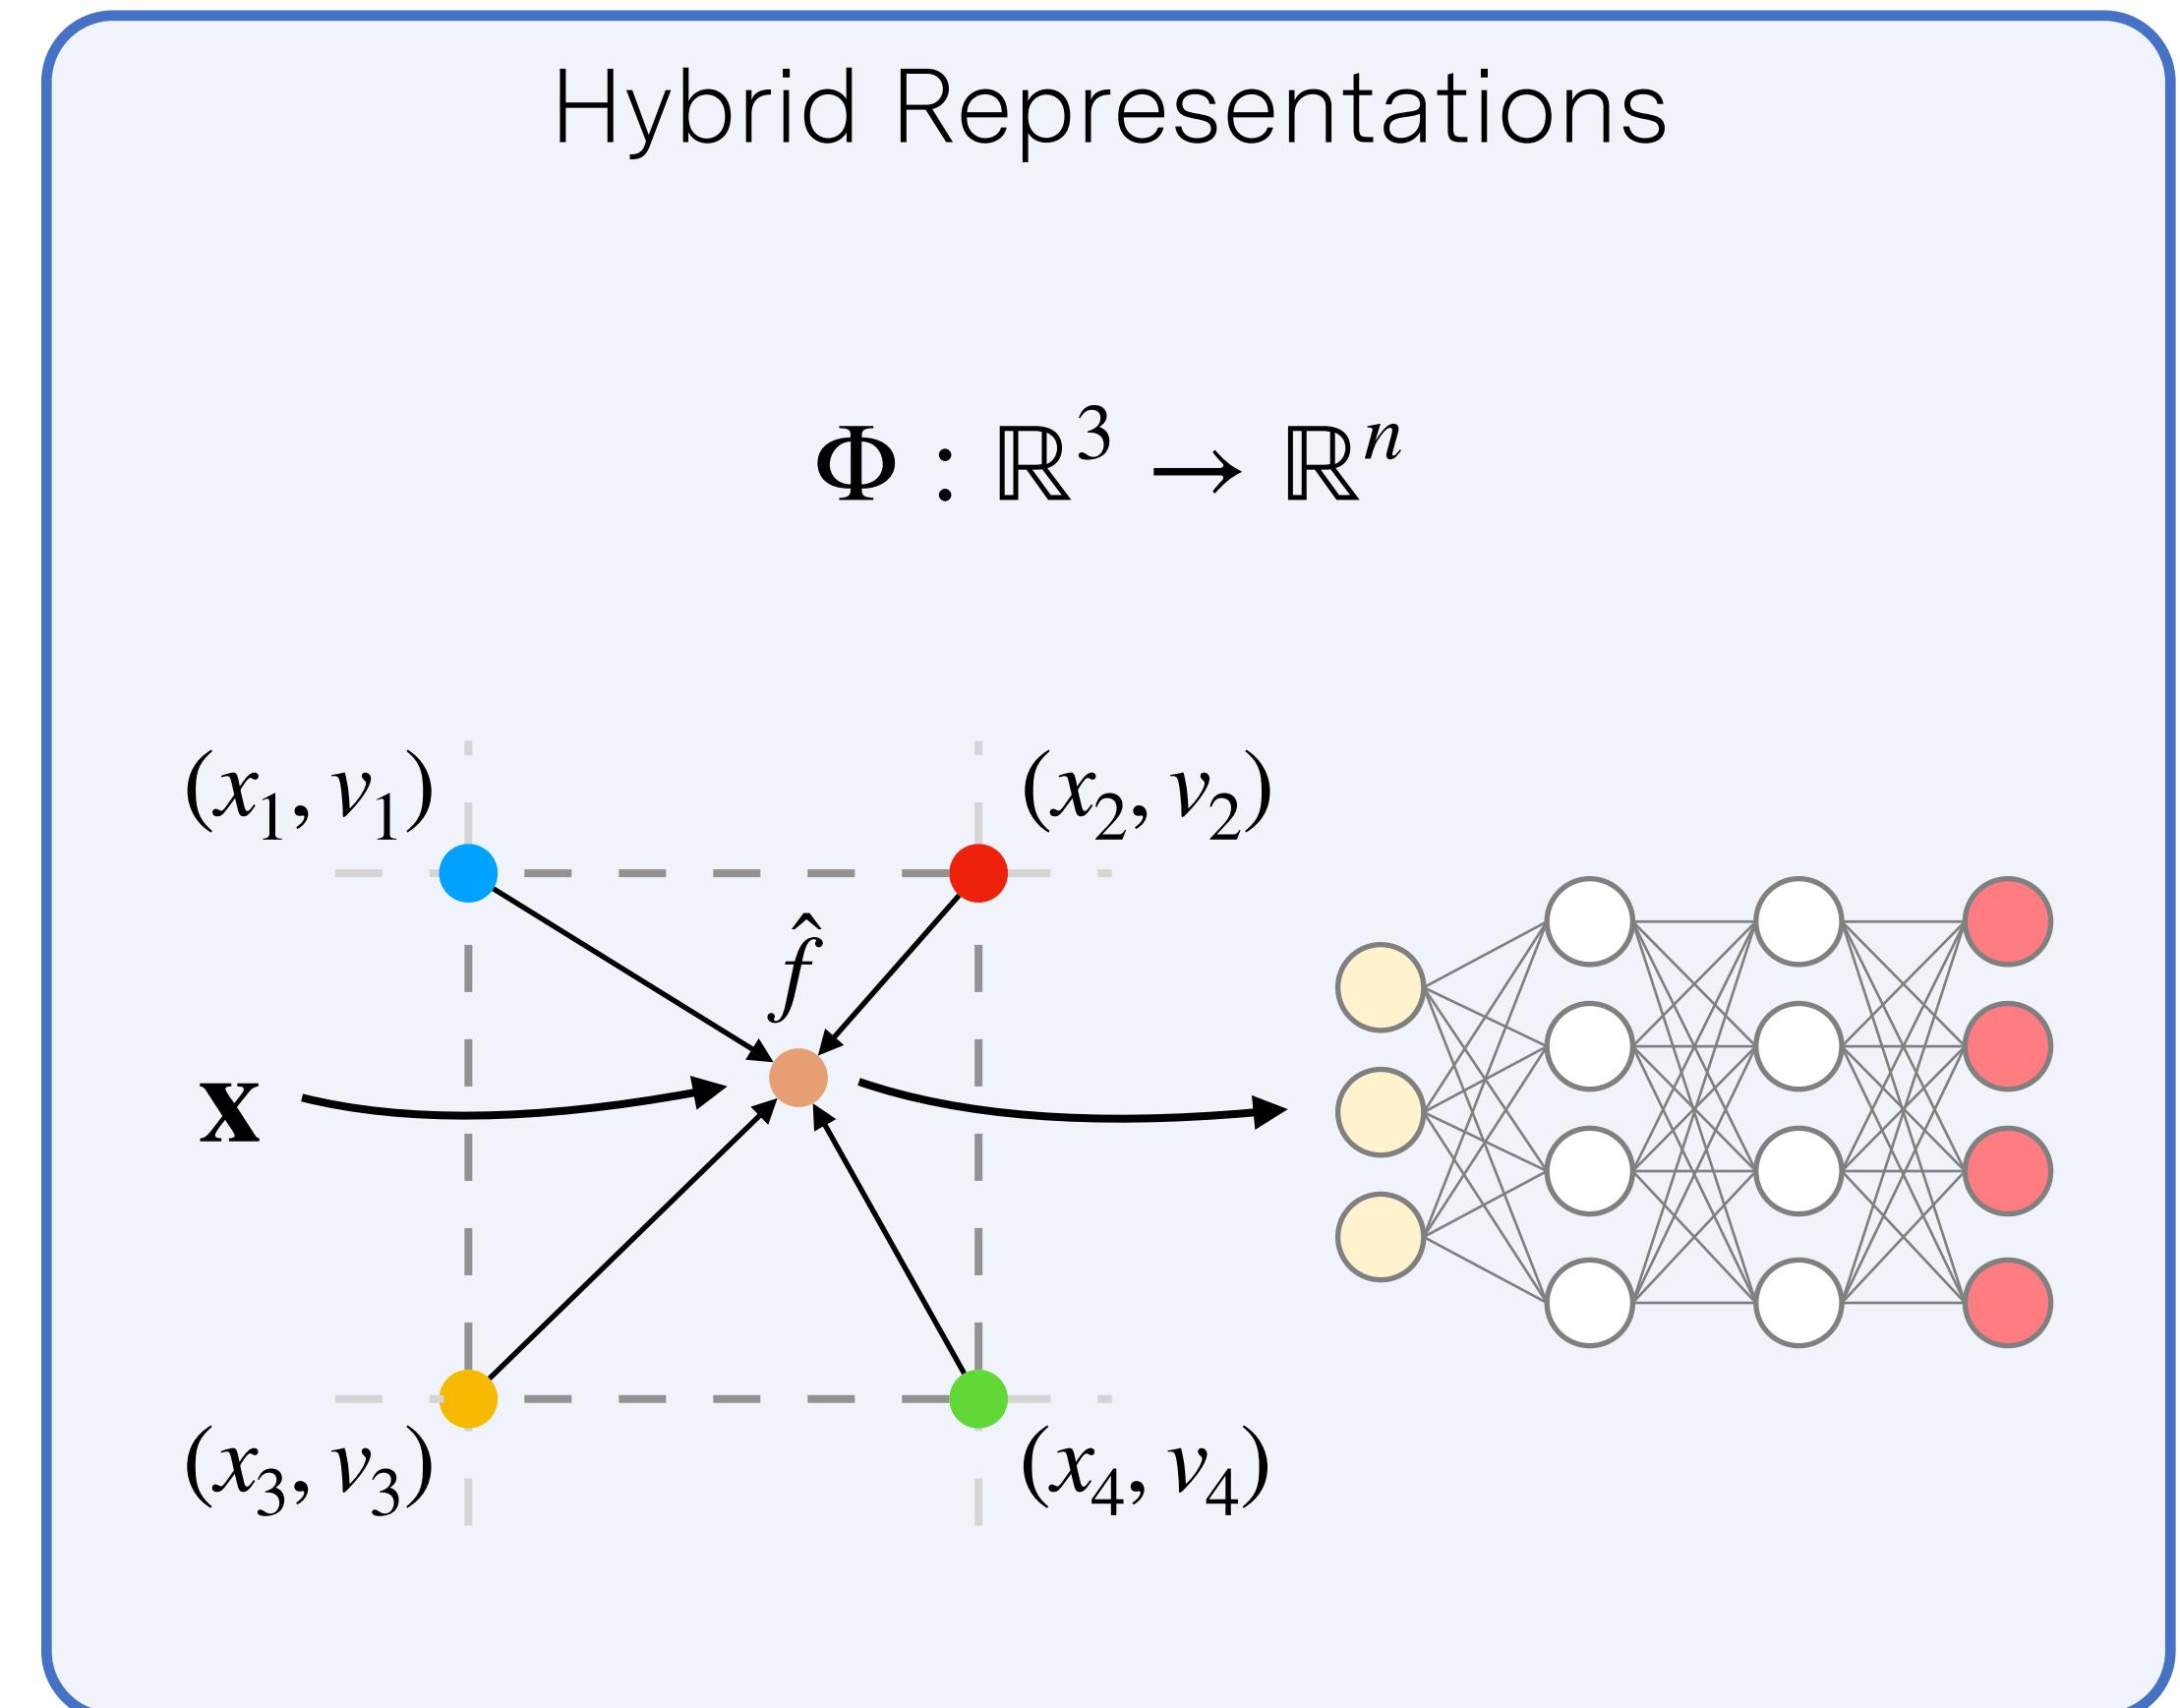


ReLU Fields: How much MLP do you need?



Hybrid Representations

- Very natural to trade off memory with compute.
- Neural network can locally “super resolve” discrete representation: Basically acts as interpolation kernel.
- Can be the best of both worlds: Fast and “continuous”.

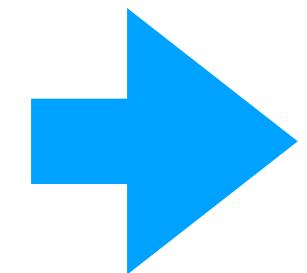


How can we represent unbounded scenes?



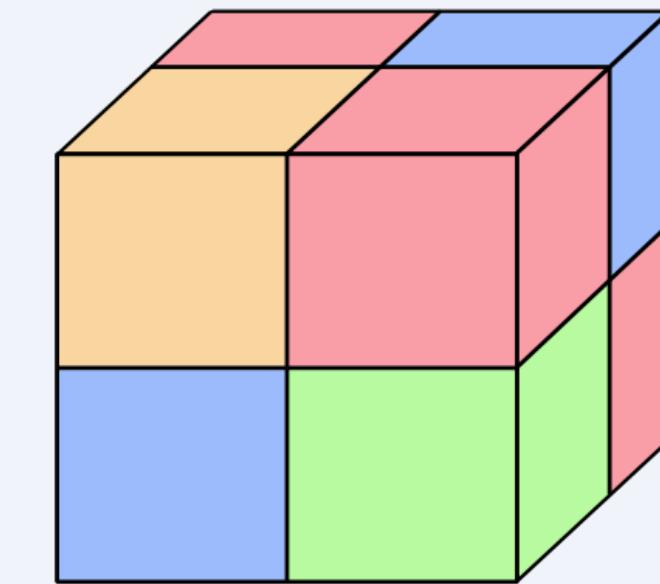
Image Source: Wikipedia

How can we represent unbounded scenes?

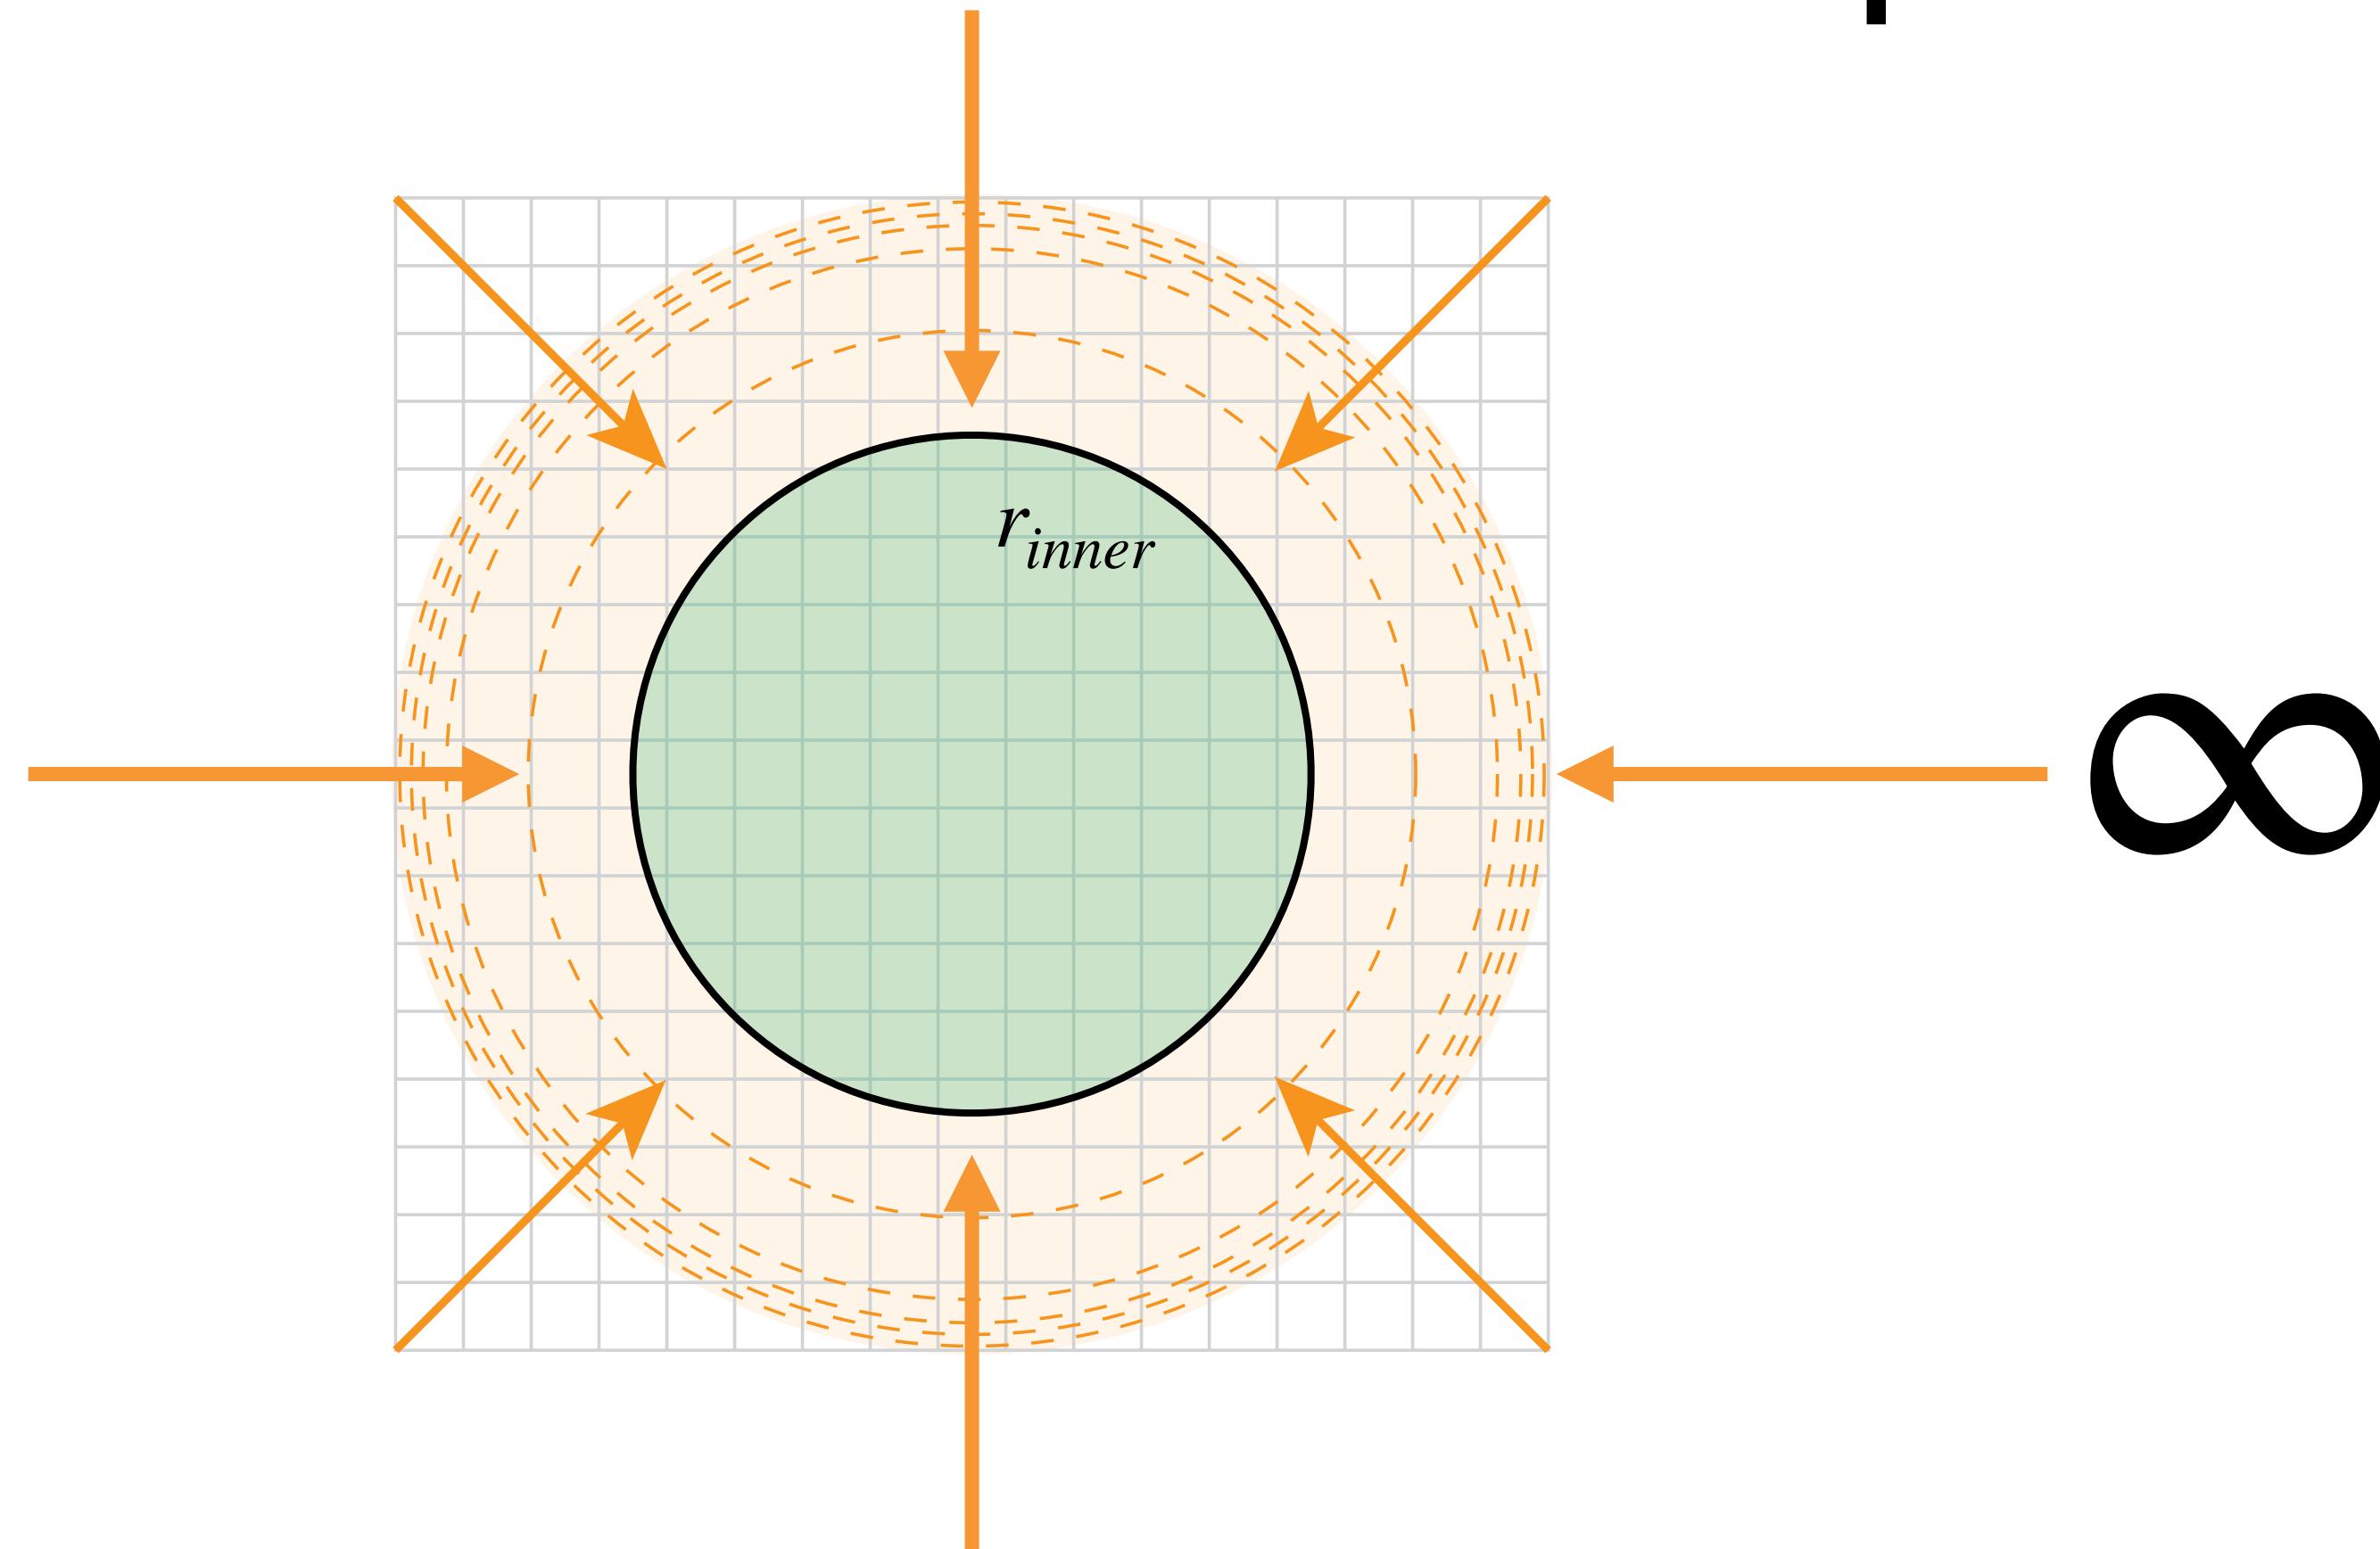


Scene
Representation

$$\Phi : \mathbb{R}^3 \rightarrow \mathbb{R}^n$$



Parameterizing Unbounded Scenes: Compactification

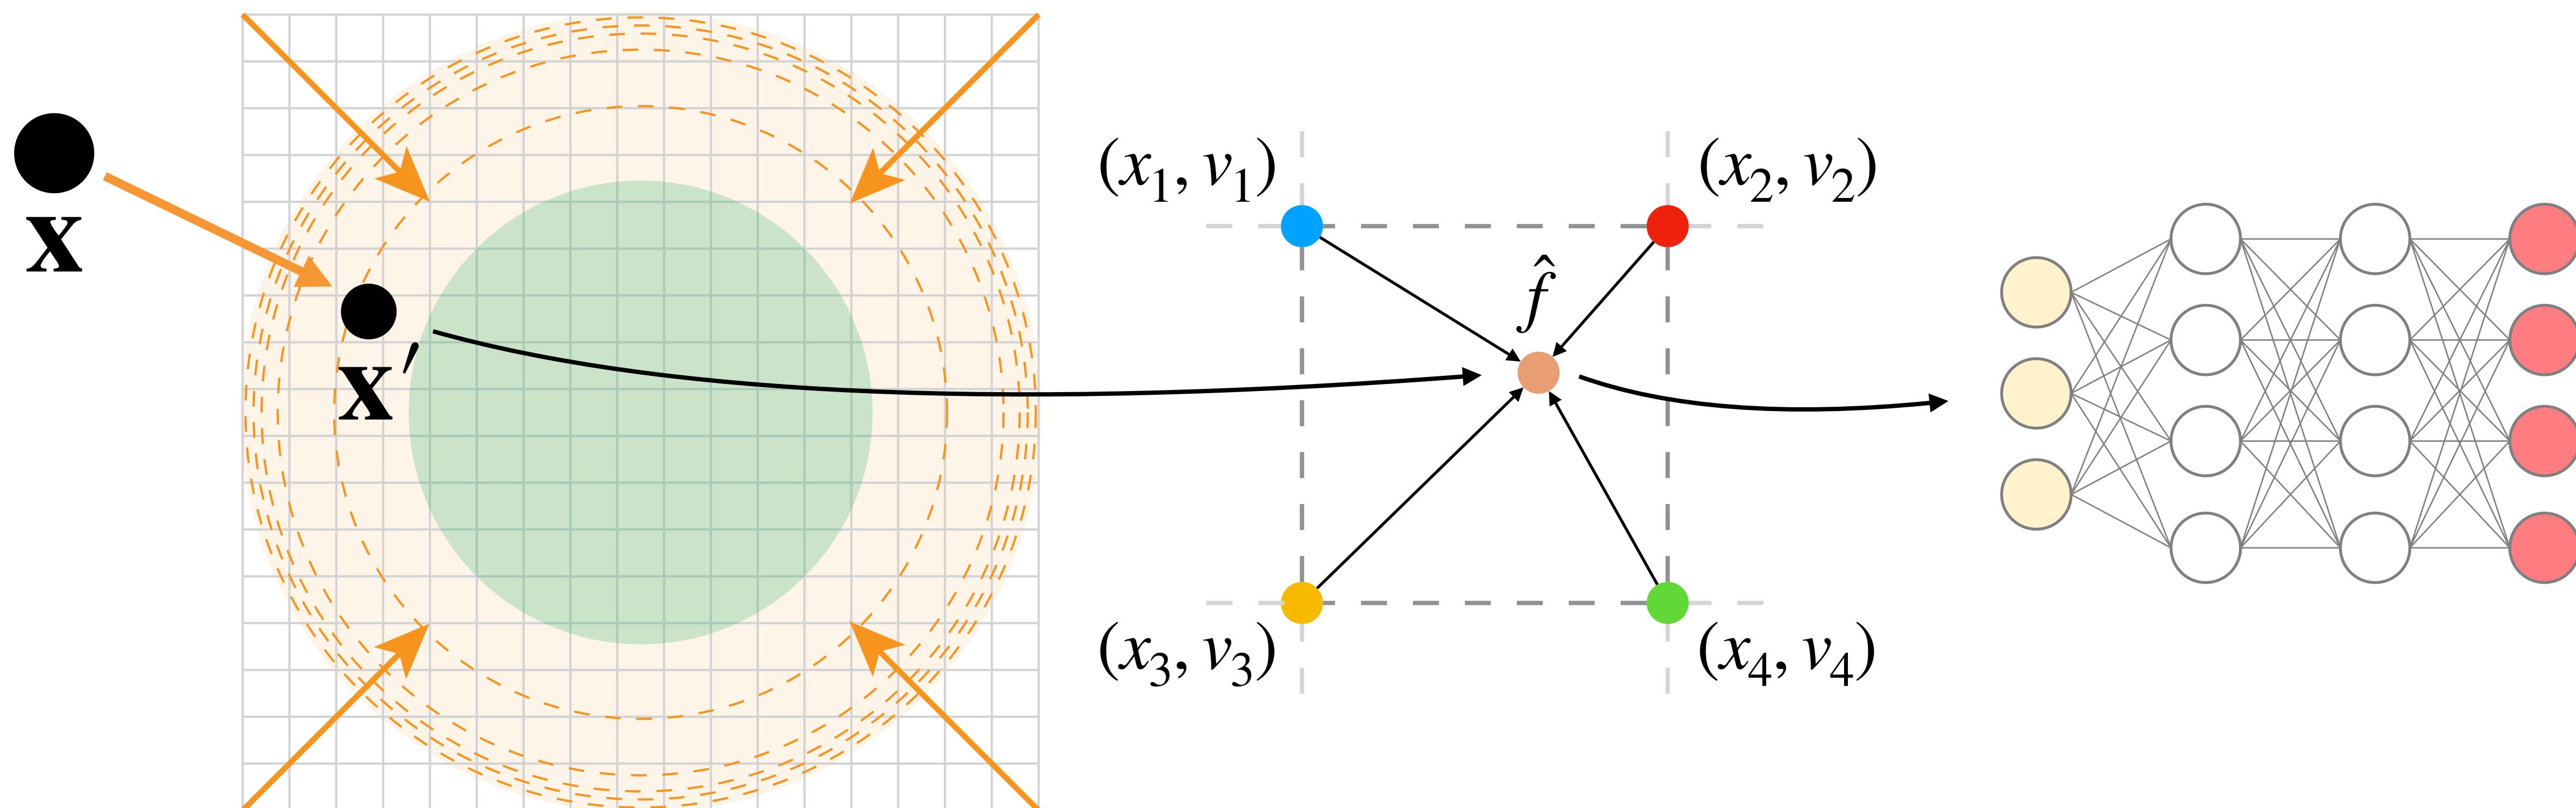


$$\mathbf{x}' = \frac{(1 + k) - k/\|\mathbf{u}\|}{\mathbf{u}/\|\mathbf{u}\|} r_{inner}$$
$$\mathbf{u} = \frac{\mathbf{x}}{r_{inner}}$$

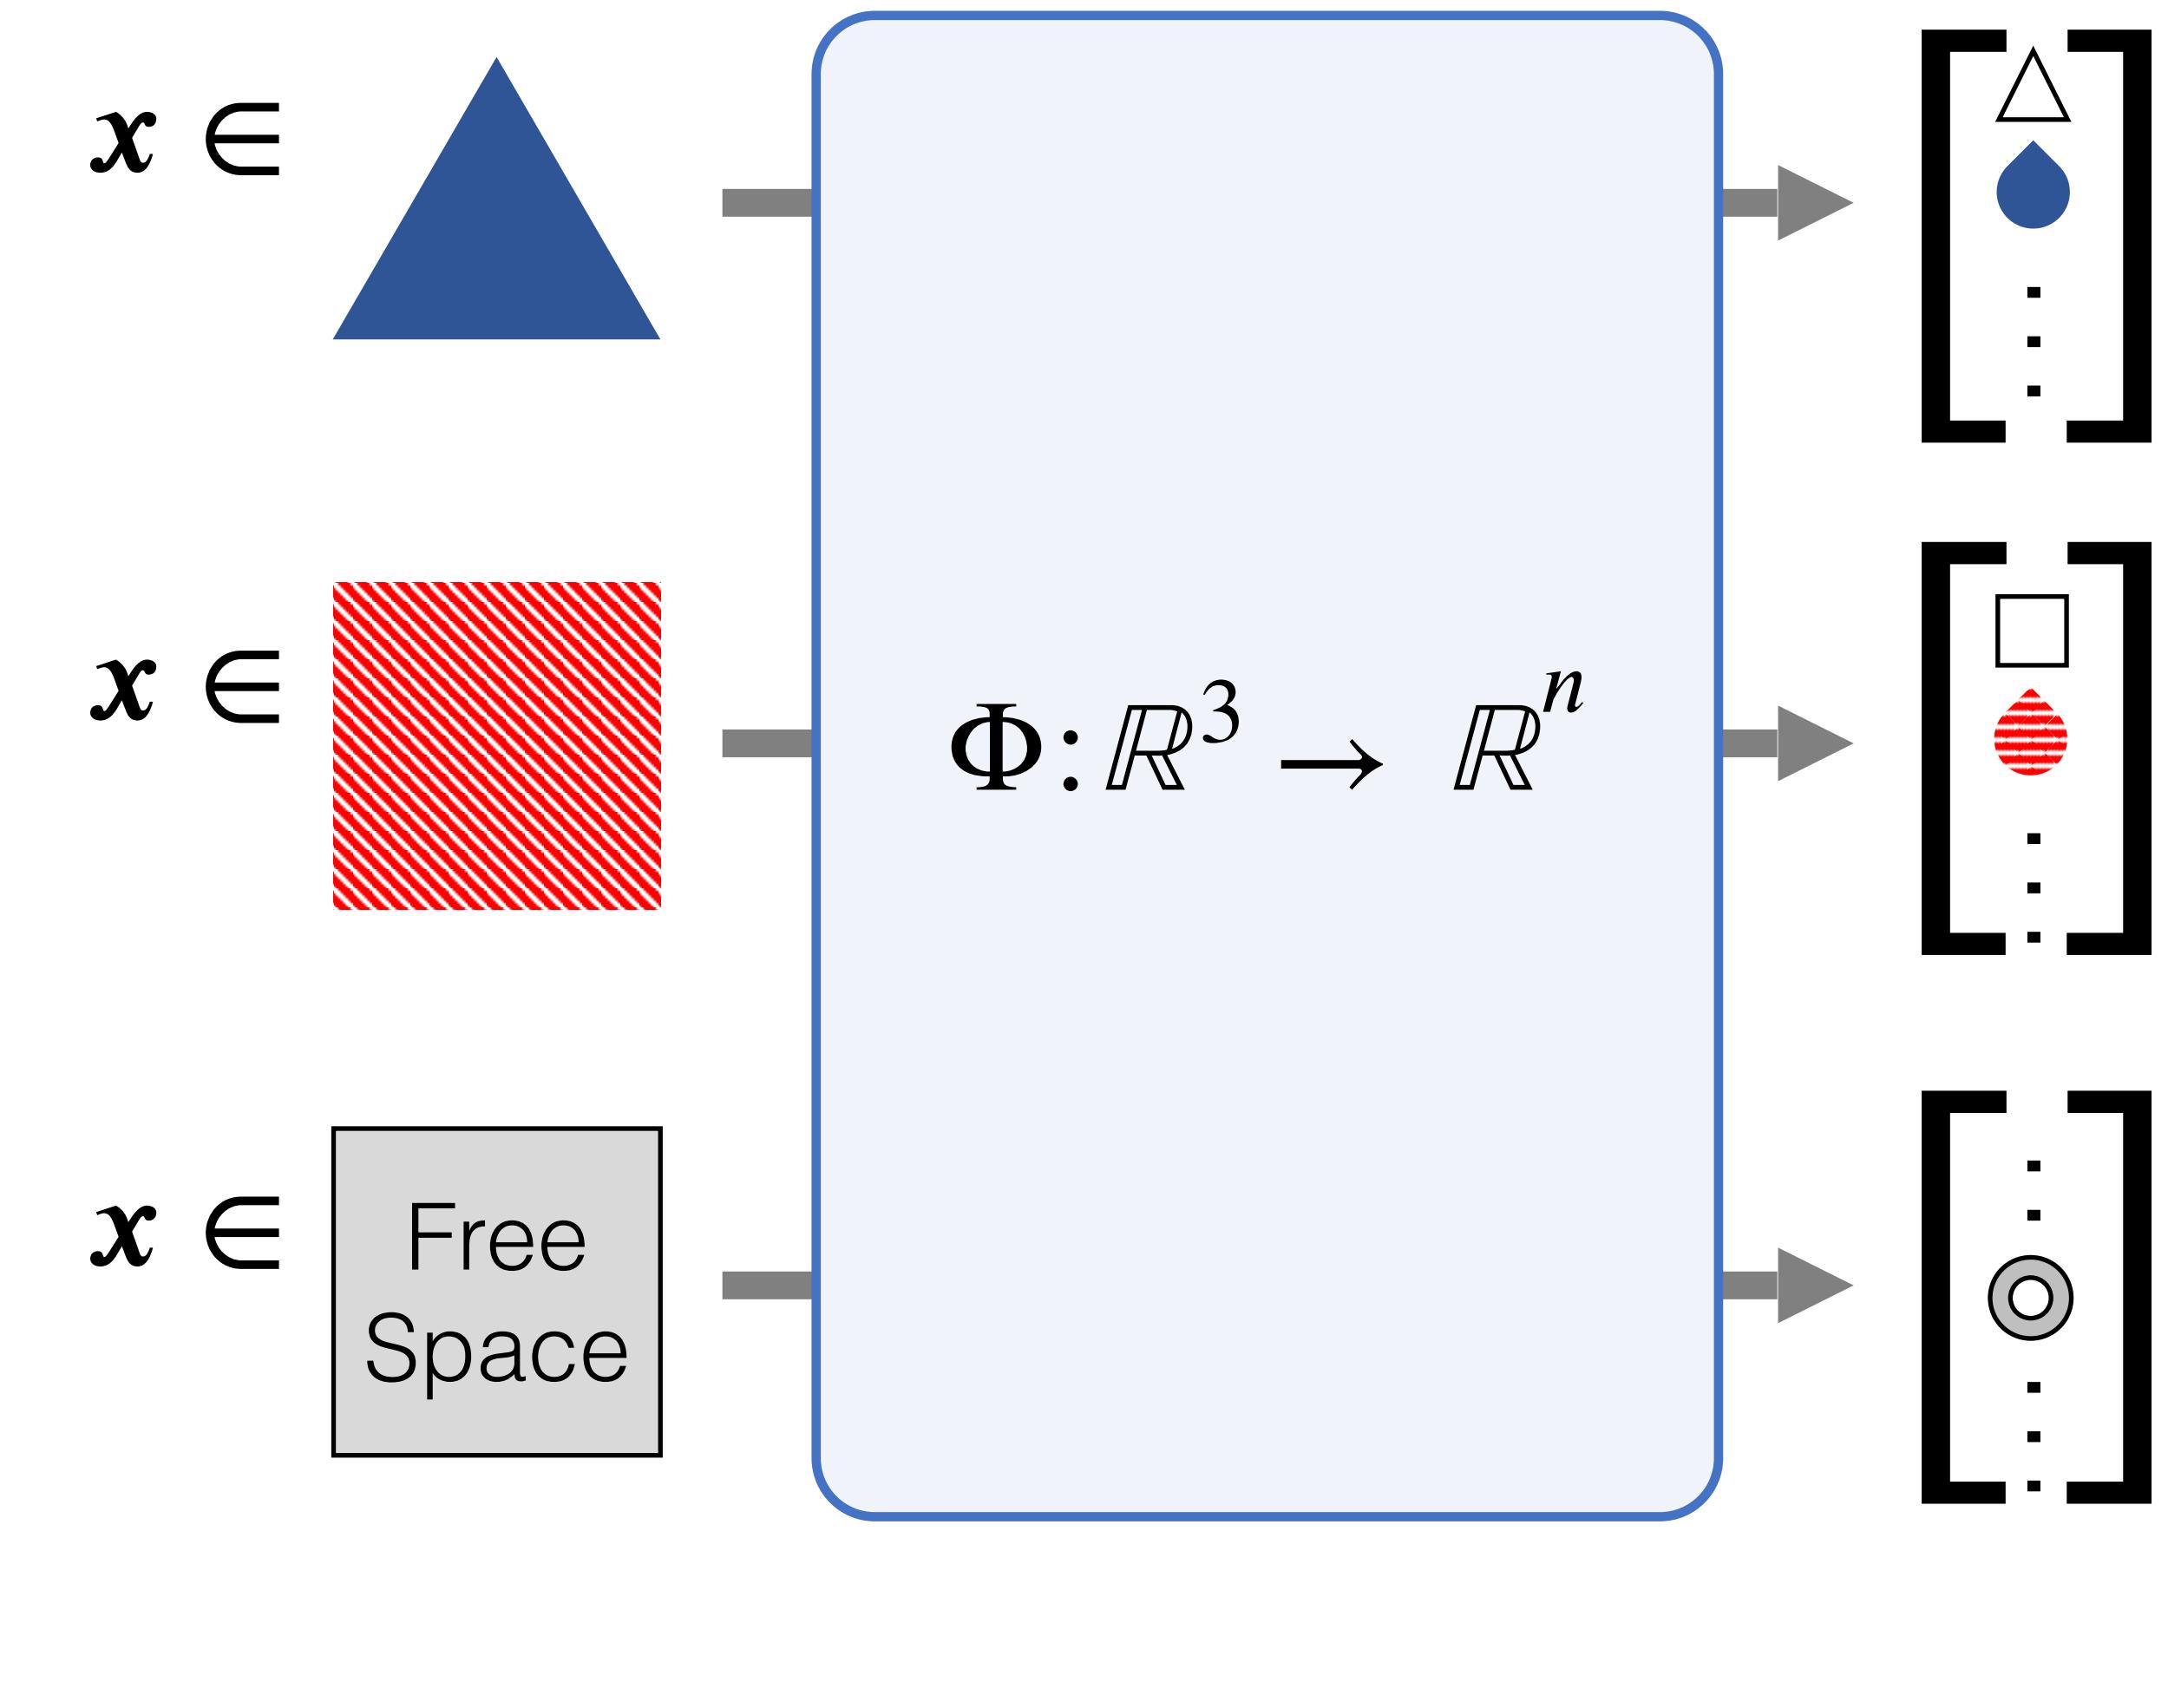
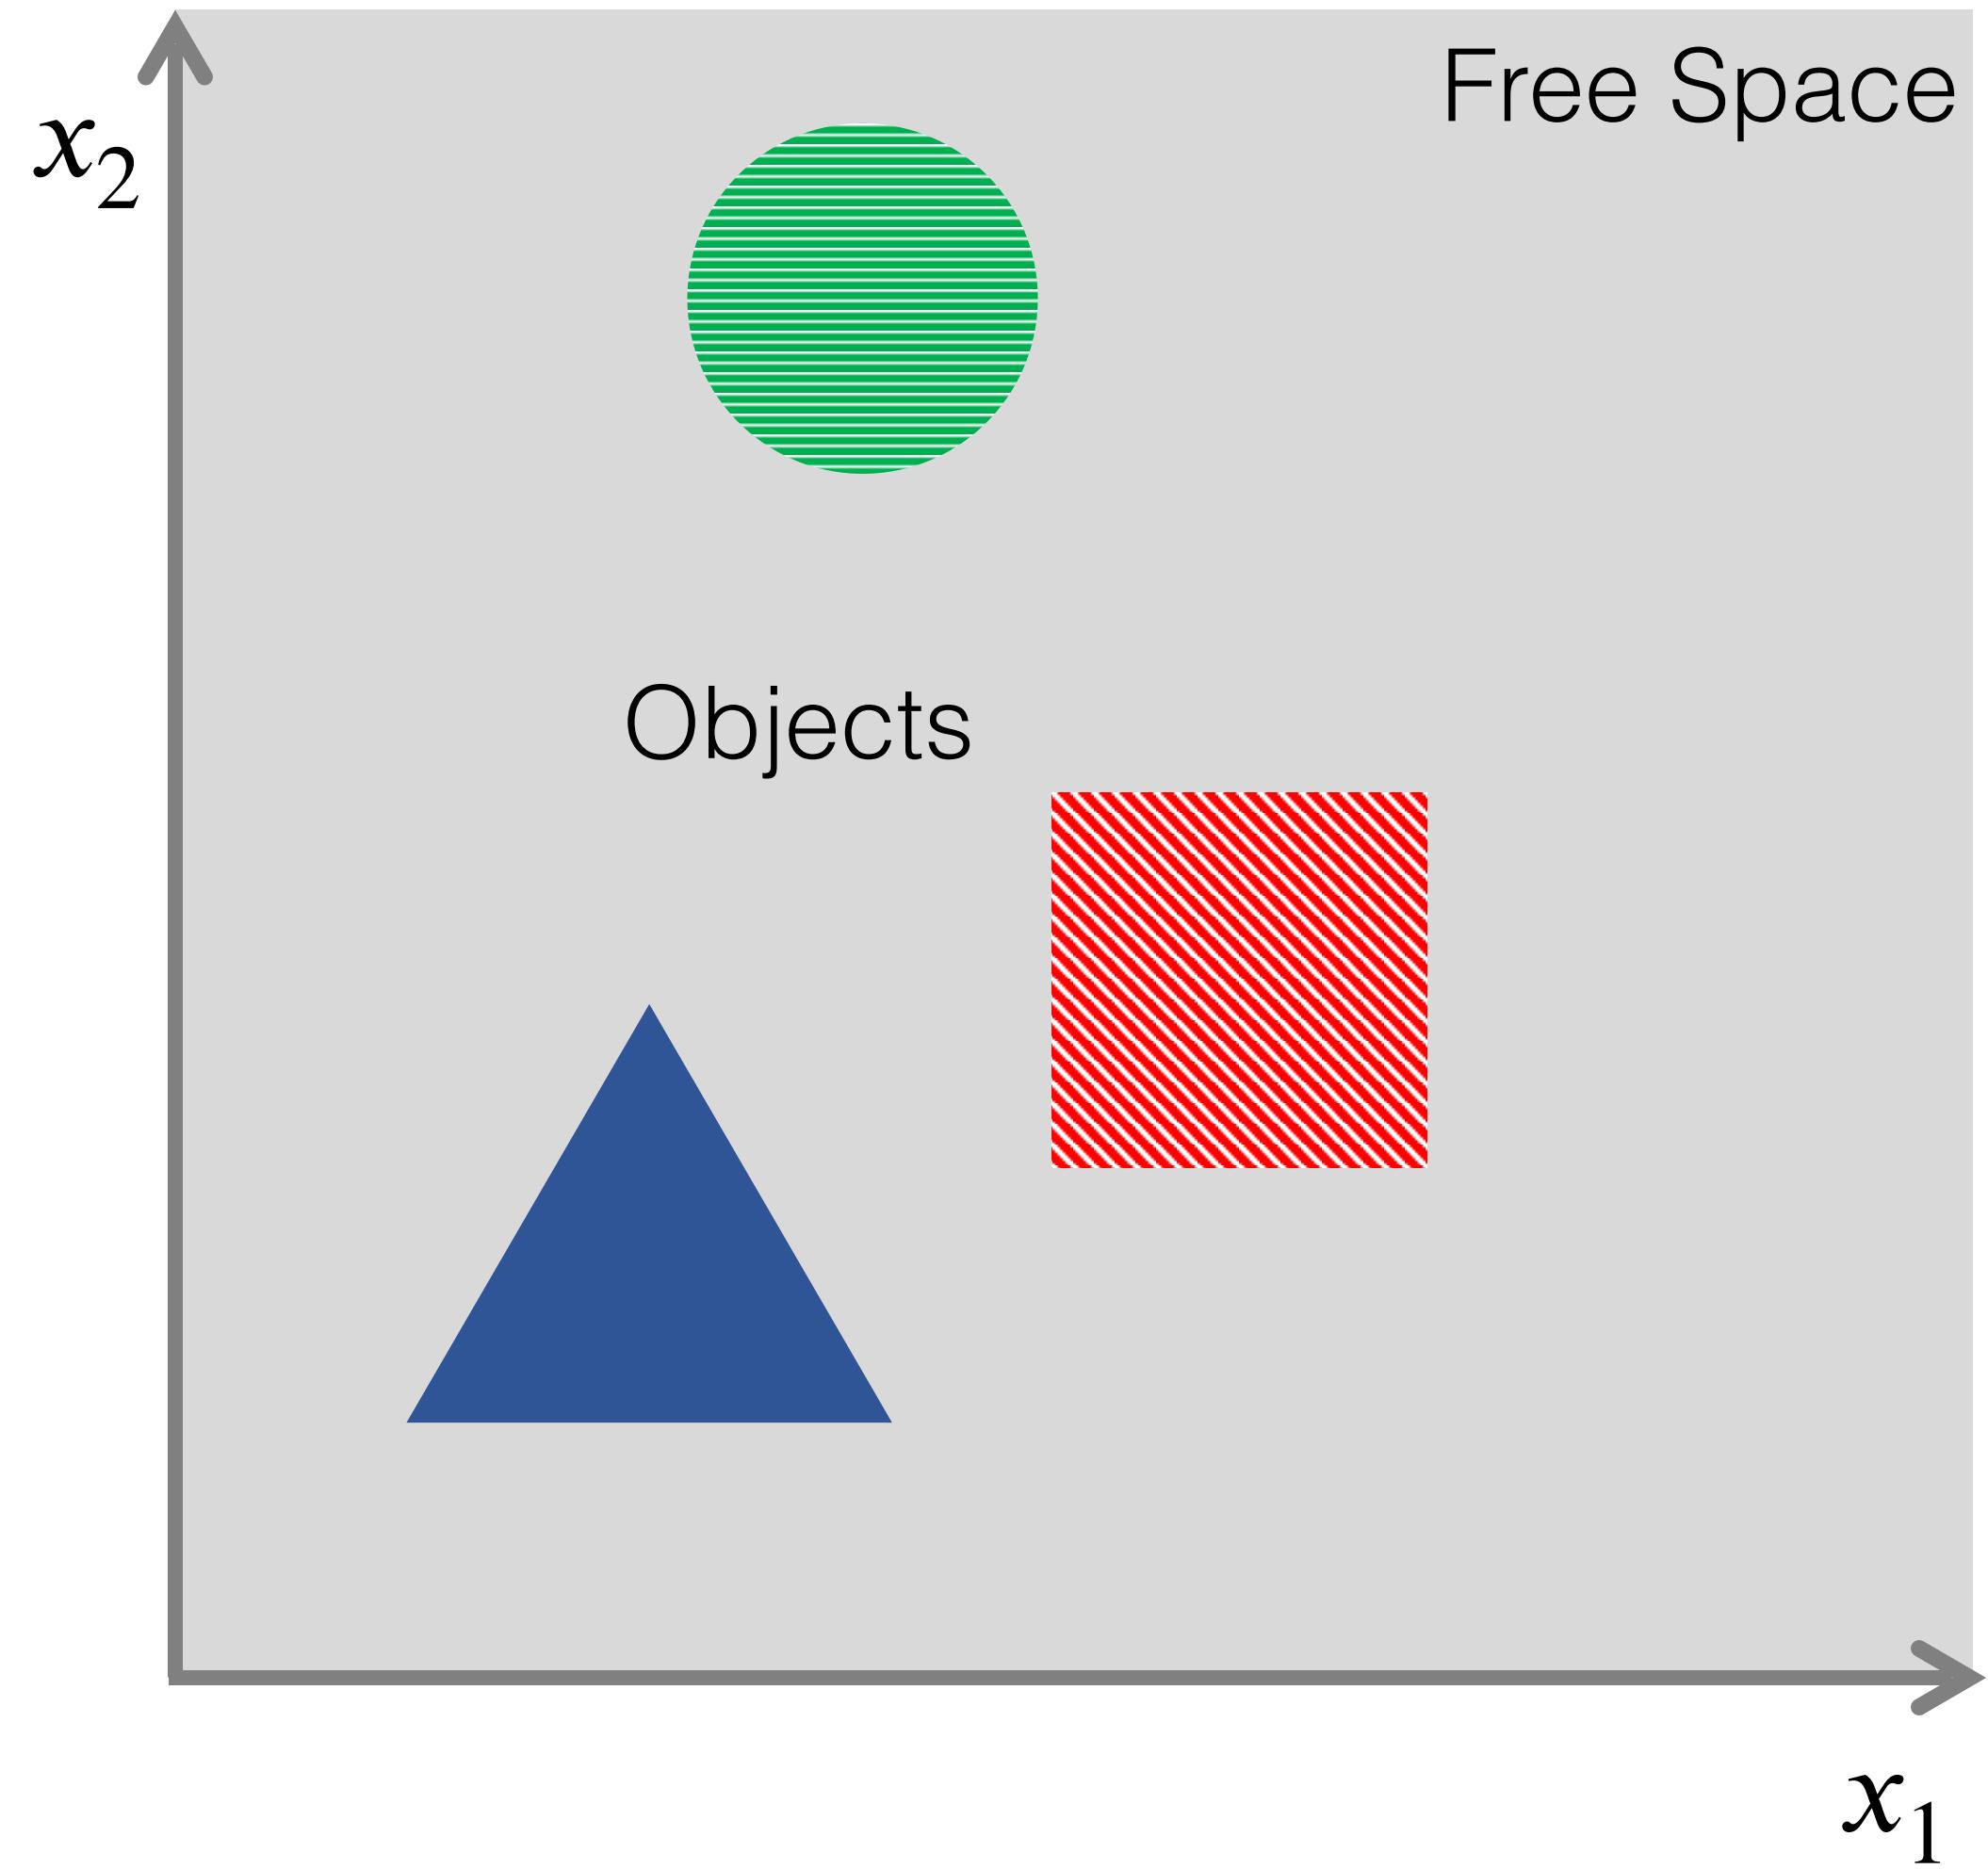
Example Compactification function from
“Mip-NeRF 360: Unbounded Anti-Aliased Neural
Radiance Fields”, Barron et al., CVPR 2022

Come up with nonlinear mapping that “squashes” space, such that infinity is mapped to some **finite** value. Aligns with intuition “further away, less resolution”.

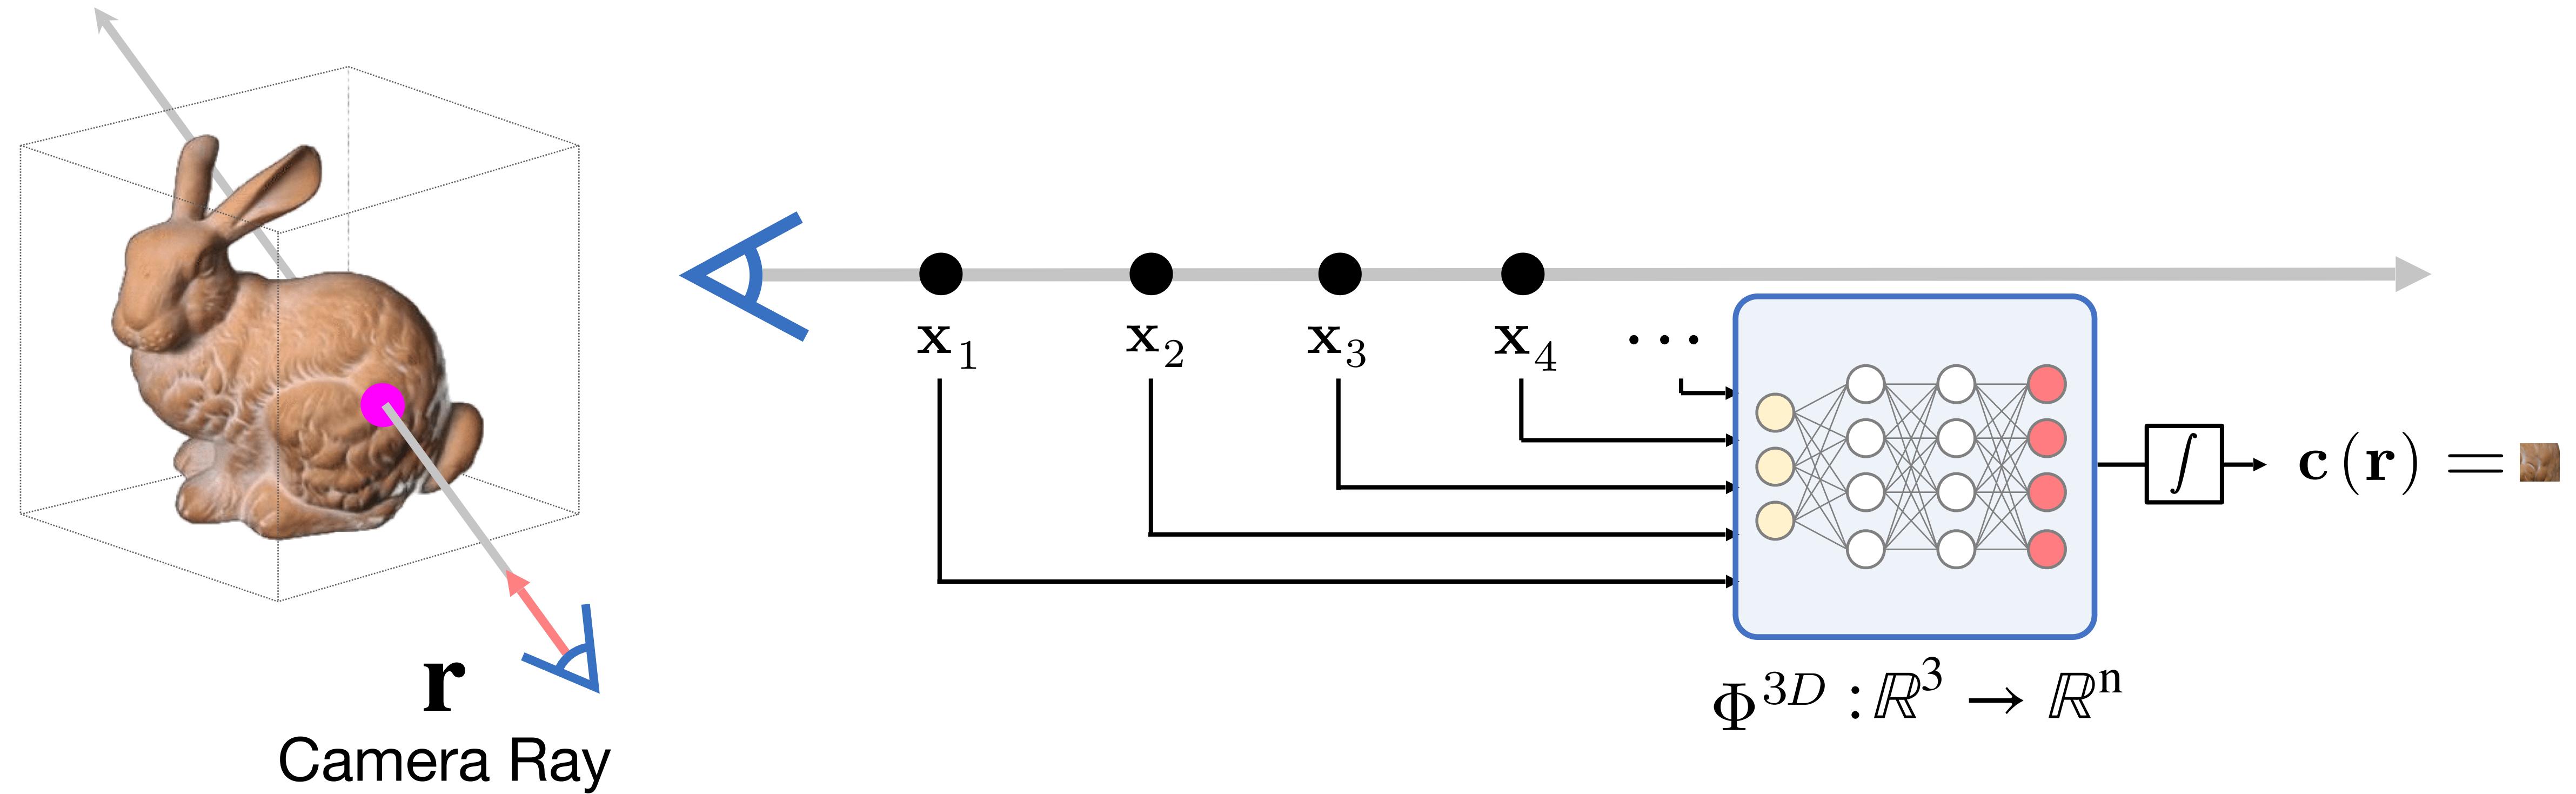
Parameterizing Unbounded Scenes: Compactification

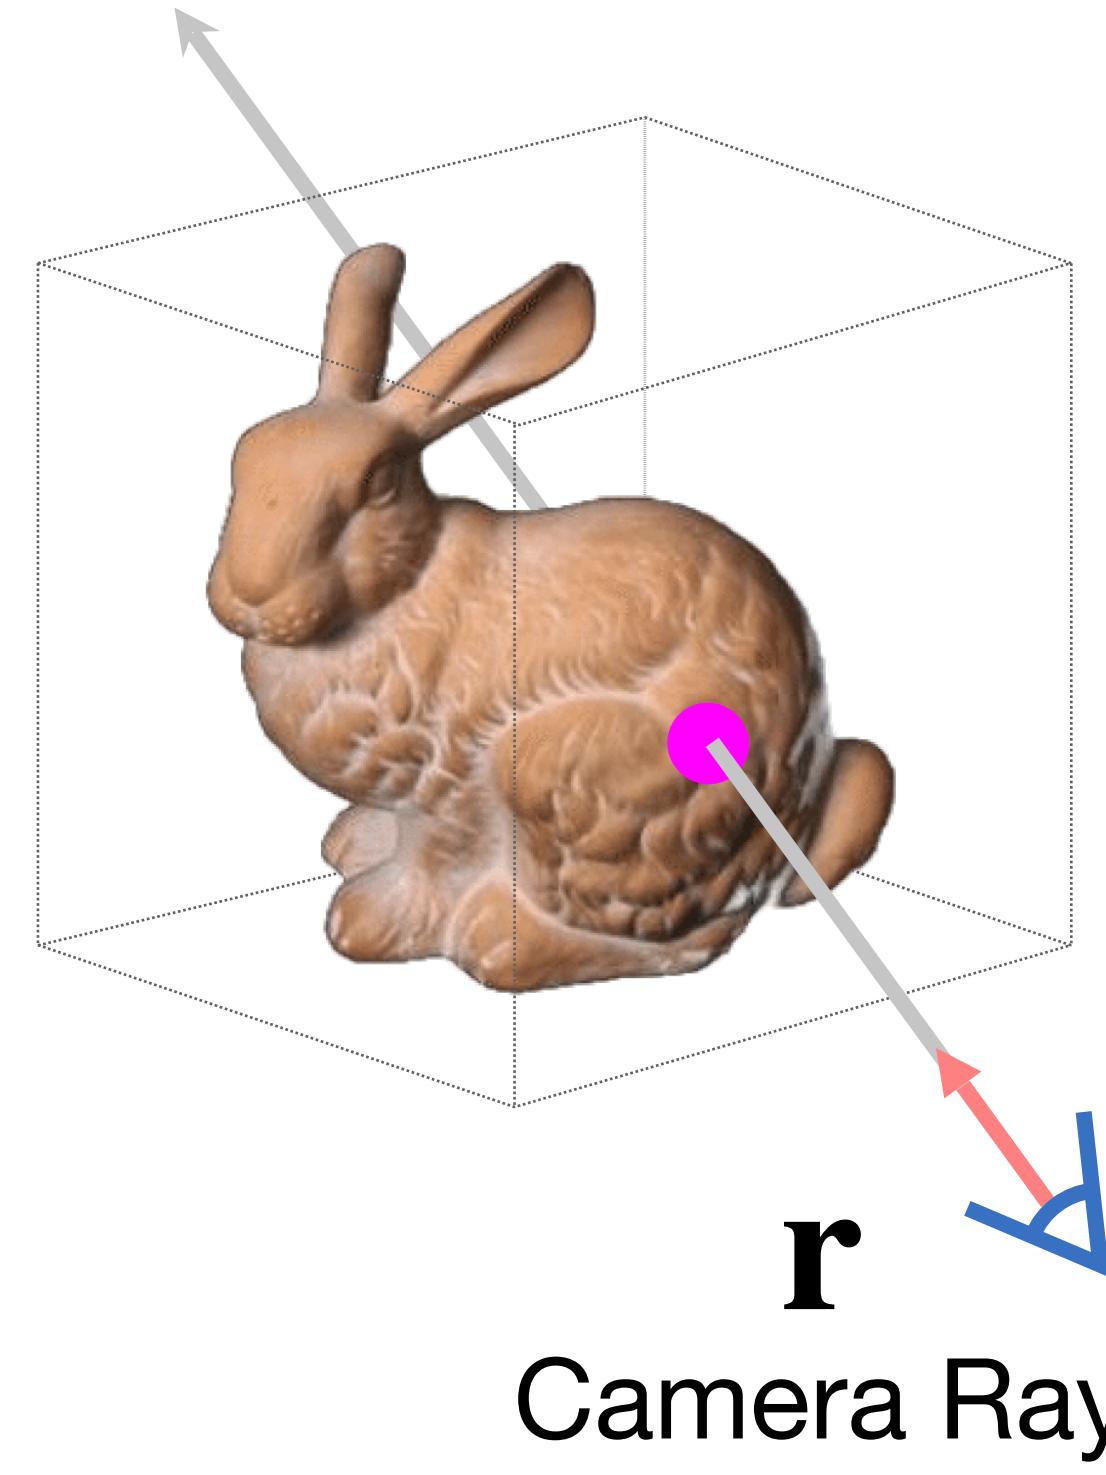


Do we have to parameterize 3D Scene as 3D Function?



3D-structured Representation





A

x_1

x_2

x_3

x_4

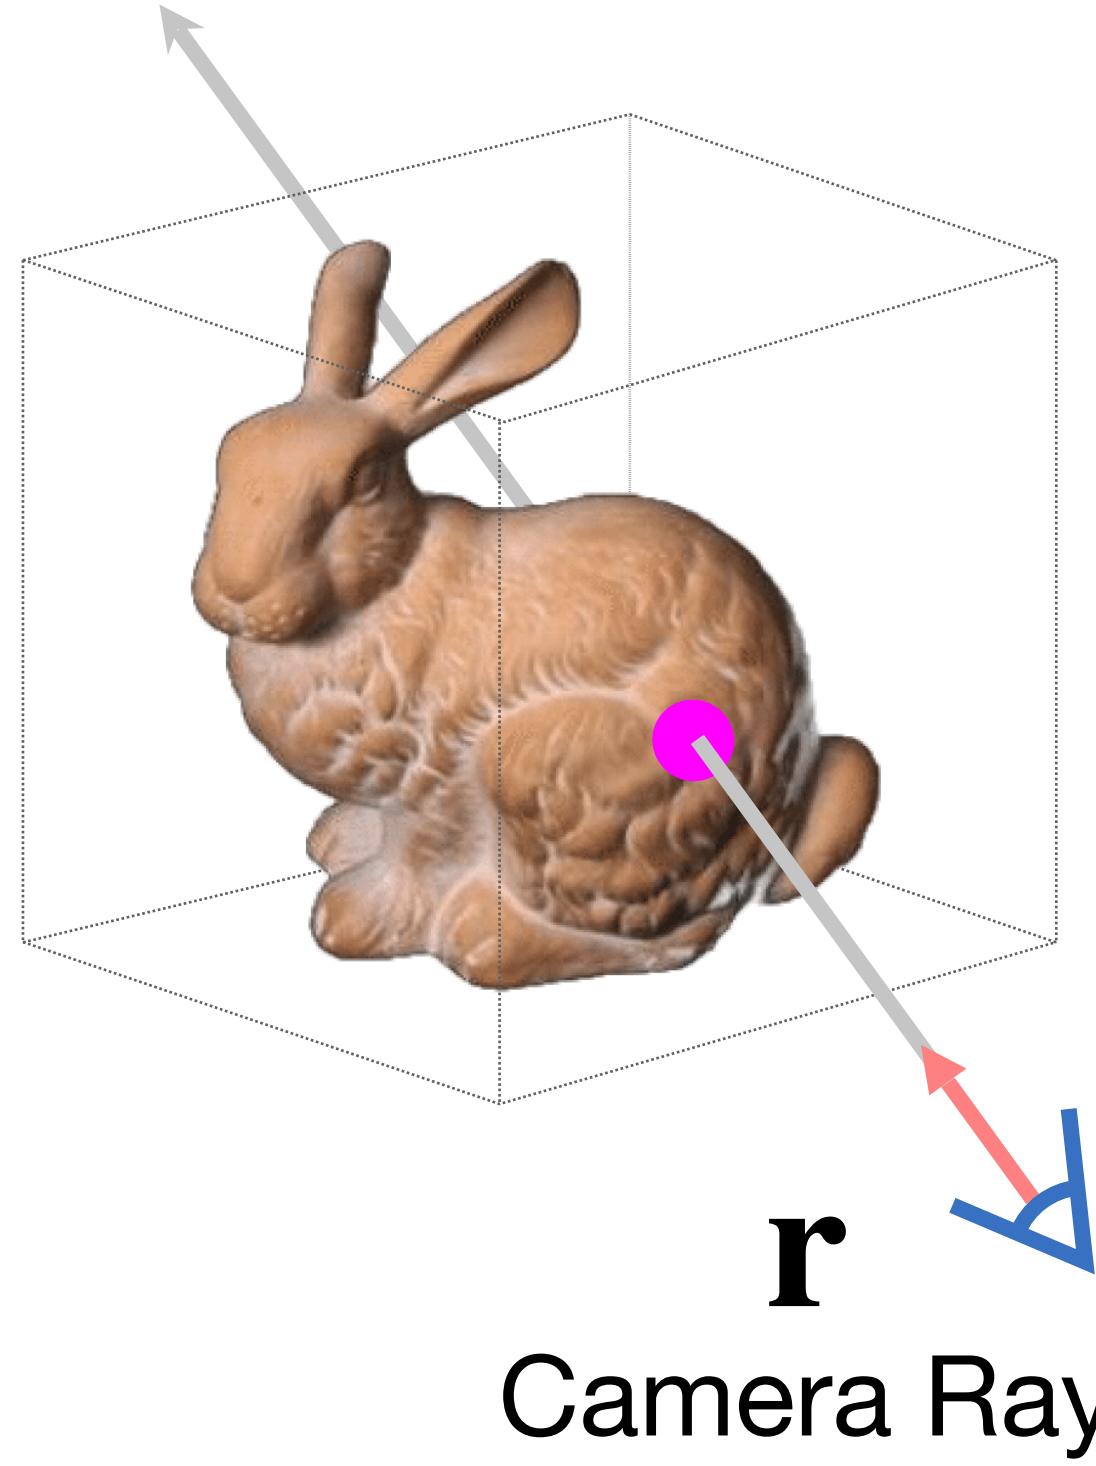
\dots

Light Field

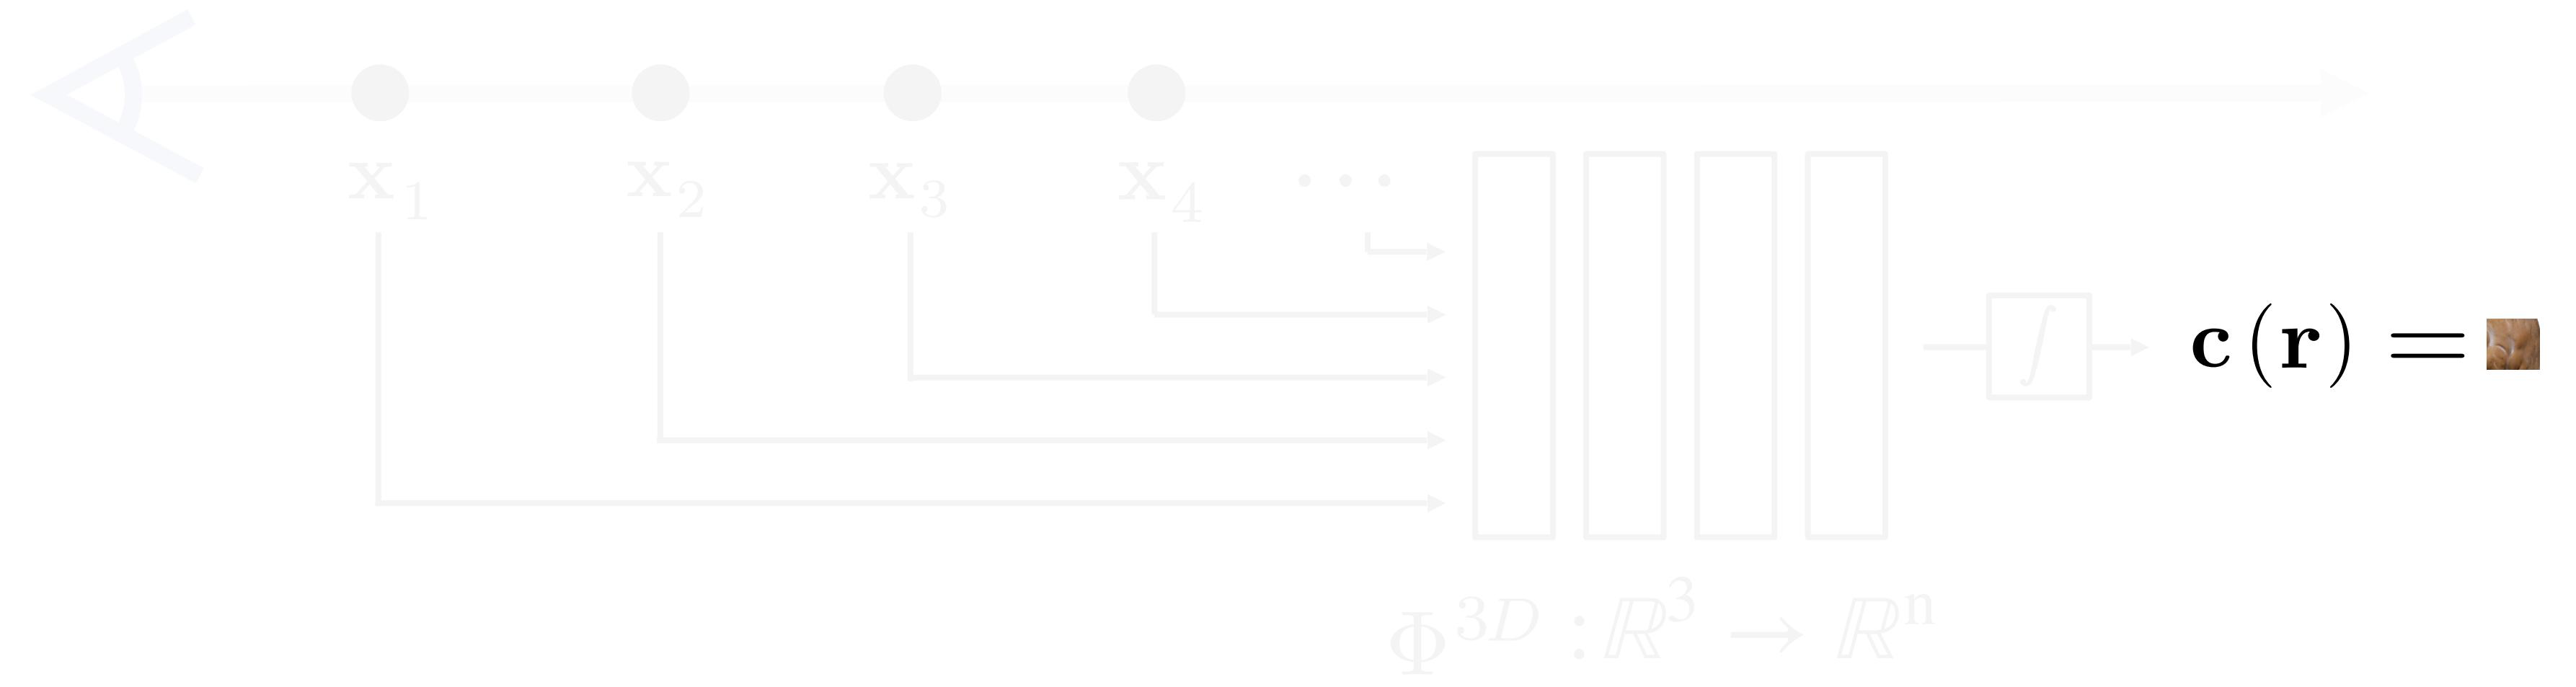
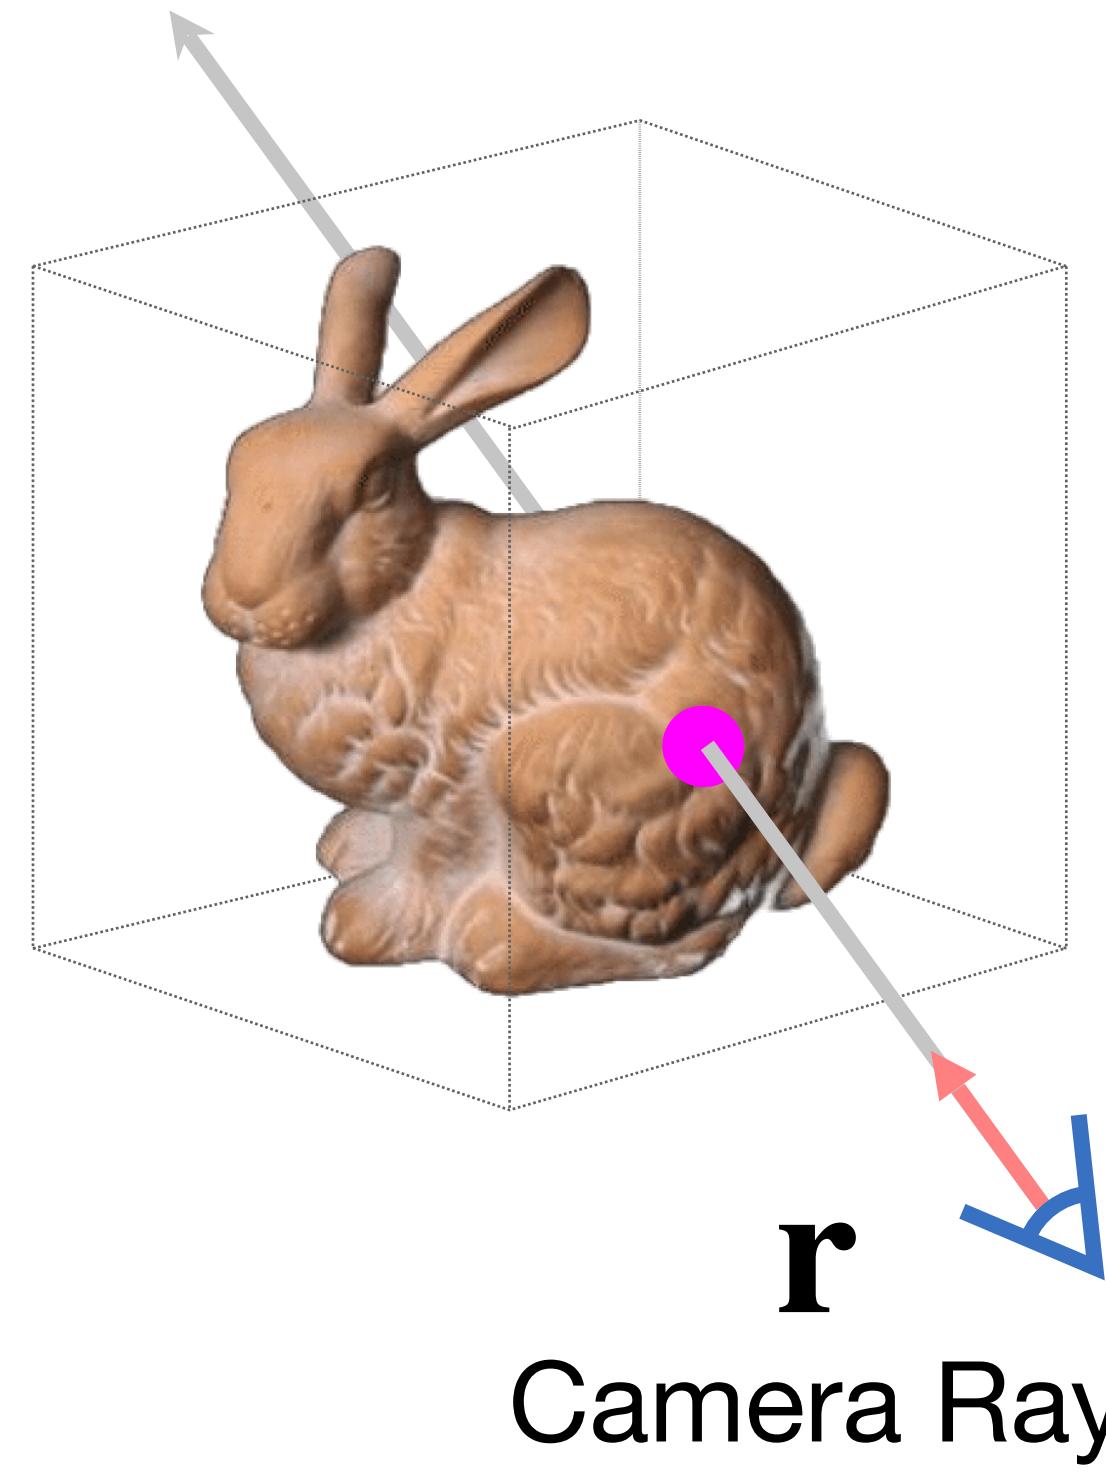
$$\Phi^{3D} : \mathbb{R}^3 \rightarrow \mathbb{R}^n$$



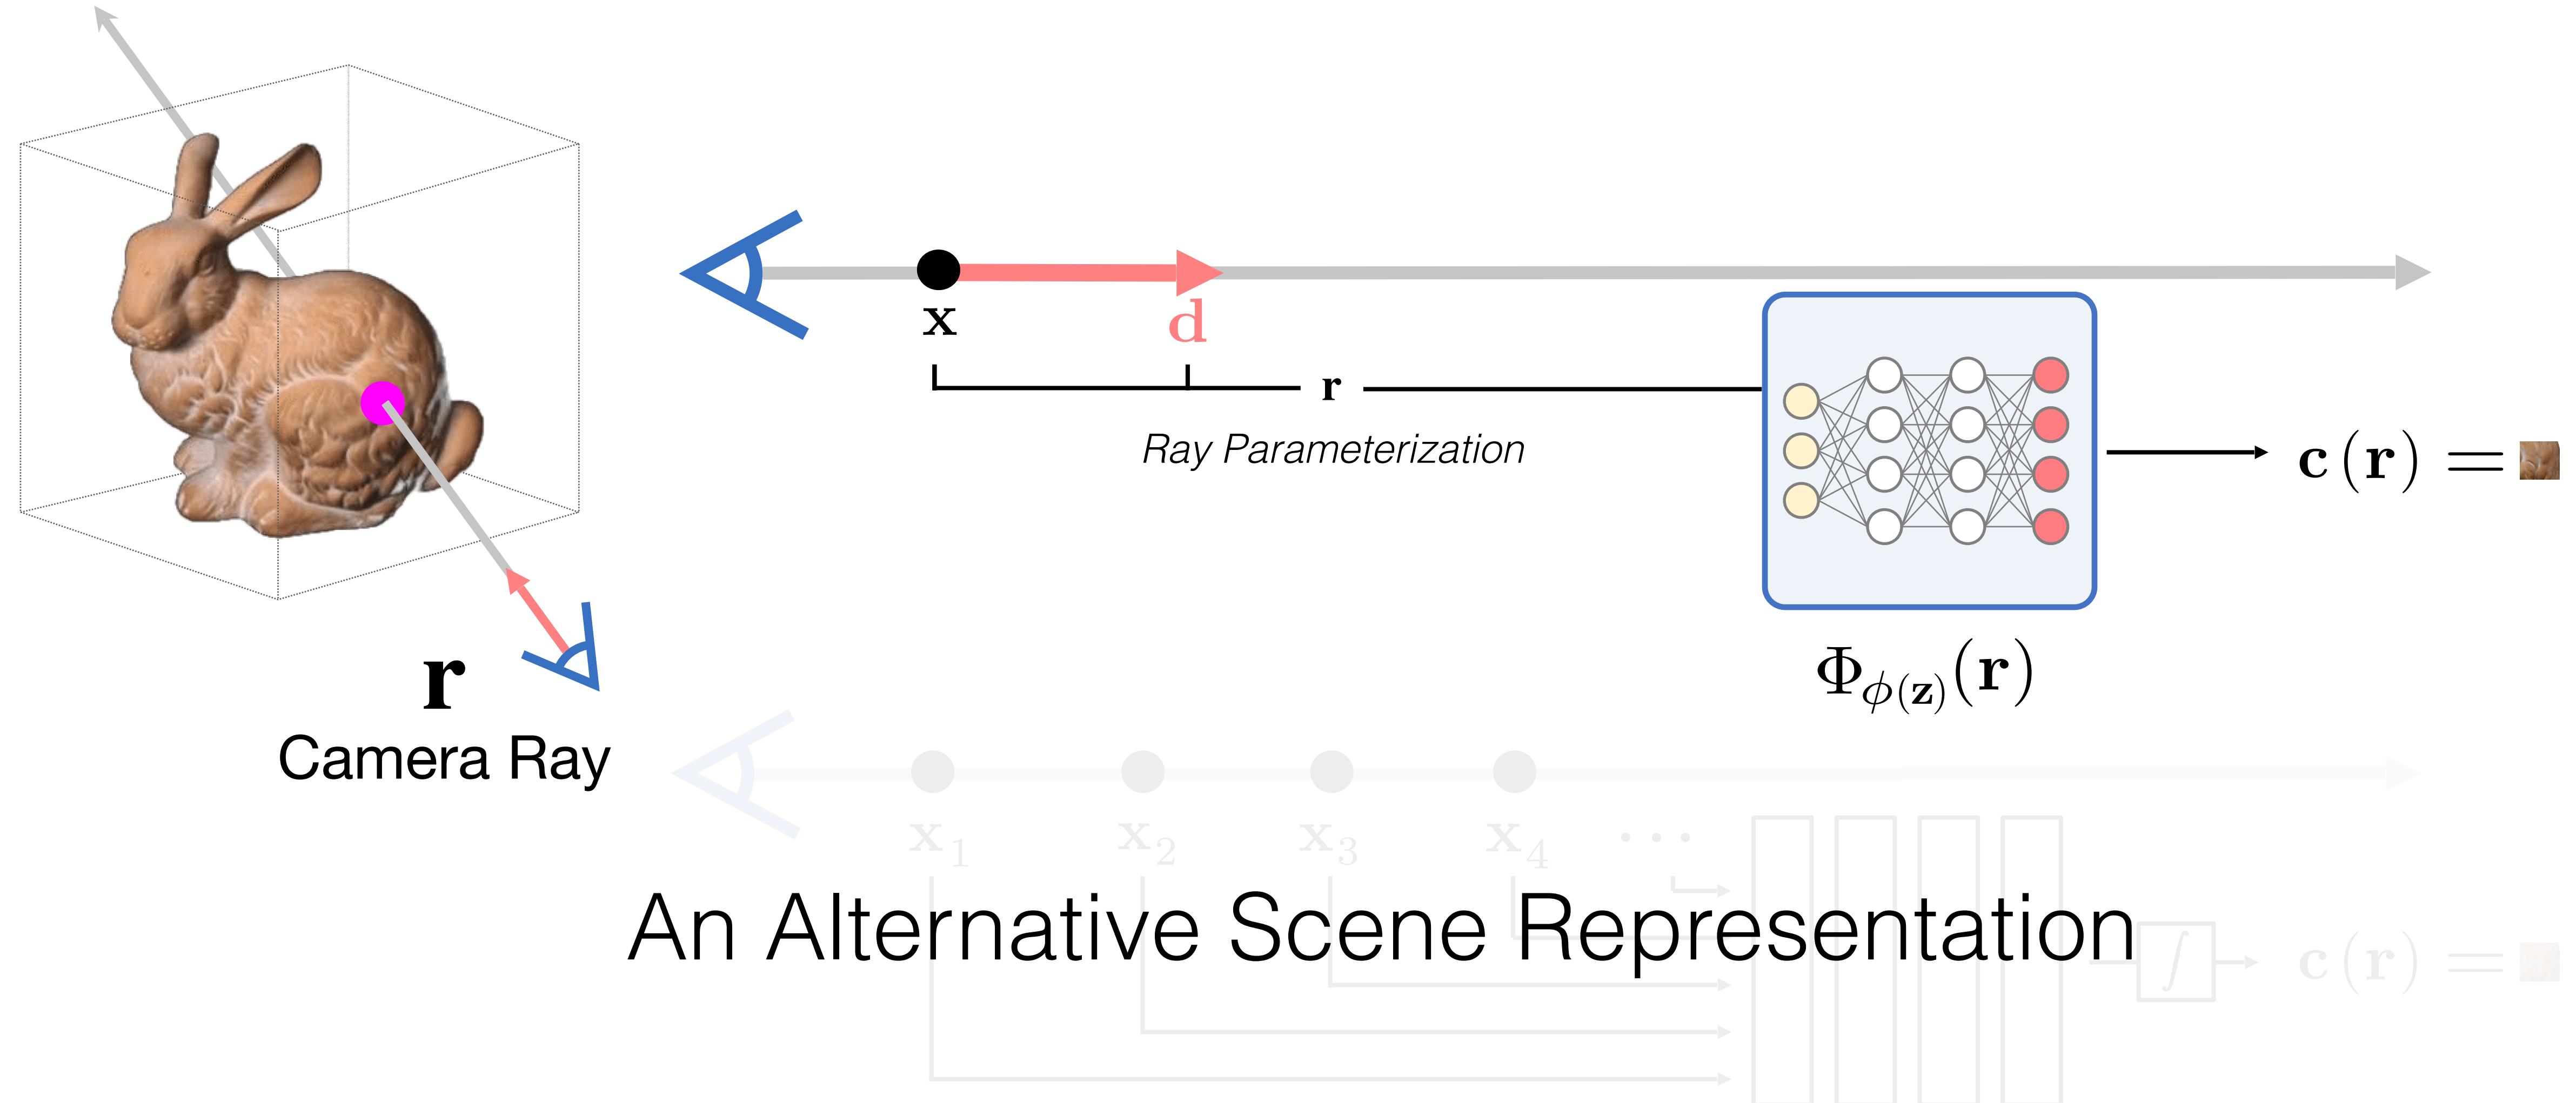
$$c(r) =$$



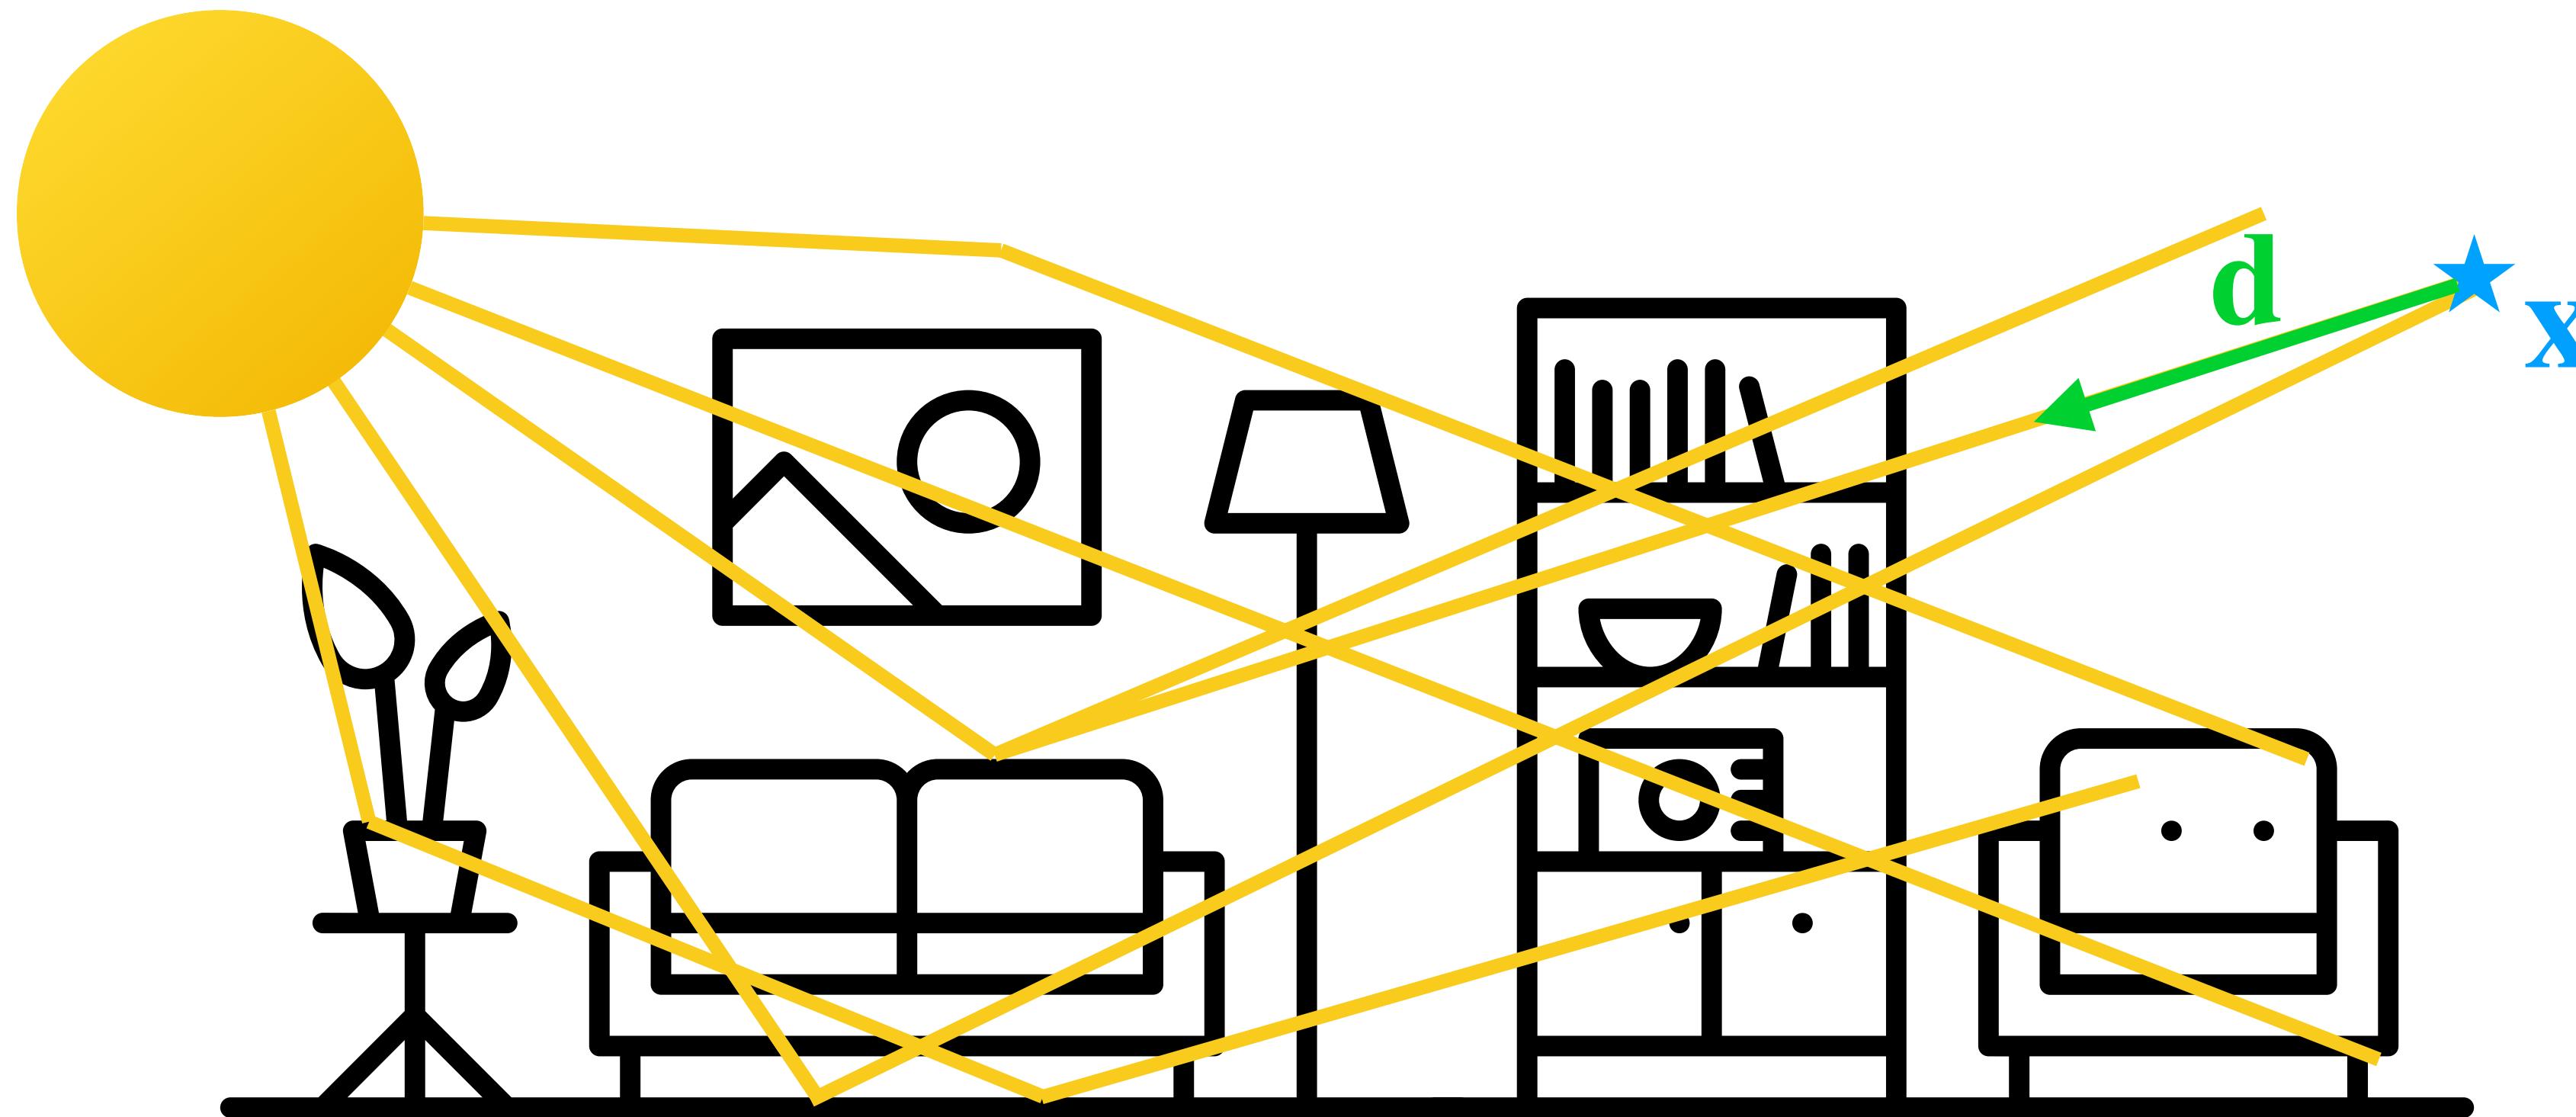
Light Field Networks



Light Field Networks



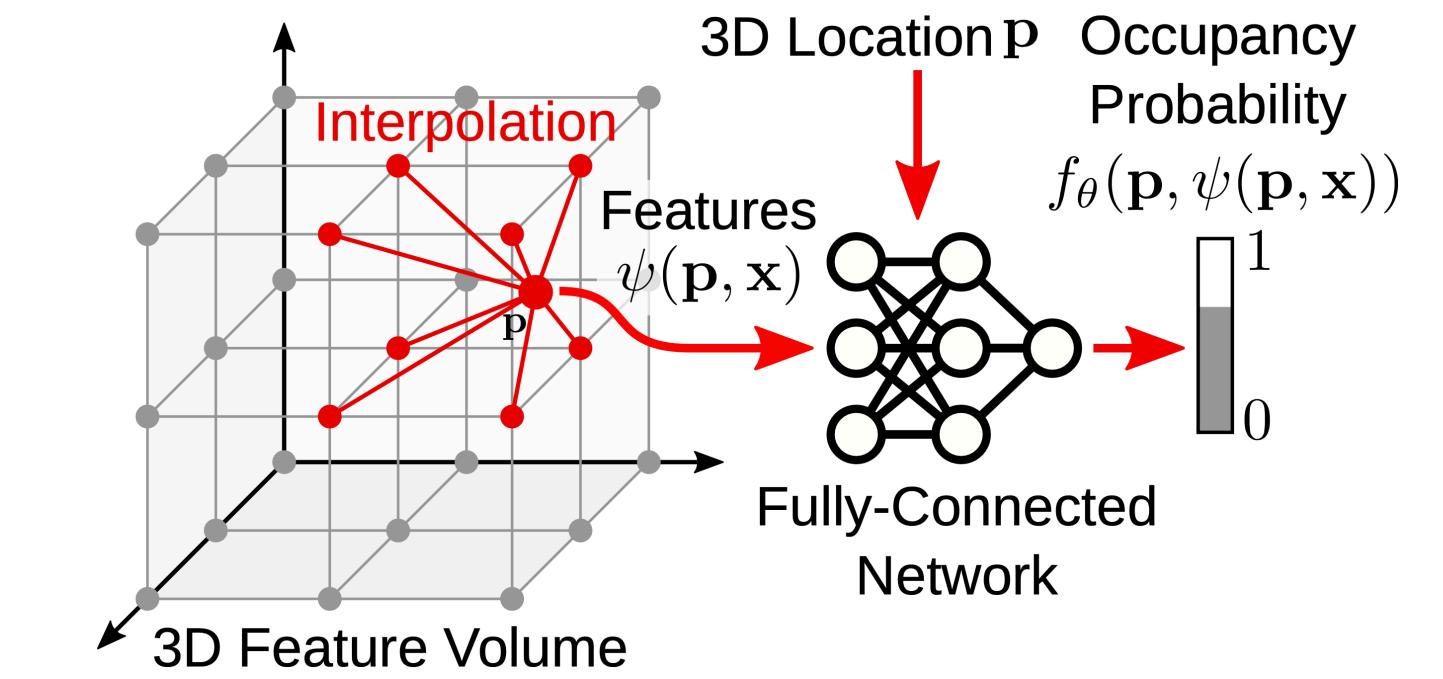
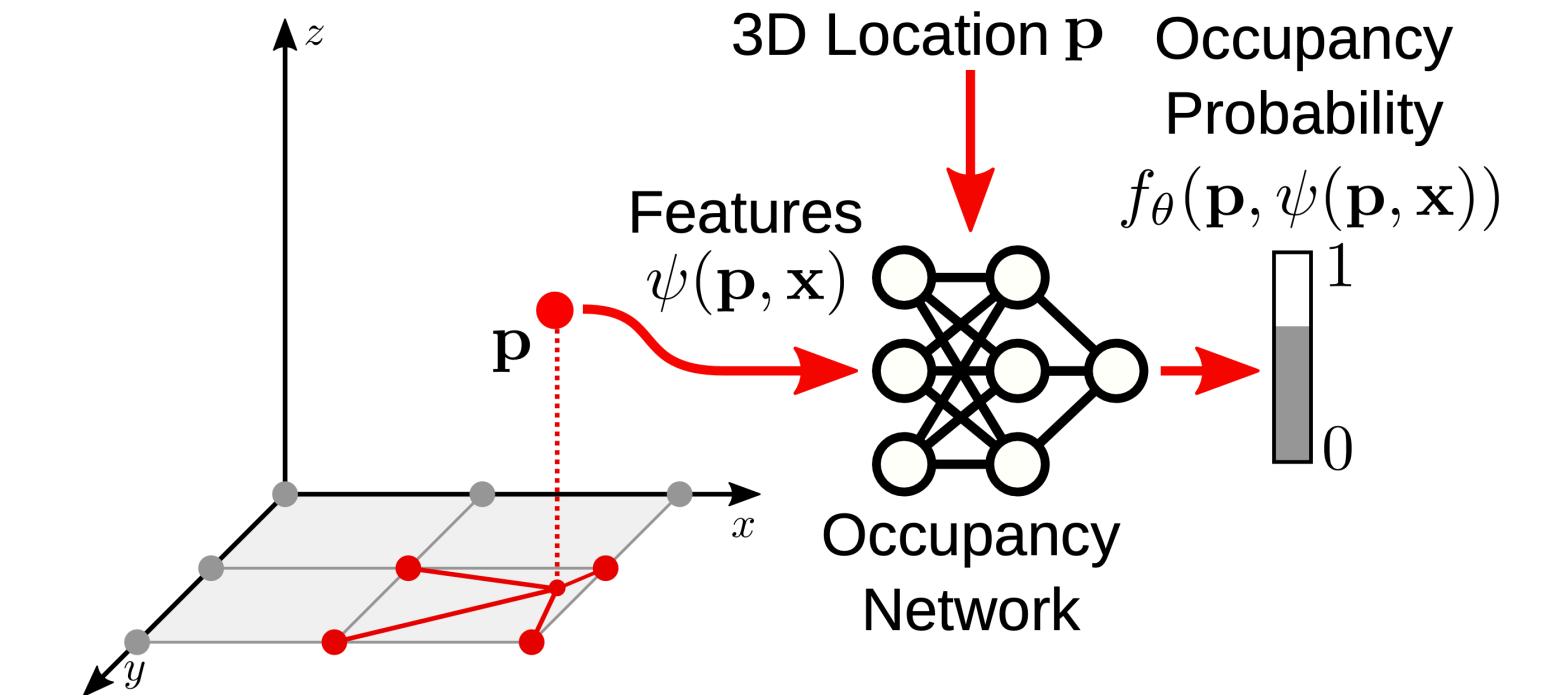
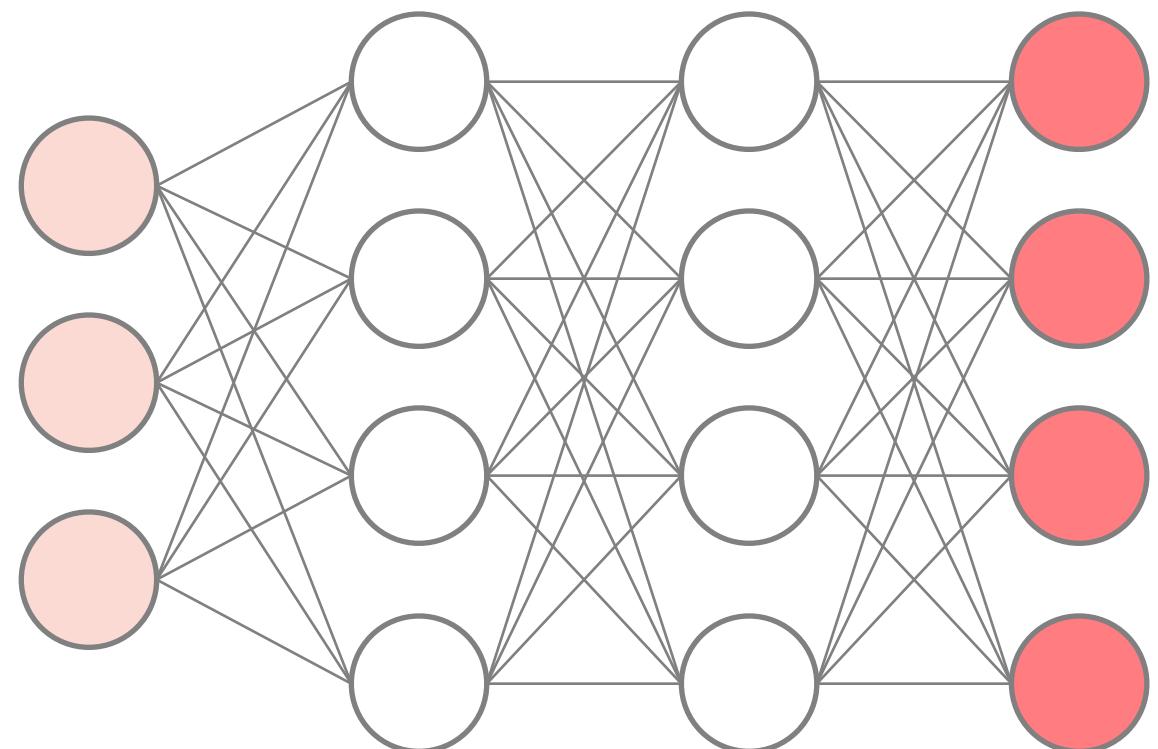
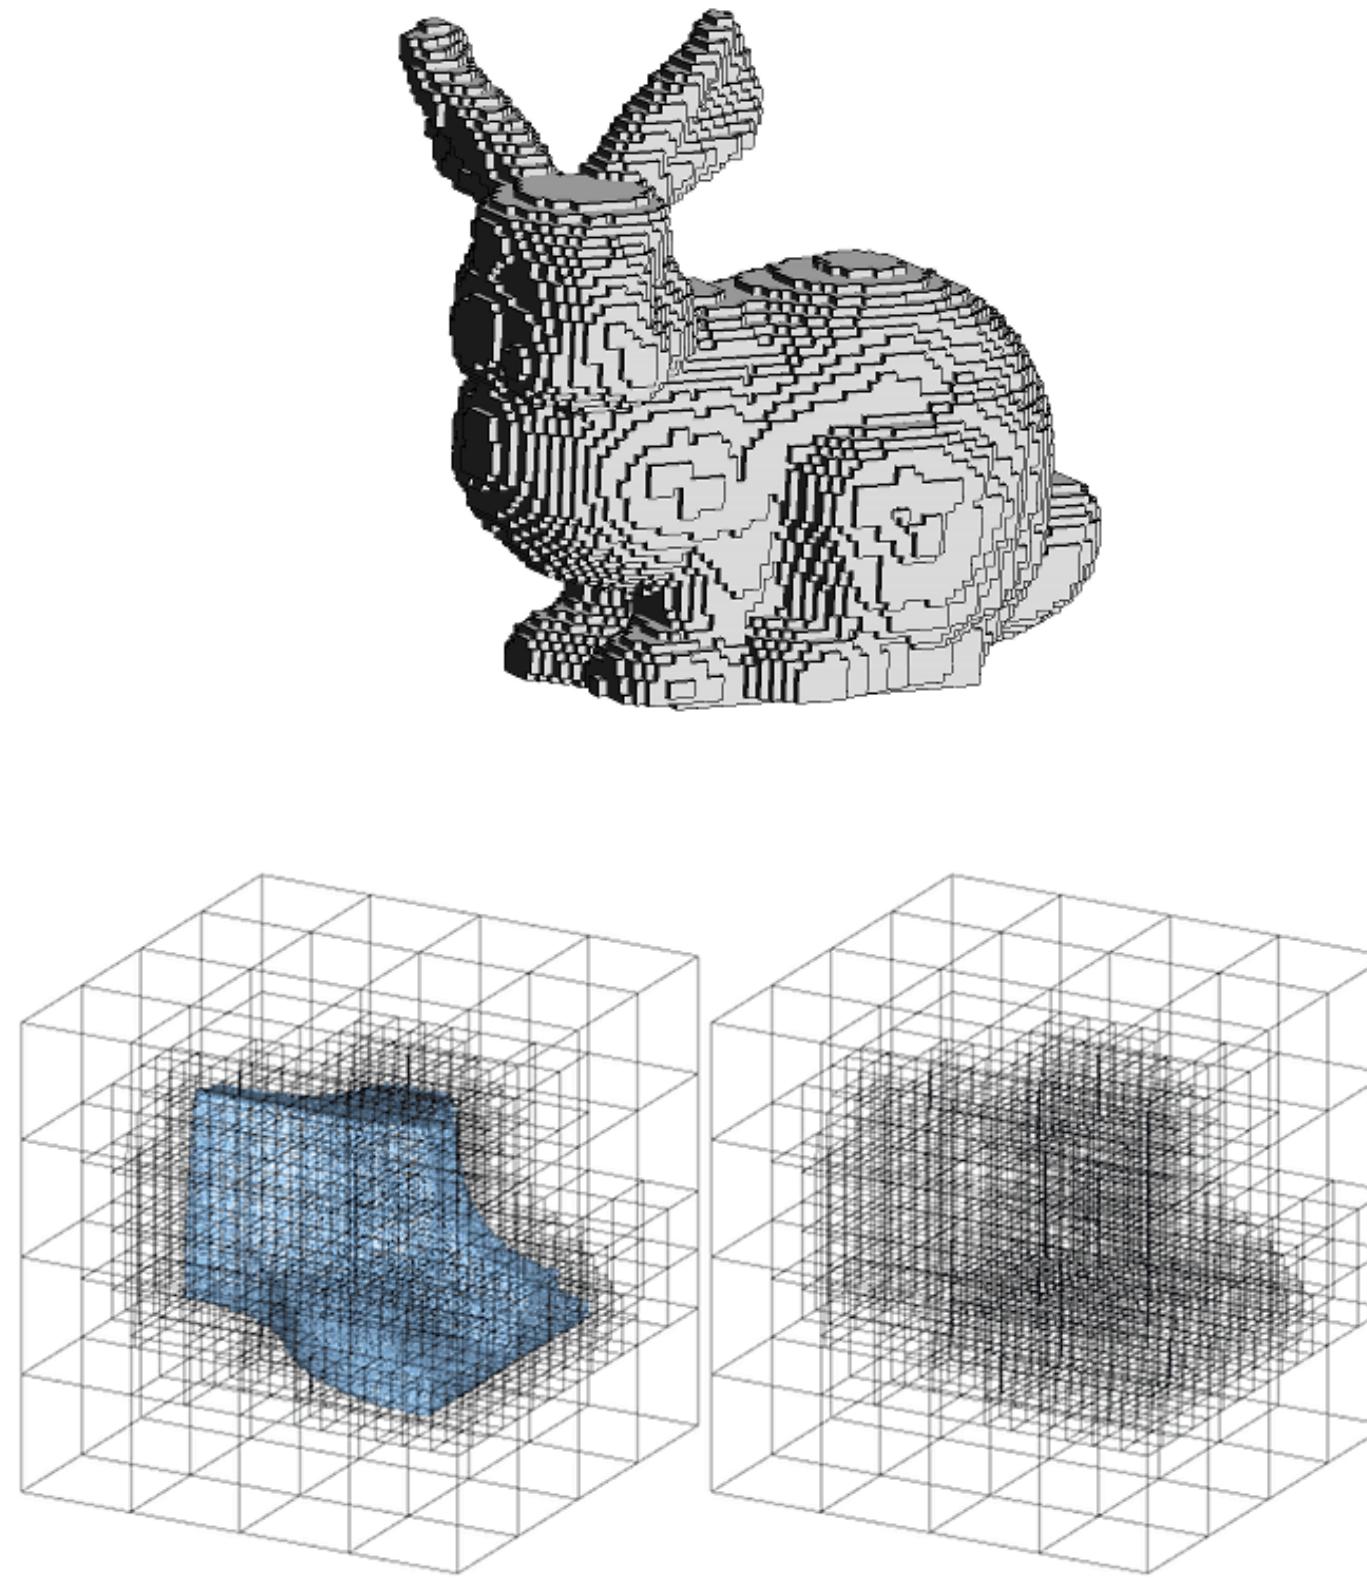
Just because Scene is 3D, don't have to store information in 3D data structure...



$$LF : \mathbb{R}^3 \times \mathbb{S}^2 \rightarrow \mathbb{R}^3, \quad LF(\mathbf{x}, \mathbf{d}) = \mathbf{c}$$

Light Field is 5-dimensional! What parameterizations are possible...?

Summary



- No “Artificial Intelligence”, “understanding” or “Machine Learning” when fitting a single representation to a single 3D thing, even if there are neural networks in them.
- It becomes *a lot* more interesting when considering *inference*, i.e., we want representation to be output of an encoder! More on that later in “Prior-based reconstruction”.